

# Terraform on Google GKE Kubernetes

Author: Nho Luong

Skill: DevOps Engineer Lead



The image shows a screenshot of the Alpine Learning Platform interface. On the left, there is a sidebar with navigation links: HOME, PROFILE, EXAM REGISTRATION, EXAM HISTORY, CERTIFICATIONS (selected), BENEFITS, DIGITAL BADGES, and SUPPORT AND FAQS. Under the CERTIFICATIONS section, there is a 'Certification status' section showing four active certifications: AWS Certified Solutions Architect - Professional (Professional, SAP), AWS Certified Security - Specialty (Specialty, SCS), AWS Certified DevOps Engineer - Professional (Professional, DOP), and AWS Certified Solutions Architect - Associate (Associate, SAA). Below this, there is a 'Digital Badges' section showing one active badge: AWS Certified Cloud Practitioner (Foundational, CLF). On the right, there is a main content area titled 'Activity' which lists '4 items'. These items are: Microsoft Certified: DevOps Engineer Expert (Certification, Active), Microsoft Certified: Azure Solutions Architect Expert (Certification, Active), Microsoft Certified: Azure Administrator Associate (Certification, Active), and Microsoft Certified: Azure Fundamentals (Certification, Active). Each item has a 'View certification details' button.



# Terraform on Google Kubernetes Engine

## Terraform Fundamentals

Terraform CLI Install

1

Terraform Commands

2

Terraform Language Basics

3

Meta-Argument Provider

4

Input Variables &  
Output Values

5

Meta-Argument count

6

Terraform Data sources

7

Meta-Argument `for_each`

8

Remote Backend  
Remote State Data source

9

40

## Real-World Demos



Step by Step  
GitHub  
Documentation

## GKE Clusters

01

GKE Public Standard Cluster

02

GKE Private Standard Cluster

03

GKE Private Standard Cluster  
(Private Endpoint)

04

GKE Private Autopilot Cluster

## GKE Scaling

01

GKE Cluster Autoscaling

02

Horizontal Pod Autoscaling

03

Vertical Pod Autoscaling

# Terraform on Google Kubernetes Engine

## GKE Storage

GCP Compute Engine  
Persistent Disk CSI Driver

1

GCP Cloud SQL

2

GCP Cloud Storage FUSE  
CSI Driver

3

GCP Filestore CSI Driver

4

## GKE Gateway API

Gateway API Basics  
(Cloud Regional Application LB)

1

Gateway API Static IP  
(GCP External IP Address)

2

Gateway API SSL  
(Kubernetes Secrets)

3

Gateway API SSL  
(GCP Certificate Manager)

4

40

## Real-World Demos



Step by Step  
GitHub  
Documentation

## GKE Gateway API

05

Gateway API SSL  
(HTTP to HTTPS Redirect)

06

Gateway API Routing  
(Context Path Routing)

07

Gateway API Routing  
(Domain Name Routing)

08

Gateway API Routing  
(Traffic Splitting)

09

Gateway API  
(Health Checks + Session Affinity)

10

Cloud Domains & Cloud DNS

11

Gateway API Prod Grade SSL  
(Certificate Manager + Cloud DNS + Cloud Domains)

12

Gateway API Prod Grade SSL  
(Certificate Manager + Cloud DNS + AWS Route53)

13

Gateway API Global  
(Cloud Global Application LB)

# Terraform on Google Kubernetes Engine

## GKE Infra DevOps

Terraform Modules

1

GKE Cluster  
(Custom Terraform Module)

2

GKE Infra DevOps  
(GitHub + Cloud Build)

3

## GKE Workload DevOps

Kubernetes Deployment  
(Custom Terraform Module)

1

GKE Workload DevOps  
(GitHub + Cloud Build)

2

GKE Continuous Integration  
(GitHub + Cloud Build + Artifact Registry)

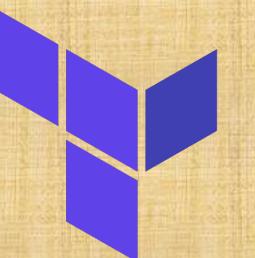
3

GKE Continuous Delivery  
(GitHub + Cloud Build)

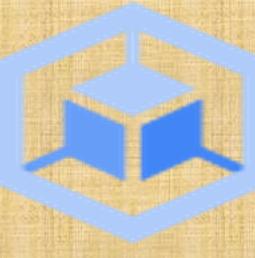
4

40

Real-World  
Demos



Terraform



GKE Cluster

Step by Step  
GitHub  
Documentation

# GitHub Step-by-Step Documentation

- ✓ [terraform-on-google-kubernetes-engine](#)
- > [01-Terraform-Install-Tools](#)
- > [02-Terraform-Commands](#)
- > [03-Terraform-Language-Basics](#)
- > [04-Terraform-MetaArgument-provider](#)
- > [05-Terraform-Variables-Output-Values](#)
- > [06-Terraform-MetaArgument-count](#)
- > [07-Terraform-Datasources](#)
- > [08-Terraform-MetaArgument-foreach](#)
- > [09-GKE-Public-Standard-Cluster](#)
- > [10-Kubernetes-Resources-yaml](#)
- > [11-Kubernetes-Resources-Terraform](#)
- > [12-GKE-Private-Standard-Cluster-Autoscaler](#)
- > [13-GKE-Horizontal-Pod-Autoscaling](#)
- > [14-GKE-Vertical-Pod-Autoscaling](#)
- > [15-GKE-Private-Standard-Cluster-private-endpoint](#)
- > [16-GKE-Private-Autopilot-cluster](#)
- > [17-GKE-Storage-Persistent-Disks](#)
- > [18-GKE-Storage-CloudSQL](#)
- > [19-GKE-Cloud-Storage-FUSE-CSl](#)
- > [20-GKE-Storage-Filestore](#)

40  
Real-World  
Demos



GKE Cluster



Terraform

- ✓ [21-GKE-Gateway-API](#)
- > [01-GKE-LB-Gateway-API-Basic](#)
- > [02-GKE-LB-Gateway-API-StaticIP](#)
- > [03-GKE-Gateway-API-Selfsigned-SSL-k8sSecrets](#)
- > [04-GKE-Gateway-API-Selfsigned-SSL-CertManager](#)
- > [05-GKE-Gateway-API-HTTP-to-HTTPS-Redirect](#)
- > [06-GKE-Gateway-API-ContextPath-Routing](#)
- > [07-GKE-Gateway-API-Domain-Routing](#)
- > [08-GKE-Gateway-API-Traffic-Splitting](#)
- > [09-GKE-Gateway-API-HealthChecks-SessionAffinity](#)
- > [10-Cloud-Domains-and-Cloud-DNS](#)
- > [11-GKE-Gateway-API-ProdSSL-CloudDNS](#)
- > [12-GKE-Gateway-API-ProdSSL-ExternalDomainProvider](#)
- > [13-GKE-Gateway-API-Global-LB](#)
- > [22-Terraform-Modules](#)
- > [23-GKE-Infra-Custom-Terraform-Modules](#)
- > [24-GKE-Infra-DevOps-CloudBuild-GitHub](#)
- > [25-GKE-Workloads-Custom-Terraform-Modules](#)
- > [26-GKE-Workloads-DevOps-CloudBuild-GitHub](#)
- > [27-GKE-App-Continuous-Integration](#)
- > [28-GKE-App-Continuous-Delivery](#)

# How are Terraform Configs organized ?

Project-1

GKE Cluster  
Terraform Manifests

- ✓ 02-GKE-LB-Gateway-API-StaticIP
  - ✓ p1-gke-autopilot-cluster-private
    - └ c1-versions.tf
    - └ c2-01-variables.tf
    - └ c2-02-local-values.tf
    - └ c3-vpc.tf
    - └ c4-01-gke-cluster.tf
    - └ c4-02-gke-outputs.tf
    - └ c5-Cloud-NAT-Cloud-Router.tf
    - └ terraform.tfvars

Project-2

Kubernetes  
YAML Manifests

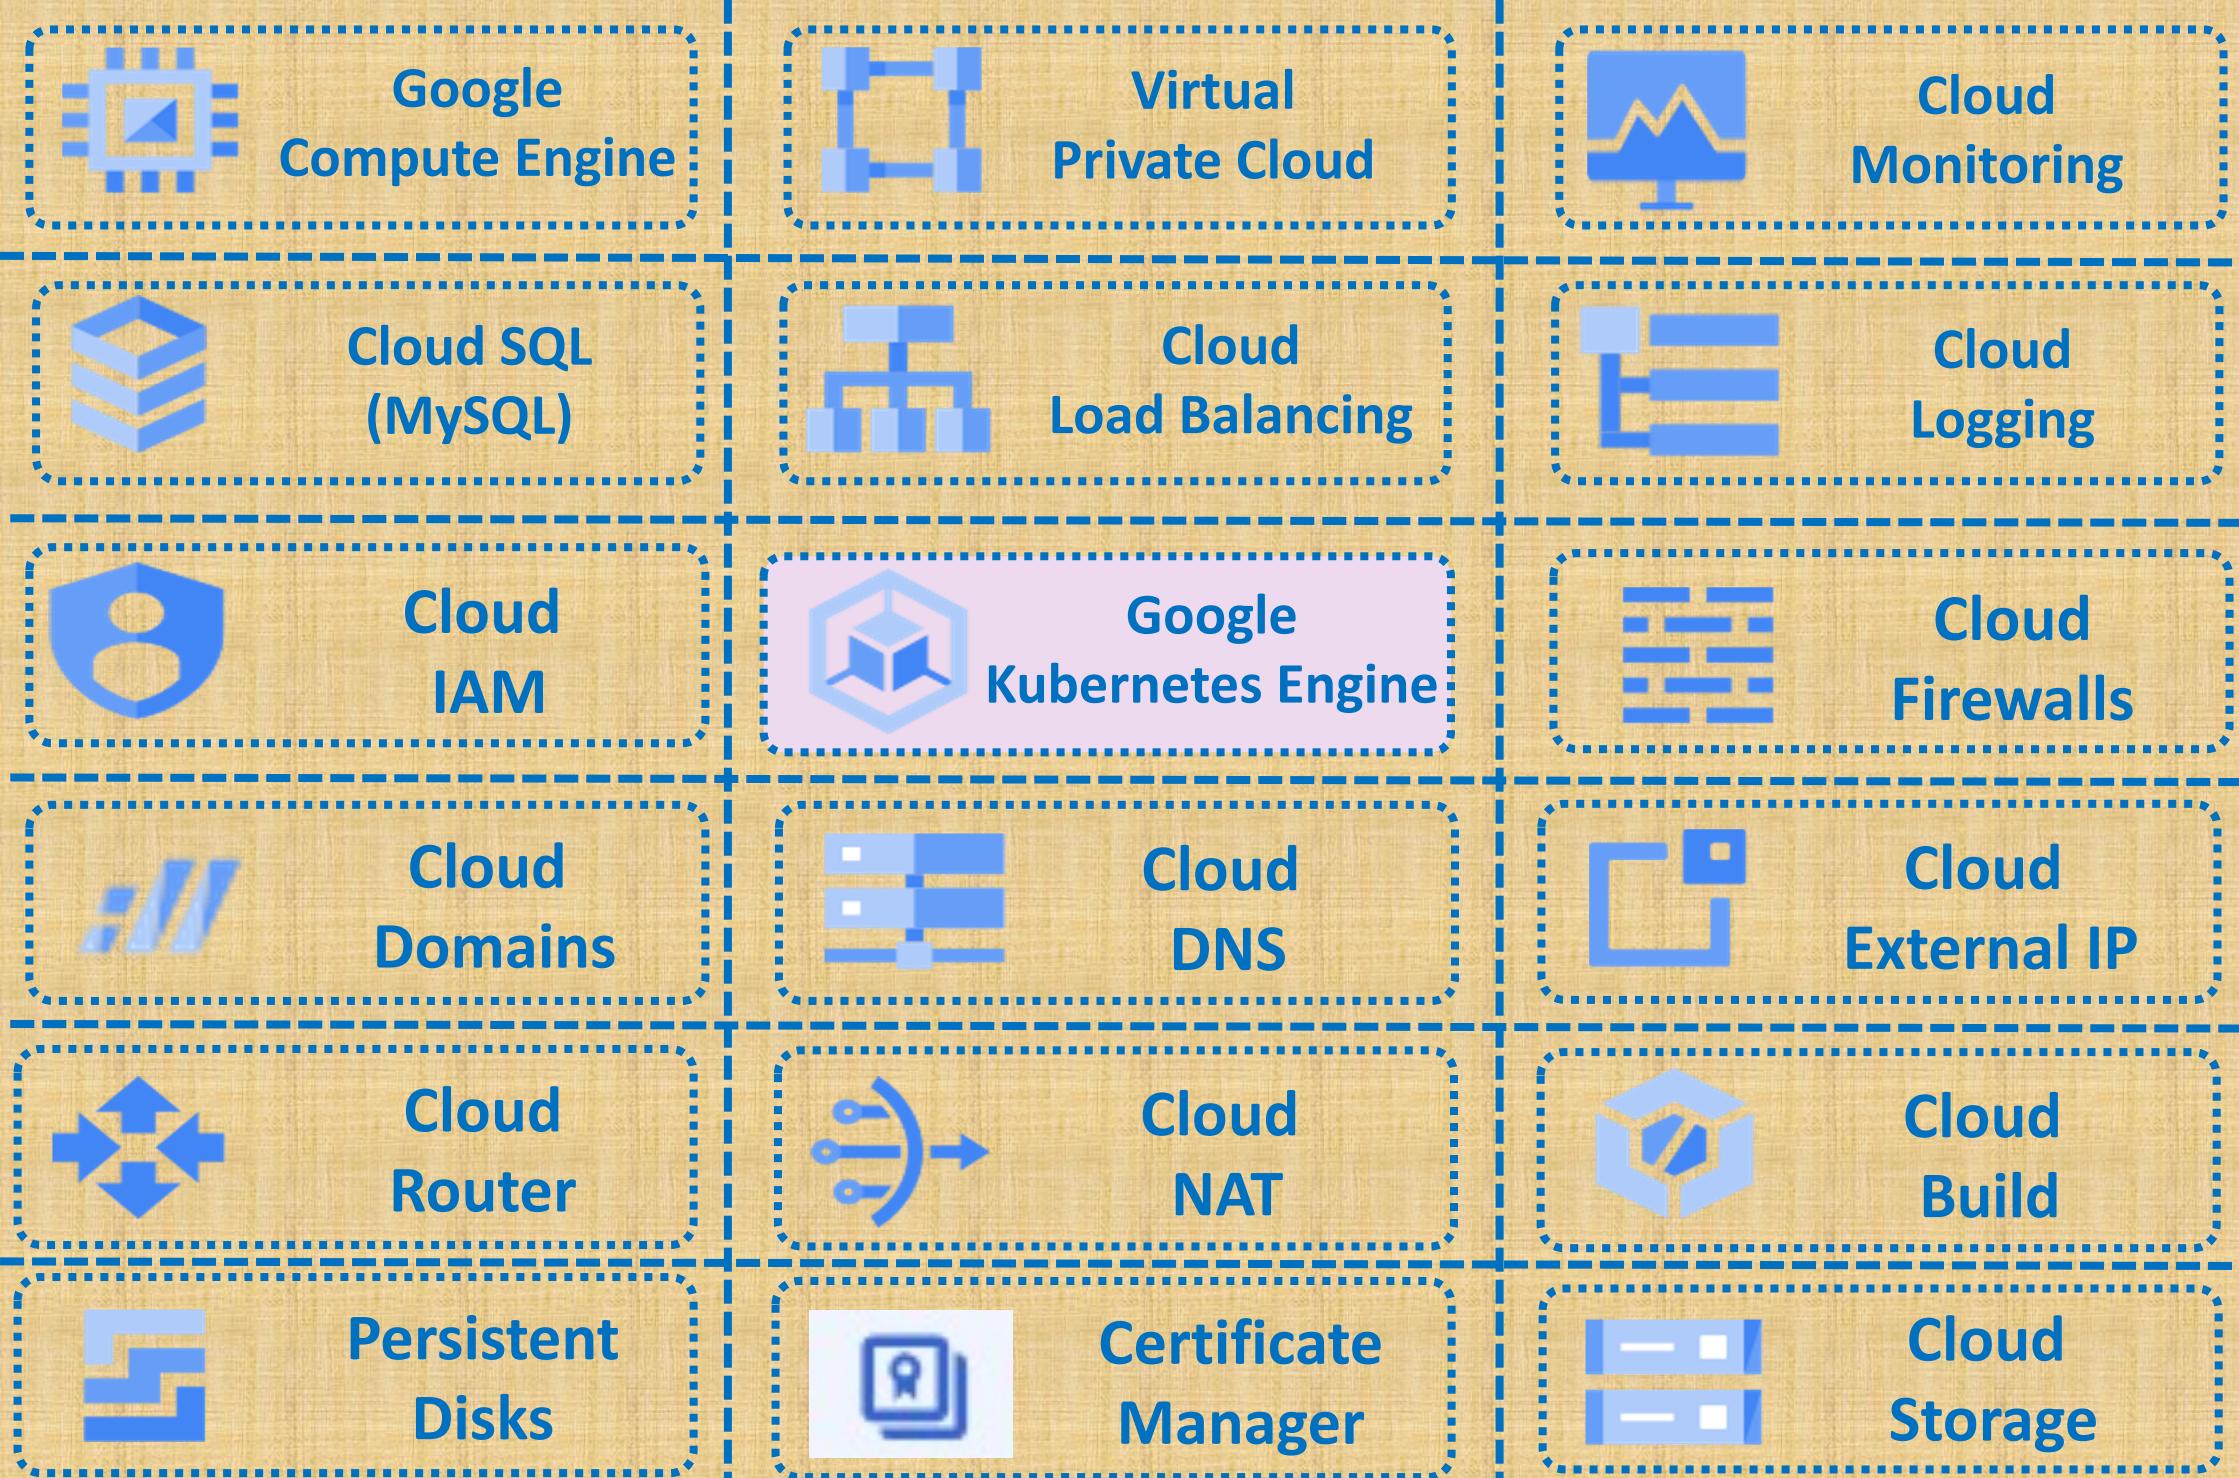
- ✓ 02-GKE-LB-Gateway-API-StaticIP
  - > p1-gke-autopilot-cluster-private
  - ✓ p2-regional-k8sresources-yaml
    - ! 01-myapp1-deployment.yaml
    - ! 02-myapp1-clusterip-service.yaml
    - ! 03-gateway.yaml
    - ! 04-gateway-http-route.yaml

Project-3

Kubernetes Terraform  
Manifests

- ✓ 02-GKE-LB-Gateway-API-StaticIP
  - > p1-gke-autopilot-cluster-private
  - > p2-regional-k8sresources-yaml
  - ✓ p3-regional-k8sresources-terraform-manifests
    - └ c1-versions.tf
    - └ c2-01-variables.tf
    - └ c2-02-local-values.tf
    - └ c3-01-remote-state-datasource.tf
    - └ c3-02-providers.tf
    - └ c4-myapp1-deployment.tf
    - └ c5-myapp1-clusterip-service.tf
    - └ c6-gateway.tf
    - └ c7-gateway-http-route.tf
    - └ c8-static-ip.tf
    - └ terraform.tfvars

# Terraform on Google Cloud



18+  
GCP  
Services

Many more  
Sub-services  
under each  
service

Automated  
with  
Terraform

# GitHub Repository

Repository Used For	Repository URL
Main Repository with step-by-step Docs	<a href="https://github.com/nholuongut/terraform-on-google-cloud">https://github.com/nholuongut/terraform-on-google-cloud</a>
Terraform Google CI Repo	<a href="https://github.com/nholuongut/terraform-google-ci">https://github.com/nholuongut/terraform-google-ci</a>
GCP GKE App DevOps Repo	<a href="https://github.com/nholuongut/terraform-gcp-gke-k8s-devops">https://github.com/nholuongut/terraform-gcp-gke-k8s-devops</a>
Terraform GCP GKE K8s Repo	<a href="https://github.com/stacksimplify/terraform-gcp-gke-k8s-devops">https://github.com/stacksimplify/terraform-gcp-gke-k8s-devops</a>
Terraform GCP GKE Infra DevOps	<a href="https://github.com/nholuongut/terraform-gcp-gke-infra-devops">https://github.com/nholuongut/terraform-gcp-gke-infra-devops</a>
Terraform GCP GKE App DevOps	git@github.com:nholuongut/terraform-gcp-gke-app-devops.git

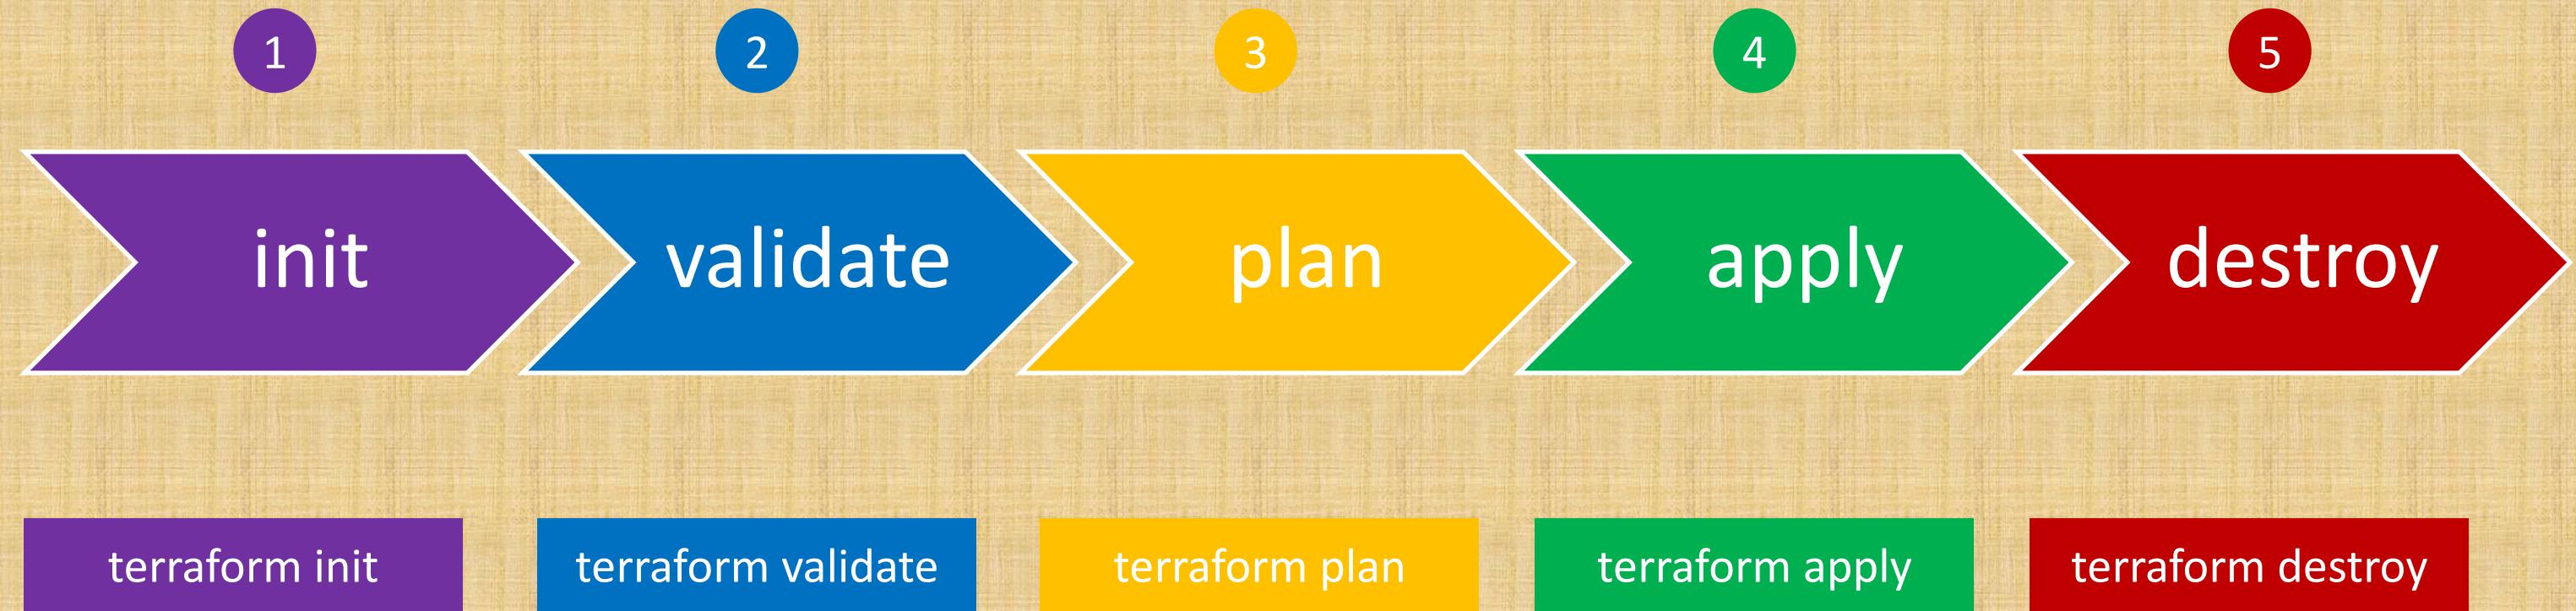
# Terraform Fundamentals

- ✓ Demo-01: Install CLI Tools: gcloud CLI, Terraform CLI, VSCode Editor  
7 lectures • 26min
- ✓ Demo-02: Terraform Commands (init, validate, plan, apply and destroy)  
3 lectures • 27min
- ✓ Demo-03: Terraform Language Basics  
0 lectures • 1hr 15min
- ✓ Demo-04: Terraform Meta-Argument - Provider (Multiple Providers Demo)  
2 lectures • 11min
- ✓ Demo-05: Terraform Input Variables and Output Values  
5 lectures • 38min
- ✓ Demo-06: Terraform Meta-argument: count  
4 lectures • 28min
- ✓ Demo-07: Terraform Datasources  
2 lectures • 11min
- ✓ Demo-08: Terraform Meta-argument: for\_each  
3 lectures • 20min



**Terraform  
Fundamentals**

# Terraform Workflow



**Terraform** language uses a **limited** number of **top-level block** types, which are **blocks** that can appear **outside** of any other **block** in a TF configuration file.

## Terraform Top-Level Blocks

Most of **Terraform's features** are implemented as **top-level** blocks.

Terraform Block

Providers Block

Resources Block

Fundamental Blocks

Input Variables Block

Output Values Block

Local Values Block

Variable Blocks

Data Sources Block

Modules Block

Calling / Referencing Blocks

Import Block

Moved Block

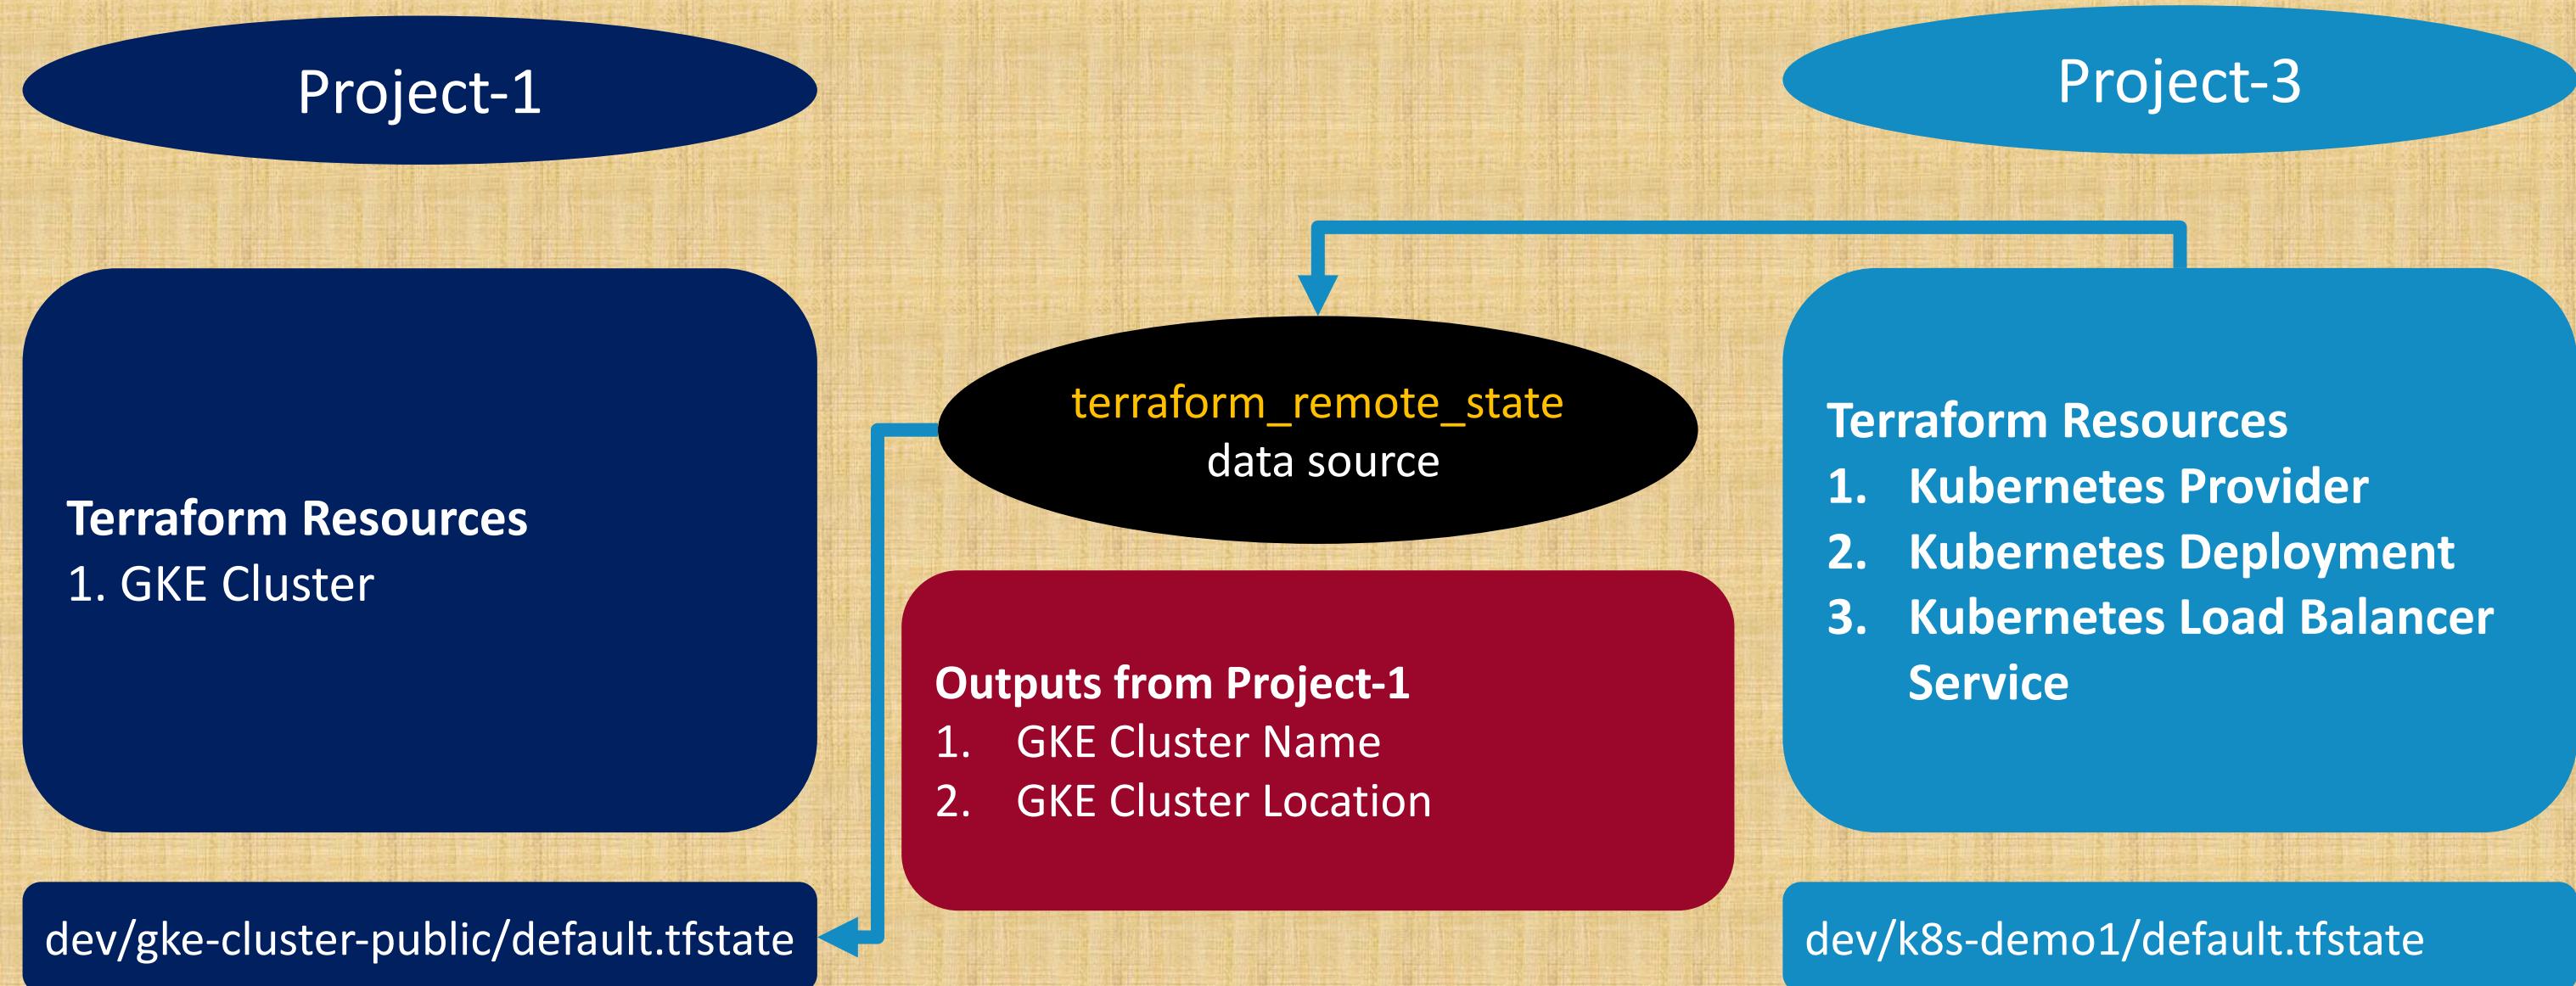
Removed Block

Check Block

**NEW BLOCKS – Added Recently**

# Terraform Remote State Datasource

The `terraform_remote_state` data source retrieves the root module output values from project-1 Terraform configuration, using the latest state snapshot from the remote backend.



# GKE Standard Public Cluster - Network Design

Customer Project: gcplearn9

Customer VPC: myvpc

Region: us-central1



GKE Cluster

Subnet: 10.128.0.0/20

Zone: us-central1-a



GKE Node-1



Load Balancer Service

Cloud Load Balancing

Zone: us-central1-b



GKE Node-2



Public IP

Private IP

Public IP

Private IP

Automated by Terraform

Google Managed Project

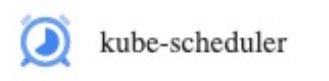
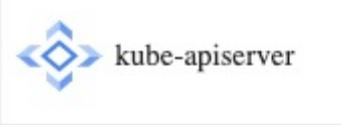


Google Managed VPC Network

Region: us-central1



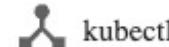
GKE Control Plane

Kube API Server  
Public IP

http://LB-IP



Users



kubectl



Admin



Docker Hub



Internet

NETWORK  
DESIGN

Download Docker Image from Docker Hub

# GKE Standard Private Cluster - Network Design

Customer Project: gcplearn9

Customer VPC: myvpc

Region: us-central1



Subnet: 10.128.0.0/20

Zone: us-central1-a



GKE Node-1



Load Balancer Service

Private IP

Cloud Load Balancing



Cloud NAT



Cloud Router

Automated by Terraform

Zone: us-central1-b



GKE Node-2



Private IP

Cluster Autoscaler

Google Managed Project

Google Managed VPC Network

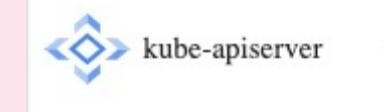
Region: us-central1



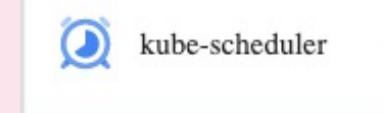
GKE Control Plane

Kube API Server Public IP

Kube API Server Private IP



kube-apiserver



kube-scheduler



Resource Controllers



Storage

VPC Network Peering  
Private Connectivity

Private Connectivity

Download Docker Image from Docker Hub

http://LB-IP



Users



kubectl



Admin



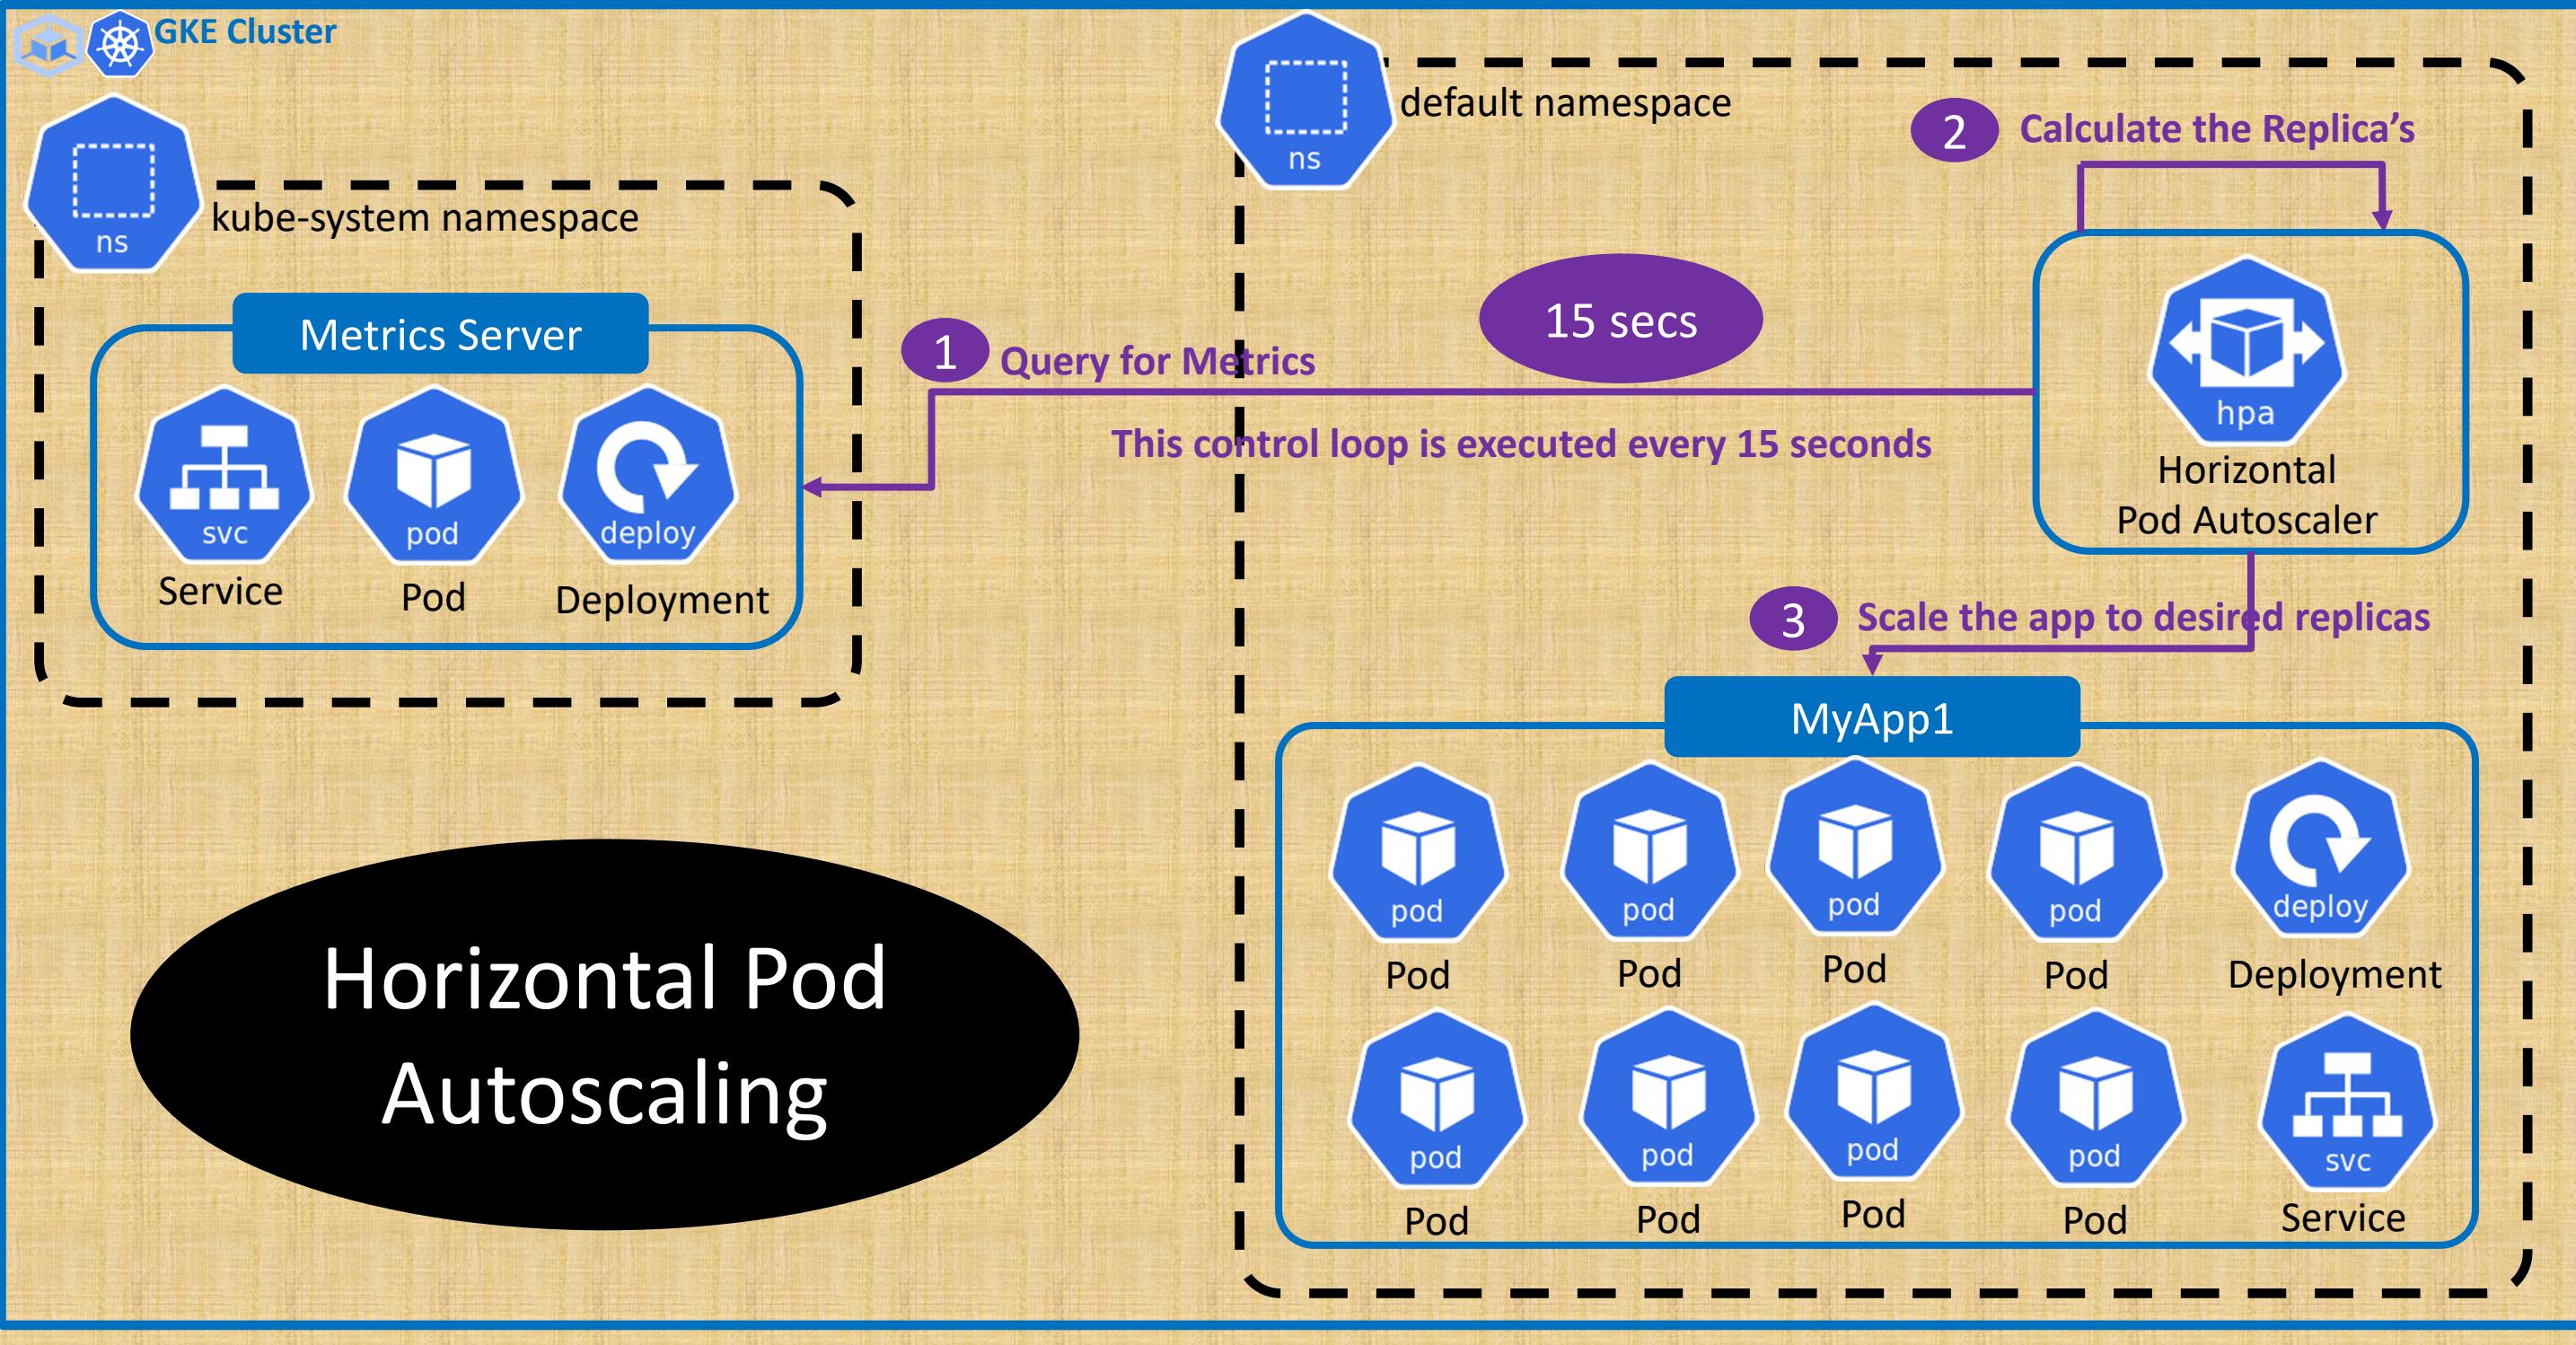
Docker Hub



Internet

NETWORK DESIGN

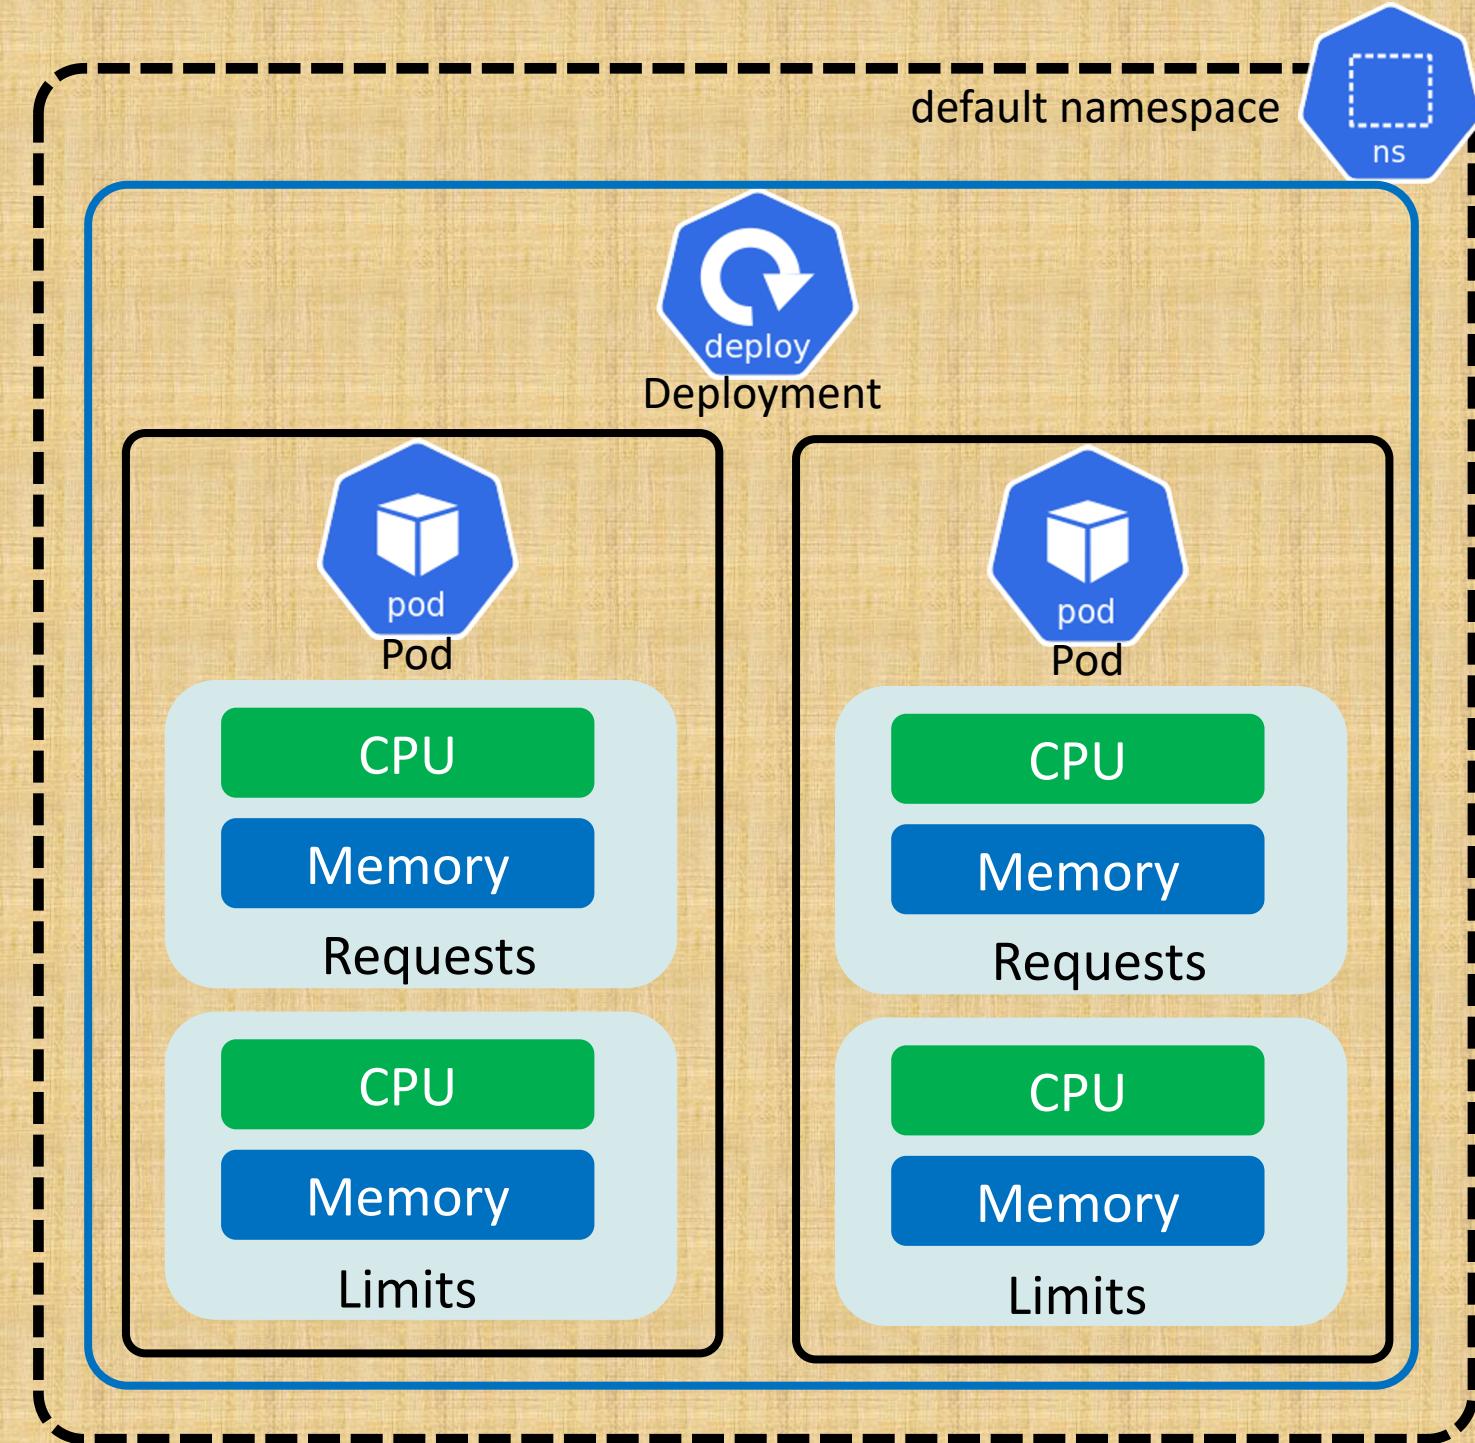
# Kubernetes Horizontal Pod Autoscaling



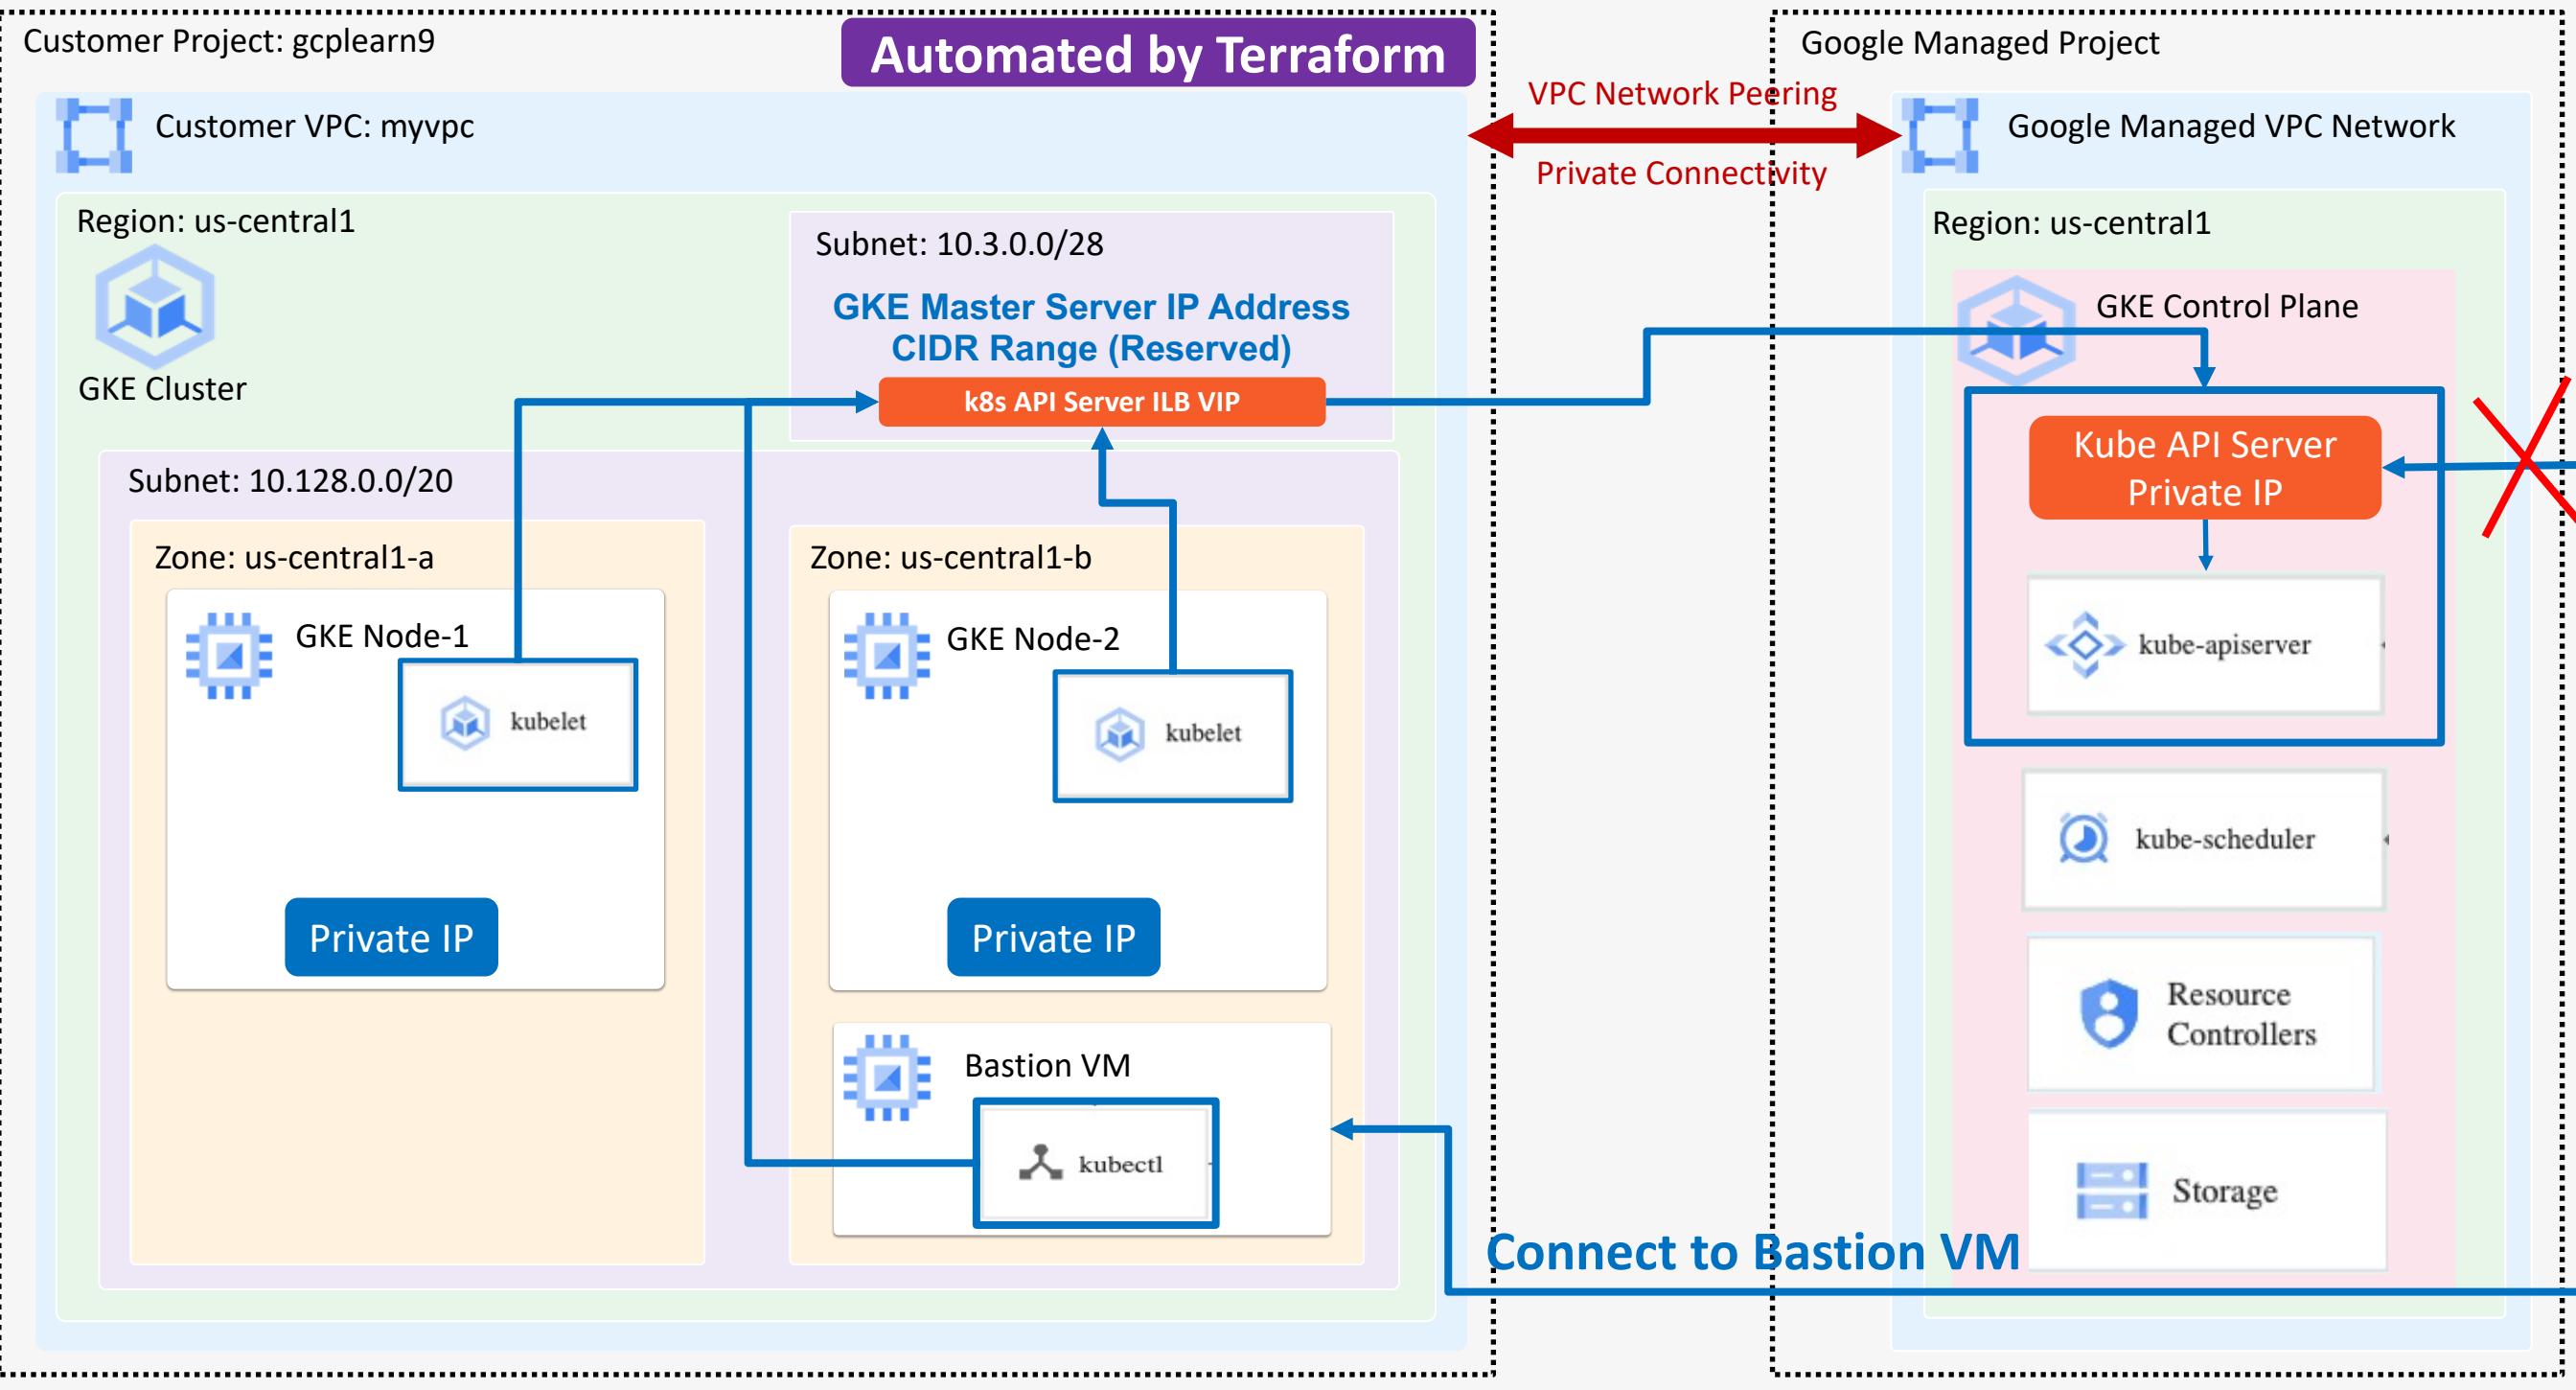
## Horizontal Pod Autoscaling

# Kubernetes Vertical Pod Autoscaling

## Vertical Pod Autoscaling



# GKE Standard Private Cluster with Private Endpoint



# GKE Autopilot Private Cluster - Network Design

Customer Project: gcplearn9

Customer VPC: myvpc

Region: us-central1

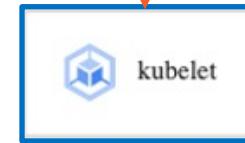


Subnet: 10.128.0.0/20

Zone: us-central1-a



GKE Node-1



Load Balancer Service

Private IP

Cloud Load Balancing

Automated by Terraform



Zone: us-central1-b



GKE Node-2



Private IP

Google Managed Project

Google Managed VPC Network

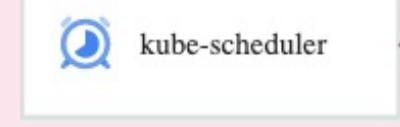
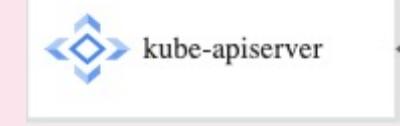
Region: us-central1



GKE Control Plane

Kube API Server Public IP

Kube API Server Private IP



VPC Network Peering

Private Connectivity

Private Connectivity

Download Docker Image from Docker Hub

http://LB-IP



Users



kubectl



Admin

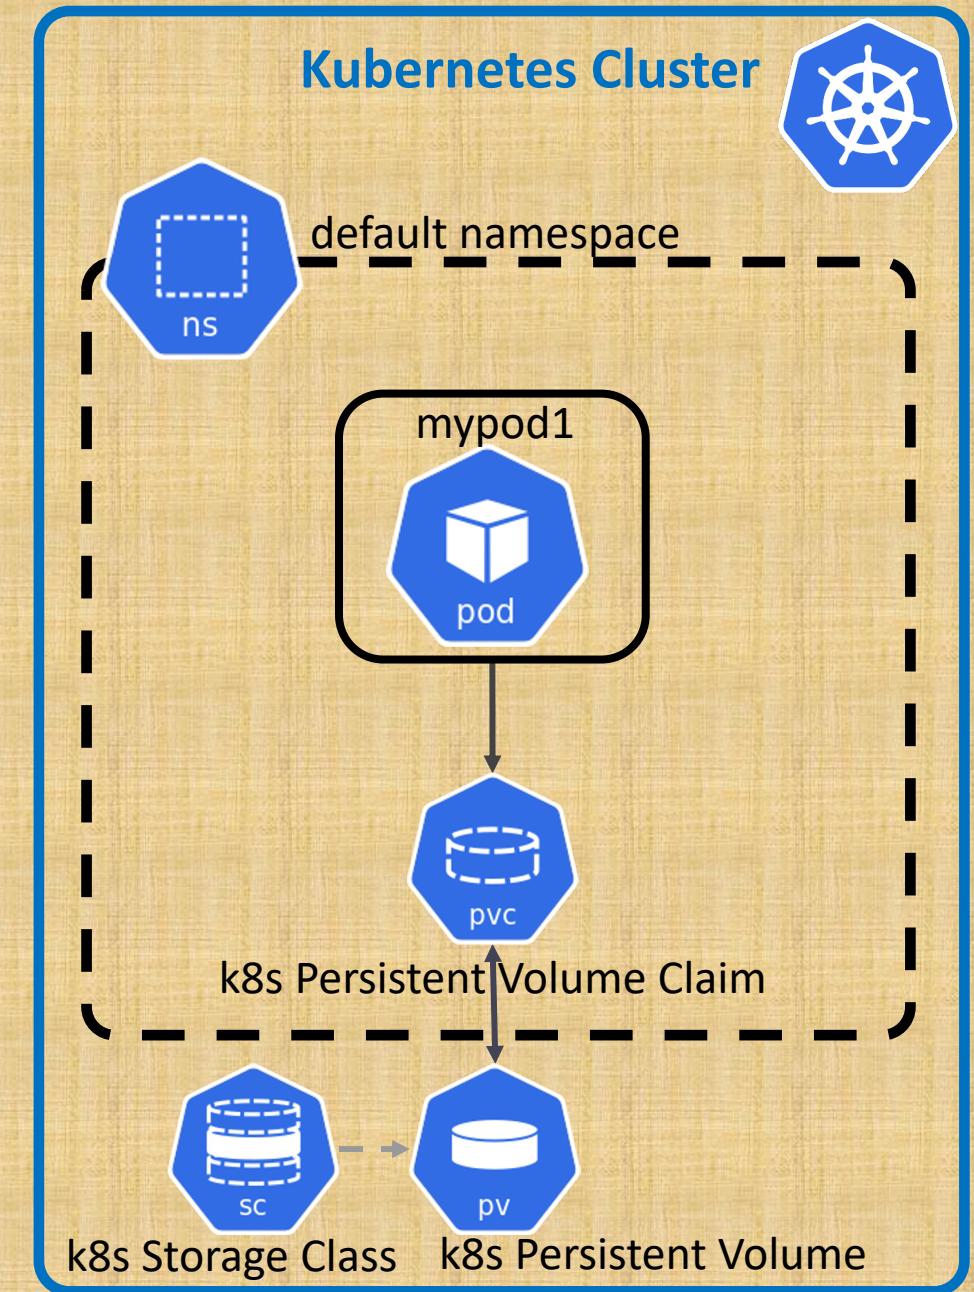
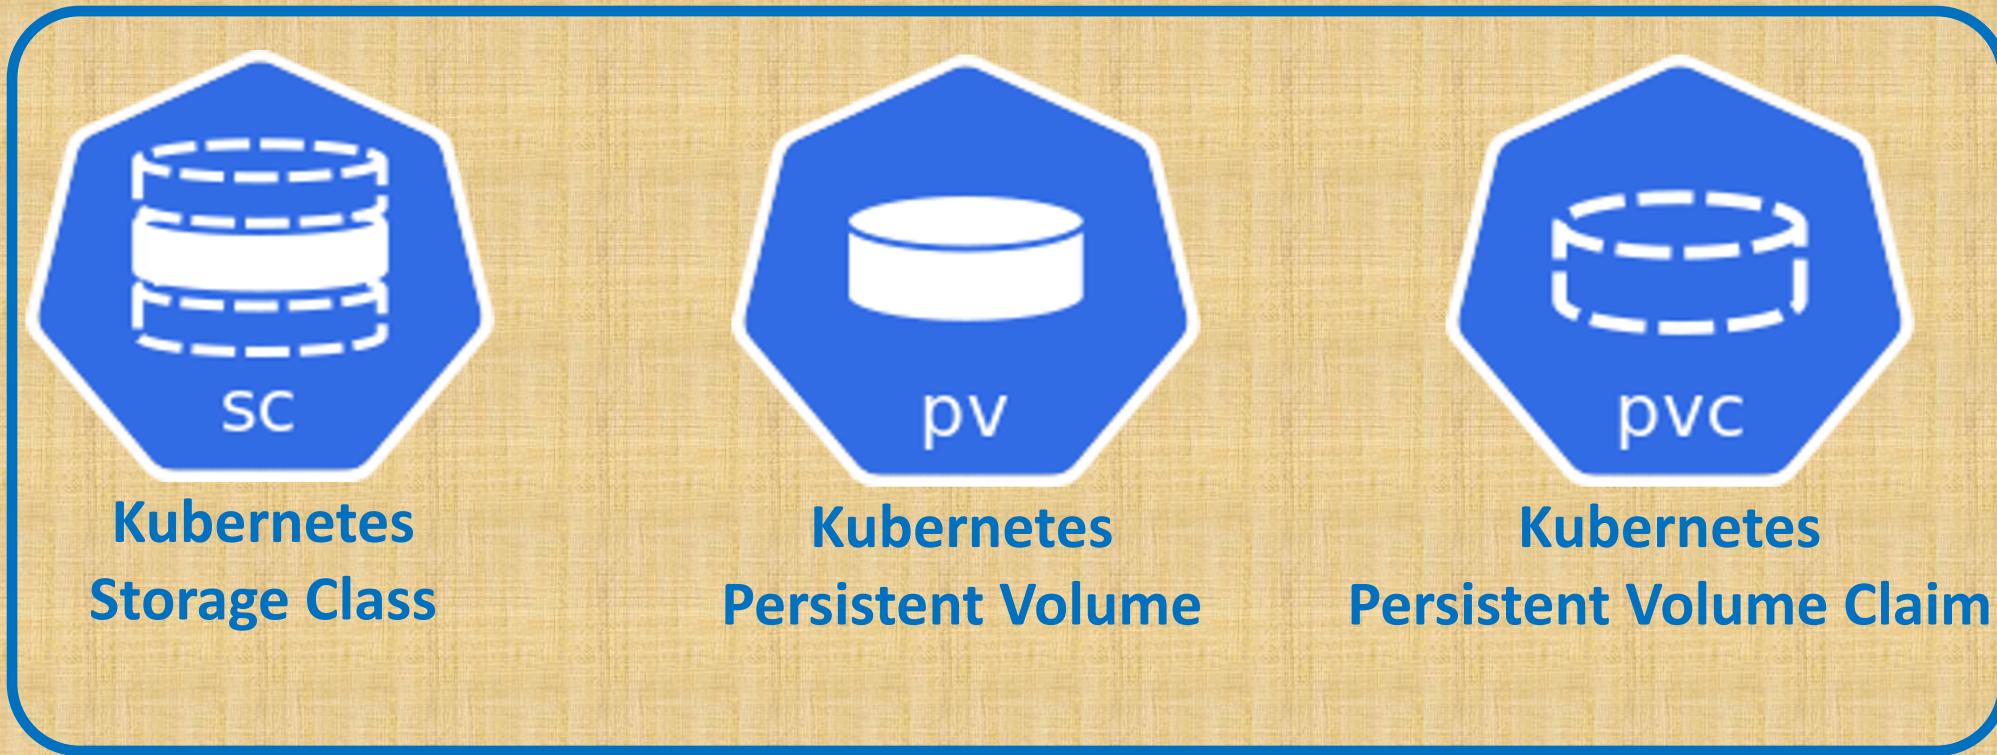


Docker Hub

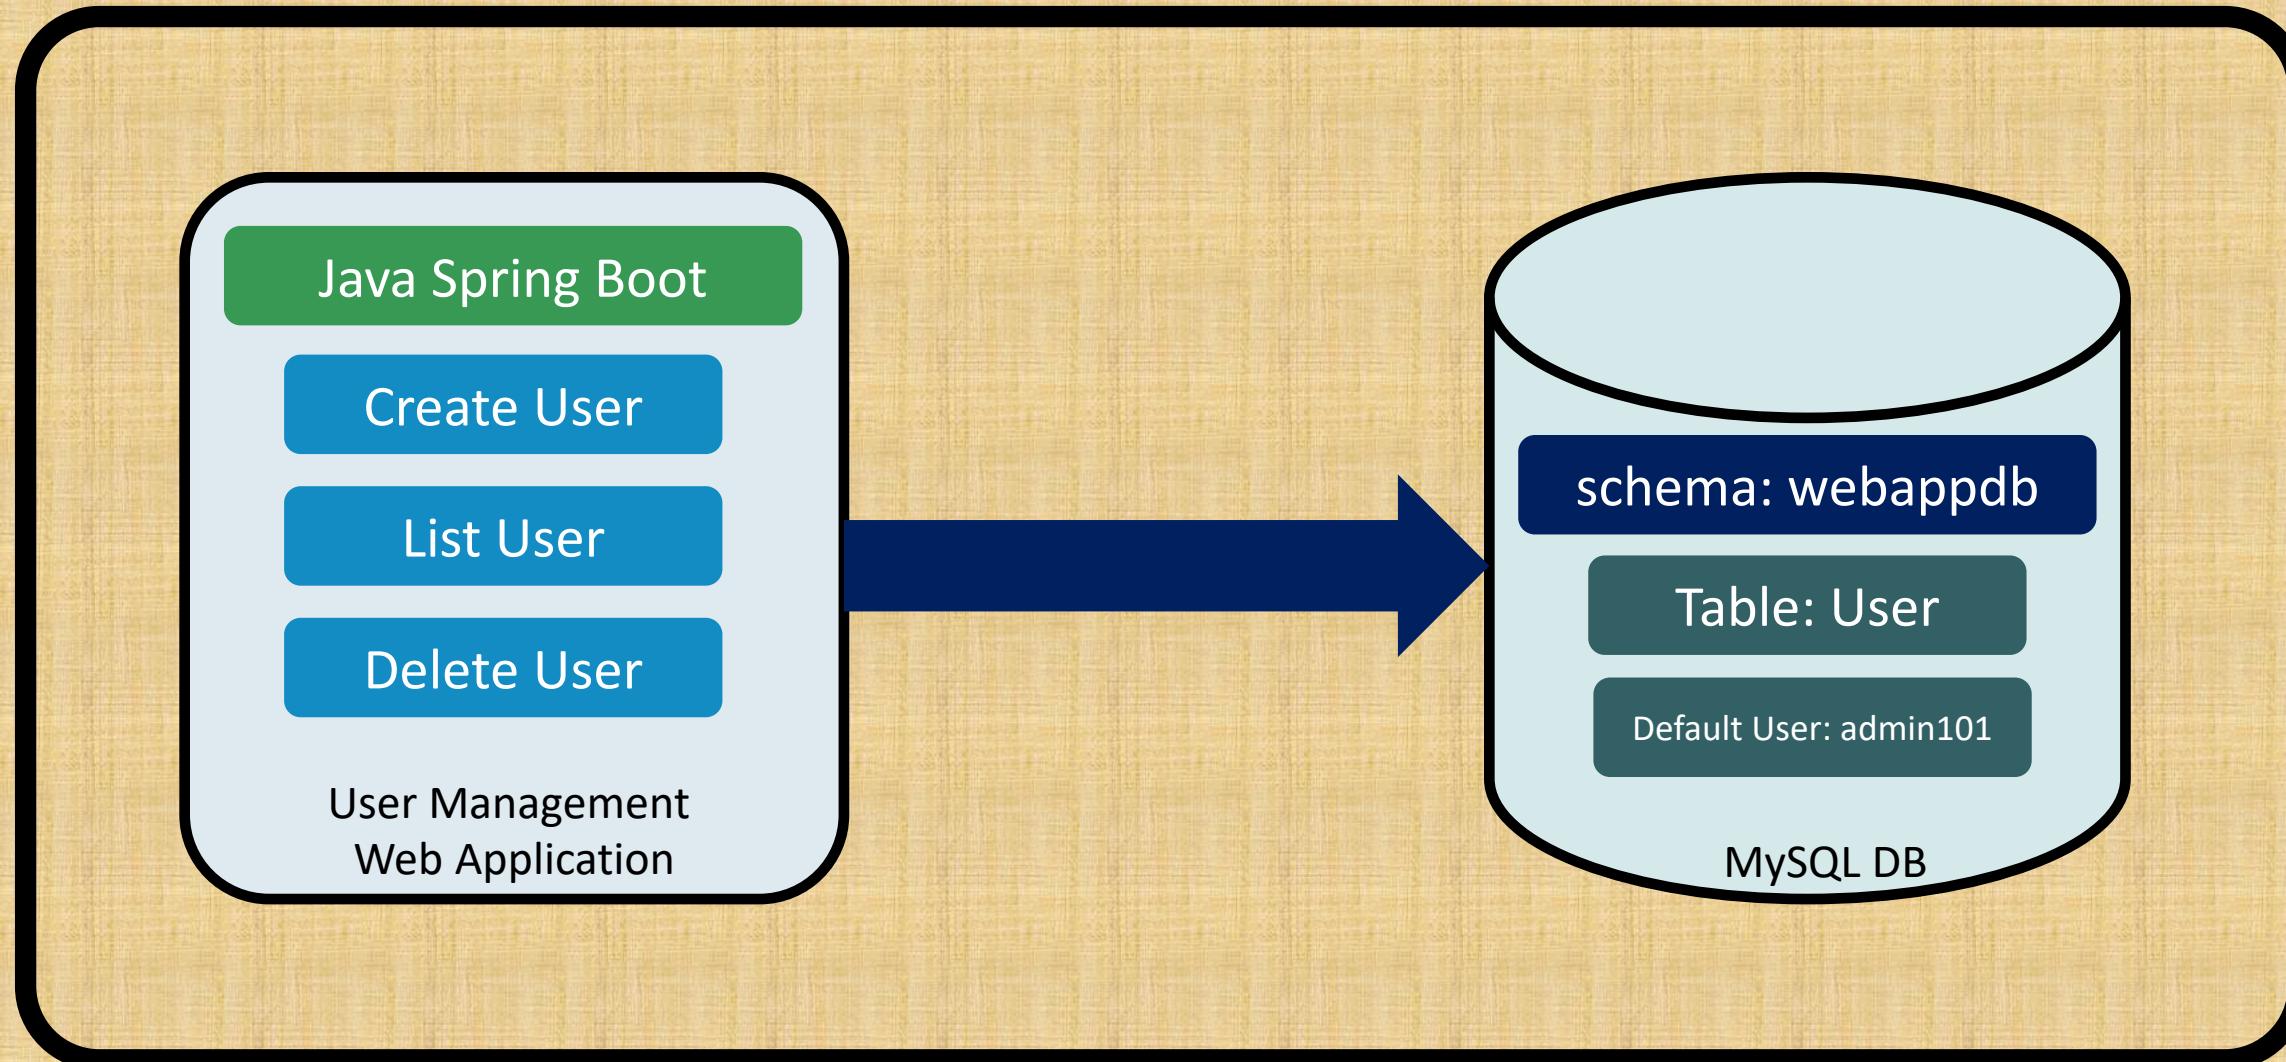


NETWORK DESIGN

# Kubernetes Storage



# User Management Web Application - Architecture

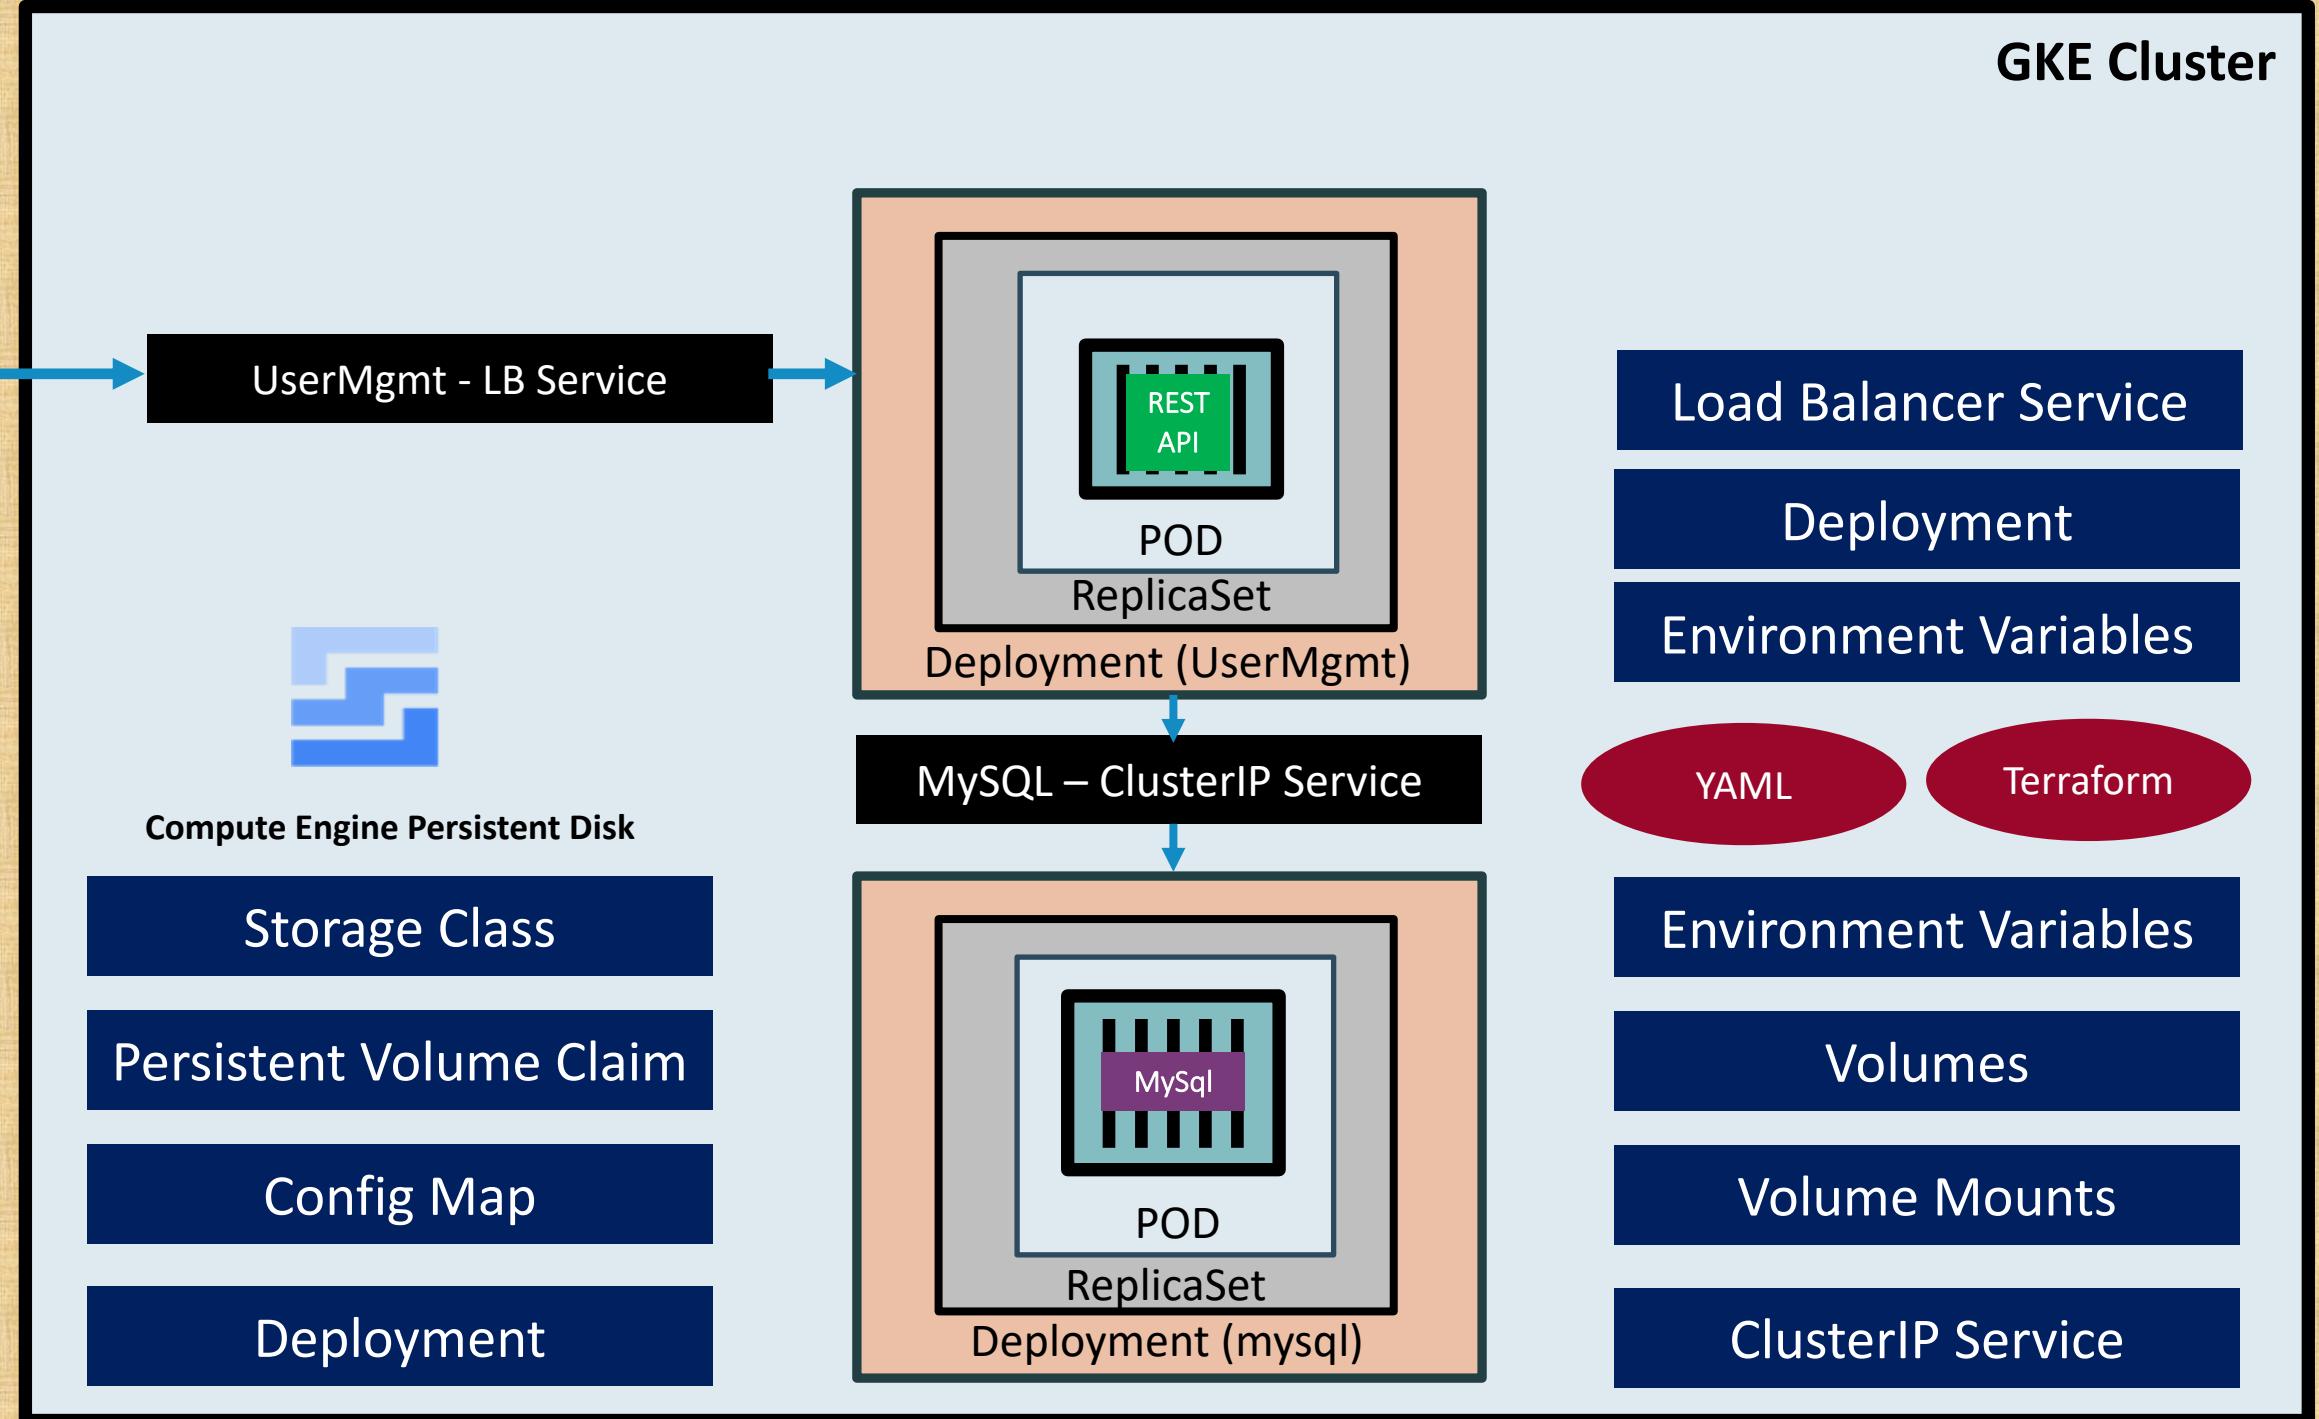


# GKE Storage with Compute Engine Persistence Disks

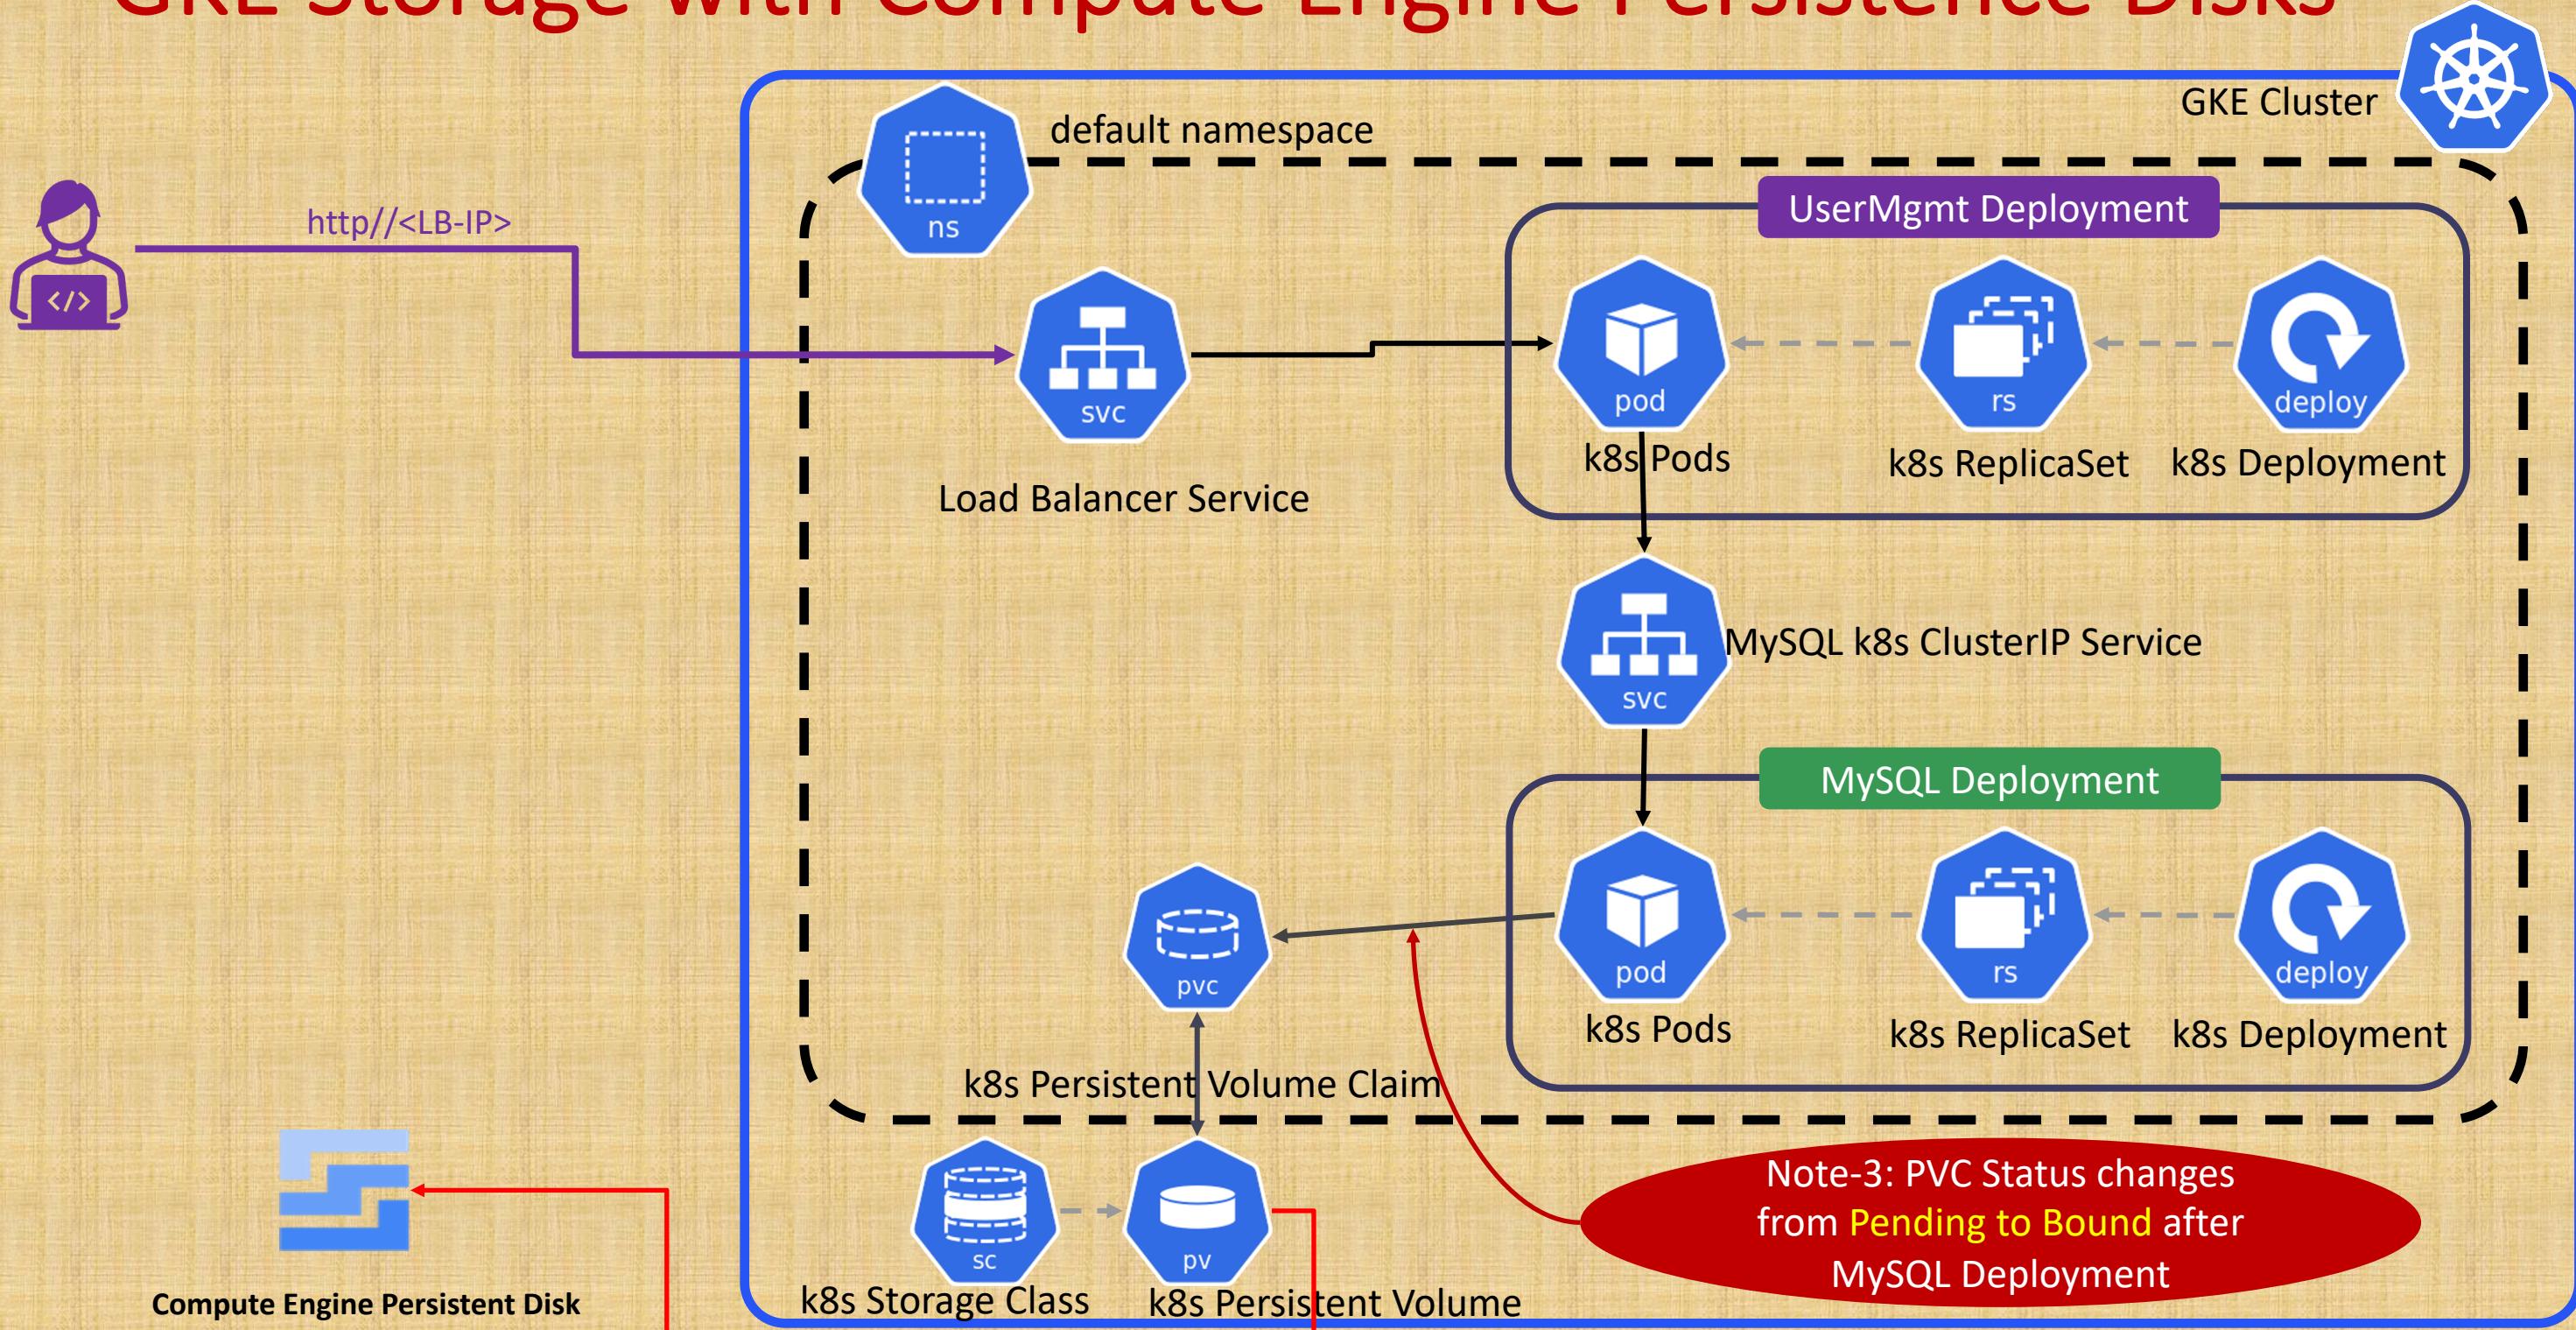


http://<LB-IP>

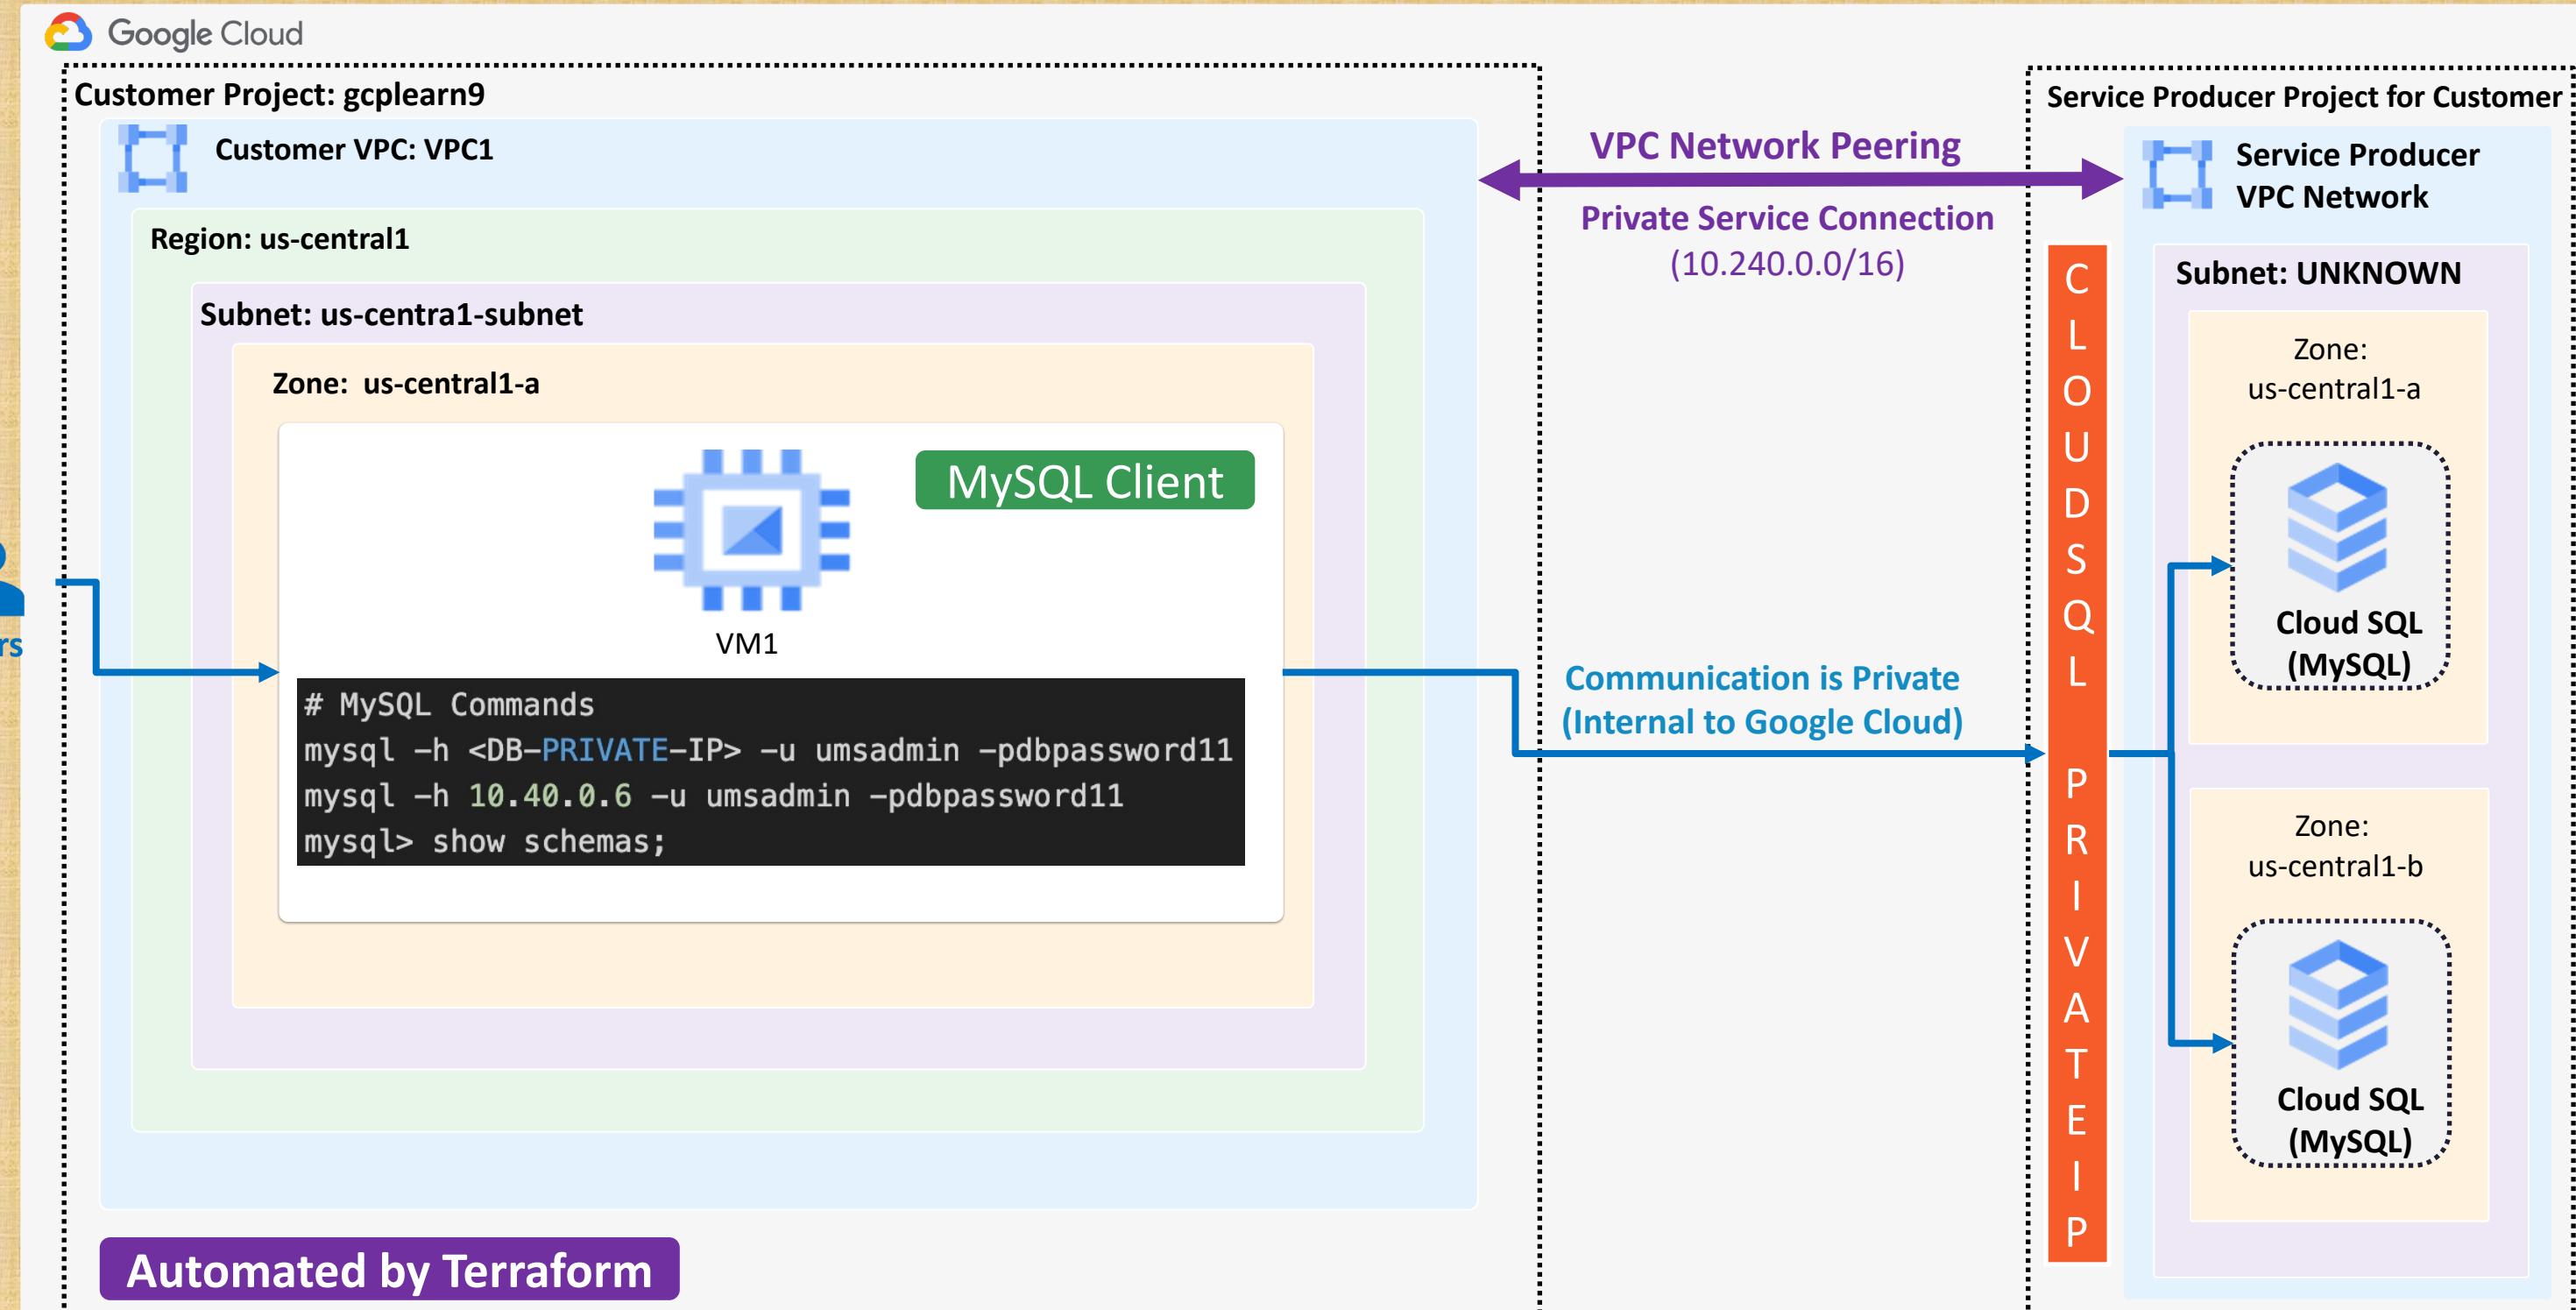
GKE Storage  
Compute  
Engine  
Persistent Disk  
CSI Driver



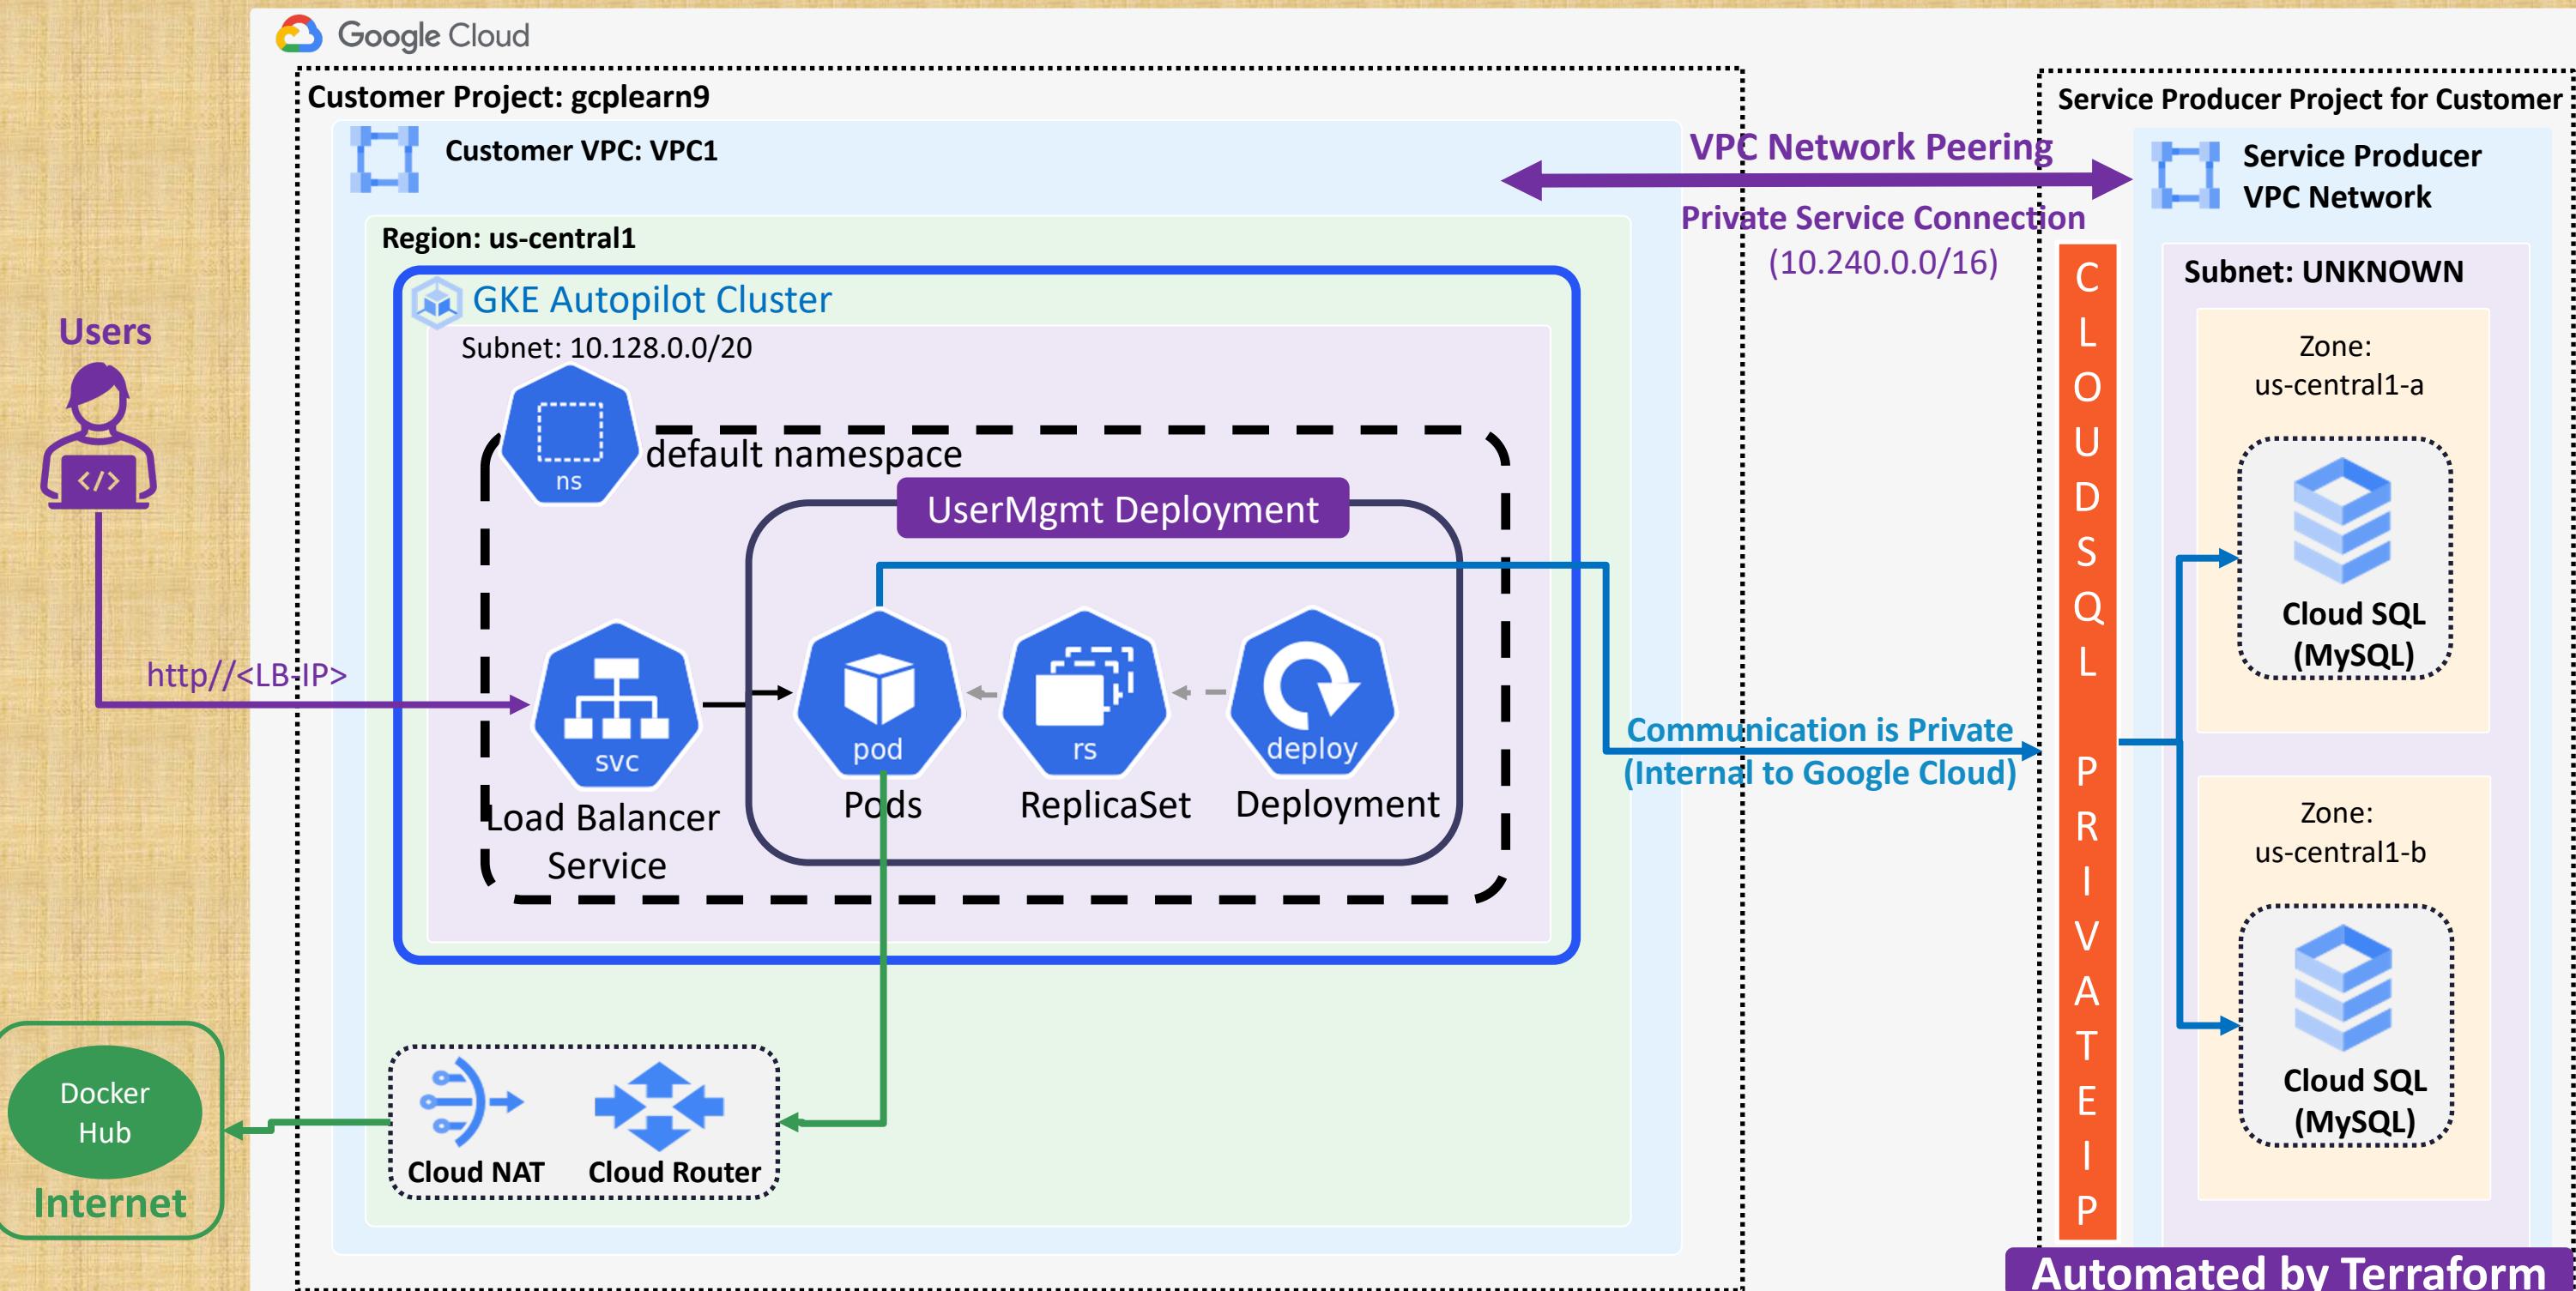
# GKE Storage with Compute Engine Persistence Disks



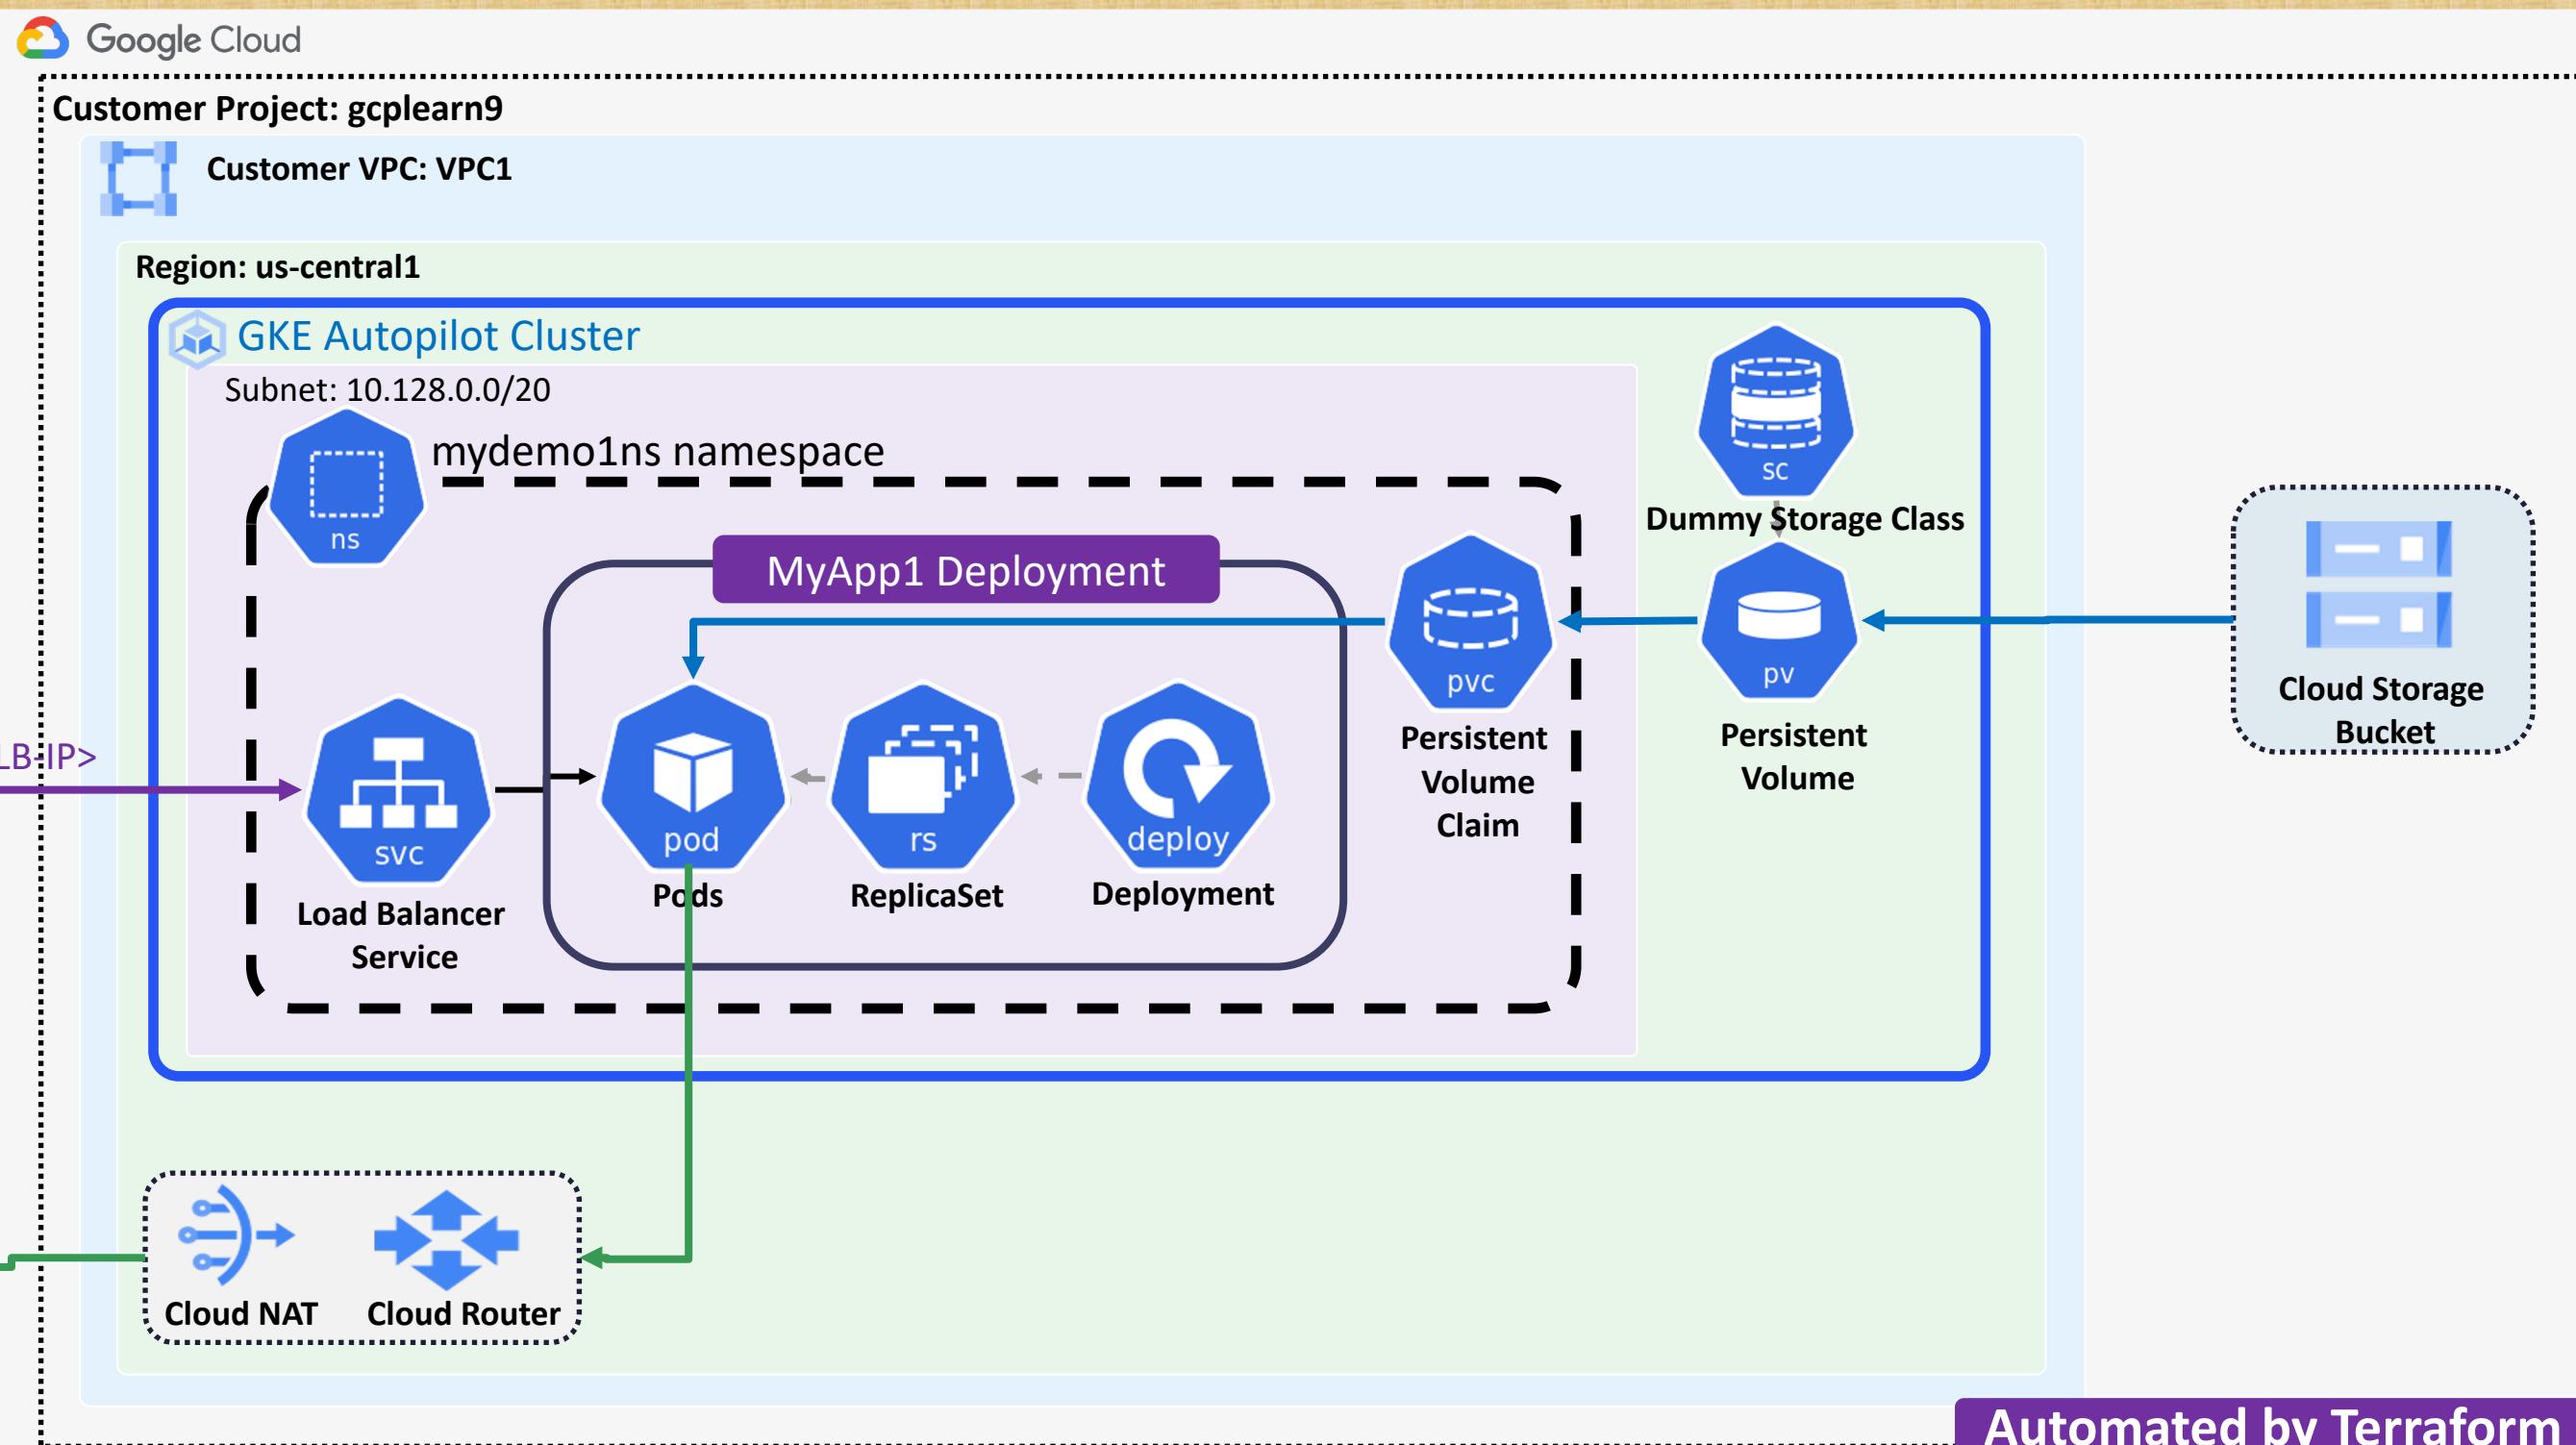
# GCP Cloud SQL Private IP



# GCP Cloud SQL Private IP + GKE Autopilot Cluster + User Management Web Application

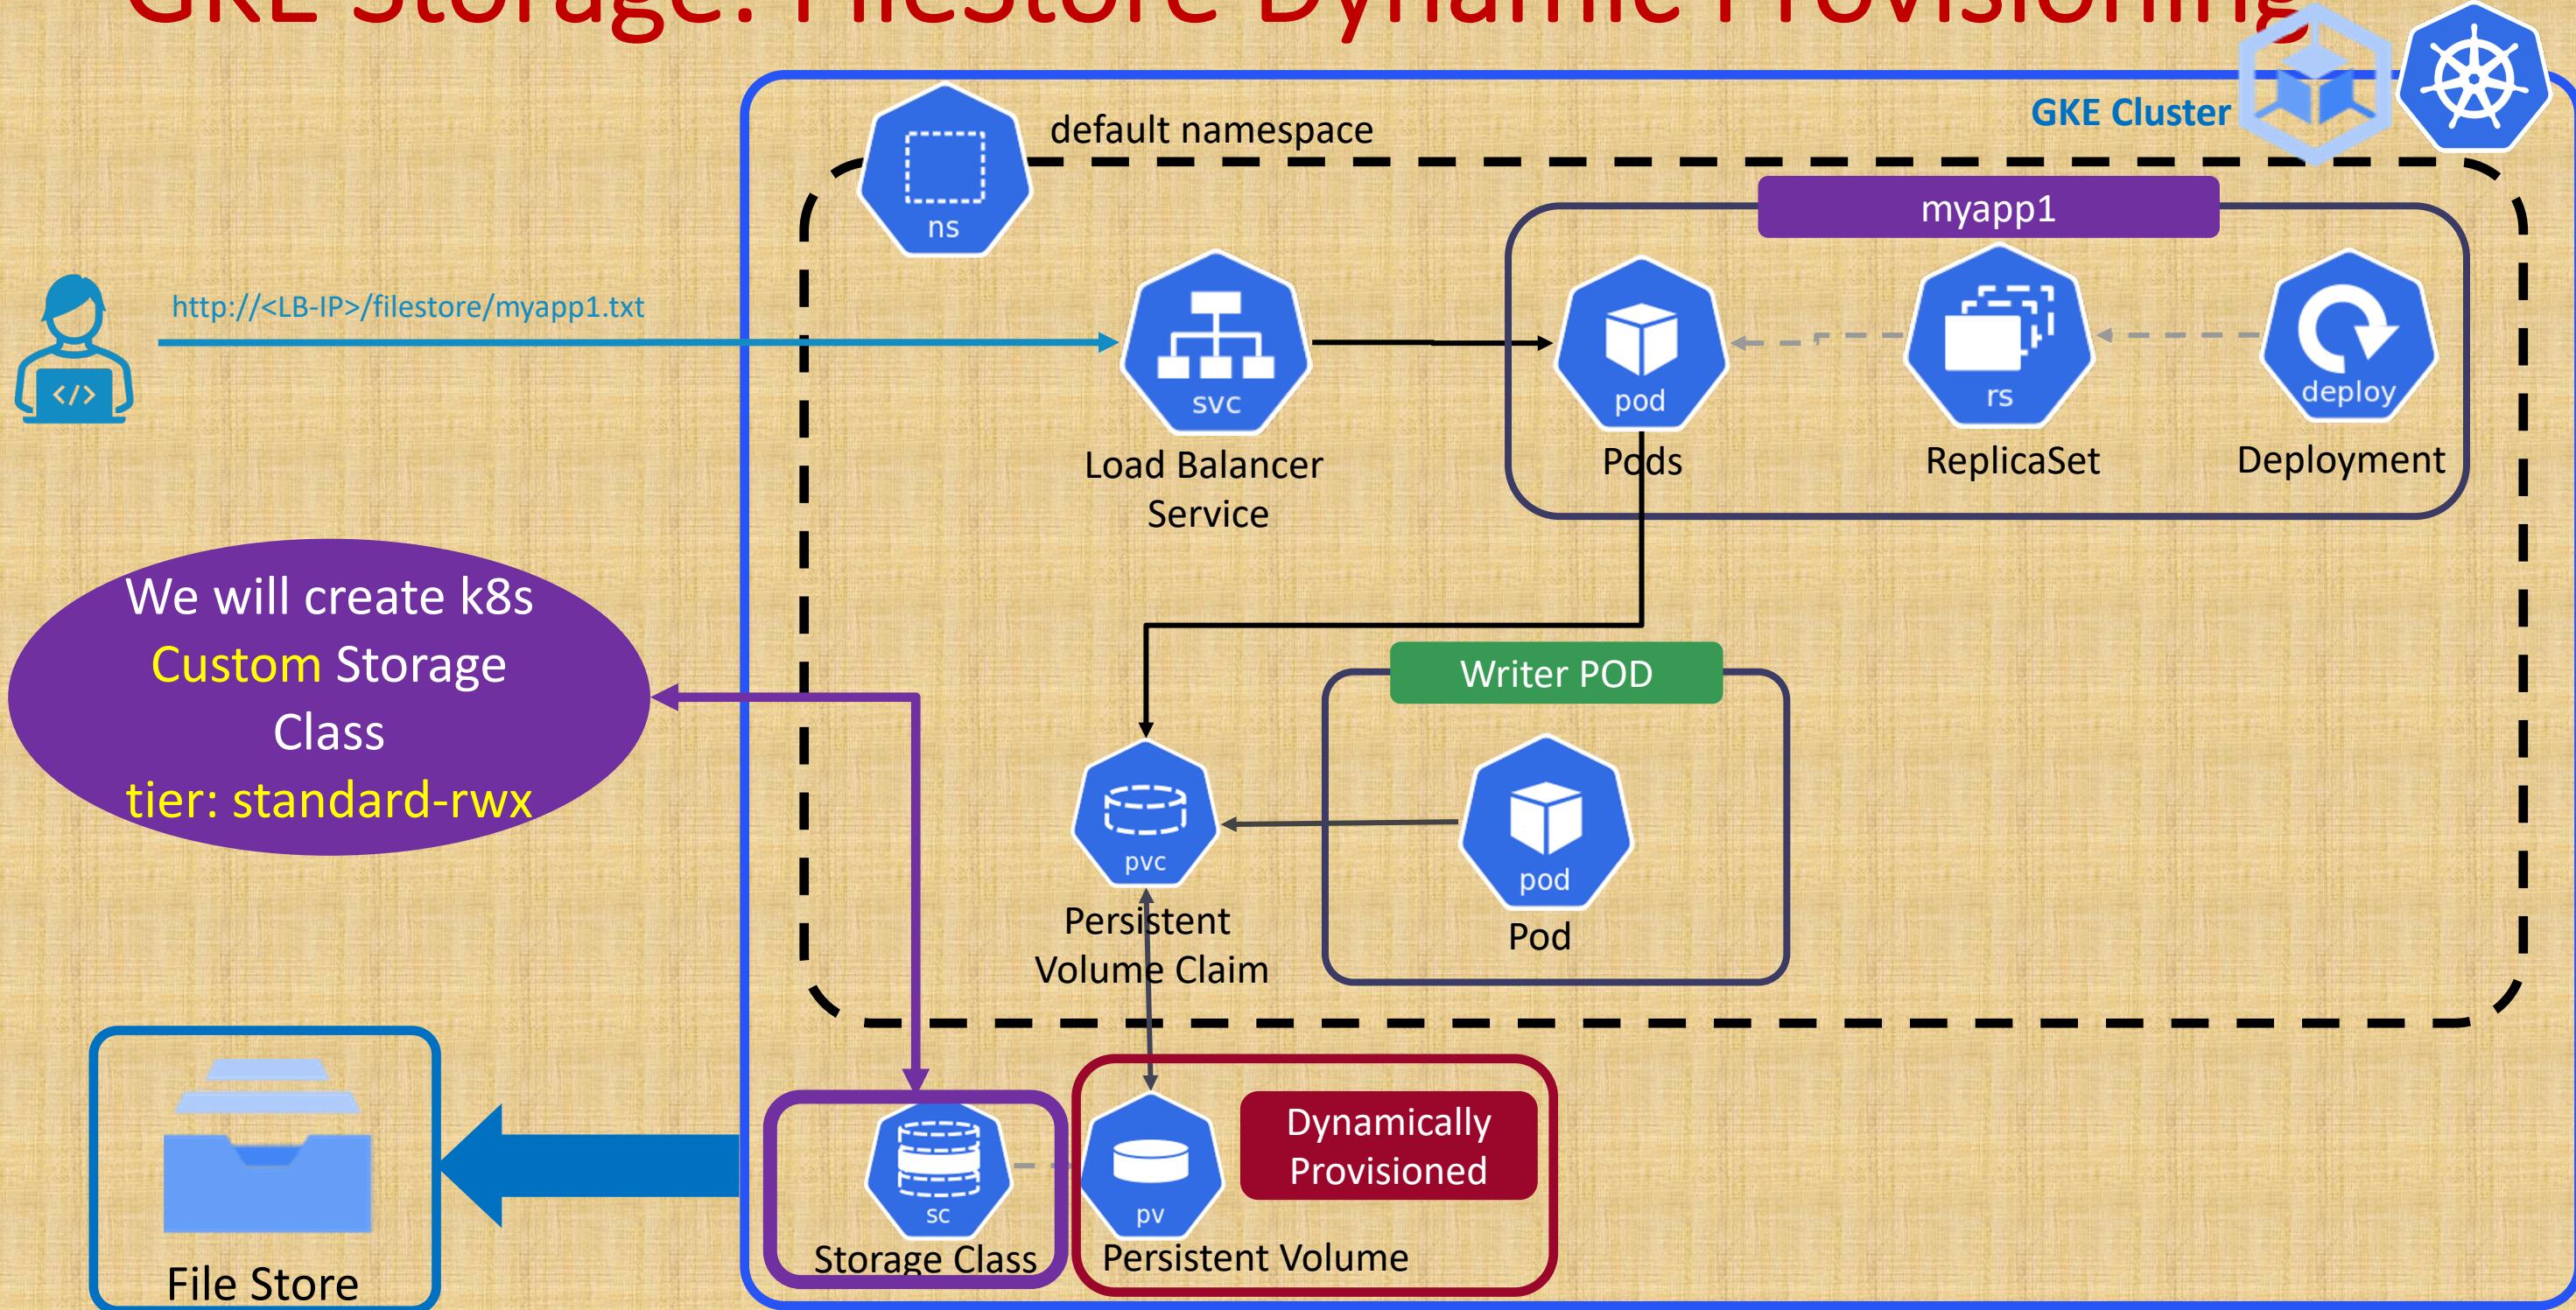


# GKE Autopilot Cluster + Cloud Storage Bucket as Volume Mount

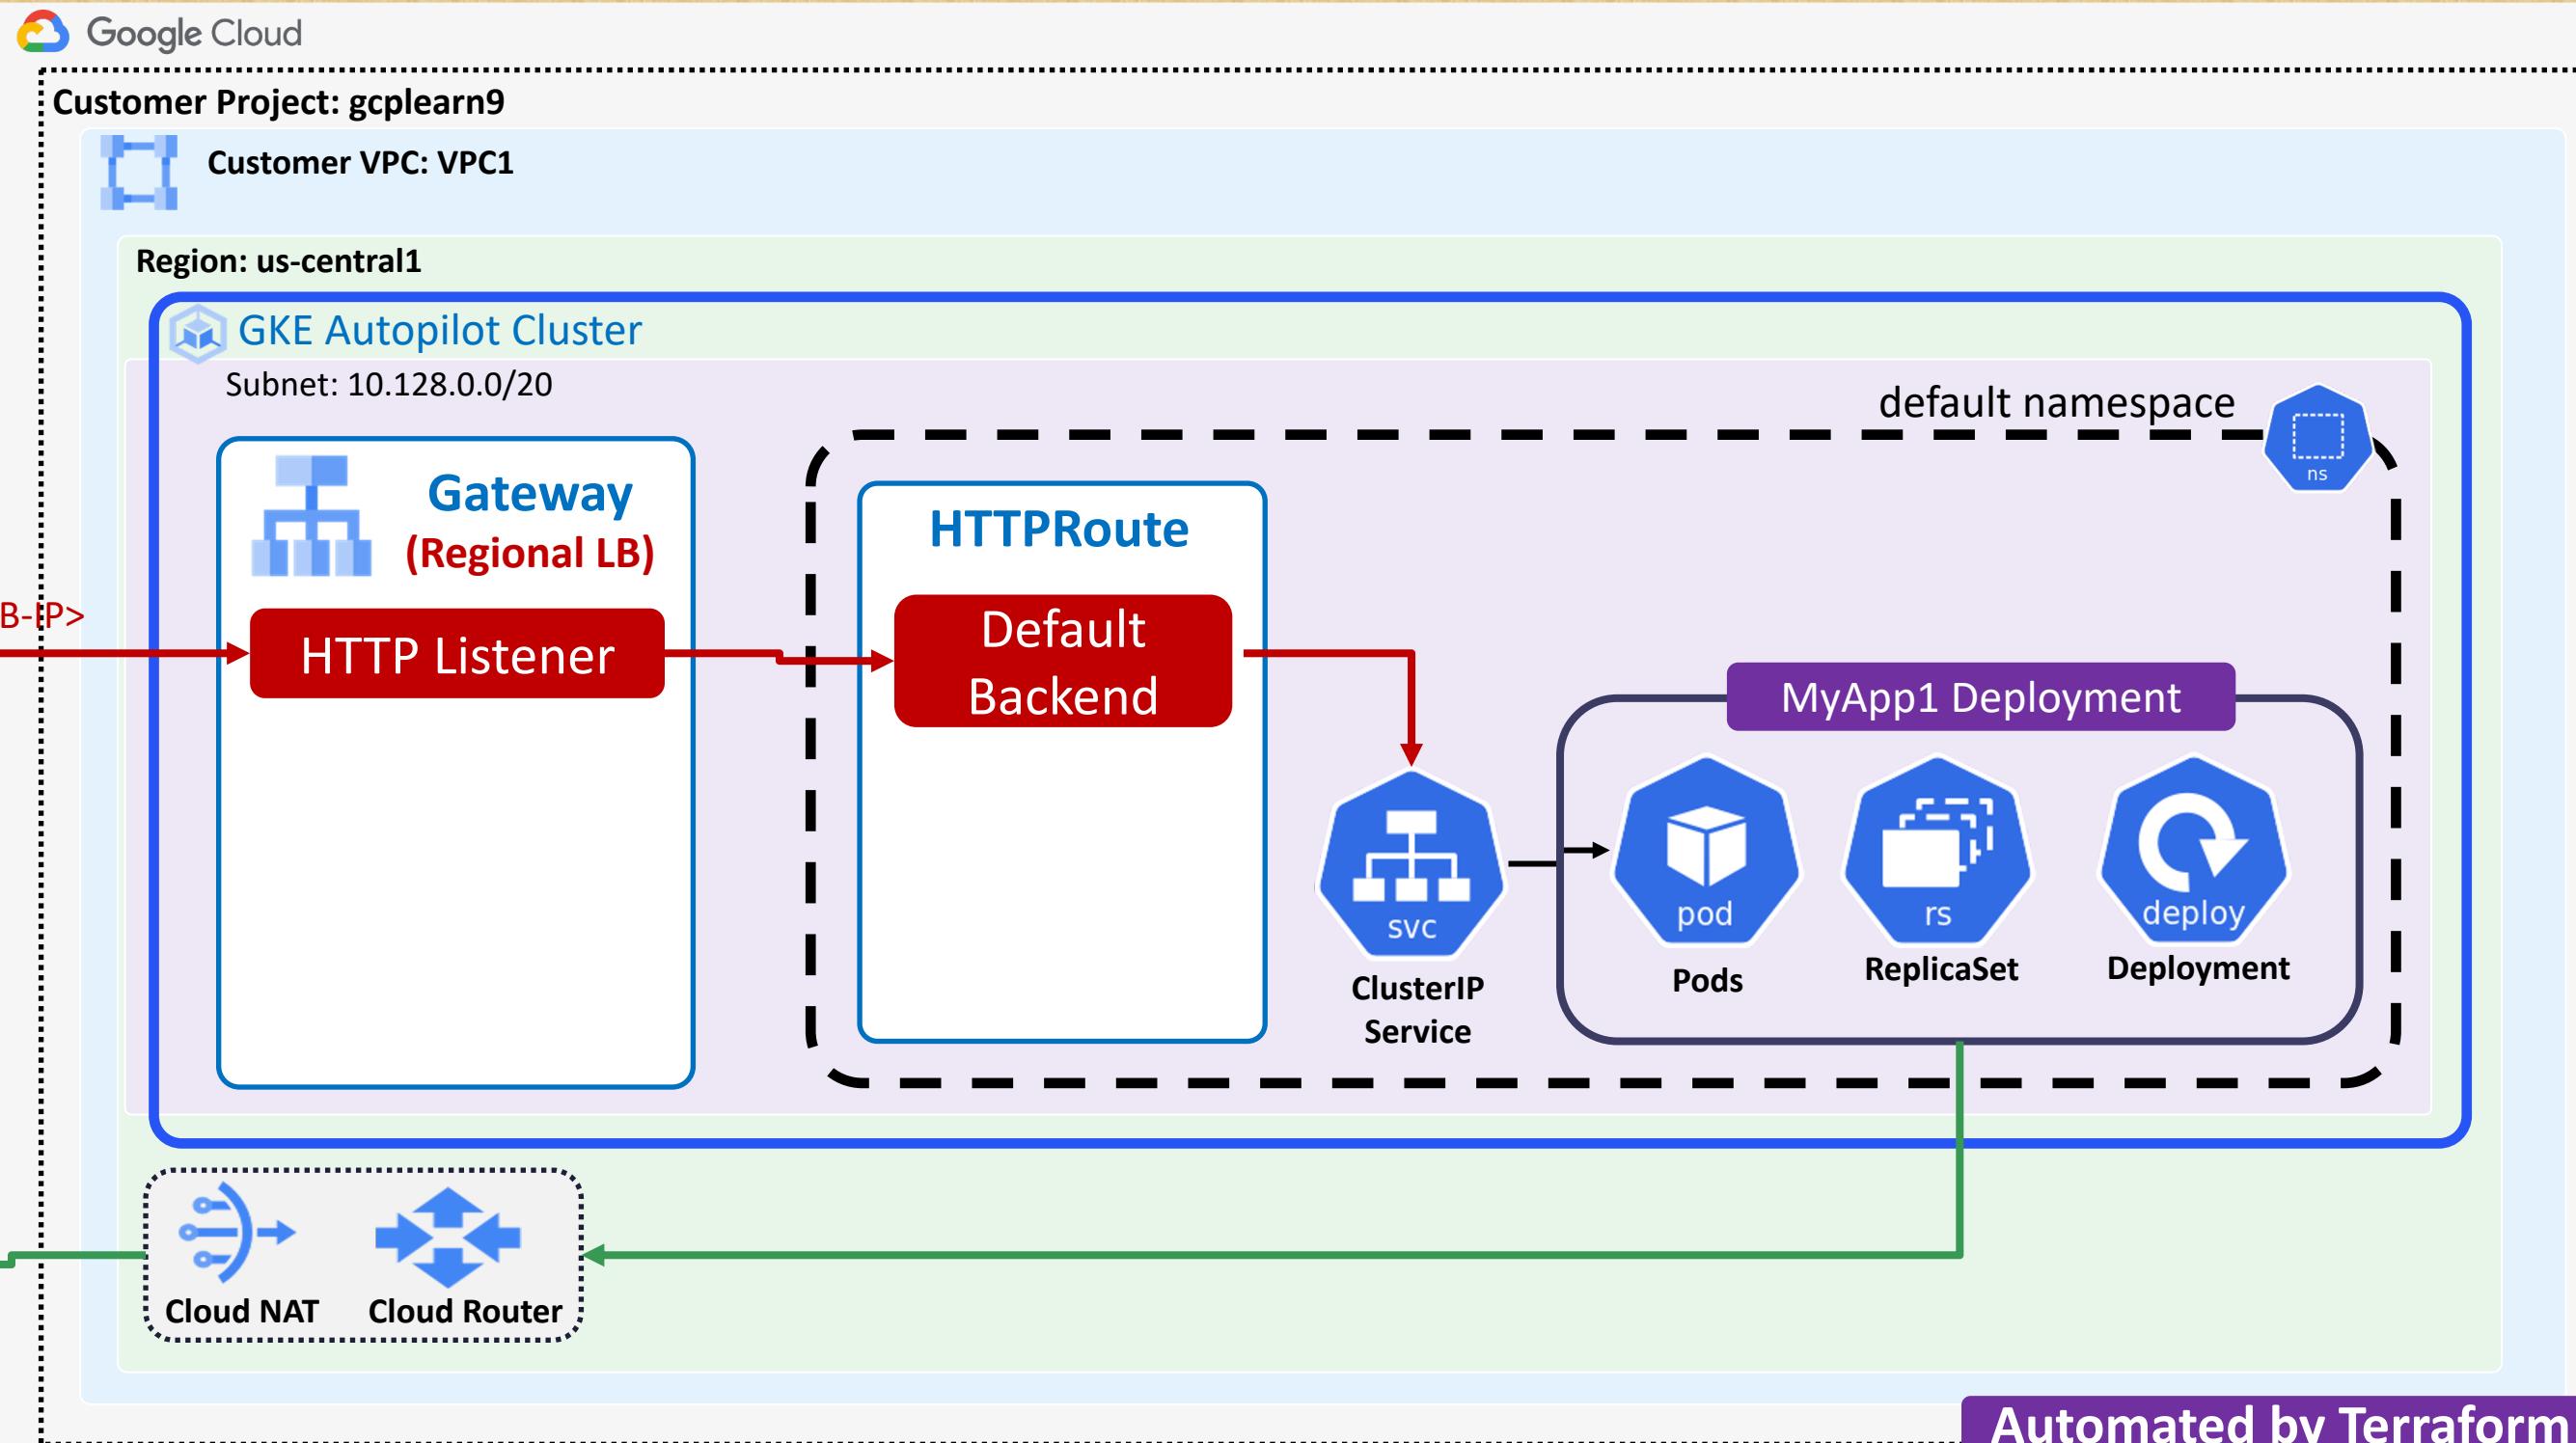


Automated by Terraform

# GKE Storage: FileStore Dynamic Provisioning

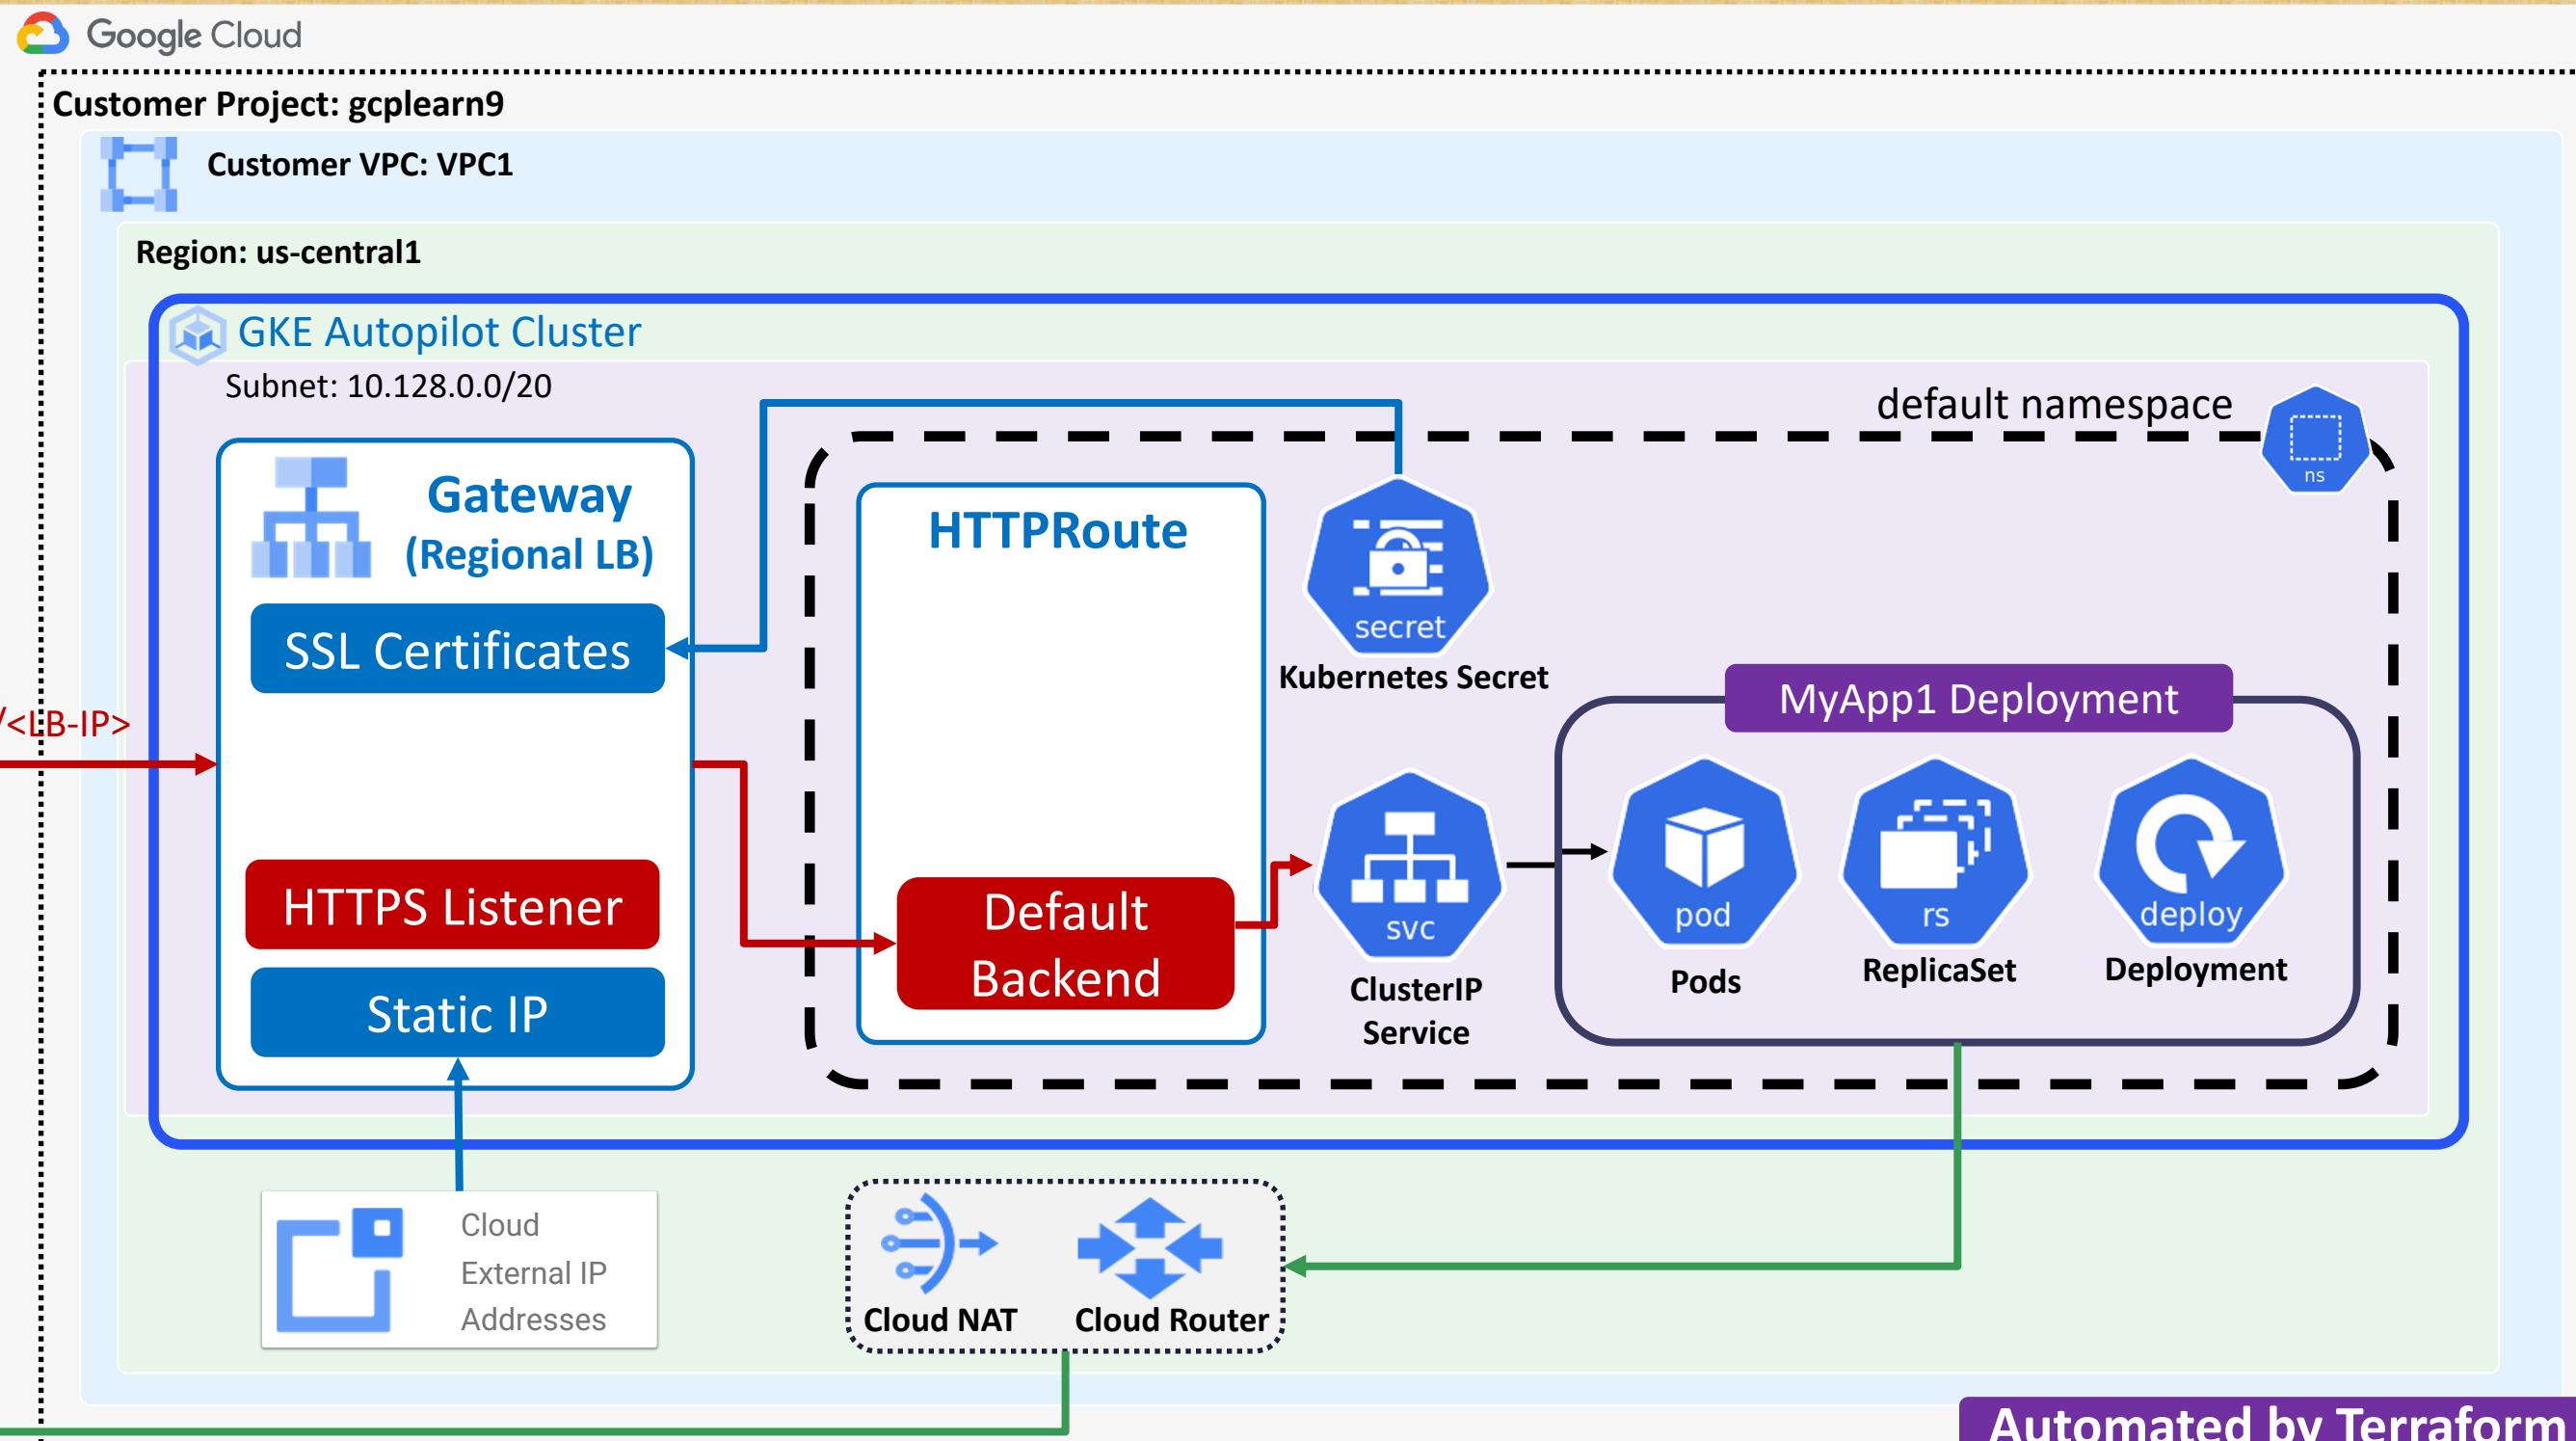


# Gateway API Basics - Regional Application Load Balancer

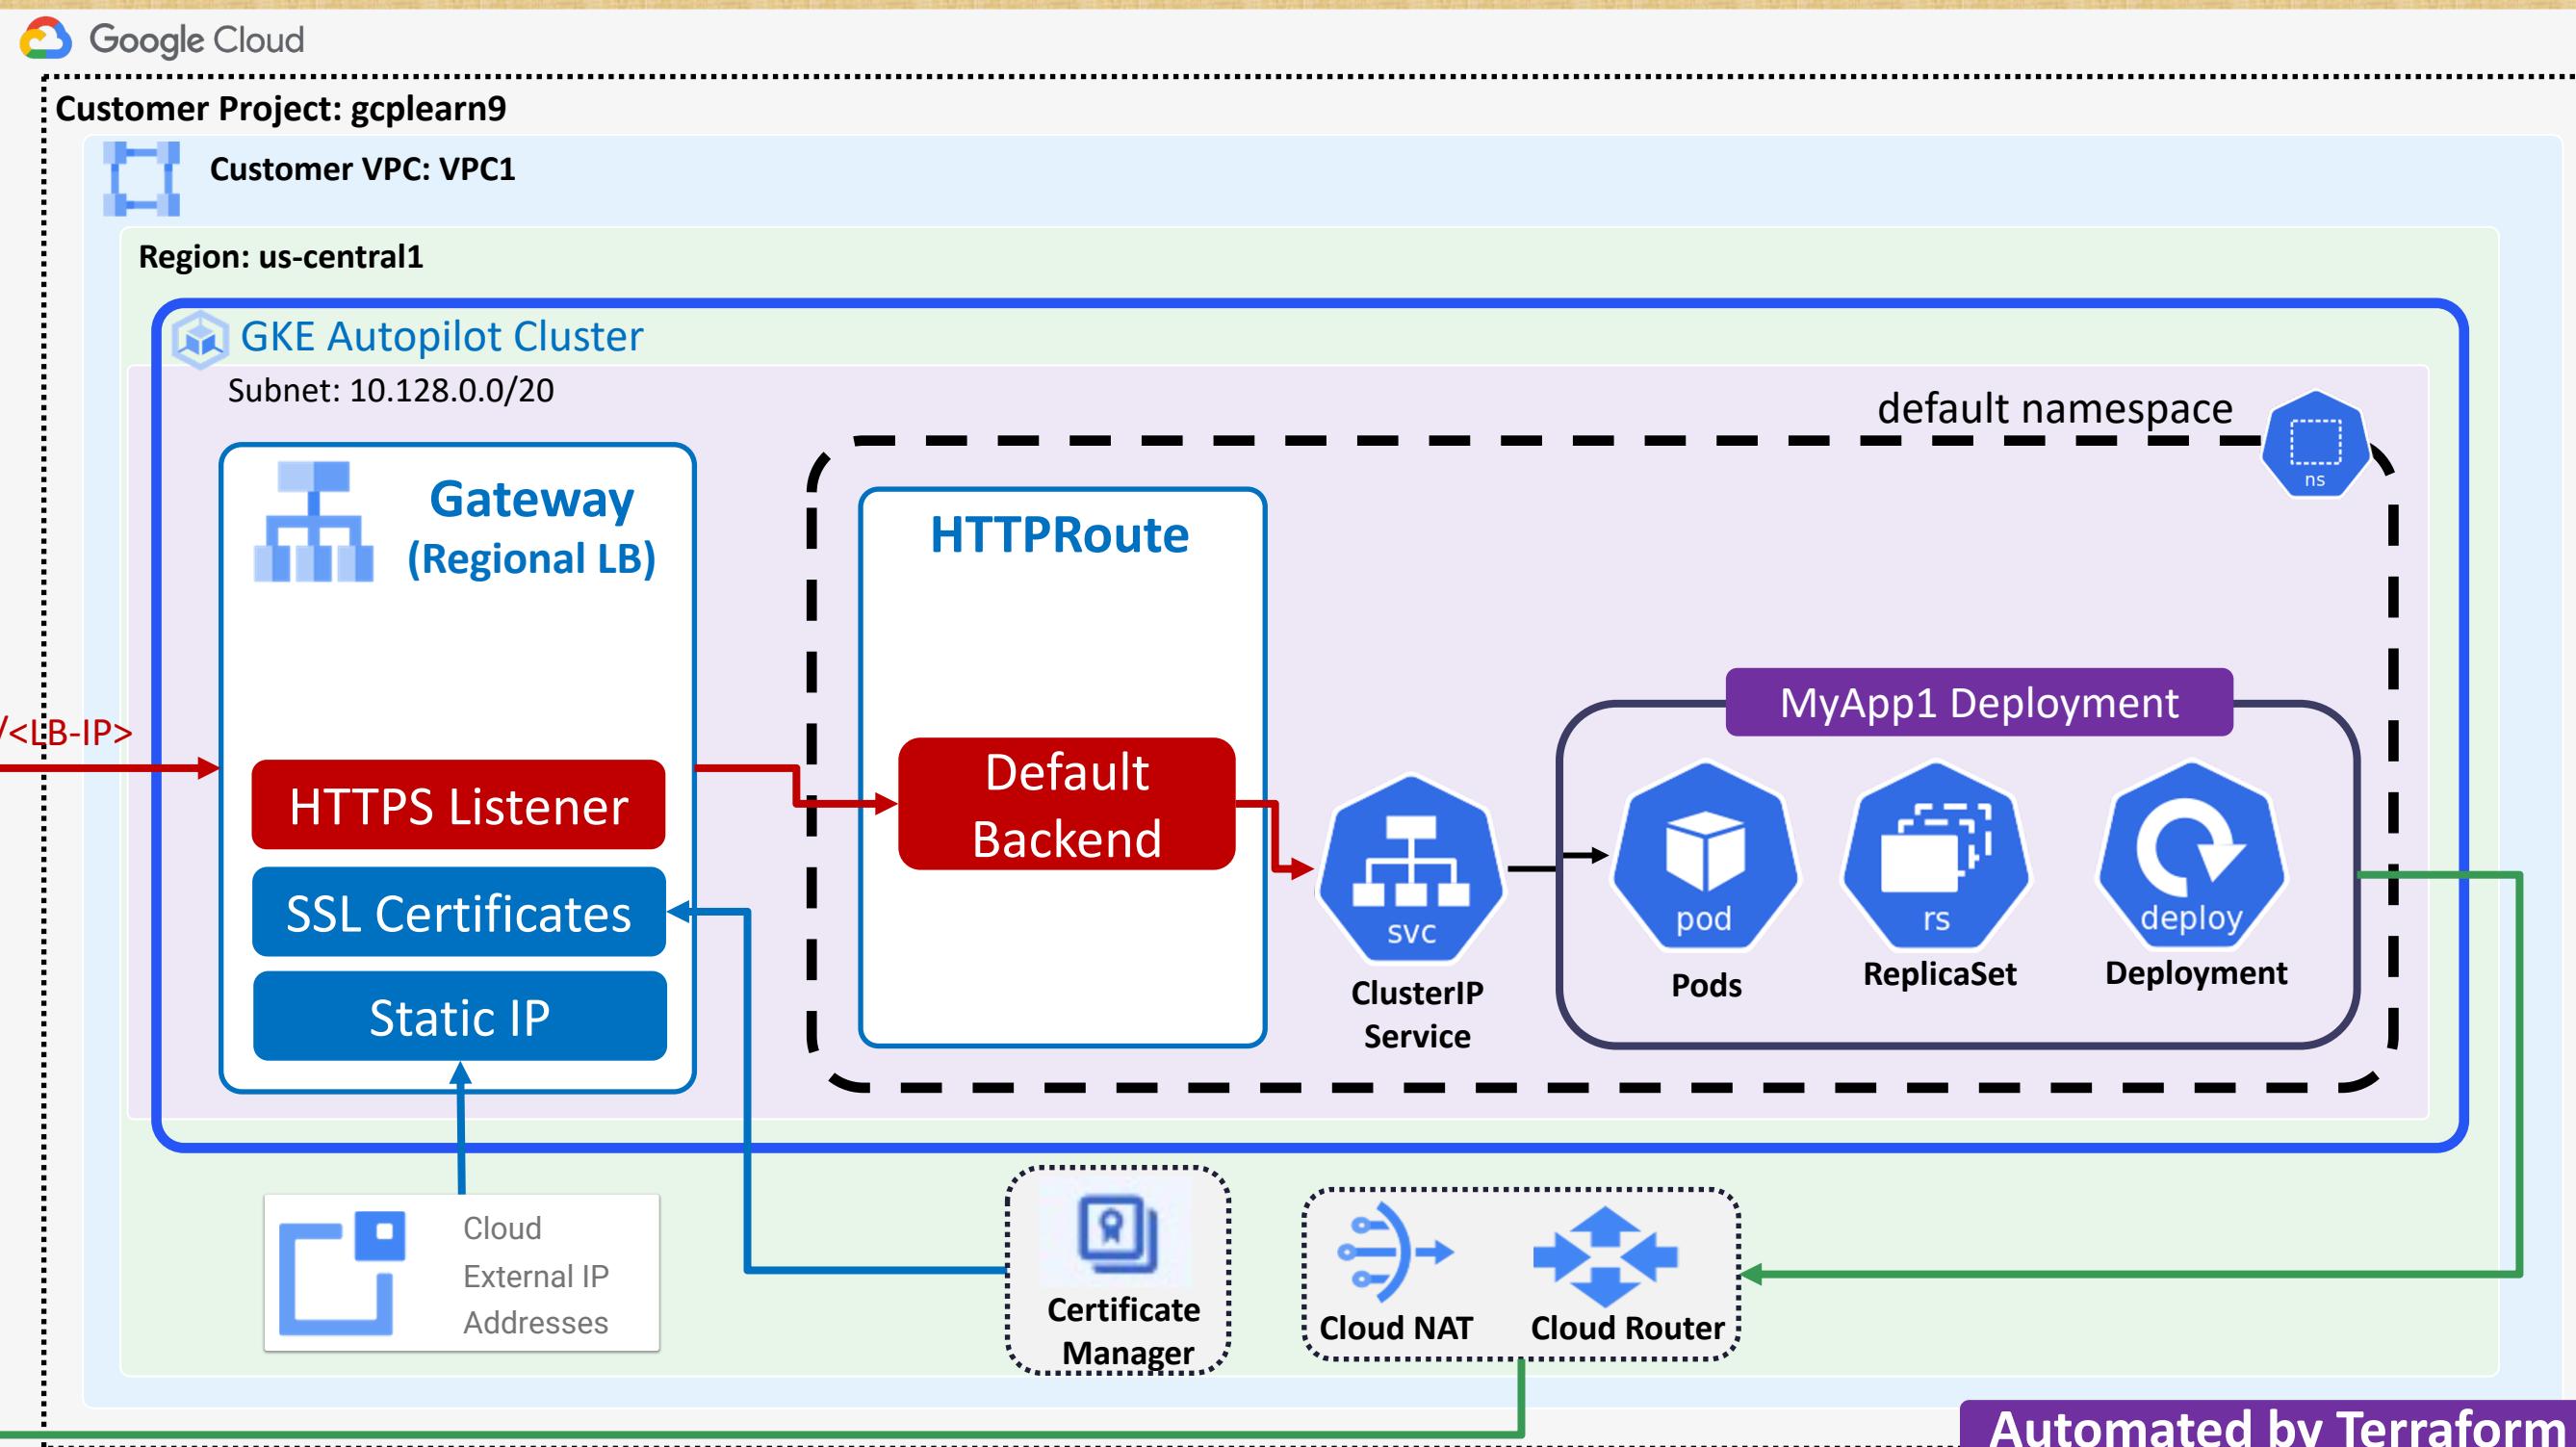


Automated by Terraform

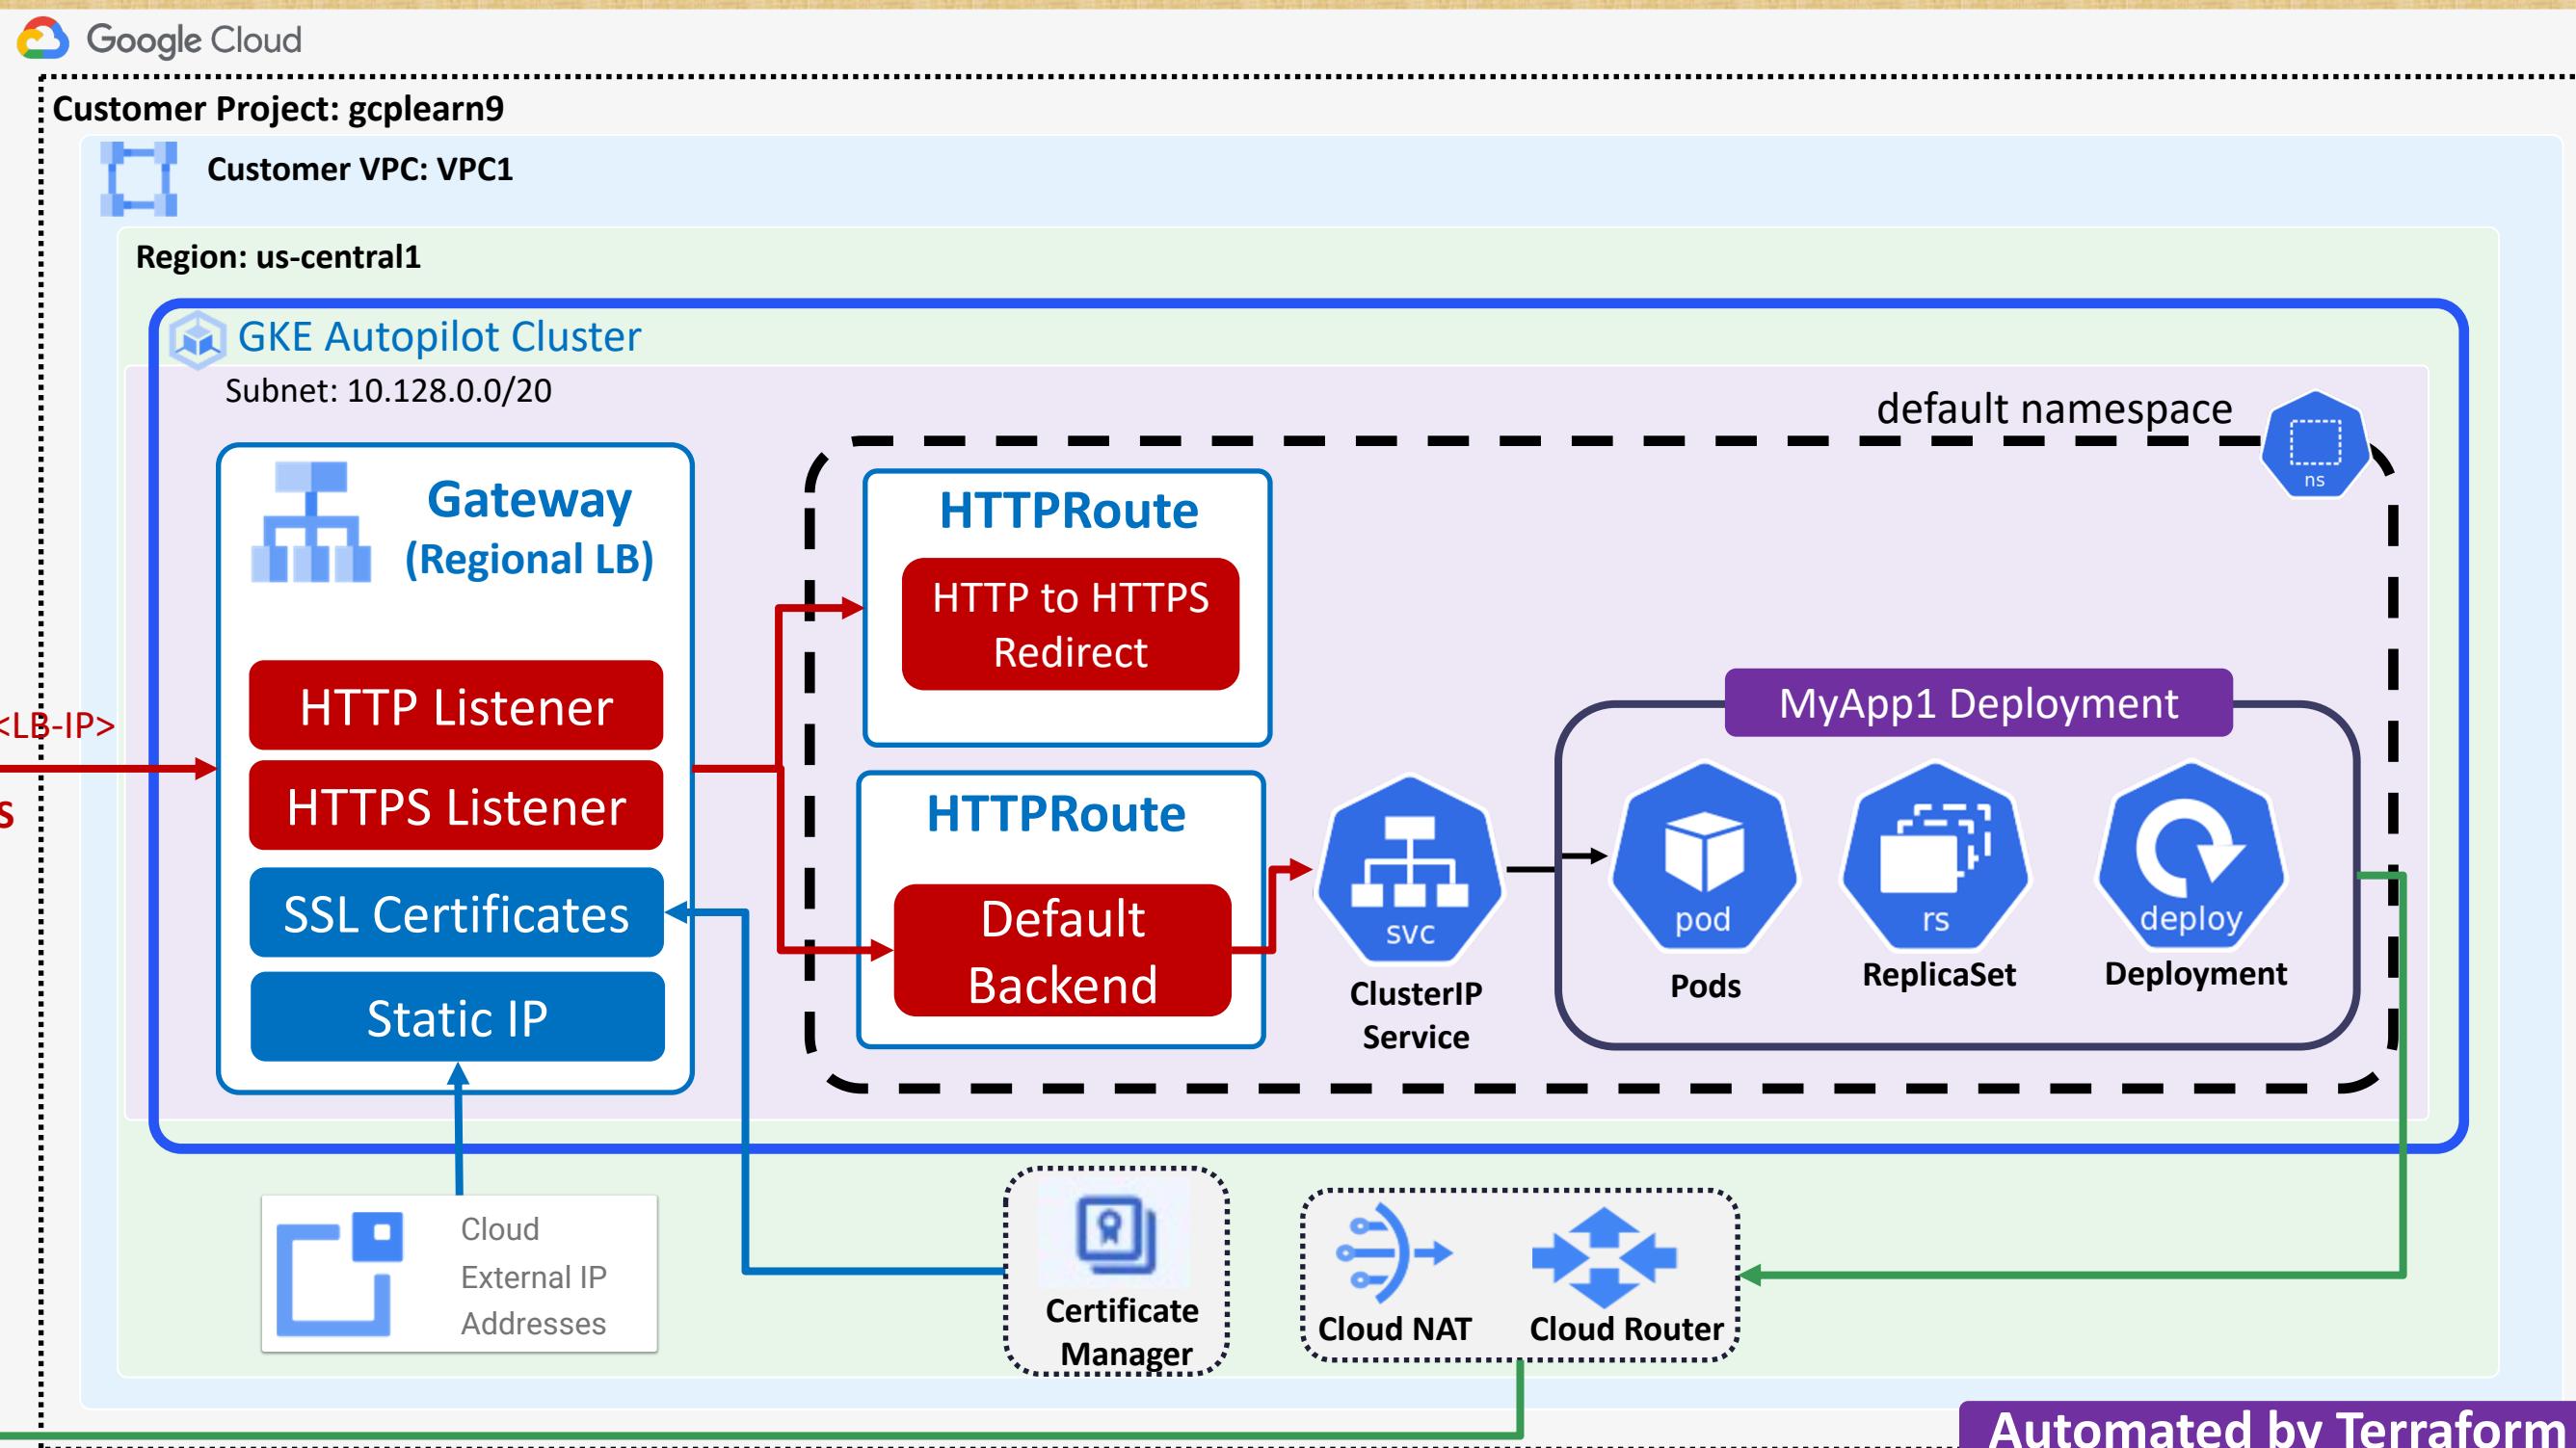
# Gateway API Basics + Static IP + Self-signed SSL with Kubernetes Secret



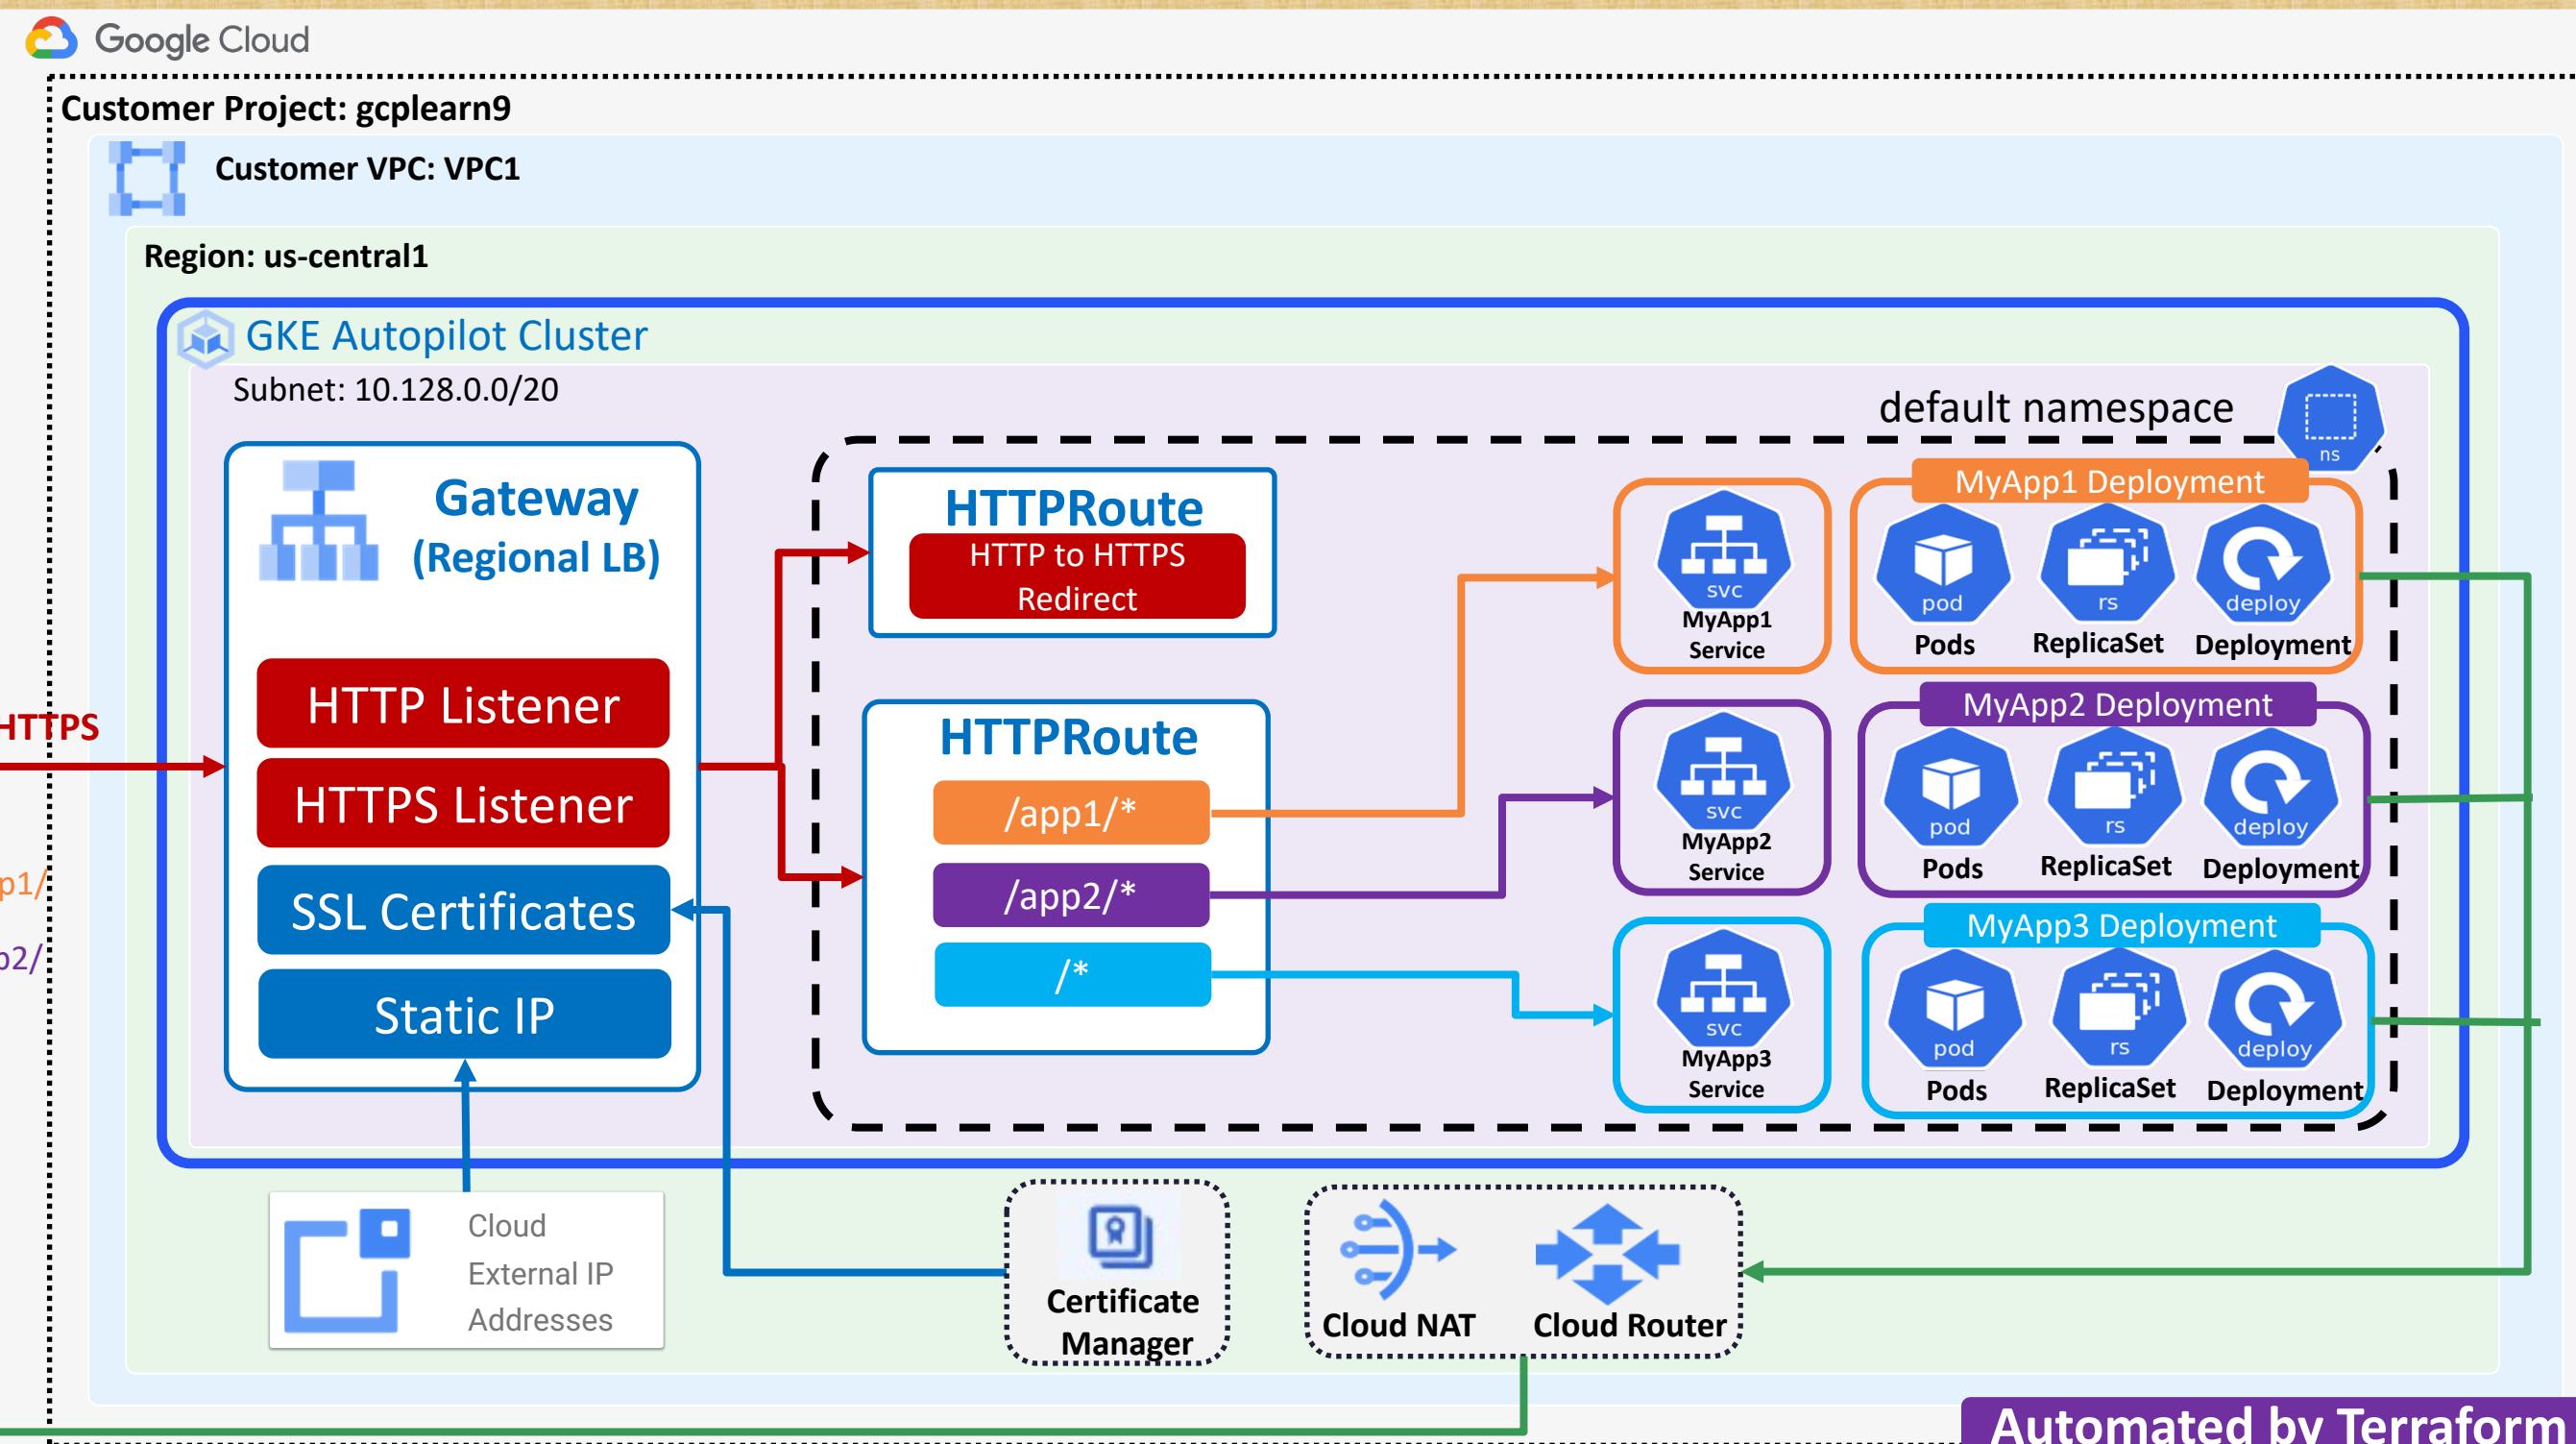
# Gateway API Basics + Static IP + Self-signed SSL with Certificate Manager



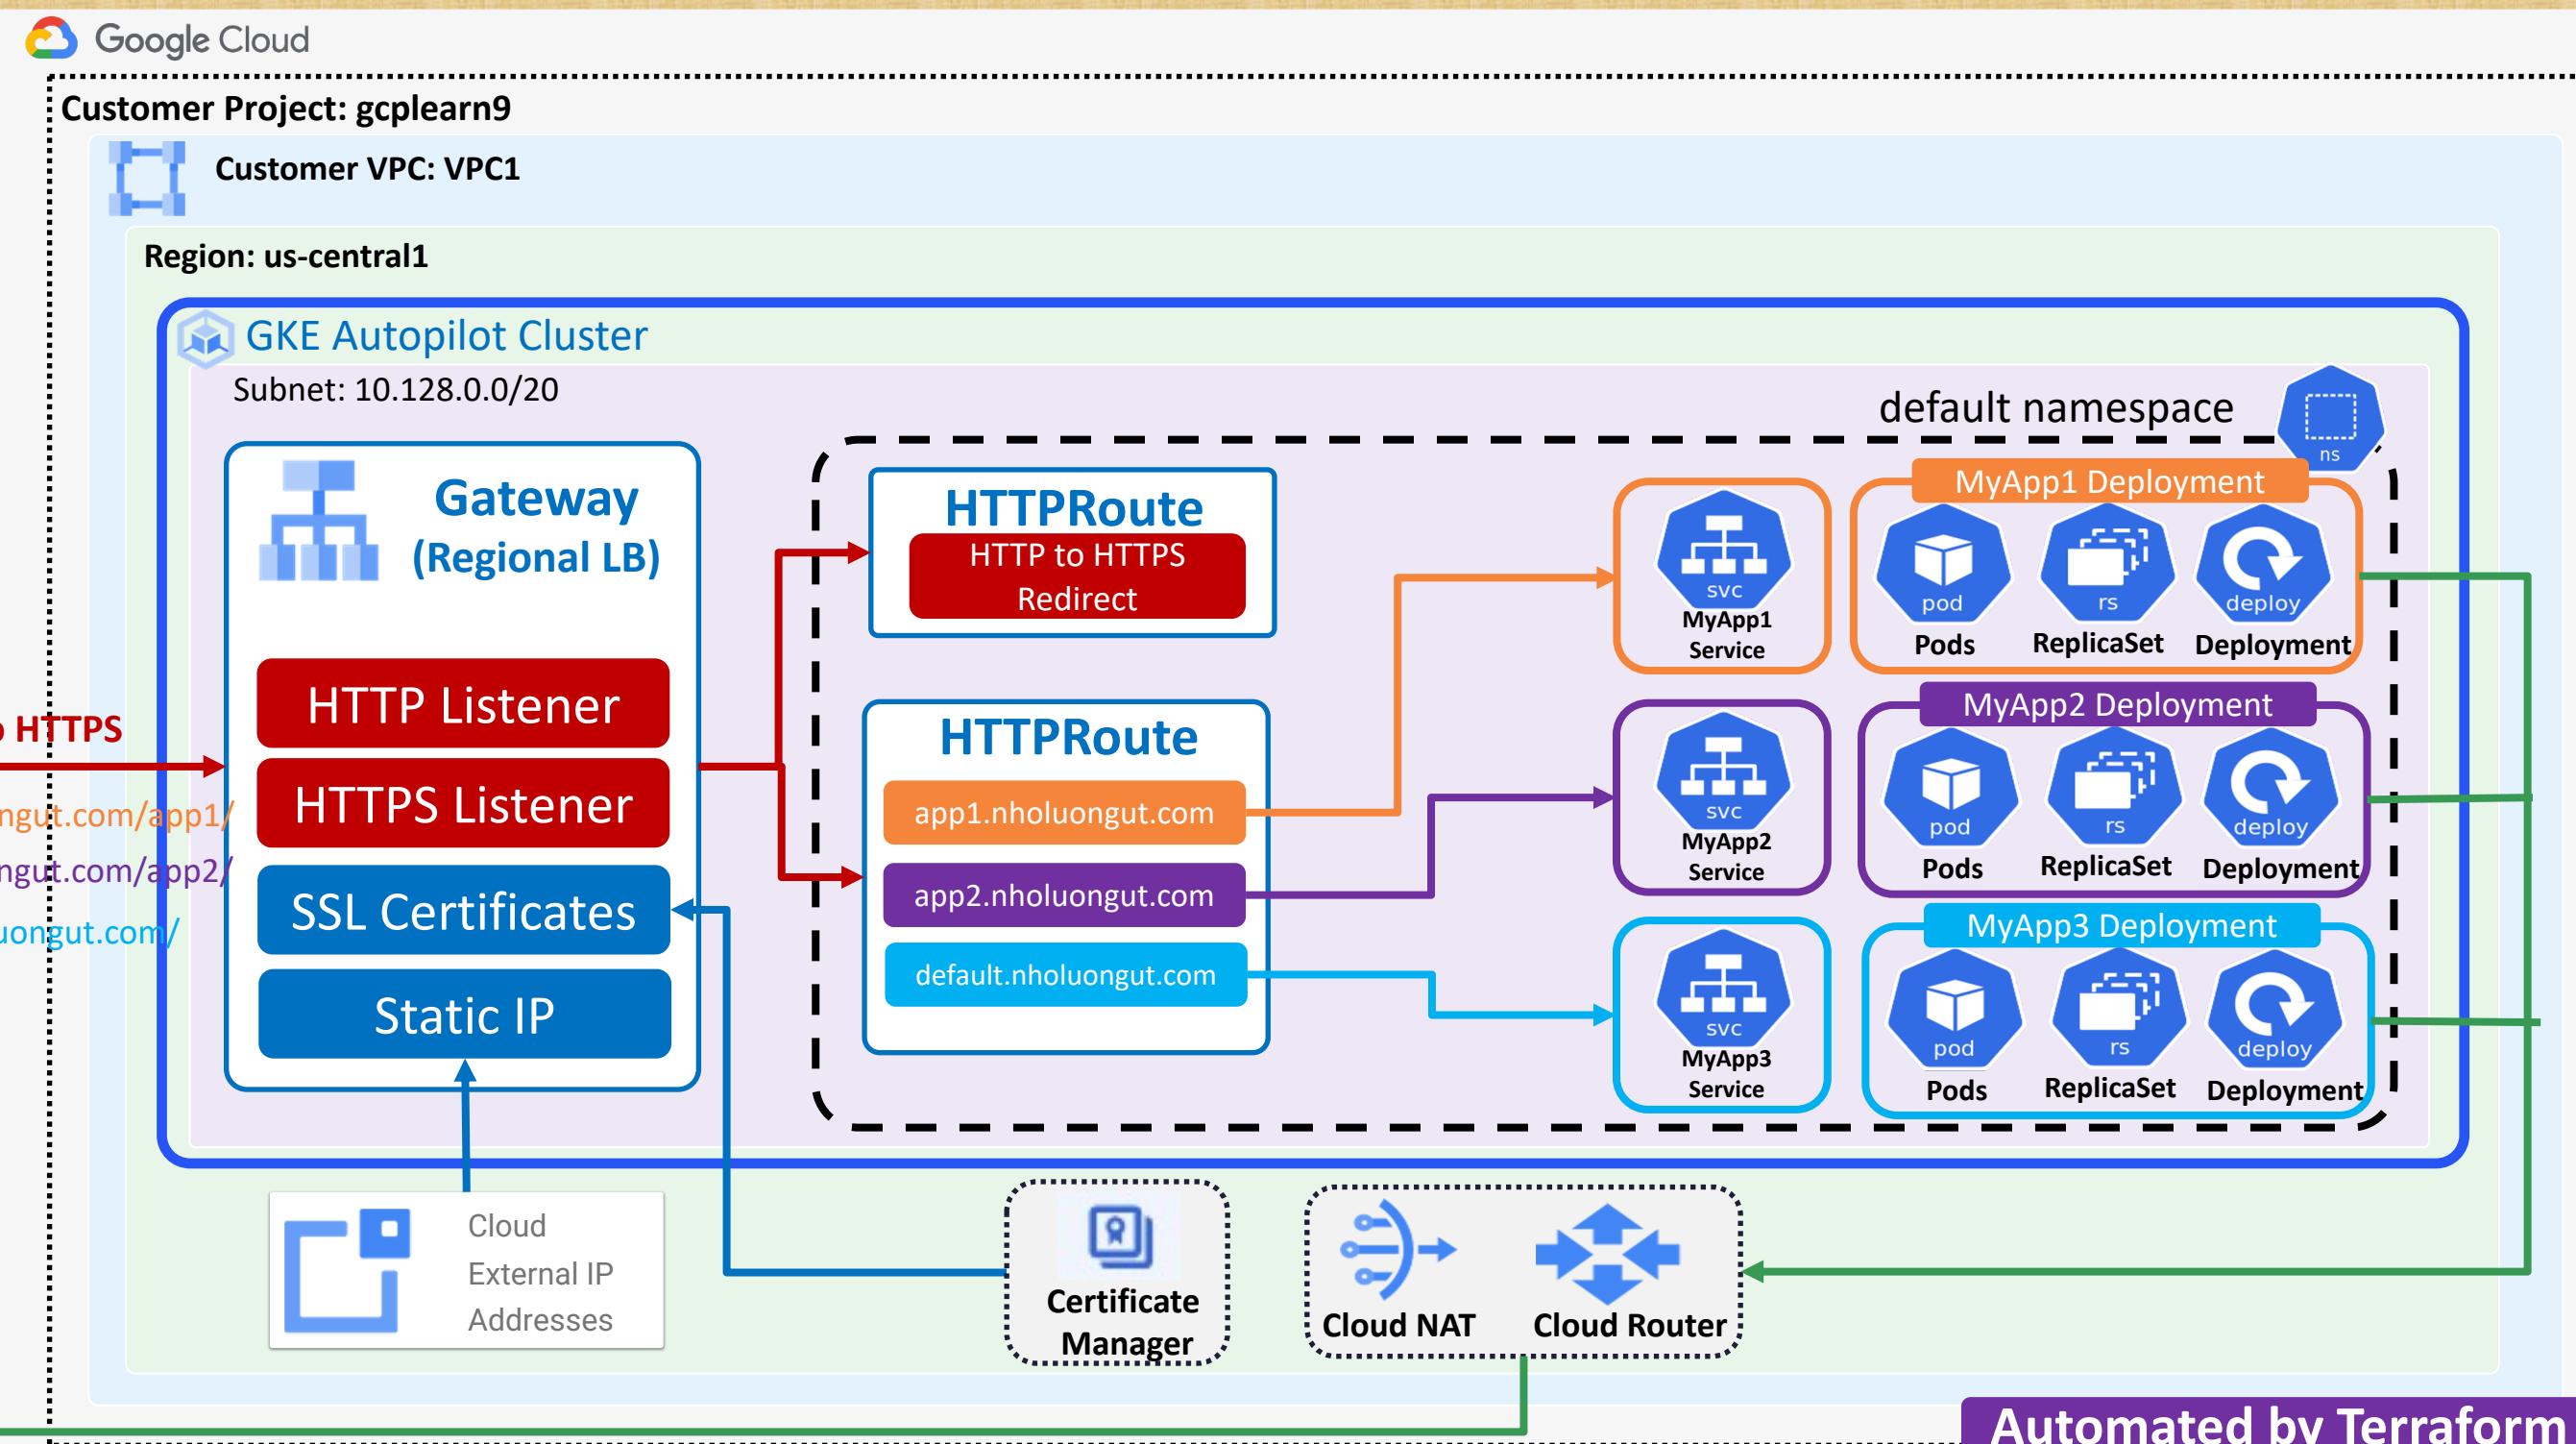
# Gateway API Basics + Static IP + SSL Certificate Manager + HTTP to HTTPS



# Gateway API Basics + Static IP + SSL Certificate Manager + HTTP to HTTPS + Context Path based Routing

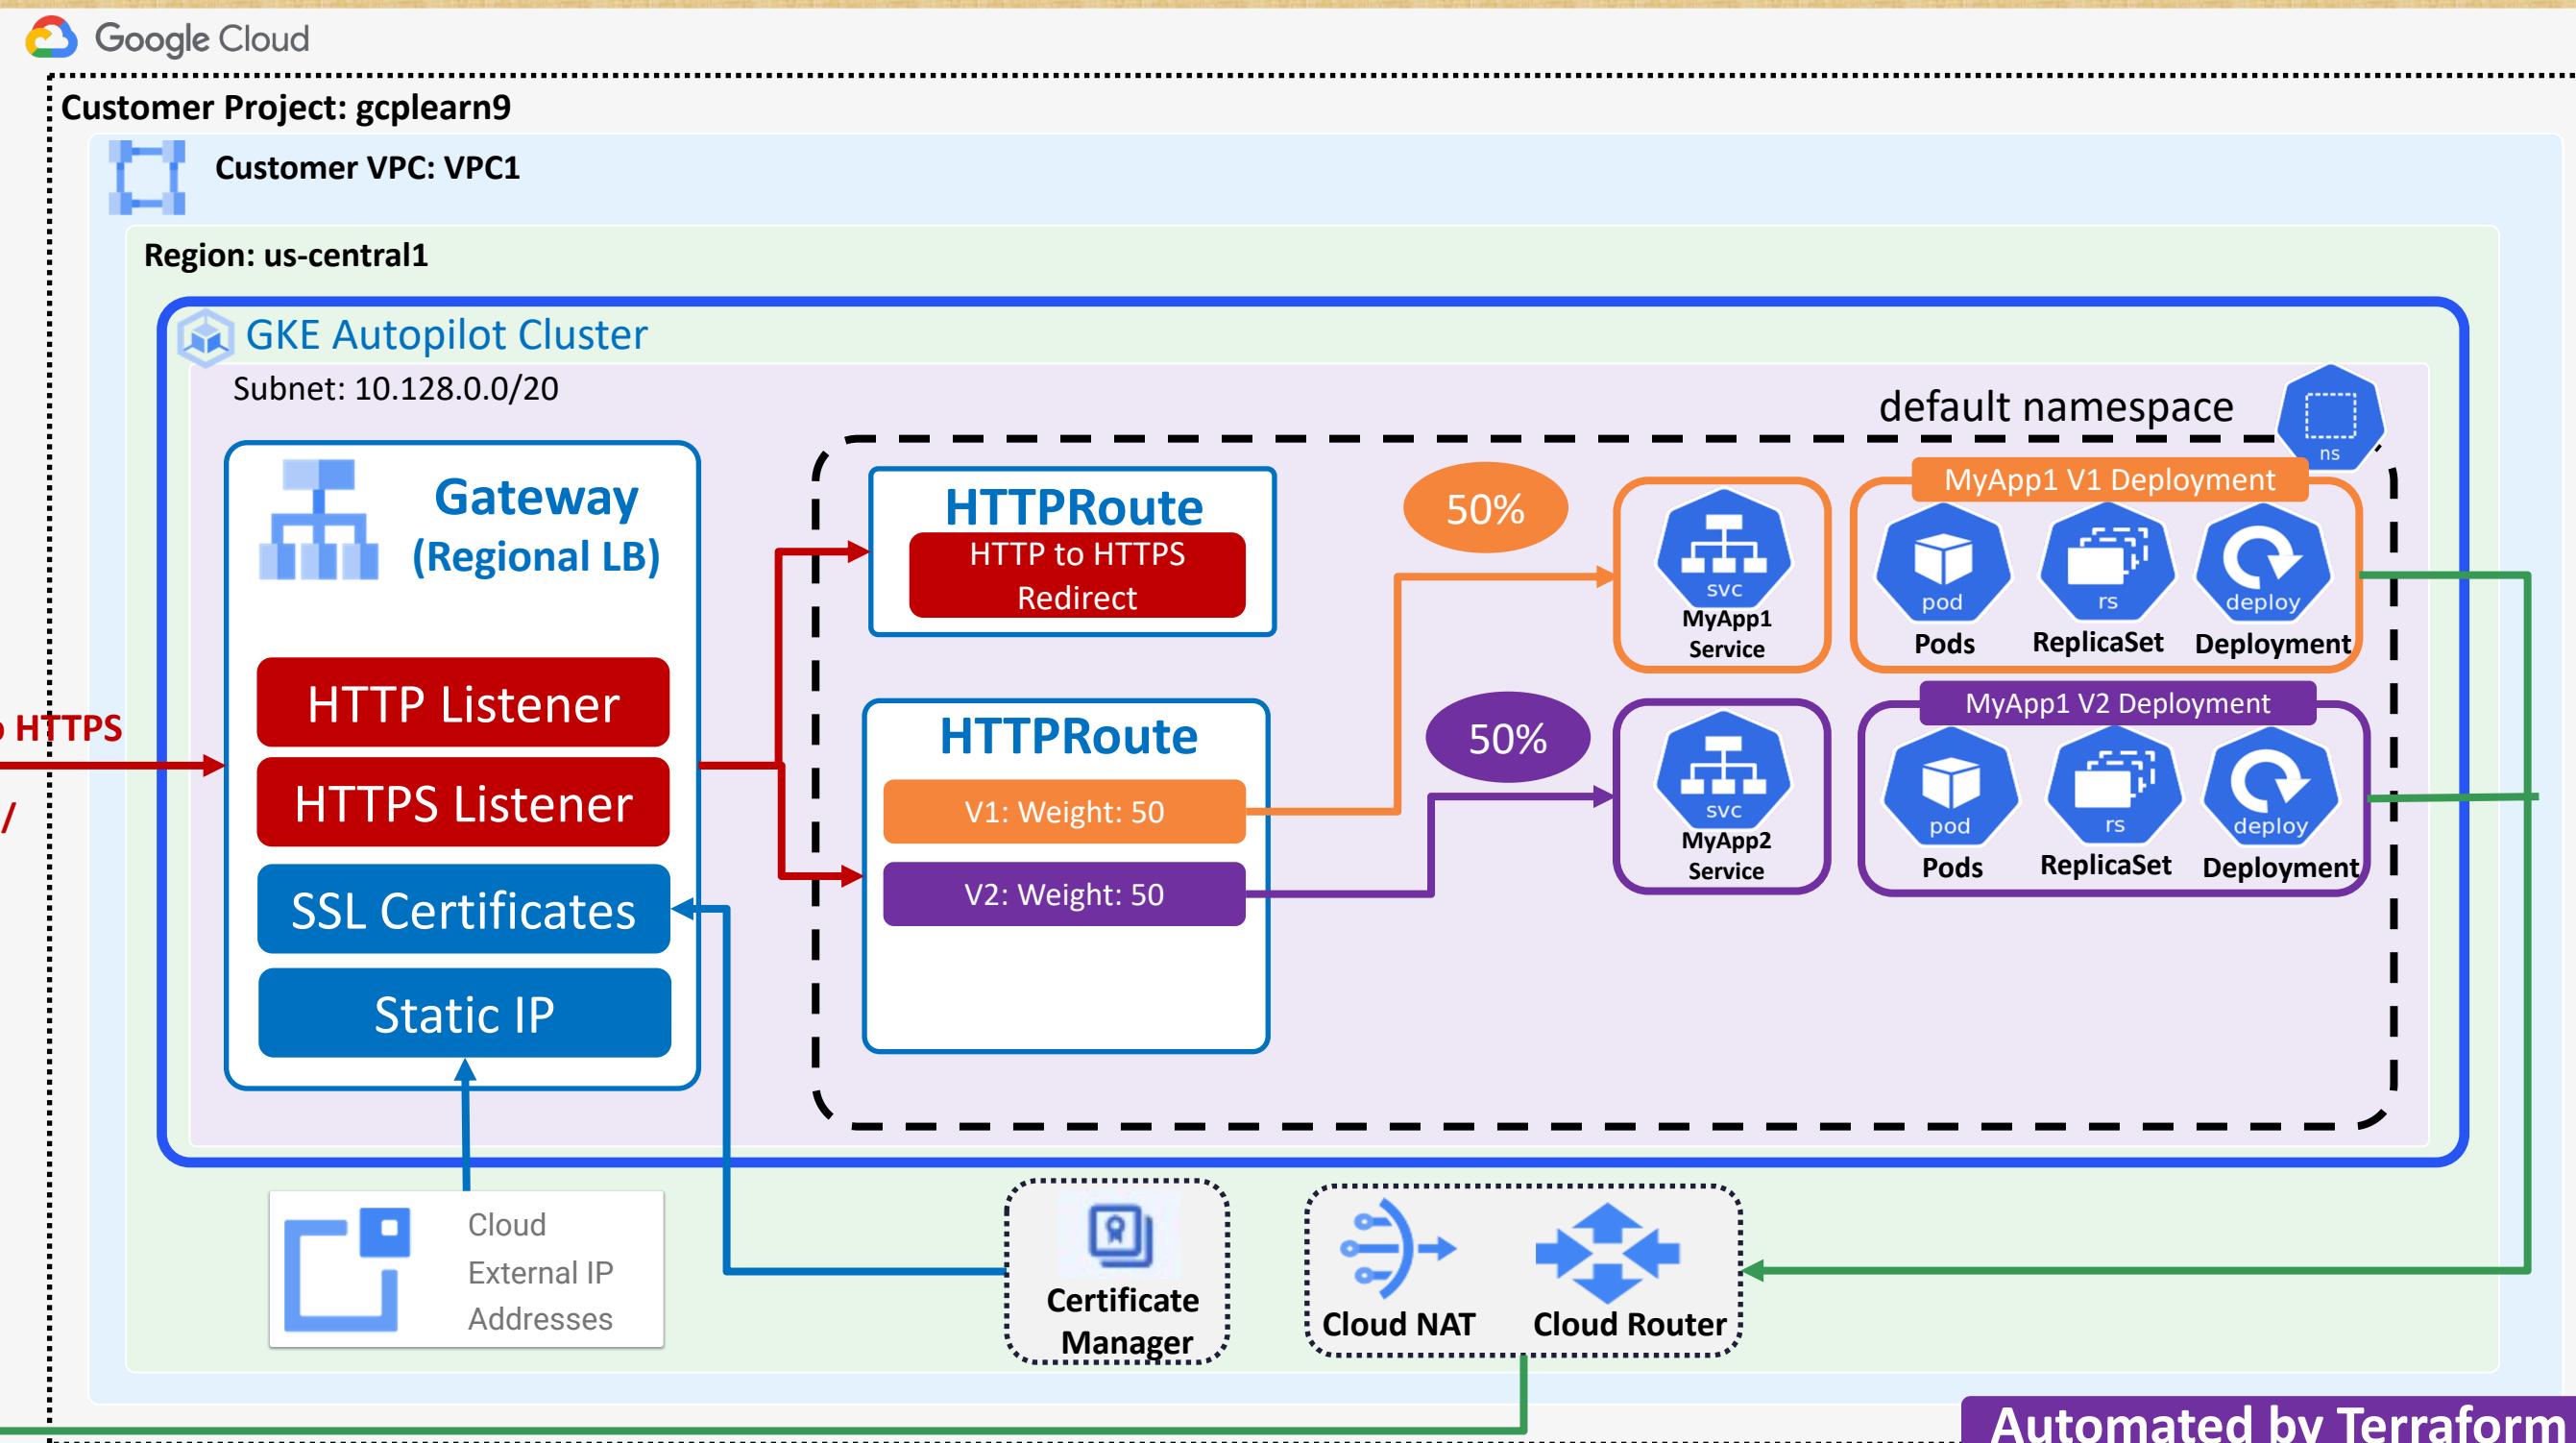


# Gateway API Basics + Static IP + SSL Certificate Manager + HTTP to HTTPS + Domain name-based Routing

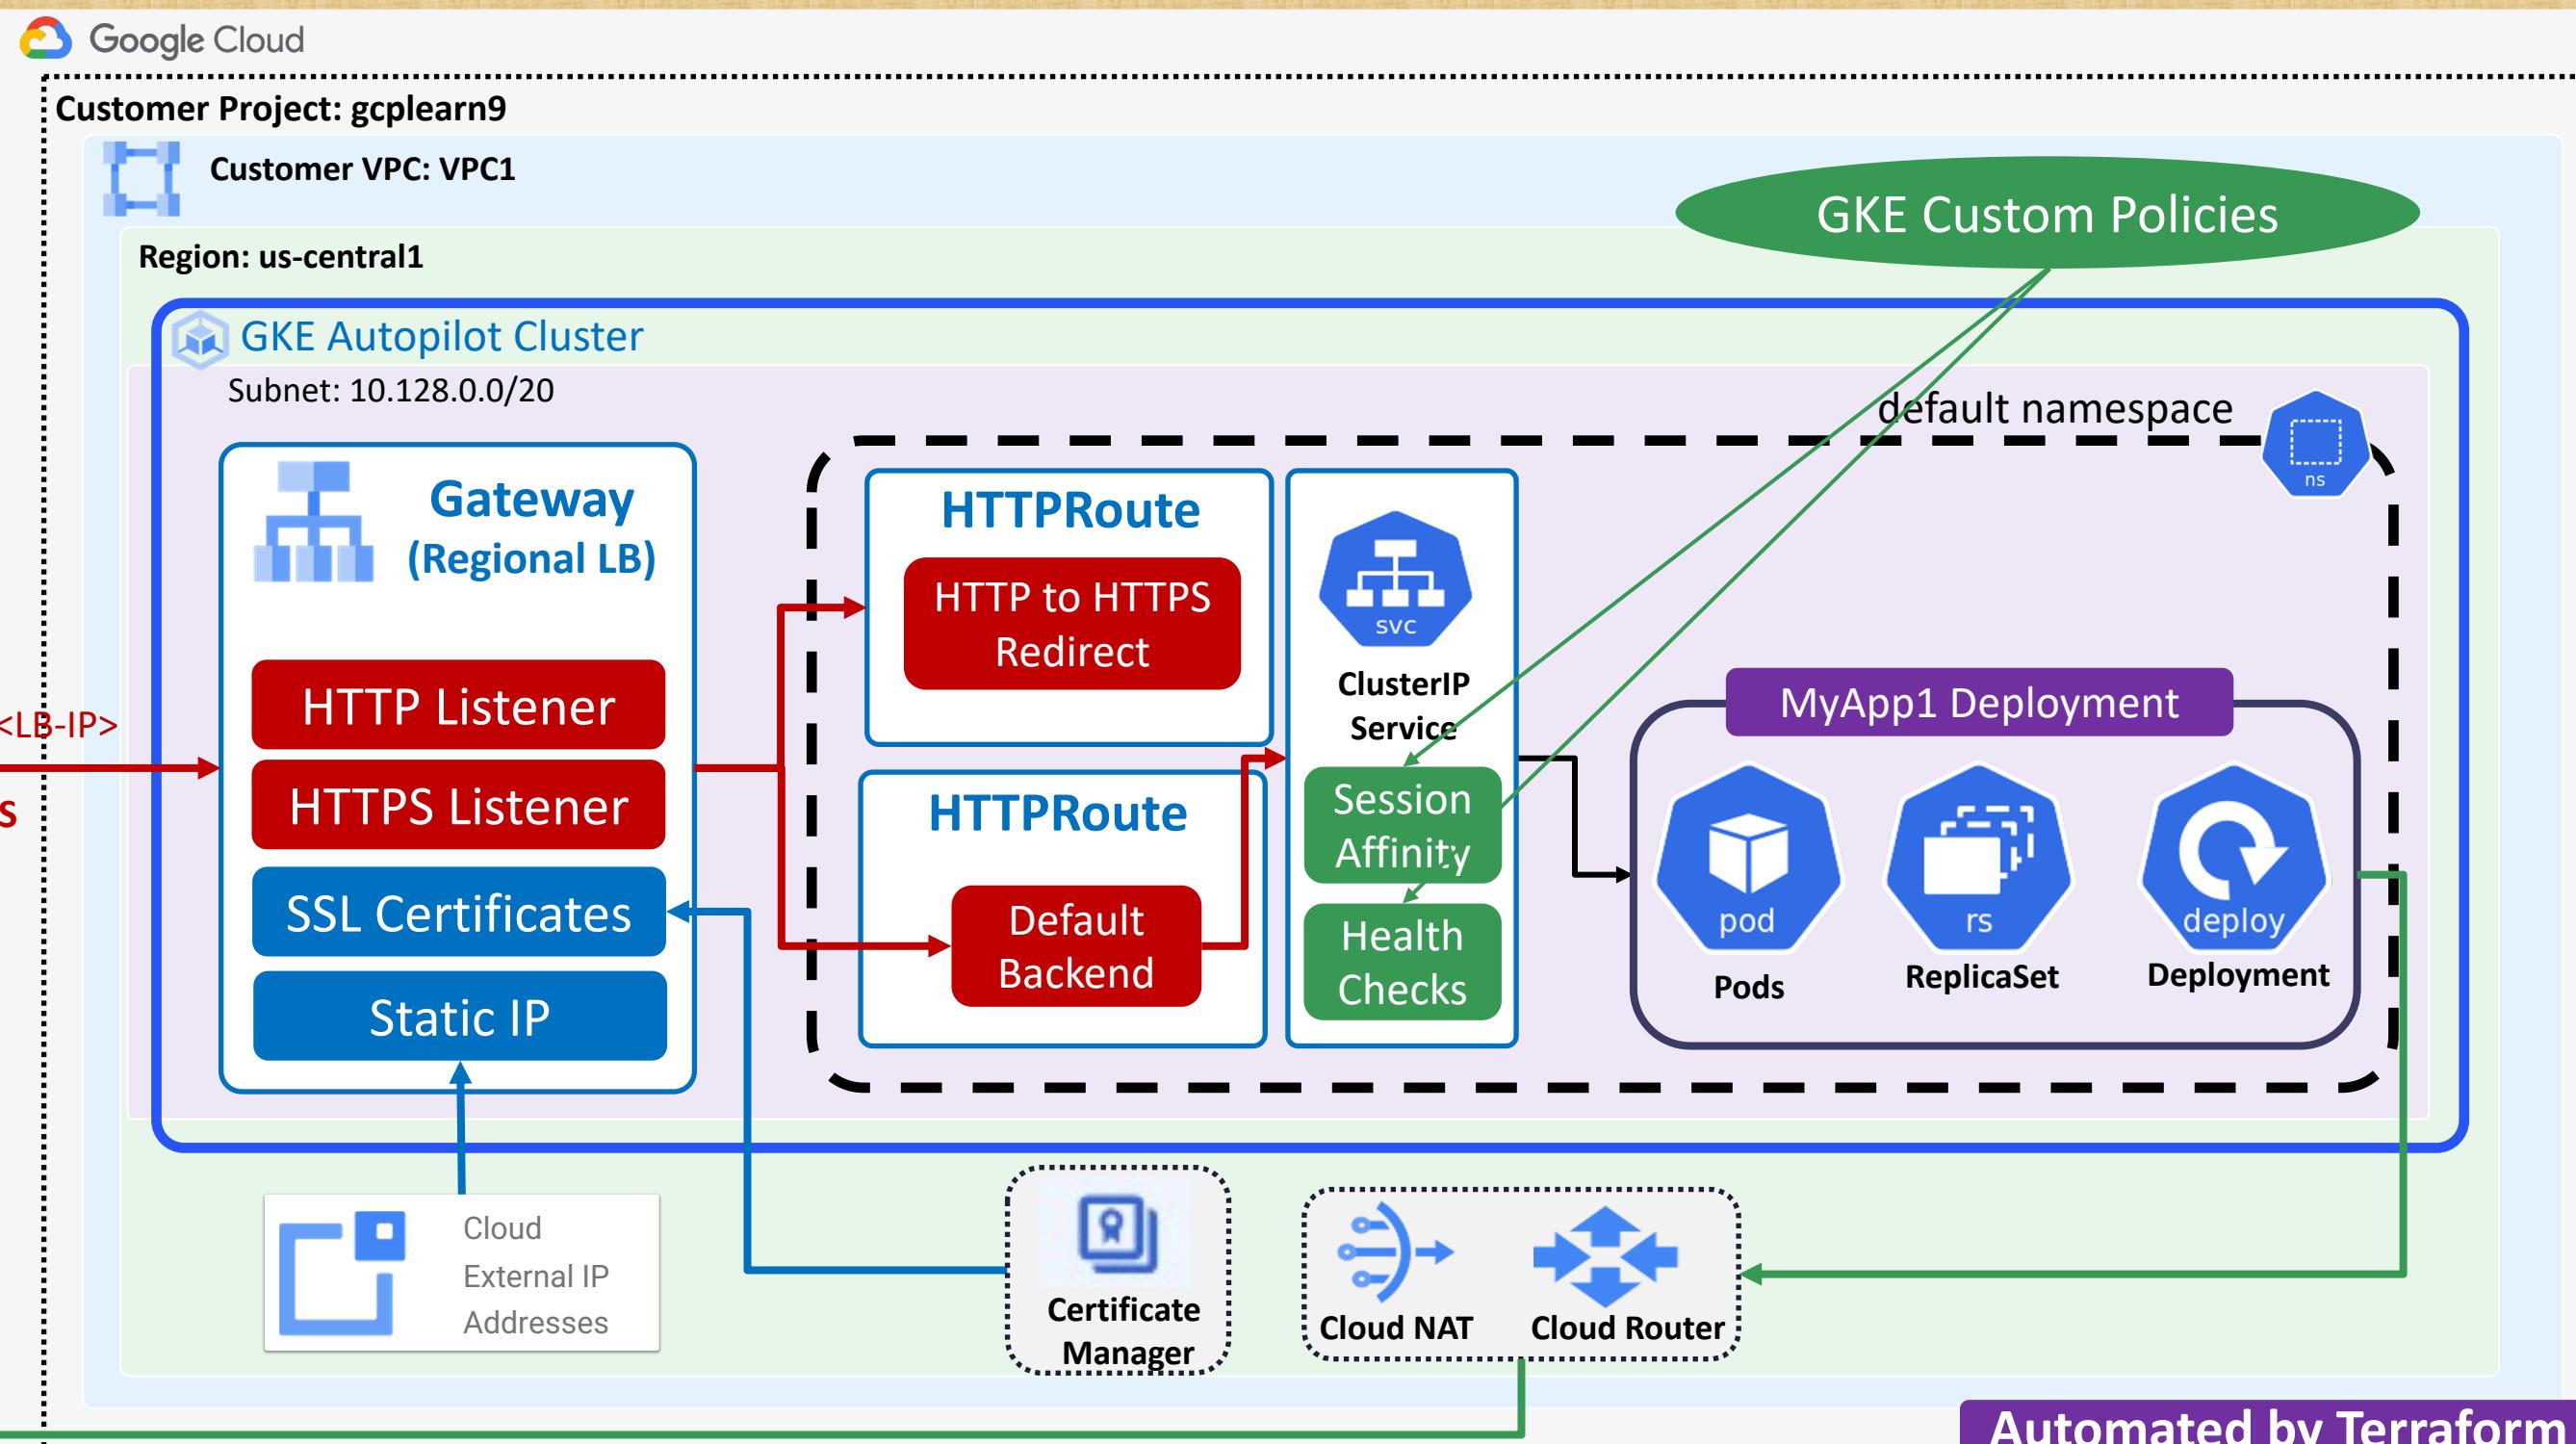


Automated by Terraform

# Gateway API Basics + Static IP + SSL Certificate Manager + HTTP to HTTPS + Traffic Splitting



# Gateway API + Static IP + SSL Certificate Manager + HTTP to HTTPS + Health Checks + Session Affinity



# Gateway API + Static IP + SSL Certificate Manager + HTTP to HTTPS + Health Checks + Session Affinity + Cloud Domains + Cloud DNS



Customer Project: gcplearn9

Automated by Terraform



Customer VPC: VPC1

Region: us-central1



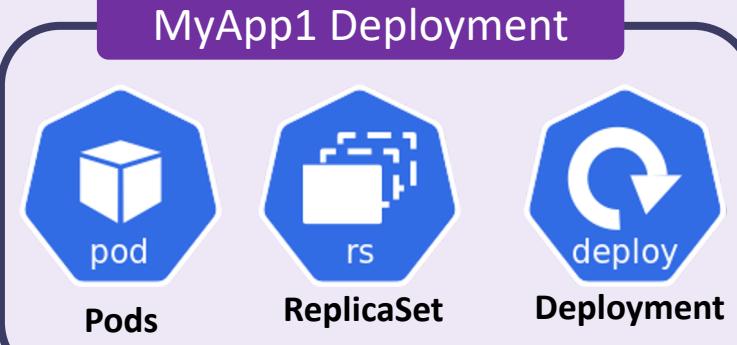
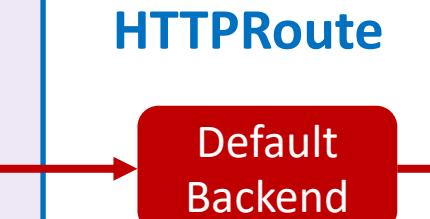
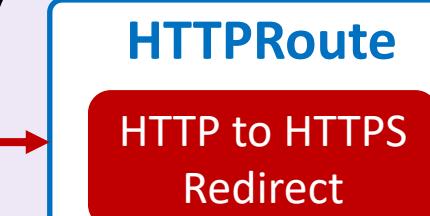
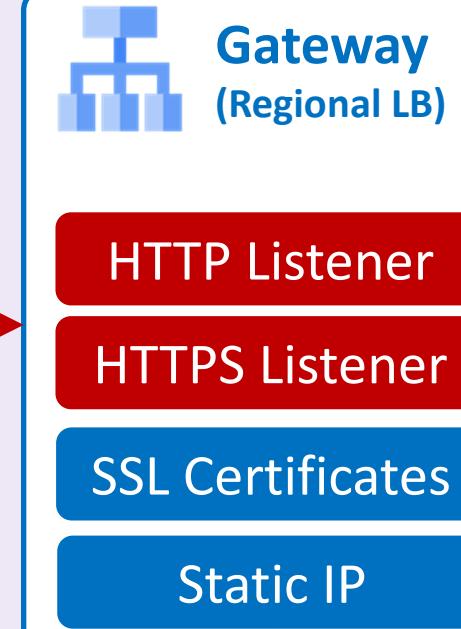
GKE Autopilot Cluster

Subnet: 10.128.0.0/20

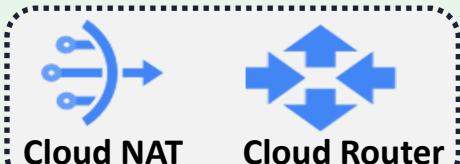
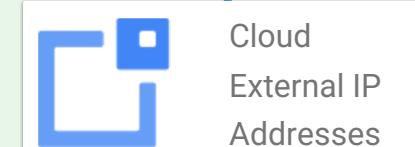
Users



http://<DNS-NAME>  
HTTP to HTTPS



default namespace





Customer Project: gcplearn9

Automated by Terraform



Customer VPC: VPC1

Region: us-central1



GKE Autopilot Cluster

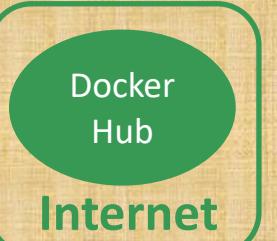
Subnet: 10.128.0.0/20

Users



http://&lt;DNS-NAME&gt;

AWS Route53



Gateway (Regional LB)

HTTP Listener

HTTPS Listener

SSL Certificates

Static IP

HTTPRoute

HTTP to HTTPS Redirect

HTTPRoute

Default Backend



ClusterIP Service

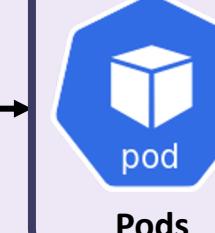
Session Affinity

Health Checks

default namespace



MyApp1 Deployment



Pods



ReplicaSet



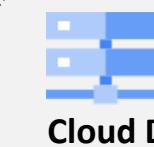
Deployment



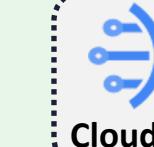
Cloud External IP Addresses



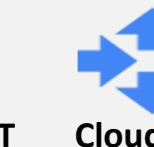
Production grade SSL certificates



Cloud DNS

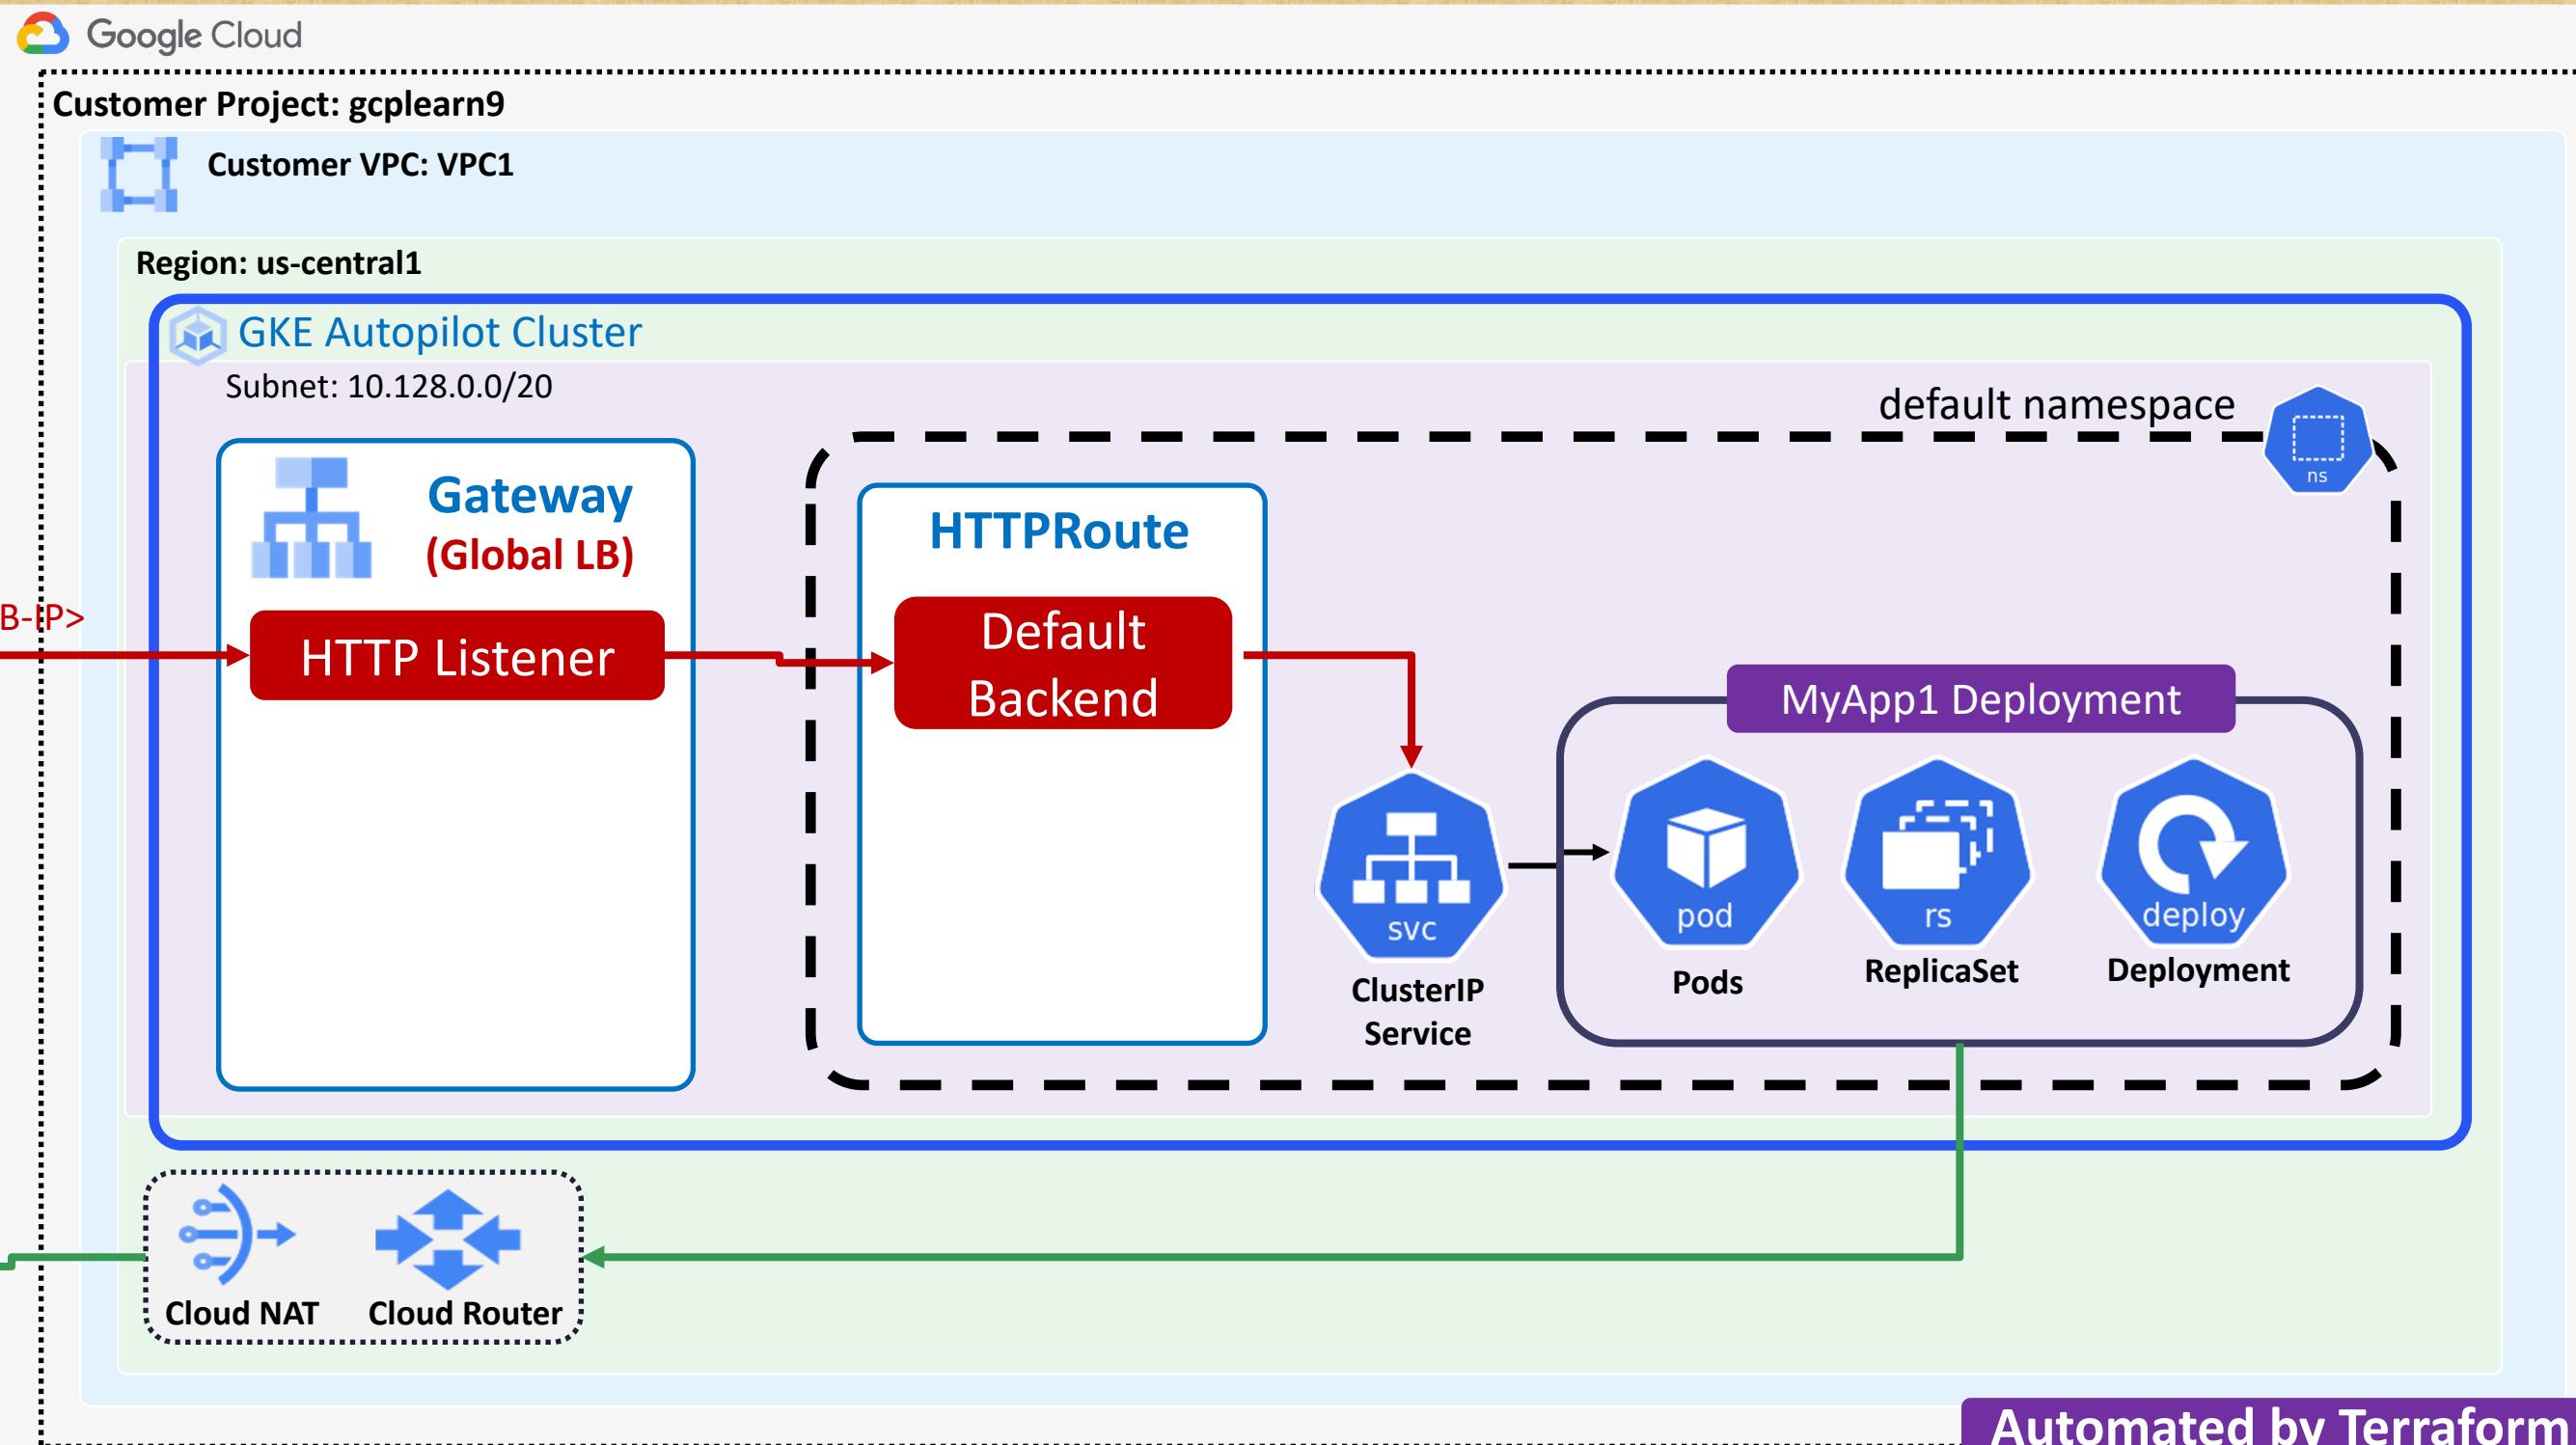


Cloud NAT



Cloud Router

# Gateway API – Global Application Load Balancer



Automated by Terraform

# Terraform Registry - Use Publicly Available Modules

The Terraform Registry hosts a broad collection of publicly available Terraform modules for configuring many kinds of common infrastructure.

## Modules Demo 1

These modules are free to use, and Terraform can download them automatically if you specify the appropriate source and version in a module call block.

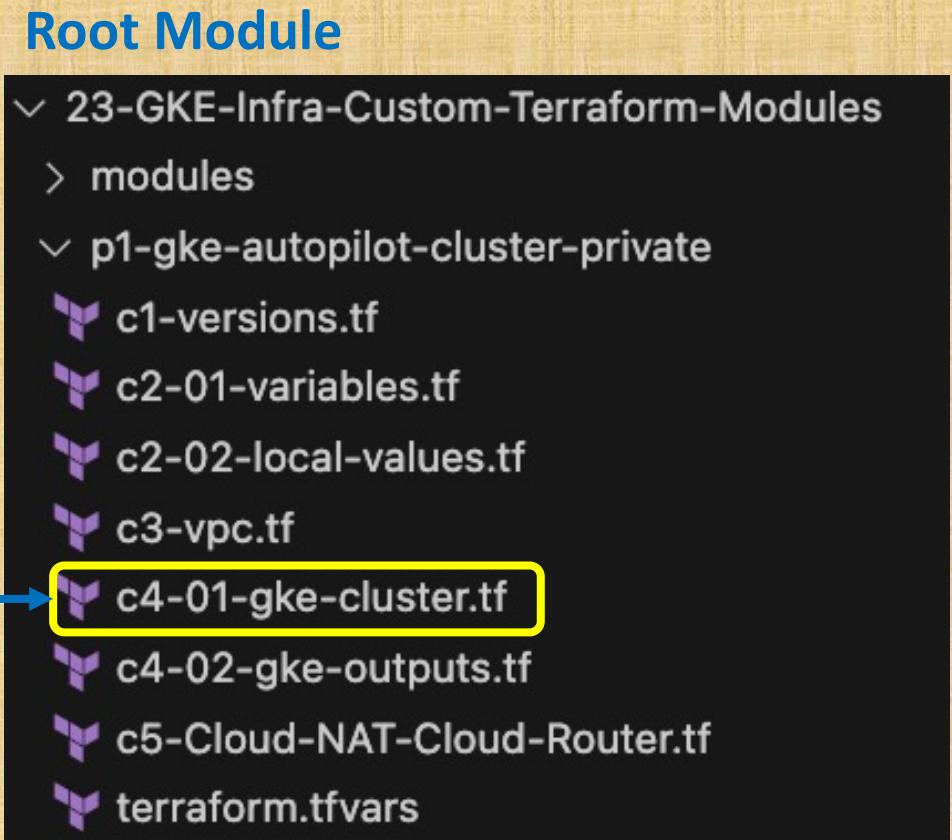
```
# Module: VPC
module "vpc" {
    source  = "terraform-google-modules/network/google"
    version = "~> 9.1"

    project_id    = var.gcp_project
    network_name = "${local.name}-vpc"
    routing_mode = "GLOBAL"
    subnets = [
        {
            subnet_name      = "${local.name}-${var.gcp_region1}-subnet"
            subnet_ip       = "10.128.0.0/20"
            subnet_region   = var.gcp_region1
        }
    ]
}
```

# Build a Local Terraform Module



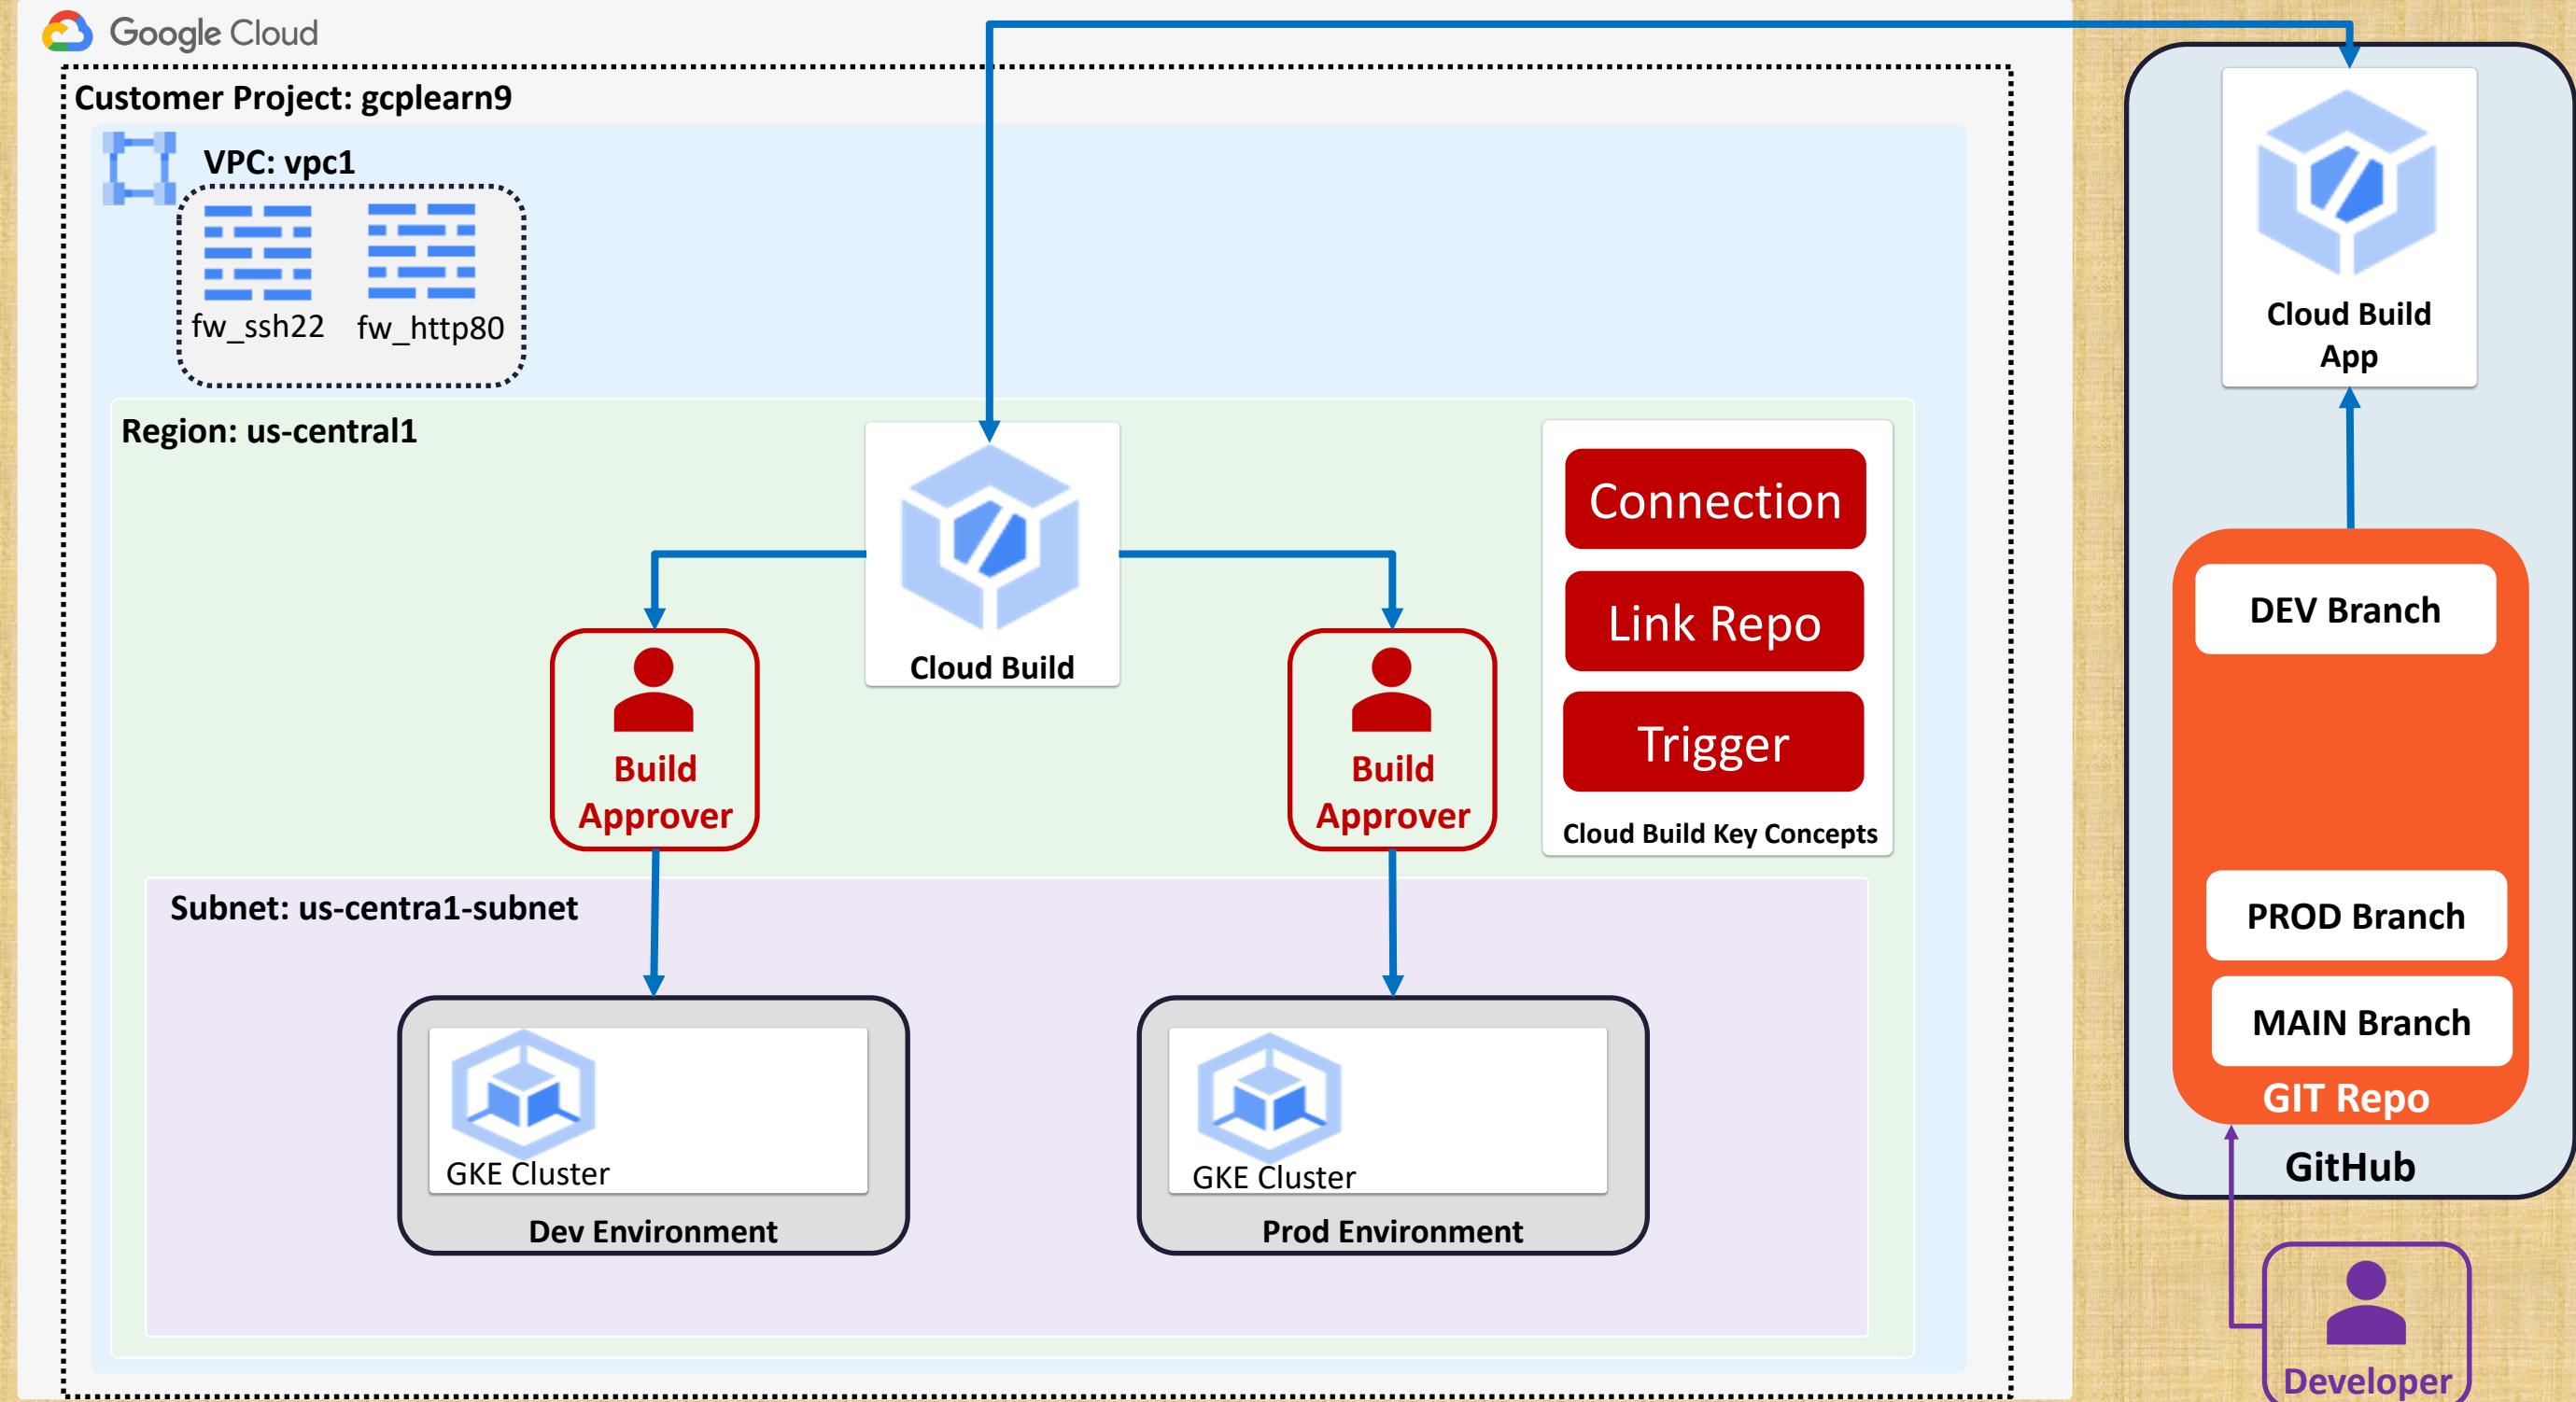
Call GKE Cluster child module in Root Module



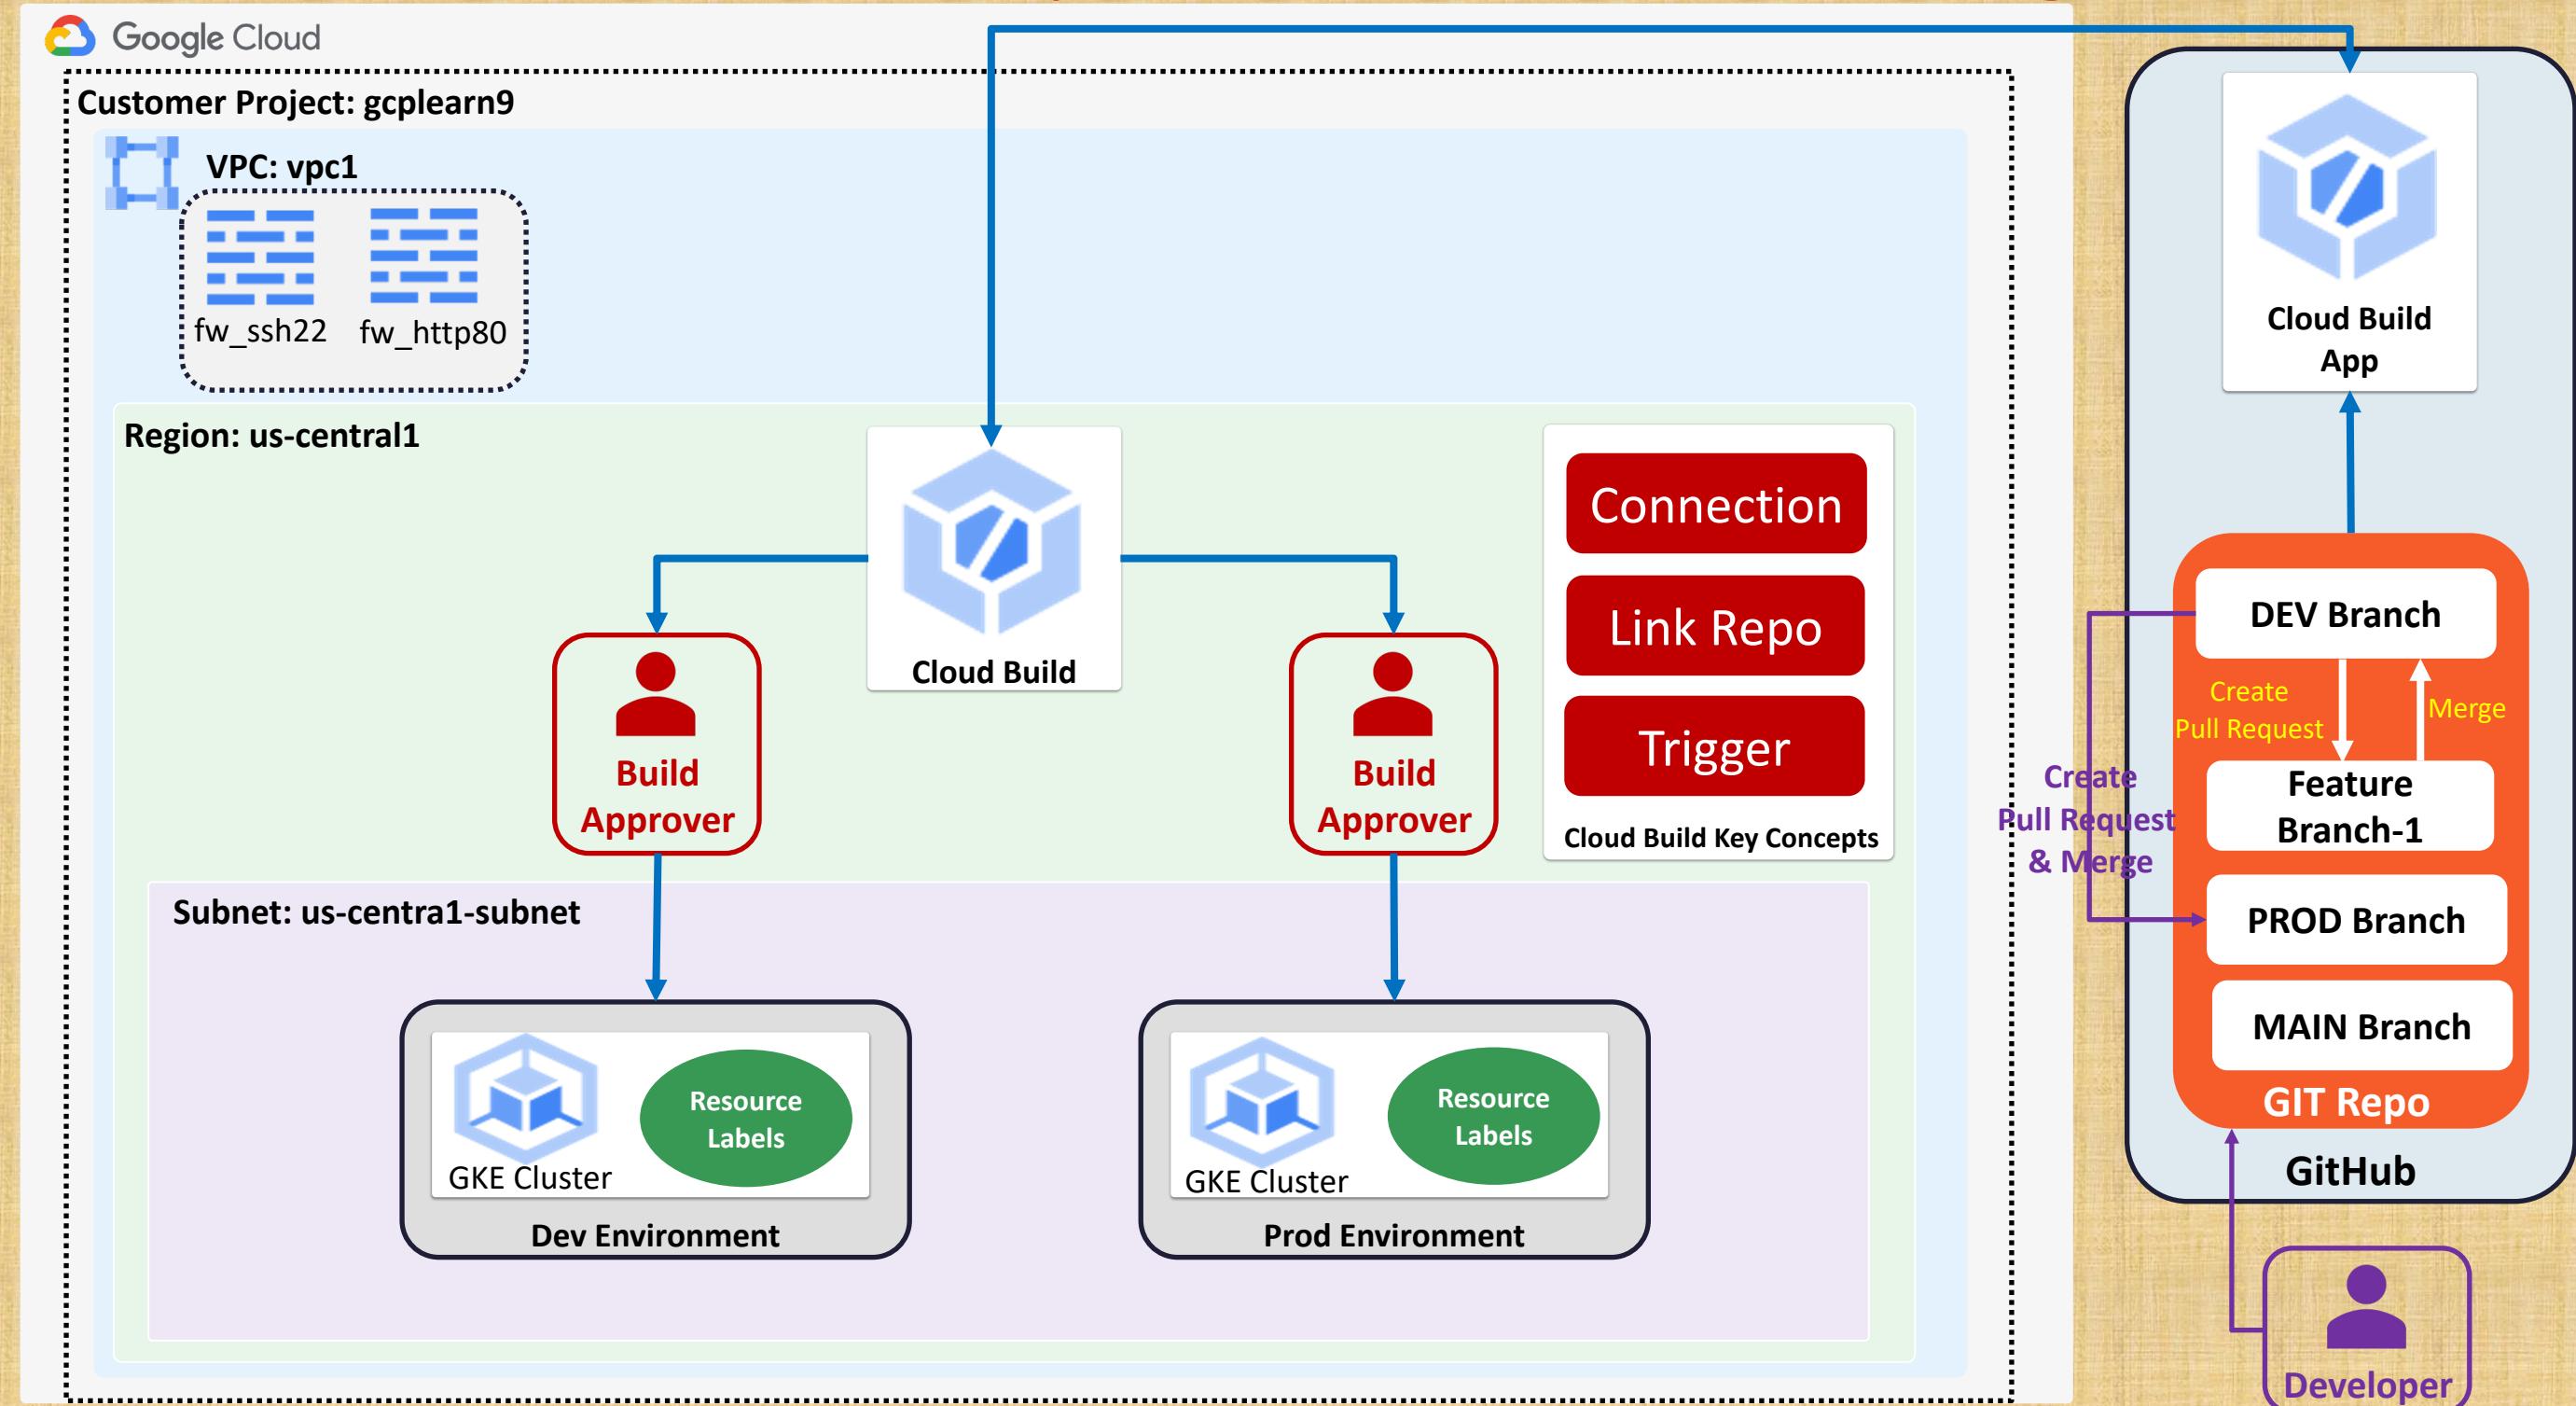
## Child Module



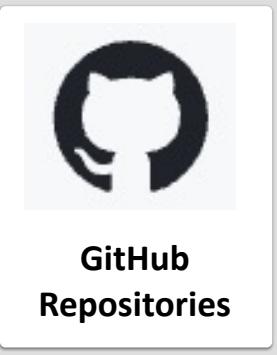
# GCP GKE INFRA DevOps for Terraform Configs



# GCP GKE INFRA DevOps for Terraform Configs



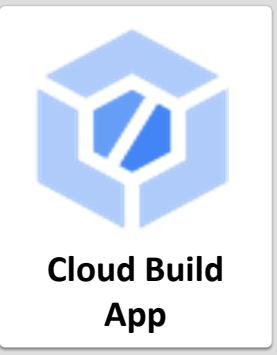
# Continuous Integration and Continuous Delivery



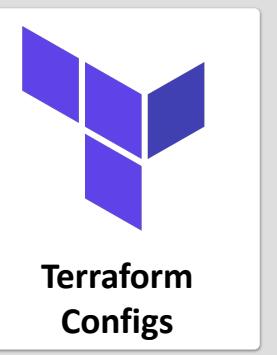
GitHub  
Repositories



Artifact  
Registry



Cloud Build  
App



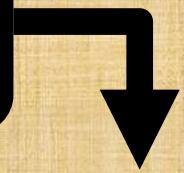
Terraform  
Configs



GKE  
Cluster

Demo  
25

Create Kubernetes Deployment Custom  
Terraform Module



Demo  
26

Implement k8s workload DevOps  
pipelines  
(Dev and Prod Environments)



Demo  
27

App: Implement Continuous Integration  
Pipeline with Artifact Registry



Demo  
28

App: Implement Continuous Integration  
and Continuous Delivery Pipeline (CI CD)

## Root Module

```
✓ 25-GKE-Workloads-Custom-Terraform-Modules
  > modules
    ✓ p2-k8sresources-terraform-manifests
      ✓ c1-versions.tf
      ✓ c2-01-variables.tf
      ✓ c2-02-local-values.tf
      ✓ c3-01-remote-state-datasource.tf
      ✓ c3-02-providers.tf
      ✓ c4-kubernetes-deployment.tf
      ✓ c5-kubernetes-loadbalancer-service.tf
    ✓ terraform.tfvars
```

## Child Module

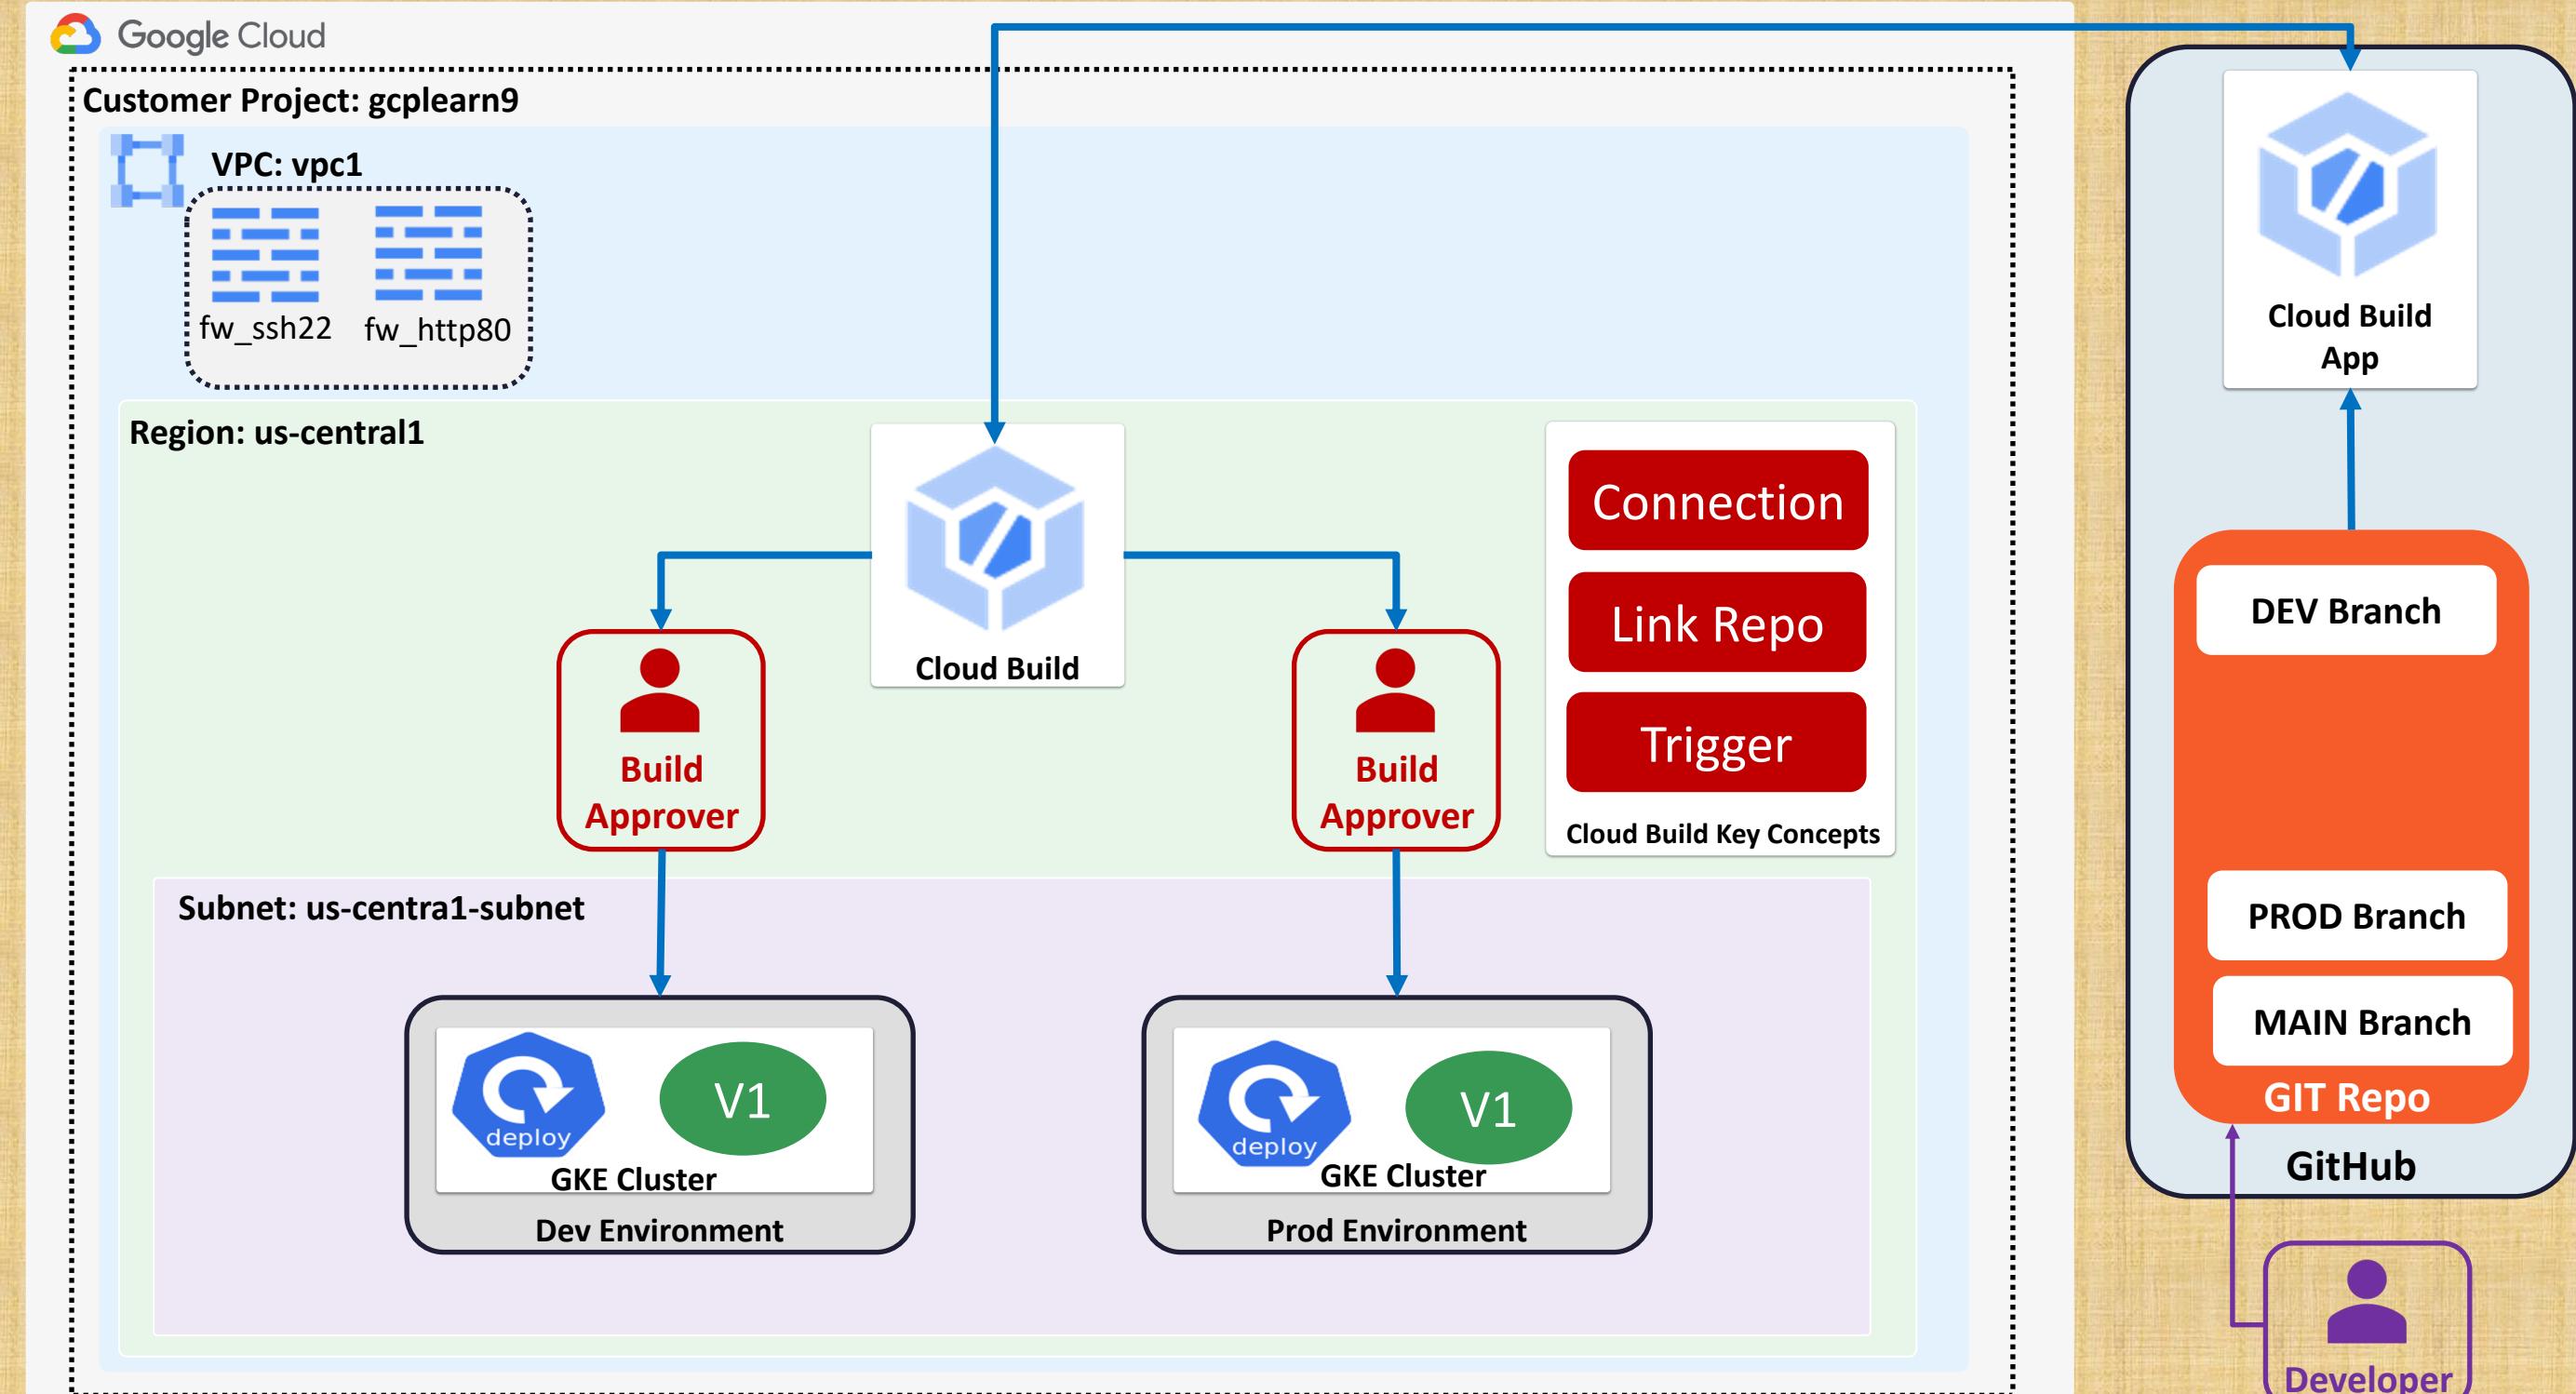
```
✓ 25-GKE-Workloads-Custom-Terraform-Modules
  ✓ modules
    > gke_cluster
      ✓ kubernetes_deployment
        ✓ main.tf
        ✓ outputs.tf
        ✓ variables.tf
        ✓ versions.tf
```

# Build a Local Terraform Module

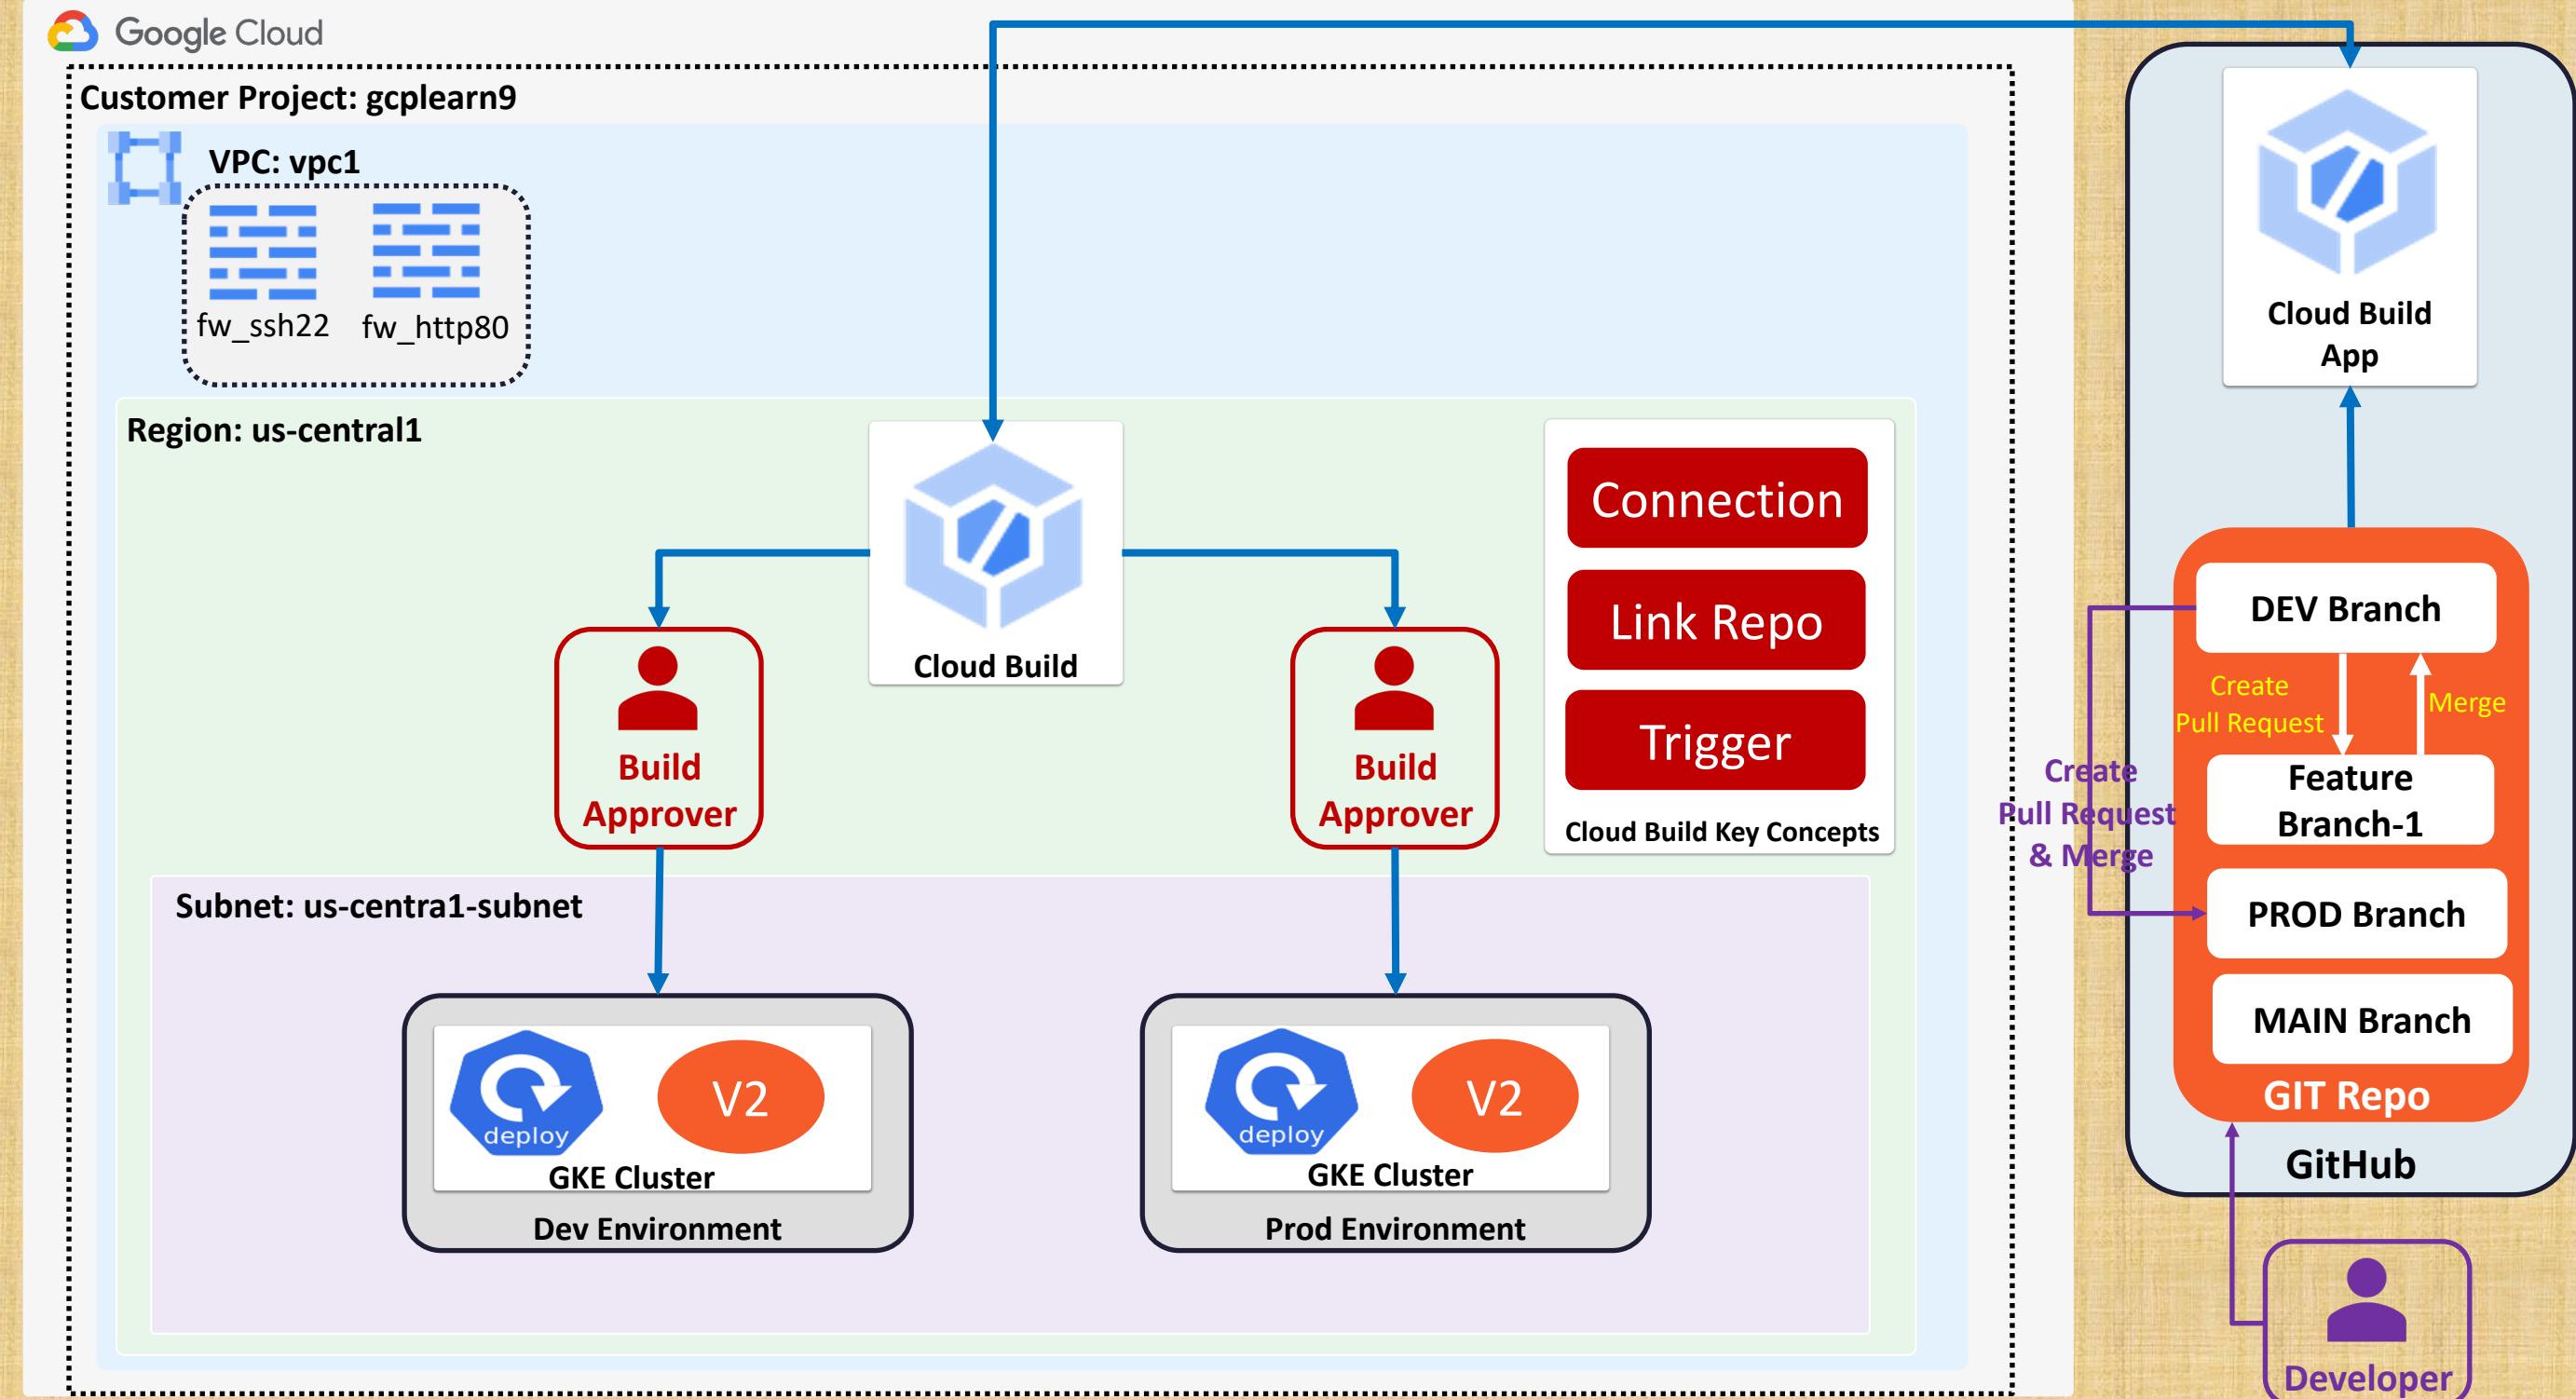
Kubernetes  
Deployment  
Custom  
Terraform Module

Call GKE Cluster child  
module in Root Module

# Kubernetes Deployment DevOps for Terraform Configs



# Kubernetes Deployment DevOps for Terraform Configs



# Continuous Integration

## Continuous Integration Pipeline



GitHub



Google  
Cloud Build



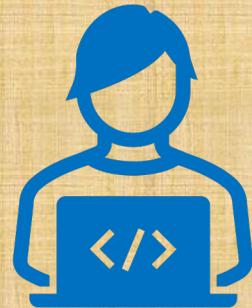
Google  
Artifact Registry

To Store App Code  
(Application Code  
+ Docker Files)

To Build Packages  
(Build Docker  
Image)

To Store Packages  
(Store Docker  
Image)

# Google Cloud Build: Continuous Integration



Admin

Check-in Code



GitHub

Google Cloud

Customer Project: gcplearn9

Region: us-central1

**Trigger the Code Build CI Pipeline**



Google  
Cloud Build

**Build and push new Docker Image**



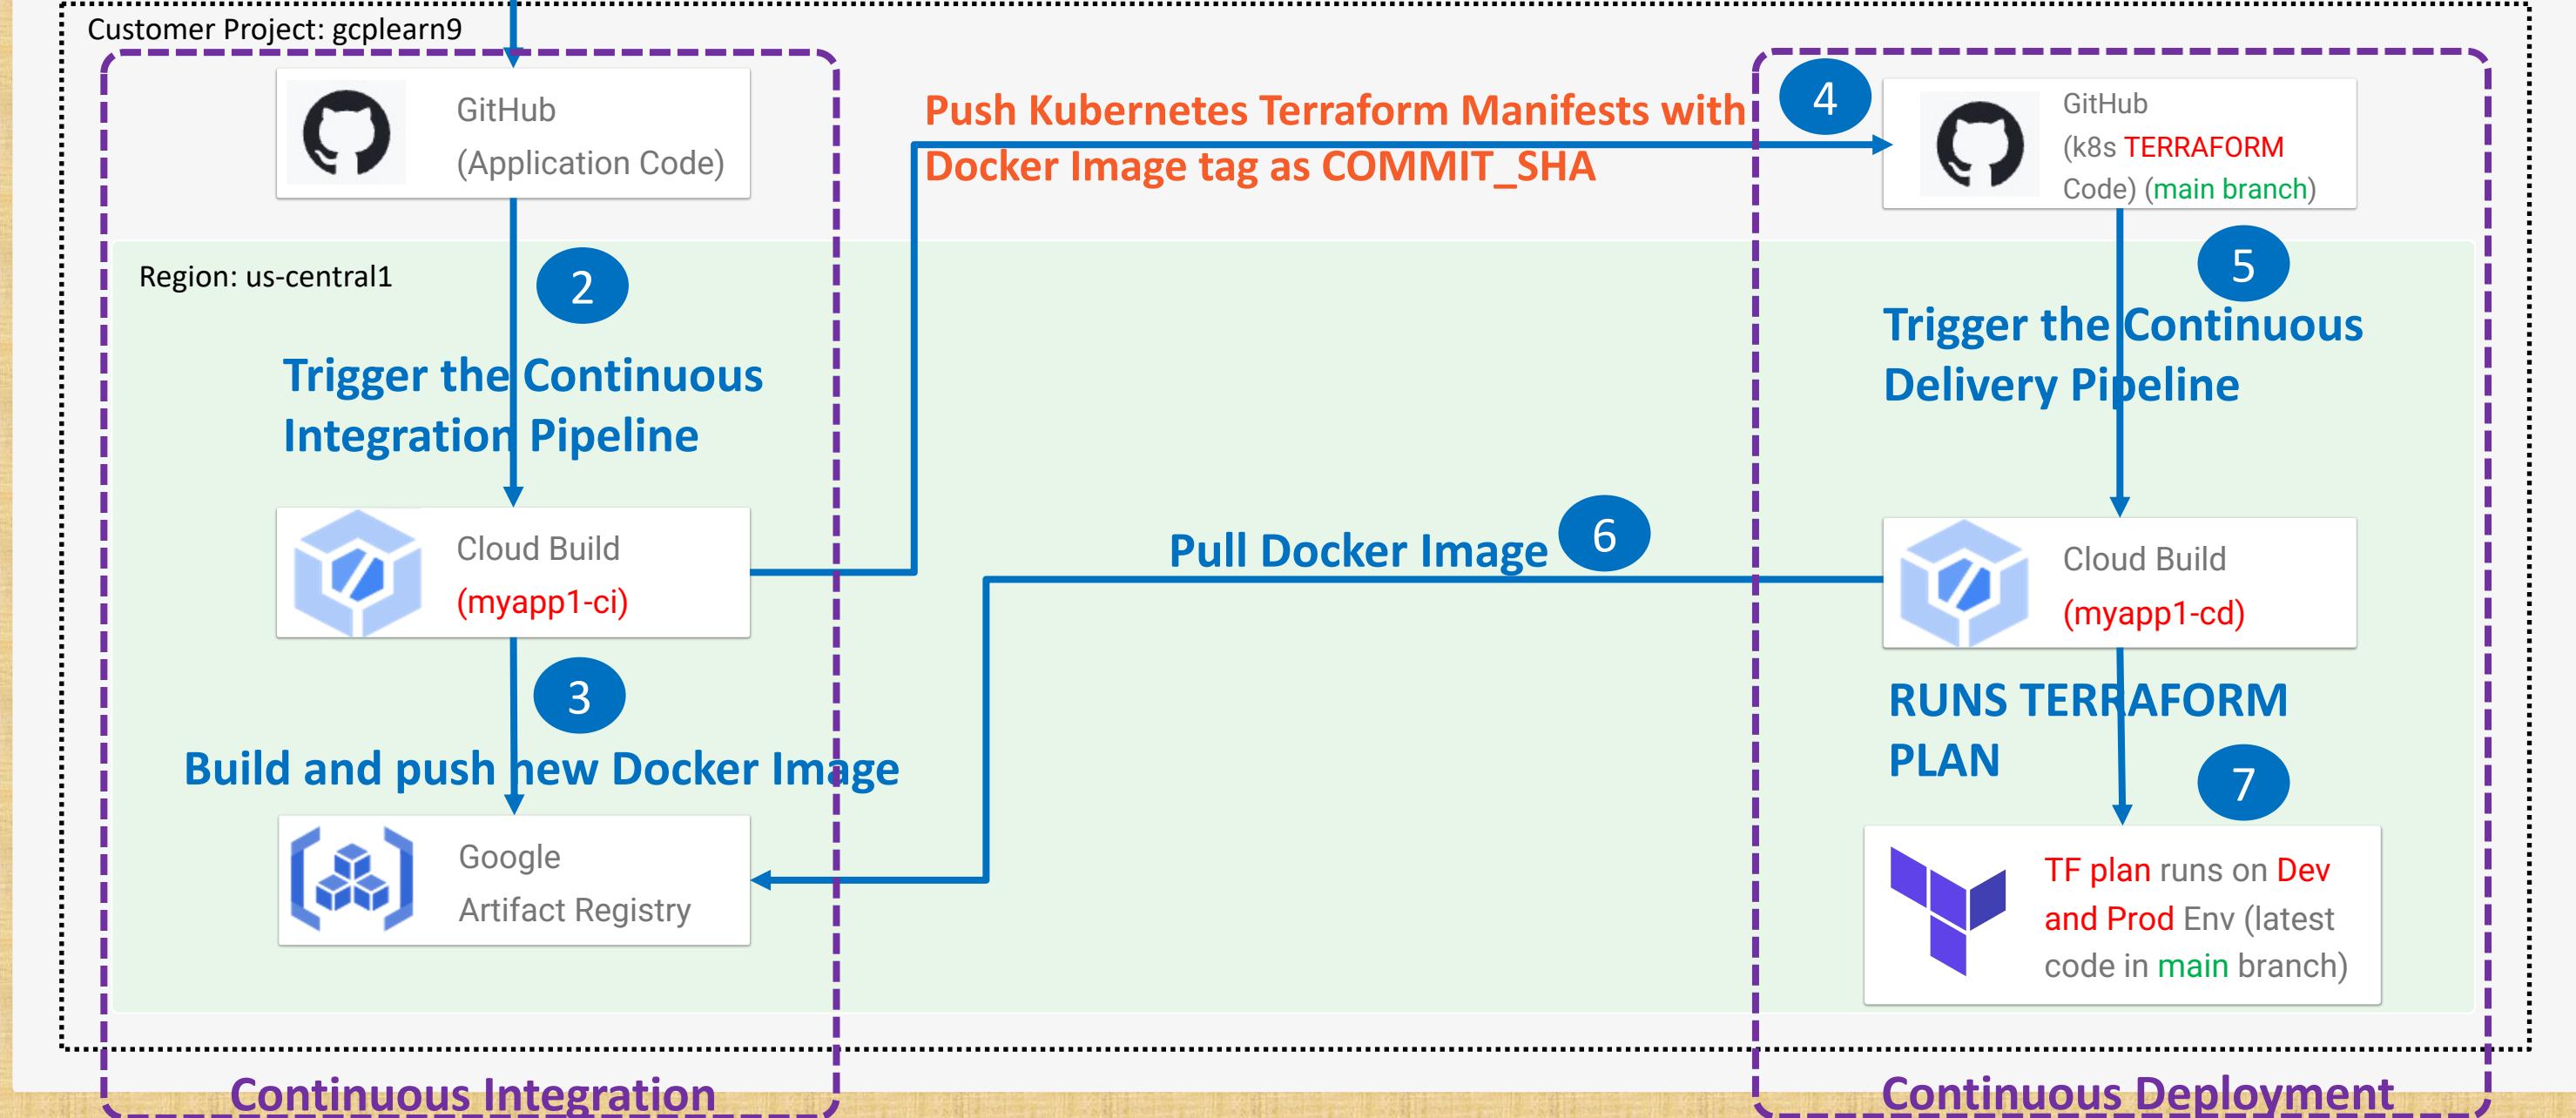
Google  
Artifact Registry



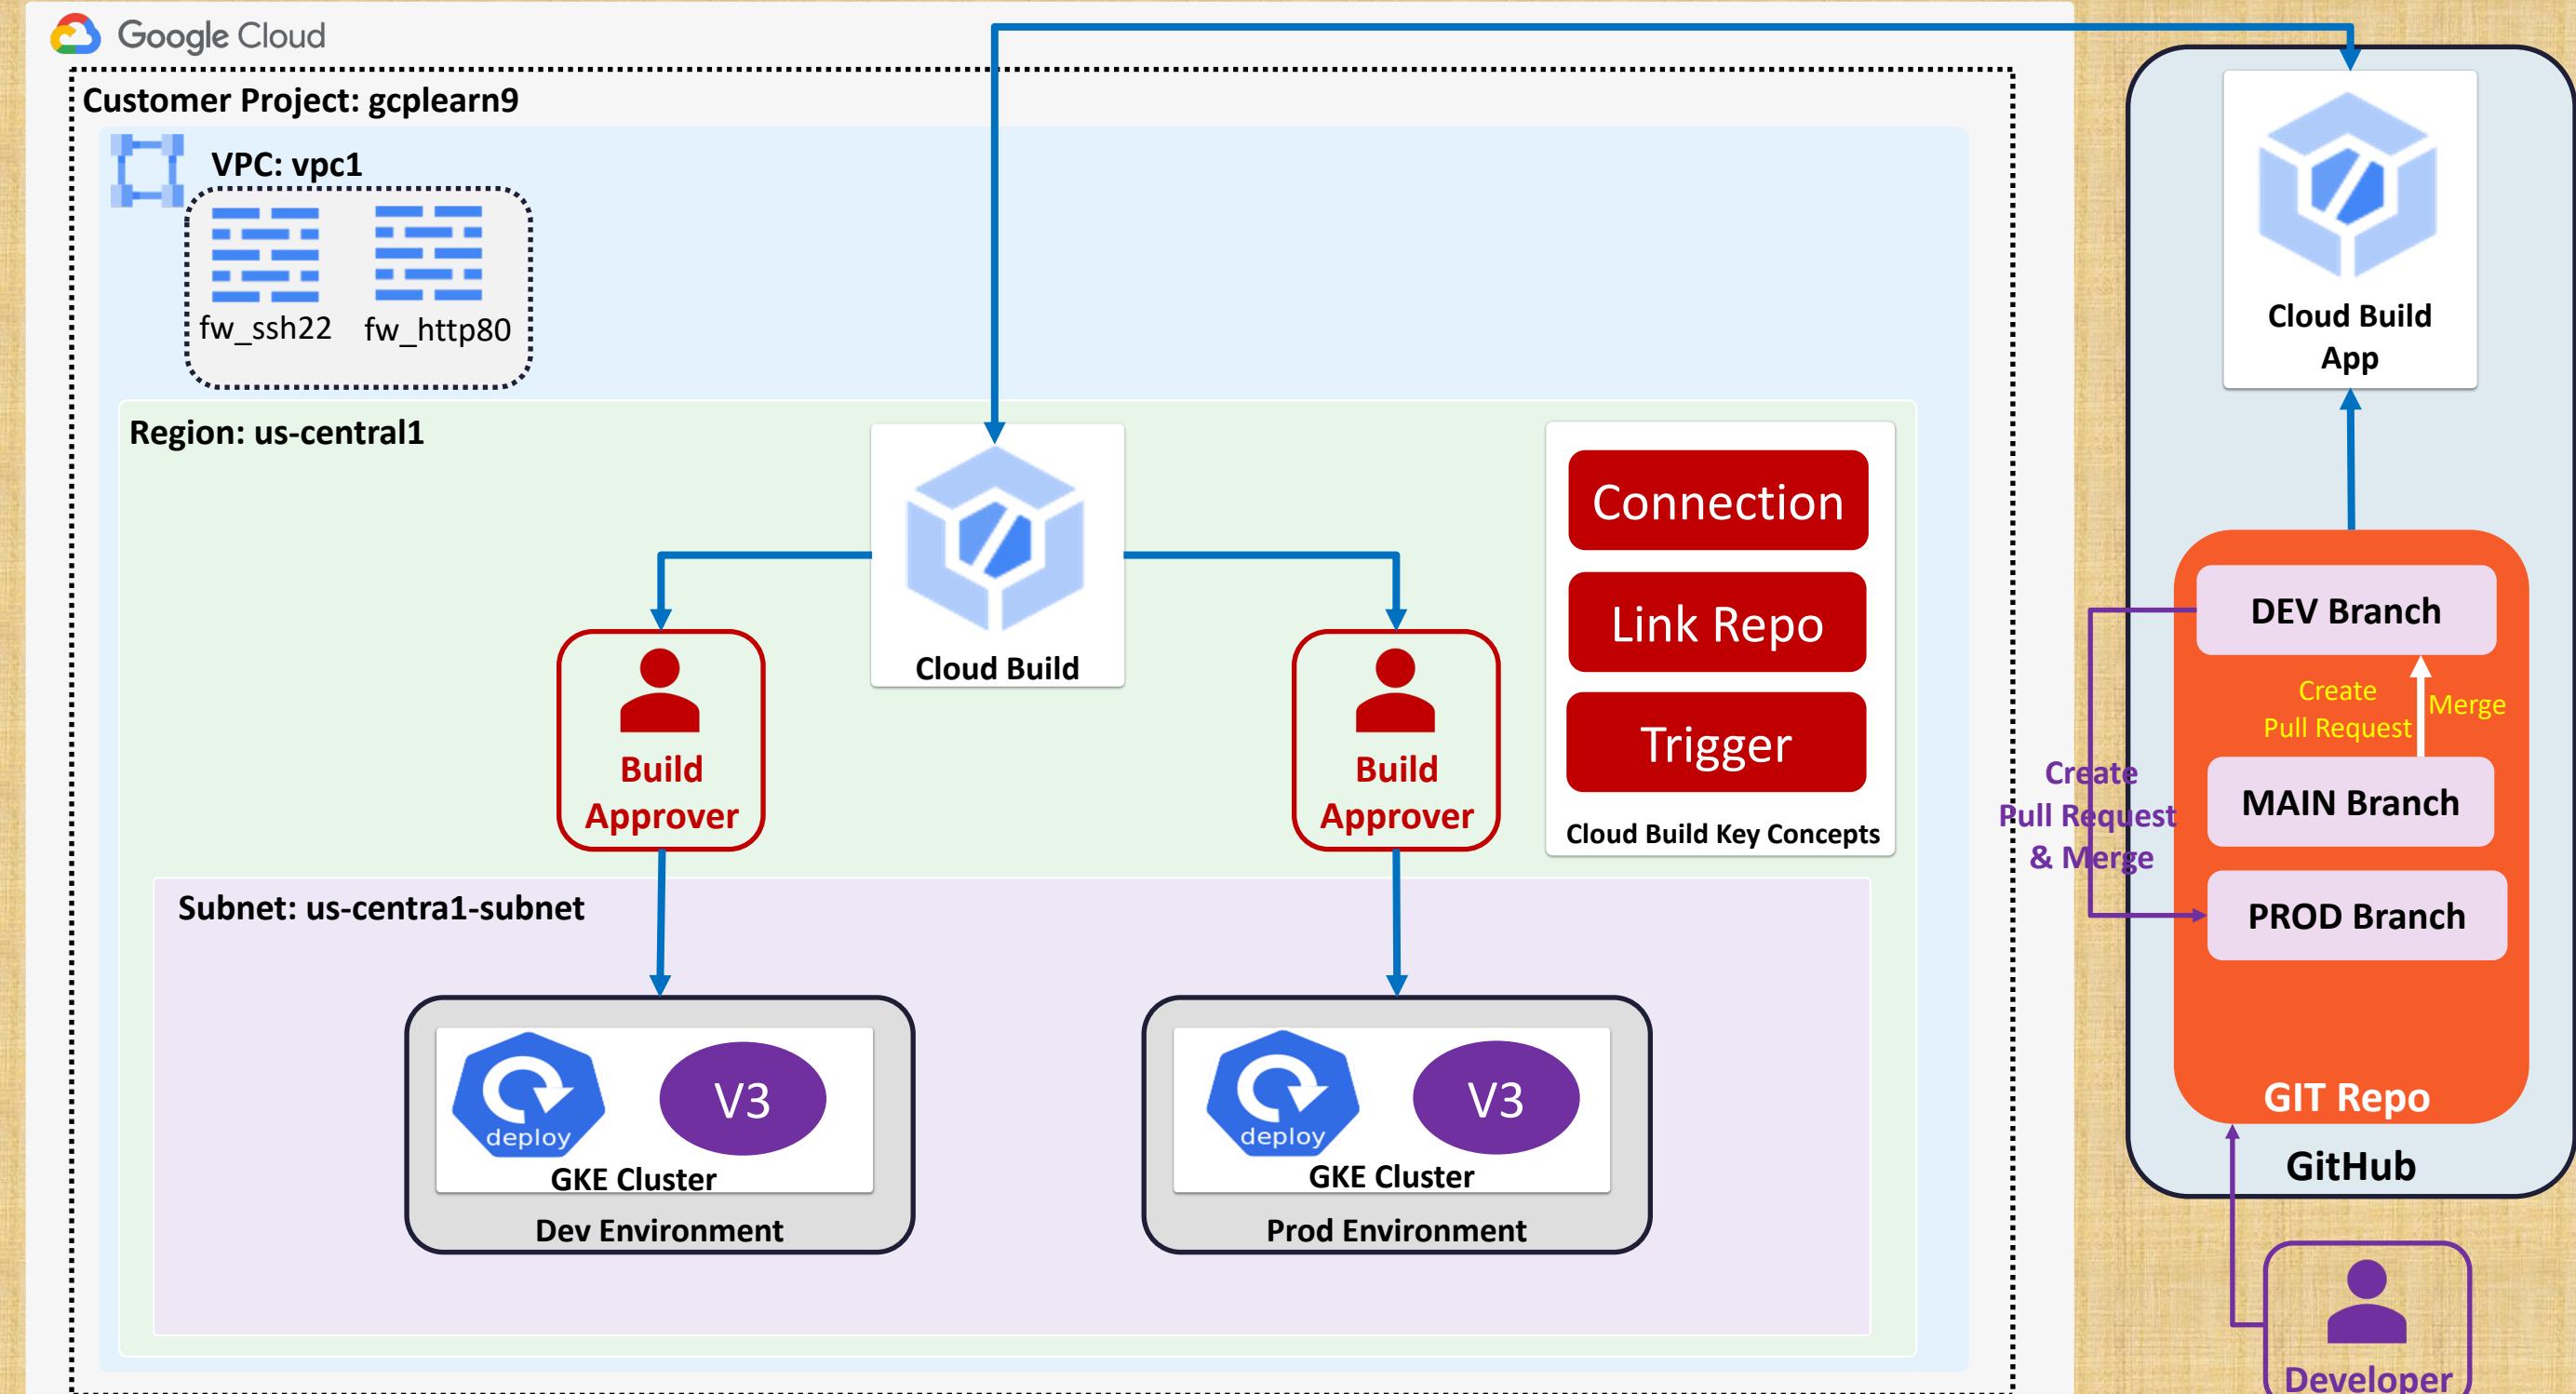
Admin



# GKE CI CD Pipelines



# Kubernetes Deployment DevOps for Terraform Configs



Author: Nho Luong

Skill: DevOps Engineer Lead

# END OF INTRODUCTION

# Terraform on Google Kubernetes Engine

## Terraform Fundamentals

Terraform CLI Install

1

Terraform Commands

2

Terraform Language Basics

3

Meta-Argument Provider

4

Input Variables &  
Output Values

5

Meta-Argument count

6

Terraform Data sources

7

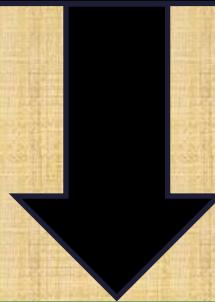
Meta-Argument for\_each

8



GCP Terraform on Google Cloud : DevOps SRE 30 Real-World Demos

Sections  
common



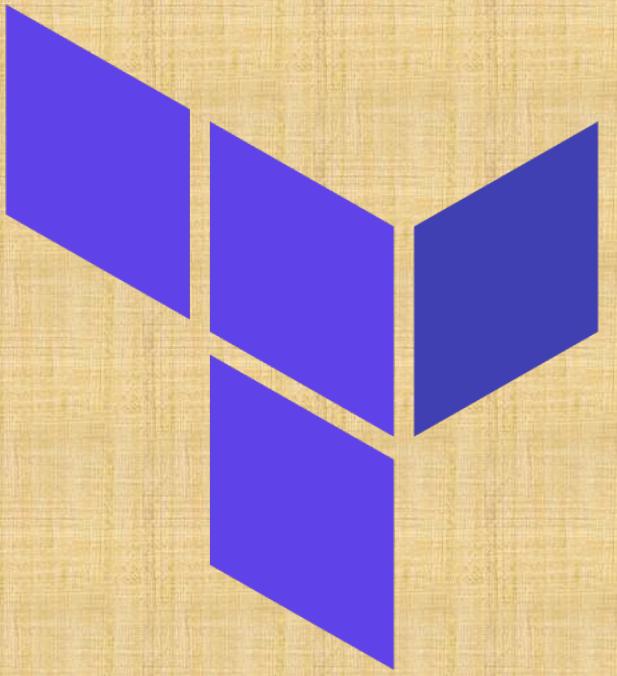
in both  
courses

GCP GKE Terraform on Google Kubernetes Engine DevOps SRE IaC



Demo

# Terraform Installation CLI Tools



# Terraform Installation

gcloud CLI

Terraform CLI

VS Code Editor

Terraform plugin  
for VS Code

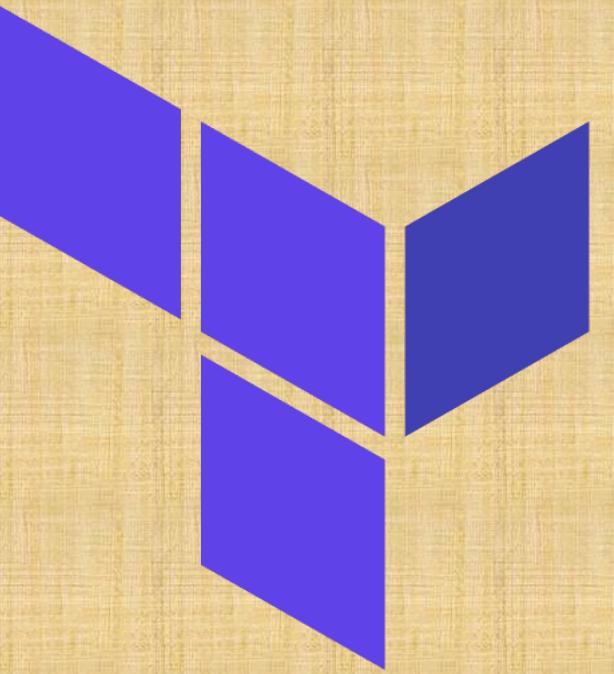
Mac OS

Windows OS

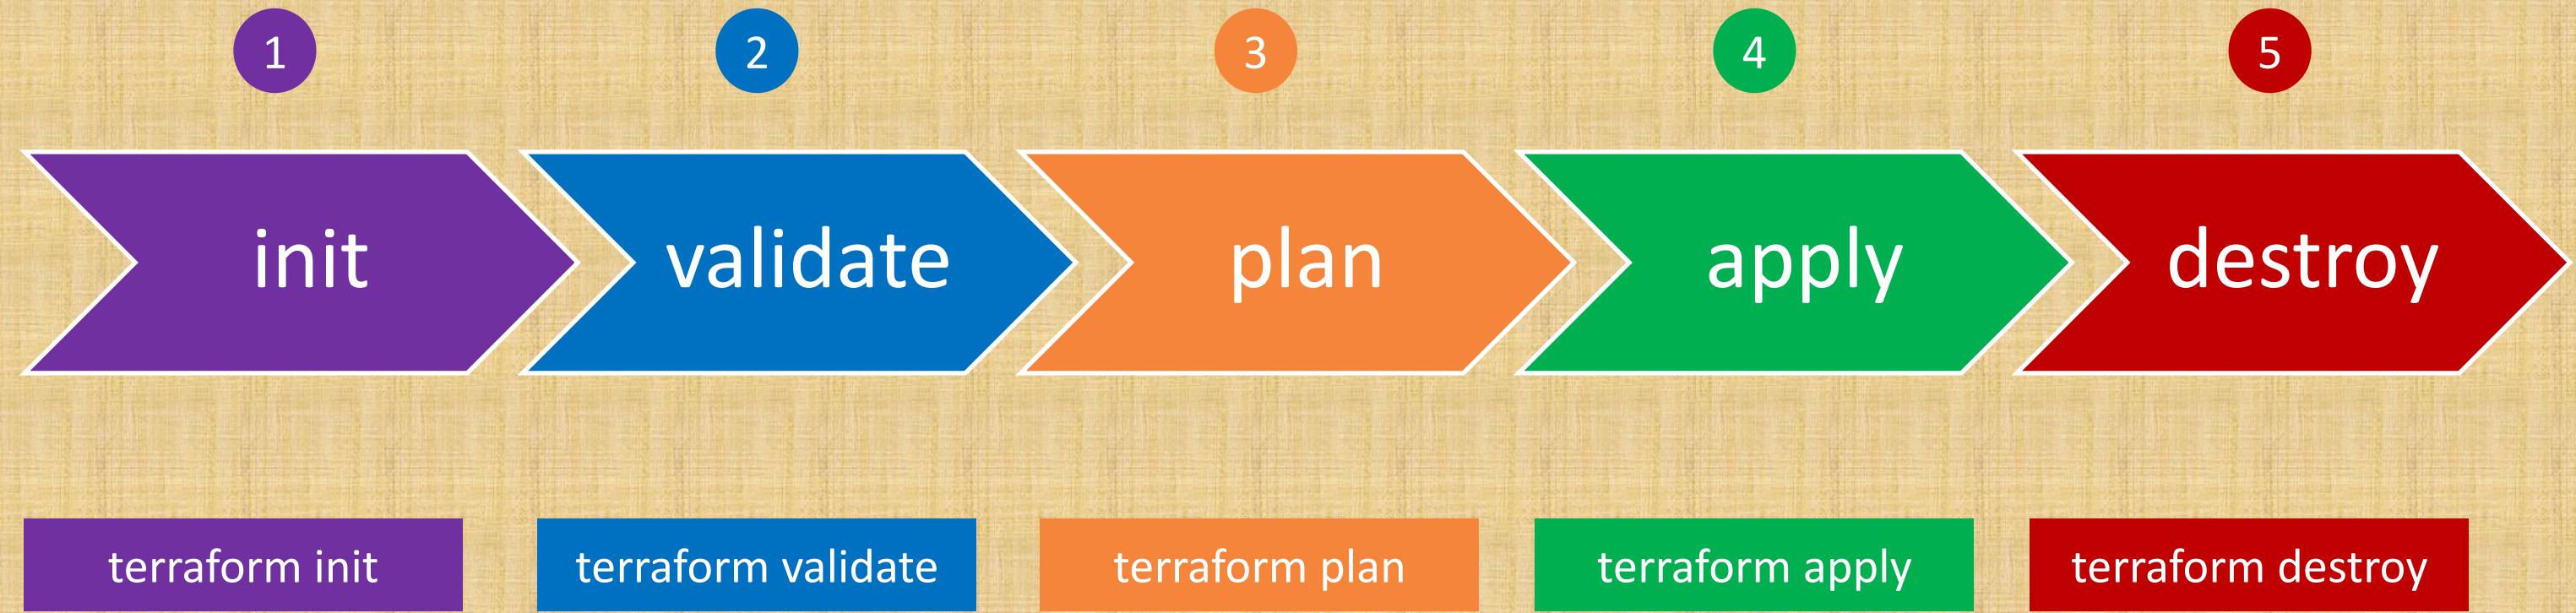


Demo

# Terraform Command Basics



# Terraform Workflow



# Terraform Workflow

1

init

2

validate

3

plan

4

apply

5

destroy

- Used to **Initialize** a working directory containing terraform config files
- This is the first command that should be run after writing a new Terraform configuration
- Downloads **Providers**

- Validates the terraform configurations files in that respective directory to ensure they are **syntactically valid** and **internally consistent**.

- Creates an **execution plan**
- Terraform performs a refresh and determines what actions are necessary to achieve the **desired state** specified in configuration files

- Used to apply the changes required to reach the **desired state** of the configuration.
- By default, apply scans the current directory for the configuration and applies the changes appropriately.

- Used to destroy the Terraform-managed infrastructure
- This will ask for confirmation before destroying.

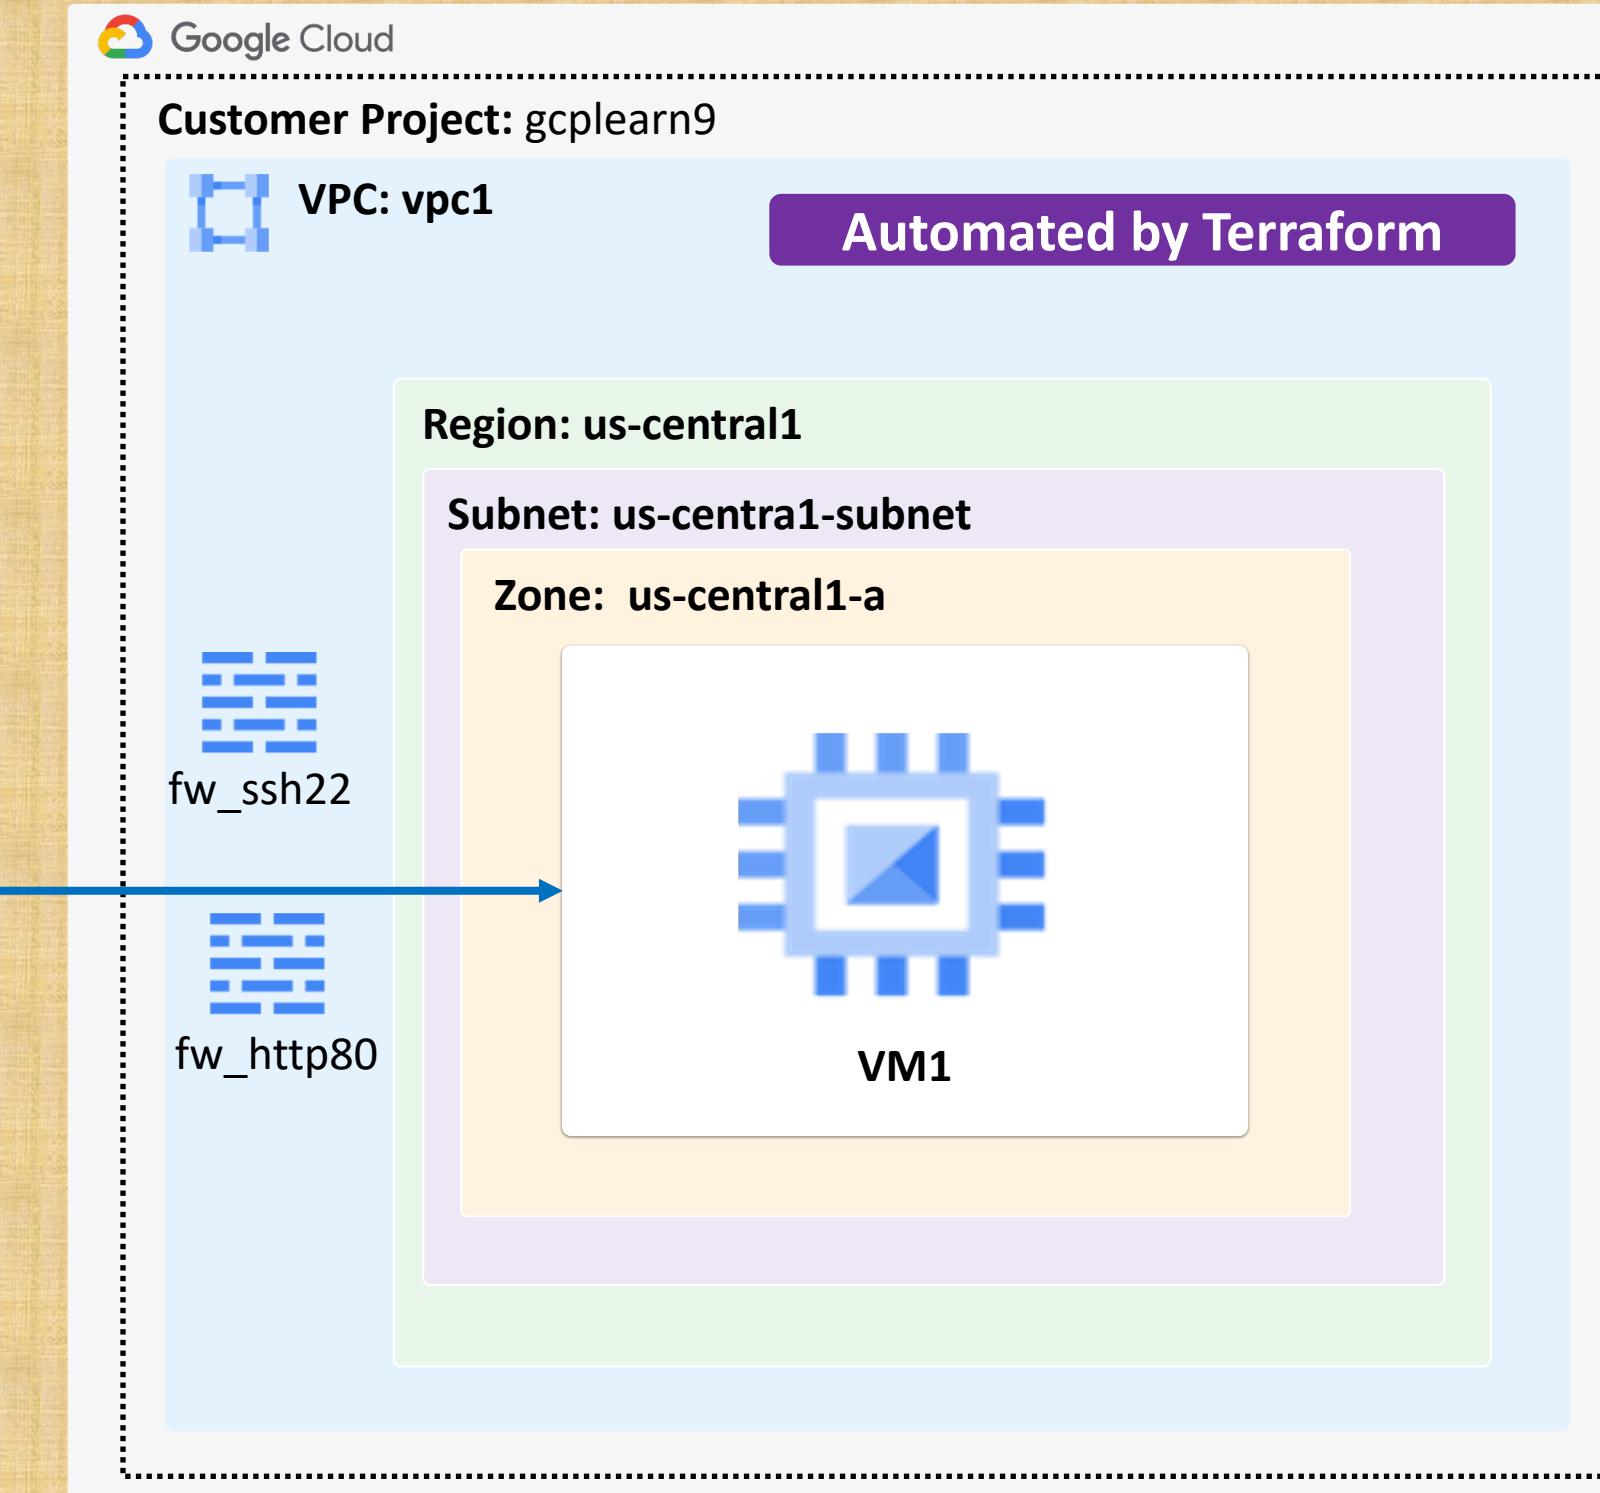


Demo

**Terraform**  
What are we going  
to automate as part of this demo  
using  
**Terraform?**



# Terraform Language Basics

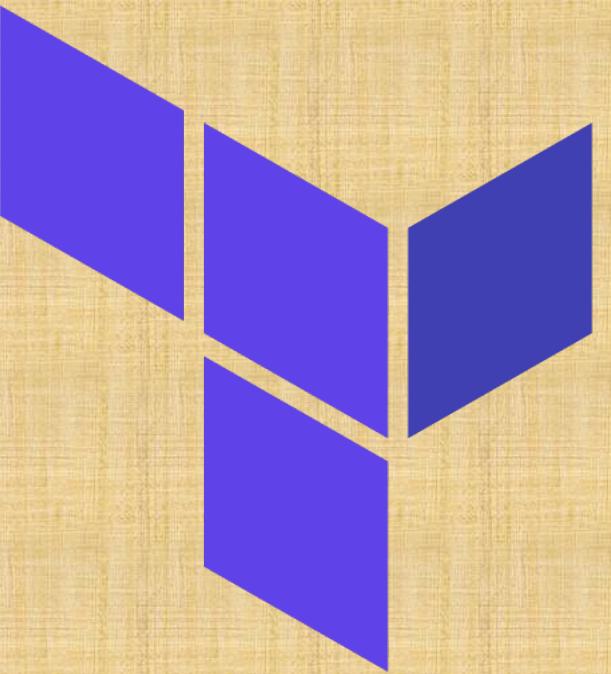




Demo

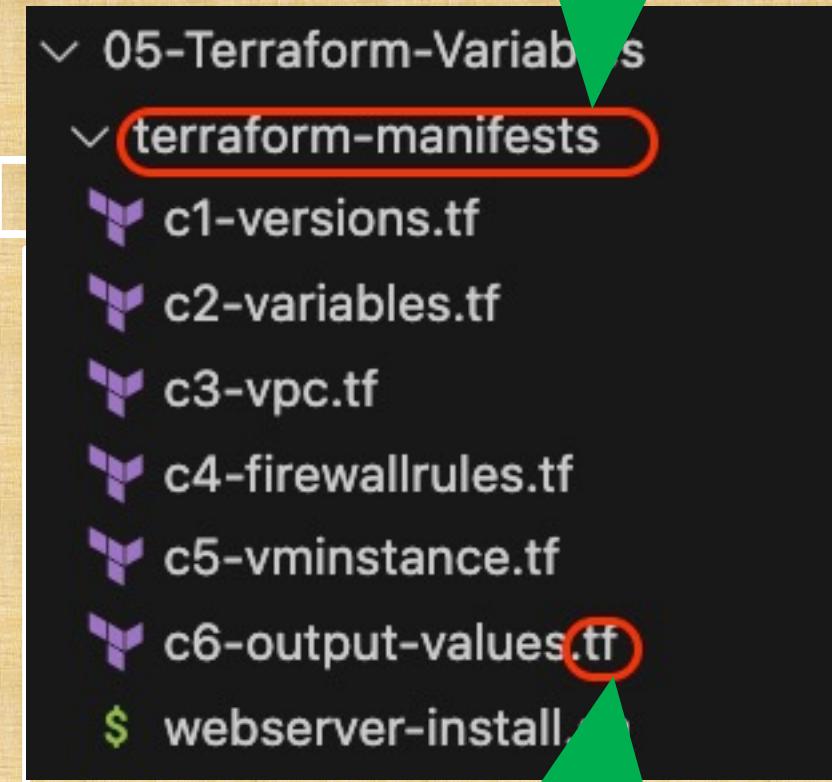
# Terraform

## Language Basics



# Terraform Language Basics - Files

- Code in the Terraform language is stored in plain text files with the **.tf** file extension.
- There is also a **JSON-based** variant of the language that is named with the **.tf.json** file extension.
- We can call the files containing terraform code as **Terraform Configuration Files** or **Terraform Manifests**
- Terraform files should be saved in a directory (**Terraform Working Directory**)

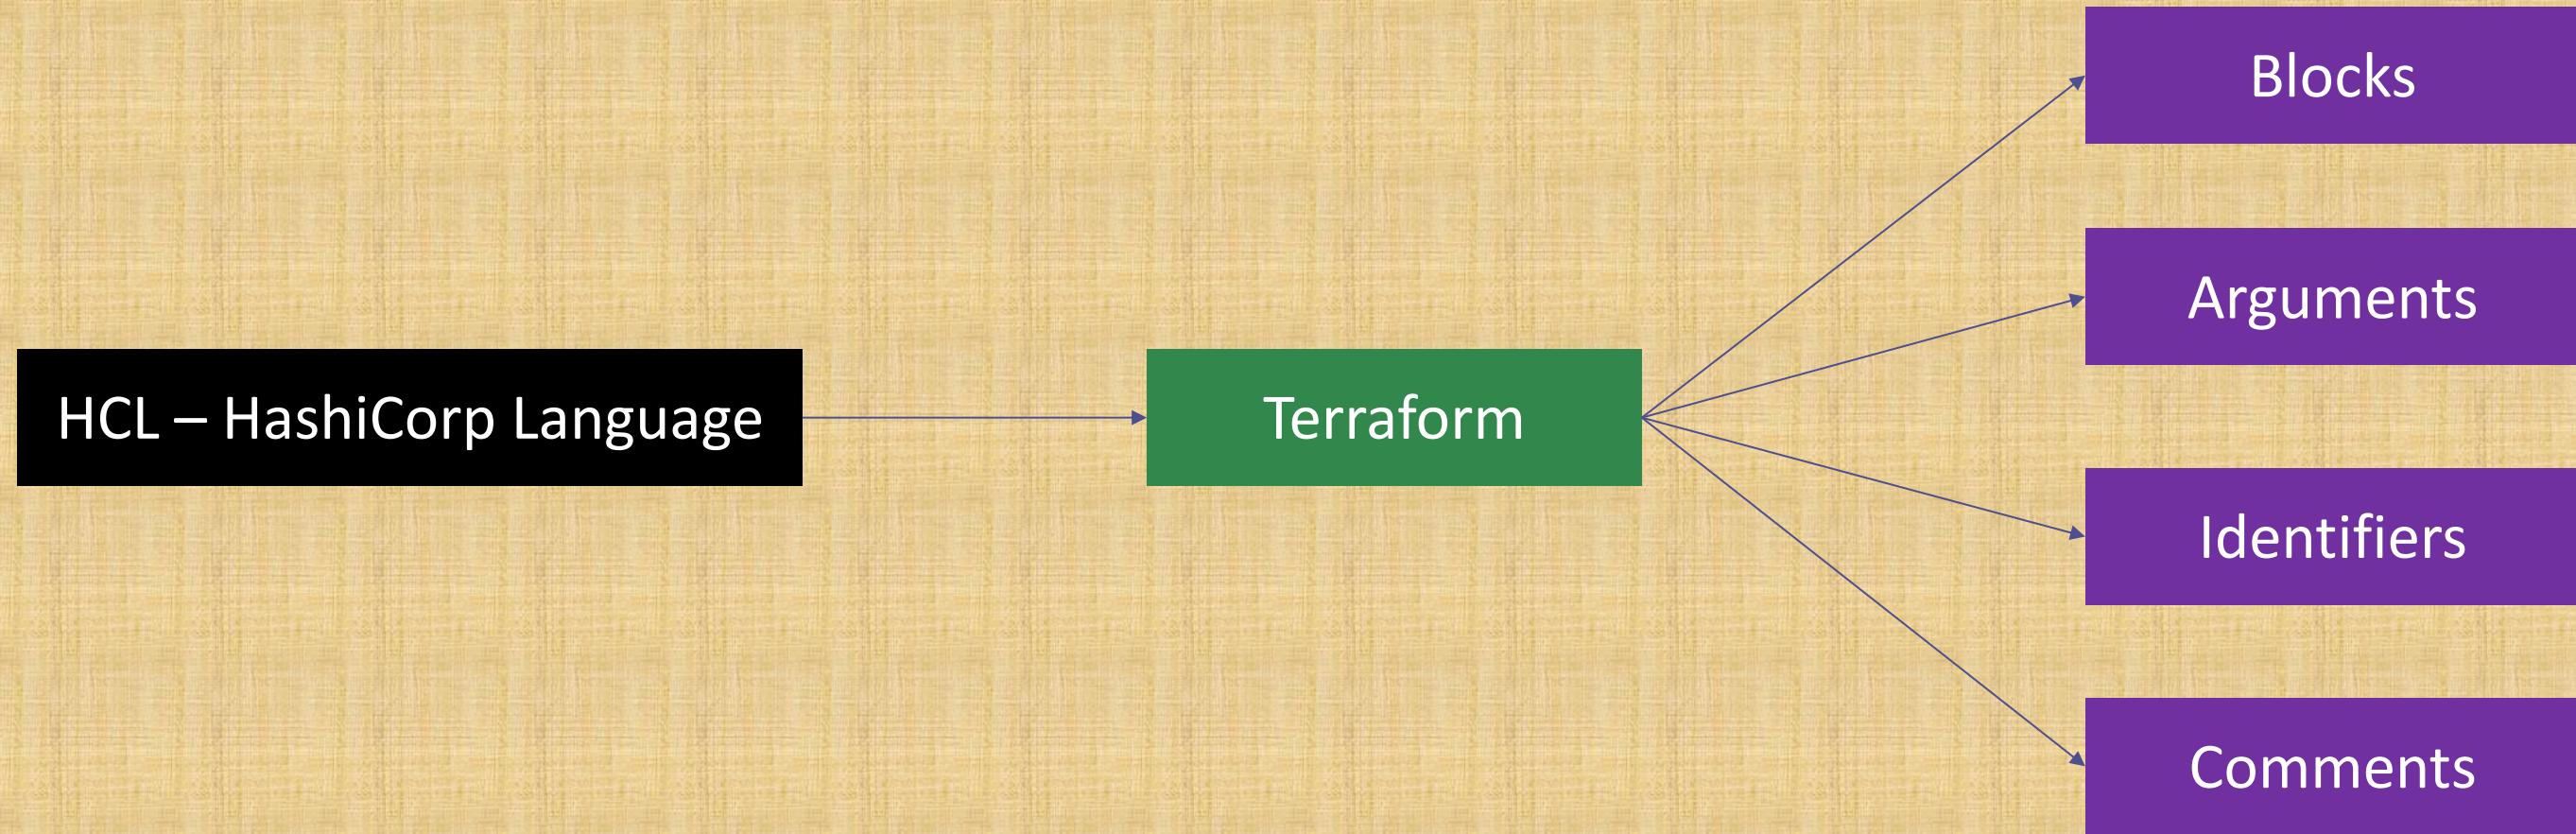


```
05-Terraform-Variab.../terraform-manifests/c1-versions.tf
05-Terraform-Variab.../terraform-manifests/c2-variables.tf
05-Terraform-Variab.../terraform-manifests/c3-vpc.tf
05-Terraform-Variab.../terraform-manifests/c4-firewallrules.tf
05-Terraform-Variab.../terraform-manifests/c5-vminstance.tf
05-Terraform-Variab.../terraform-manifests/c6-output-values.tf
$ webserver-install...
```

Terraform Working  
Directory

Terraform Configuration Files  
ending with **.tf** as extension

# Terraform Language Basics - Configuration Syntax



# Terraform Language Basics - Configuration Syntax

```
# Template
<BLOCK TYPE> "<BLOCK LABEL>" "<BLOCK LABEL>" {
  # Block Body
  <IDENTIFIER> = <EXPRESSION> # Argument
}
```

# Resource: Subnet

```
resource "google_compute_subnetwork" "mysubnet" {
  name = "subnet1"
  region = "us-central1"
  ip_cidr_range = "10.128.0.0/20"
  network = google_compute_network.myvpc.id
}
```

Block Type

Top Level & Block inside Blocks

Block Labels

Arguments

Based on Block Type,  
block labels can be 1 or 2  
**Example:**  
Resources: 2 block labels  
Variables: 1 block label

**Top Level Blocks:** resource, provider, terraform, variable, data, module, output

**Block Inside Block:** provisoners, resource specific blocks like tags

# Terraform Language Basics - Configuration Syntax

```
# Template
<BLOCK TYPE> "<BLOCK LABEL>" "<BLOCK LABEL>" {
  # Block Body
  <IDENTIFIER> = <EXPRESSION> # Argument
}

# Resource: Subnet
resource "google_compute_subnetwork" "mysubnet" {
  name      = "subnet1"
  region    = "us-central1"
  ip_cidr_range = "10.128.0.0/20"
  network   = google_compute_network.myvpc.id
}
```

Argument  
Name  
[OR]  
Identifier

Argument  
Value  
[OR]  
Expression

# Terraform Language Basics - Configuration Syntax

Multi-line comment /\* ..... \*/

```
/*
# Template
<BLOCK TYPE> "<BLOCK LABEL>" "<BLOCK LABEL>" {
  # Block Body
  <IDENTIFIER> = <EXPRESSION> # Argument
}
*/
```

Single Line Comments with # or //

```
# Resource: Subnet
resource "google_compute_subnetwork" "mysubnet" {
  name        = "subnet1"
  region      = "us-central1"
  ip_cidr_range = "10.128.0.0/20"
  network     = google_compute_network.myvpc.id // GET VPC ID
}
```

**Terraform** language uses a **limited** number of **top-level block** types, which are **blocks** that can appear **outside** of any other **block** in a TF configuration file.

## Terraform Top-Level Blocks

Most of **Terraform's features** are implemented as **top-level** blocks.

Terraform Block

Providers Block

Resources Block

Fundamental Blocks

Input Variables Block

Output Values Block

Local Values Block

Variable Blocks

Data Sources Block

Modules Block

Calling / Referencing Blocks

Import Block

Moved Block

Removed Block

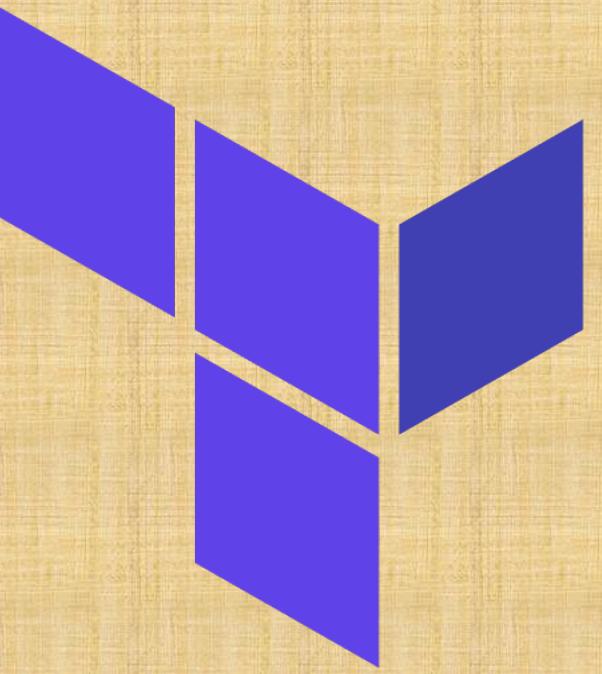
Check Block

**NEW BLOCKS – Added Recently**



Demo

# Terraform Fundamental Blocks



# Terraform Basic Blocks

## Terraform Block

Special block used to configure some **behaviors**

Specifying a **required Terraform Version**

Specifying **Provider Requirements**

Configuring a Terraform Backend (**Terraform State**)

## Provider Block

**HEART** of Terraform

Terraform relies on providers to **interact** with Remote Systems

Declare providers for Terraform to **install** providers & use them

Provider configurations belong to **Root Module**

## Resource Block

Each Resource Block describes one or more Infrastructure Objects

**Resource Syntax:** How to declare Resources?

**Resource Behavior:** How Terraform handles resource declarations?

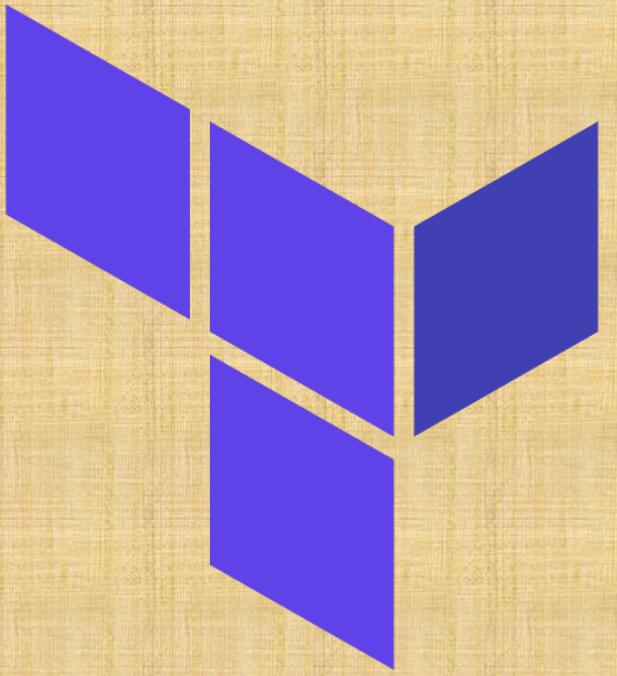
**Provisioners:** We can configure Resource post-creation actions



Demo

# Terraform

## Terraform Block



# Terraform Block

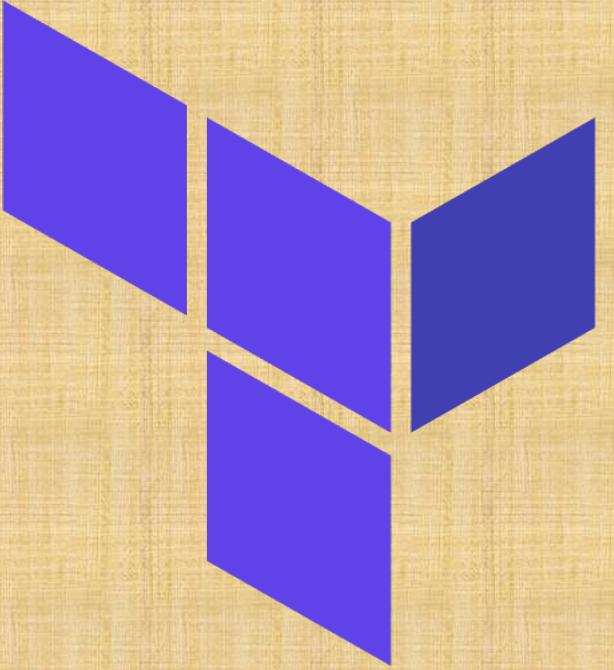
- This block can be called in 3 ways. All means the same.
  - Terraform Block
  - Terraform Settings Block
  - Terraform Configuration Block
- Each terraform block can contain a number of settings related to Terraform's behavior.
- Within a terraform block, **only constant values can be used**; arguments **may not refer** to named objects such as resources, input variables, etc, and **may not use any** of the Terraform language built-in functions.

```
# Terraform Settings Block
terraform {
    required_version = ">= 1.8"
    required_providers {
        google = {
            source  = "hashicorp/google"
            version = ">= 5.26.0"
        }
    }
    backend "gcs" {
        bucket  = "kalyanbucket201"
        prefix  = "terraform/state"
    }
}
```



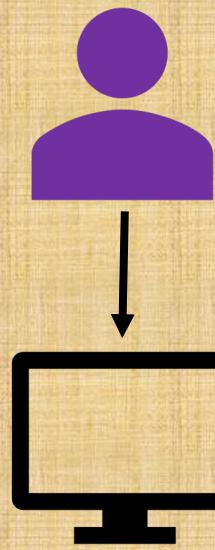
Demo

# Terraform Providers



# Terraform Providers

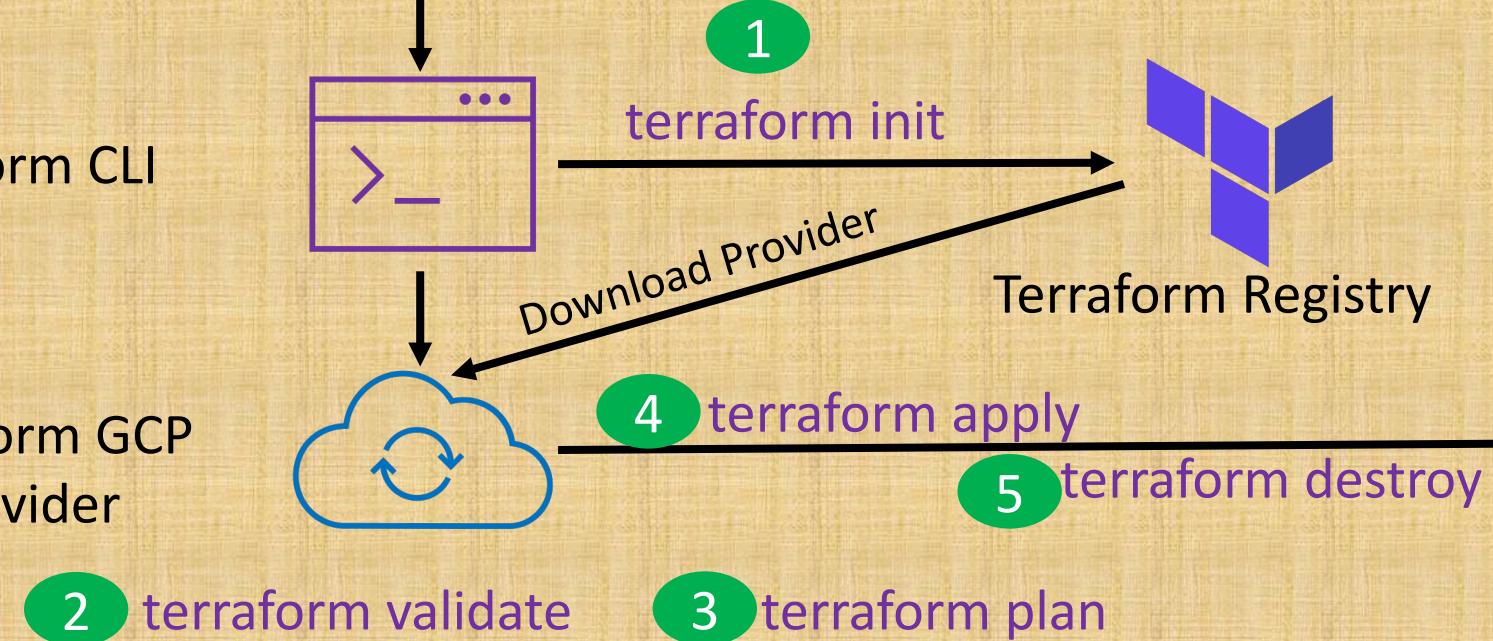
Terraform Admin



Local Desktop

Terraform CLI

Terraform GCP Provider



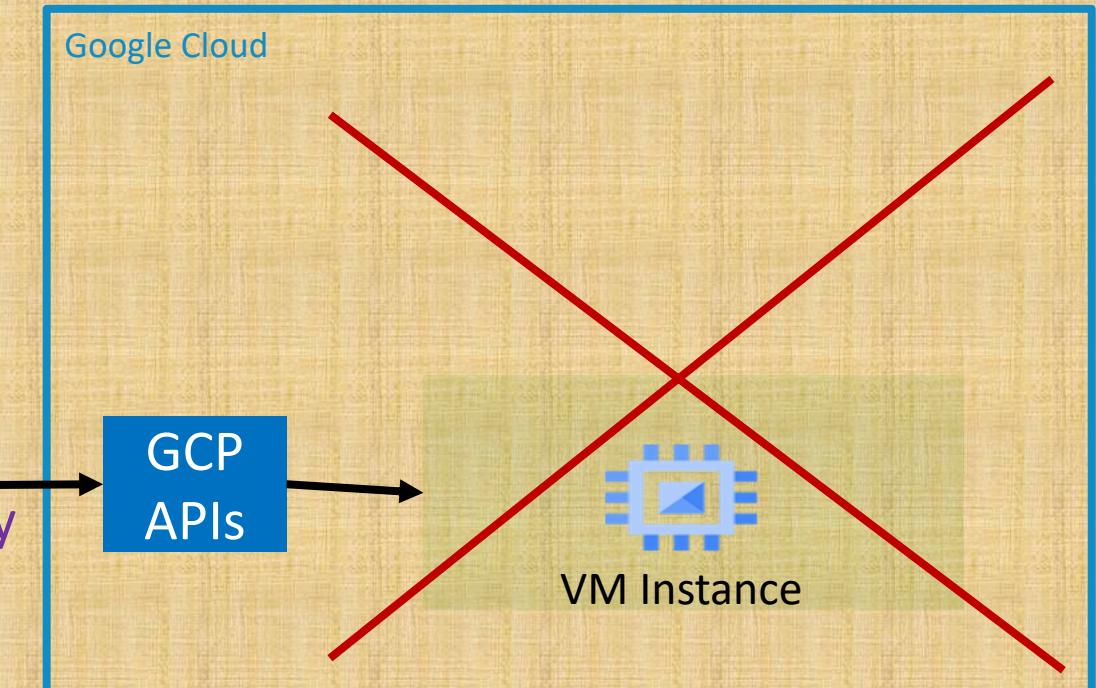
Providers are **HEART** of Terraform

Every **Resource Type** (example: VM Instance), is implemented by a Provider

Without Providers Terraform **cannot** manage any infrastructure.

Providers are distributed separately from Terraform and each provider has its own **release cycles** and **Version Numbers**

Terraform **Registry** is publicly available which contains many Terraform Providers for most **major** Infra Platforms



## Provider Requirements

```
# Terraform Settings Block
terraform {
  required_version = ">= 1.8"
  required_providers {
    google = {
      source = "hashicorp/google"
      version = ">= 5.26.0"
    }
  }
}
```

## Terraform Providers

## Provider Configuration

```
# Terraform Provider Block
provider "google" {
  project = "gcplearn9"
  region = "us-central1"
}
```

## Dependency Lock File

```
✓ 03-Terraform-Language-Basics
  ✓ terraform-manifests
    > .terraform
      ≡ .terraform.lock.hcl
    ✎ c1-versions.tf
    ✎ c2-variables.tf
    ✎ c3-vpc.tf
    ✎ c4-firewallrules.tf
    ✎ c5-vminstance.tf
    $ webserver-install.sh
```

# Dependency Lock File

```
terraform-on-google-cloud > 03-Terraform-Language-Basics > terraform-manifests > .terraform.lock.hcl
1 # This file is maintained automatically by "terraform init".
2 # Manual edits may be lost in future updates.
3
4 provider "registry.terraform.io/hashicorp/google" {
5   version      = "5.26.0"
6   constraints = ">= 5.26.0"
```

# Required Providers

```
# Terraform Settings Block
terraform {
  required_version = ">= 1.8"
  required_providers {
    google = {
      source = "hashicorp/google"
      version = ">= 5.26.0"
    }
  }
}

# Terraform Provider Block
provider "google" {
  project = "gcplearn9"
  region = "us-central1"
}
```

## Local Names

Local Names are **Module specific** and should be **unique per-module**

Terraform configurations always refer to **local name** of provider **outside required\_provider block**

Users of a provider can choose **any local name** for it (mygoogle, google1).

Recommended way of choosing local name is to use preferred local name of that provider (For Google Provider: hashicorp/google, **preferred local name** is google)

## Source

It is the **primary location** where we can download the Terraform Provider

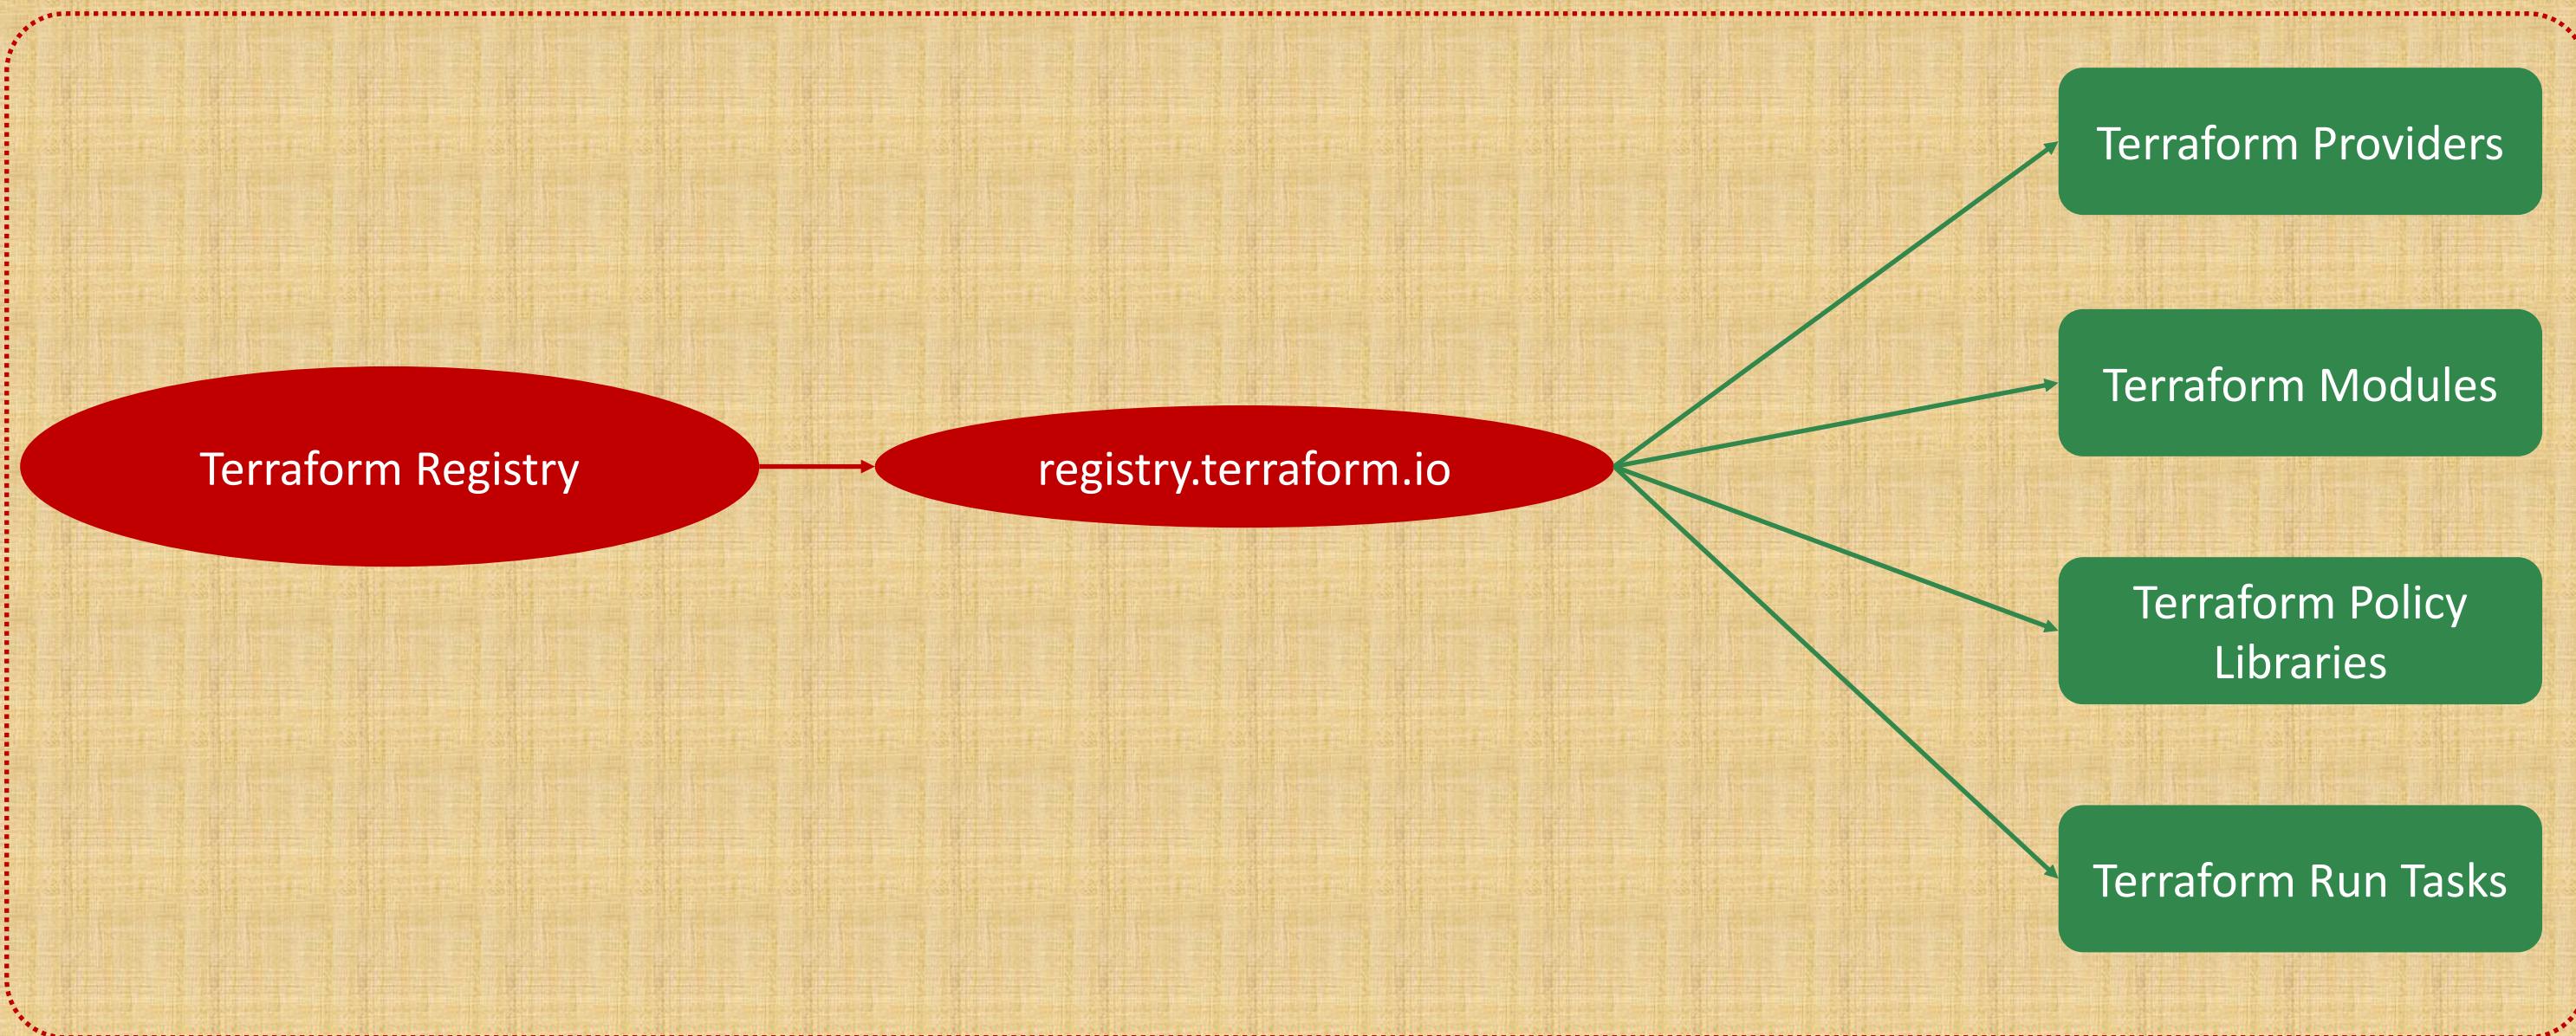
Source addresses consist of **three parts** delimited by **slashes (/)**

[<HOSTNAME>/]<NAMESPACE>/<TYPE>

registry.terraform.io/hashicorp/google

Registry Name is **optional** as default is going to be Terraform Public Registry

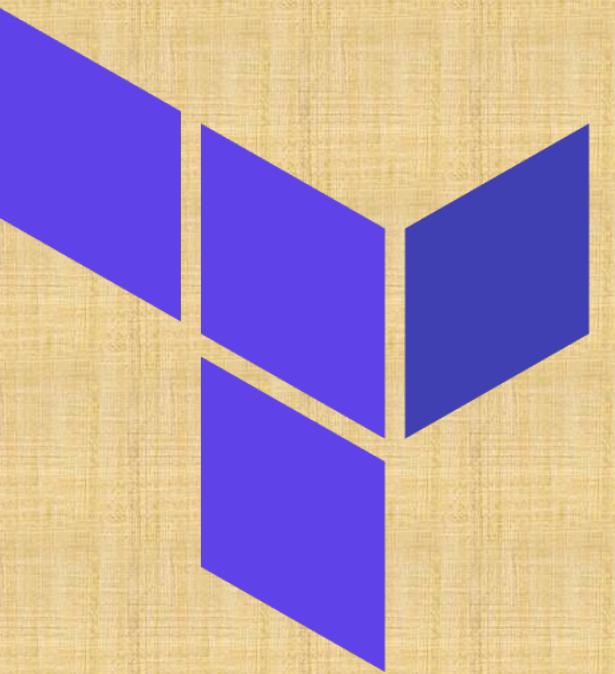
# Terraform Registry





Demo

# Terraform State

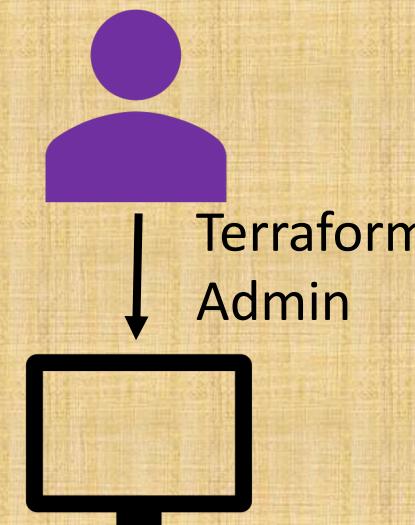


# Terraform State

Local Desktop

Terraform CLI

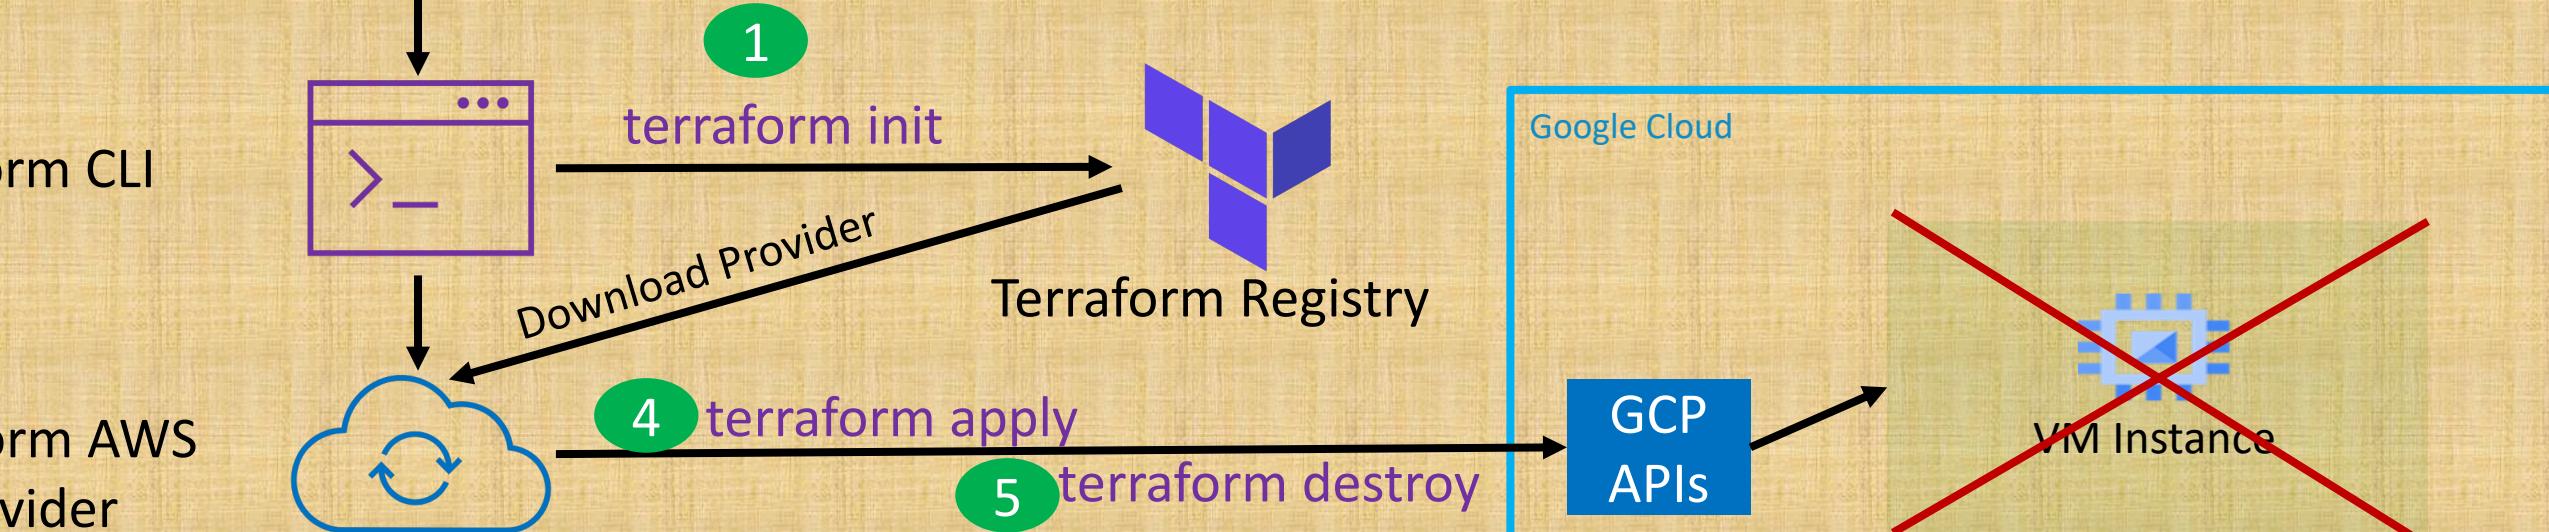
Terraform AWS Provider



Terraform must **store state** about your managed infrastructure and configuration

This state is used by Terraform to map **real world resources** to your **configuration (.tf files)**, keep track of metadata, and to **improve performance** for large **infrastructures**.

This state is stored by default in a local file named "**terraform.tfstate**", but it can also be stored **remotely**, which works better in a **team** environment.



**Terraform State**  
File : **terraform.tfstate**



# Desired & Current Terraform States

Terraform Configuration Files

- c1-versions.tf
- c2-variables.tf
- c3-vpc.tf
- c4-firewallrules.tf
- c5-vminstance.tf

Desired State

Real World Resource  
VM Instance, VPC, Subnet,  
Firewall Rules



vpc1	1	1460	Custom
subnet1	us-central1	IPv4	10.128.0.0/20
<a href="#">fwrule-allow-ssh22</a>			Ingress firewall rule
<a href="#">fwrule-allow-http80</a>			Ingress firewall rule
<input type="button"/> Filter Enter property name or value			
<input type="checkbox"/>	Status	Name ↑	Zone Machine type
<input checked="" type="checkbox"/>		<a href="#">myapp1</a>	us-central1-a e2-micro

Current State

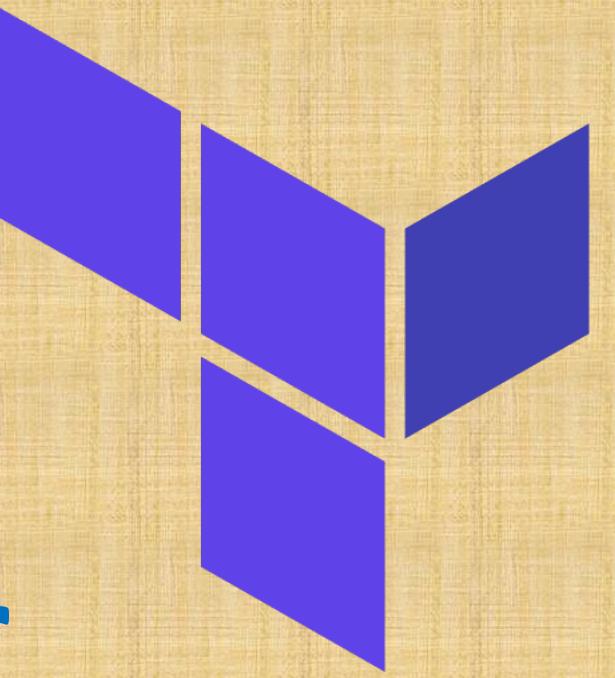
Demo-4



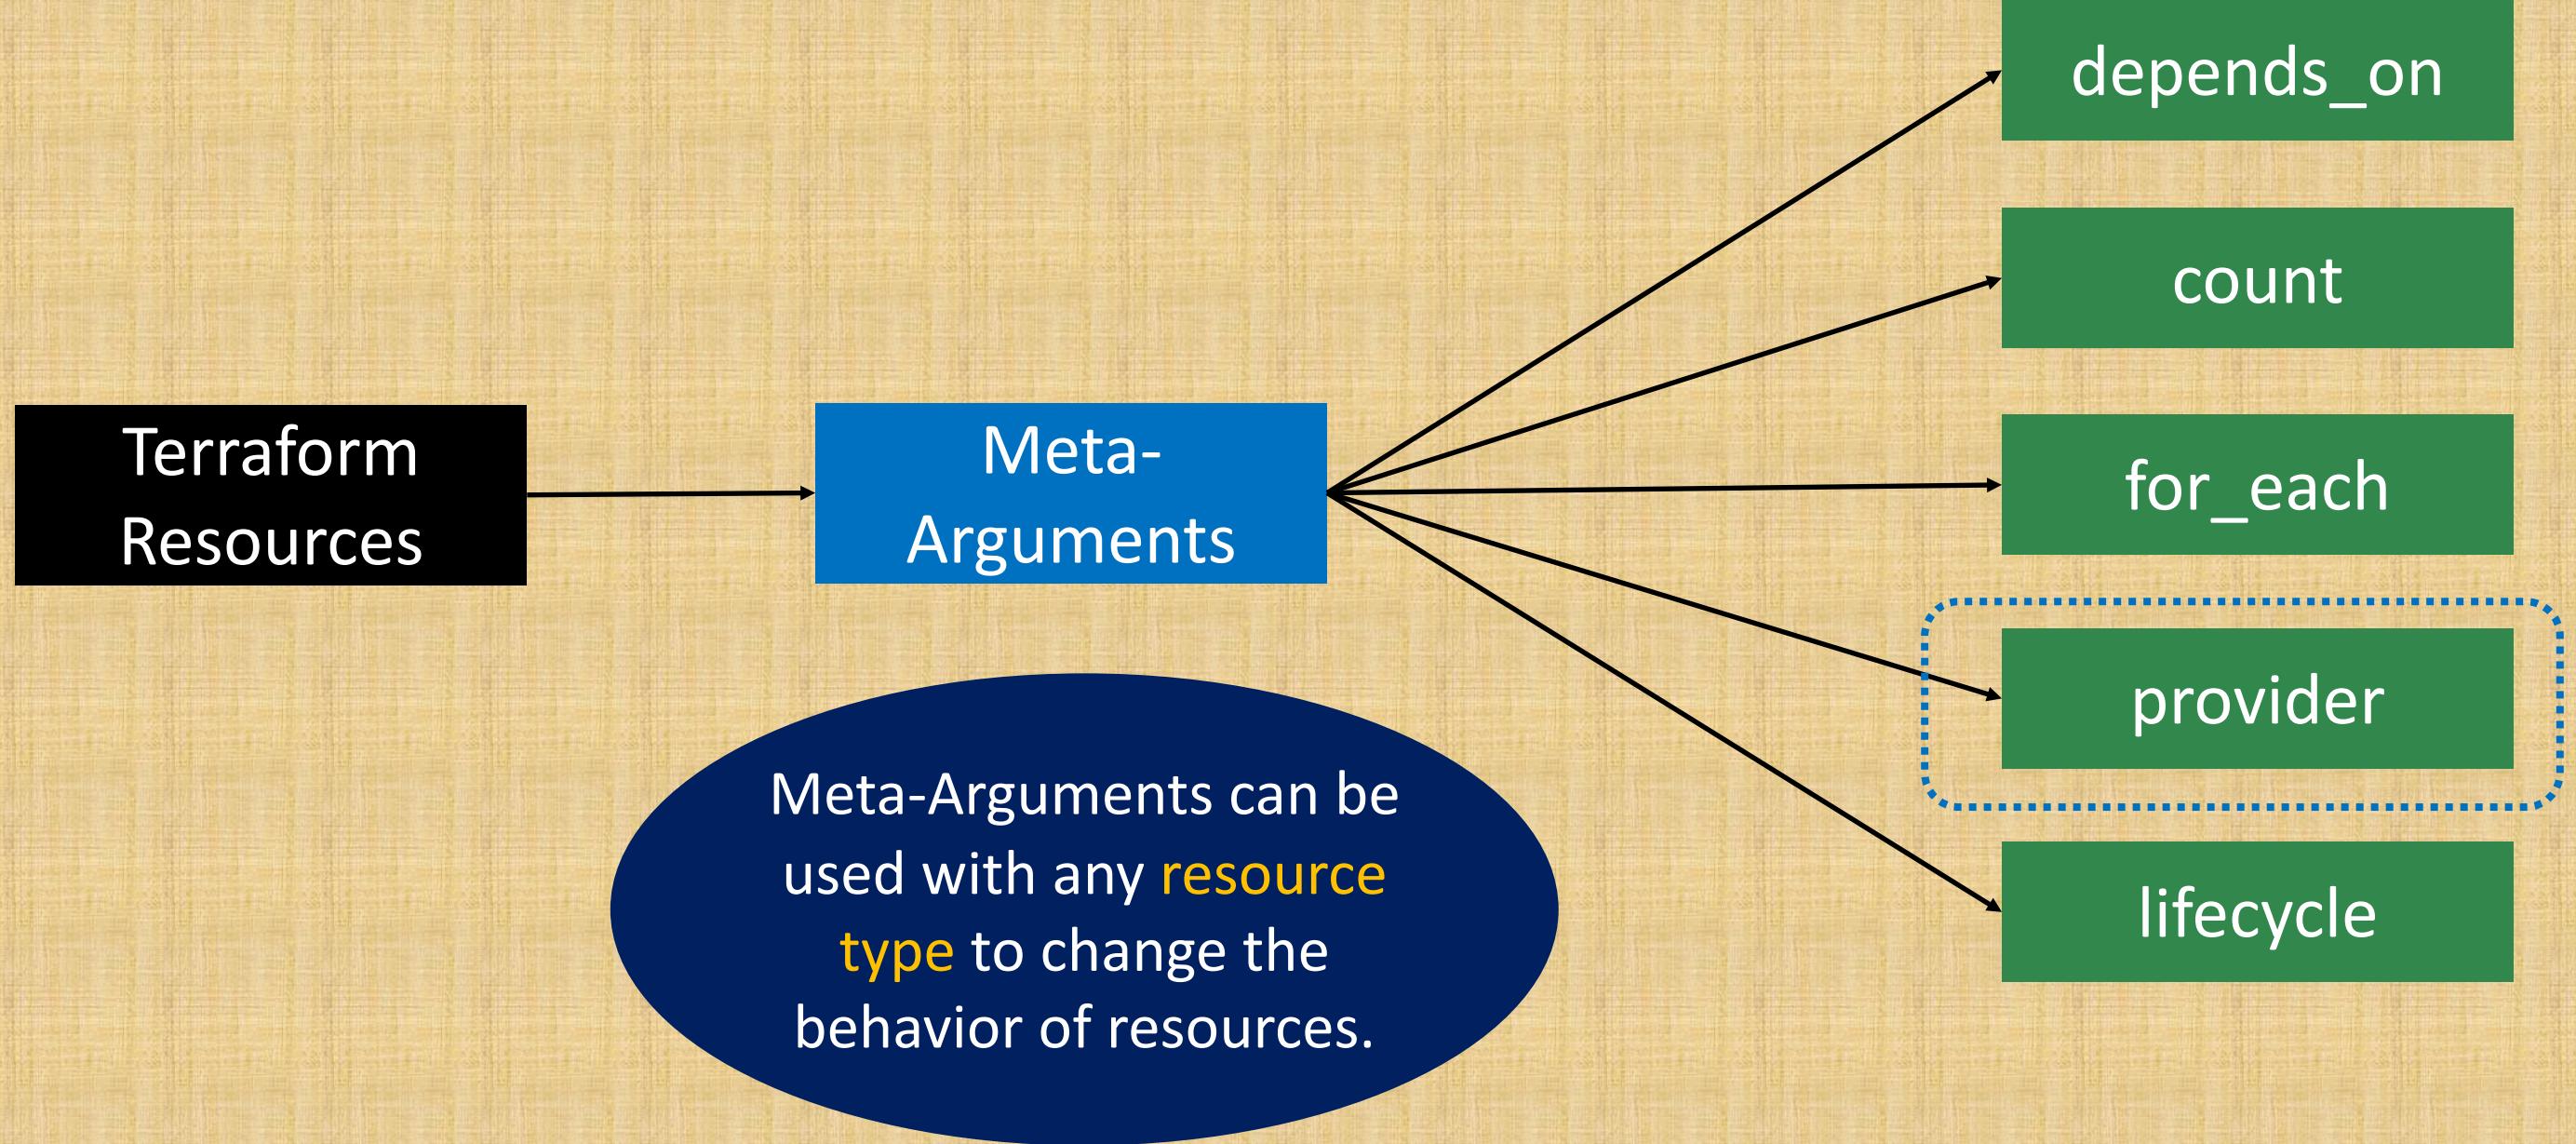
# Terraform

## Multiple Providers

### Meta-Argument: provider



# Resource Meta-Arguments



# Terraform Multiple Providers

- We can define **multiple configurations** for the same provider (Ex: google) and select which one to use **on a per-resource or per-module basis**.
- The primary reason for this is to **support multiple regions for a cloud platform**, additionally for **multiple environments** (dev, prod) considering each environment has different provider configurations (**Example: different authentication credentials for dev and prod**)
- We can use the alternate provider in a resource, data or module by referencing it as **<PROVIDER NAME>.<ALIAS>**

```
# Terraform Provider-1: us-central1
provider "google" {
  project = "gcplearn9"
  region = "us-central1"
  alias = "us-central1"
}

# Terraform Provider-2: europe-west1
provider "google" {
  project = "gcplearn9"
  region = "europe-west1"
  alias = "europe-west1"
}
```

```
# Resource: Subnet1
resource "google_compute_subnetwork" "mysubnet1" {
  provider = google.us-central1 # Define provider to use
  name = "subnet1"
  ip_cidr_range = "10.128.0.0/20"
  network = google_compute_network.myvpc.id
}

# Resource: Subnet2
resource "google_compute_subnetwork" "mysubnet2" {
  provider = google.europe-west1 # Define provider to use
  name = "subnet2"
  ip_cidr_range = "10.132.0.0/20"
  network = google_compute_network.myvpc.id
}
```

# Terraform Multiple Providers

- We can define **multiple configurations** for the same provider (Ex: google) and select which one to use **on a per-resource or per-module basis**.
- **Usage**
  - To support **multiple regions** for a cloud platform (we will practically implement this)
  - For supporting **multiple environments** (dev, prod) considering each environment has different provider configurations
    - **Example:** Different authentication credentials for dev and prod configured in their respective provider blocks
  - We can use the alternate provider in a resource, data or module by referencing it as <PROVIDER NAME>.<ALIAS>

```
# Terraform Provider-1: us-central1
provider "google" {
  project = "gcplearn9"
  region  = "us-central1"
  alias   = "us-central1"
}

# Terraform Provider-2: europe-west1
provider "google" {
  project = "gcplearn9"
  region  = "europe-west1"
  alias   = "europe-west1"
}

# Resource: Subnet1
resource "google_compute_subnetwork" "mysubnet1" {
  provider = google.us-central1 # Define provider to use
  name     = "subnet1"
  ip_cidr_range = "10.128.0.0/20"
  network  = google_compute_network.myvpc.id
}

# Resource: Subnet2
resource "google_compute_subnetwork" "mysubnet2" {
  provider = google.europe-west1 # Define provider to use
  name     = "subnet2"
  ip_cidr_range = "10.132.0.0/20"
  network  = google_compute_network.myvpc.id
}
```

# Terraform Multiple Providers



Customer Project: gcplearn9



VPC: vpc1

Automated by Terraform

```
# Terraform Provider-1: us-central1
provider "google" {
  project = "gcplearn9"
  region  = "us-central1"
  alias   = "us-central1"
}
```

```
# Terraform Provider-2: europe-west1
provider "google" {
  project = "gcplearn9"
  region  = "europe-west1"
  alias   = "europe-west1"
}
```

Region: us-central1

Subnet: us-central1-subnet

Region: europe-west1

Subnet: europe-west1-subnet

# What are we going to learn?

```
✓ terraform-manifests  
$ app1-webserver-install.sh  
  c1-versions.tf  
  c2-variables.tf  
  c3-vpc.tf
```

Demo: Terraform Concept: Multiple Providers

C1

Update multiple providers

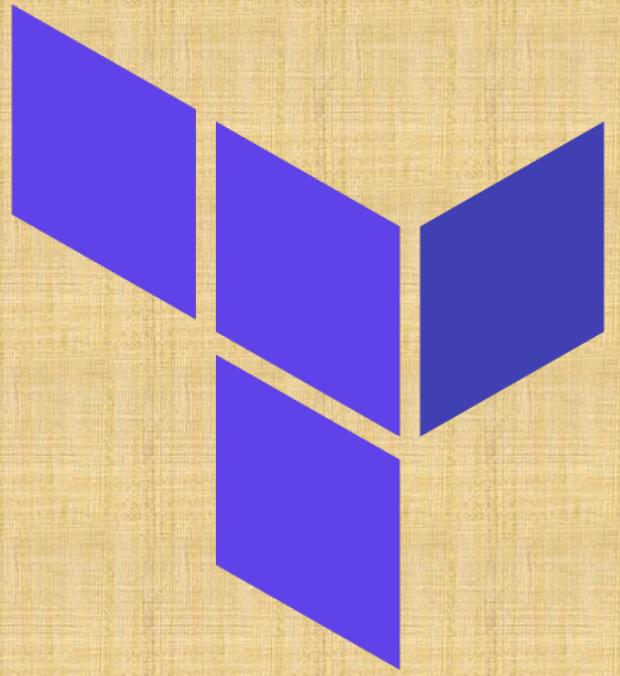
C2

Use provider alias in subnet Terraform resources



Demo-5

# Terraform Input Variables Output Values



**Demo: Terraform Concept: Input Variables and Output Values**

# Terraform Input Variables

- Input variables allow us to **customize** the Terraform Resources or modules without altering the source code.
- **Key Benefit:** Allows us to make the **code reusable** by just changing variables
- We have **multiple options** to declare variables
  - variables.tf (Default value)
  - terraform.tfvars
  - vm.auto.tfvars
  - vm.tfvars
  - Environment Variables
  - --var-file flag
  - --var flag

variables.tf

```
# GCP Compute Engine Machine Type
variable "machine_type" {
  description = "Compute Engine Machine Type"
  type        = string
  default     = "e2-small"
}
```

terraform.tfvars

```
terraform.tfvars ×
terraform-on-google-cloud > 05-Terraform-Variables-Output-Values >
  1 gcp_project      = "gcplearn9"
  2 gcp_region1     = "us-central1"
  3 machine_type    = "e2-micro"
```

# What are we going to learn?

## ✓ terraform-manifests

\$ app1-webserver-install.sh

- c1-versions.tf
- c2-variables.tf
- c3-vpc.tf
- c4-firewallrules.tf
- c5-vminstance.tf
- c6-output-values.tf
- terraform.tfvars
- vm.auto.tfvars
- vm.tfvars

### Demo: Terraform Concept: Input Variables

C2

Understand and define Terraform **Input Variables**

C1

Update **Provider** block with Terraform input variables

C3

Update **subnet region** with input variable

C5

Update **VM Instance machine type** using input variable

TF  
VARS

Play with **terraform.tfvars**, **vm.auto.tfvars** and **vm.tfvars**

## Output Values - Attributes

```
# Terraform Output Values
## ATTRIBUTES
output "vm_instanceid" {
  description = "VM Instance ID"
  value = google_compute_instance.myapp1.instance_id
}

output "vm_selflink" {
  description = "VM Instance Self link"
  value = google_compute_instance.myapp1.self_link
}
```

## Output Values - Arguments

```
## ARGUMENTS
output "vm_name" {
  description = "VM Name"
  value = google_compute_instance.myapp1.name
}

output "vm_machine_type" {
  description = "VM Machine Type"
  value = google_compute_instance.myapp1.machine_type
}
```

## Output Values

### Outputs:

```
vm_external_ip = "34.45.238.213"
vm_id = "projects/gcplearn9/zones/us-central1-a/instances/myapp1"
vm_instanceid = "3515727697506246169"
vm_machine_type = "e2-micro"
vm_name = "myapp1"
vm_selflink = "https://www.googleapis.com/compute/v1/projects/gcplearn9/zones/u
s-central1-a/instances/myapp1"
```

# What are we going to learn?

## ✓ terraform-manifests

\$ app1-webserver-install.sh

Y c1-versions.tf

Y c2-variables.tf

Y c3-vpc.tf

Y c4-firewallrules.tf

Y c5-vminstance.tf

Y c6-output-values.tf

Y terraform.tfvars

Y vm.auto.tfvars

Y vm.tfvars

Demo: Terraform Concept: Output Values

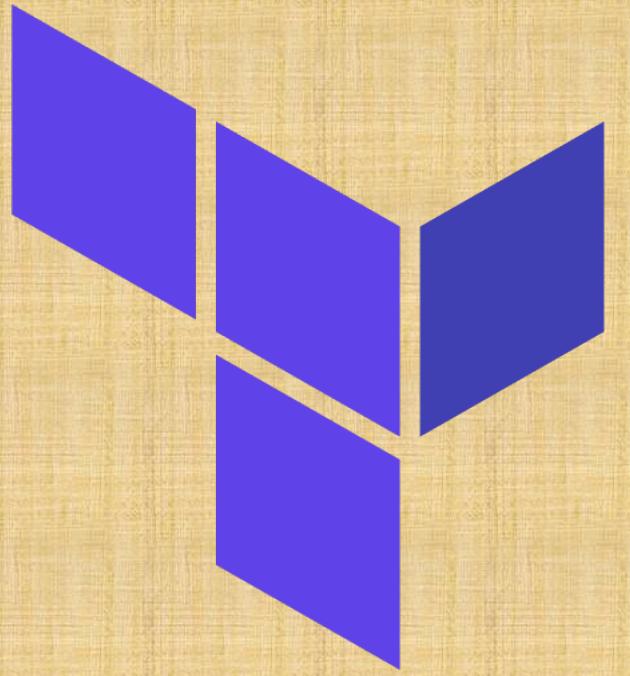
C6

Understand and define Terraform Output values



Demo-6

# Terraform Meta-Argument count



**Demo: Terraform Concept: Meta-argument count**

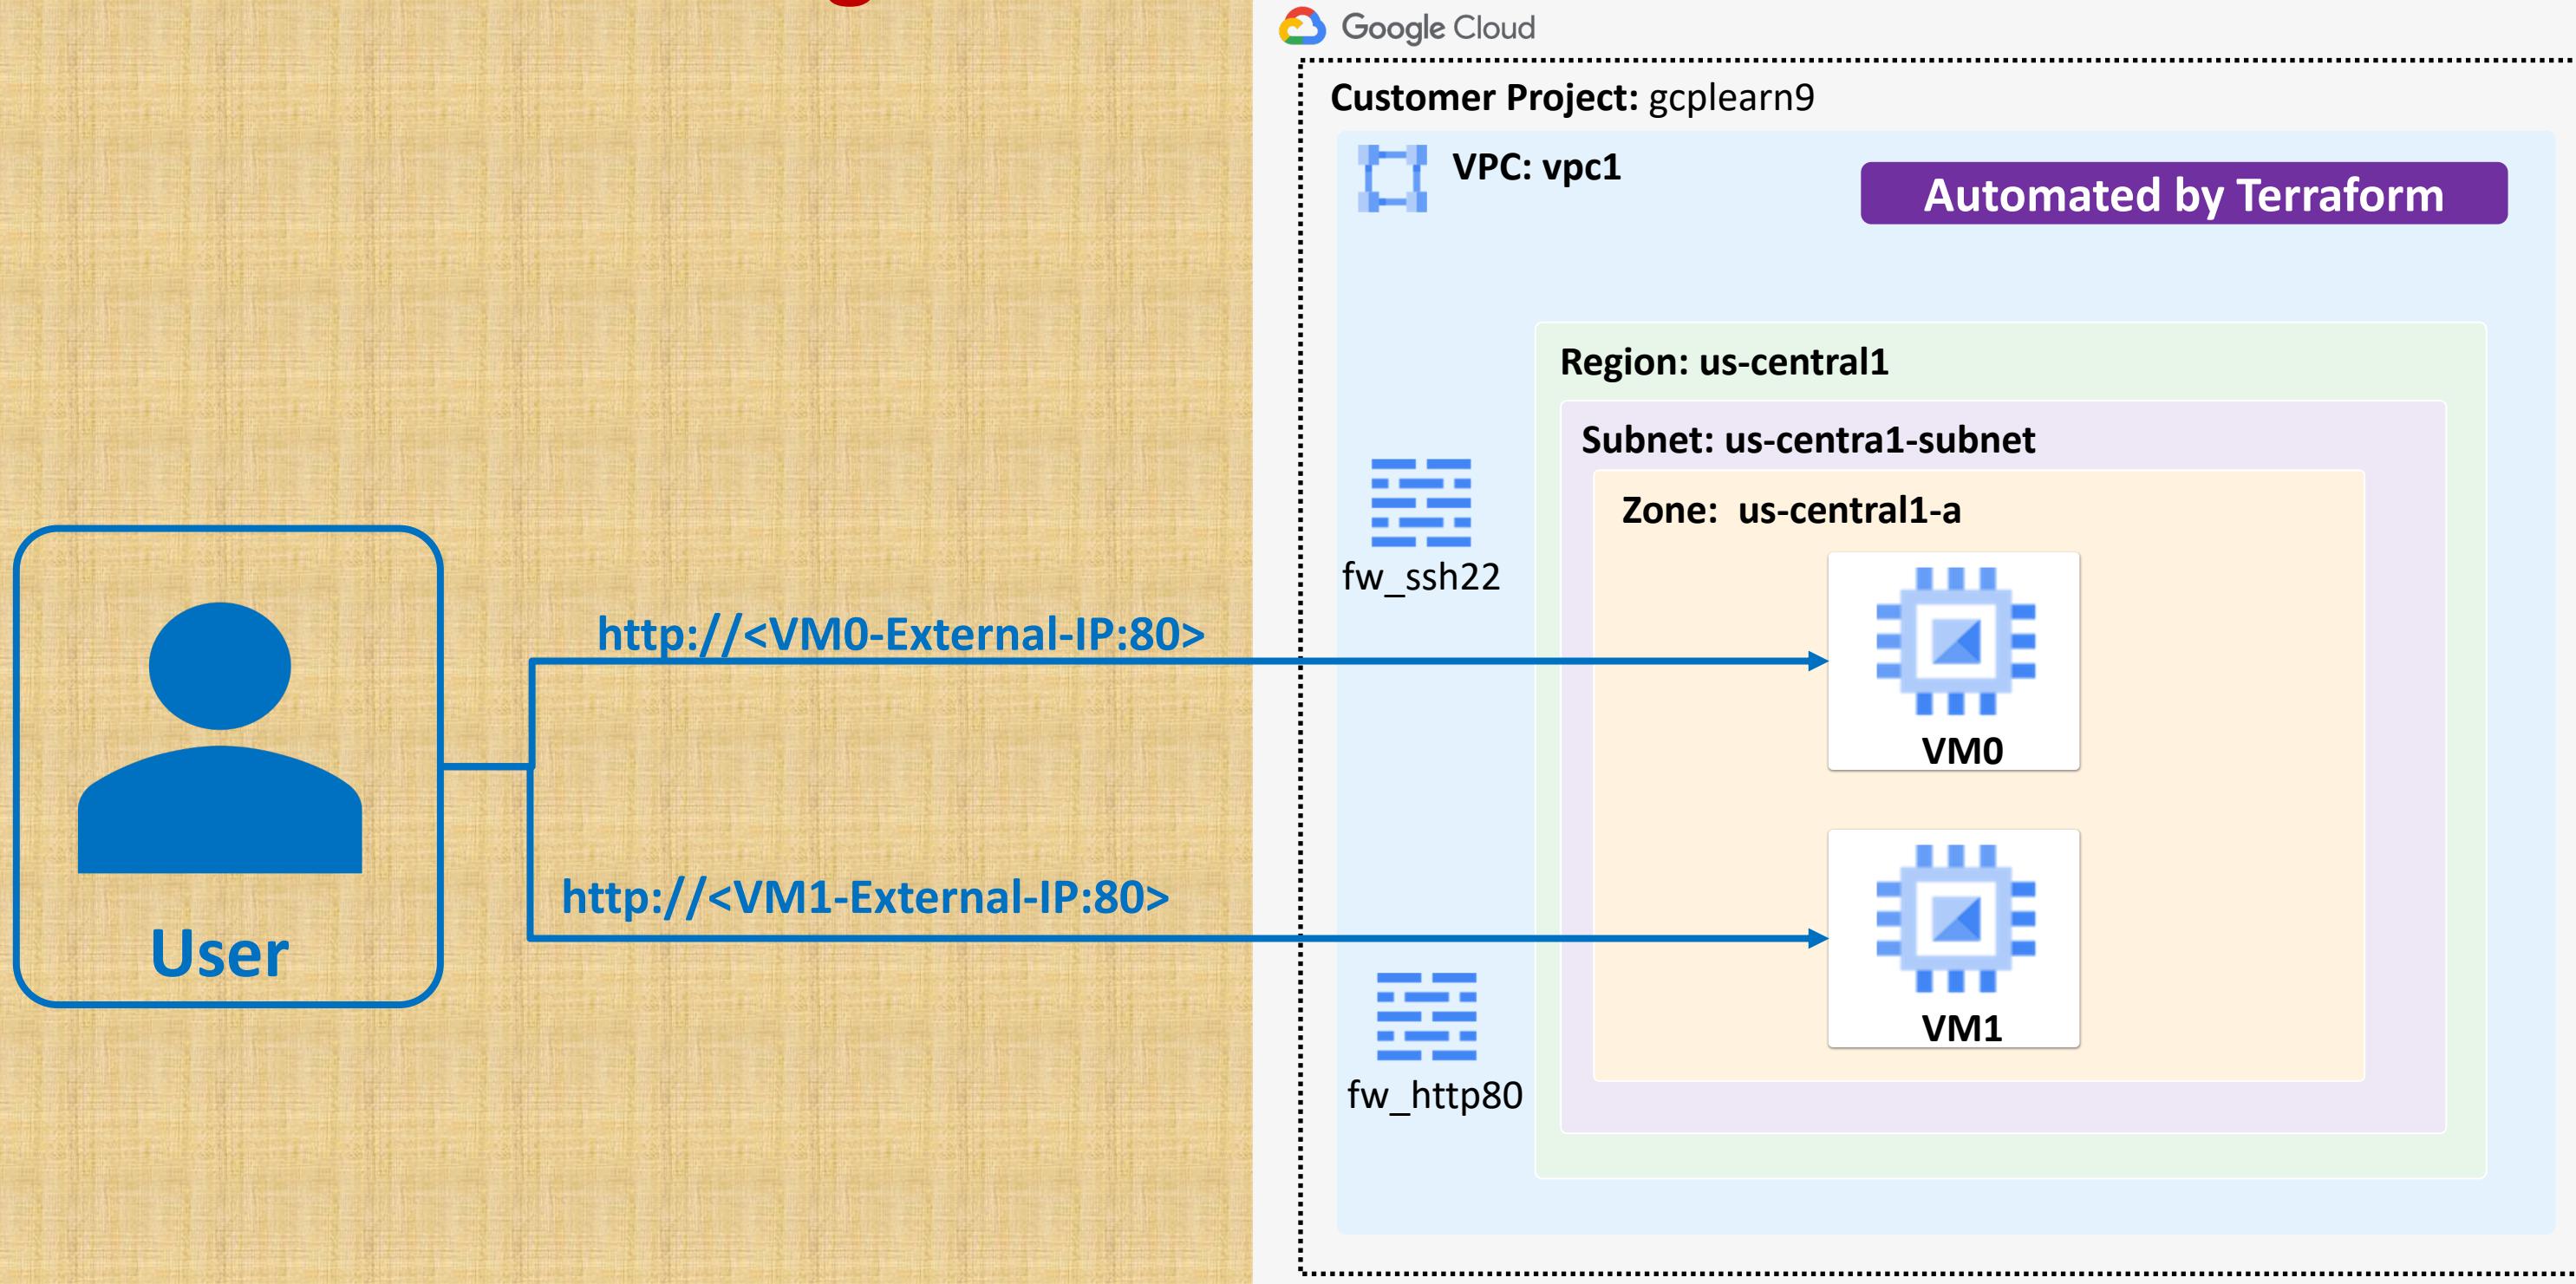
# Terraform Meta-Argument: count

- By default, resource block creates a **single instance** of a resource
  - **For example:** 1 VM Instance
- With the meta-argument **count** we can create **specified number of similar instances** of a resource
  - **For example:** `count = 2` (will create **2 VM instances** from the same defined resource block)
  - **count.index:** The distinct index number (**starting with 0**) corresponding to this instance.

## Resource Meta-Argument: count

```
# Resource Block: Create a Compute Engine instance
resource "google_compute_instance" "myapp1" {
    # Meta-Argument: count
    count = 2
    name      = "myapp1-vm-${count.index}"
    machine_type = var.machine_type
    zone       = "us-central1-a"
```

# Terraform Meta-argument: count



# Terraform: For expression

## For expression

- **For Expression:** Enables efficient management of multiple resources by [looping](#) through lists, maps, or sets
- **Examples:**
  - For loop with [list](#)
  - For loop with [map](#)
  - For loop with [map advanced](#)
  - Legacy [splat operator](#)
  - Latest generalized [splat operator](#)

```
# Output - For Loop with List
output "for_output_list" {
  description = "For Loop with List"
  value = [for instance in google_compute_instance.myapp1: instance.name]
}

# Output - For Loop with Map
output "for_output_map1" {
  description = "For Loop with Map"
  value = {for instance in google_compute_instance.myapp1: instance.name => instance.instance_id}
}

# Output - For Loop with Map Advanced
output "for_output_map2" {
  description = "For Loop with Map - Advanced"
  value = {for c, instance in google_compute_instance.myapp1: c => instance.name}
}
```

```
# Output Legacy Splat Operator (Legacy) - Returns the List
output "legacy_splat_instance" {
  description = "Legacy Splat Operator"
  value = google_compute_instance.myapp1.*.name
}

# Output Latest Generalized Splat Operator - Returns the List
output "latest_splat_instance" {
  description = "Generalized latest Splat Operator"
  value = google_compute_instance.myapp1[*].name
}
```

# What are we going to learn?

## ✓ **terraform-manifests**

\$ **app1-webserver-install.sh**

Y **c1-versions.tf**

Y **c2-variables.tf**

Y **c3-vpc.tf**

Y **c4-firewallrules.tf**

Y **c5-vminstance.tf**

Y **c6-output-values.tf**

Y **terraform.tfvars**

### Demo: Terraform Concept: Meta-argument count

NC

No Changes from C1 to C4

C5

Define VM Instance with Meta-argument count

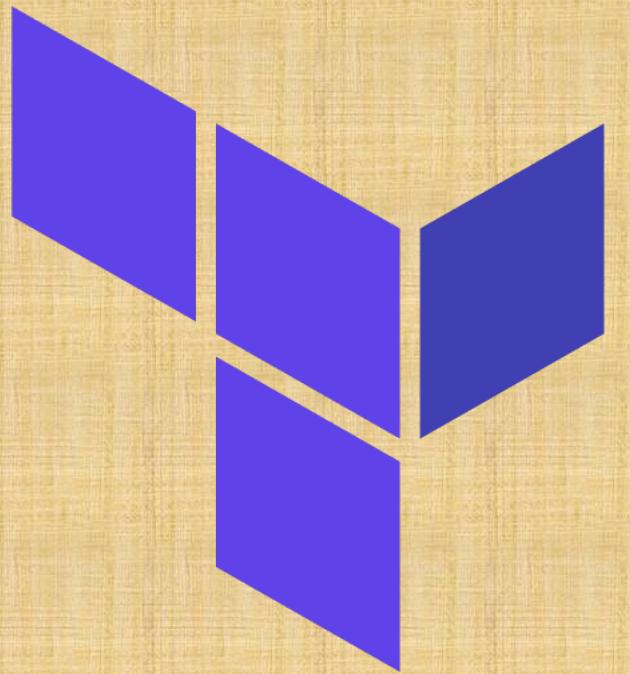
C6

Define Terraform output values with Terraform **for** expression and also learn about **Splat** operator



Demo-7

# Terraform Data sources



**Demo: Terraform Concept: Datasources**

# Terraform Data sources

- **Data sources:** allow Terraform to **use** information defined **outside** of Terraform
- A data source is accessed via a special kind of resource known as a **data** resource, declared using a **data block**:
- **Use case:**
  - **Step-01:** Get the **available compute zones** in that respective region using **google\_compute\_zone** datasource
  - **Step-02:** Create **two VM Instances** in two available compute zones using the **Meta-argument count**

## Terraform Data source

```
# Terraform Datasources
/* Datasource: Get a list of Google
Compute zones that are UP in a region */
data "google_compute_zones" "available" {
  status = "UP"
}

# Output value
output "compute_zones" {
  description = "List of compute zones"
  value = data.google_compute_zones.available.names
}
```

# What are we going to learn?

✓ **terraform-manifests**

\$ **app1-webserver-install.sh**

└ **c1-versions.tf**

└ **c2-variables.tf**

└ **c3-vpc.tf**

└ **c4-firewallrules.tf**

└ **c5-datasource.tf**

└ **c6-01-vminstance.tf**

└ **c6-02-vminstance-outputs.tf**

└ **terraform.tfvars**

**Demo: Terraform Concept: Datasources**

NC

C5

C6-01

No Changes from C1 to C4

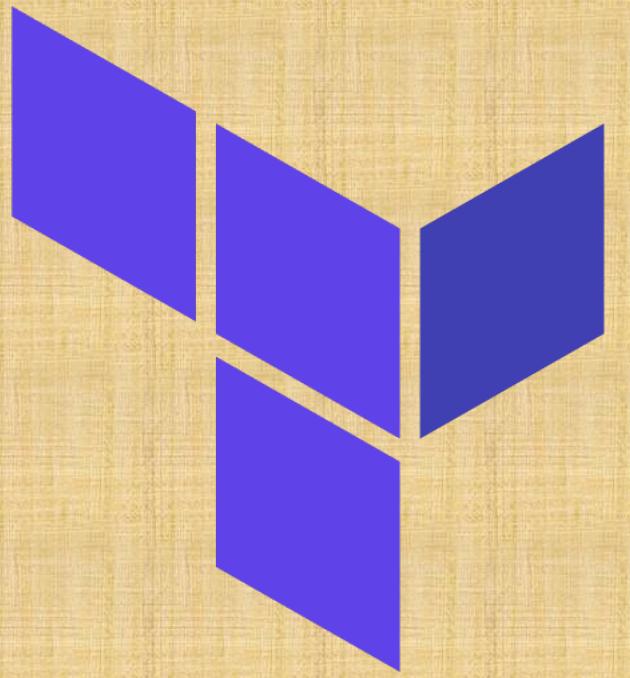
Define a data source to get a **list of all available zones** in that region

Update VM Instance with **zone value** updated from Terraform data source



Demo-8

# Terraform Meta-Argument `for_each`



**Demo: Terraform Concept: Meta-Argument `for_each`**

# Terraform Meta-Argument: `for_each`

- By default, resource block creates a **single instance** of a resource
  - **For example:** 1 VM Instance
- The **`for_each`** meta-argument accepts a **map or a set of strings** and **creates an instance** for each **item** in that **map or set**
  - **For example:** Creates 1 VM instance for each compute zone
- **Each instance has a**
  - **distinct** infrastructure object associated with it
  - **each** is **separately** created, updated, or destroyed when the configuration is applied

## Resource Meta-Argument: `for_each`

```
# Resource Block: Create a Compute Engine instance
resource "google_compute_instance" "myapp1" {
    # Meta-Argument: for_each
    for_each = toset(data.google_compute_zones.available.names)

    name      = "myapp1-vm-${each.key}"
    machine_type = var.machine_type
    zone      = each.key # You can also use each.value because for list
    tags       = [tolist(google_compute_firewall.fw_ssh.target_tags)[0],
```

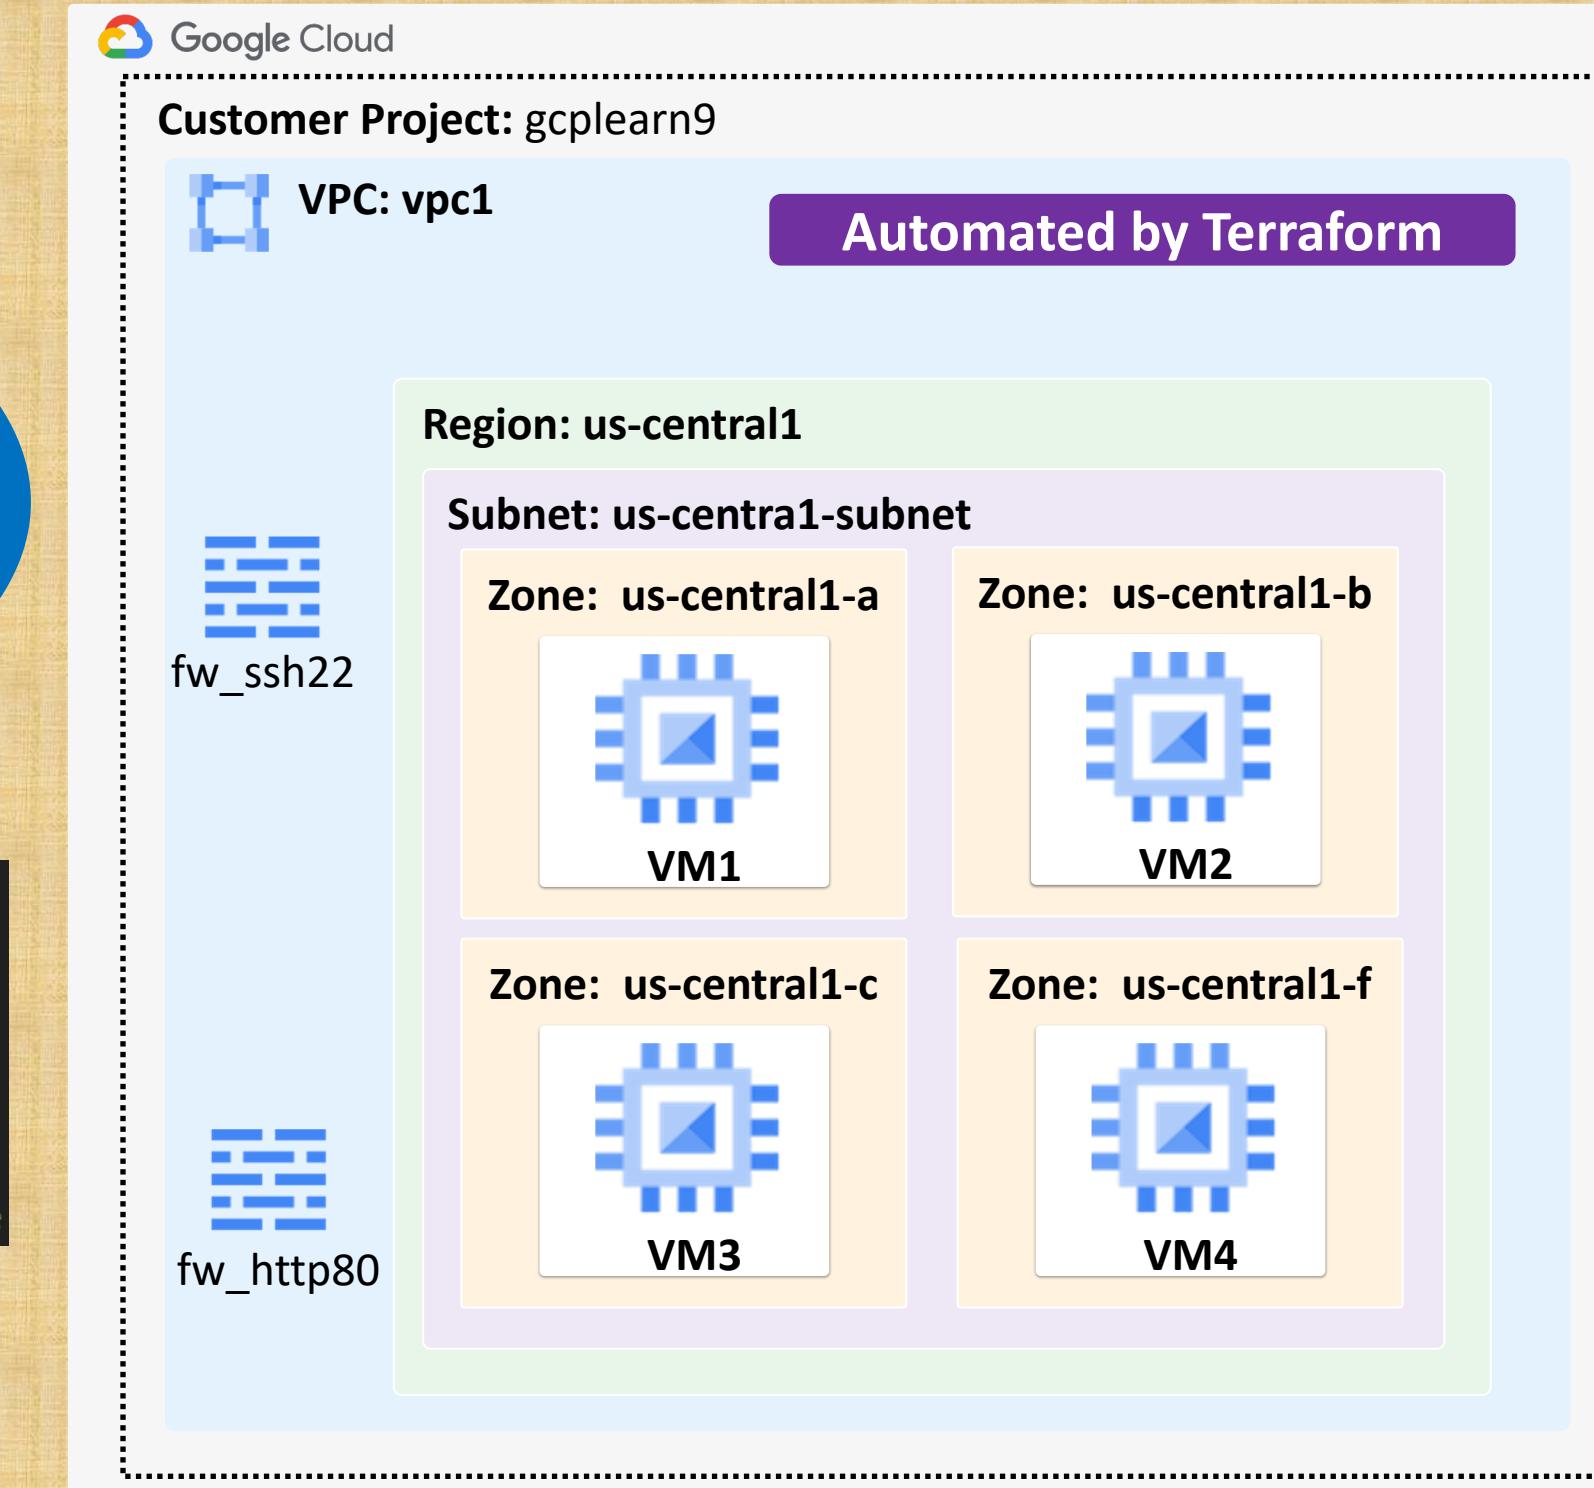
- **`each.key`:** The **map key** (or set member) corresponding to this instance.
- **`each.value`:** The **map value** corresponding to this instance (If a **set** was provided, **`each.key == each.value`**)

# Terraform Meta-argument: for each

## Demo-1: for\_each with Set values

Creates VM instance  
for each available  
compute zone in a  
region

```
# Resource Block: Create a Compute Engine instance
resource "google_compute_instance" "myapp1" {
  # Meta-Argument: for_each
  for_each = toset(data.google_compute_zones.available.names)
  name      = "myapp1-vm-${each.key}"
  machine_type = var.machine_type
  zone       = each.key # You can also use each.value because
```



# What are we going to learn?

```
✓ terraform-manifests
$ app1-webserver-install.sh
└ c1-versions.tf
└ c2-variables.tf
└ c3-vpc.tf
└ c4-firewallrules.tf
└ c5-datasource.tf
└ c6-01-vminstance.tf
└ c6-02-vminstance-outputs.tf
└ terraform.tfvars
```

## Demo-01: Terraform Concept: Meta-argument for\_each

NC

No Changes from C1 to C5

C6-01

VM Instance with **for\_each** Meta-argument with Set Values

C6-02

Terraform outputs for VM Instance

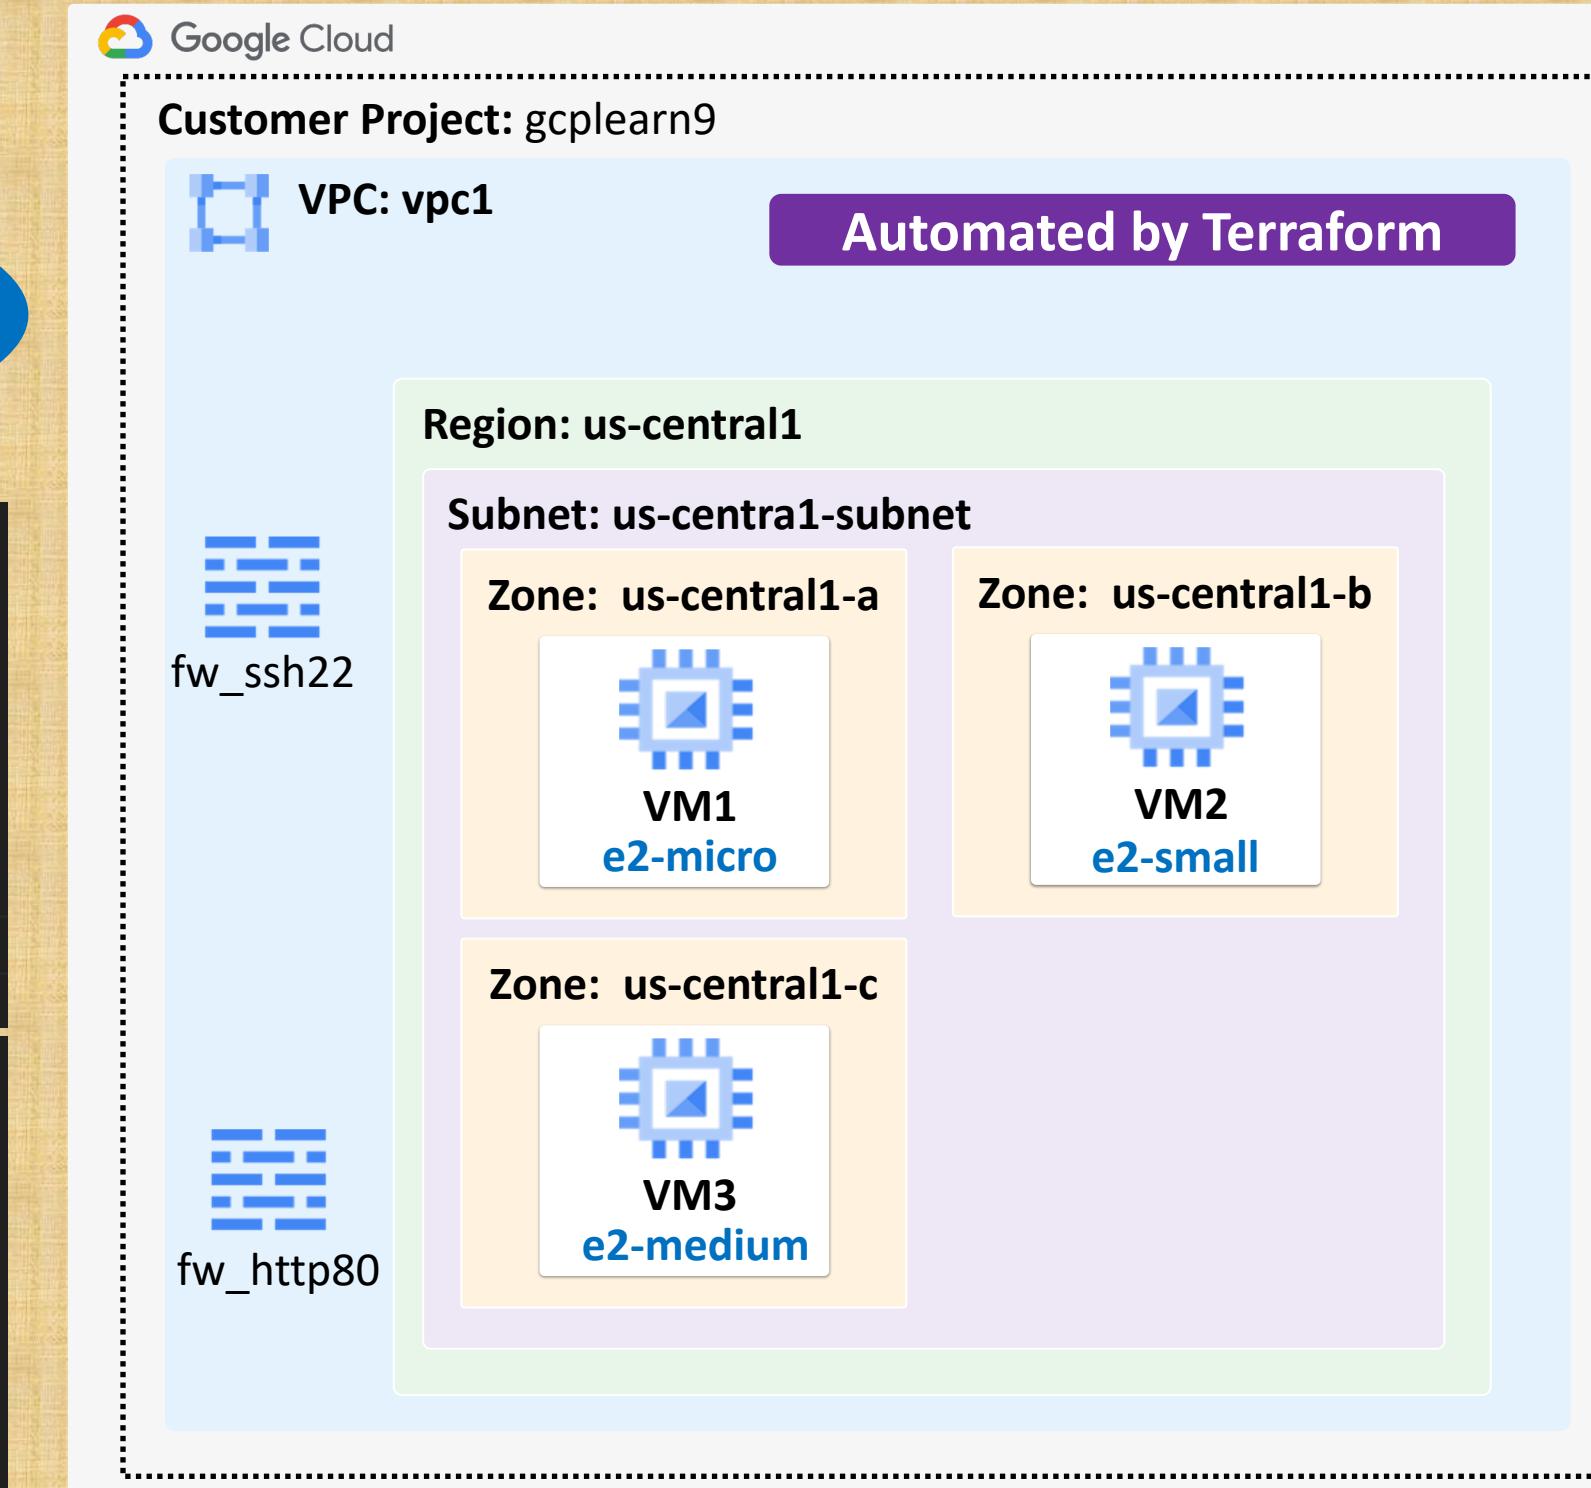
# Terraform Meta-argument: for each

## Demo-2: for\_each with Map values

Creates VM instance with different machine type in each compute zone

```
# Define a map with zone as key and machine_type as value
variable "zone_machine_map" {
  type = map(string)
  default = {
    "us-central1-a" = "e2-micro"
    "us-central1-b" = "e2-small"
    "us-central1-c" = "e2-medium"
  }
}

# Resource Block: Create a Compute Engine instance
resource "google_compute_instance" "myapp1" {
  # Meta-Argument: for_each
  for_each = var.zone_machine_map
  name      = "myapp1-vm-${each.key}"
  machine_type = each.value
  zone      = each.key
}
```



# What are we going to learn?

```
✓ terraform-manifests
$ app1-webserver-install.sh
└── c1-versions.tf
└── c2-variables.tf
└── c3-vpc.tf
└── c4-firewallrules.tf
└── c5-datasource.tf
└── c6-01-vminstance.tf
└── c6-02-vminstance-outputs.tf
└── terraform.tfvars
```

Demo-02: Terraform Concept: Meta-argument for\_each

NC

No Changes from C1 to C5

C6-01

VM Instance with for\_each Meta-argument with  
Map Values

Demo-9

Google

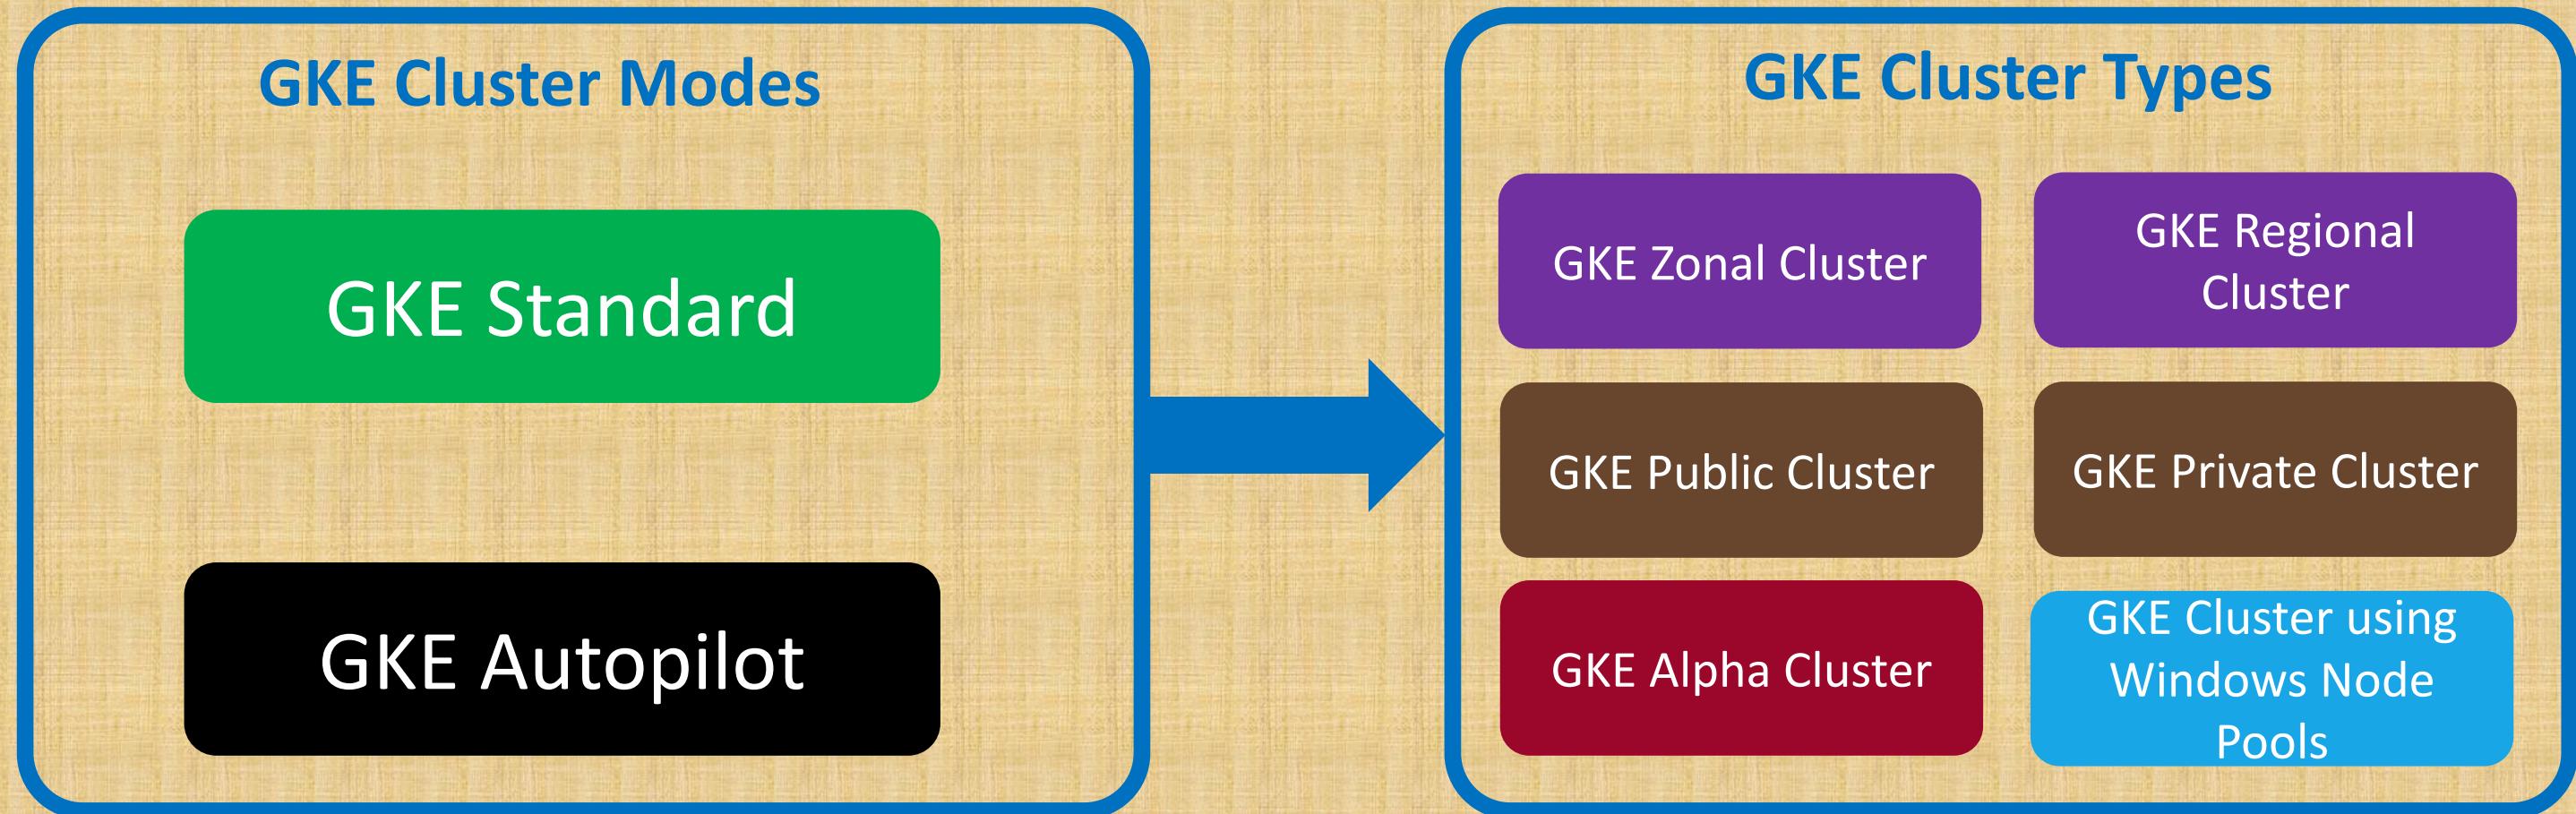
Kubernetes Engine

Create GKE

Standard Public Cluster

Terraform New Concept: Terraform Remote Backend

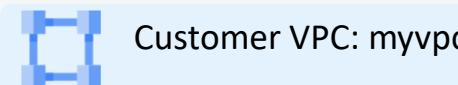
# GKE Cluster Modes & Types



# GKE Standard Public Cluster - Network Design

Customer Project: gcplearn9

Automated by Terraform



Region: us-central1



Subnet: 10.128.0.0/20

Zone: us-central1-a



Public IP

Private IP

Zone: us-central1-b



Public IP

Private IP

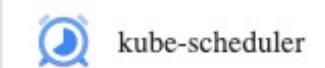
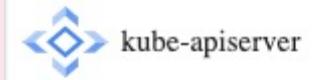
Google Managed Project



Region: us-central1



Kube API Server  
Public IP

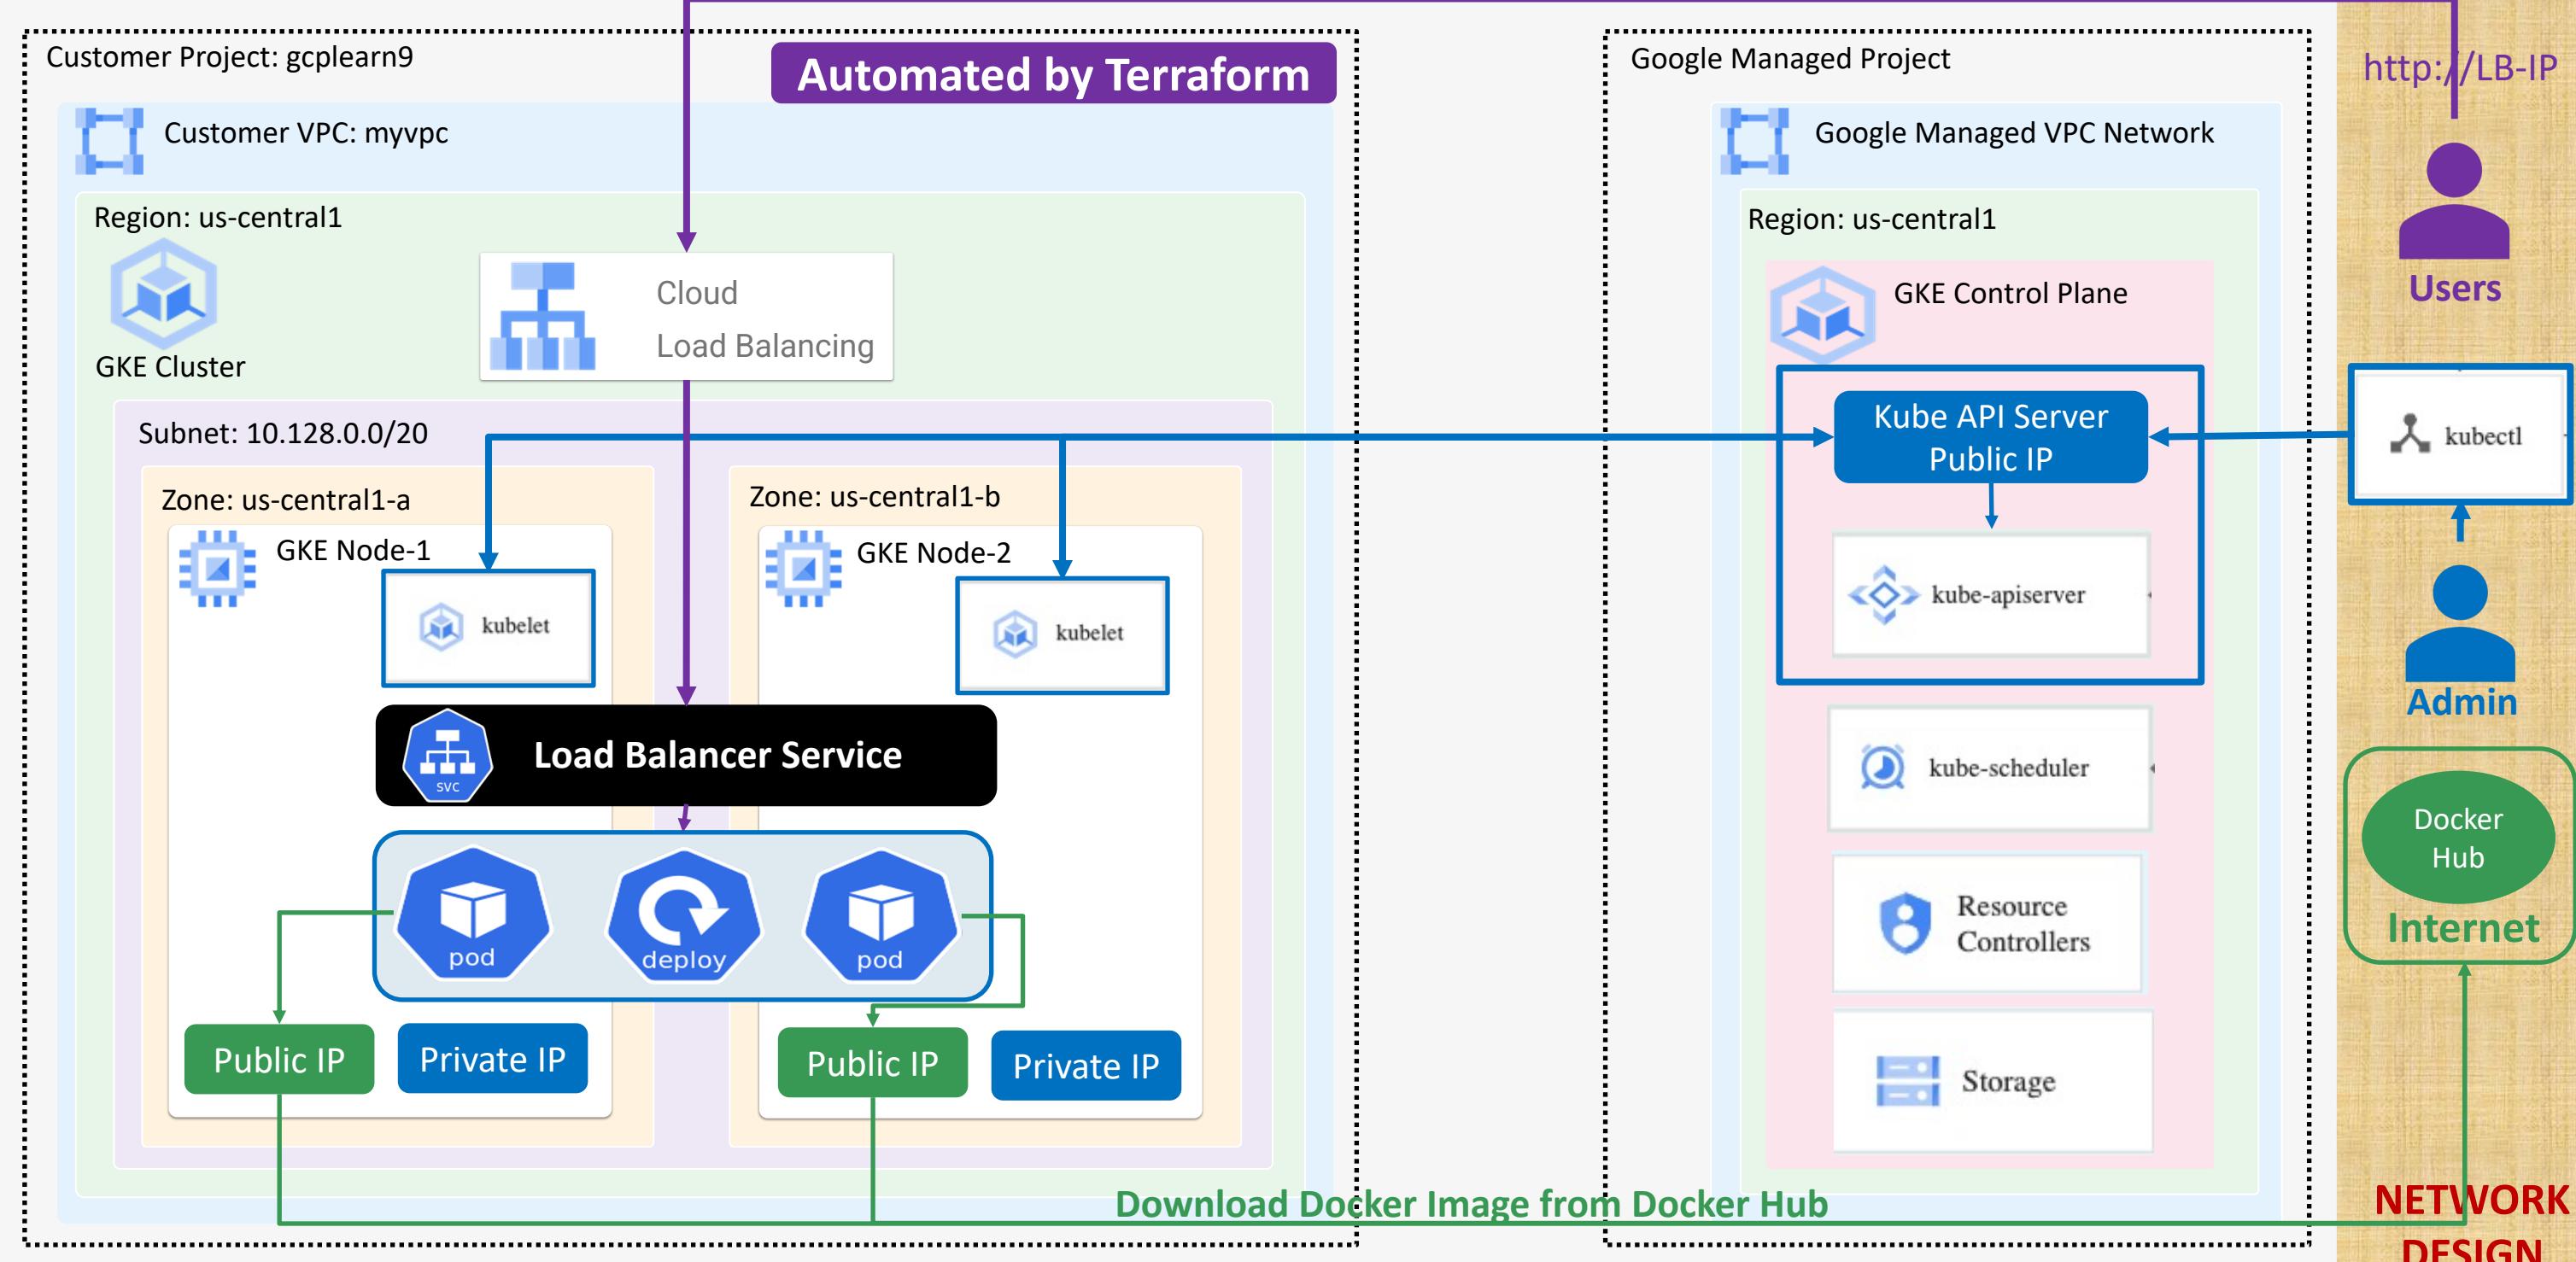


kubectl

Admin

**NETWORK  
DESIGN**

# GKE Standard Public Cluster - Network Design



# What is Terraform Backend ?

Backends are responsible for storing state and providing an API for state locking.

Local Backend

Terraform  
Local  
State  
Storage

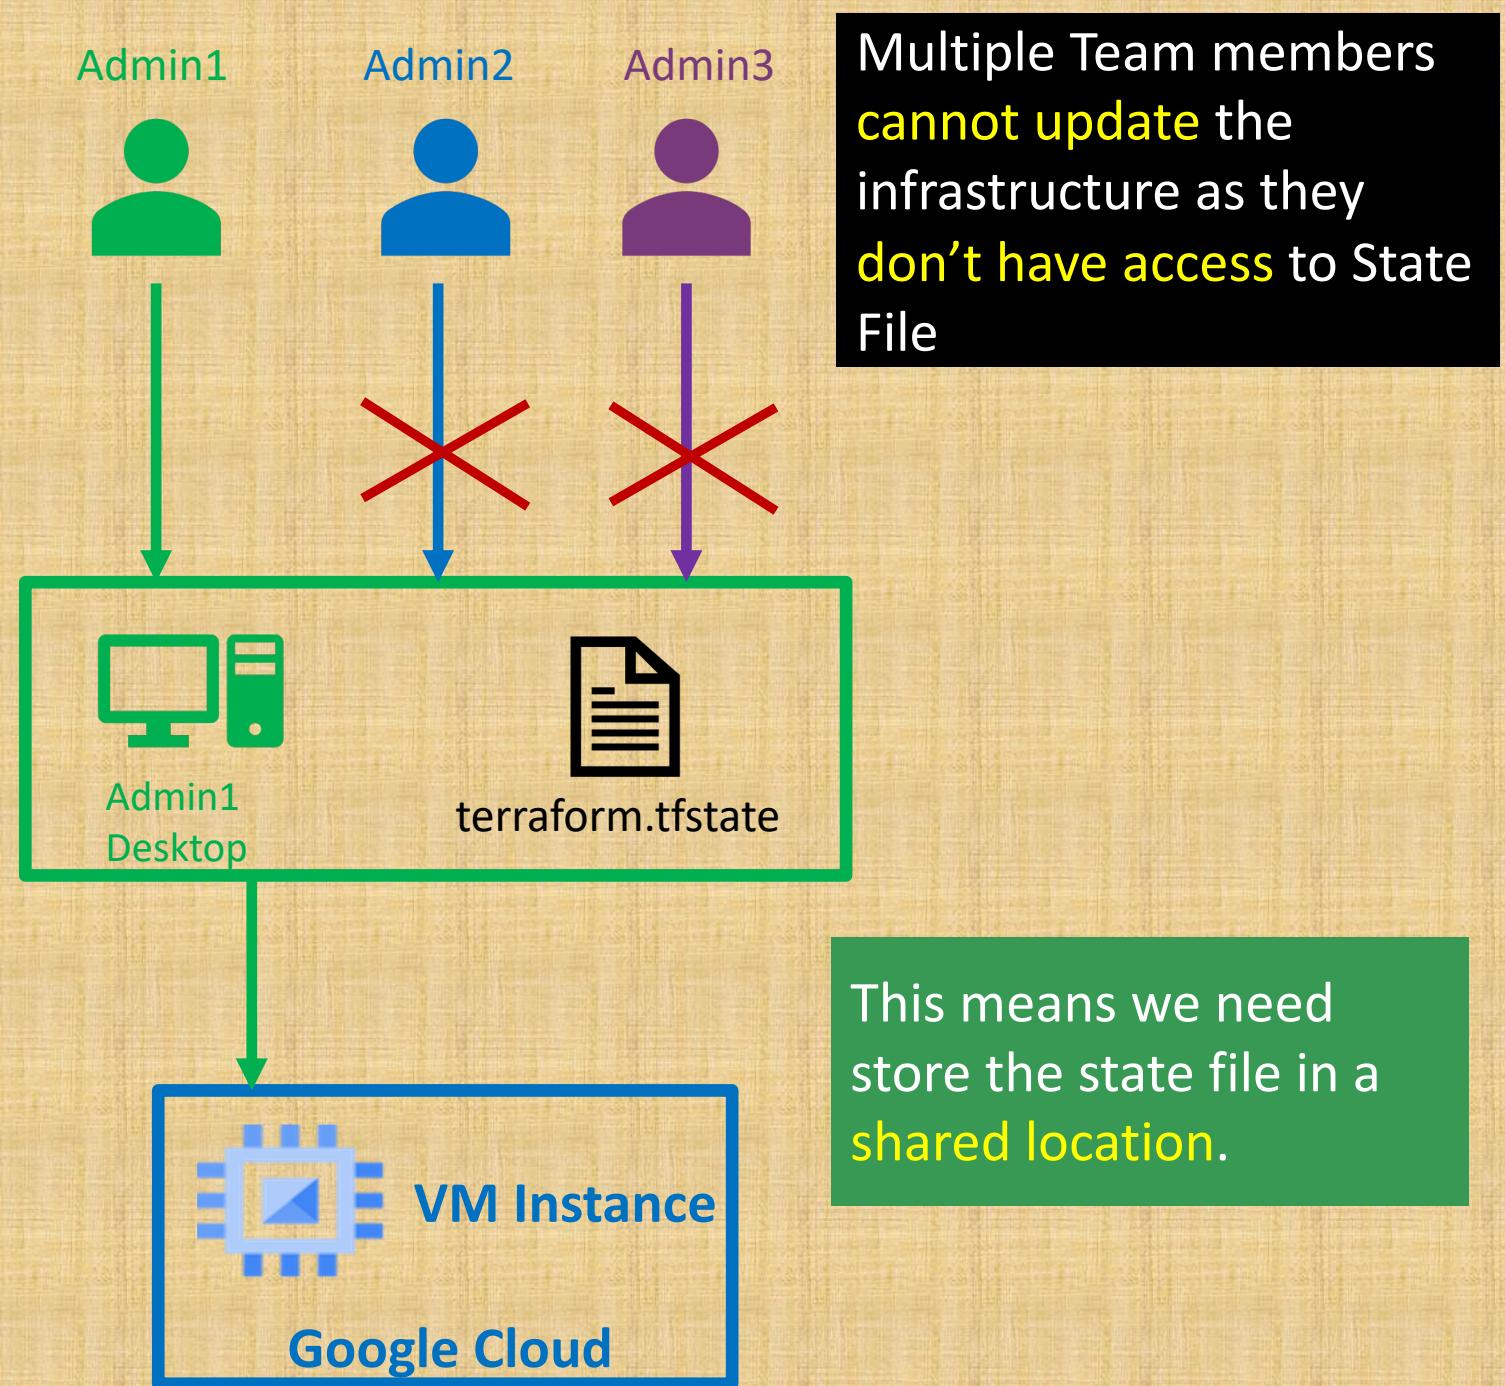
Remote Backend

Terraform  
Remote  
State  
Storage

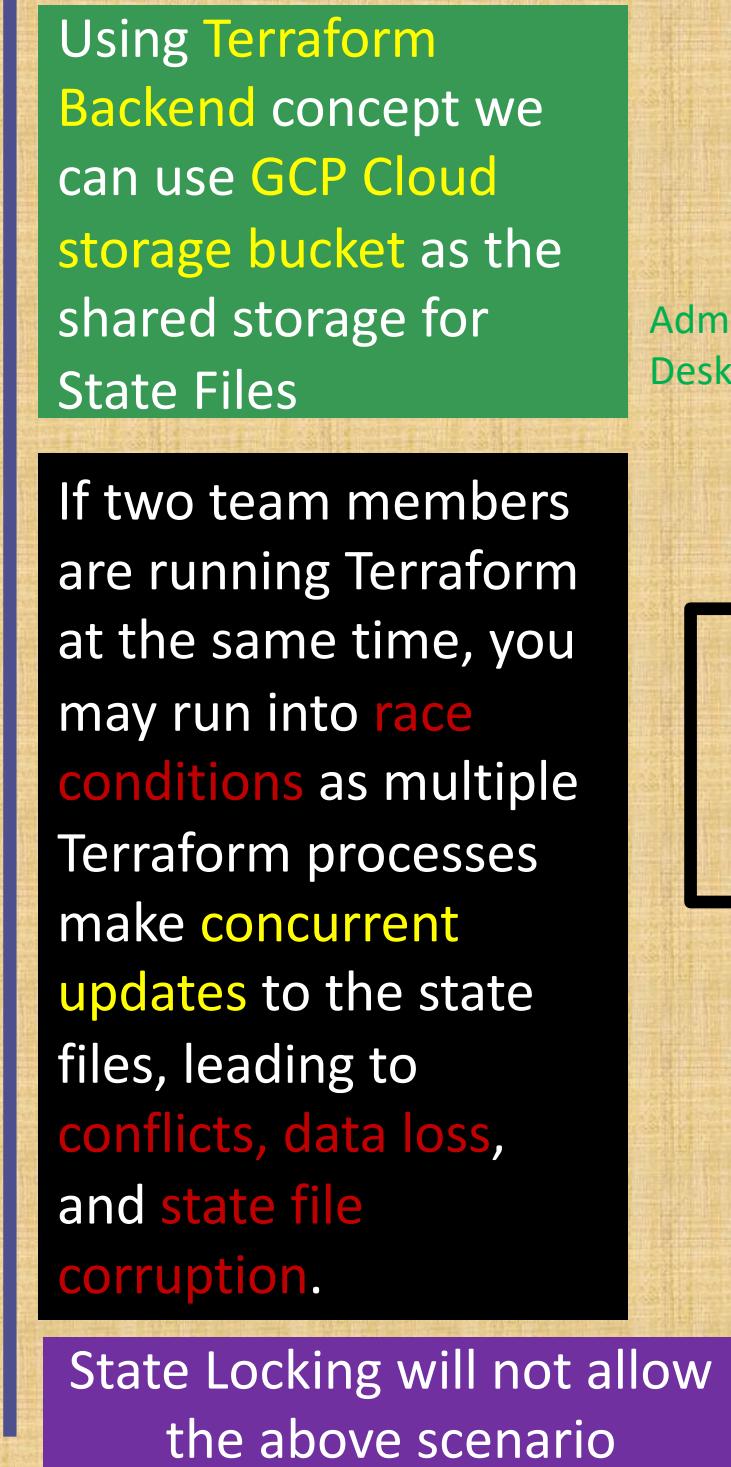
Inside Terraform Local Working  
Directory

Google Cloud Storage Buckets

# Local State File

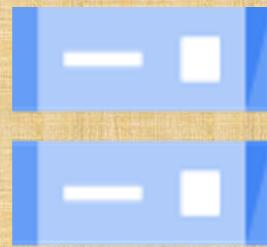


# Remote State File



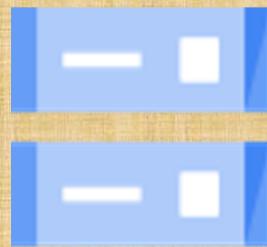
# What is Terraform State Storage and State Locking?

Terraform  
State Storage

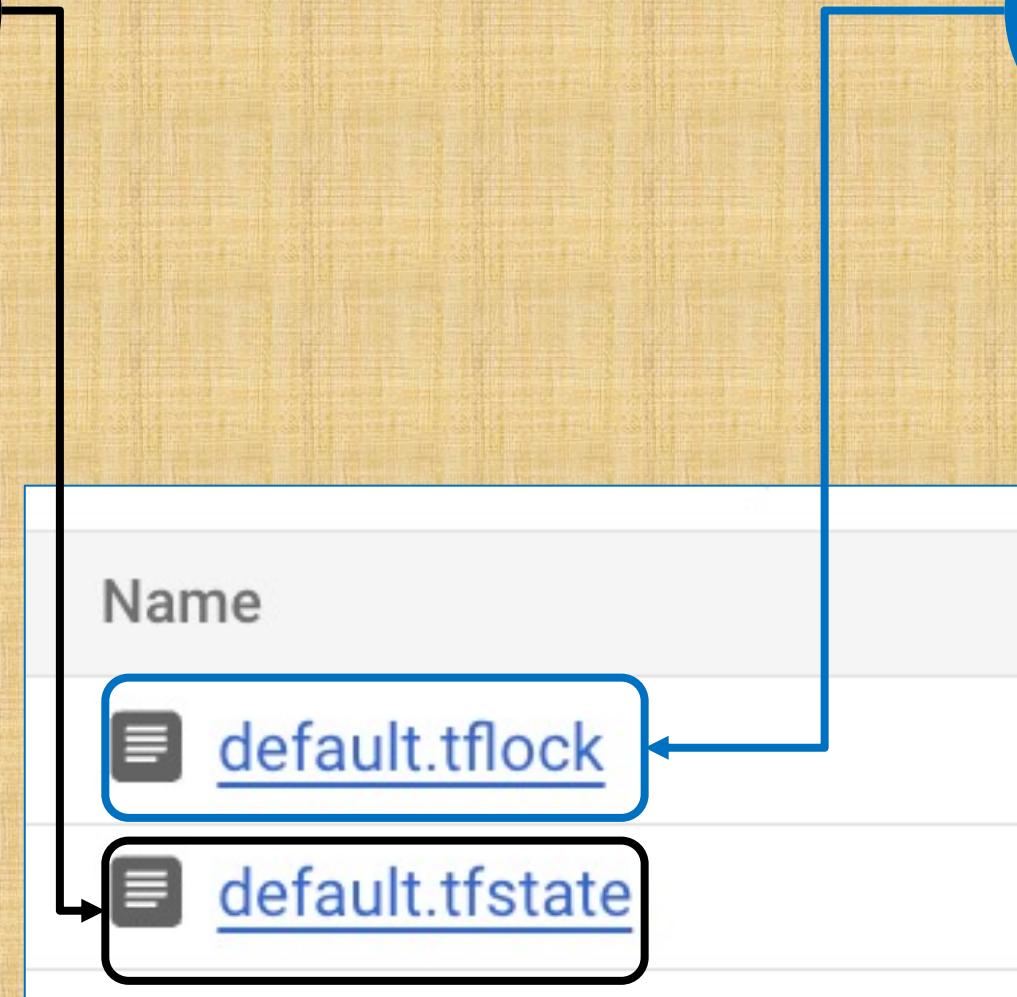


GCP Cloud Storage Bucket

Terraform  
State Locking



GCP Cloud Storage Bucket



# Terraform Remote Backend for GCP Terraform Resources

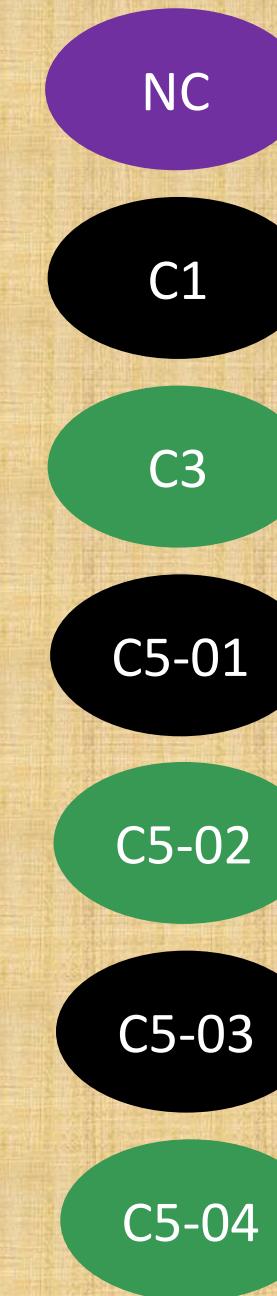
Terraform Remote Backend as GCP Cloud Storage Bucket

```
# Terraform Settings Block
terraform {
    required_version = ">= 1.9"
    required_providers {
        google = {
            source = "hashicorp/google"
            version = ">= 5.38.0"
        }
    }
    backend "gcs" {
        bucket = "terraform-on-gcp-gke"
        prefix = "dev/gke-cluster-public"
    }
}
```

# What are we going to learn?

Demo: GKE Standard Public Cluster + Terraform Remote Backend

- ✓ p1-gke-public-cluster
- ✗ c1-versions.tf
- ✗ c2-01-variables.tf
- ✗ c2-02-local-values.tf
- ✗ c3-vpc.tf
- ✗ c4-firewallrules.tf
- ✗ c5-01-gke-service-account.tf
- ✗ c5-02-gke-cluster.tf
- ✗ c5-03-gke-linux-nodepool.tf
- ✗ c5-04-gke-outputs.tf
- ✗ terraform.tfvars



No Changes from C2-01, C2-02, C4

Add Remote Backend as Cloud Storage Bucket

Standard VPC and Subnet, no major changes

IAM Service Account

GKE Cluster

GKE Linux Node Pool

GKE Outputs

Demo-10

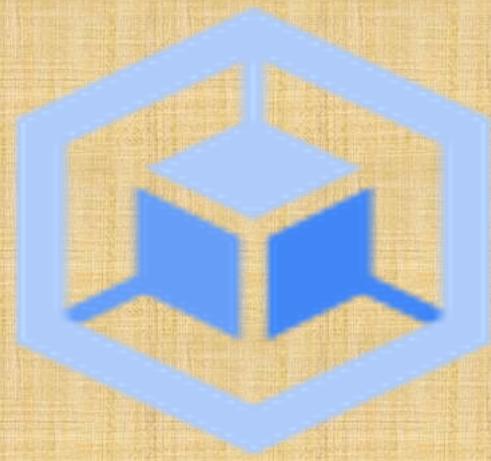
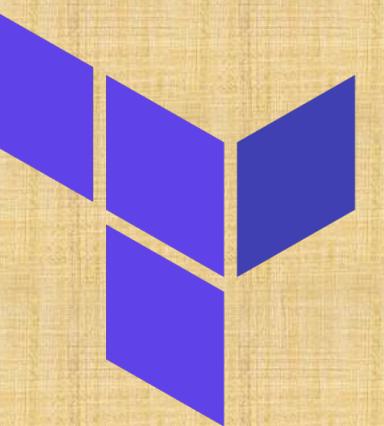
# Google Kubernetes Engine

## Kubernetes Deployment

## Kubernetes Load Balancer Service

YAML  
Manifests

Demo: Kubernetes YAML Manifests



# GKE Standard Public Cluster - Network Design

Customer Project: gcplearn9

Customer VPC: myvpc

Region: us-central1



GKE Cluster

Subnet: 10.128.0.0/20

Zone: us-central1-a



GKE Node-1



Load Balancer Service

Cloud Load Balancing

Zone: us-central1-b



GKE Node-2



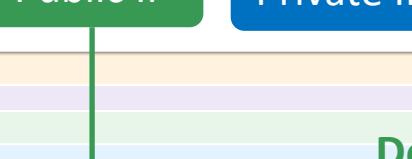
k8s  
YAML  
Manifests

Public IP

Private IP

Public IP

Private IP



Download Docker Image from Docker Hub

Automated by Terraform

Google Managed Project



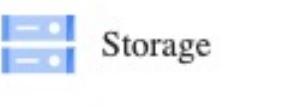
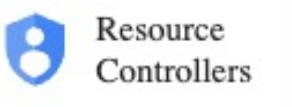
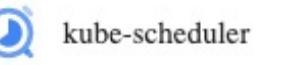
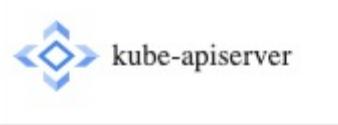
Google Managed VPC Network

Region: us-central1



GKE Control Plane

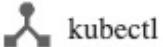
Kube API Server  
Public IP



http://LB-IP



Users



kubectl



Admin



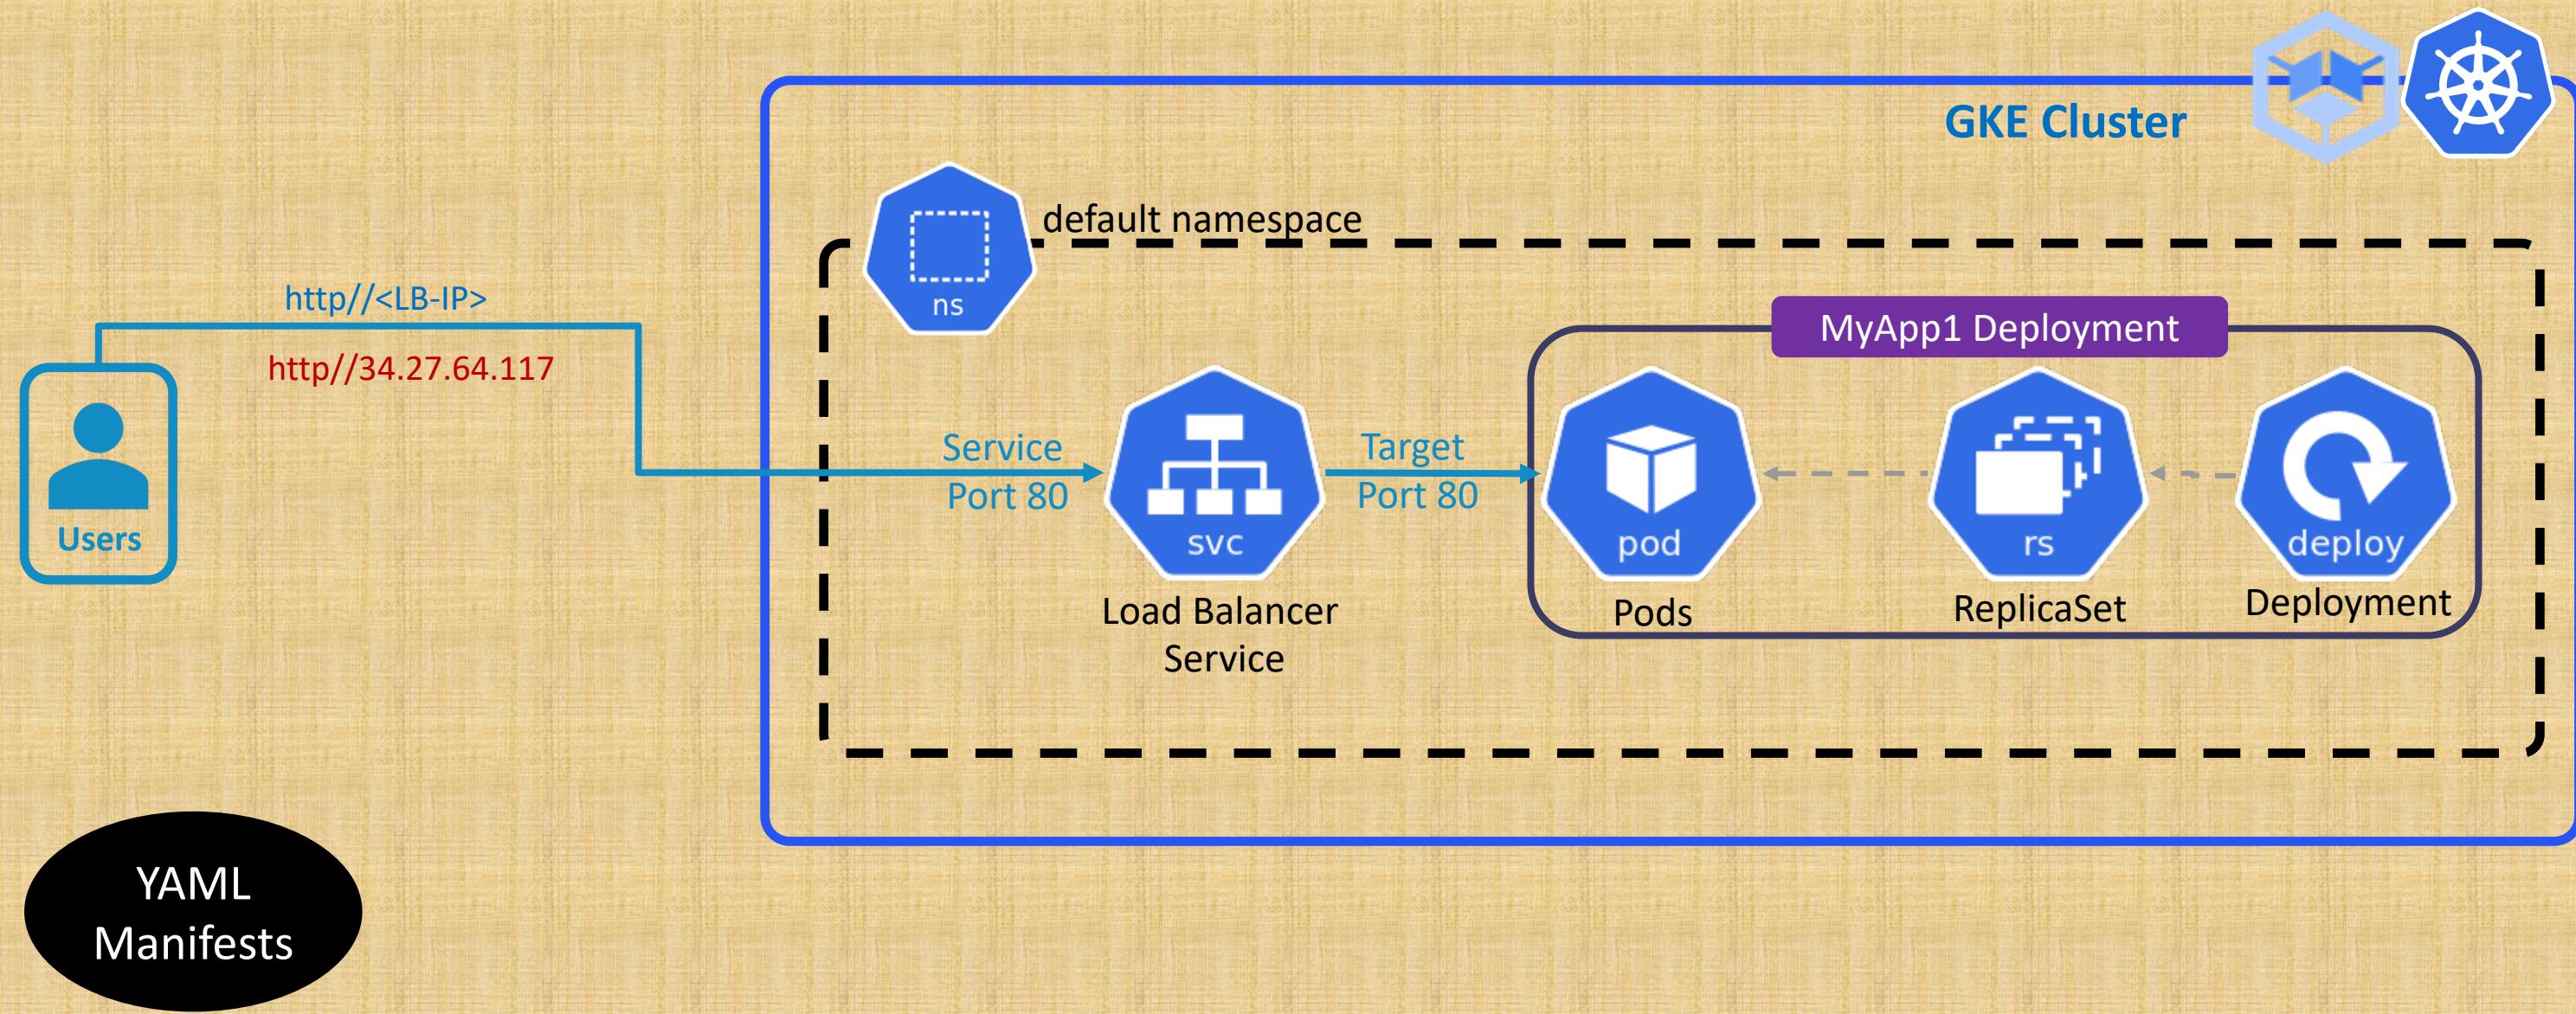
Docker Hub



Internet

**NETWORK  
DESIGN**

# GKE- Kubernetes Load Balancer Service



# What are we going to learn?

Demo: Deploy Kubernetes Deployment and Service in GKE Standard Public Cluster

- ✓ kube-manifests-yaml
- ! 01-kubernetes-deployment.yaml
- ! 02-kubernetes-loadbalancer-service.yaml

01

Kubernetes Deployment

02

Kubernetes Load Balancer Service

Demo-11

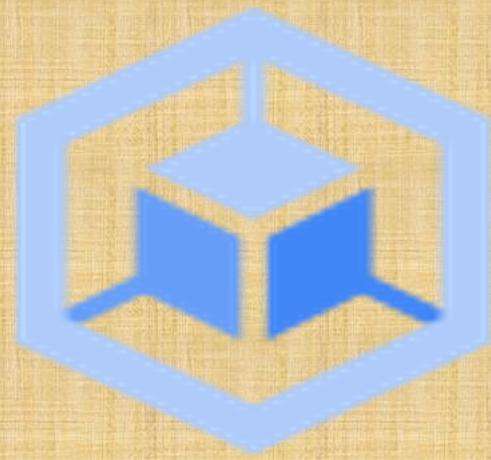
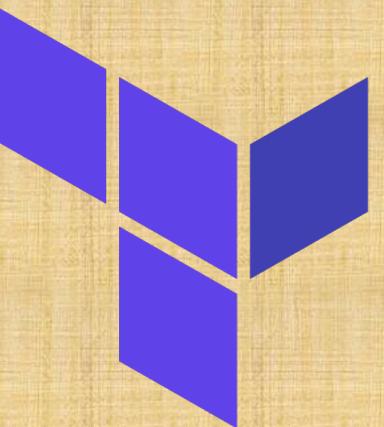
# Google Kubernetes Engine

## Kubernetes Deployment

## Kubernetes Load Balancer Service

**Demo: Kubernetes Deployment and Service using  
Terraform + Terraform Remote State Data source  
concept**

k8s  
Terraform  
Manifests



# GKE Standard Public Cluster - Network Design

Customer Project: gcplearn9

Automated by Terraform

Customer VPC: default

Region: us-central1



GKE Cluster

Subnet: 10.128.0.0/20

Zone: us-central1-a



GKE Node-1



Load Balancer Service

Cloud Load Balancing

Zone: us-central1-b



GKE Node-2

k8s  
Terraform  
Manifests

Public IP

Private IP

Public IP

Private IP

Download Docker Image from Docker Hub

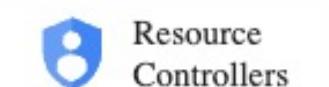
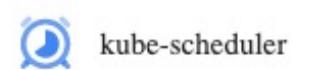
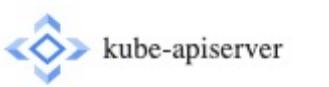
Google Managed Project

Google Managed VPC Network

Region: us-central1



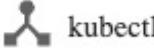
GKE Control Plane

Kube API Server  
Public IP

http://LB-IP



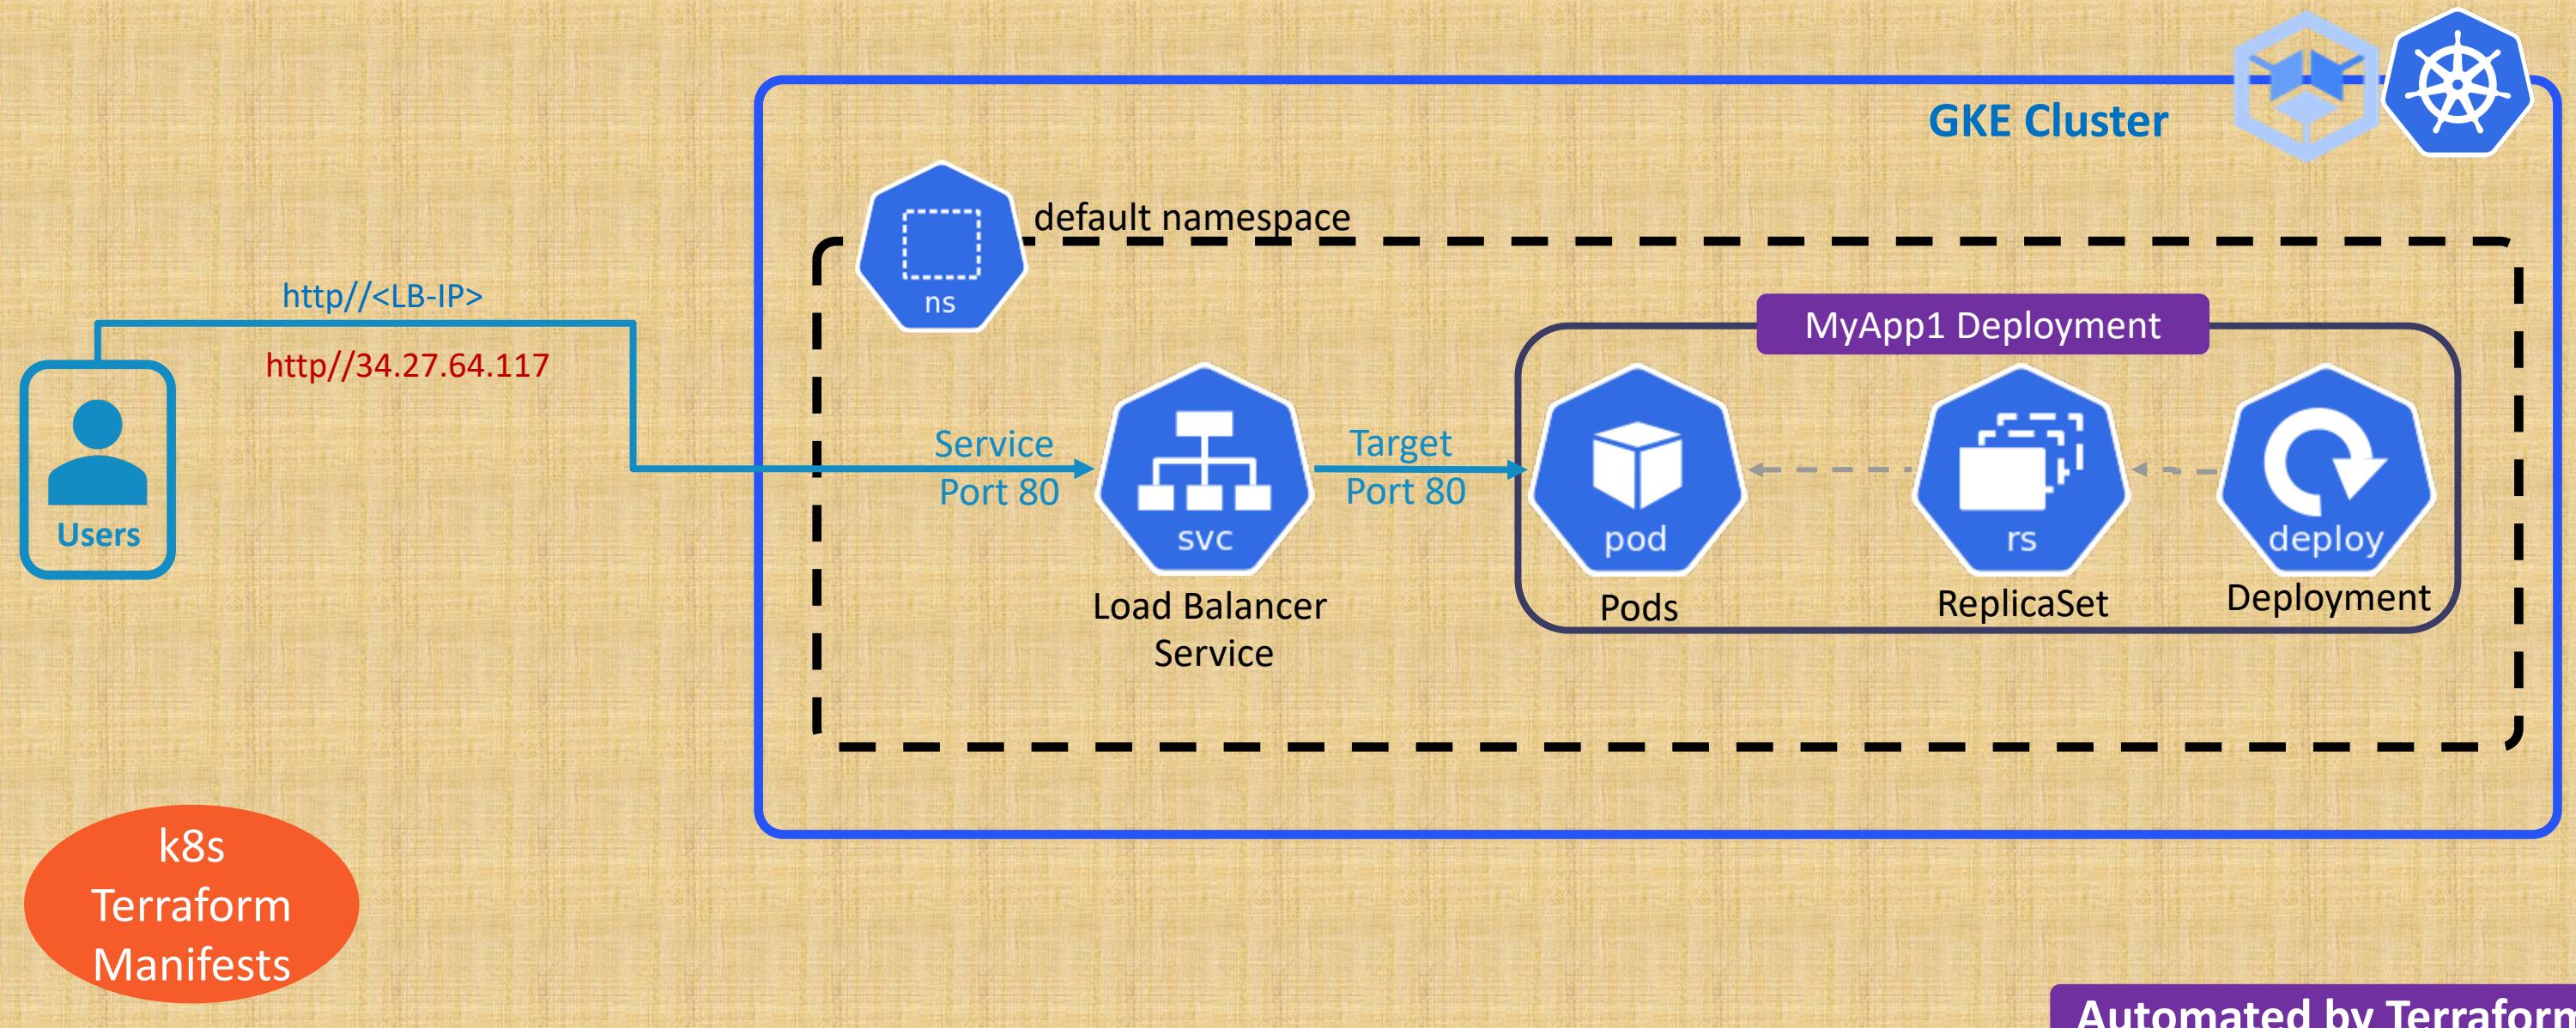
Users



Internet

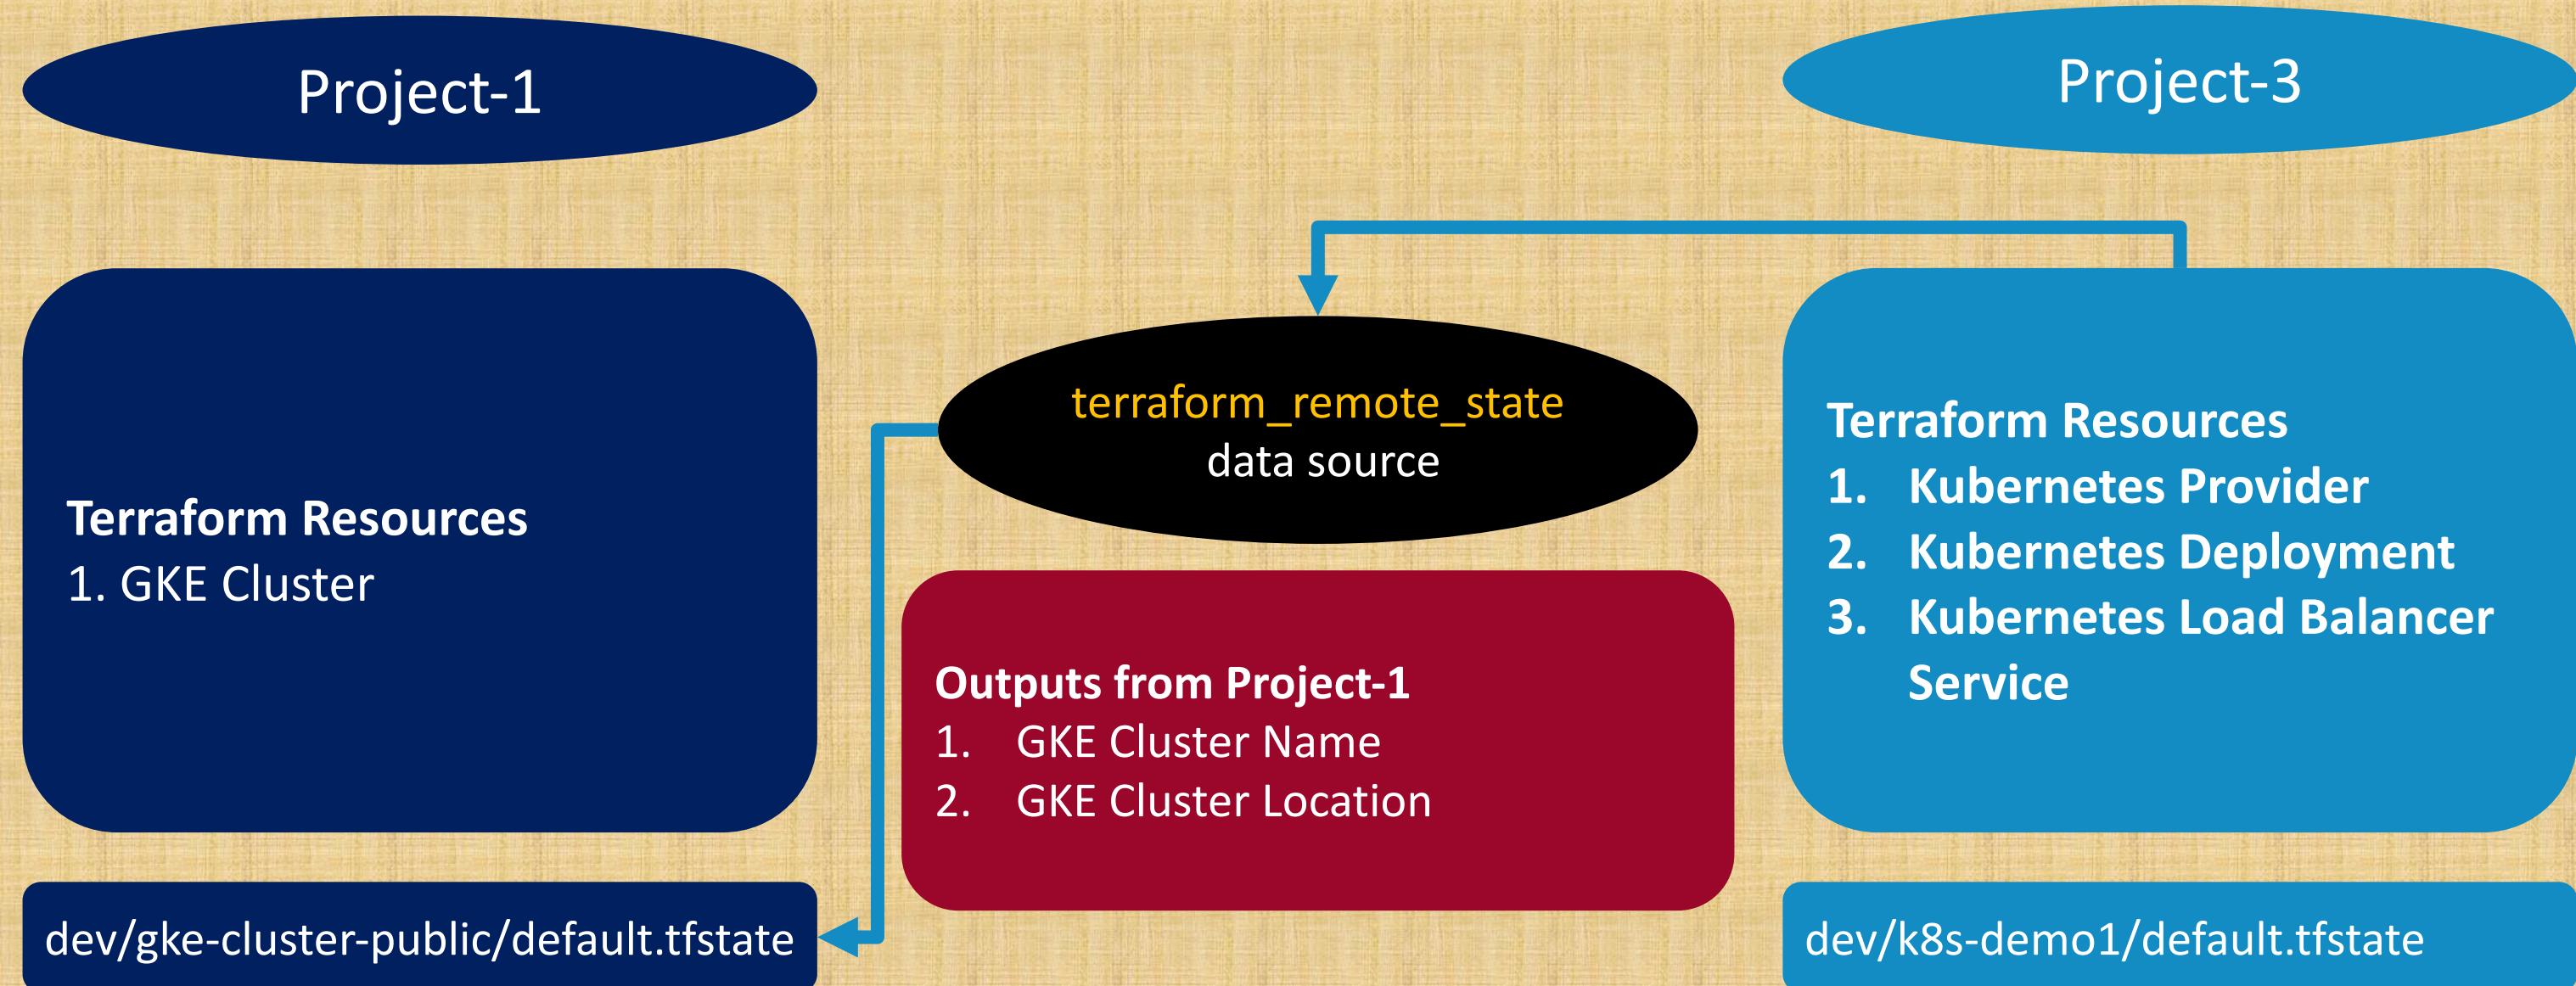
NETWORK  
DESIGN

# GKE- Kubernetes Load Balancer Service



# Terraform Remote State Datasource

The `terraform_remote_state` data source retrieves the root module output values from project-1 Terraform configuration, using the latest state snapshot from the remote backend.



# What are we going to learn?

**Demo: Kubernetes Deployment and Service using Terraform + Terraform Remote State Data source concept**

- ✓ p3-k8sresources-terraform-manifests
  - └ c1-versions.tf
  - └ c2-01-variables.tf
  - └ c2-02-local-values.tf
  - └ c3-01-remote-state-datasource.tf
  - └ c3-02-providers.tf
  - └ c4-kubernetes-deployment.tf
  - └ c5-kubernetes-loadbalancer-service.tf
  - └ terraform.tfvars

- C1
- C2-01
- C2-02
- C3-01
- C3-02
- C4
- C5

Add **Remote Backend** as Cloud Storage Bucket

Standard **Input Variables** as project-1

Standard **Local values** as project-1

Terraform **Remote State Data source** concept to access outputs from Project-1

**Kubernetes Provider** to connect to GKE Cluster created using Project-1

**Kubernetes Deployment** using Terraform

**Kubernetes Load Balancer Service** using Terraform

Demo-12

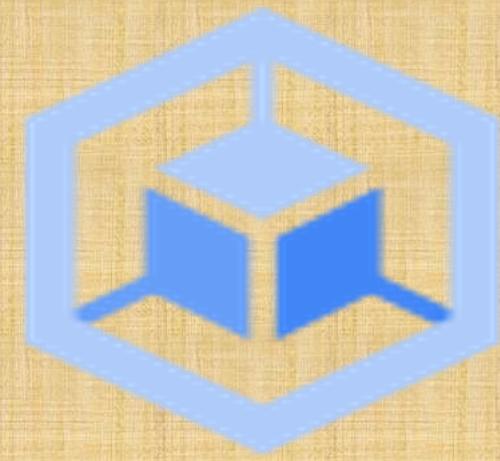
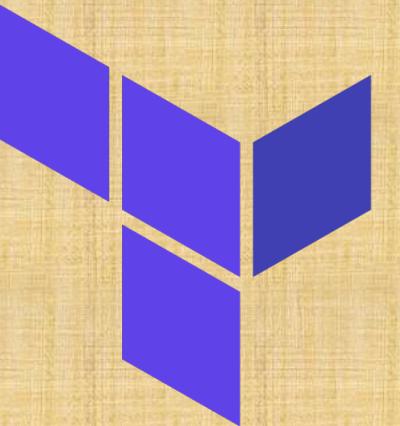


# Google Kubernetes Engine

Create GKE

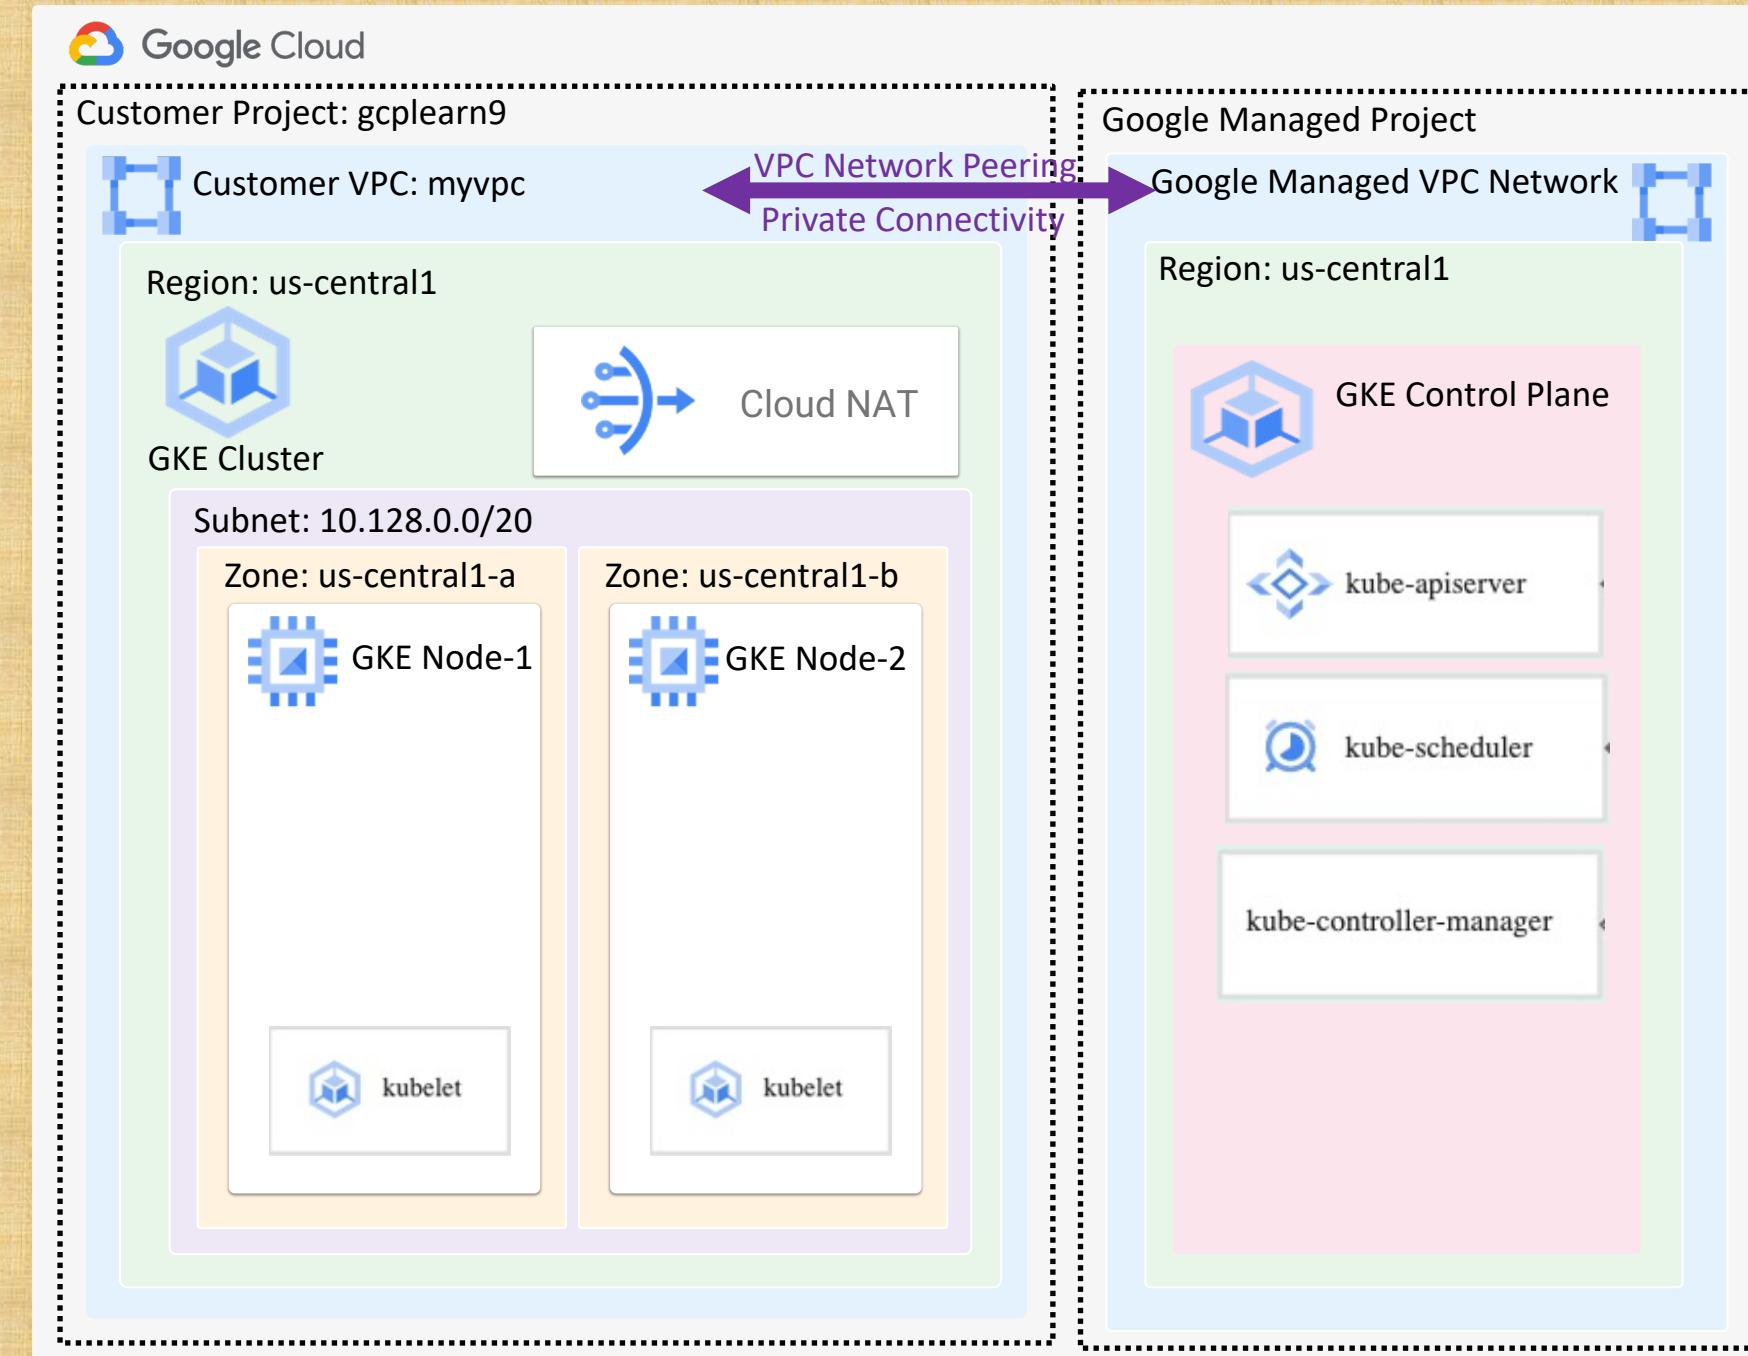
## Standard Private Cluster

**Demo:** GKE Standard Private Cluster with  
Public Kubernetes API Server Endpoint + Cluster  
Autoscaler



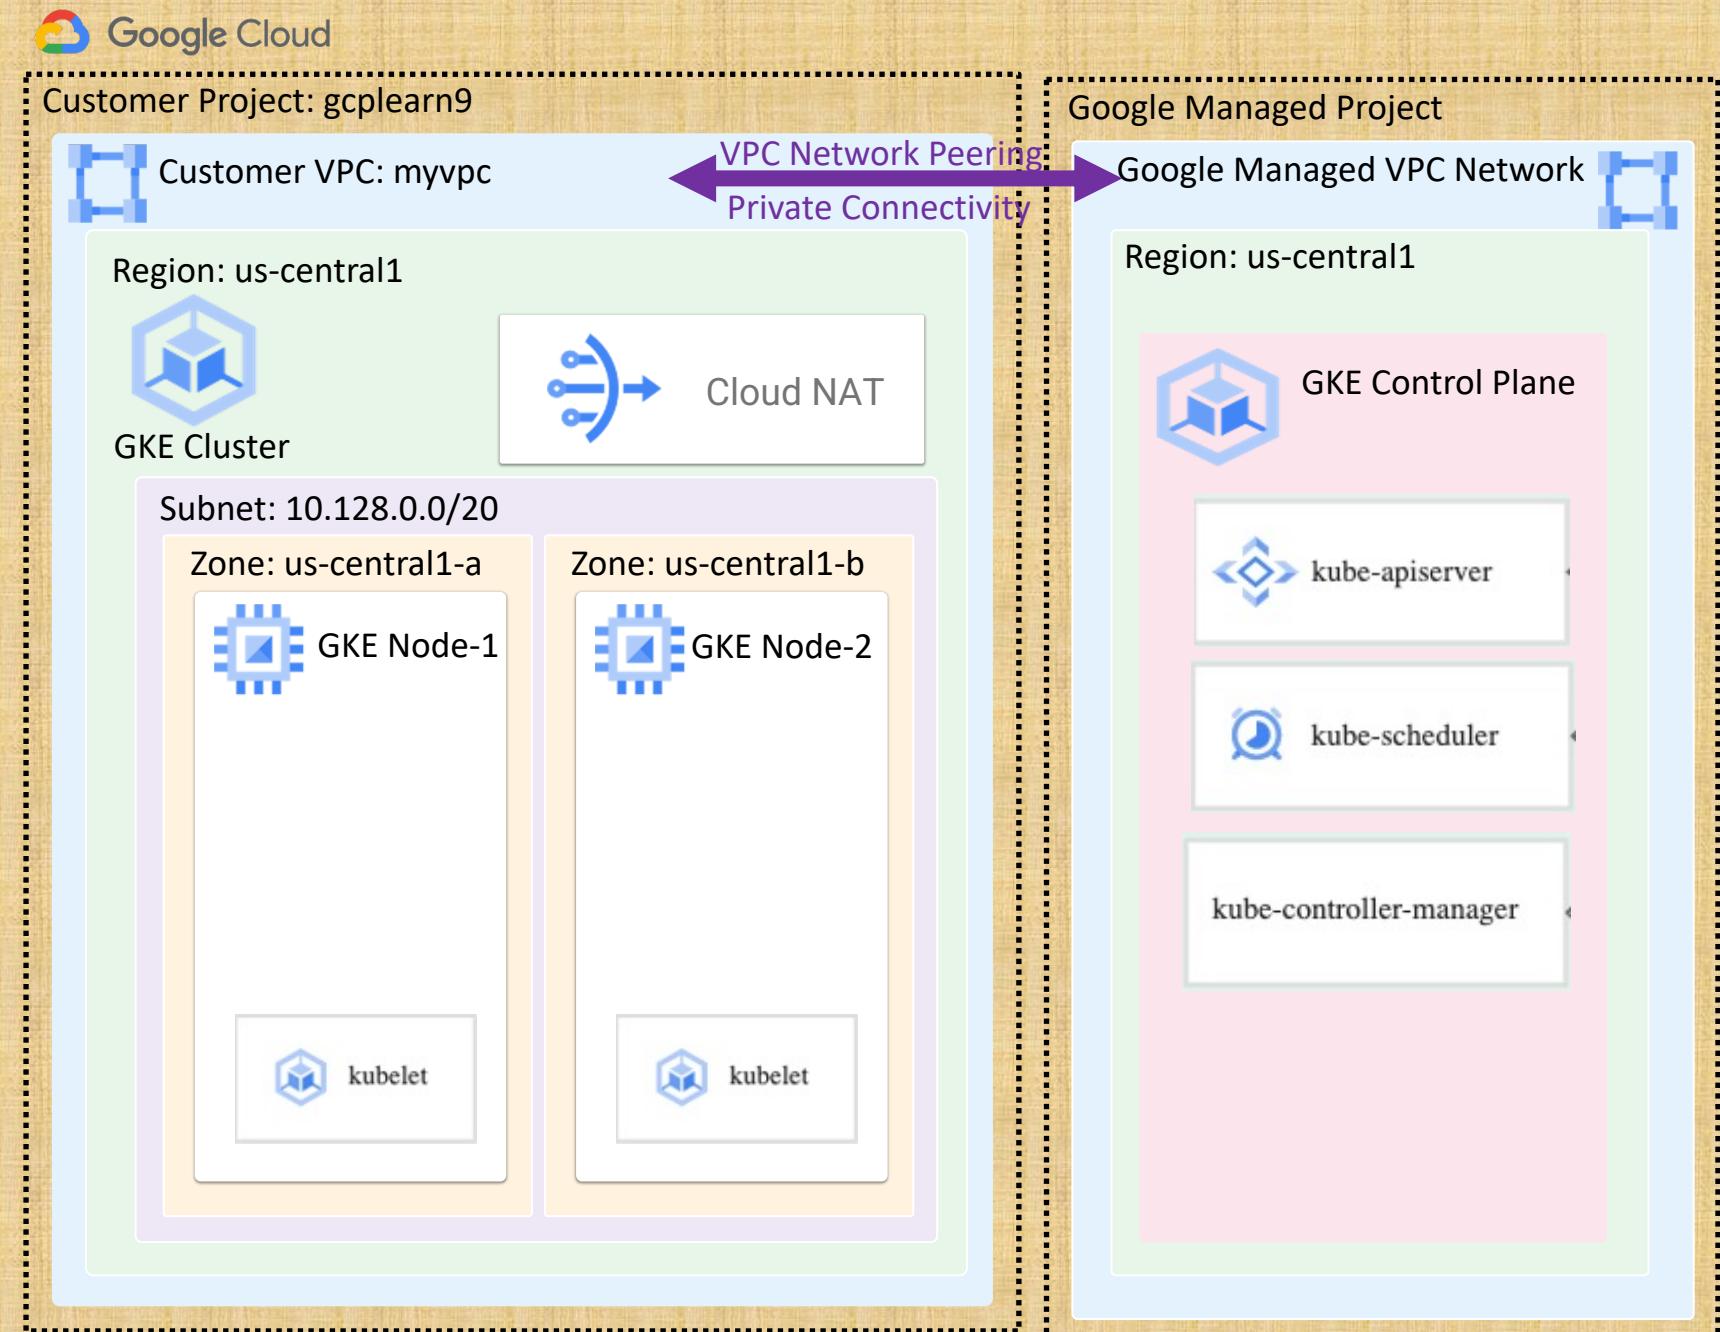
# GKE Private Cluster - Network Design

- **Private Clusters:** Control Plane VPC network is connected to customer (our) VPC network using **VPC Network peering**
- Our VPC network contains **GKE nodes**
- Google managed VPC network contains **GKE cluster Control plane**
- Traffic between nodes and control plane is routed using **internal IP addresses only**



# GKE Private Clusters

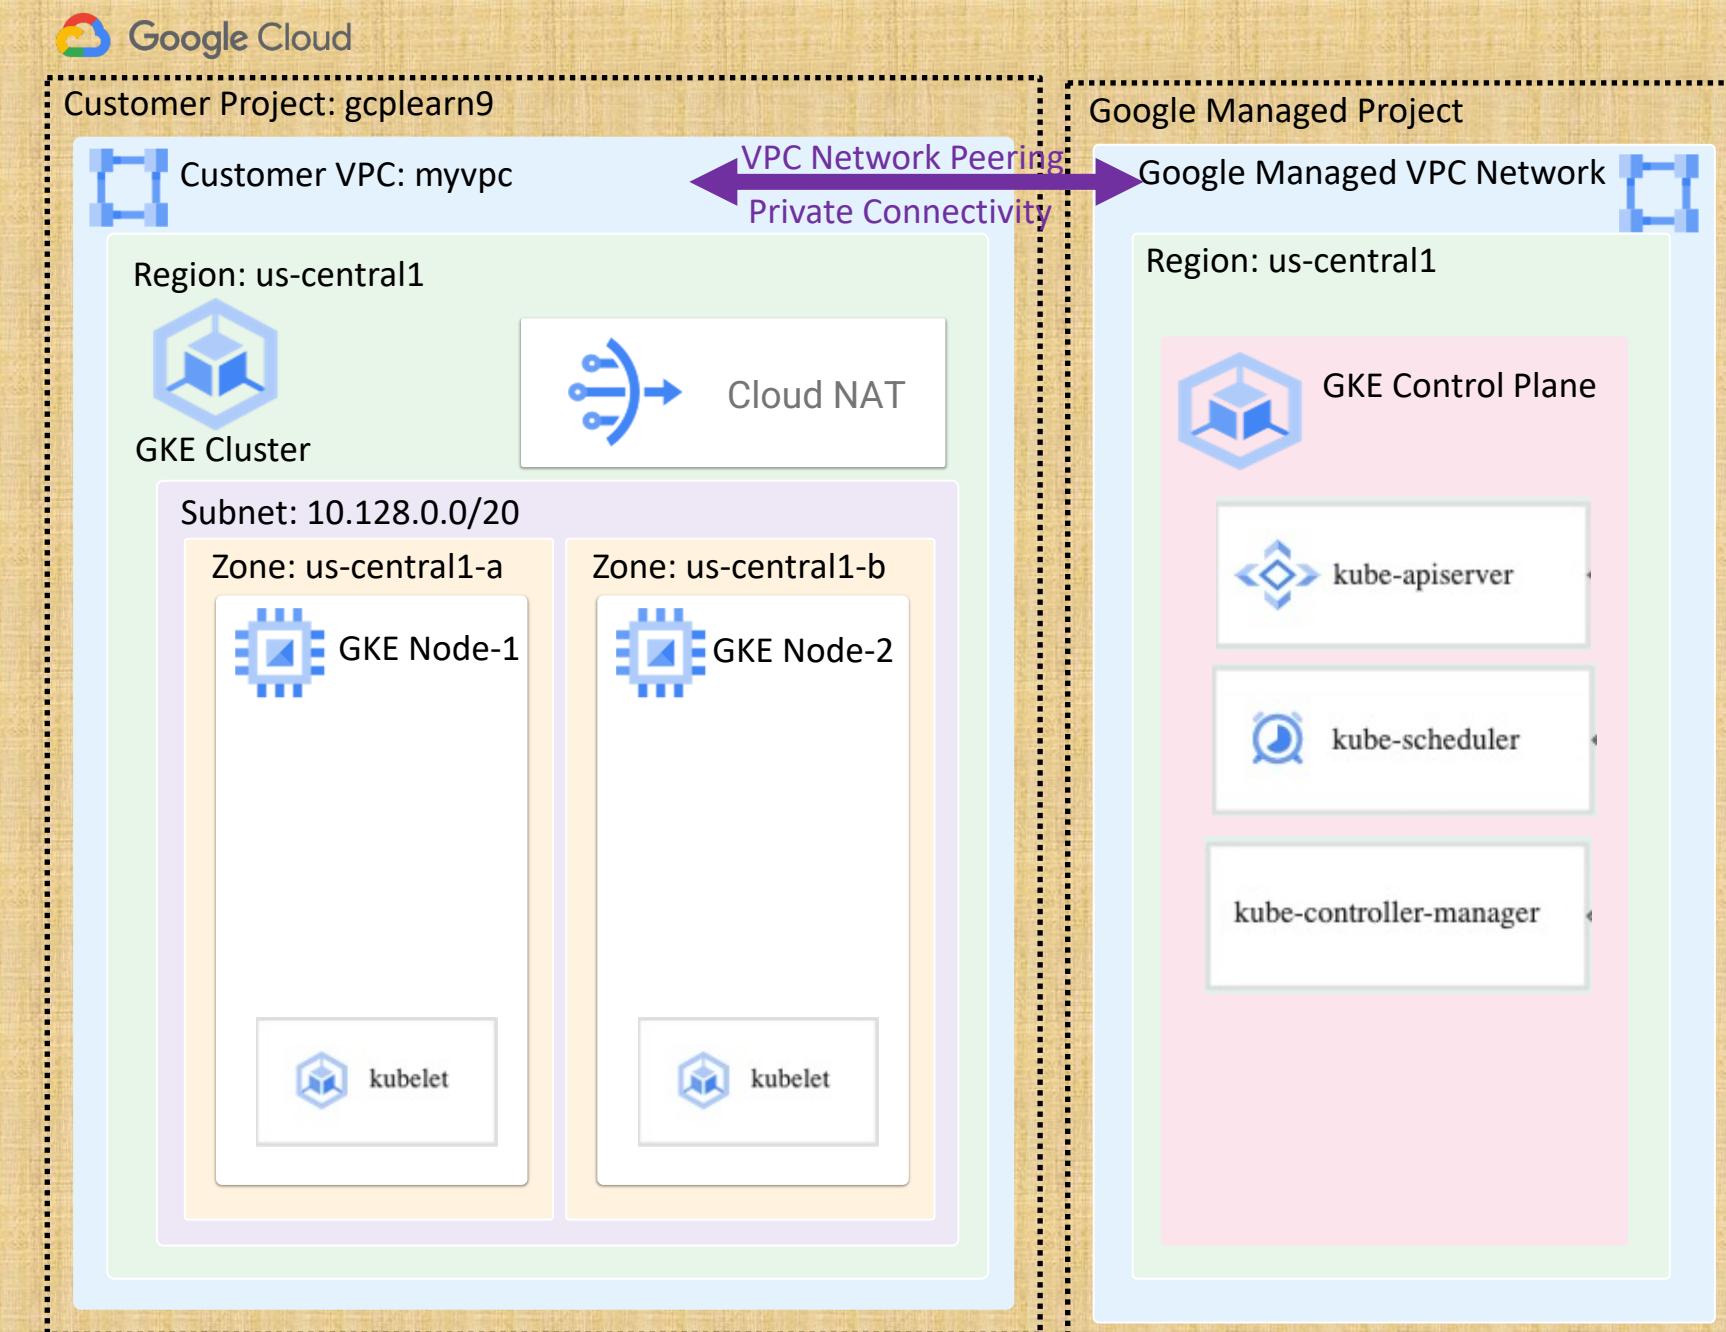
- You can create and configure private clusters in **Standard or Autopilot**.
- Private clusters will have nodes that **do not have** external IP addresses
- If you want to provide outbound internet access for certain private nodes, you can use **Cloud NAT** and **Cloud Router**



<https://cloud.google.com/kubernetes-engine/docs/concepts/alias-ips#benefits>

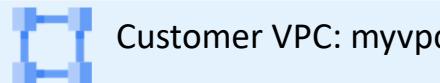
# GKE Private Clusters

- In private clusters, **private google access enabled by default on VPC subnet** to access other Google cloud APIs and services using private network
  - **Example-1:** To access container images from [Artifact Registry](#)
  - **Example-2:** To send logs to [Cloud Logging](#)



# GKE Standard Private Cluster - Network Design

Customer Project: gcplearn9



Region: us-central1



Subnet: 10.128.0.0/20

Zone: us-central1-a



Load Balancer Service

Private IP

Automated by Terraform

Cloud Load Balancing

Cloud NAT

Cloud Router

Zone: us-central1-b



Private IP

Google Managed Project

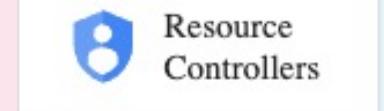
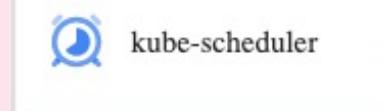
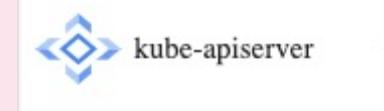


Region: us-central1



Kube API Server Public IP

Kube API Server Private IP



VPC Network Peering

Private Connectivity

Private Connectivity

Download Docker Image from Docker Hub

<http://LB-IP>



kubectl



Internet

NETWORK DESIGN

# GKE Private Cluster - Access using kubectl

- **Least secure Option**

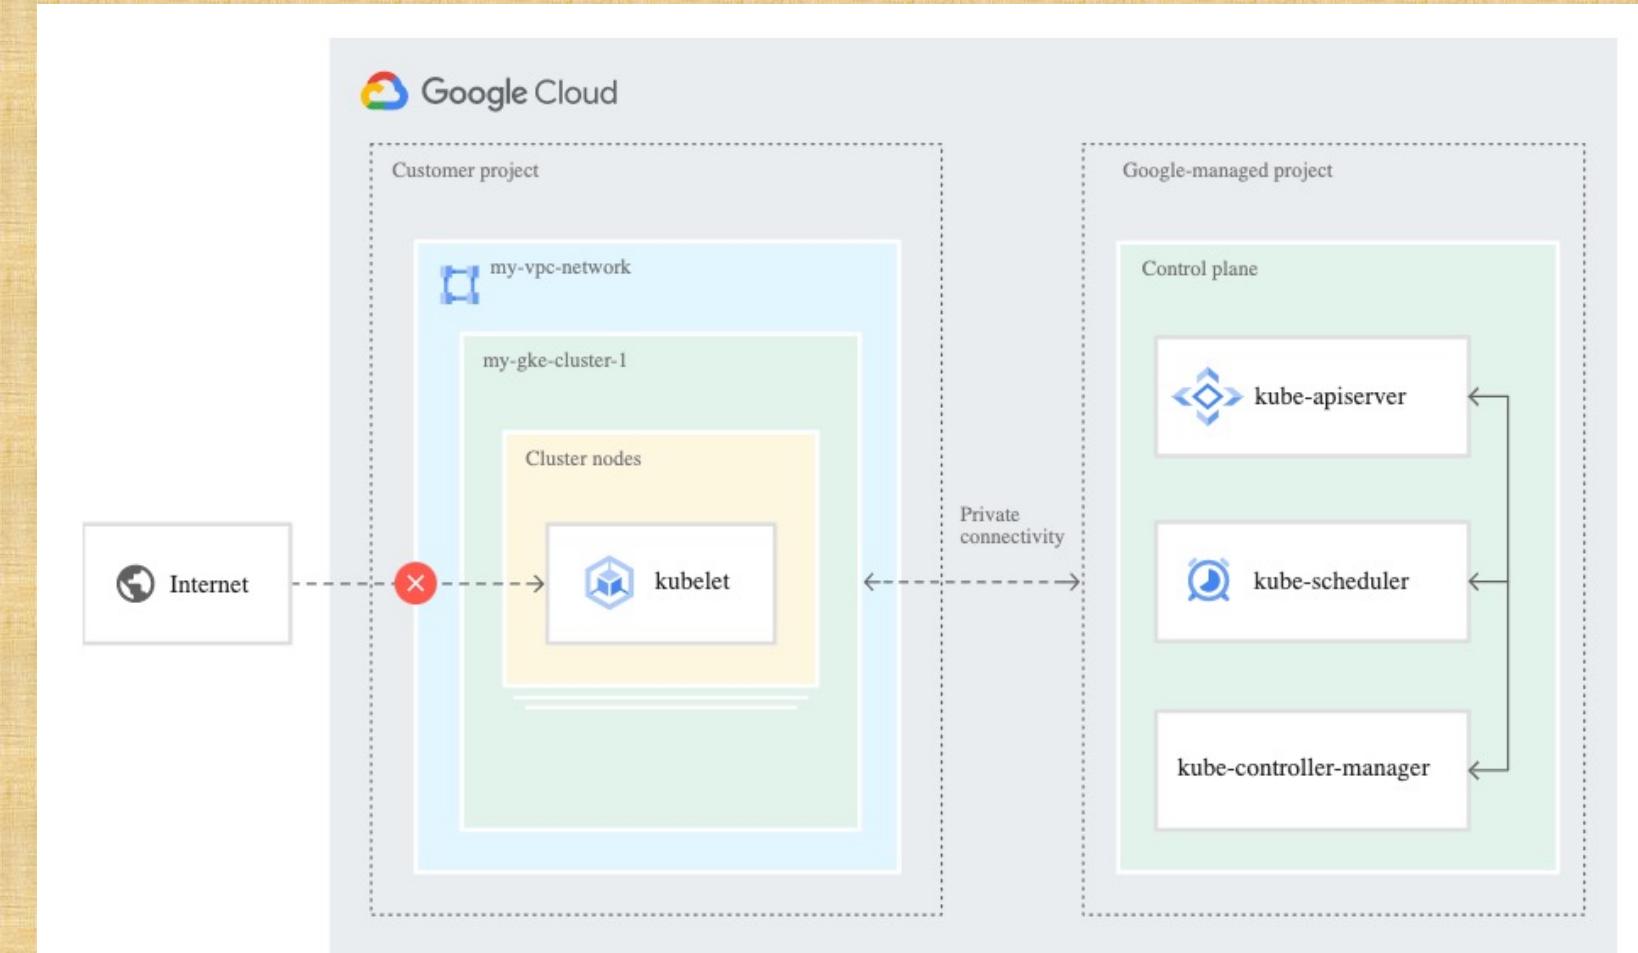
- Public endpoint access: **Enabled**
- Authorized Networks: **Disabled**
- Accessible via **Internet**

- **Medium secure option: Demo-12**

- Public endpoint access: **Enabled**
- Authorized Networks: **Enabled**
- Accessible via **authorized internet IP ranges** (Example: CloudShell, local desktop, from specified network in authorized network)

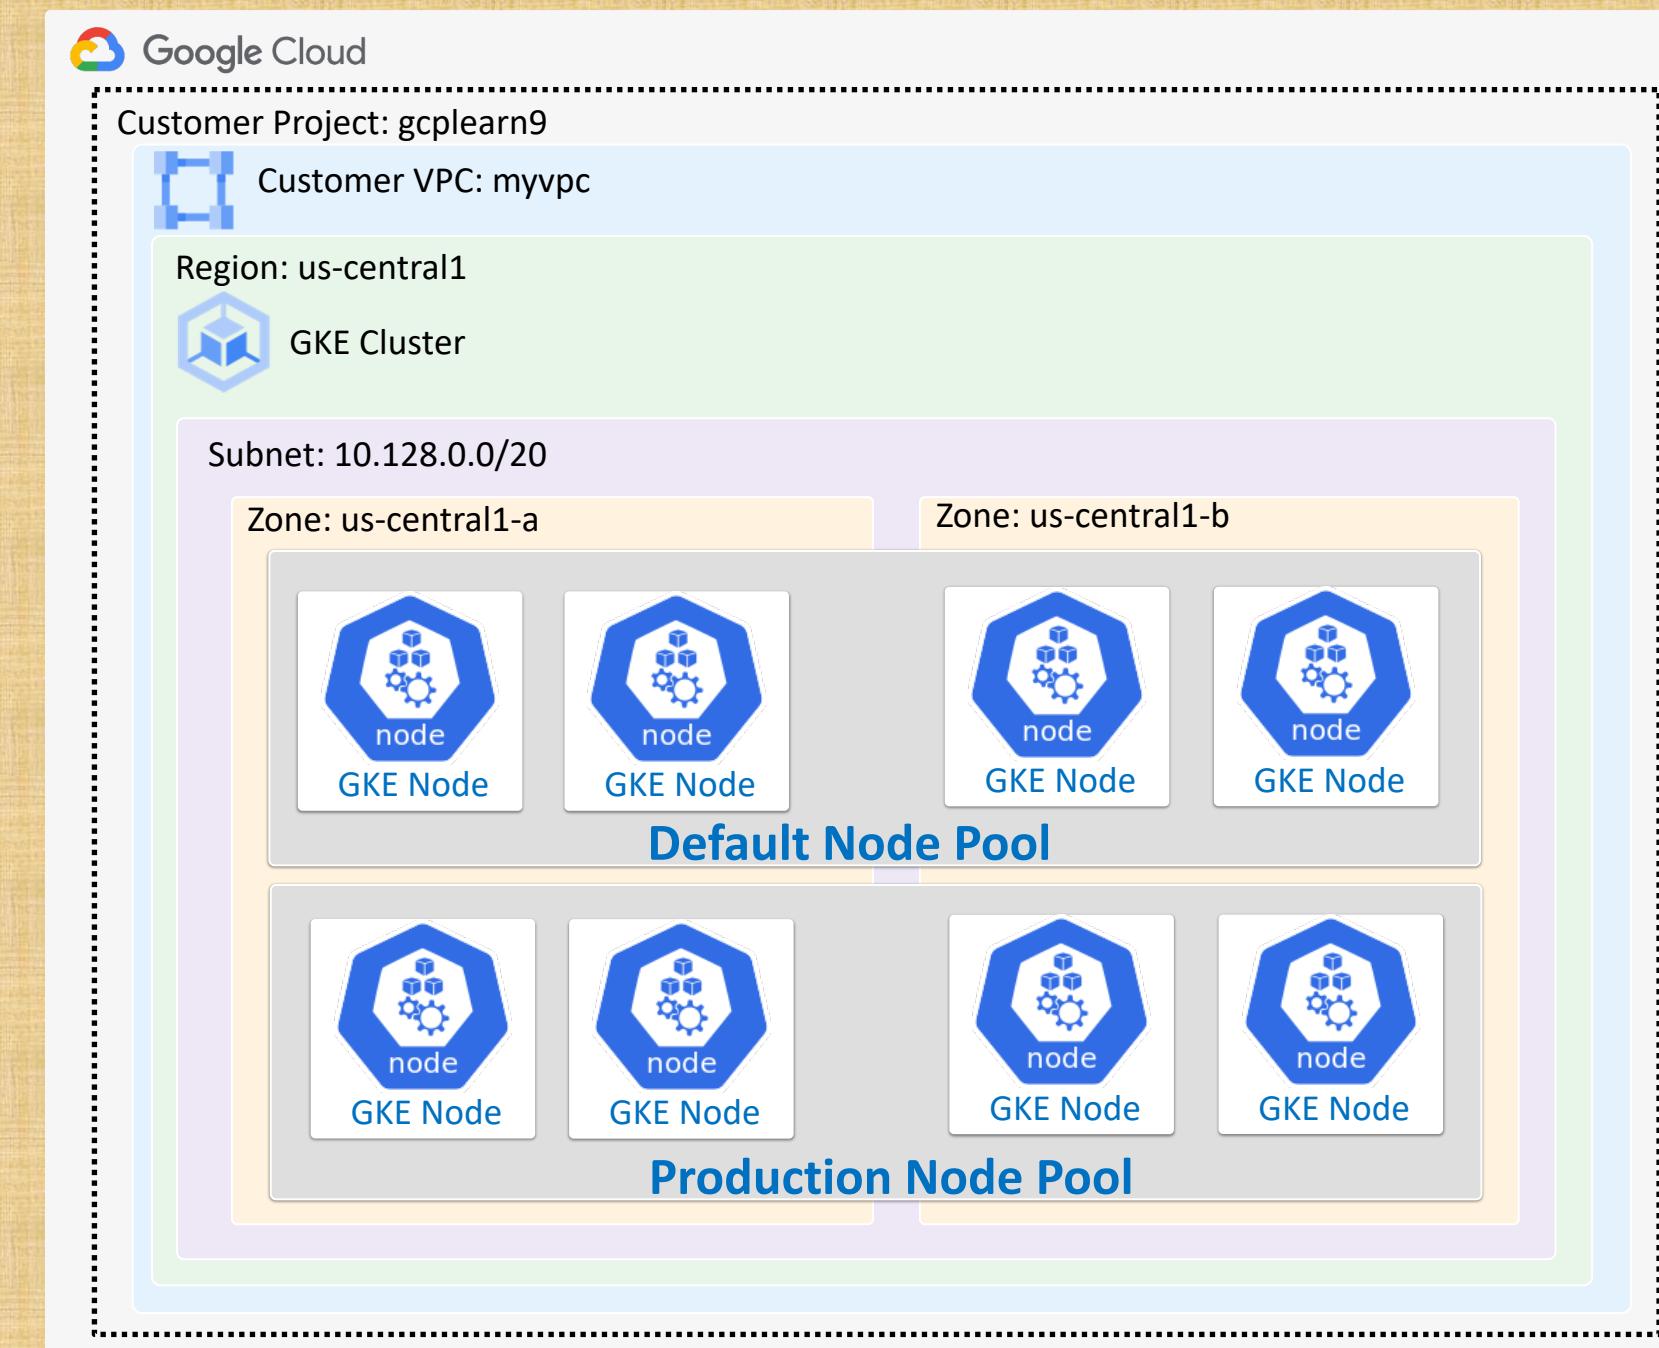
- **High secure option: Demo-15**

- Public endpoint access: **Disabled**
- Accessible via
  - VM in google cloud VPC network
  - **On-premise network** provided Cloud VPN or Cloud Interconnect is configured



# Kubernetes Cluster Autoscaler

- **Cluster Autoscaler:** **Automatically resizes** the number of nodes in a node pool, based on the **demands of your workloads**.
  - Minimum Size: 3
  - Maximum Size: 12
- **Scale-Out:** **Gradually increases** nodes to maximum size based on need (When pods are unschedulable)
- **Scale-In:** Decreases nodes to minimum size when nodes are idle (no workloads scheduled on them)



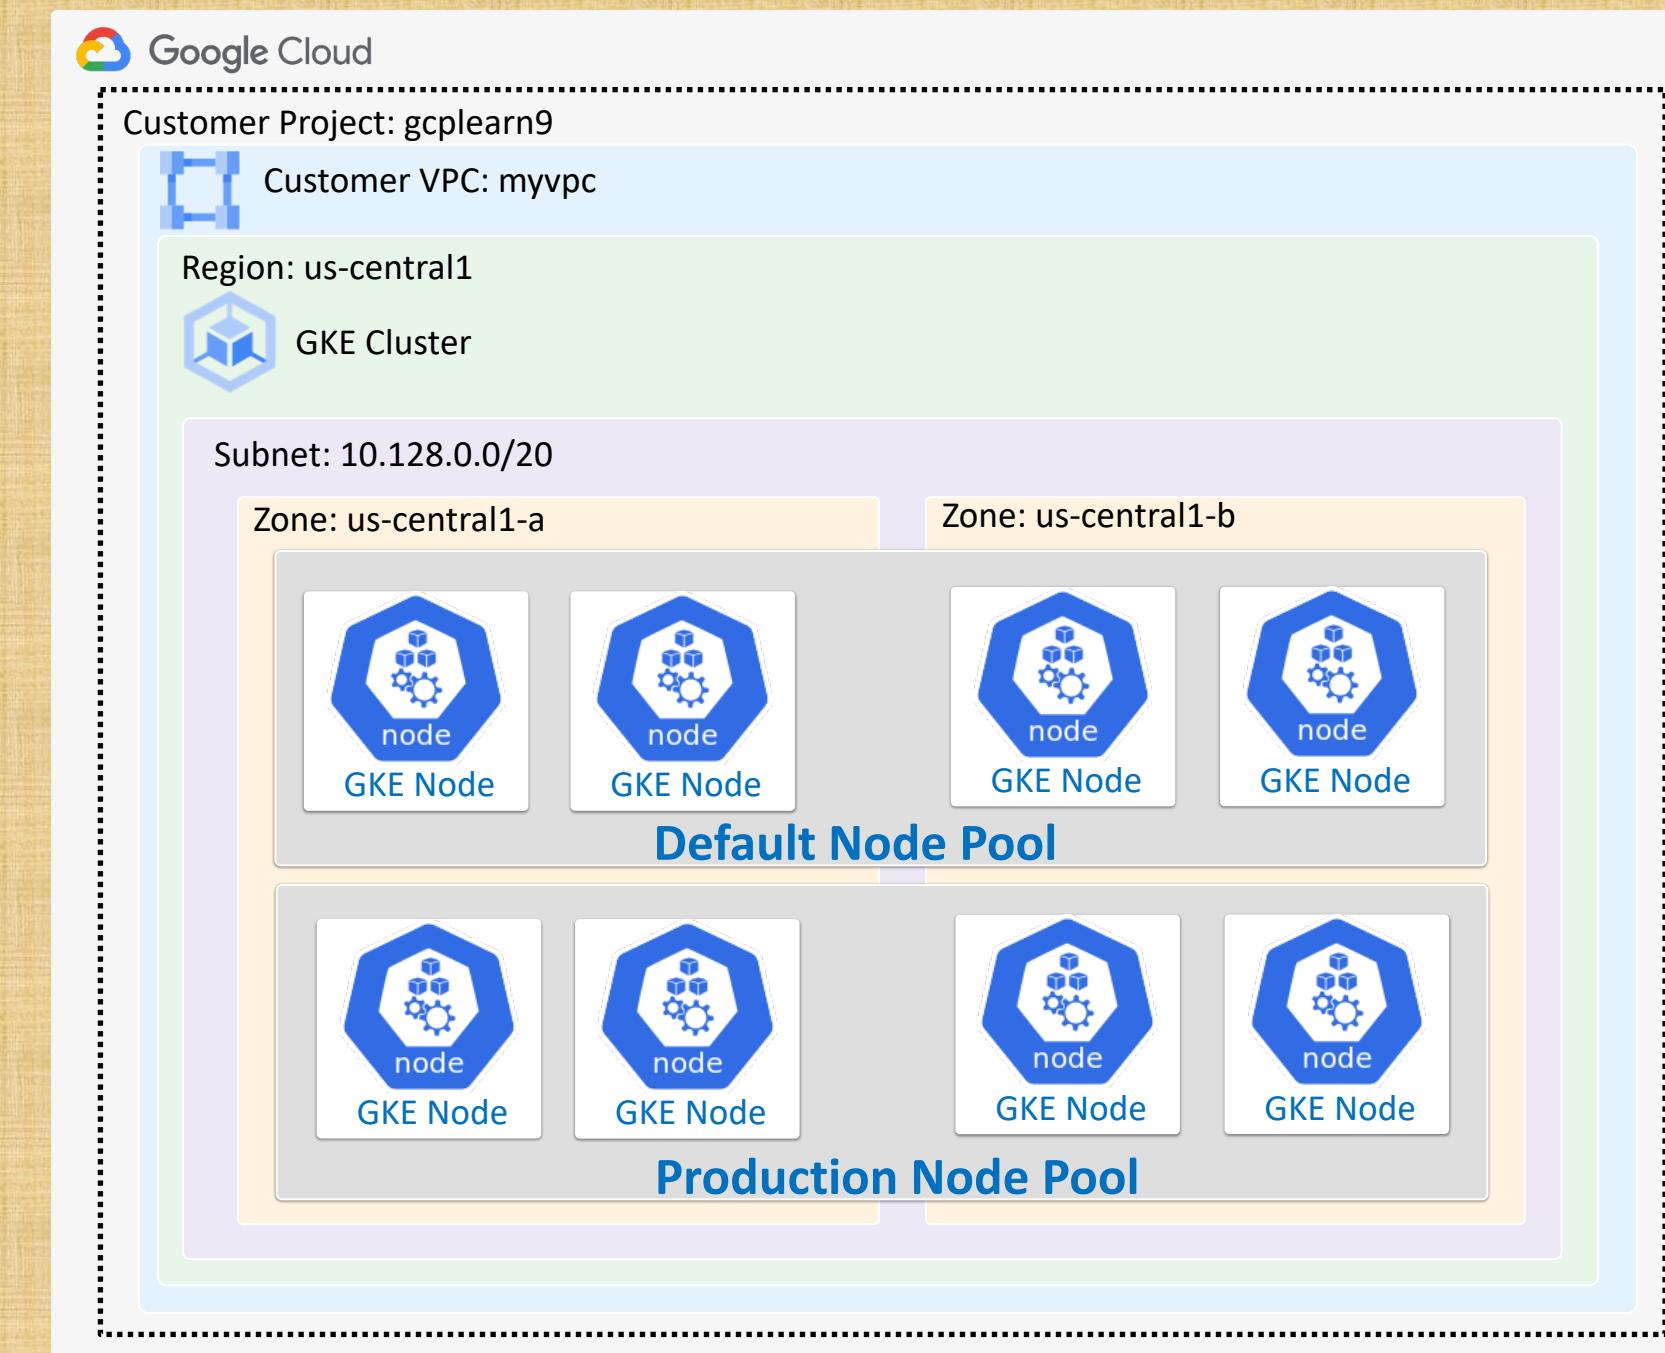
# Kubernetes Cluster Autoscaler

- Benefits

- Increases the **availability** of your workloads
- **Control** the costs
- **Effective** use of system resources (CPU, memory)

- Standard vs Autopilot Clusters

- Standard Cluster uses **Cluster Autoscaler**
- Autopilot cluster uses **Node auto-provisioning**
  - GKE automatically takes care of nodes and manage node pools.

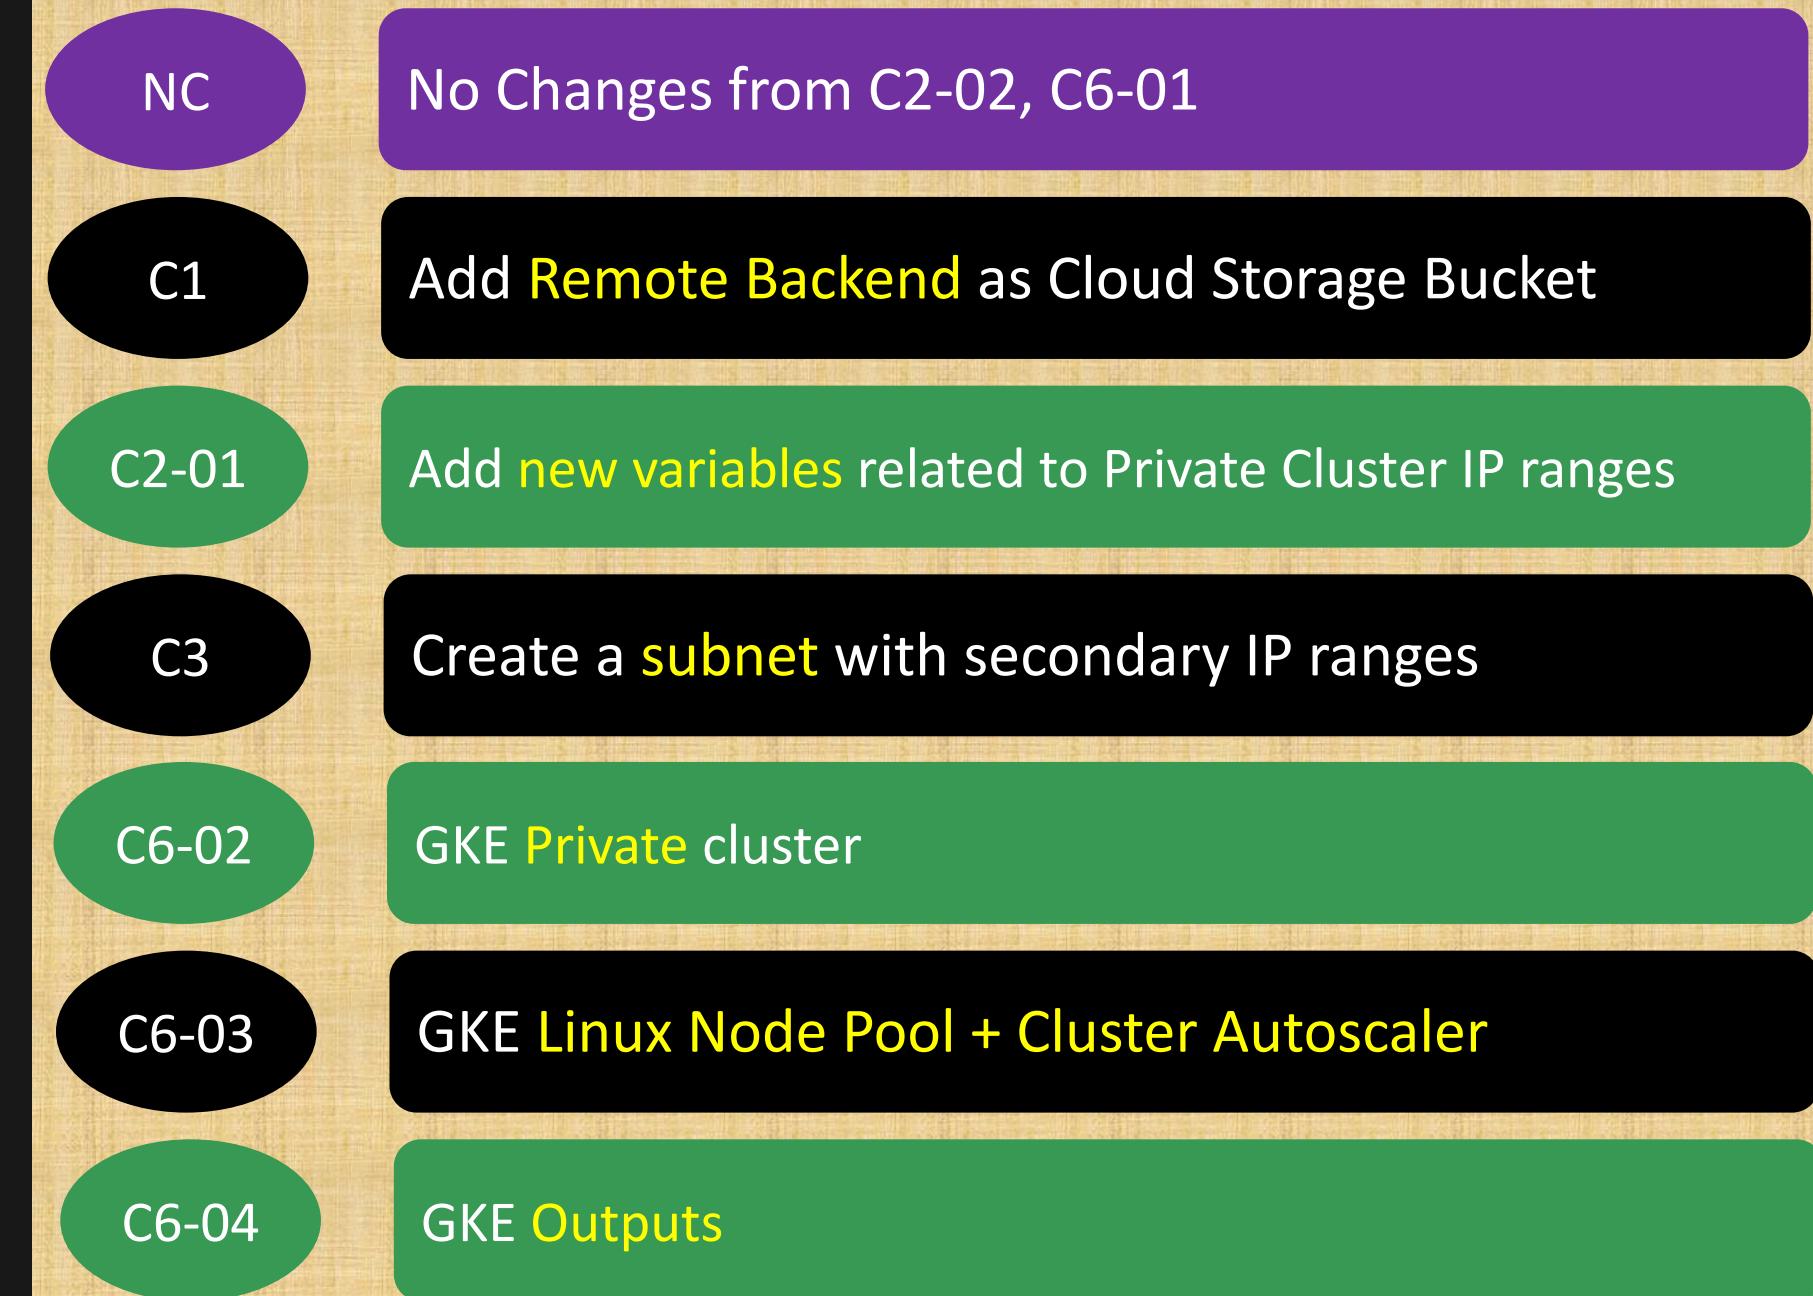


## Project-1: GKE Cluster

- ✓ p1-gke-private-cluster-autoscaler
- └ c1-versions.tf
- └ c2-01-variables.tf
- └ c2-02-local-values.tf
- └ c3-vpc.tf
- └ c4-firewallrules.tf
- └ c5-datasource.tf
- └ c6-01-gke-service-account.tf
- └ c6-02-gke-cluster.tf
- └ c6-03-gke-linux-nodepool.tf
- └ c6-04-gke-outputs.tf
- └ c7-Cloud-NAT-Cloud-Router.tf
- └ terraform.tfvars

# What are we going to learn?

Demo: GKE Standard Private Cluster + k8s API Server Public Endpoint + Cluster Autoscaler



## Project-1: GKE Cluster

- ✓ p1-gke-private-cluster-autoscaler
- └ c1-versions.tf
- └ c2-01-variables.tf
- └ c2-02-local-values.tf
- └ c3-vpc.tf
- └ c4-firewallrules.tf
- └ c5-datasource.tf
- └ c6-01-gke-service-account.tf
- └ c6-02-gke-cluster.tf
- └ c6-03-gke-linux-nodepool.tf
- └ c6-04-gke-outputs.tf
- └ c7-Cloud-NAT-Cloud-Router.tf
- └ terraform.tfvars

# What are we going to learn?

Demo: GKE Standard Private Cluster + k8s API Server Public Endpoint

C7

TFVARS

Cloud NAT and Cloud Router Resources

Update Input variable values: terraform.tfvars related to Private Cluster CIDR Ranges

# What are we going to learn?

**Demo: Kubernetes Deployment and Service using Terraform + Terraform Remote State Data source concept**

- ✓ p3-k8sresources-terraform-manifests
  - └ c1-versions.tf
  - └ c2-01-variables.tf
  - └ c2-02-local-values.tf
  - └ c3-01-remote-state-datasource.tf
  - └ c3-02-providers.tf
  - └ c4-kubernetes-deployment.tf
  - └ c5-kubernetes-loadbalancer-service.tf
  - └ terraform.tfvars

NC

C3-01

No Changes: c1, c2-01, c2-02, c3-02, c4, c5

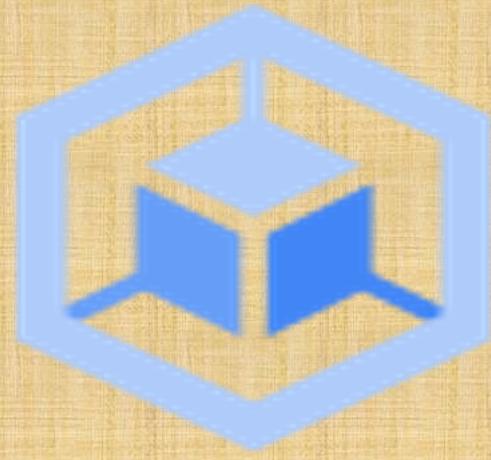
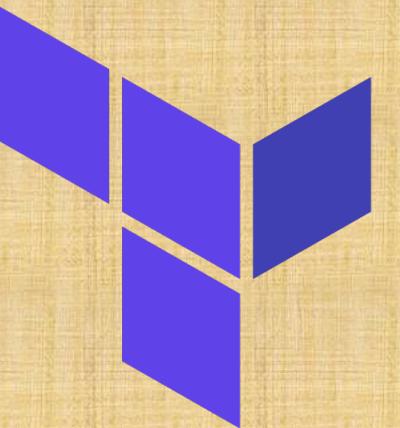
Update **project-1** Terraform state file **bucket prefix** in  
Remote State Data source



Demo-13

# Google Kubernetes Engine Horizontal Pod Autoscaler

**Demo: GKE Horizontal Pod Autoscaler**



# Kubernetes Horizontal Pod Autoscaling

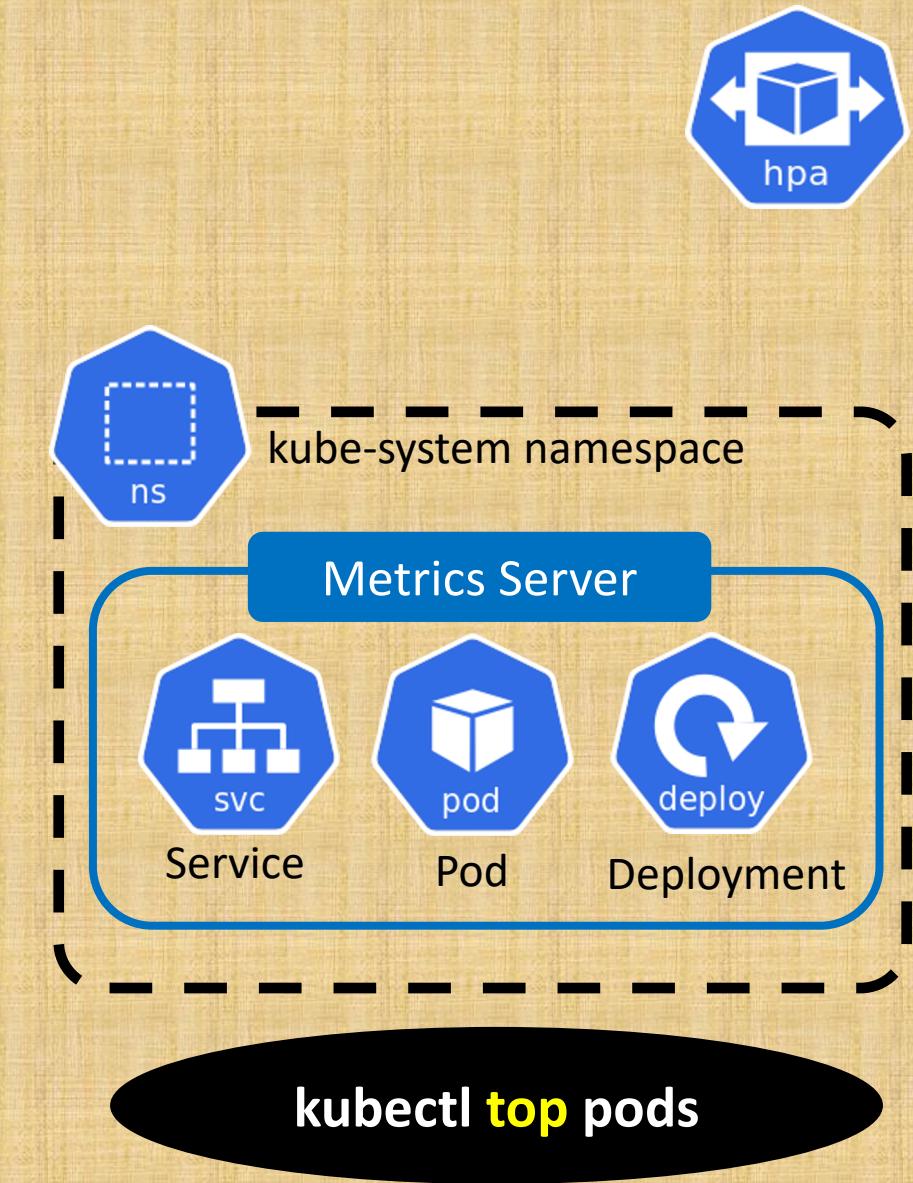


- Automatically increase or decrease number of pods in response to
  - Workloads CPU and Memory Utilization
  - Custom metrics reported from within Kubernetes Cluster
  - External metrics (Load Balancers, Messaging Services ...)
  - Custom metrics using Managed Service for Prometheus
- HPA automatically scales the pods in workload types
  - Kubernetes ReplicaSet
  - Kubernetes Replication Controller
  - Kubernetes Deployment
  - Kubernetes StatefulSet
- This can help our applications
  - Scale out to meet increased demand or
  - Scale in when resources are not needed, thus freeing up your worker nodes for other applications

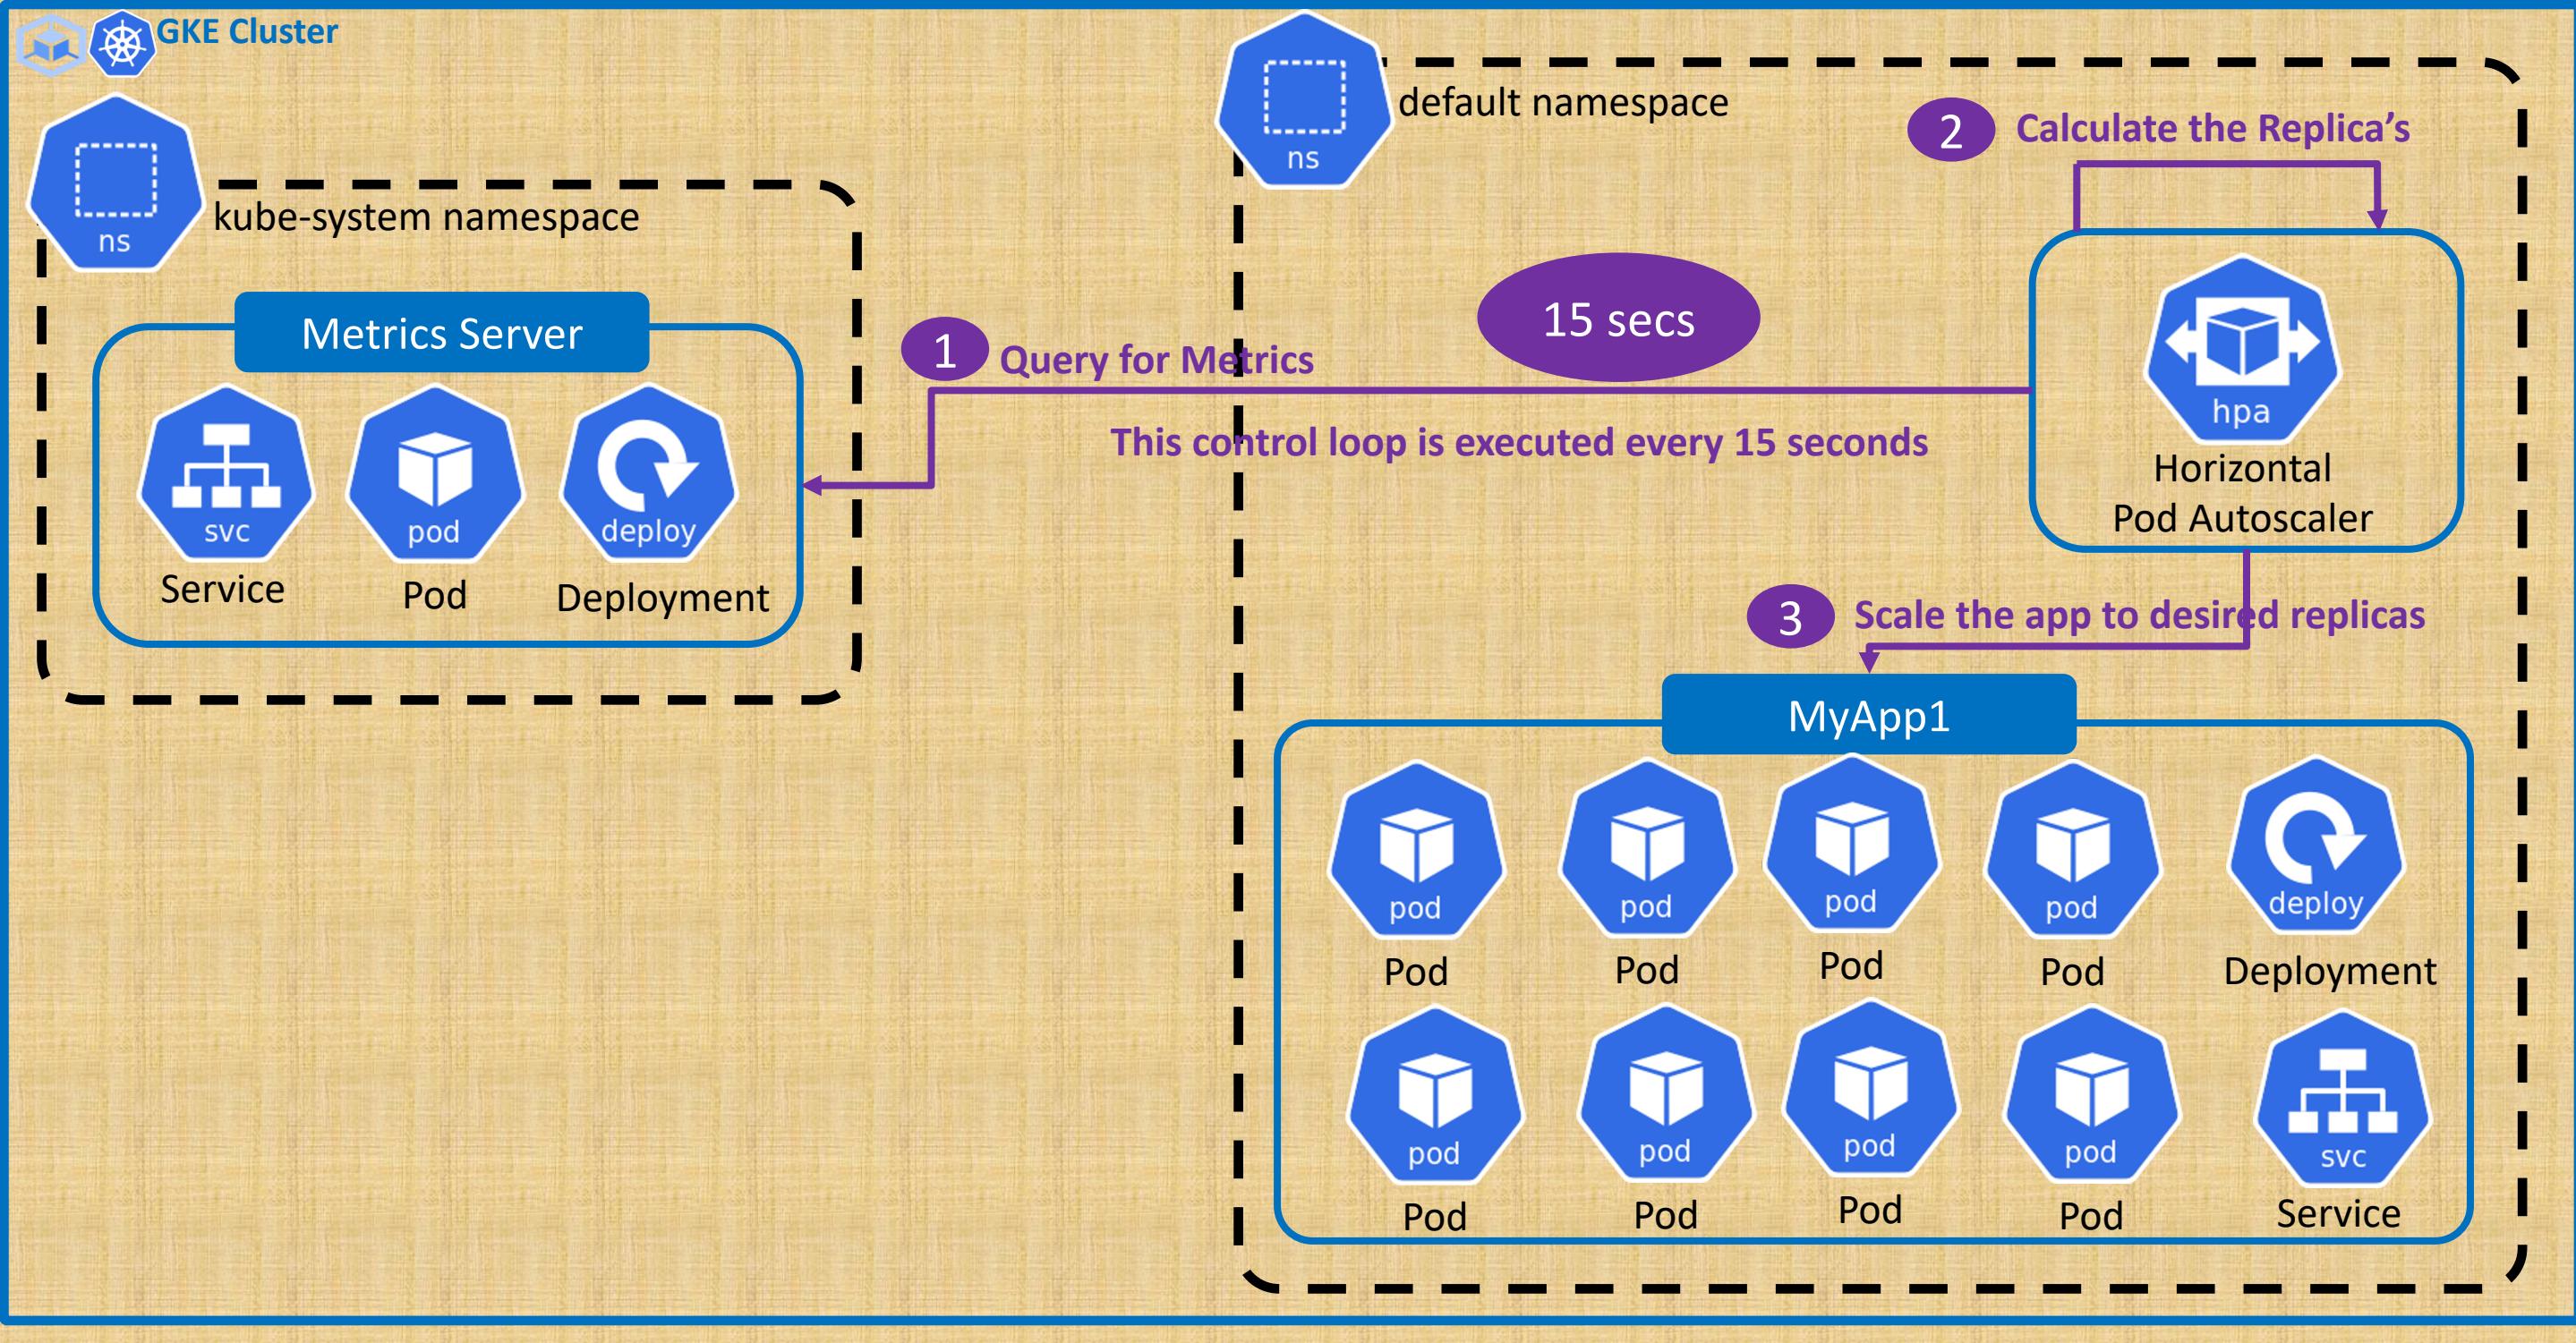
**HPA Imperative Command:** kubectl autoscale deployment my-app --max 6 --min 4 --cpu-percent 50

# Kubernetes Metrics Server

- Metrics Server collects resource metrics from Kubelets and exposes them in Kubernetes apiserver through Metrics API
- Metrics API can also be accessed by `kubectl top`, making it easier to debug autoscaling pipelines.
- Fast Autoscaling solution, collecting metrics every 15 seconds
- Metrics Server is not meant for non-autoscaling purposes. For example, don't use it to forward metrics to monitoring solutions.
- Resource efficiency, using 1 mili core of CPU and 2 MB of memory for each node in a cluster
- Metrics Server used for CPU/Memory based horizontal autoscaling
- Metrics Server used for automatically adjusting/suggesting resources needed by containers (Vertical Autoscaling)



# Kubernetes Horizontal Pod Autoscaling



# Kubernetes Horizontal Pod Autoscaling

Autoscaling v1

```
apiVersion: autoscaling/v1
kind: HorizontalPodAutoscaler
metadata:
  name: hpa-myapp1
spec:
  scaleTargetRef:
    apiVersion: apps/v1
    kind: Deployment
    name: myapp1-deployment
  minReplicas: 1
  maxReplicas: 10
  targetCPUUtilizationPercentage: 50
```

Autoscaling v2

```
apiVersion: autoscaling/v2
kind: HorizontalPodAutoscaler
metadata:
  name: cpu
spec:
  scaleTargetRef:
    apiVersion: apps/v1
    kind: Deployment
    name: myapp1-deployment
  minReplicas: 1
  maxReplicas: 10
  metrics:
  - type: Resource
    resource:
      name: cpu
    target:
      type: Utilization
      averageUtilization: 30
```



Horizontal  
Pod Autoscaler

## Project-2: HPA V1

- ✓ p2-k8sresources-yaml-hpa-v1
  - ! 01-kubernetes-deployment.yaml
  - ! 02-kubernetes-cip-service.yaml
  - ! 03-kubernetes-hpa.yaml

# What are we going to learn?

Demo: HPA using Autoscaler V1 API Version

01

Kubernetes Deployment

02

Kubernetes ClusterIP Service

03

Kubernetes HPA Autoscaler V1 API Version

## Project-3: HPA V2

- ✓ p3-k8sresources-yaml-hpa-v2
  - ! 01-kubernetes-deployment.yaml
  - ! 02-kubernetes-cip-service.yaml
  - ! 03-kubernetes-hpa.yaml

Demo: HPA using Autoscaler V1 API Version

01

Kubernetes Deployment

02

Kubernetes ClusterIP Service

03

Kubernetes HPA Autoscaler V2 API Version

- ✓ p4-k8sresources-terraform-manifests
  - └ c1-versions.tf
  - └ c2-01-variables.tf
  - └ c2-02-local-values.tf
  - └ c3-01-remote-state-datasource.tf
  - └ c3-02-providers.tf
  - └ c4-kubernetes-deployment.tf
  - └ c5-kubernetes-clusterip-service.tf
  - └ c6-kubernetes-hpa.tf
  - └ terraform.tfvars

# What are we going to learn?

Demo: GKE Standard Private Cluster + HPA V2

NC

C5

C6

No Changes from C1 to C4

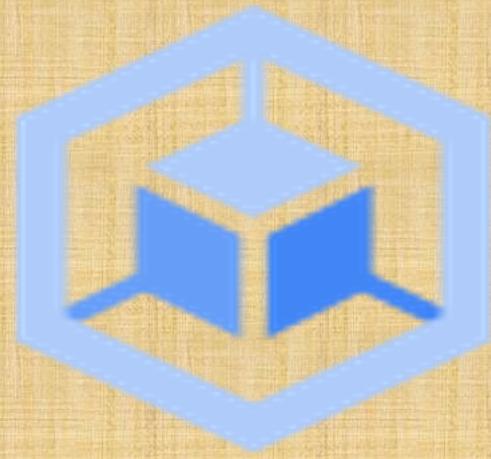
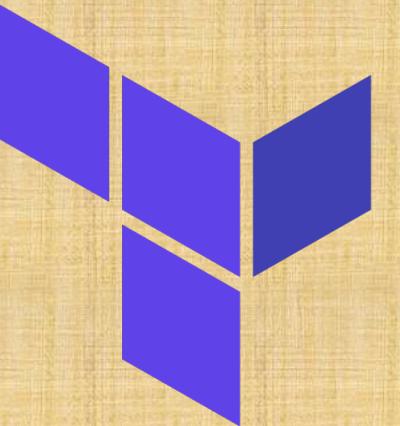
Kubernetes ClusterIP Service

Kubernetes HPA Autoscaler V2 API Version

Demo-14

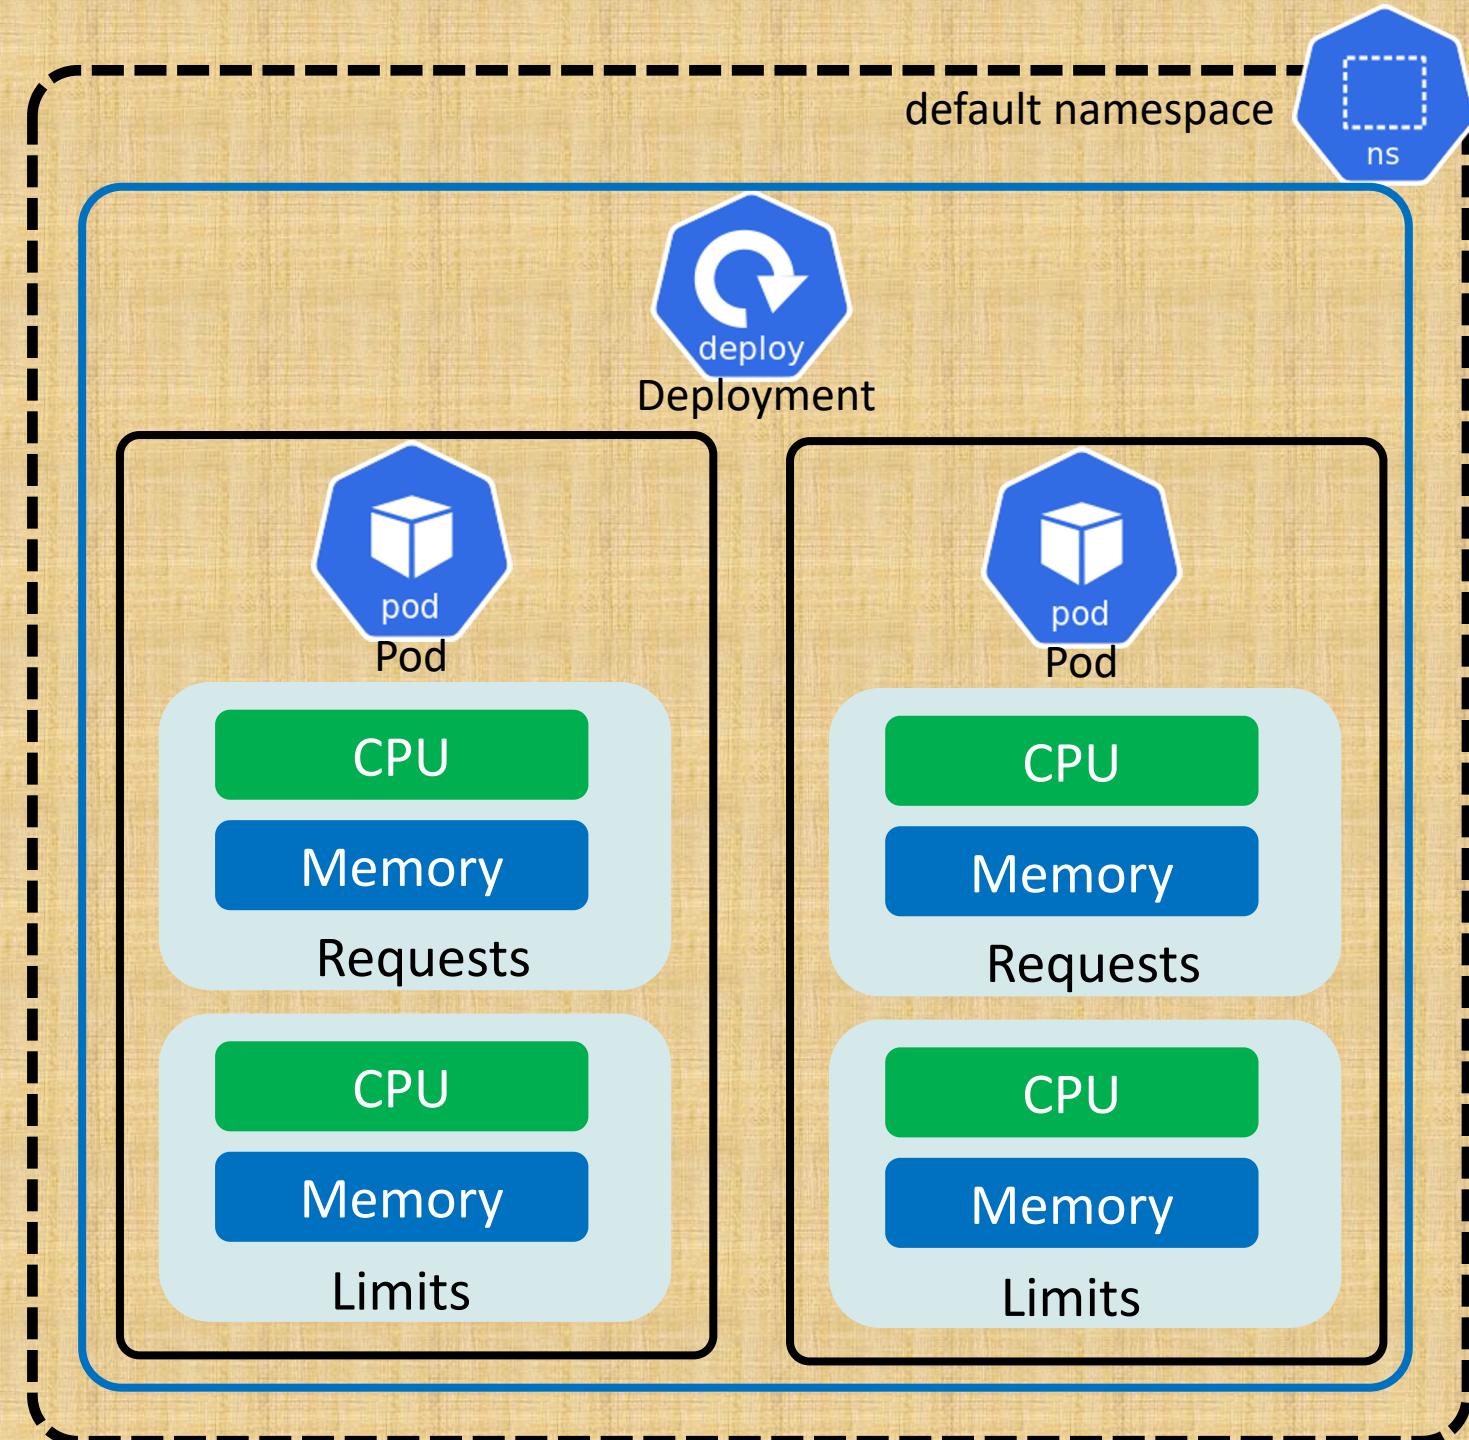
# Google Kubernetes Engine Vertical Pod Autoscaler

Demo: GKE Vertical Pod Autoscaler



# Kubernetes Vertical Pod Autoscaling

- **Vertical Pod autoscaling:** Lets you **analyze and set** CPU and memory resources required by Pods
- **Manual**
  - VPA **recommends** cpu and memory requests and limits
  - We can review the recommendation and update the values **manually**
- **Automatic**
  - VPA **recommends and automatically updates** the cpu and memory requests and limits
  - VPA **evicts and recreates** the pod during the resource request updates



# Kubernetes Vertical Pod Autoscaling

- **Standard vs Autopilot clusters**

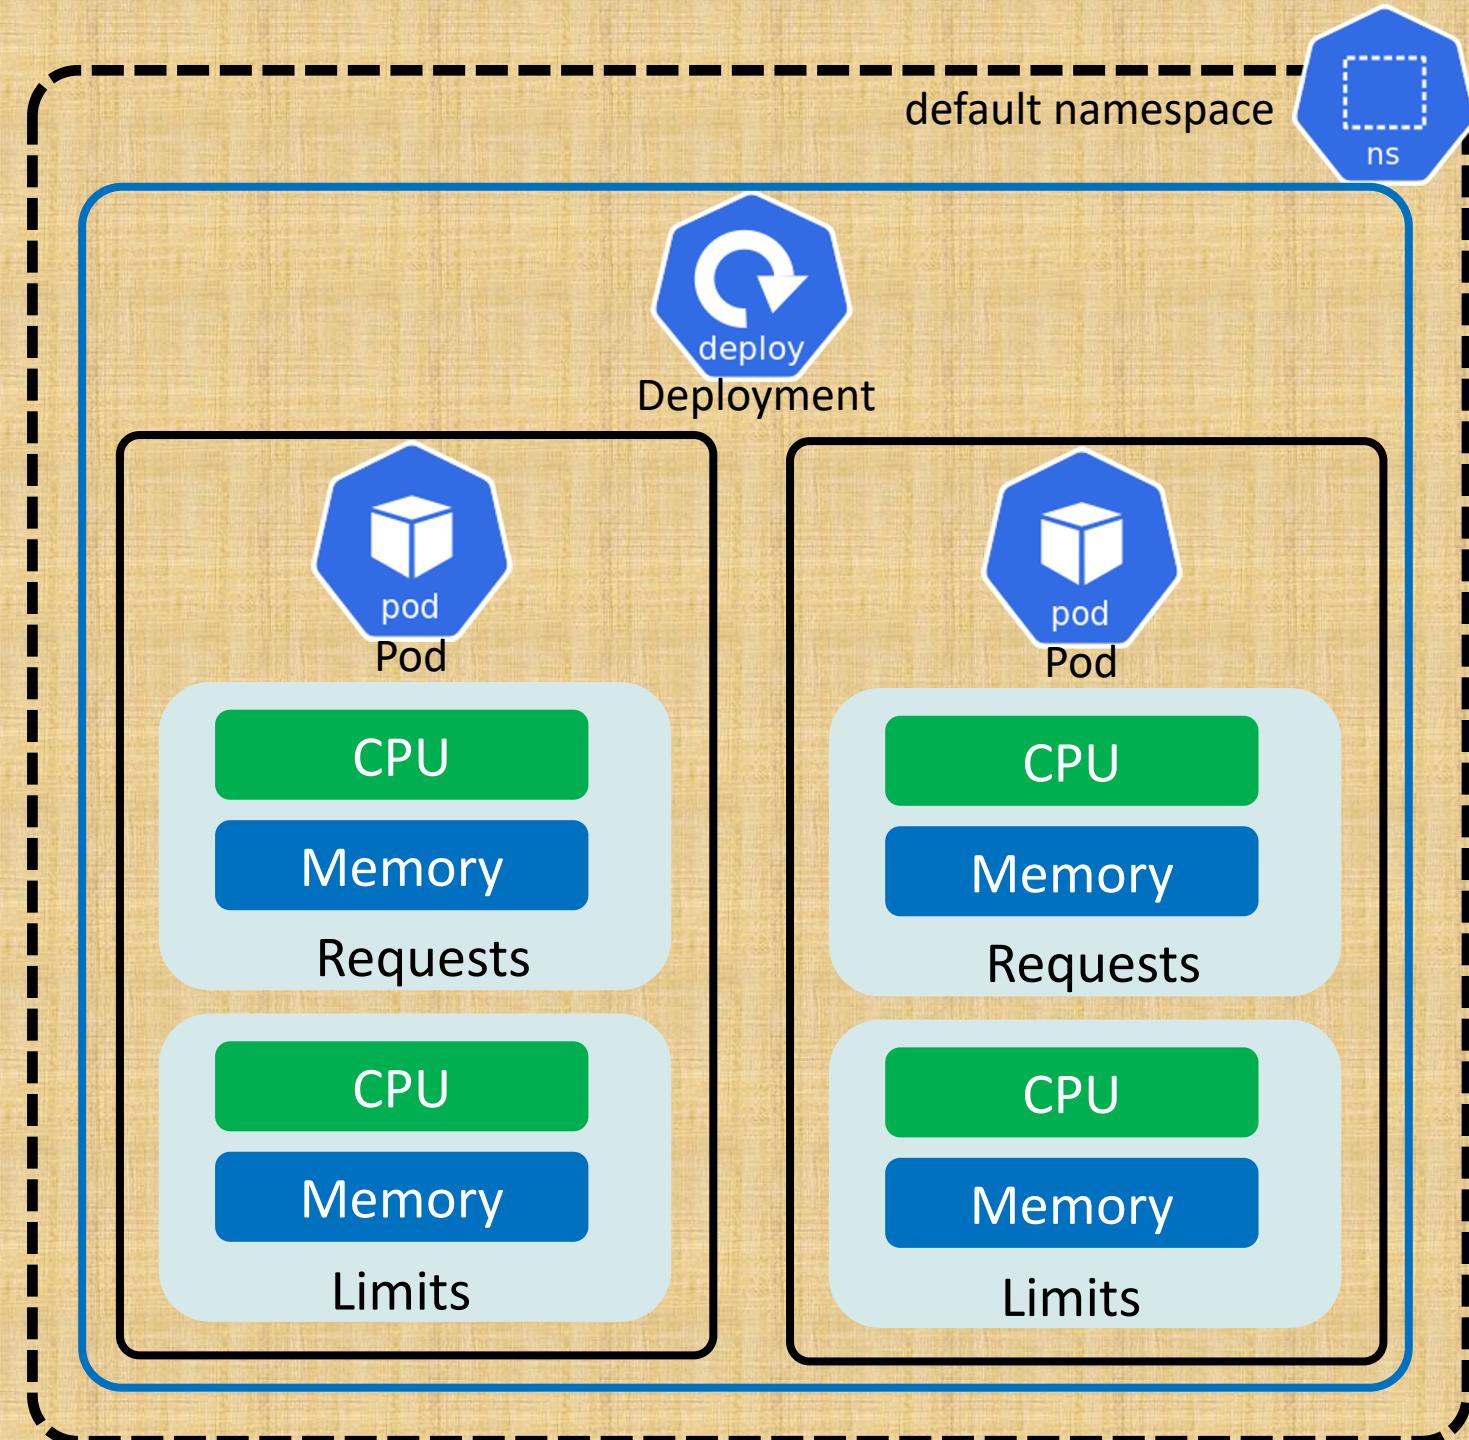
- VPA need to be enabled at [workload level](#) for standard clusters
- VPA is [by default enabled](#) in Autopilot clusters

- **Benefits**

- Cluster nodes are used [efficiently](#) because Pods use exactly what they need
- You don't have to run time-consuming [benchmarking tasks to determine](#) the correct values for CPU and memory requests
- Pods are scheduled onto nodes that have the [appropriate resources available](#)
- Runs Vertical Pod Autoscaler Pods as [control plane processes](#) instead of Deployment on worker nodes ([GKE special feature](#))

- **Important Note:**

- [Enable Cluster Autoscaler](#) before enabling VPA
- VPA [can talk to Cluster Autoscaler](#) to increase the number of nodes if needed for the VPA enabled workloads



# What are we going to learn?

- ✓ p2-k8sresources-yaml-vpa
  - ! 01-kubernetes-deployment.yaml
  - ! 02-kubernetes-cip-service.yaml
  - ! 03-vpa-manifest.yaml
- ✓ p3-k8sresources-terraform-manifests
  - ✓ c1-versions.tf
  - ✓ c2-01-variables.tf
  - ✓ c2-02-local-values.tf
  - ✓ c3-01-remote-state-datasource.tf
  - ✓ c3-02-providers.tf
  - ✓ c4-kubernetes-deployment.tf
  - ✓ c5-kubernetes-clusterip-service.tf
  - ✓ c6-kubernetes-vpa.tf
  - ✓ terraform.tfvars

## Demo: Project-2: VPA YAML Manifests

NC

No Changes – 01, 02

03

Kubernetes Vertical Pod Autoscaler

## Demo: Project-3: VPA Terraform Manifests

NC

No Changes from C1 to C5

C6

Kubernetes Vertical Pod Autoscaler

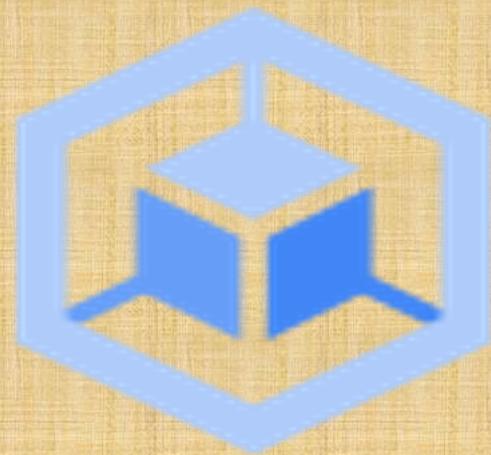
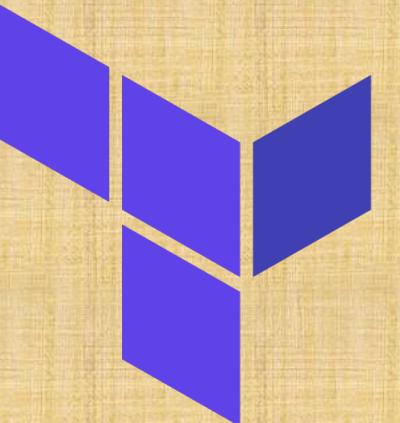


Demo-15

# Google Kubernetes Engine

Create GKE Standard Private Cluster  
with k8s API Server Private Endpoint

**Demo: GKE Standard Private Cluster with  
Private Kubernetes API Server Endpoint**



# GKE Private Cluster - Access using kubectl

- **Least secure Option**

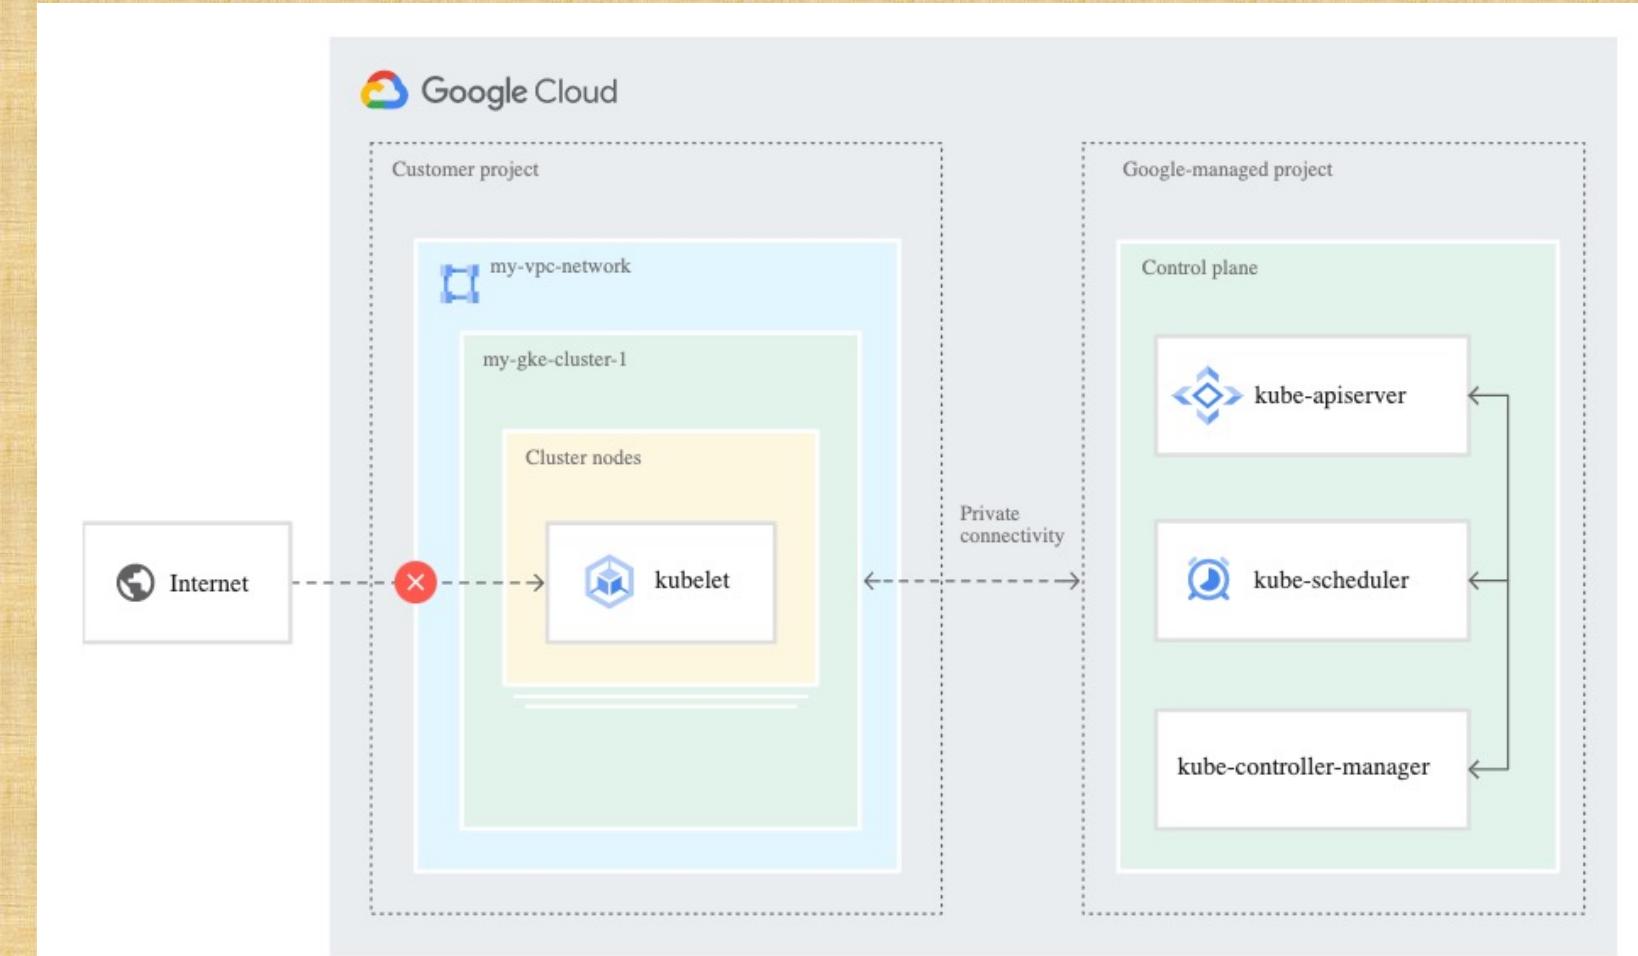
- Public endpoint access: **Enabled**
- Authorized Networks: **Disabled**
- Accessible via **Internet**

- **Medium secure option: Demo-12**

- Public endpoint access: **Enabled**
- Authorized Networks: **Enabled**
- Accessible via **authorized internet IP ranges** (Example: CloudShell, local desktop, from specified network in authorized network)

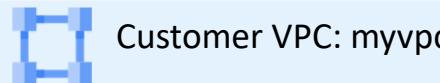
- **High secure option: Demo-15**

- Public endpoint access: **Disabled**
- Accessible via
  - VM in google cloud VPC network
  - **On-premise network** provided Cloud VPN or Cloud Interconnect is configured



# GKE Standard Private Cluster - Network Design

Customer Project: gcplearn9

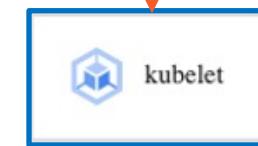


Region: us-central1



Subnet: 10.128.0.0/20

Zone: us-central1-a

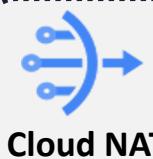


Load Balancer Service

Private IP

Automated by Terraform

Cloud Load Balancing

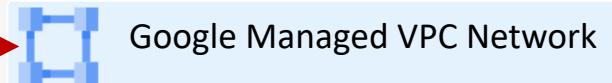


Zone: us-central1-b



Private IP

Google Managed Project

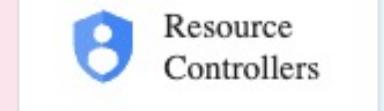
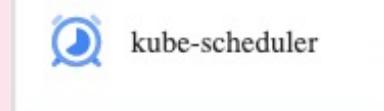
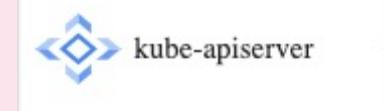


Region: us-central1



Kube API Server Public IP

Kube API Server Private IP



VPC Network Peering

Private Connectivity

Private Connectivity

Download Docker Image from Docker Hub

<http://LB-IP>



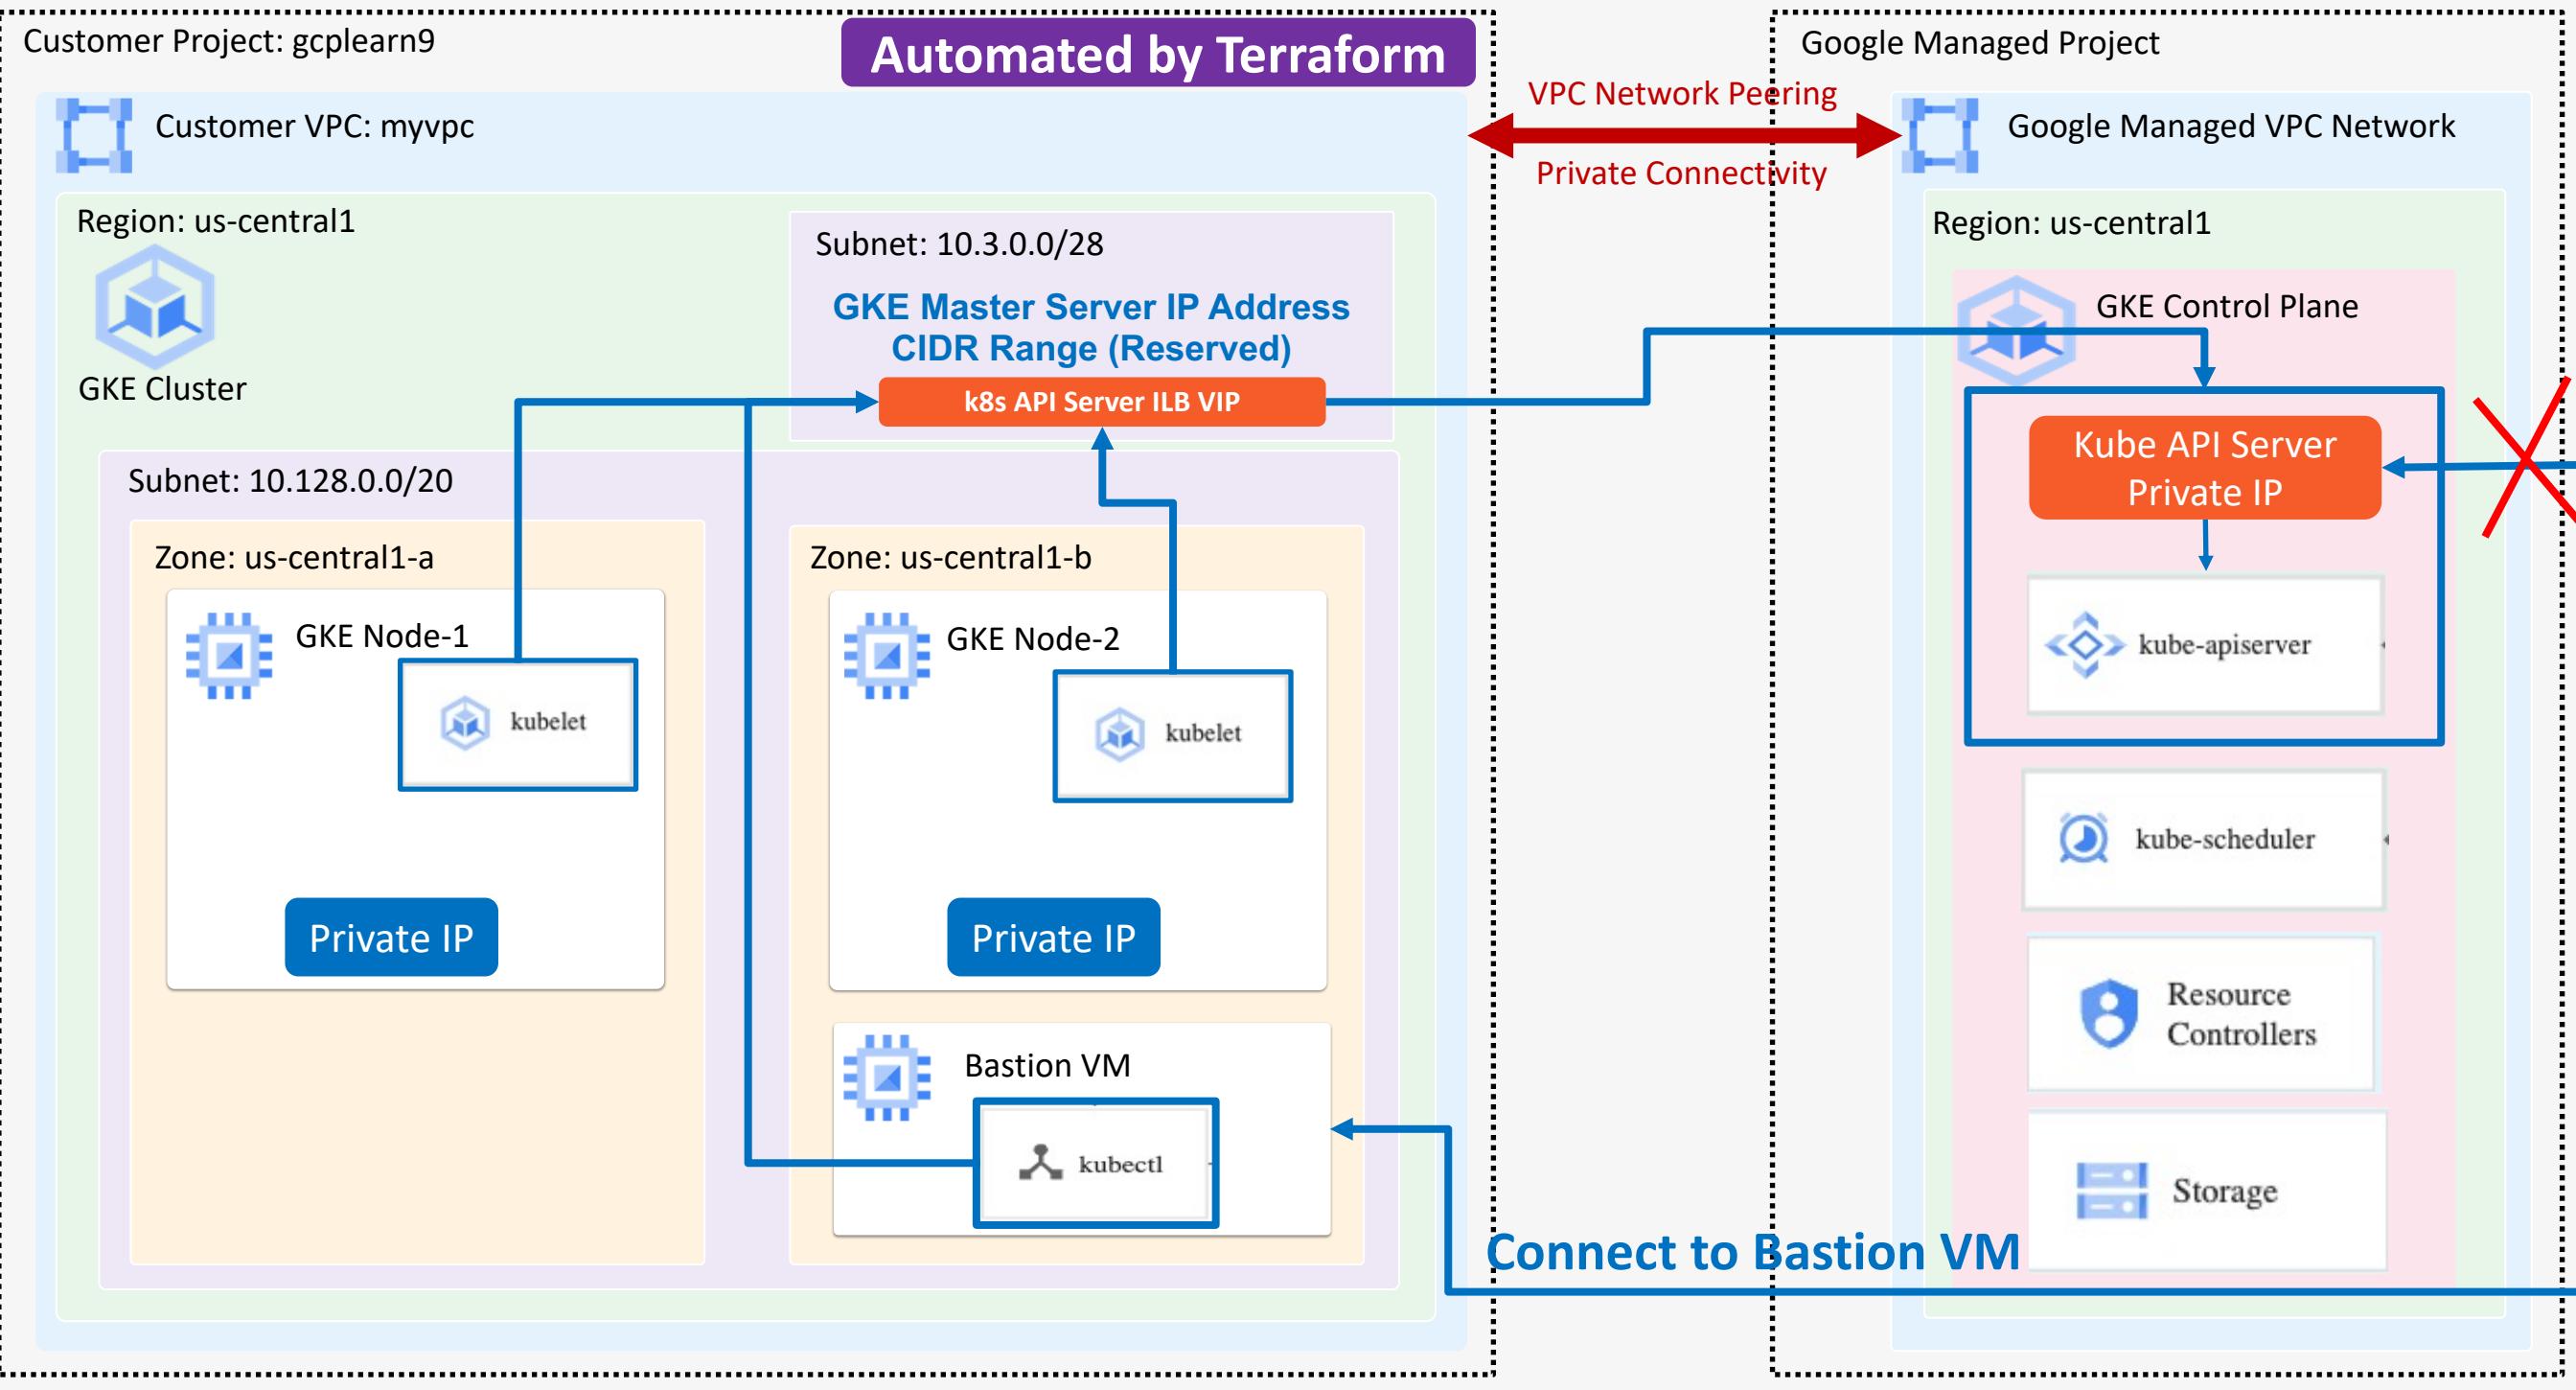
kubectl



Internet

NETWORK DESIGN

# GKE Standard Private Cluster with Private Endpoint



- ✓ p1-gke-private-cluster-private-endpoint
  - └ c1-versions.tf
  - └ c2-01-variables.tf
  - └ c2-02-local-values.tf
  - └ c3-vpc.tf
  - └ c4-firewallrules.tf
  - └ c5-datasource.tf
  - └ c6-01-gke-service-account.tf
  - └ c6-02-gke-cluster.tf
  - └ c6-03-gke-linux-nodepool.tf
  - └ c6-04-gke-outputs.tf
  - └ c7-Cloud-NAT-Cloud-Router.tf
  - └ c8-bastion-vm.tf
  - └ terraform.tfvars

# What are we going to learn?

Demo: GKE Standard Private Cluster + Private k8s API Server  
EnpPoint

NC

No Changes from C1 to C3

C4

Create Firewall rule with source as IAP IP Range  
(Bastion VM accessible using IAP only)

C6-02

GKE cluster – Enable Private Endpoint setting

C6-03

Comment Tags where earlier SSH firewall referenced

C8

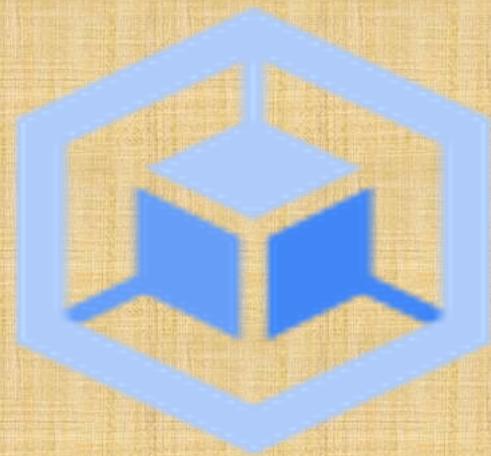
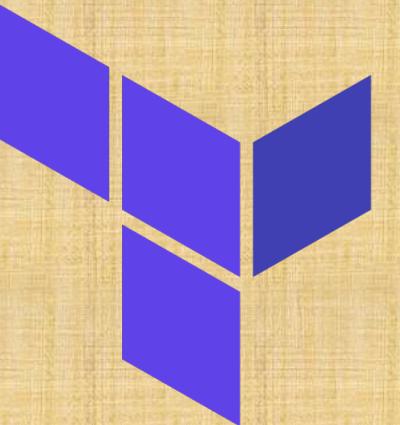
Bastion VM Instance

Demo-16

# Google Kubernetes Engine

Create GKE Autopilot Private cluster  
with k8s API Server Public endpoint

**Demo: GKE Autopilot Private Cluster + k8s API Server  
Public API endpoint**



# GKE Autopilot Cluster

Autopilot mode		Standard mode
Optimized Kubernetes cluster with a hands-off experience		Kubernetes cluster with node configuration flexibility
<a href="#">CONFIGURE</a>	<a href="#">TRY THE DEMO</a>	<a href="#">CONFIGURE</a>
<a href="#">TRY THE DEMO</a>		
Scaling	Automatic based on workload	You configure scaling
Nodes	Google manages and configures your nodes	You manage and configure your nodes >
Configuration	Streamlined configuration ready to use	You can configure all options
Workloads supported	Most workloads except <a href="#">these limitations</a>	All Kubernetes workloads
Billing method	<a href="#">Pay per pod</a>	<a href="#">Pay per node (VM)</a>
SLA	<a href="#">Kubernetes API and node availability</a>	<a href="#">Kubernetes API availability</a>

Complete Comparison: <https://cloud.google.com/kubernetes-engine/docs/resources/autopilot-standard-feature-comparison>

Author: Nho Luong

Skill: DevOps Engineer Lead

# GKE Autopilot Cluster

**Pricing:** Pay per pod, Pay for the CPU, memory and storage used by our workloads

**Pricing:** Not charged for system workloads, unused capacity on nodes, Operating System Costs

**Security:** Clusters have a super hardened configs with many security settings enabled by default.

**Security:** GKE automatically applies the security patches.

**Node Management:** Google manages Worker Nodes (Creation, Upgrades, Repairs)

**Networking:** very advanced. All pod traffic passes through VPC Firewalls.

GKE  
Autopilot  
Cluster

**Scaling:** When high load, GKE provisions new nodes for those pods (Node Autoprovisioning)

**Resource Management:** If we don't specify resource values in workloads, Autopilot sets pre-configured default values

**Release Management:** All Autopilot clusters are enrolled in GKE release channel which ensures control plane and nodes run on latest versions in that channel

**Managed Flexibility:** Autopilot supports specific hardware requirements (Compute Classes: General Purpose, Balanced, Scale-Out)

**Reduced Operational Complexity:** Autopilot reduces platform administration overhead by removing need to continuously monitor nodes, scaling and scheduling operations

# Autopilot Cluster

OVERVIEW      OBSERVABILITY      COST OPTIMIZATION

Filter Enter property name or value      ?      III

Status	Name ↑	Location	Mode	Number of nodes	Total vCPUs	Total memory
✓	autopilot-cluster-private-1	us-central1	Autopilot		0	0 GB

Observation: Once the Autopilot cluster is created, if we don't deploy any workloads, after some time Number of nodes, Total vCPUs and Total Memory all will come to zero.

- ✓ p1-gke-autopilot-cluster-private
- ✗ c1-versions.tf
- ✗ c2-01-variables.tf
- ✗ c2-02-local-values.tf
- ✗ c3-vpc.tf
- ✗ c4-01-gke-cluster.tf
- ✗ c4-02-gke-outputs.tf
- ✗ c5-Cloud-NAT-Cloud-Router.tf
- ✗ terraform.tfvars

# What are we going to learn?

Demo: GKE Autopilot Private Cluster + k8s API Server Public Endpoint

NC

No Changes from C1 to C3, C5, C4-02

C4-01

Create GKE Autopilot private cluster with  
`enable_autopilot = true`

# What are we going to learn?

- ✓ p2-k8resources-yaml
  - ! 01-kubernetes-deployment.yaml
  - ! 02-kubernetes-loadbalancer-service.yaml
- ✓ p3-k8resources-terraform-manifests
  - ⚡ c1-versions.tf
  - ⚡ c2-01-variables.tf
  - ⚡ c2-02-local-values.tf
  - ⚡ c3-01-remote-state-datasource.tf
  - ⚡ c3-02-providers.tf
  - ⚡ c4-kubernetes-deployment.tf
  - ⚡ c5-kubernetes-loadbalancer-service.tf
  - ⚡ terraform.tfvars

## Demo: Project-2: YAML Manifests

NC

01

No Changes in 02

Kubernetes Deployment: Add Container Resource Request and Limits (Optional but recommended)

## Demo: Project-3: Terraform Manifests

NC

C4

No Changes from C1 to C3-02 and C5

Kubernetes Deployment: Add Container Resource Request and Limits (Optional but recommended)

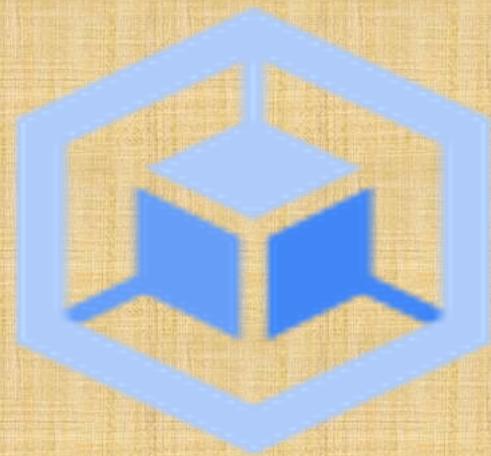
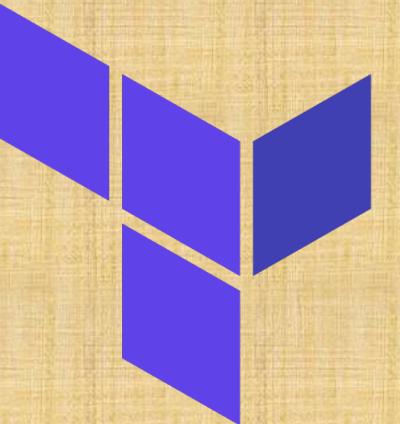
Demo-17

# Google Kubernetes Engine

Kubernetes Storage

Compute Engine Persistent Disks

**Demo: GKE Autopilot cluster + Compute Engine  
persistent disks as storage**



# Kubernetes Storage - Key Resources



**Kubernetes  
Storage Class**



**Kubernetes  
Persistent Volume**



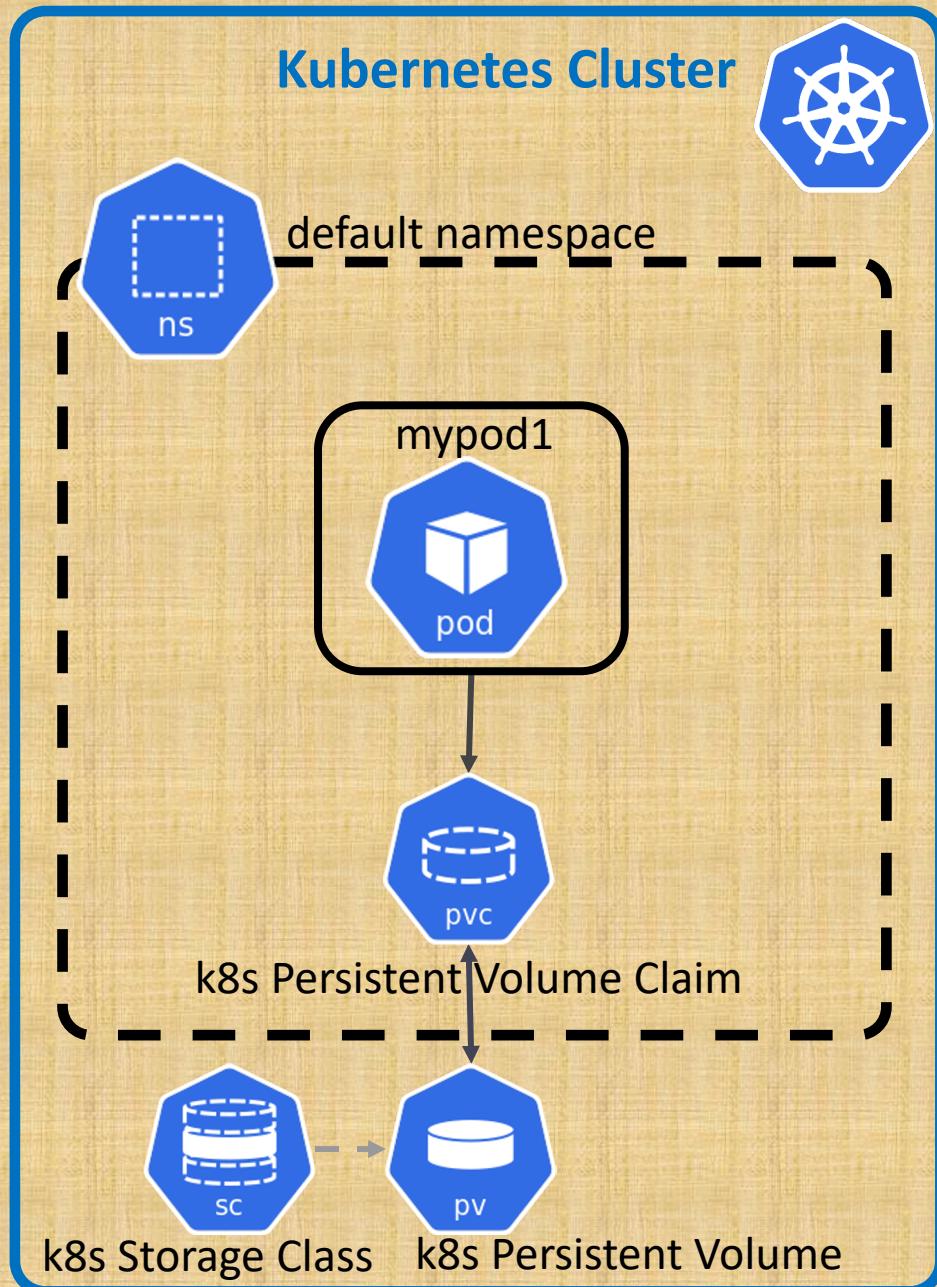
**Kubernetes  
Persistent Volume Claim**

# Kubernetes Storage Class

- **Storage Class:** Provides a way for administrators to define the **classes of storage they offer**
  - **standard-rwo:** Provides balanced disks
  - **premium-rwo:** Provides SSD disks

```
Kalyans-Mac-mini:google-kubernetes-engine kalyanreddy$ kubectl describe sc premium-rwo
Name:          premium-rwo
IsDefaultClass: No
Annotations:   components.gke.io/component-name=pdcsi,components.gke.io/component-version=0.13.2,components.gke.io/layer=addon
Provisioner:   pd.csi.storage.gke.io
Parameters:    type=pd-ssd
AllowVolumeExpansion: True
MountOptions:  <none>
ReclaimPolicy: Delete
VolumeBindingMode: WaitForFirstConsumer
Events:        <none>

Kalyans-Mac-mini:google-kubernetes-engine kalyanreddy$ kubectl describe sc standard-rwo
Name:          standard-rwo
IsDefaultClass: Yes
Annotations:   components.gke.io/layer=addon,storageclass.kubernetes.io/is-default-class=true
Provisioner:   pd.csi.storage.gke.io
Parameters:    type=pd-balanced
AllowVolumeExpansion: True
MountOptions:  <none>
ReclaimPolicy: Delete
VolumeBindingMode: WaitForFirstConsumer
Events:        <none>
```



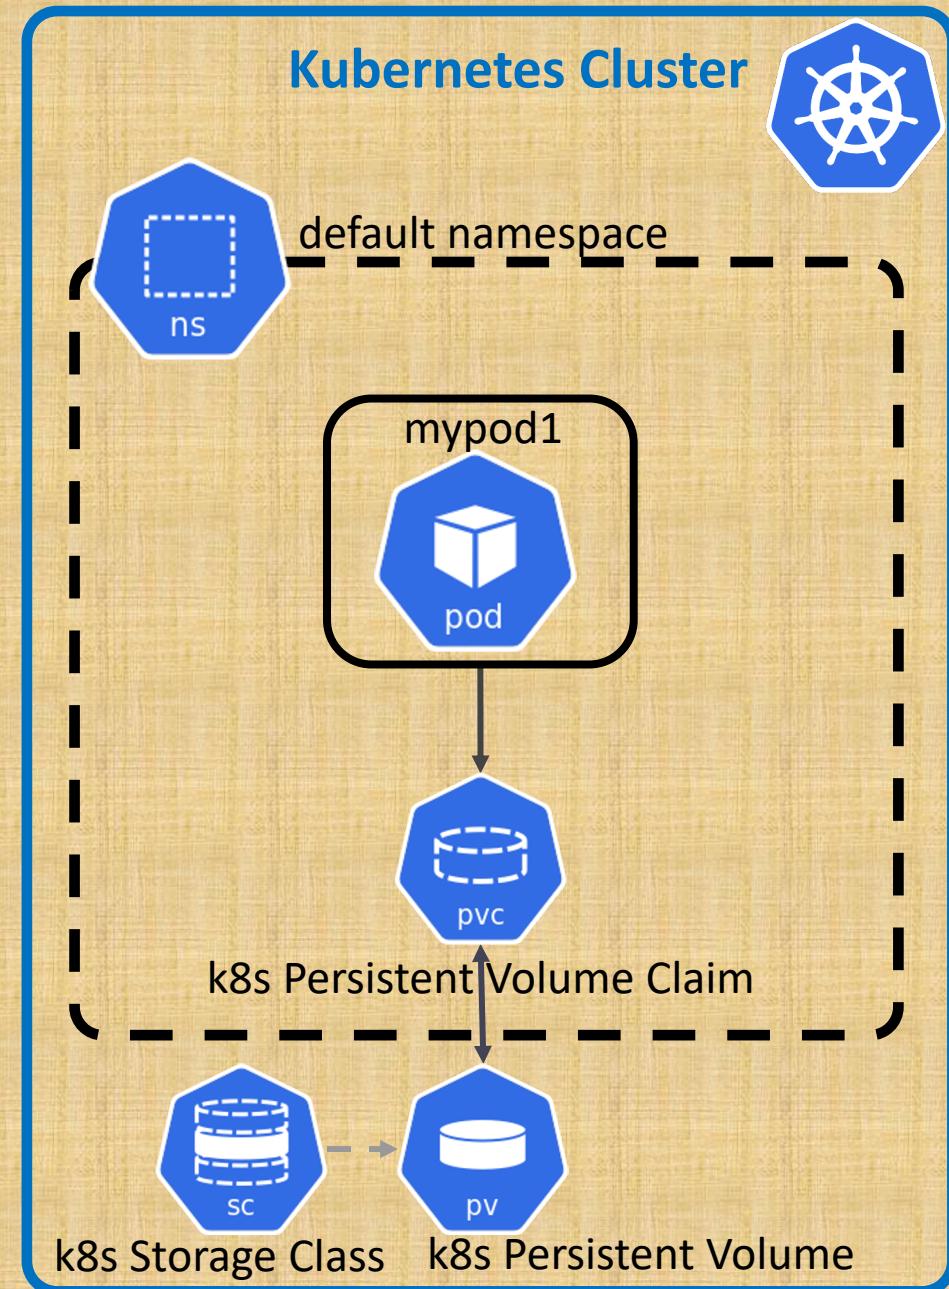
# Kubernetes Persistent Volume Claim

- **Persistent Volume Claim (PVC)**: PVC objects request a **specific size, access mode, and StorageClass** for the PersistentVolume
- If a PersistentVolume that satisfies the request **exists or can be provisioned**, the PVC is bound to that PersistentVolume.
- **Access Modes**
  - **ReadWriteOnce**: The volume can be mounted as **read-write by a single node**.
  - **ReadOnlyMany**: The volume can be mounted **read-only by many nodes**.
  - **ReadWriteMany**: The volume can be mounted as **read-write by many nodes**.
    - PersistentVolume resources that are backed by **Compute Engine persistent disks don't support** this access mode.

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: mypvc1
spec:
  accessModes:
    - ReadWriteOnce
  storageClassName: standard-rwo
  resources:
    requests:
      storage: 1Gi
```

# Kubernetes Persistent Volume

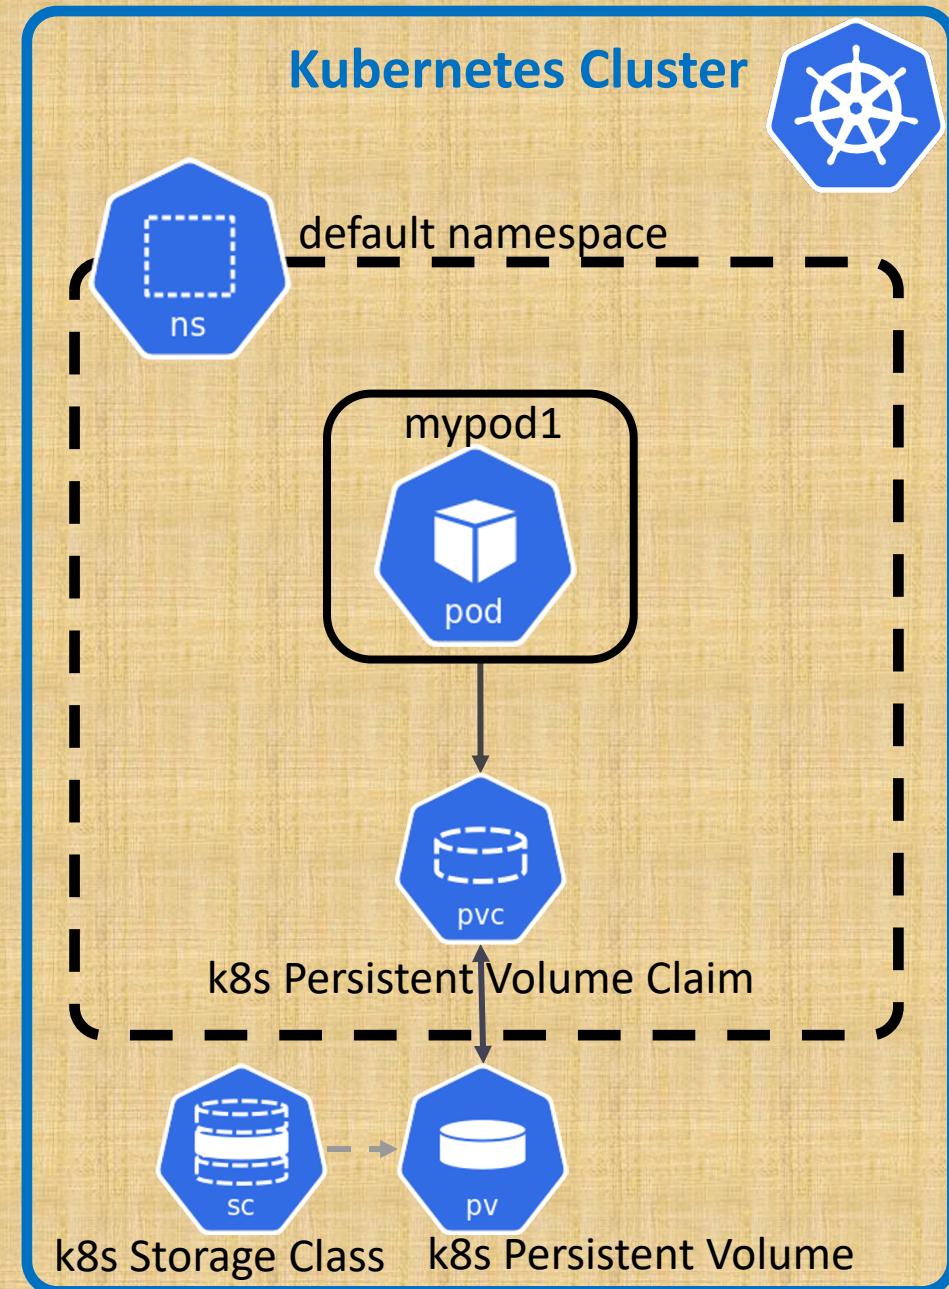
- **Persistent Volume:** In GKE, a `PersistentVolume` is typically backed by a [Google persistent disk](#)
- A `PersistentVolume` can be [dynamically provisioned](#), we do not have to [manually](#) create and delete the backing storage
- `PersistentVolume` resources are cluster resources that exist [independently](#) of Pods.
  - Doesn't impact due to [cluster changes](#)
  - Doesn't impact due to [pods deleted or recreated](#)
- Persistent Volume will be [provisioned based on configurations](#) defined in Storage Class and Persistent Volume Claim
  - **Storage Class:** Storage type (pd-balanced, pd-ssd)
  - **PVC:**
    - **Size:** 1GB, 2GB
    - **Access Modes:** ReadWriteOnce, ReadOnlyMany, ReadWriteMany



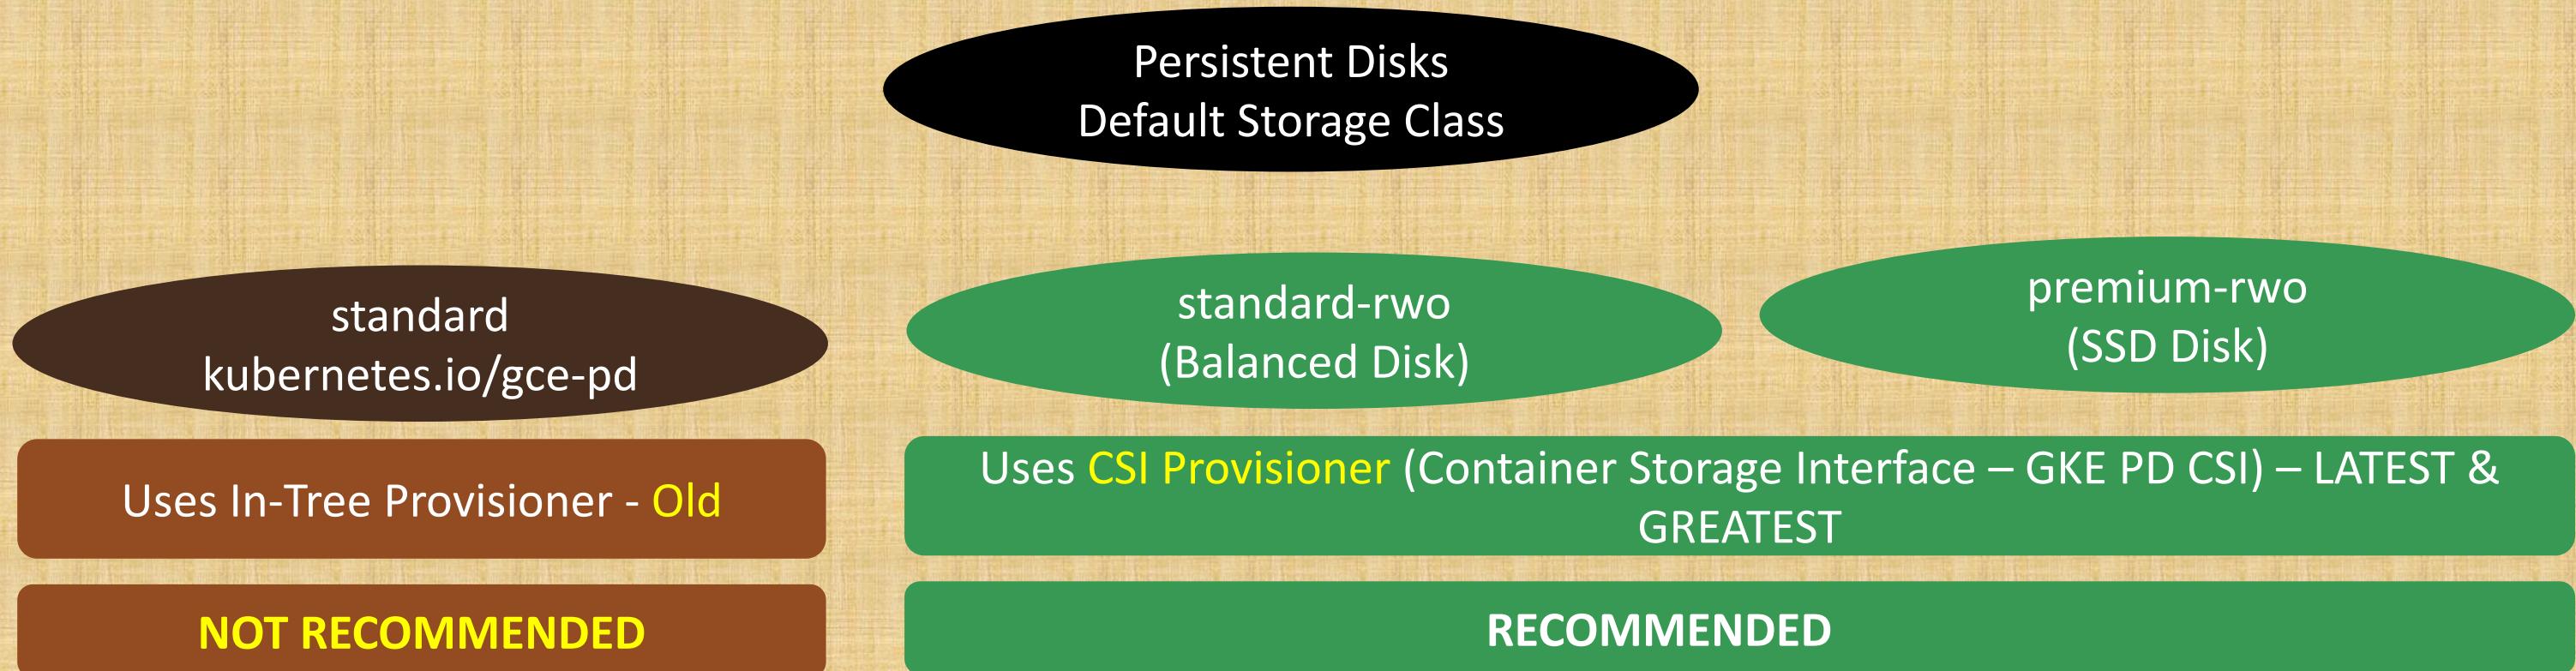
# Kubernetes Persistent Volume

- **Resource Types [IMPORTANT]**

- Storage Class and Persistent Volume or **cluster level resources**
- Persistent Volume Claim (PVC) is a **namespace level** resource



# GKE Cluster - Default Kubernetes Storage Classes



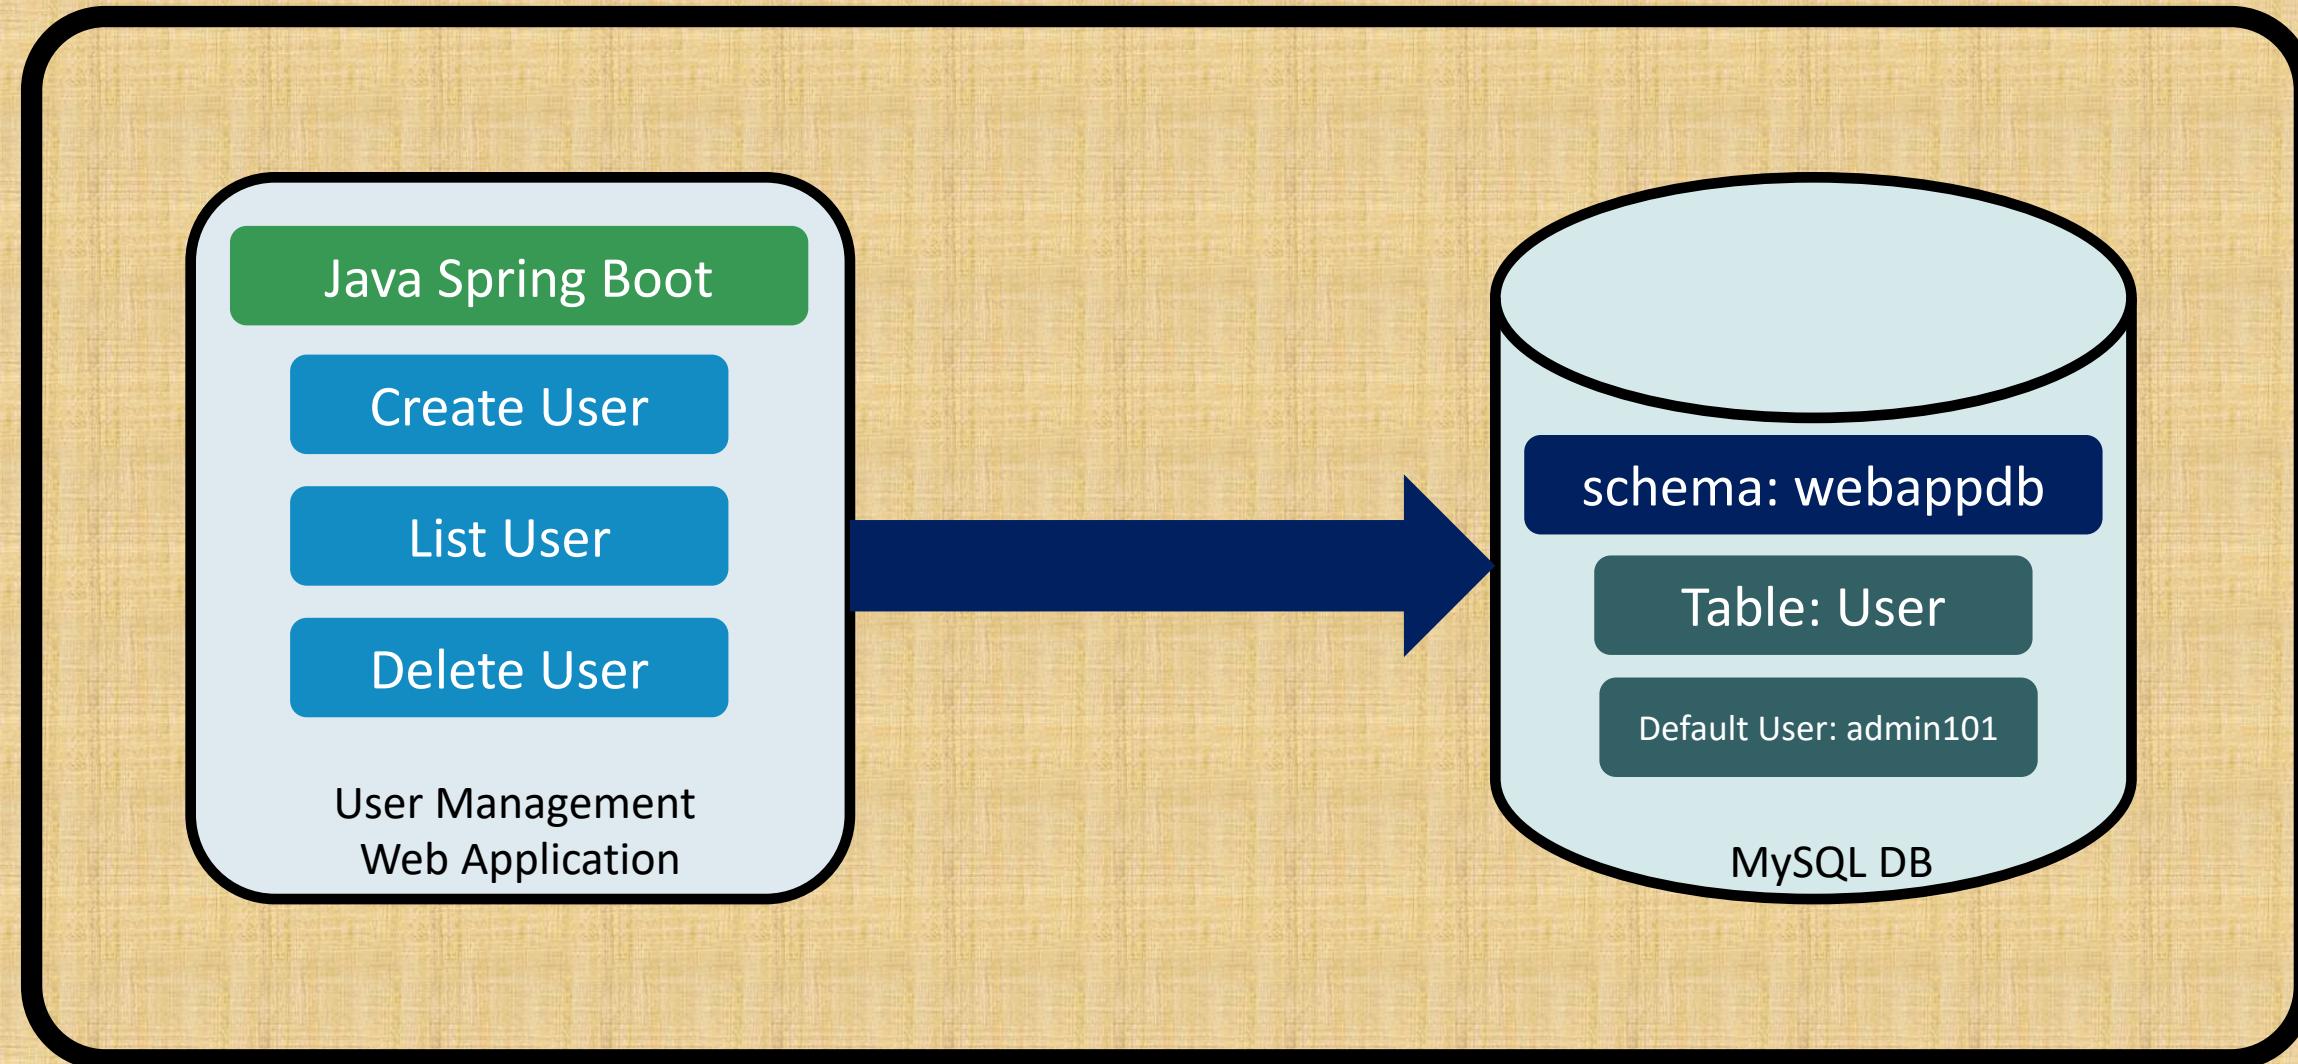
NAME	PROVISIONER	RECLAIMPOLICY	VOLUMEBINDINGMODE	ALLOWVOLUMEEXPANSION	AGE
premium-rwo	pd.csi.storage.gke.io	Delete	WaitForFirstConsumer	true	27h
standard	kubernetes.io/gce-pd	Delete	Immediate	true	27h
standard-rwo (default)	pd.csi.storage.gke.io	Delete	WaitForFirstConsumer	true	27h

Cluster Feature  
should be enabled

Compute Engine persistent disk CSI Driver

Enabled

# User Management Web Application - Architecture



# GKE Storage with Compute Engine Persistence Disks

## GKE Storage Compute Engine Persistent Disk CSI Driver



Compute Engine Persistent Disk

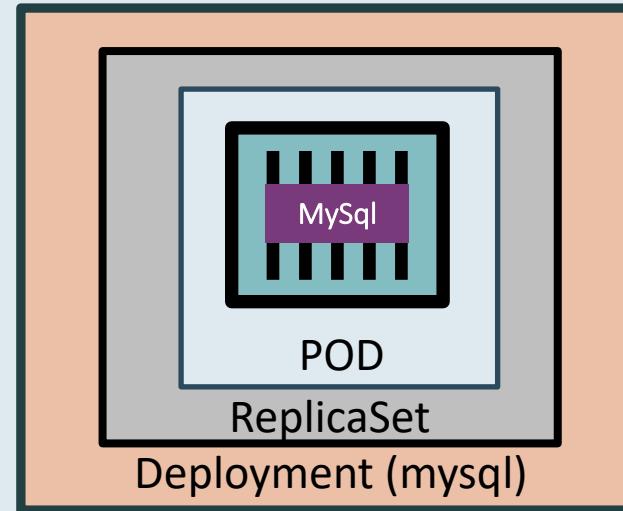
Storage Class

Persistent Volume Claim

Config Map

Deployment

MySQL – ClusterIP Service



YAML

Terraform

Environment Variables

Volumes

Volume Mounts

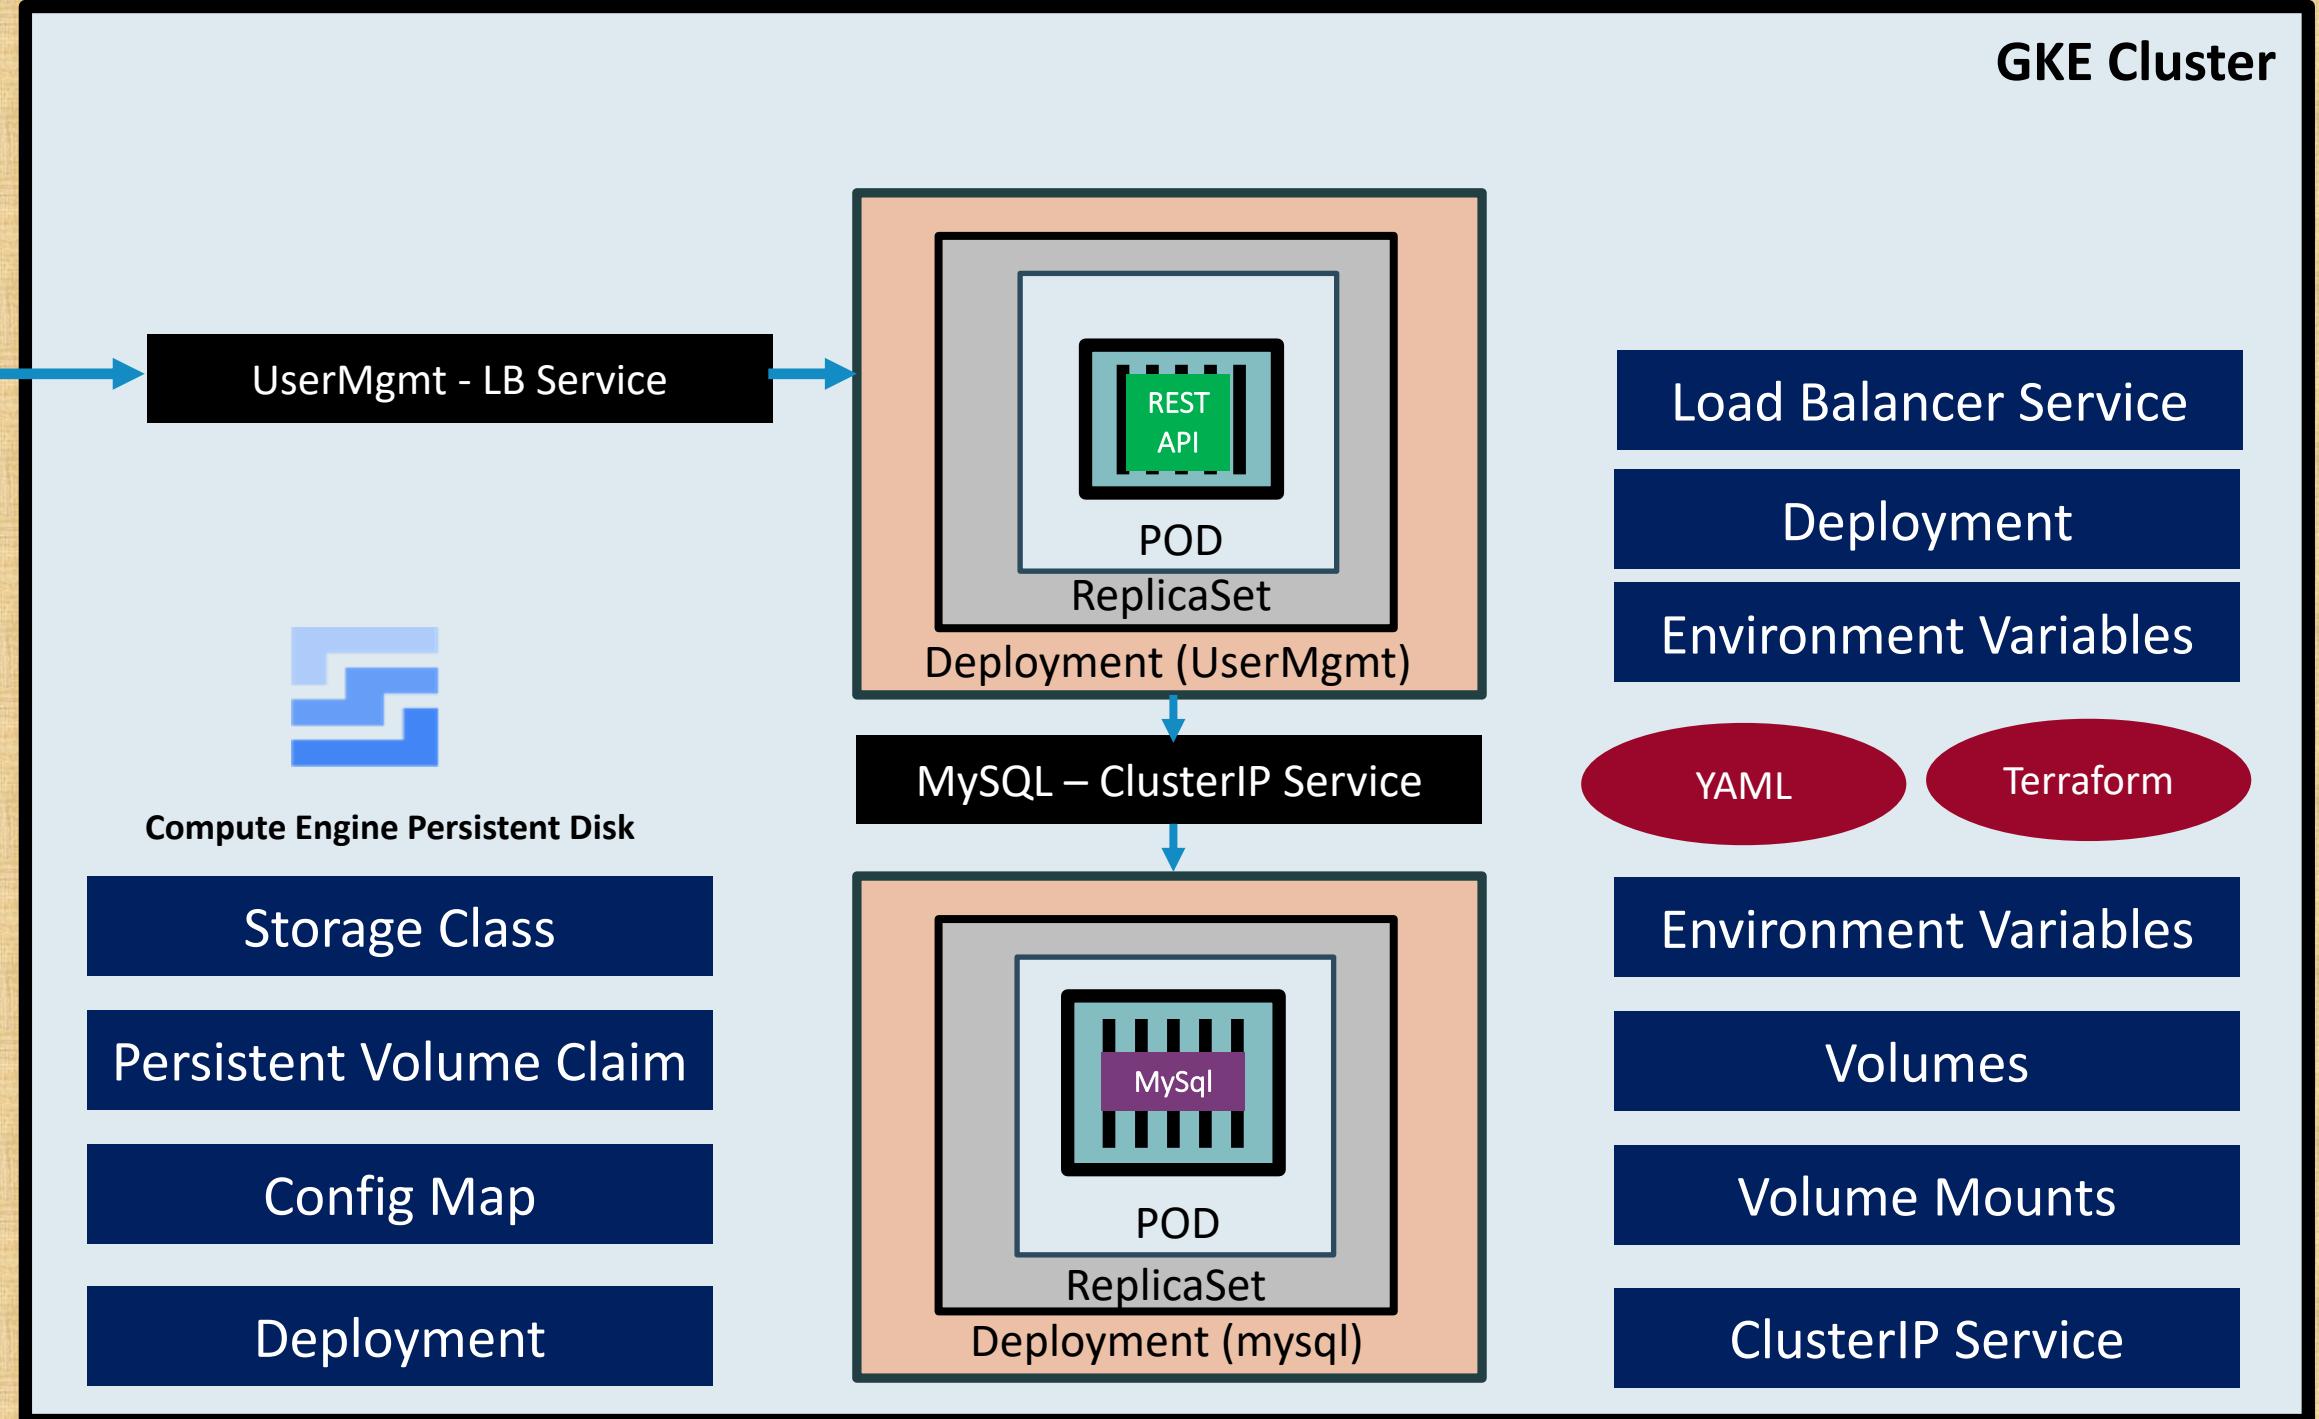
ClusterIP Service

# GKE Storage with Compute Engine Persistence Disks



http://<LB-IP>

GKE Storage  
Compute  
Engine  
Persistent Disk  
CSI Driver



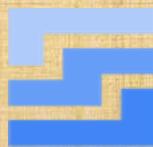
# GKE Storage with Compute Engine Persistence Disks



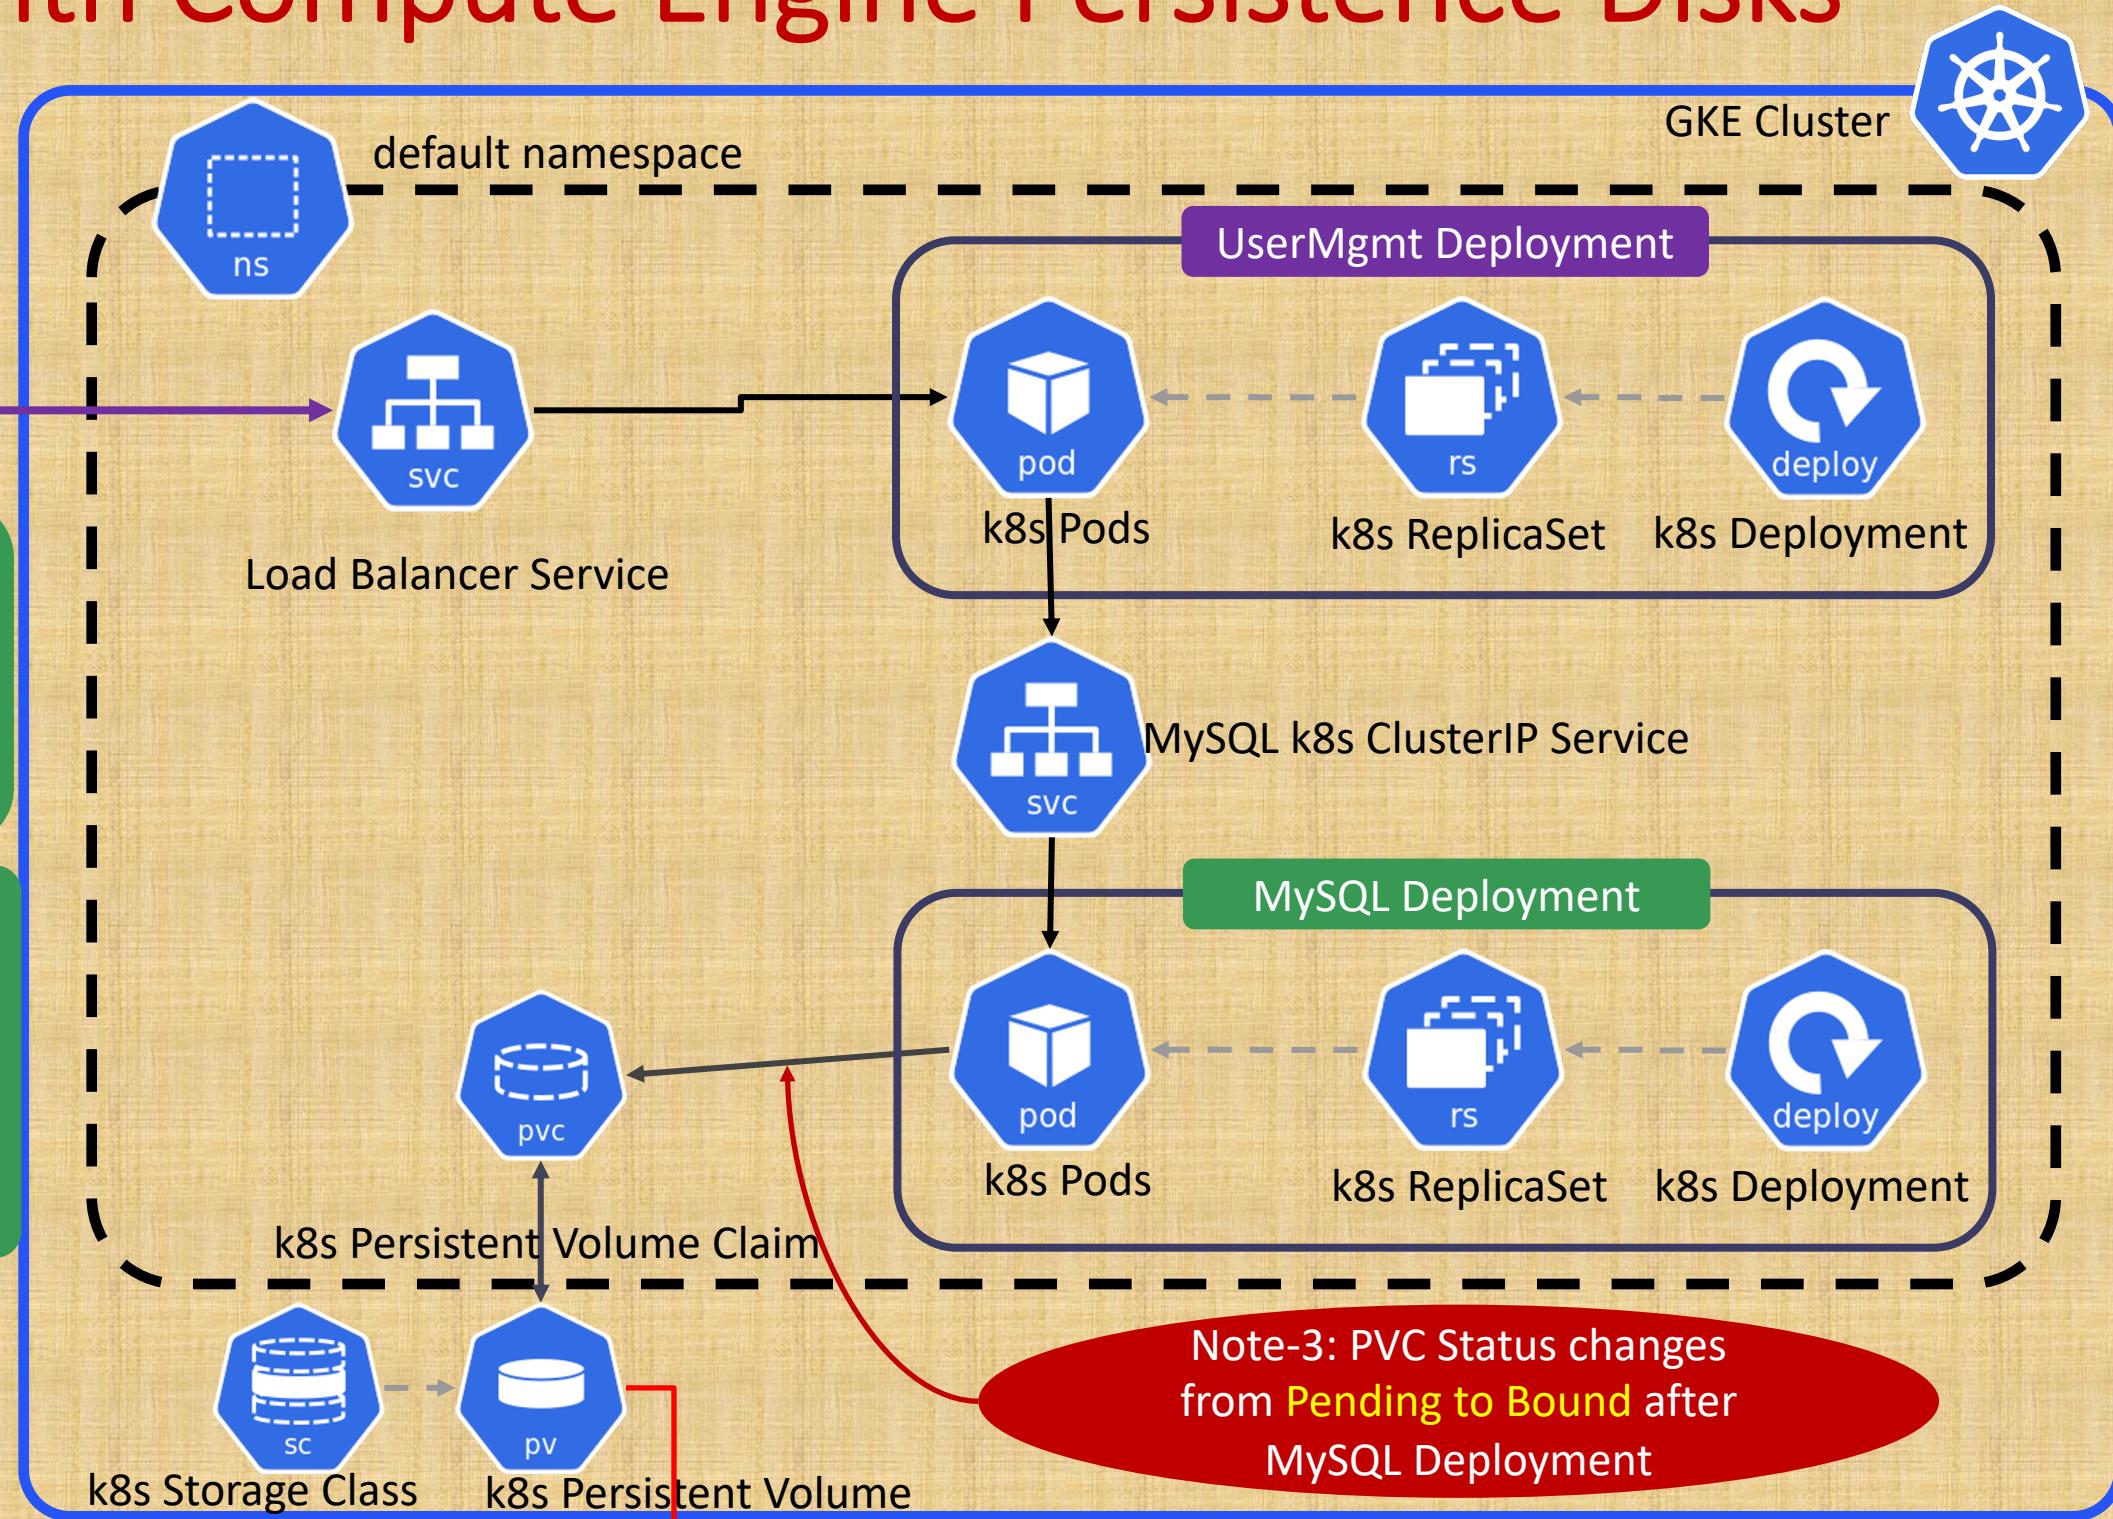
http://<LB-IP>

**Note-1:** PVC Claims must exist in the same namespace as the Pod using the claim. The cluster finds the claim in the Pod's namespace and uses it to get the PersistentVolume backing the claim

**Note-2:** A PersistentVolume (PV) is a piece of storage in the cluster that has been provisioned by an administrator or dynamically provisioned using Storage Classes. It is a resource in the cluster just like a node is a cluster resource



Compute Engine Persistent Disk



# What are we going to learn?

```
✓ p2-k8sresources-yaml
! 00-storage-class.yaml
! 01-persistent-volume-claim.yaml
! 02-UserManagement-ConfigMap.yaml
! 03-mysql-deployment.yaml
! 04-mysql-clusterip-service.yaml
! 05-UserMgmtWebApp-Deployment.yaml
! 06-UserMgmtWebApp-LoadBalancer-Service.yaml
```

## Demo: Project-2: YAML Manifests

00

Custom Kubernetes Storage Class

01

Kubernetes Persistent Volume Claim

02

Kubernetes ConfigMap

03

Kubernetes Deployment for MySQL Database

04

Kubernetes ClusterIP Service with clusterIP:None

05

User Management Web Application Kubernetes Deployment

06

User Management Web Application Kubernetes Load Balancer Service

```
✓ p3-k8sresources-terraform-manifests
  ✓ c1-versions.tf
  ✓ c2-01-variables.tf
  ✓ c2-02-local-values.tf
  ✓ c3-01-remote-state-datasource.tf
  ✓ c3-02-providers.tf
  ✓ c4-01-storage-class.tf
  ✓ c4-02-persistent-volume-claim.tf
  ✓ c4-03-UserMgmtWebApp-ConfigMap.tf
  ✓ c4-04-mysql-deployment.tf
  ✓ c4-05-mysql-clusterip-service.tf
  ✓ c4-06-UserMgmtWebApp-deployment.tf
  ✓ c4-07-UserMgmtWebApp-loadbalancer-service.tf
  ✓ terraform.tfvars
  ⚡ webappdb.sql
```

# What are we going to learn?

## Demo: Project-3: Terraform Manifests

NC

c4-01

c4-02

c4-03

c4-04

c4-05

c4-06

No Changes from C1 to C3-02

Custom Kubernetes Storage Class

Kubernetes Persistent Volume Claim

Kubernetes ConfigMap

Kubernetes Deployment for MySQL Database

Kubernetes ClusterIP Service with clusterIP:None

User Management Web Application Kubernetes Deployment

```
✓ p3-k8sresources-terraform-manifests
  ✓ c1-versions.tf
  ✓ c2-01-variables.tf
  ✓ c2-02-local-values.tf
  ✓ c3-01-remote-state-datasource.tf
  ✓ c3-02-providers.tf
  ✓ c4-01-storage-class.tf
  ✓ c4-02-persistent-volume-claim.tf
  ✓ c4-03-UserMgmtWebApp-ConfigMap.tf
  ✓ c4-04-mysql-deployment.tf
  ✓ c4-05-mysql-clusterip-service.tf
  ✓ c4-06-UserMgmtWebApp-deployment.tf
  ✓ c4-07-UserMgmtWebApp-loadbalancer-service.tf
  terraform.tfvars
  webappdb.sql
```

# What are we going to learn?

## Demo: Project-3: Terraform Manifests

NC

c4-07

SQL

No Changes from C1 to C3-02

User Management Web Application Kubernetes  
Load Balancer Service

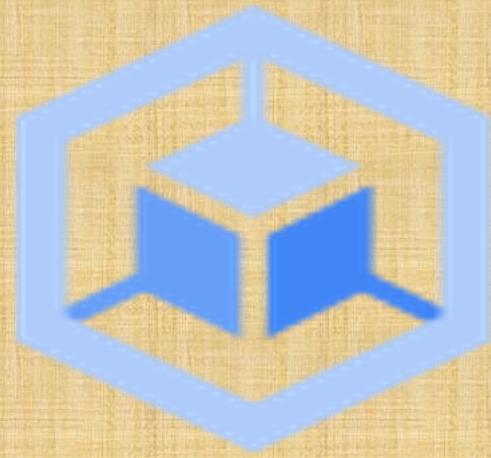
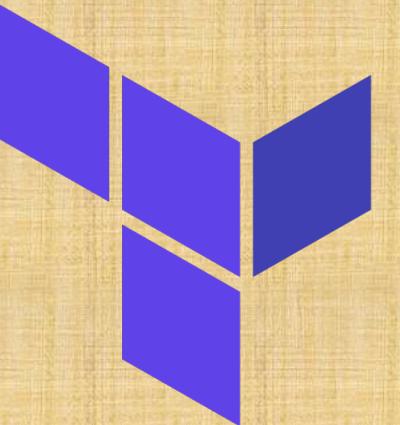
webappdb.sql file to create base database schema

Demo-18

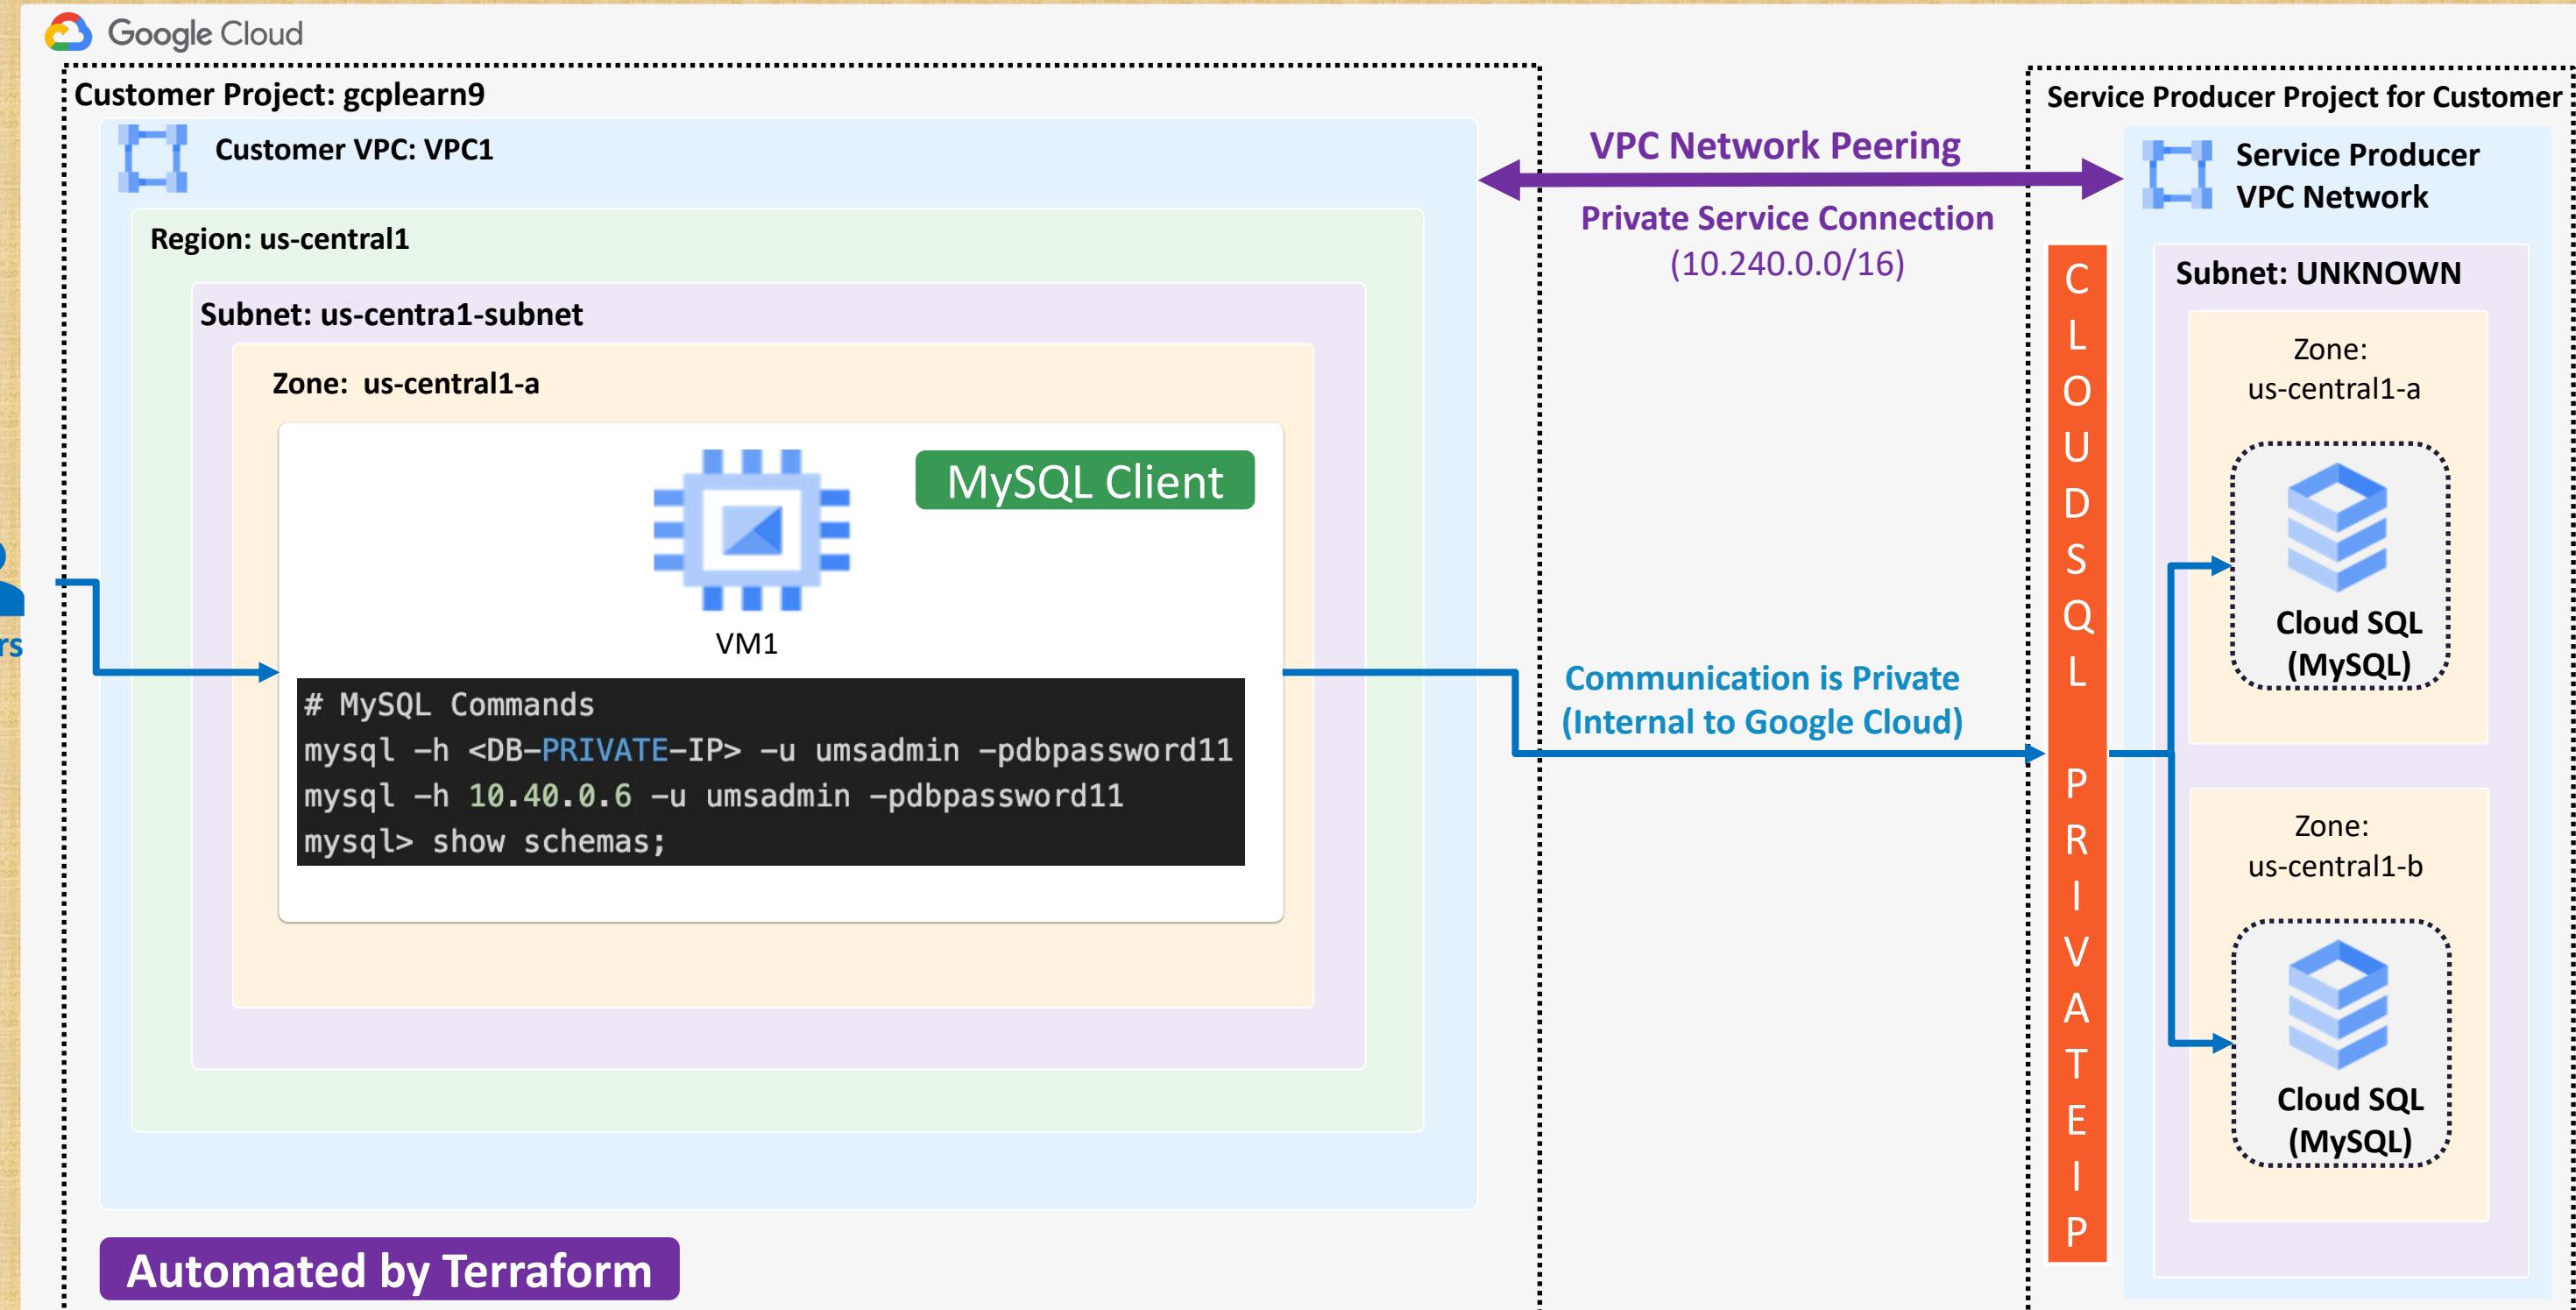
# Google Kubernetes Engine

## Cloud SQL MySQL Database as Storage for Applications Deployed on GKE Cluster

**Demo: GKE Autopilot cluster + Cloud SQL MySQL DB**



# GCP Cloud SQL Private IP



```
└─ p2-cloudsql-privatedb
    └─ c1-versions.tf
    └─ c2-01-variables.tf
    └─ c2-02-local-values.tf
    └─ c3-01-remote-state-datasource.tf
    └─ c3-02-private-service-connection.tf
    └─ c4-01-cloudsql.tf
    └─ c4-02-cloudsql-outputs.tf
    └─ c5-vminstance.tf
    $ mysql-client-install.sh
    └─ terraform.tfvars
```

# What are we going to learn?

## Demo: Project-2: Cloud SQL Private Database

c1

Define Remote Backend in Cloud Storage Bucket for this demo

c2-01

Define Cloud SQL MySQL Database version variable

c3-01

Define Remote State data source to access GKE cluster from Project-1

c3-02

Create VPC Network peering between our VPC and Google Managed Project VPC

c4-01

Create Cloud SQL Terraform Resources

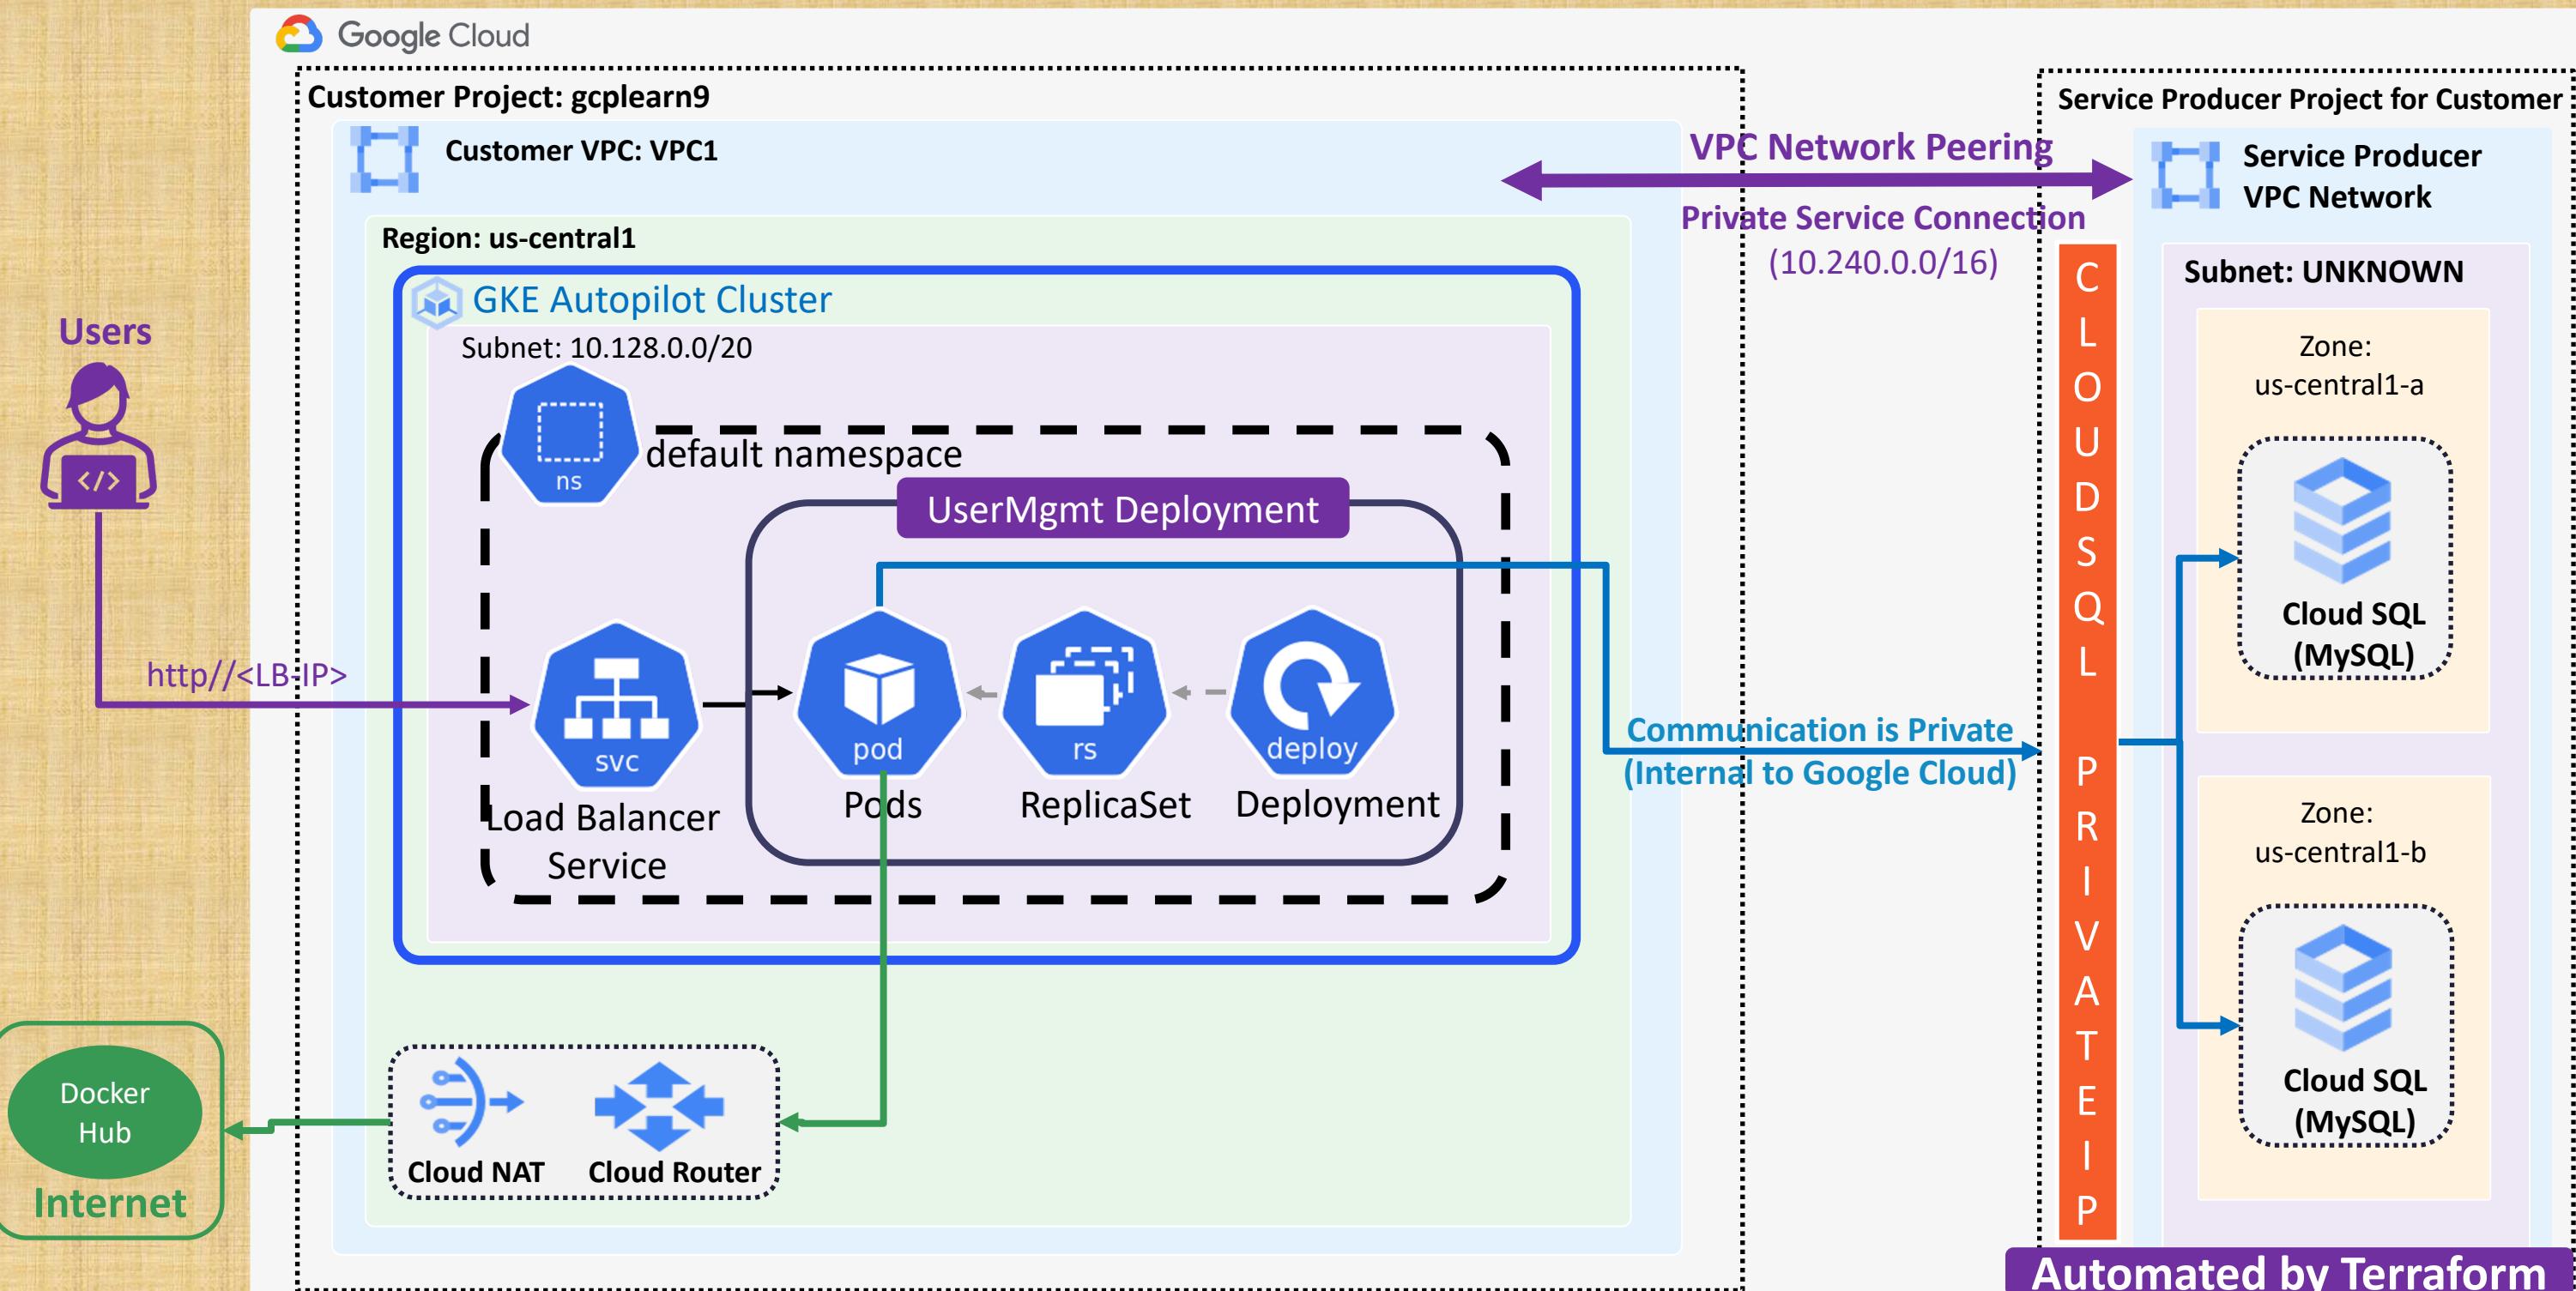
c4-02

Define Cloud SQL Outputs mainly Database IP, so that it can accessed in Project-3 UMS Web App

C5

VM Instance to test the DB created is accessible privately from our customer VPC

# GCP Cloud SQL Private IP + GKE Autopilot Cluster + User Management Web Application



## Project-3: Terraform Manifests

```
✓ p3-k8sresources-terraform-manifests
  ✓ c1-versions.tf
  ✓ c2-01-variables.tf
  ✓ c2-02-local-values.tf
  ✓ c3-01-remote-state-datasource.tf
  ✓ c3-02-providers.tf
  ✓ c4-06-UserMgmtWebApp-deployment.tf
  ✓ c4-07-UserMgmtWebApp-loadbalancer-service.tf
  ✓ terraform.tfvars
```

# What are we going to learn?

## Demo: Project-3: UMS Web Application using Cloud SQL MySQL Database



Define Remote Backend in Cloud Storage Bucket for this demo

Define Input variable

Define Remote State data source to access GKE cluster from Project-1 and Cloud SQL IP from Project-2

Define Kubernetes Provider

UMS WebApp k8s Deployment pointing to Cloud SQL MySQL Database

UMS WebApp Load Balancer Service

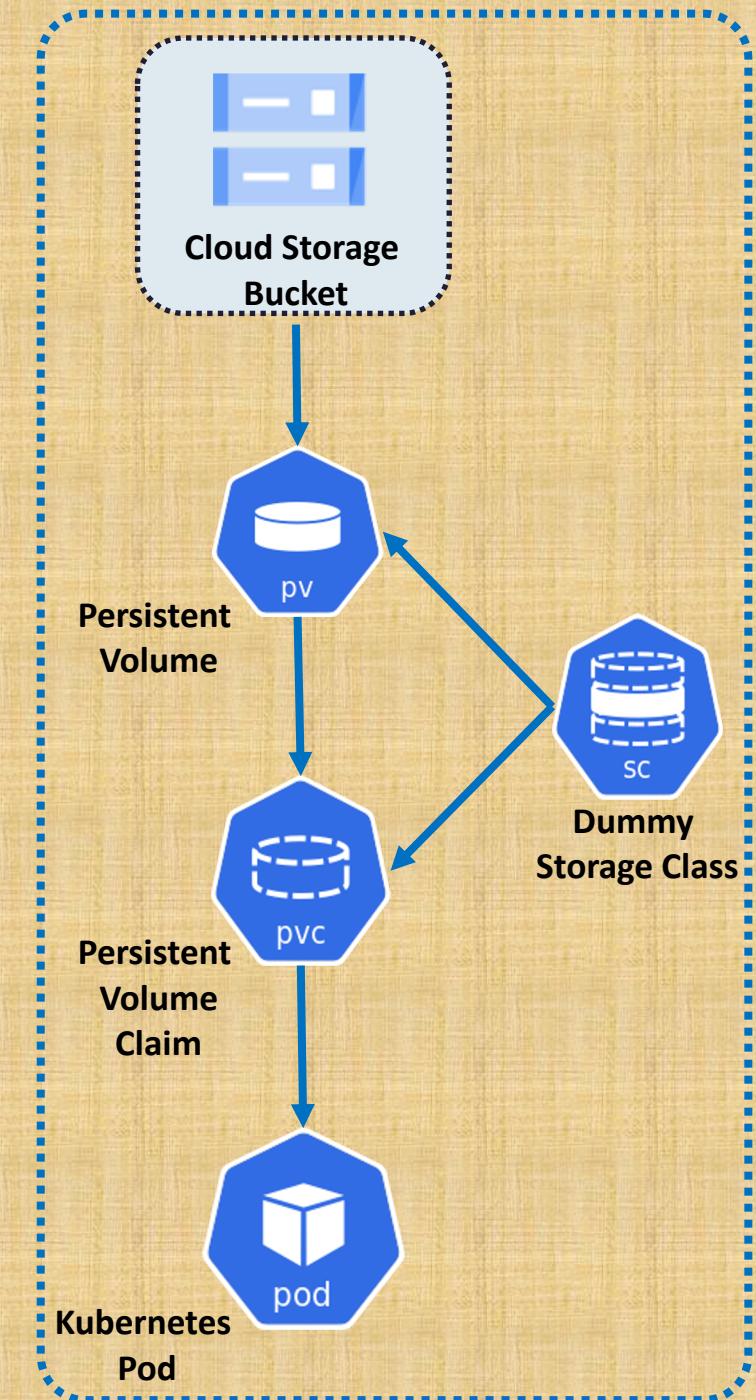
Demo-19

# Google Kubernetes Engine Cloud Storage FUSE CSI Driver

Demo: GKE Autopilot cluster + Cloud Storage Buckets

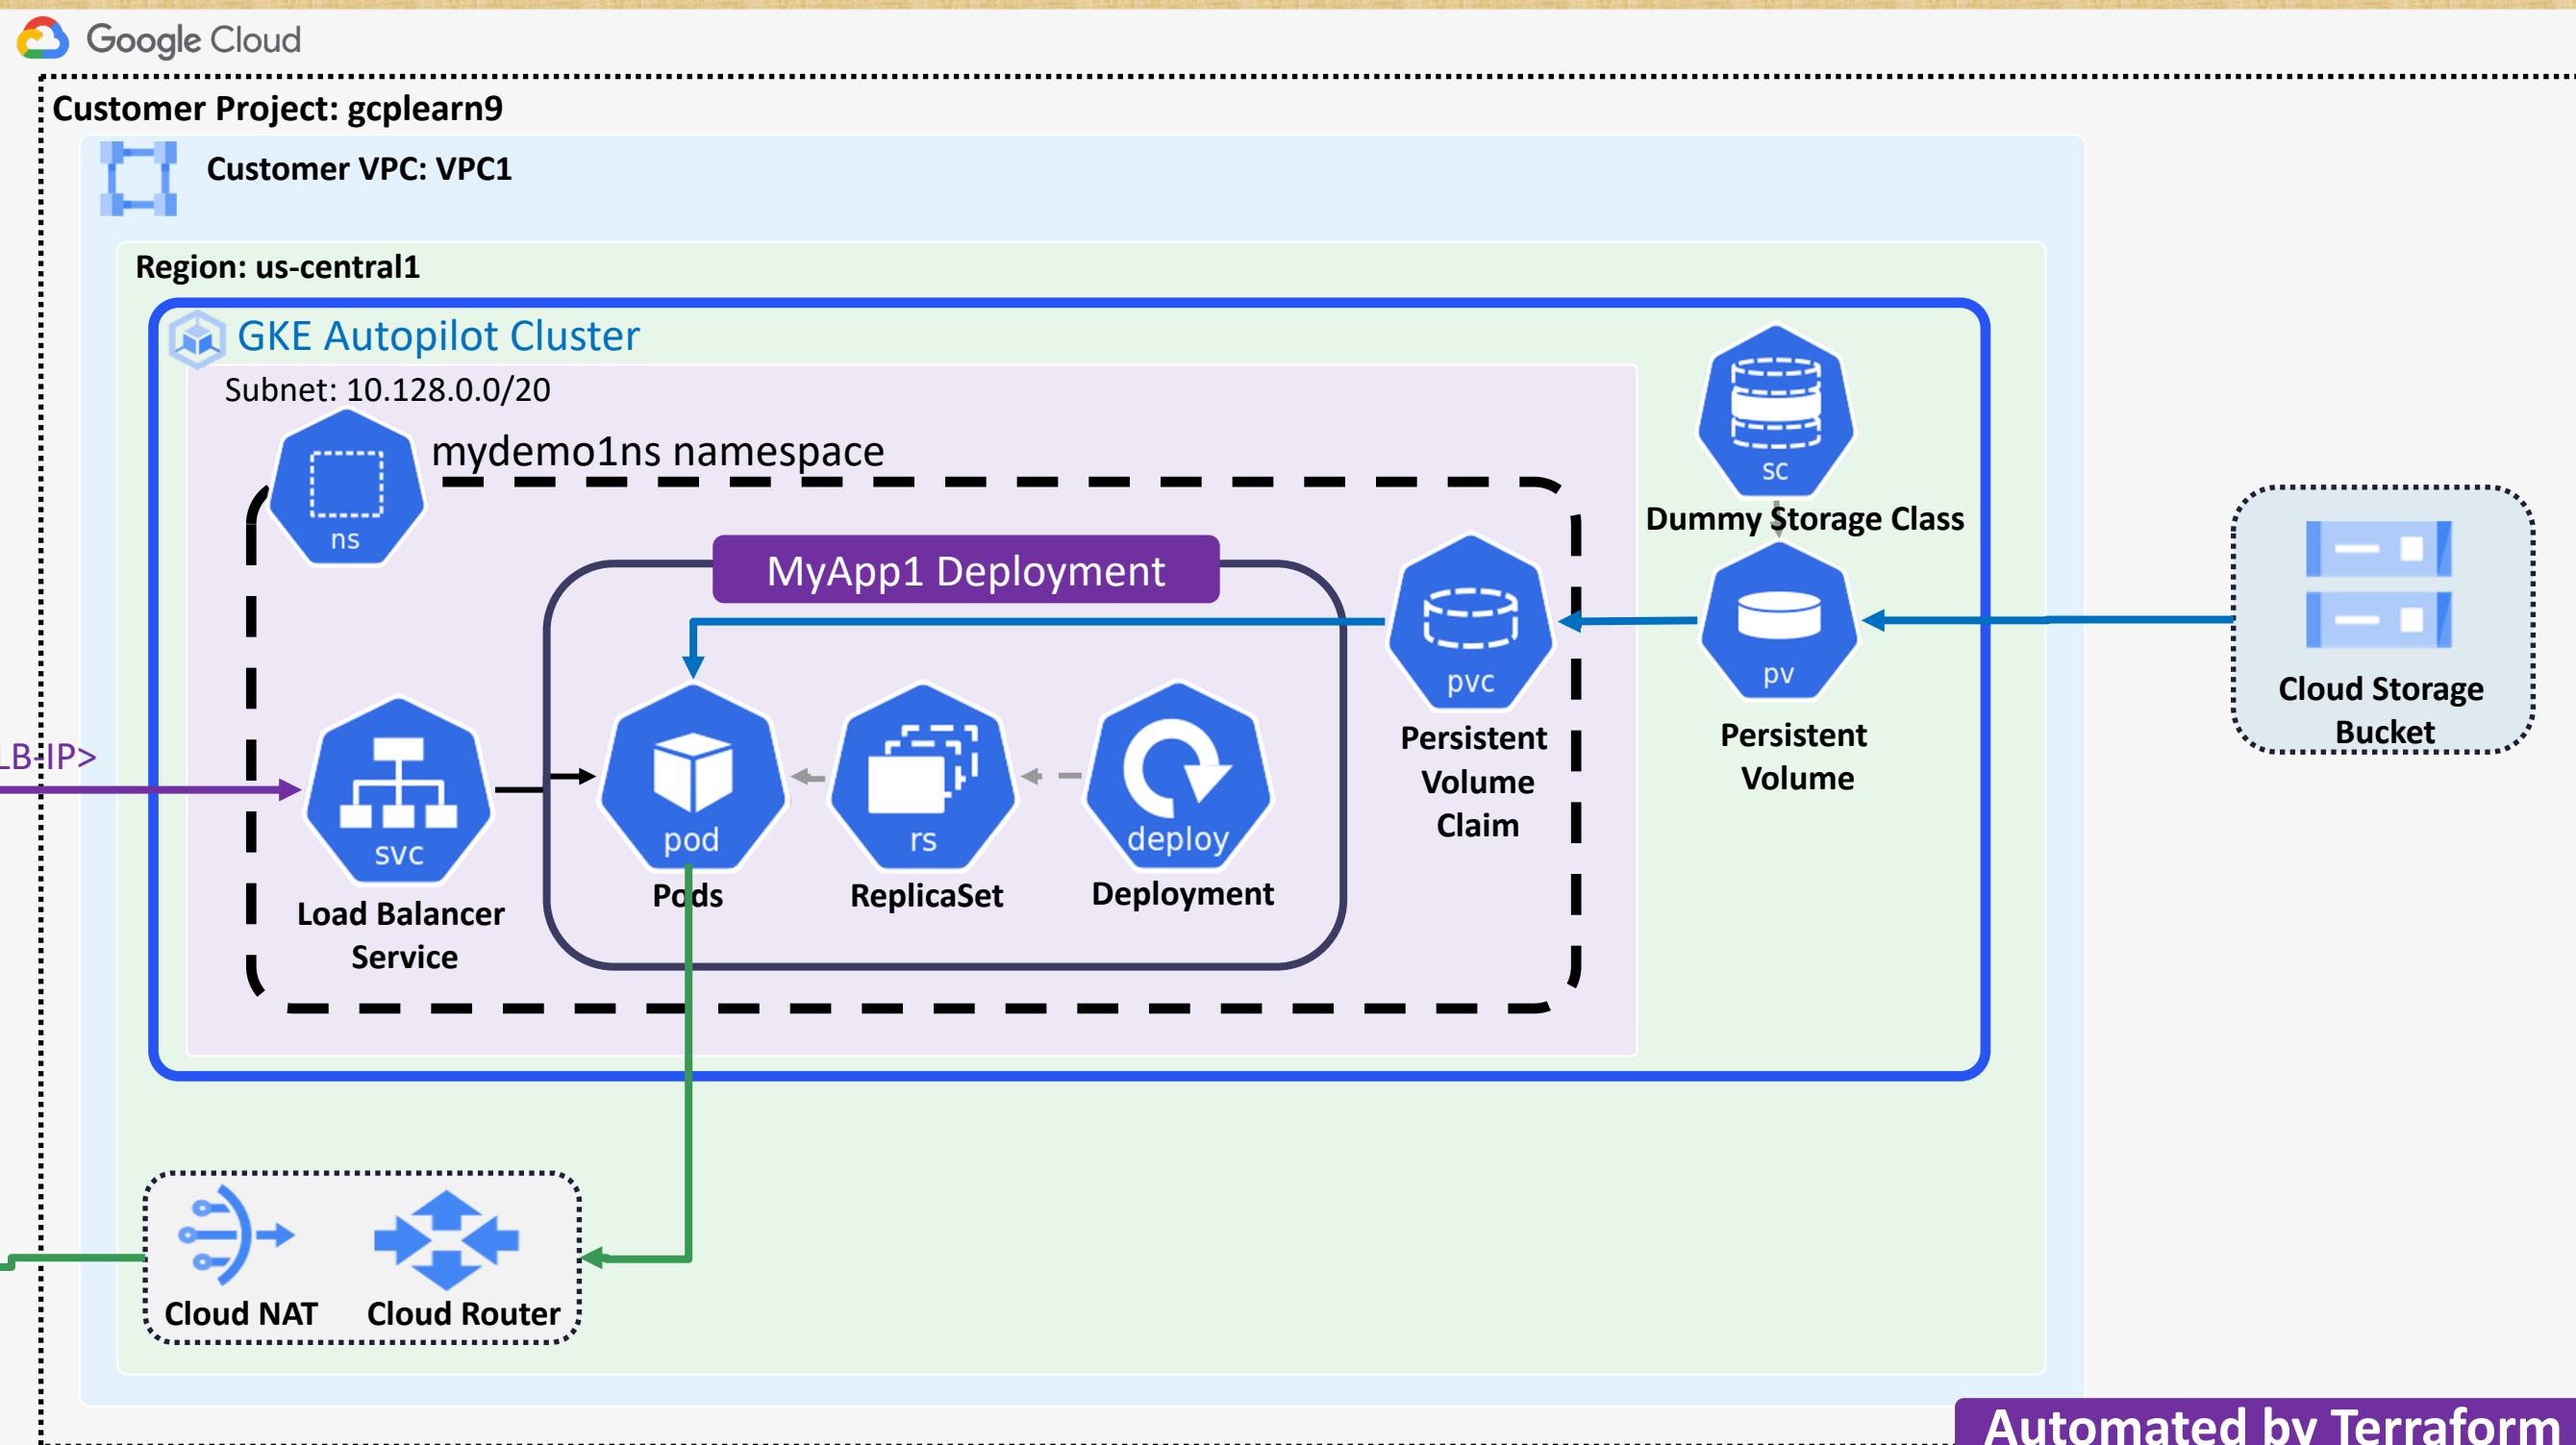
# Cloud Storage FUSE CSI Driver

- **FUSE: Filesystem in Userspace**
  - Fuse is an interface primarily used to **export a filesystem to the Linux kernel**
- **Cloud Storage FUSE CSI Driver**
  - Allow you to **mount Cloud Storage Buckets** as a file system
  - Applications can access the objects in a bucket using **common File IO operations** (e.g. open, read, write, close) rather than using cloud-specific APIs
  - In short, it lets you use the **Kubernetes API** to consume pre-existing Cloud Storage buckets as volumes
  - Driver can be deployed **automatically** just using **enable/disable** options
  - **Supports GKE Standard and Auto-pilot clusters**



Reference: <https://cloud.google.com/kubernetes-engine/docs/how-to/persistent-volumes/cloud-storage-fuse-csi-driver>

# GKE Autopilot Cluster + Cloud Storage Bucket as Volume Mount



Automated by Terraform

## Project-2: YAML Manifests

```
✓ p2-k8sresources-yaml
! c1-k8s-namespace.yaml
! c2-k8s-serviceaccount.yaml
! c3-k8s-pv.yaml
! c4-k8s-pvc.yaml
! c5-kubernetes-deployment.yaml
! c6-kubernetes-loadbalancer-service.yaml
```

Manual Steps  
(Required)

Cloud Storage Bucket

IAM Policy Binding

# What are we going to learn?

Demo: Project-2: GKE Autopilot cluster + Cloud Storage Buckets  
(GCS FUSE CSI)

c1

Kubernetes Namespace

c2

Kubernetes Service Account

c3

Kubernetes Persistent Volume using Cloud Storage Bucket

c4

Kubernetes Persistent Volume Claim

c5

Kubernetes Deployment mounted with Cloud Storage bucket as volume

c6

Kubernetes Load Balancer Service

## Project-3: Terraform Manifests

- ✓ p3-k8sresources-terraform-manifests
  - └ c1-versions.tf
  - └ c2-01-variables.tf
  - └ c2-02-local-values.tf
  - └ c3-01-remote-state-datasource.tf
  - └ c3-02-providers.tf
  - └ c4-01-cloud-storage-bucket.tf
  - └ c4-02-bucket-iam-binding.tf
  - └ c5-01-persistent-volume.tf
  - └ c5-02-namespace.tf
  - └ c5-03-persistent-volume-claim.tf
  - └ c5-04-service-account.tf
  - └ c5-05-myapp1-deployment.tf
  - └ c5-06-myapp1-loadbalancer-service.tf
  - └ terraform.tfvars

# What are we going to learn?

Demo: Project-3: GKE Autopilot cluster + Cloud Storage Buckets  
(GCS FUSE CSI)

NC

No changes c1 to c3

c4-01

Create Cloud Storage Bucket

c4-02

Cloud Storage Bucket IAM Binding

c5-01

Kubernetes Persistent Volume using Cloud Storage Bucket

c5-02

Kubernetes Namespace

c5-03

Kubernetes Persistent Volume Claim

c5-04

Kubernetes Service Account

## Project-3: Terraform Manifests

- ✓ p3-k8sresources-terraform-manifests
  - └ c1-versions.tf
  - └ c2-01-variables.tf
  - └ c2-02-local-values.tf
  - └ c3-01-remote-state-datasource.tf
  - └ c3-02-providers.tf
  - └ c4-01-cloud-storage-bucket.tf
  - └ c4-02-bucket-iam-binding.tf
  - └ c5-01-persistent-volume.tf
  - └ c5-02-namespace.tf
  - └ c5-03-persistent-volume-claim.tf
  - └ c5-04-service-account.tf
  - └ c5-05-myapp1-deployment.tf
  - └ c5-06-myapp1-loadbalancer-service.tf
  - └ terraform.tfvars

# What are we going to learn?

Demo: Project-3: GKE Autopilot cluster + Cloud Storage Buckets  
(GCS FUSE CSI)

c5-05

Kubernetes Deployment mounted with Cloud Storage bucket as volume

c5-06

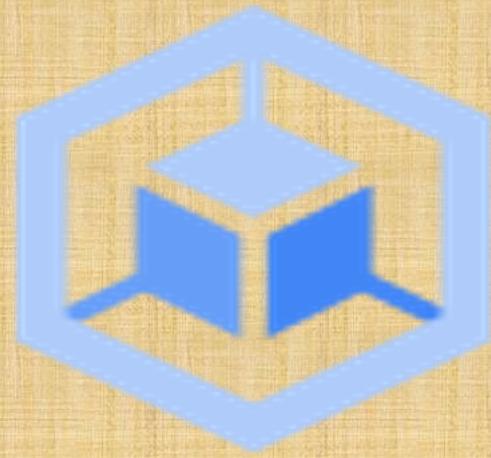
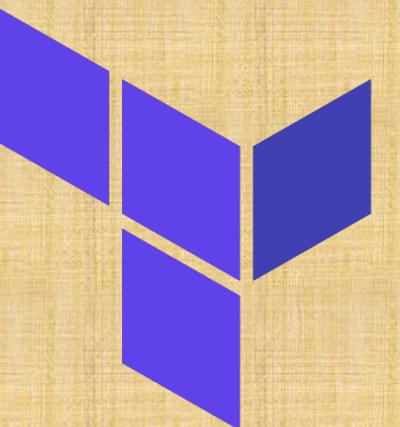
Kubernetes Load Balancer Service



Demo-20

# Google Kubernetes Engine File Store CSI Driver

**Demo: GKE Autopilot cluster + File Store**

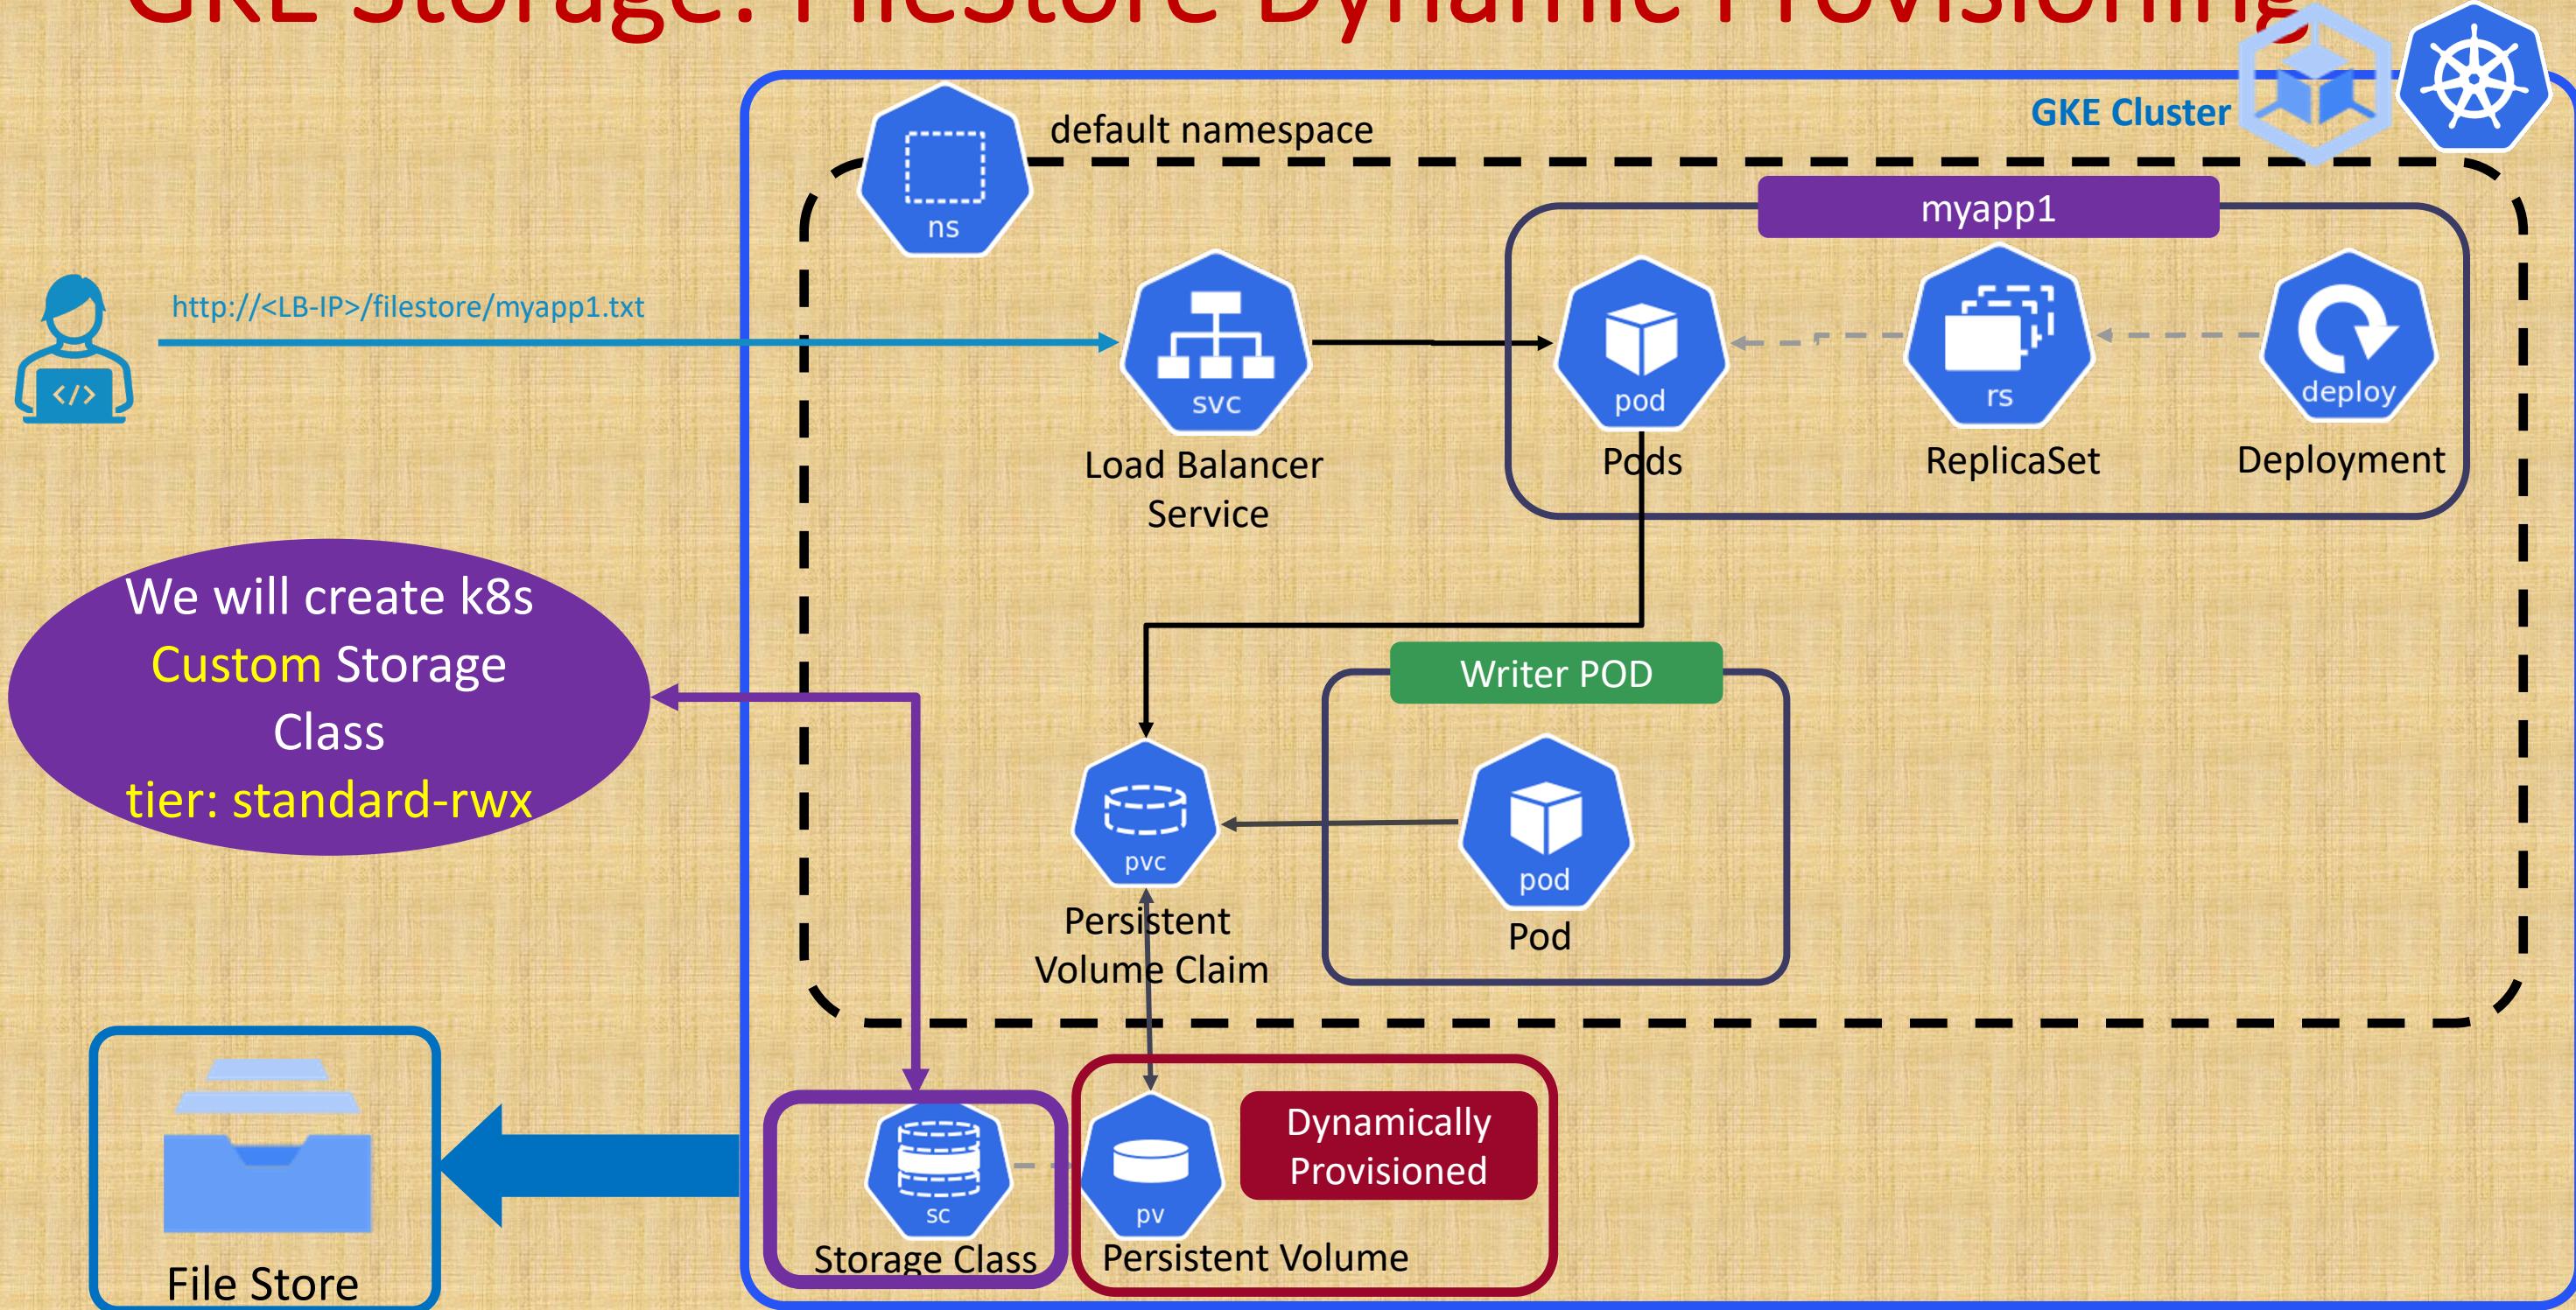


# File Store CSI Driver - Storage Classes

Service tier	GKE StorageClass	Number of shares per instance	GKE PVC size	Filestore capacity	Incremental change	Deployment	Access mode
Basic HDD <b>Standard</b>	<code>standard-rwx</code>	One	1 TiB to 64 TiB	1 TiB to 64 TiB	1 GiB	Zonal	read/write many
Basic SSD <b>Premium</b>	<code>premium-rwx</code>	One	2.5 TiB to 64 TiB	2.5 TiB to 64 TiB	1 GiB	Zonal	read/write many
Enterprise <b>Enterprise</b>	<code>enterprise-rwx</code>	One	1 TiB to 10 TiB	1 TiB to 10 TiB	256 GiB	Regional	read/write many
Enterprise with the multishares feature <b>Enterprise</b>	<code>enterprise-multishare-rwx</code>	Up to 80	10 GiB to 1 TiB	Pool of Filestore instances, each 1 TiB to 10 TiB in size	256 GiB per instance 1 GiB per PVC (share)	Regional	read/write many

Reference: [https://cloud.google.com/filestore/docs/multishares#storageclasses\\_and\\_dynamic\\_volume\\_provisioning](https://cloud.google.com/filestore/docs/multishares#storageclasses_and_dynamic_volume_provisioning)

# GKE Storage: FileStore Dynamic Provisioning



# What are we going to learn?

## ✓ p2-k8sresources-yaml

- ! 00-filestore-storage-class.yaml
- ! 01-filestore-pvc.yaml
- ! 02-write-to-filestore-pod.yaml
- ! 03-myapp1-deployment.yaml
- ! 04-loadBalancer-service.yaml

Demo: Project-2: GKE Autopilot cluster + File Store

00

Kubernetes Custom Storage Class

01

Kubernetes Persistent Volume Claim

02

Kubernetes Pod - Write to File store

03

Kubernetes MyApp1 Deployment (Read content from file store by mounting it)

04

Kubernetes Load Balancer Service

- ✓ p3-k8sresources-terraform-manifests
  - └ c1-versions.tf
  - └ c2-01-variables.tf
  - └ c2-02-local-values.tf
  - └ c3-01-remote-state-datasource.tf
  - └ c3-02-providers.tf
  - └ c4-01-storage-class.tf
  - └ c4-02-persistent-volume-claim.tf
  - └ c5-01-write-to-filestore-pod.tf
  - └ c5-02-myapp1-deployment.tf
  - └ c5-03-myapp1-loadbalancer-service.tf
  - └ terraform.tfvars

# What are we going to learn?

Demo: Project-3: GKE Autopilot cluster + File Store

NC

No changes c1 to c3

c4-01

Create Custom Storage Class

c4-02

Kubernetes Persistent Volume Claim

c5-01

Kubernetes Pod – Write to File Store instance

c5-02

Kubernetes MyApp1 Deployment (Read content from file store by mounting)

c5-03

Kubernetes Load Balancer Service

# Kubernetes Gateway API

## Introduction



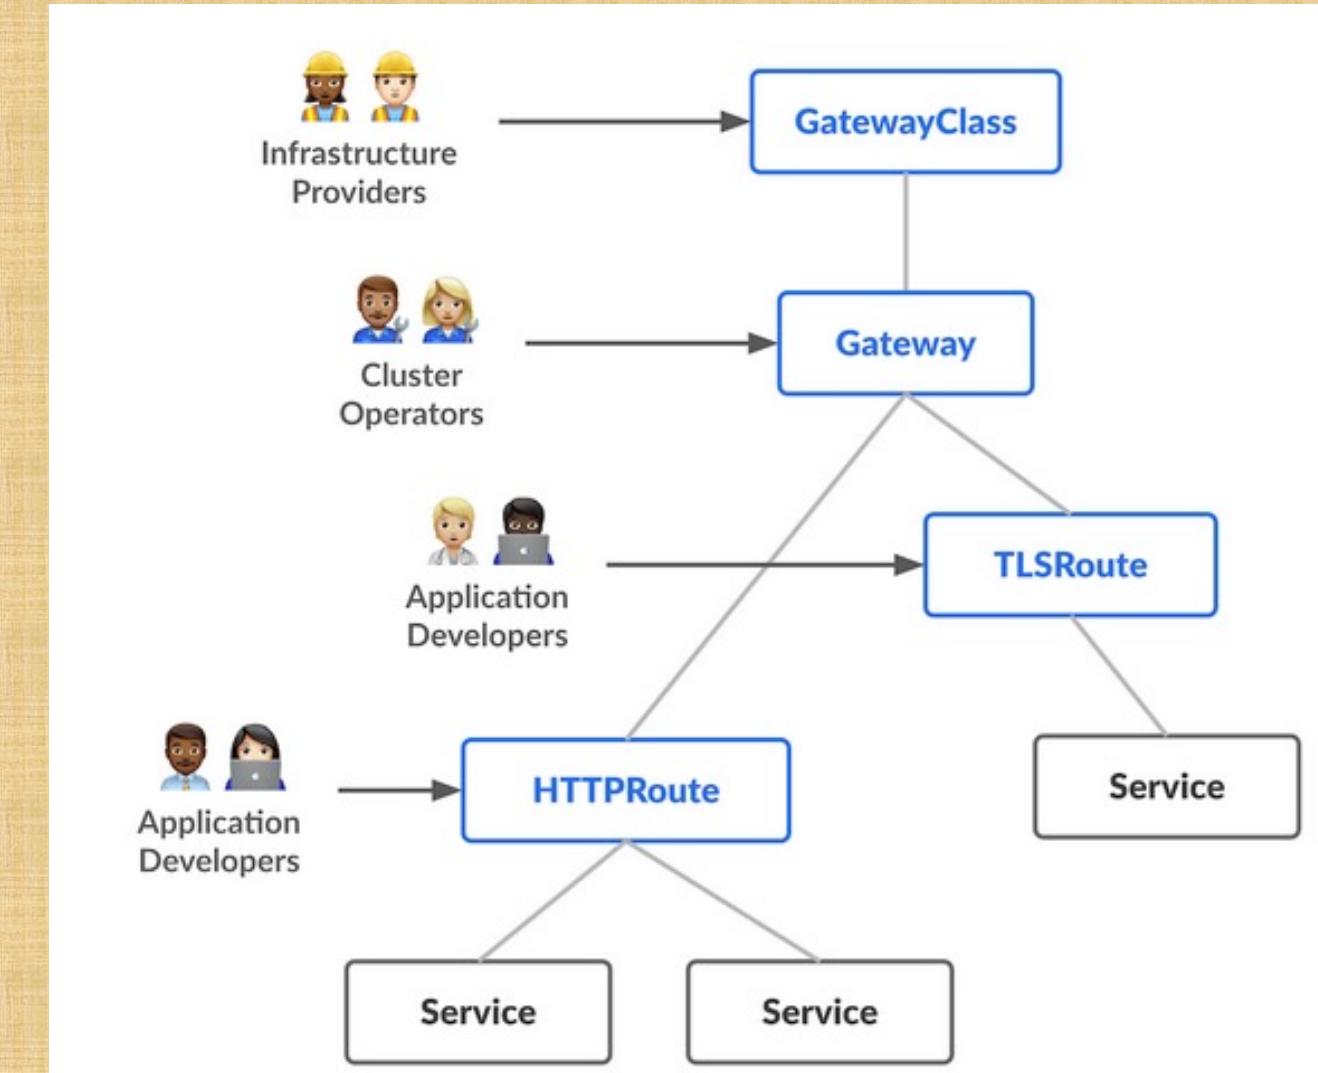
# Kubernetes Gateway API

- Why do we need to use Gateway API ?

- Kubernetes Ingress is frozen and new features are being added to the Gateway API
- Gateway API is going to be the replacement for Ingress going forward

- What is Gateway API ?

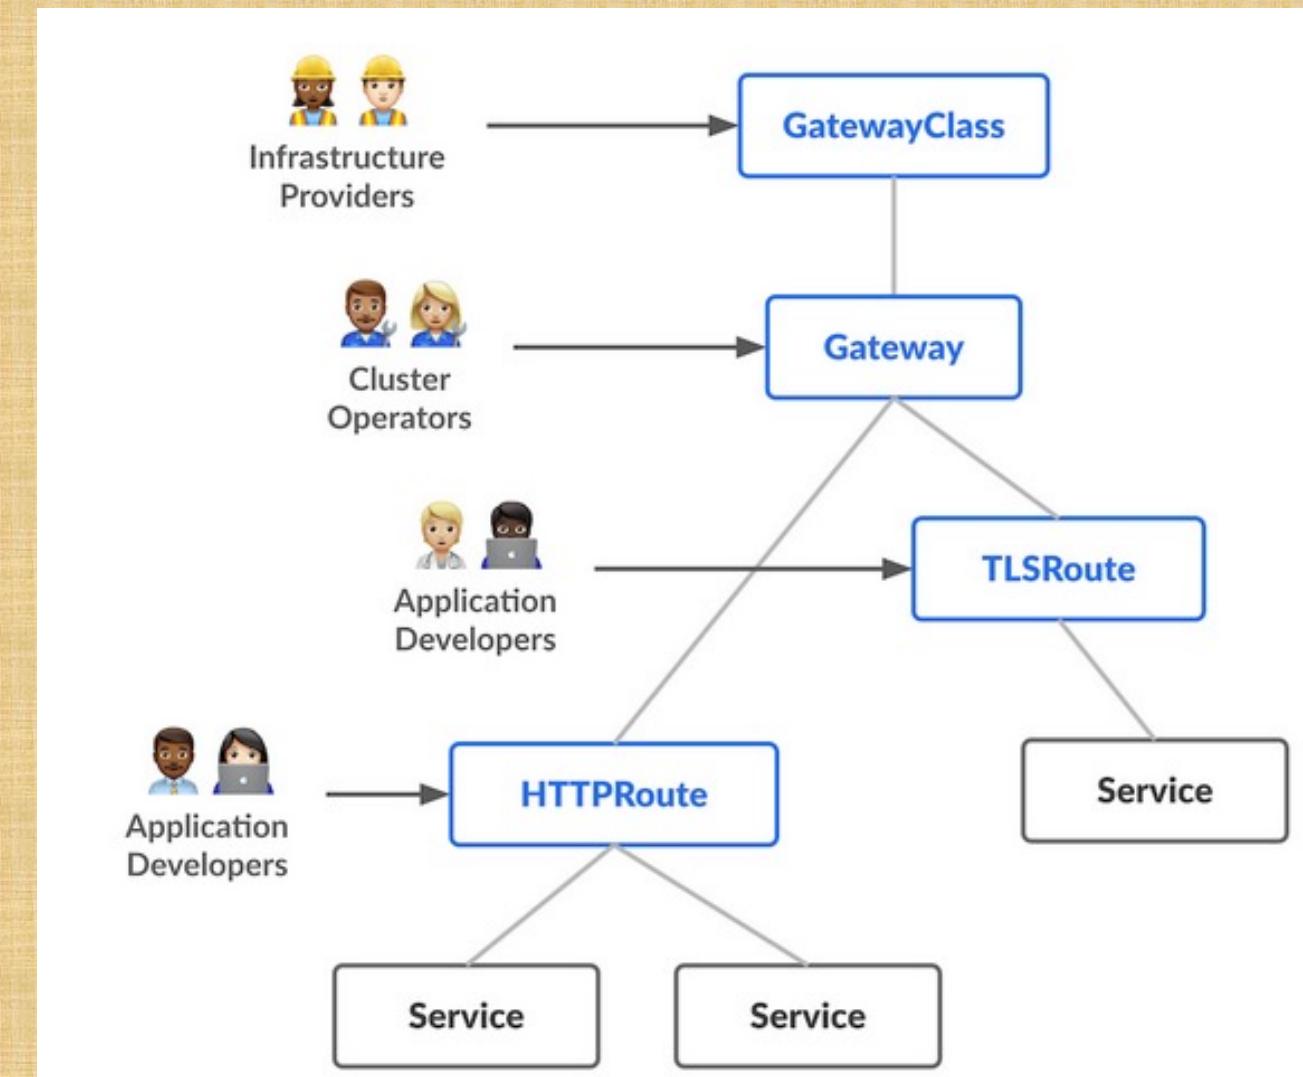
- Gateway API is an advanced version of Ingress
- Ingress supports only Layer7 protocols (HTTP, HTTPS)
- Gateway API supports
  - Layer7 Protocols: HTTP and HTTPS
  - Layer4 Protocols: TCP and UDP
  - gRPC Protocol



Reference: <https://gateway-api.sigs.k8s.io/>

# Kubernetes Gateway API

- **Role-oriented Design:** The Gateway API is a role-oriented resource model
  - Users with **specific roles** can manage those **specific Gateway API Resources**
  - **Example:**
    - Infrastructure Providers can manage **GatewayClass**
    - Cluster Operators can manage **Gateways**
    - Application Developers can define **TLSRoute**, **HTTPRoute**



Reference: <https://gateway-api.sigs.k8s.io/>

# Kubernetes Gateway API Terminology

- **GatewayClass**

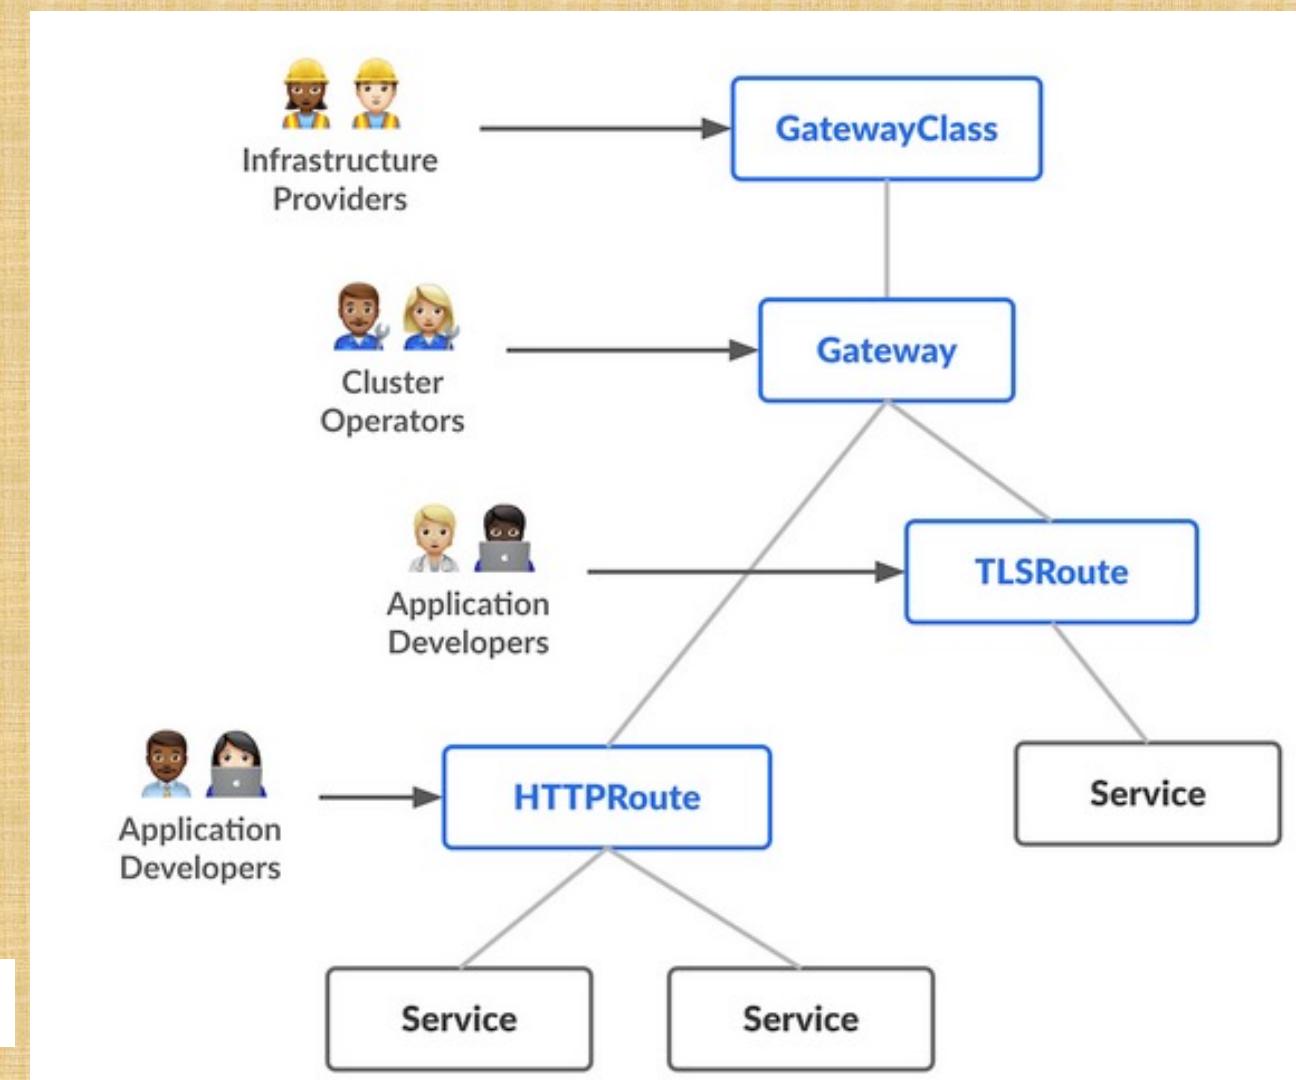
- Template for creating load balancers in a cluster.
- It's a **cluster-scoped resource**

- **Gateway**

- Defines **where and how** the load balancers **listen for traffic**.
- **Cluster operators** create Gateways in their clusters based on a GatewayClass.

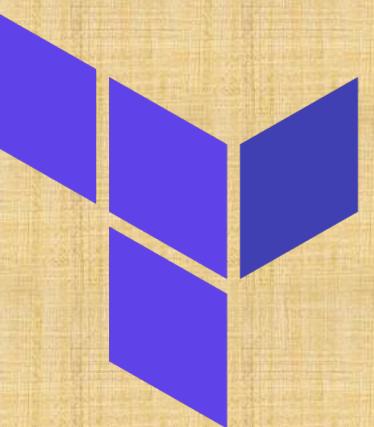
- **HTTPRoute**

- Defines **protocol-specific rules** for **routing requests** from a Gateway to Kubernetes services.
- **Application developers** create HTTPRoutes to expose their HTTP applications using Gateways.



Reference: <https://gateway-api.sigs.k8s.io/>

Demo-01

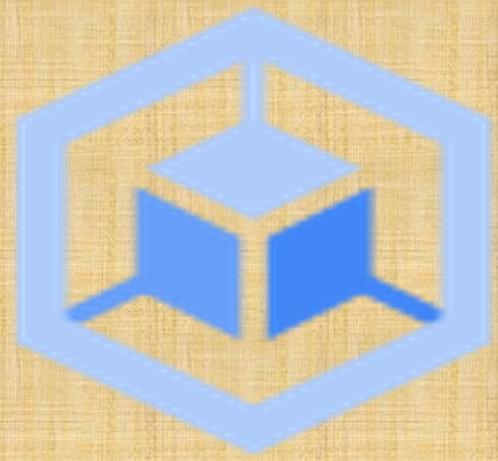


# Google Kubernetes Engine

## Gateway API Basics

## Cloud Application Load Balancer

**Demo: GKE Autopilot cluster + Gateway API  
(Cloud Application Load Balancer)**



# Kubernetes Gateway API

- In GKE, Gateway API is implemented using **GKE Gateway controller**
- **GKE Gateway classes**

Gatewayclass	GCP Load Balancer	Load Balancer Scope
gke-l7-global-external-managed	Global External Application Load Balancer	Global
gke-l7-regional-external-managed	Regional External Application Load Balancer	Regional
gke-l7-rlb	Internal Application Load Balancer	Regional
gke-l7-gxlb-mc	Classic Application Load Balancer	Global

Reference: <https://cloud.google.com/kubernetes-engine/docs/concepts/gateway-api>

# Gateway API Basics



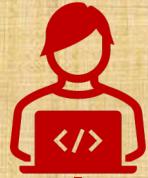
Customer Project: gcplearn9



Customer VPC: VPC1

Region: us-central1

Users



http://<LB-IP>

GKE Autopilot Cluster

Subnet: 10.128.0.0/20

Gateway  
(Regional LB)

HTTP Listener

HTTPRoute

Default  
Backend

ClusterIP  
Service



MyApp1 Deployment



Pods



ReplicaSet



Deployment

default namespace



Docker  
Hub  
Internet



Cloud NAT



Cloud Router

Automated by Terraform

## Project-2: YAML Manifests

- ✓ p2-regional-k8sresources-yaml
  - ! 01-myapp1-deployment.yaml
  - ! 02-myapp1-clusterip-service.yaml
  - ! 03-gateway.yaml
  - ! 04-gateway-http-route.yaml

# What are we going to learn?

Demo: Project-2: GKE Autopilot cluster + Gateway API + basic scenario

NC

No changes 01 and 02

03

Gateway Manifest

04

Gateway **HTTPRoute** Manifest

## Project-3: Terraform Manifests

- ✓ p3-regional-k8sresources-terraform-manifests
- ✗ c1-versions.tf
- ✗ c2-01-variables.tf
- ✗ c2-02-local-values.tf
- ✗ c3-01-remote-state-datasource.tf
- ✗ c3-02-providers.tf
- ✗ c4-myapp1-deployment.tf
- ✗ c5-myapp1-clusterip-service.tf
- ✗ c6-gateway.tf
- ✗ c7-gateway-http-route.tf
- ✗ terraform.tfvars

# What are we going to learn?

Demo: Project-3: GKE Autopilot cluster + Gateway API + basic scenario

NC

No changes c1 to c5

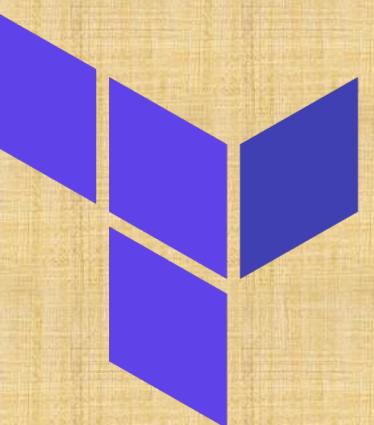
c6

Create **Gateway Manifest** which should create  
Regional Application Load Balancer

c7

Create Gateway **HTTPRoute**

Demo-02

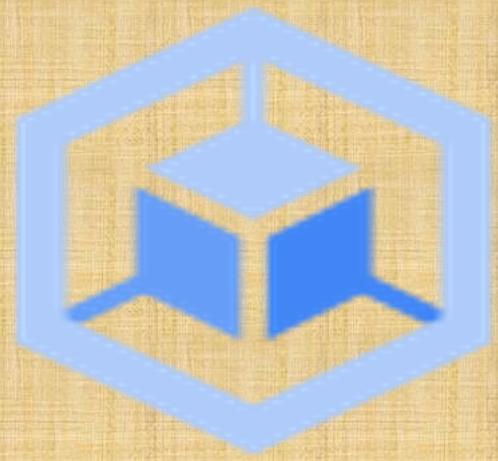


# Google Kubernetes Engine

Gateway API

## Cloud Application Load Balancer

Demo: GKE Autopilot cluster + Gateway + Static IP



# Gateway API Basics + Static IP



Customer Project: gcplearn9



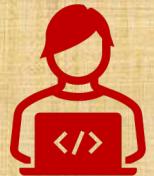
Customer VPC: VPC1

Region: us-central1

GKE Autopilot Cluster

Subnet: 10.128.0.0/20

Users



http://<LB-IP>

Gateway  
(Regional LB)

HTTP Listener

Static IP

HTTPRoute

Default  
Backend

ClusterIP  
Service



MyApp1 Deployment



Pods



ReplicaSet

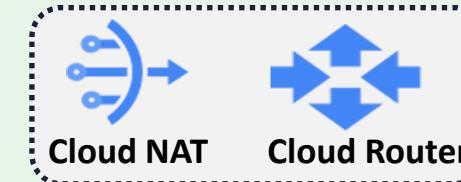


Deployment

default namespace



Cloud  
External IP  
Addresses



Automated by Terraform

- ✓ p2-regional-k8sresources-yaml
  - ! 01-myapp1-deployment.yaml
  - ! 02-myapp1-clusterip-service.yaml
  - ! 03-gateway.yaml
  - ! 04-gateway-http-route.yaml

# What are we going to learn?

Demo: Project-2: GKE Autopilot cluster + Gateway + Regional Static IP

NC

No changes 01, 02 and 04

03

Create Static IP using gcloud and Update Gateway Manifest with Static IP

## Project-3: Terraform Manifests

- ✓ p3-regional-k8sresources-terraform-manifests
- ✓ c1-versions.tf
- ✓ c2-01-variables.tf
- ✓ c2-02-local-values.tf
- ✓ c3-01-remote-state-datasource.tf
- ✓ c3-02-providers.tf
- ✓ c4-myapp1-deployment.tf
- ✓ c5-myapp1-clusterip-service.tf
- ✓ c6-gateway.tf
- ✓ c7-gateway-http-route.tf
- ✓ c8-static-ip.tf
- ✓ terraform.tfvars

# What are we going to learn?

Demo: Project-3: GKE Autopilot cluster + Gateway + Regional Static IP

NC

No changes c1 to c5, c7

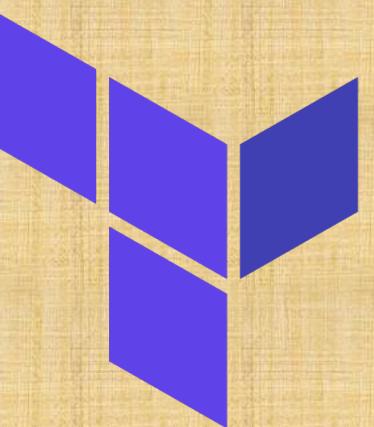
c8

Create Regional Static IP

c6

Update Gateway manifest with Static IP

Demo-03

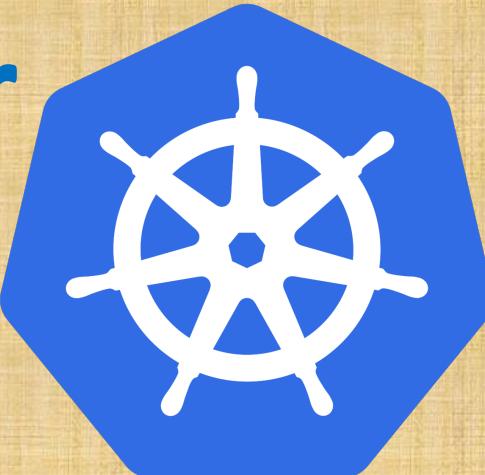
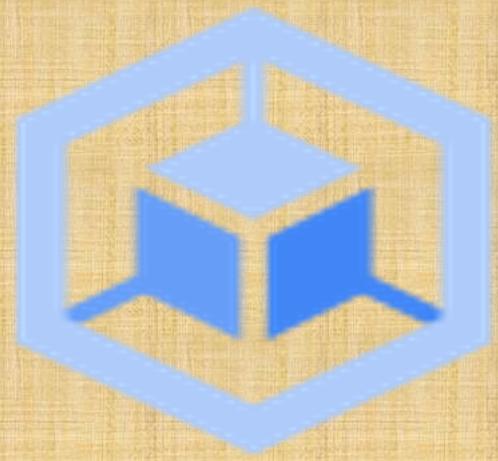


# Google Kubernetes Engine

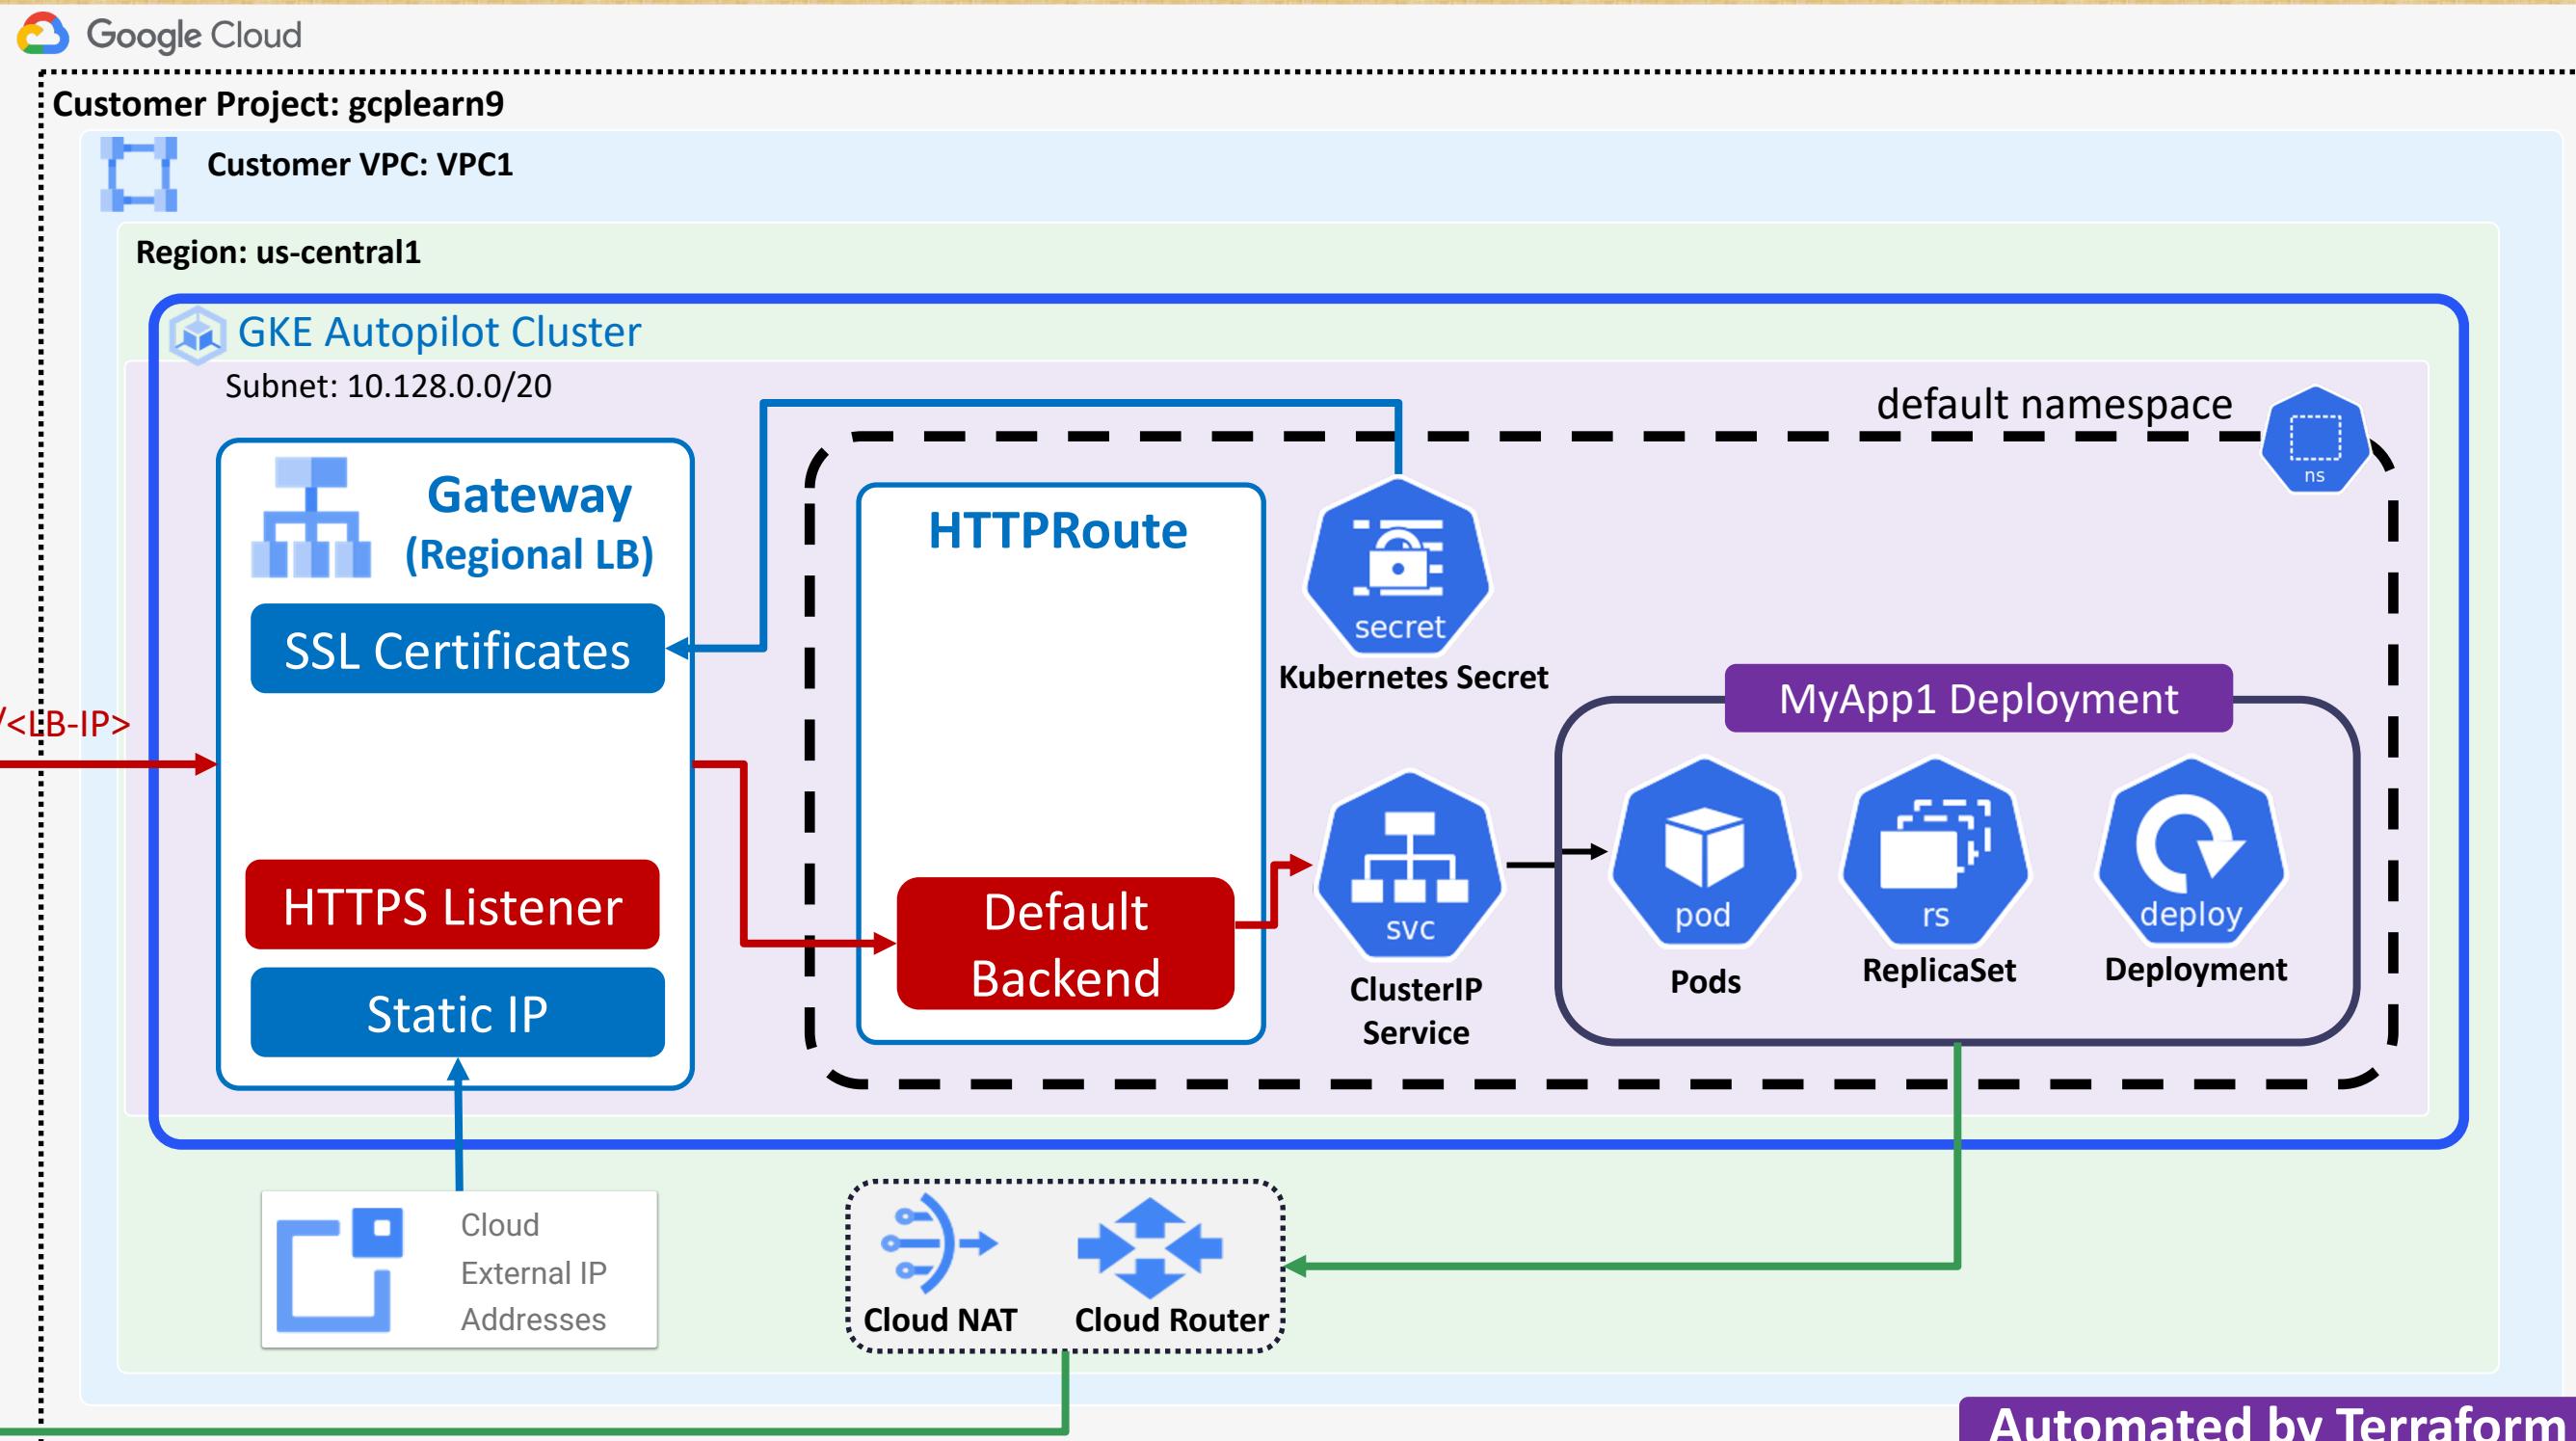
Gateway API

## Cloud Application Load Balancer

**Demo: GKE Autopilot cluster + Gateway + Static IP +  
Self-signed SSL with k8s Secret**



# Gateway API Basics + Static IP + Self-signed SSL with Kubernetes Secret



- ✓ p2-regional-k8sresources-yaml
  - > self-signed-ssl
  - ! 01-myapp1-deployment.yaml
  - ! 02-myapp1-clusterip-service.yaml
  - ! 03-gateway.yaml
  - ! 04-gateway-http-route.yaml

# What are we going to learn?

Demo: Project-2: GKE Autopilot cluster + Gateway + Regional Static IP + Self-signed SSL k8s secret

NC

No changes 01, 02 and 04

03

Create SSL Cert, Kubernetes secret and map the k8s secret in Gateway manifest

03

Update Listener to HTTPS and Port to 443

## Project-3: Terraform Manifests

- ✓ p3-regional-k8sresources-terraform-manifests
  - > self-signed-ssl
  - └ c1-versions.tf
  - └ c2-01-variables.tf
  - └ c2-02-local-values.tf
  - └ c3-01-remote-state-datasource.tf
  - └ c3-02-providers.tf
  - └ c4-myapp1-deployment.tf
  - └ c5-myapp1-clusterip-service.tf
  - └ c6-gateway.tf
  - └ c7-gateway-http-route.tf
  - └ c8-static-ip.tf
  - └ c9-kubernetes-secret.tf
  - └ terraform.tfvars

# What are we going to learn?

Demo: Project-3: GKE Autopilot cluster + Gateway API + Regional Static IP + Self-signed SSL k8s secret

NC

No changes c1 to c5, c7, c8

c9

Create Kubernetes Secret with cert and key

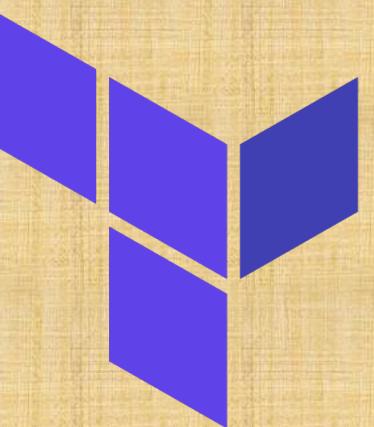
c6

Update Gateway manifest with Kubernetes secret

c6

Update Listener to HTTPS and Port to 443

Demo-04

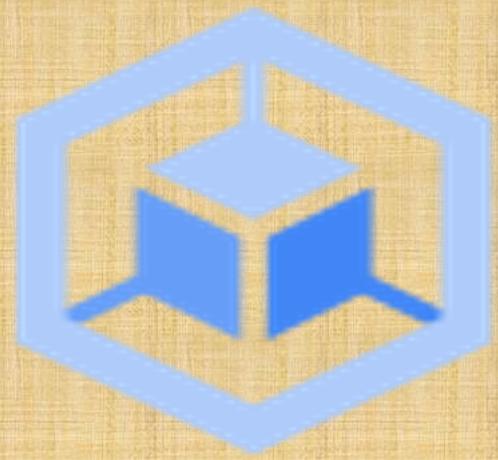


# Google Kubernetes Engine

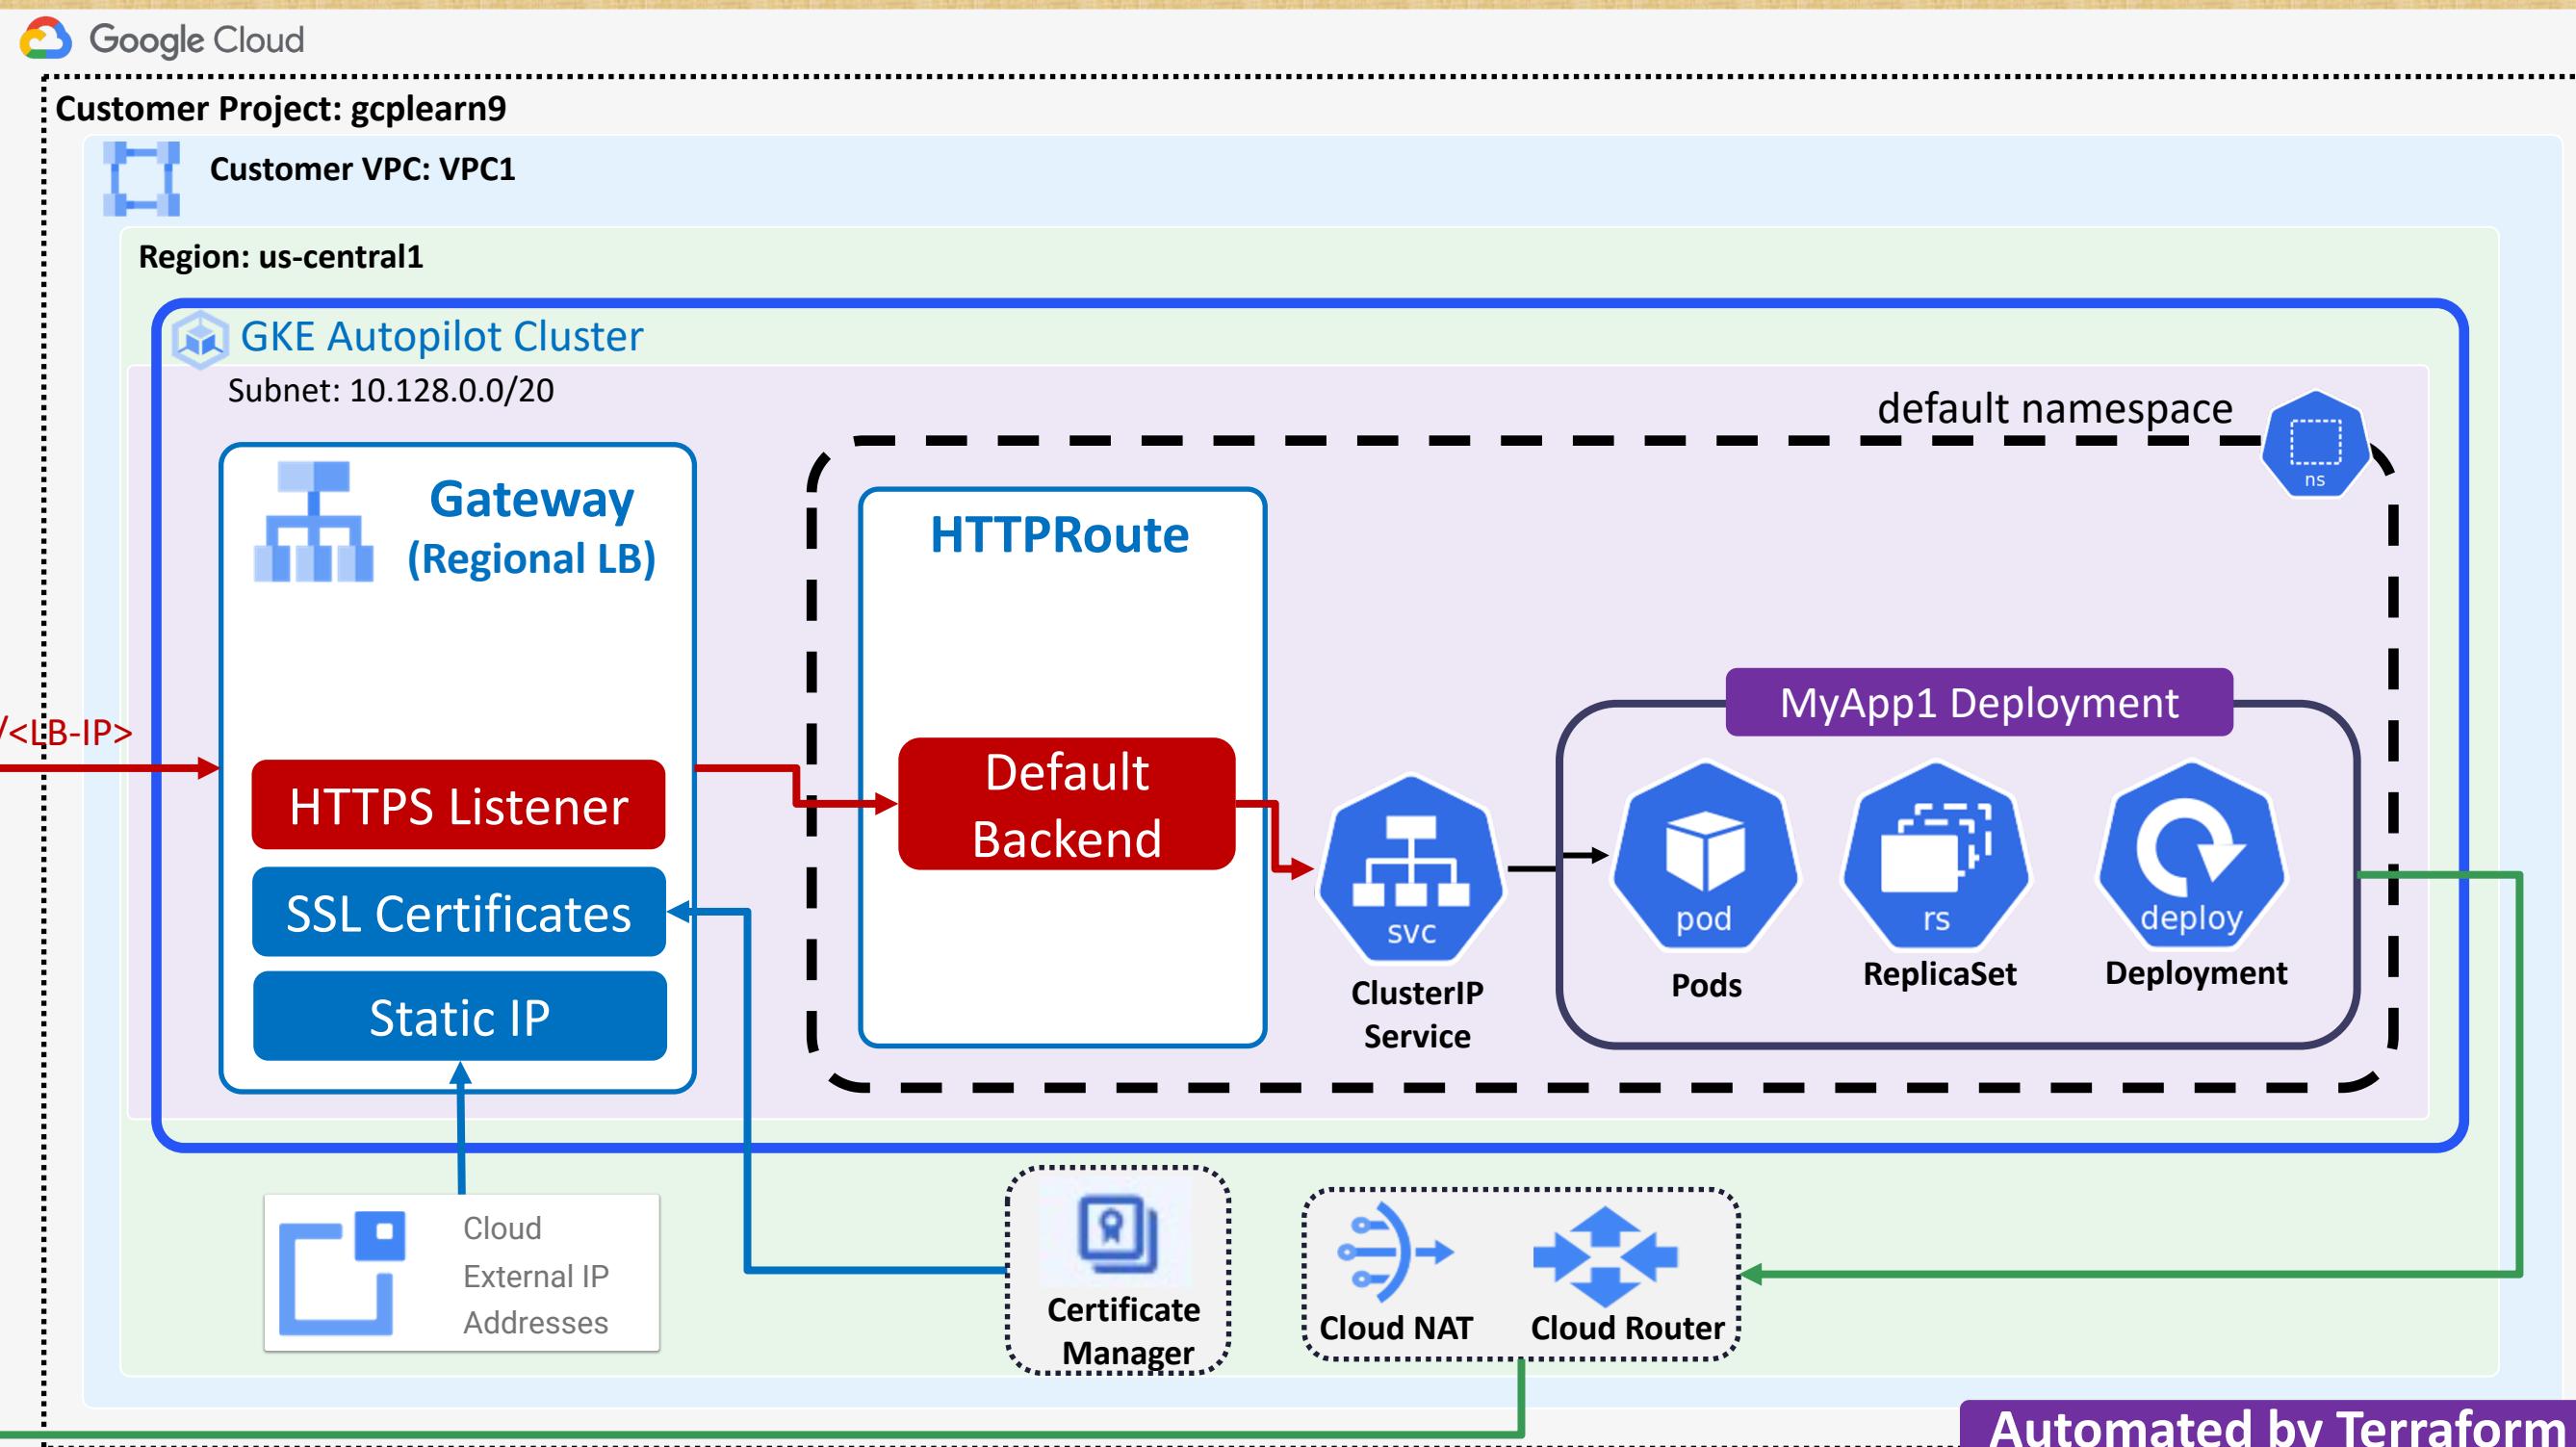
Gateway API

## Cloud Application Load Balancer

**Demo: GKE Autopilot cluster + Gateway + Static IP +  
Self-signed SSL with Certificate Manager**



# Gateway API Basics + Static IP + Self-signed SSL with Certificate Manager



- ✓ p2-regional-k8sresources-yaml
  - > self-signed-ssl
  - ! 01-myapp1-deployment.yaml
  - ! 02-myapp1-clusterip-service.yaml
  - ! 03-gateway.yaml
  - ! 04-gateway-http-route.yaml

# What are we going to learn?

Demo: Project-2: GKE Autopilot cluster + Gateway + Regional Static IP + Self-signed SSL Certificate Manager

NC

No changes 01, 02 and 04

03

Create SSL Cert, Certificate Manager Object and reference it in Gateway Manifest

## Project-3: Terraform Manifests

- ✓ p3-regional-k8sresources-terraform-manifests
  - > self-signed-ssl
  - └ c1-versions.tf
  - └ c2-01-variables.tf
  - └ c2-02-local-values.tf
  - └ c3-01-remote-state-datasource.tf
  - └ c3-02-providers.tf
  - └ c4-myapp1-deployment.tf
  - └ c5-myapp1-clusterip-service.tf
  - └ c6-gateway.tf
  - └ c7-gateway-http-route.tf
  - └ c8-static-ip.tf
  - └ c9-certificate-manager.tf
  - └ terraform.tfvars

# What are we going to learn?

Demo: Project-3: GKE Autopilot cluster + Gateway + Regional Static IP + Self-signed SSL Certificate Manager

NC

No changes c1 to c5, c7, c8

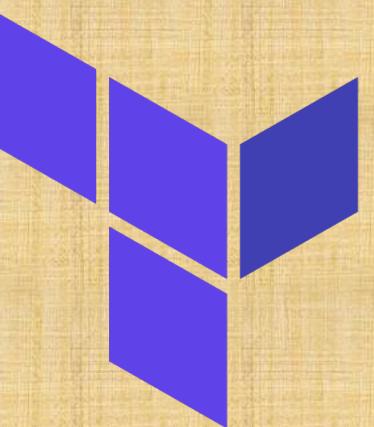
c9

Create Self-signed SSL certificate using Certificate Manager

c6

Update Gateway manifest with Certificate Manager object

Demo-05

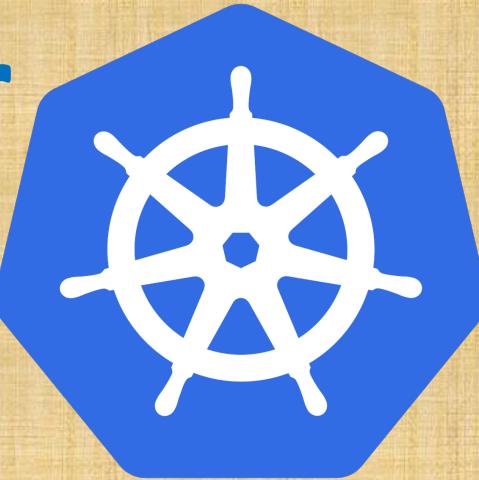
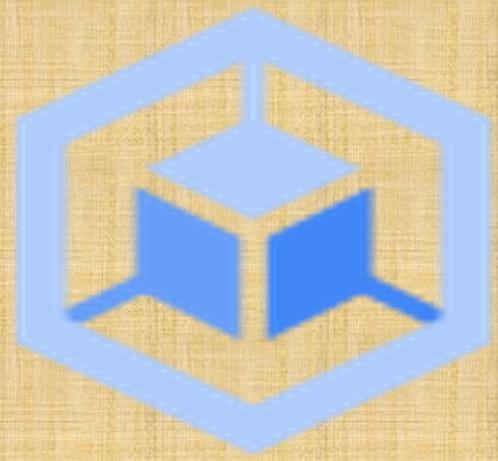


# Google Kubernetes Engine

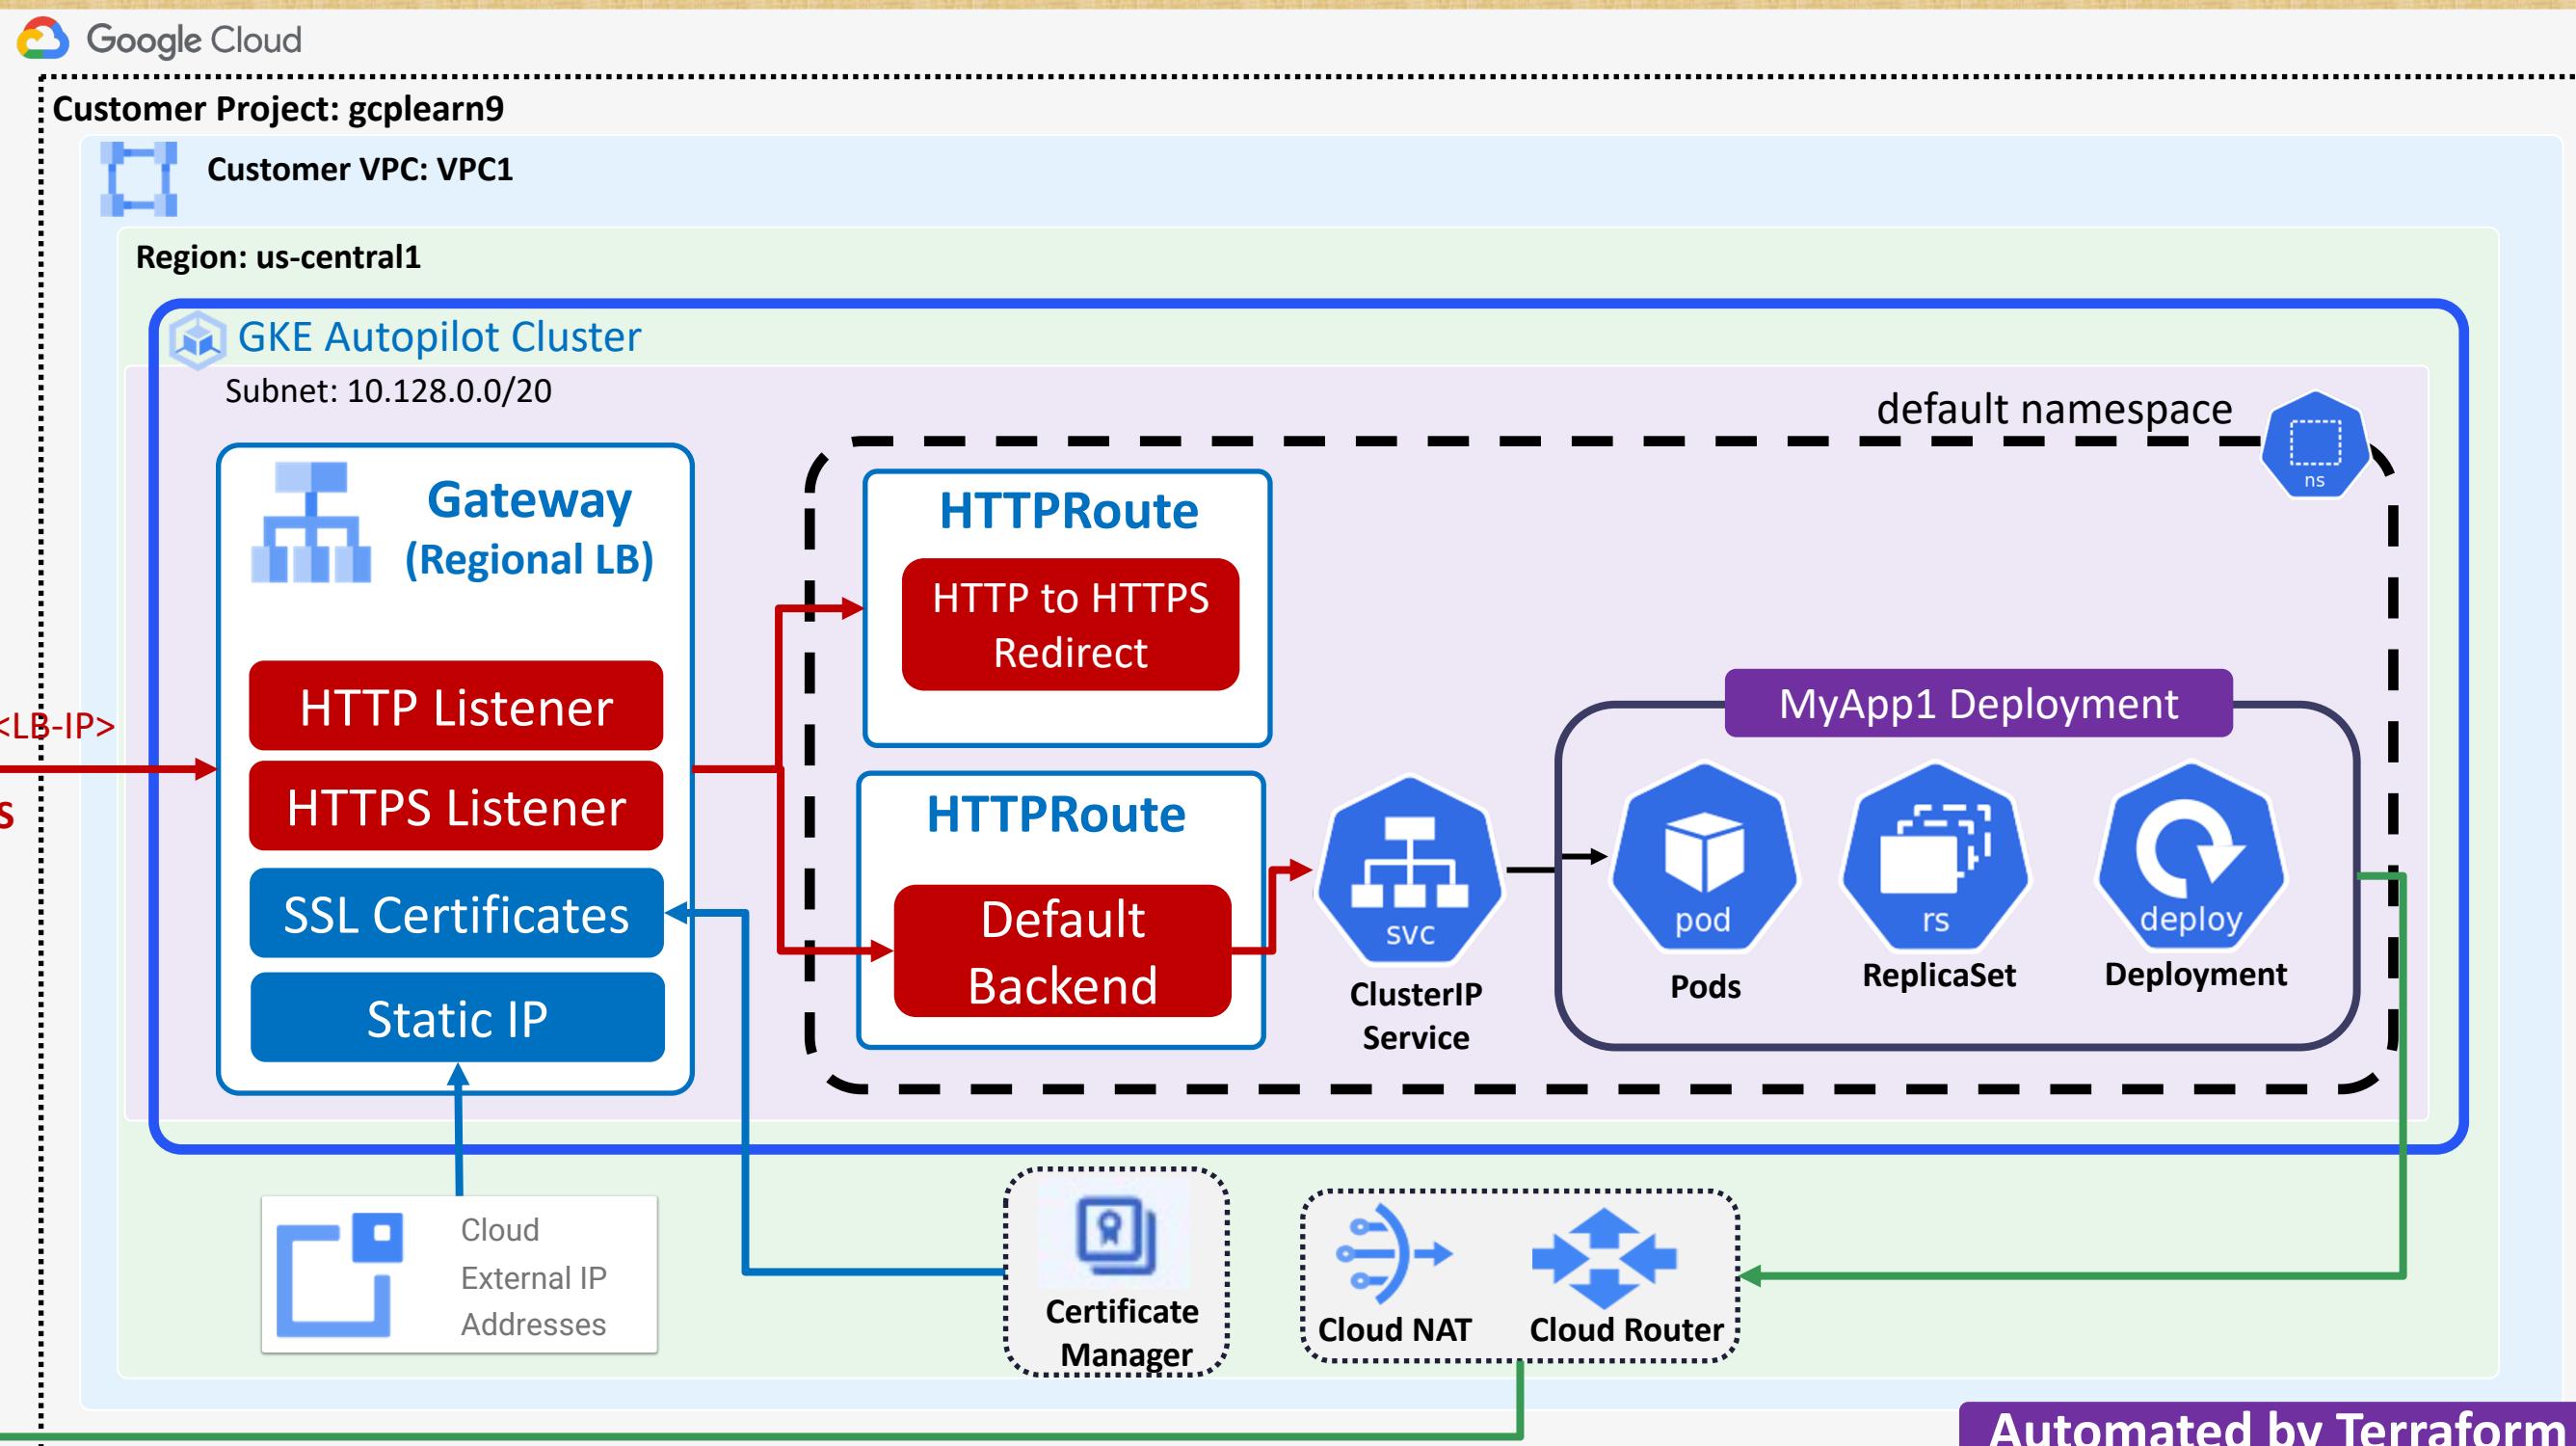
Gateway API

## Cloud Application Load Balancer

**Demo:** GKE Autopilot cluster + Gateway + Static IP +  
Self-signed SSL with Certificate Manager + HTTP to  
HTTPS Redirect



# Gateway API Basics + Static IP + SSL Certificate Manager + HTTP to HTTPS



- ✓ p2-regional-k8sresources-yaml
  - > self-signed-ssl
    - ! 01-myapp1-deployment.yaml
    - ! 02-myapp1-clusterip-service.yaml
    - ! 03-gateway.yaml
    - ! 04-gateway-http-route.yaml
    - ! 05-gateway-http-to-https-route.yaml

# What are we going to learn?

Demo: Project-2: Gateway + Regional Static IP + Self-signed SSL Certificate Manager + HTTP to HTTPS Redirect

NC

No changes 01, 02 and 04

03

Add both **HTTP** and **HTTPS** listeners

04

Update route with “**sectionName: https**”, in short map this route to HTTPS listener in gateway

05

Create new **HTTP to HTTPS redirect Route**

## Project-3: Terraform Manifests

- ✓ p3-regional-k8sresources-terraform-manifests
  - > self-signed-ssl
  - c1-versions.tf
  - c2-01-variables.tf
  - c2-02-local-values.tf
  - c3-01-remote-state-datasource.tf
  - c3-02-providers.tf
  - c4-myapp1-deployment.tf
  - c5-myapp1-clusterip-service.tf
  - c6-01-gateway.tf
  - c6-02-gateway-http-route.tf
  - c6-03-gateway-http-to-https-route.tf
  - c7-static-ip.tf
  - c8-certificate-manager.tf
  - terraform.tfvars

# What are we going to learn?

Demo: Project-3: Gateway + Regional Static IP + Self-signed SSL Certificate Manager + HTTP to HTTPS Redirect

NC

No changes c1 to c5, c7, c8

C6-01

Create both **HTTP and HTTPS Listeners** in Gateway manifest

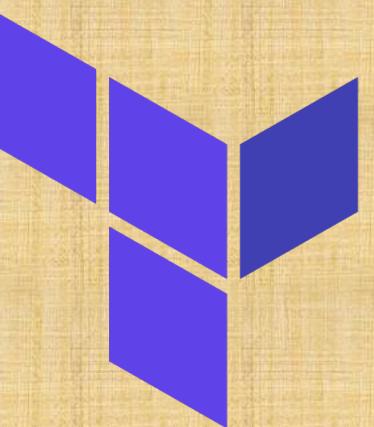
C6-02

Update route with “**sectionName: https**”, in short map this route to HTTPS listener in gateway

C6-03

Create new **HTTP to HTTPS redirect Route**

Demo-06

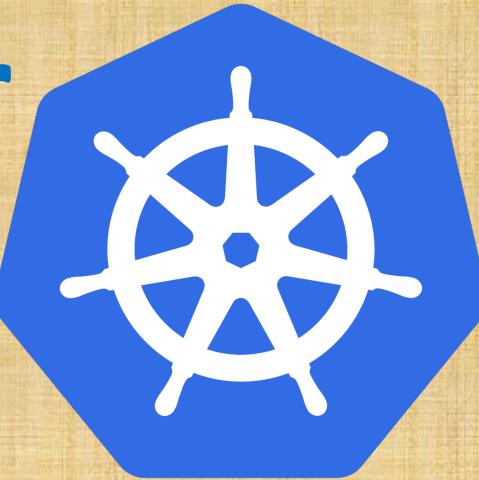
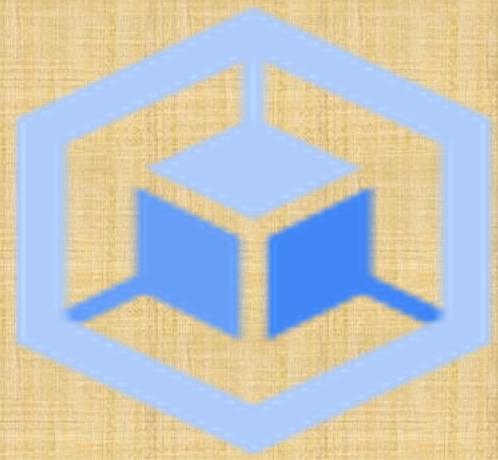


# Google Kubernetes Engine

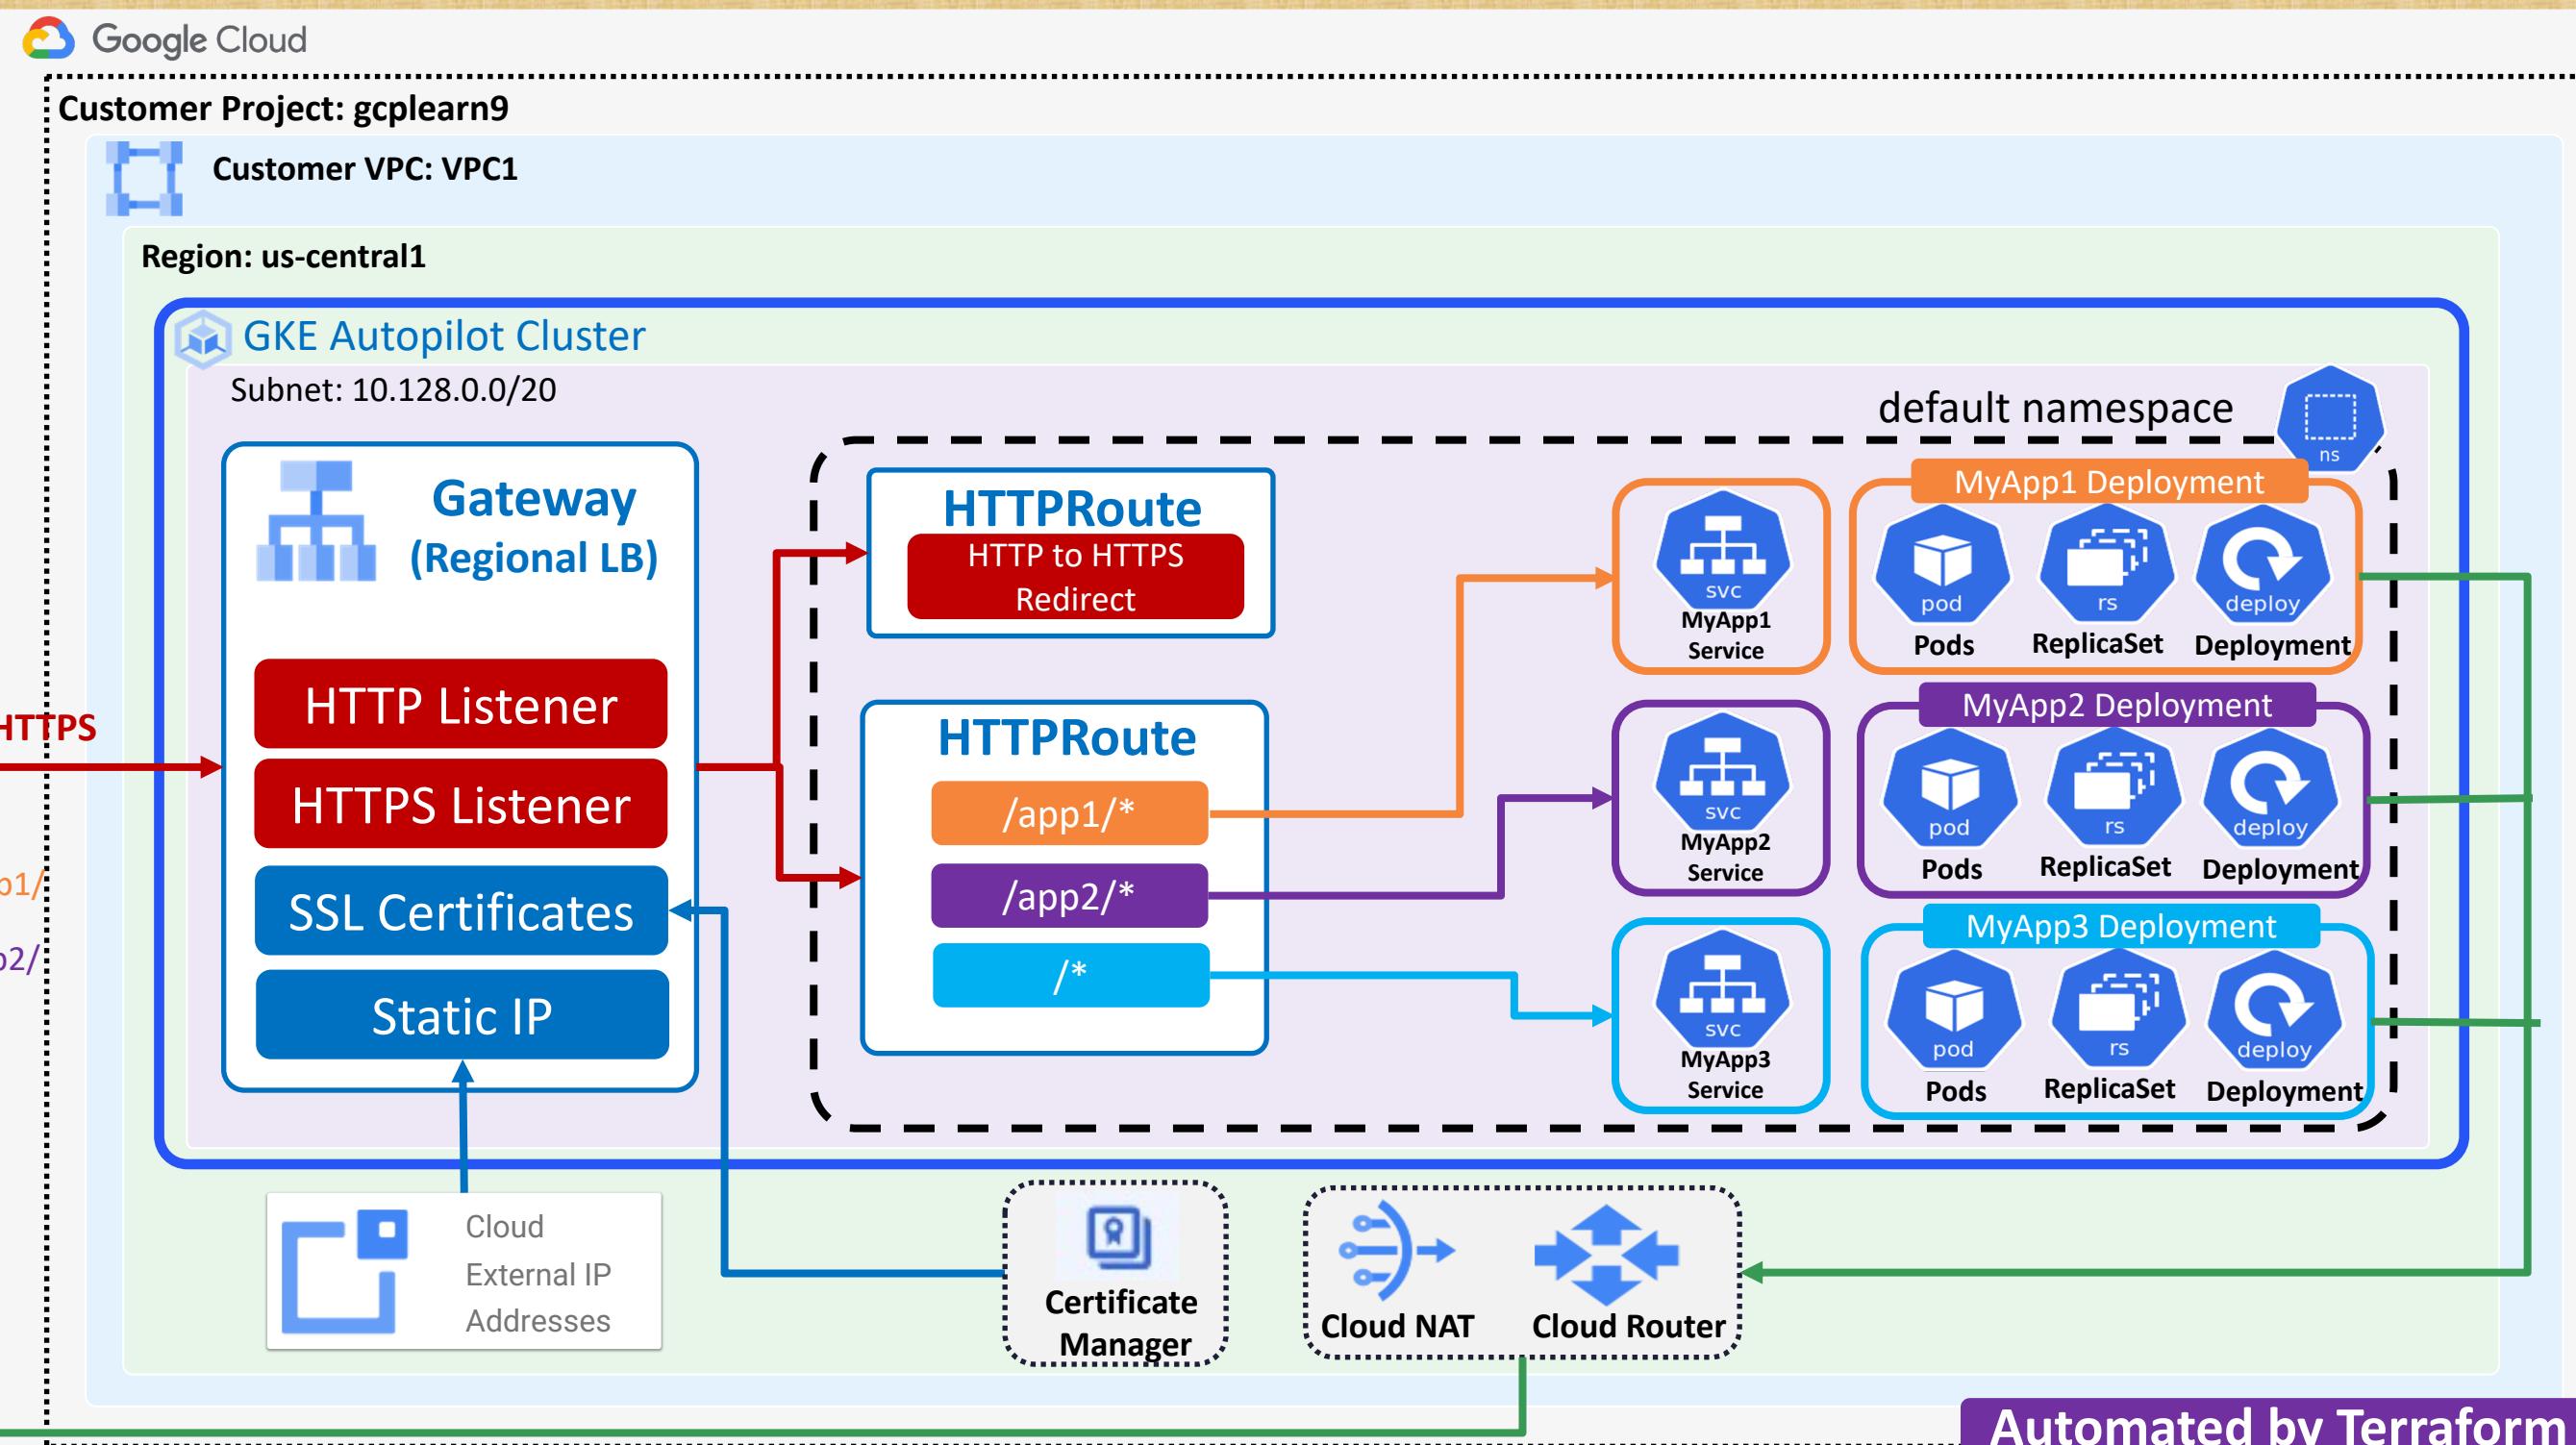
Gateway API

## Cloud Application Load Balancer

**Demo:** GKE Autopilot cluster + Gateway + Static IP +  
Self-signed SSL with Certificate Manager + HTTP to  
HTTPS Redirect + Context Path based Routing



# Gateway API Basics + Static IP + SSL Certificate Manager + HTTP to HTTPS + Context Path based Routing



## Project-2: YAML Manifests

- ✓ p2-regional-k8sresources-yaml
  - > self-signed-ssl
  - ! 01-myapp1-deployment.yaml
  - ! 02-myapp1-clusterip-service.yaml
  - ! 03-myapp2-deployment.yaml
  - ! 04-myapp2-clusterip-service.yaml
  - ! 05-myapp3-deployment.yaml
  - ! 06-myapp3-clusterip-service.yaml
  - ! 07-gateway.yaml
  - ! 08-gateway-http-route.yaml
  - ! 09-gateway-http-to-https-route.yaml

# What are we going to learn?

Demo: Project-2: Gateway + Regional Static IP + Self-signed SSL Certificate Manager + HTTP to HTTPS + Context Path Routing

01 to  
06

07

09

08

MyApp1, 2, 3 Deployments and Cluster IP Services

Gateway – No changes

HTTP to HTTPS redirect – No changes

Define Context path-based routing

## Project-3: Terraform Manifests

- ✓ p3-regional-k8sresources-terraform-manifests
  - > self-signed-ssl
  - ✓ c1-versions.tf
  - ✓ c2-01-variables.tf
  - ✓ c2-02-local-values.tf
  - ✓ c3-01-remote-state-datasource.tf
  - ✓ c3-02-providers.tf
  - ✓ c4-01-myapp1-deployment.tf
  - ✓ c4-02-myapp1-clusterip-service.tf
  - ✓ c4-03-myapp2-deployment.tf
  - ✓ c4-04-myapp2-clusterip-service.tf
  - ✓ c4-05-myapp3-deployment.tf
  - ✓ c4-06-myapp3-clusterip-service.tf
  - ✓ c5-01-gateway.tf
  - ✓ c5-02-gateway-http-route.tf
  - ✓ c5-03-gateway-http-to-https-route.tf
  - ✓ c6-static-ip.tf
  - ✓ c7-certificate-manager.tf
  - ✓ terraform.tfvars

# What are we going to learn?

**Demo: Project-3: Gateway + Regional Static IP + Self-signed SSL Certificate Manager + HTTP to HTTPS + Context Path Routing**

NC

No changes c1 to c3-02, c6, c7

C4

MyApp1, 2, 3 Deployments and Cluster IP Services

C5-01

No changes in Gateway manifest

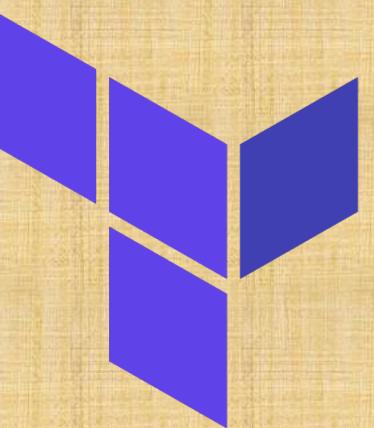
C5-03

No changes in HTTP to HTTPS redirect

C5-02

Define **Context path-based routing**

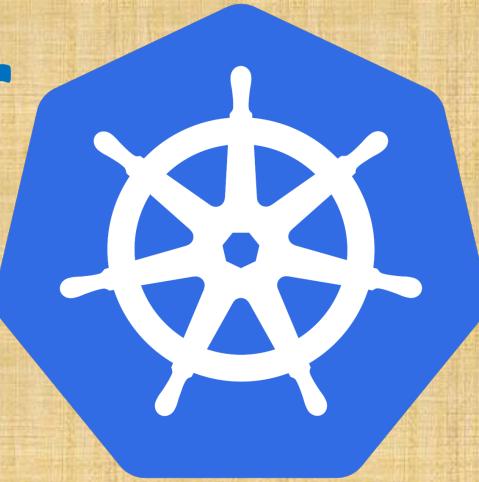
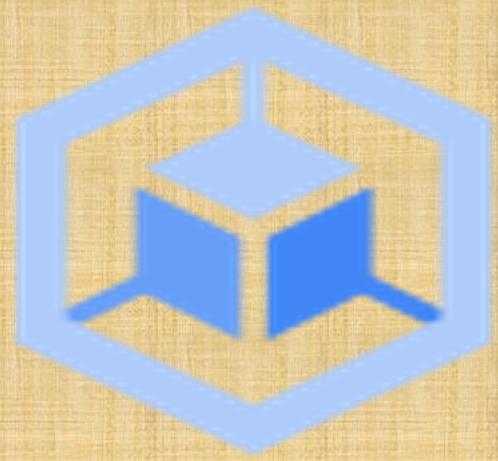
Demo-07



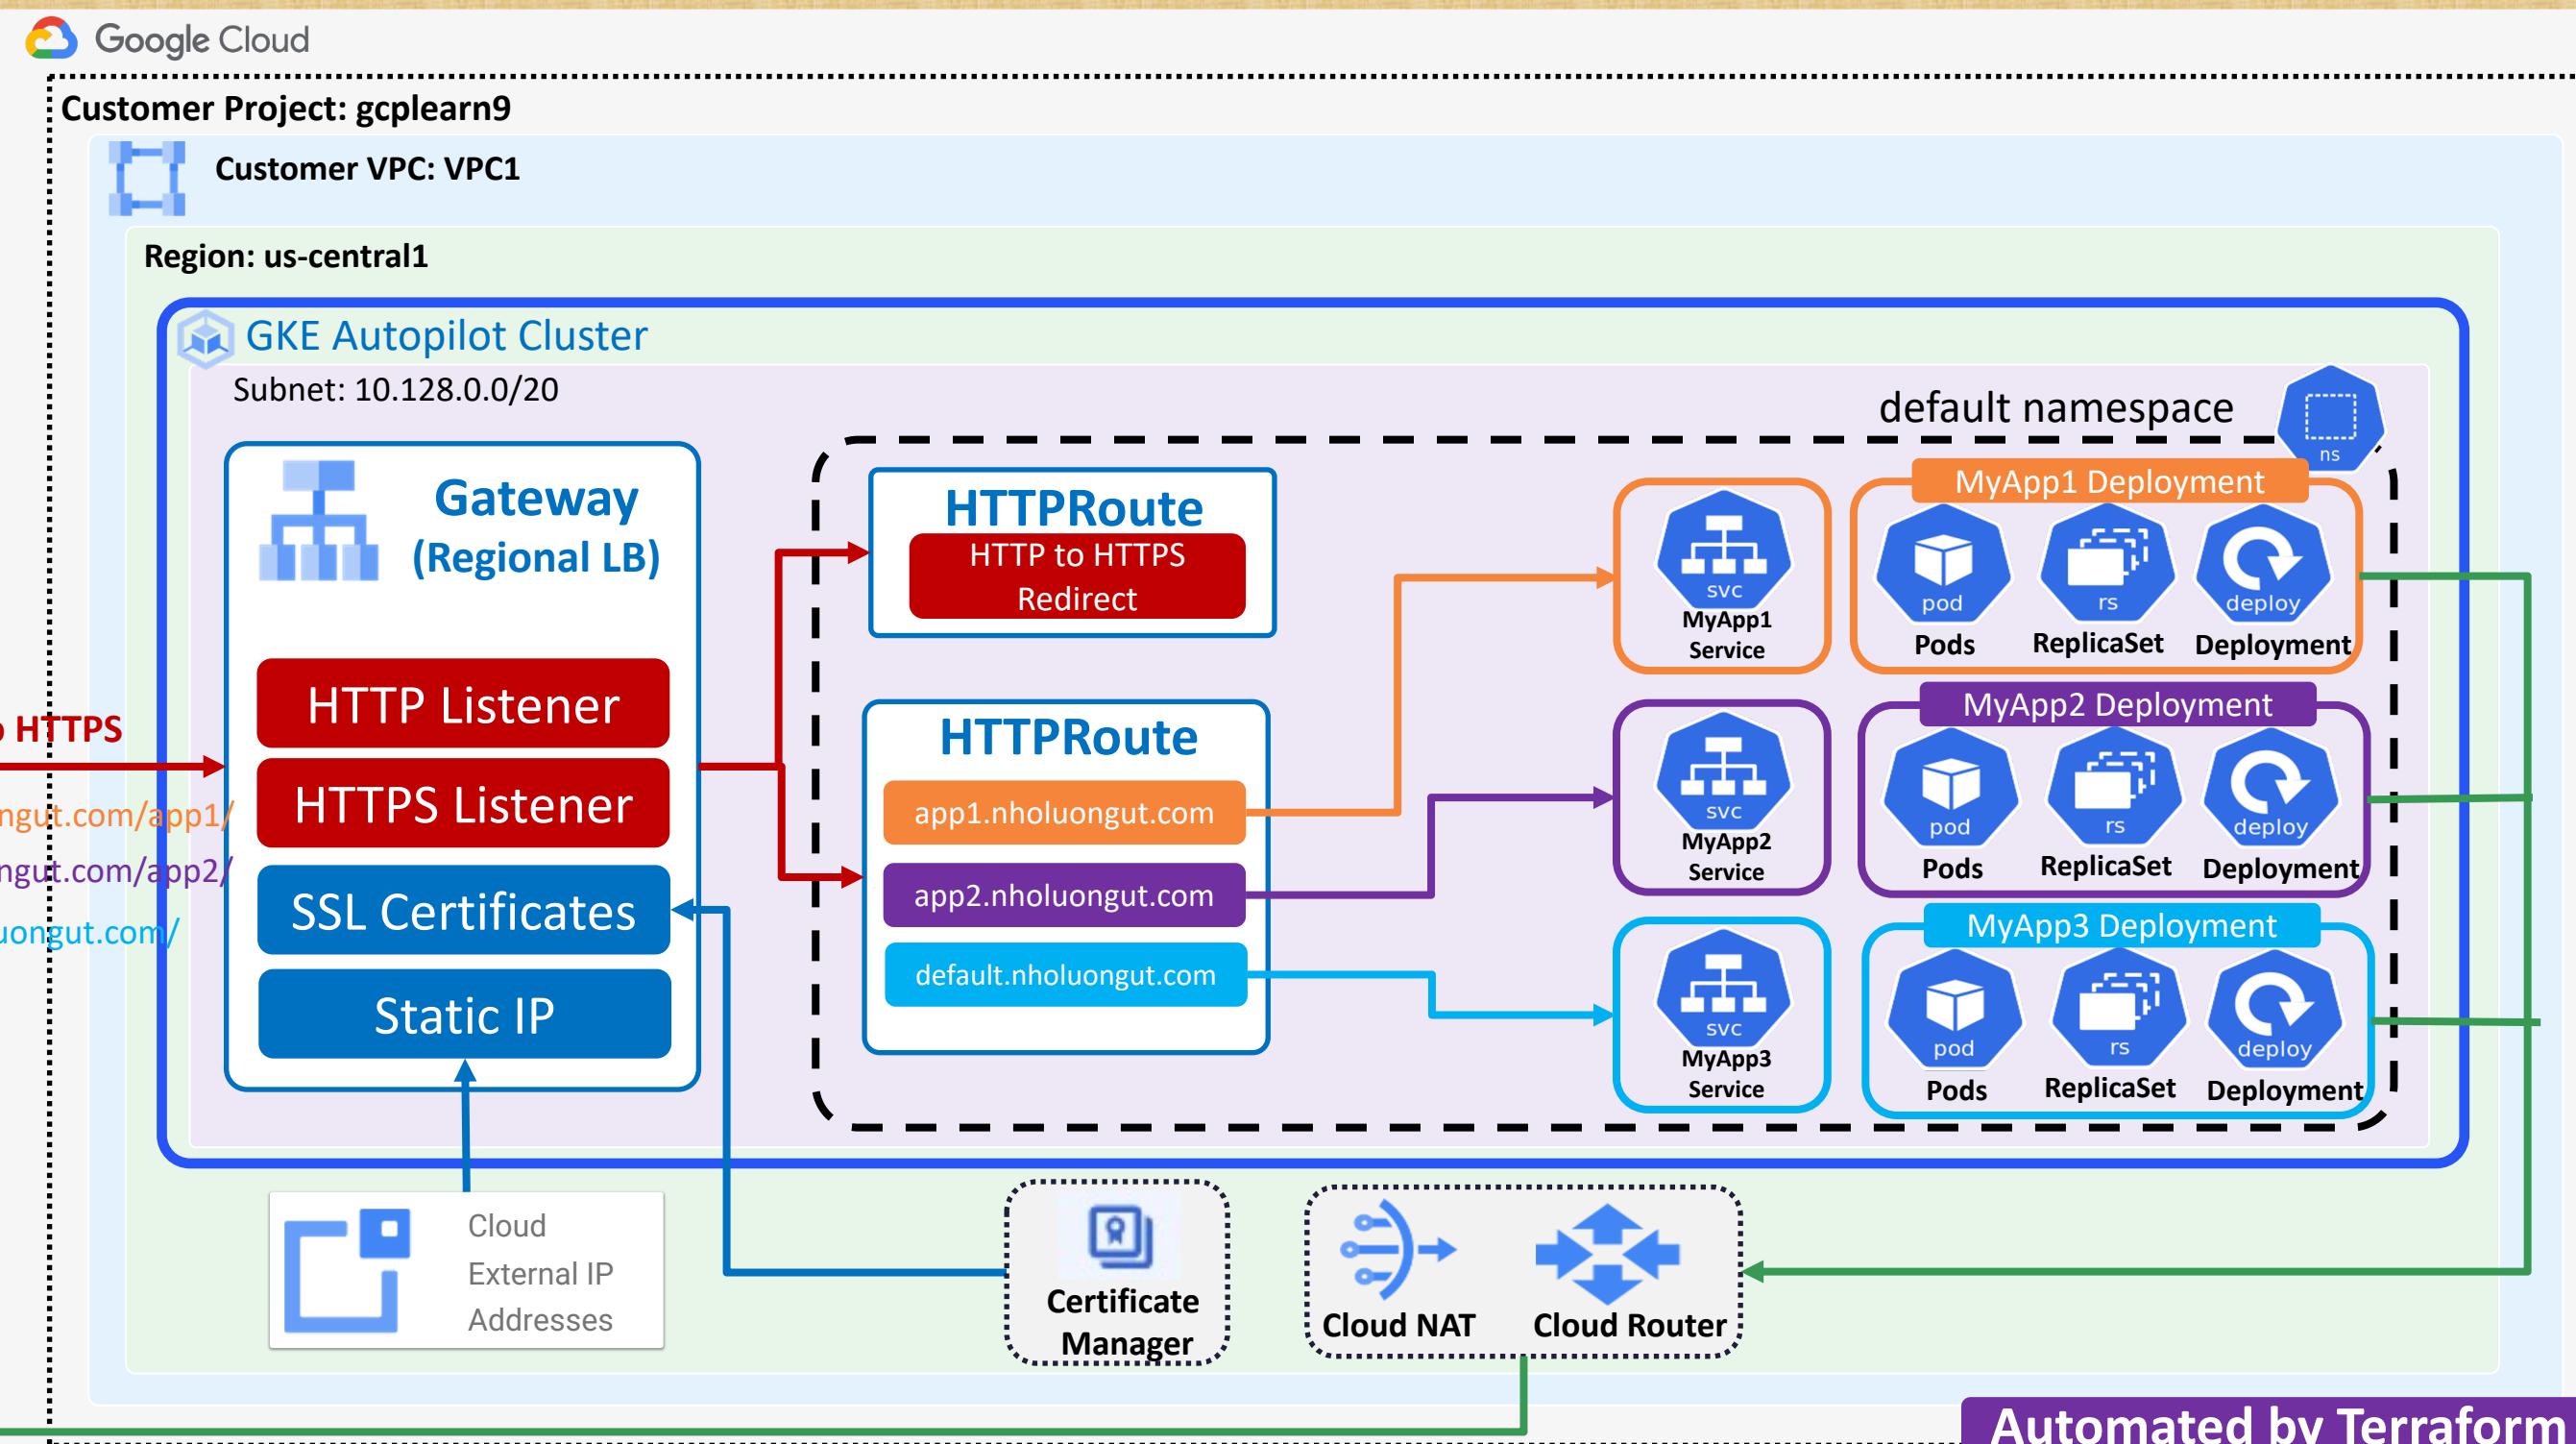
# Google Kubernetes Engine Gateway API

## Cloud Application Load Balancer

**Demo:** GKE Autopilot cluster + Gateway + Static IP +  
Self-signed SSL with Certificate Manager + HTTP to  
HTTPS Redirect + Domain name based Routing



# Gateway API Basics + Static IP + SSL Certificate Manager + HTTP to HTTPS + Domain name-based Routing



- ✓ p2-regional-k8sresources-yaml
  - > self-signed-ssl
    - ! c1-01-myapp1-deployment.yaml
    - ! c1-02-myapp1-clusterip-service.yaml
  - ! c2-01-myapp2-deployment.yaml
  - ! c2-02-myapp2-clusterip-service.yaml
  - ! c3-01-myapp3-deployment.yaml
  - ! c3-02-myapp3-clusterip-service.yaml
  - ! c4-01-gateway.yaml
  - ! c4-02-gateway-http-to-https-route.yaml
  - ! c4-03-gateway-app1-http-route.yaml
  - ! c4-04-gateway-app2-http-route.yaml
  - ! c4-05-gateway-app3-http-route.yaml

# What are we going to learn?

**Demo: Project-2: Gateway + Regional Static IP + Self-signed SSL Certificate Manager + HTTP to HTTPS + Domain Name Routing**

C1 to  
C4-02

c4-01

c4-02

c4-03

c4-04

c4-05

MyApp1, 2, 3 Deployments and Cluster IP Services

Gateway – No changes

HTTP to HTTPS redirect – No changes

App1 HTTP Route with domain name as  
**app1.stackimplify.com**

App2 HTTP Route with domain name as  
**app2.stackimplify.com**

App3 HTTP Route with domain name as  
**default.stackimplify.com**

## Project-3: Terraform Manifests

- ✓ p3-regional-k8sresources-terraform-manifests
  - > self-signed-ssl
  - ✗ c1-versions.tf
  - ✗ c2-01-variables.tf
  - ✗ c2-02-local-values.tf
  - ✗ c3-01-remote-state-datasource.tf
  - ✗ c3-02-providers.tf
  - ✗ c4-01-myapp1-deployment.tf
  - ✗ c4-02-myapp1-clusterip-service.tf
  - ✗ c4-03-myapp2-deployment.tf
  - ✗ c4-04-myapp2-clusterip-service.tf
  - ✗ c4-05-myapp3-deployment.tf
  - ✗ c4-06-myapp3-clusterip-service.tf
  - ✗ c5-01-gateway.tf
  - ✗ c5-02-gateway-http-to-https-route.tf
  - ✗ c5-03-gateway-app1-http-route.tf
  - ✗ c5-04-gateway-app2-http-route.tf
  - ✗ c5-05-gateway-app3-http-route.tf
  - ✗ c6-static-ip.tf
  - ✗ c7-certificate-manager.tf
  - ✗ terraform.tfvars

# What are we going to learn?

Demo: Project-3: Gateway + Regional Static IP + Self-signed SSL  
Certificate Manager + HTTP to HTTPS + Domain Name Routing

NC

No changes c1 to c5-02, c6, c7

c5-03

App1 HTTP Route with domain name as  
**app1.stackimplify.com**

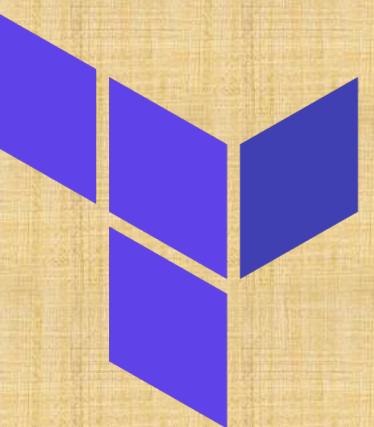
c5-04

App2 HTTP Route with domain name as  
**app2.stackimplify.com**

c5-05

App3 HTTP Route with domain name as  
**default.stackimplify.com**

Demo-08

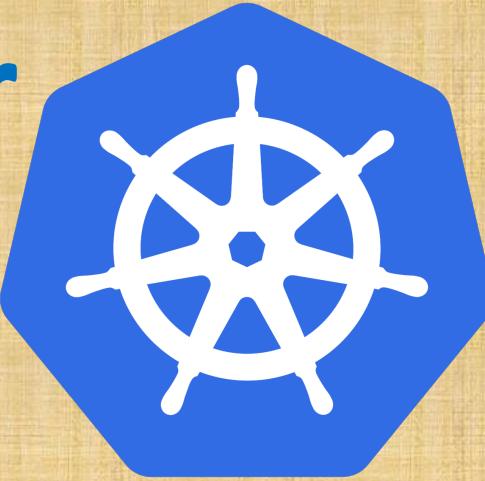
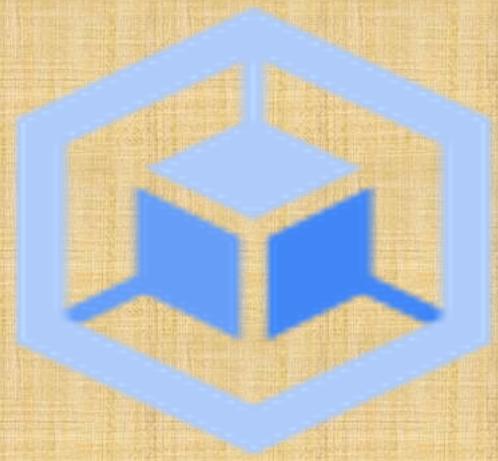


# Google Kubernetes Engine

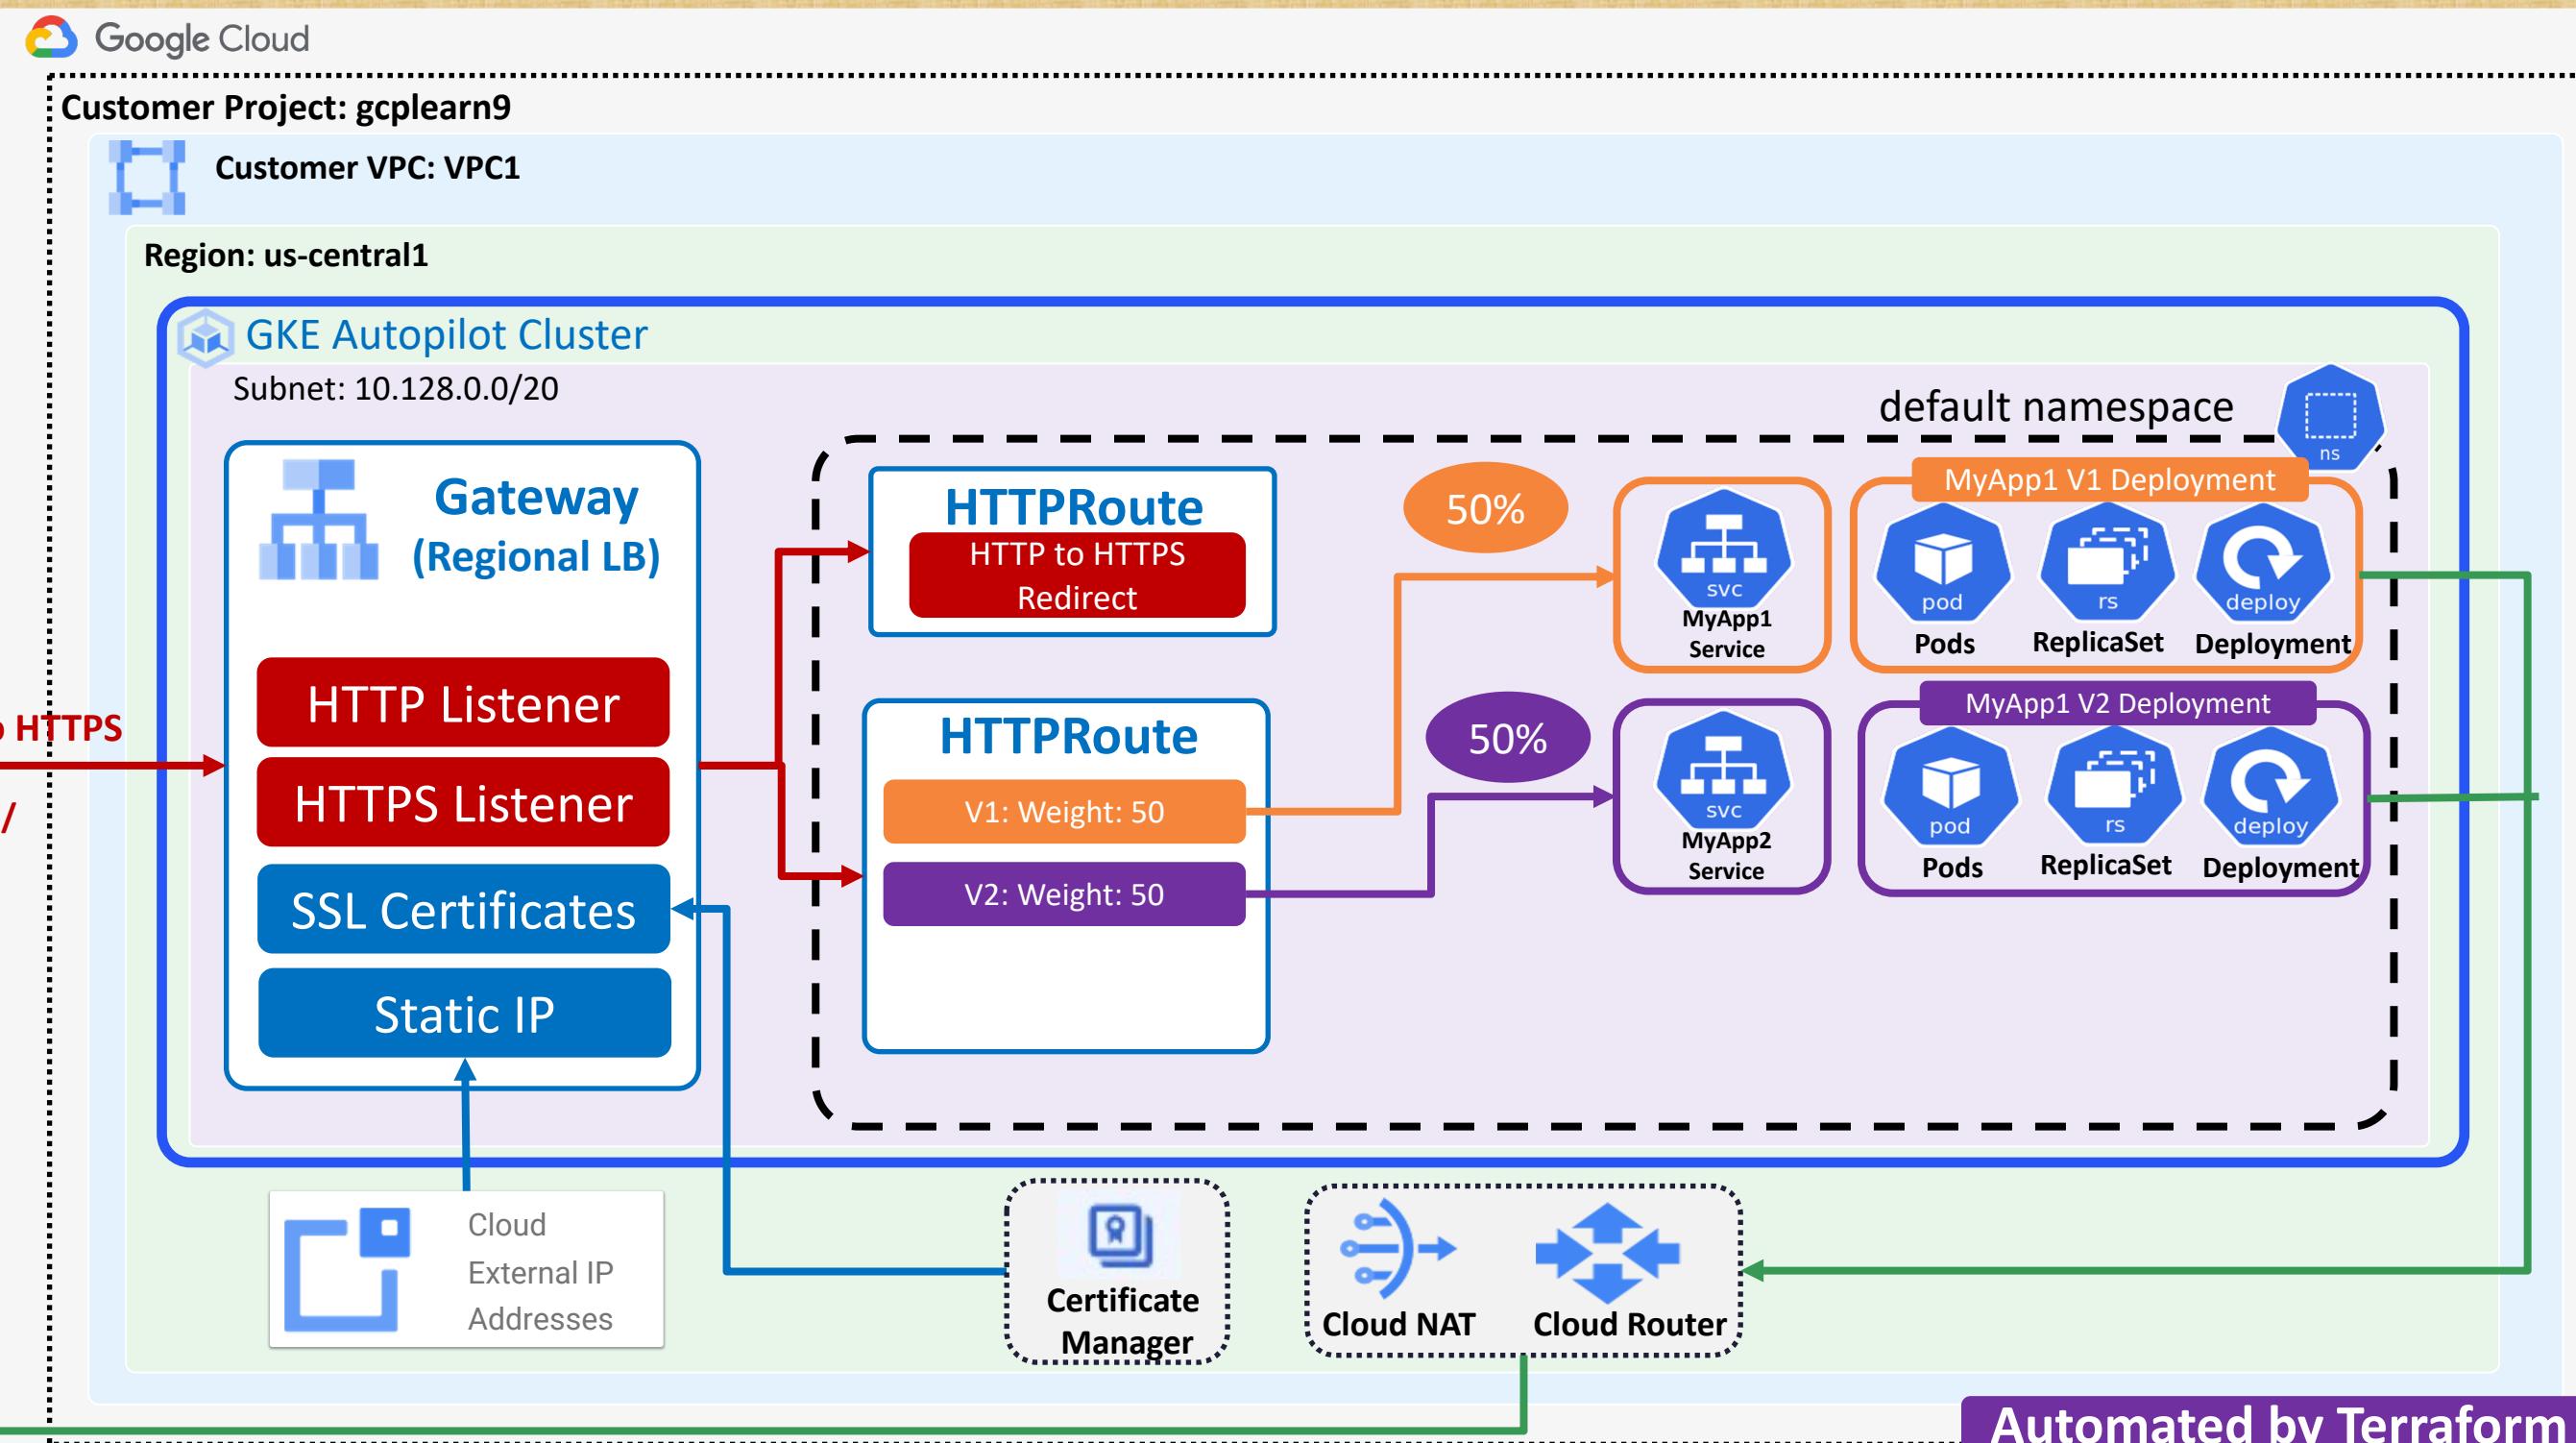
Gateway API

## Cloud Application Load Balancer

**Demo:** GKE Autopilot cluster + Gateway + Static IP +  
Self-signed SSL with Certificate Manager + HTTP to  
HTTPS Redirect + Traffic Splitting



# Gateway API Basics + Static IP + SSL Certificate Manager + HTTP to HTTPS + Traffic Splitting



## Project-2: YAML Manifests

✓ p2-regional-k8sresources-yaml

> self-signed-ssl

- ! c1-01-myapp1-v1-deployment.yaml
- ! c1-02-myapp1-v1-clusterip-service.yaml
- ! c2-01-myapp1-v2-deployment.yaml
- ! c2-02-myapp1-v2-clusterip-service.yaml
- ! c4-01-gateway.yaml
- ! c4-02-gateway-http-to-https-route.yaml
- ! c4-03-gateway-app1-http-route.yaml

# What are we going to learn?

Demo: Project-2: Gateway + Regional Static IP + Self-signed SSL  
Certificate Manager + HTTP to HTTPS + Traffic Splitting

NC

NO Changes c4-01, c4-02,

c1-01

MyApp1 V1 Kubernetes Deployment

c1-02

MyApp1 V1 Cluster IP Service

c2-01

MyApp1 V2 Kubernetes Deployment

c2-02

MyApp1 V2 Cluster IP Service

c4-03

HTTP Route: V1 and V2 Traffic Splitting

## Project-3: Terraform Manifests

- ✓ p3-regional-k8sresources-terraform-manifests
  - > self-signed-ssl
  - ↳ c1-versions.tf
  - ↳ c2-01-variables.tf
  - ↳ c2-02-local-values.tf
  - ↳ c3-01-remote-state-datasource.tf
  - ↳ c3-02-providers.tf
  - ↳ c4-01-myapp1-v1-deployment.tf
  - ↳ c4-02-myapp1-v1-clusterip-service.tf
  - ↳ c4-03-myapp2-v2-deployment.tf
  - ↳ c4-04-myapp2-v2-clusterip-service.tf
  - ↳ c5-01-gateway.tf
  - ↳ c5-02-gateway-http-to-https-route.tf
  - ↳ c5-03-gateway-app1-http-route.tf
  - ↳ c6-static-ip.tf
  - ↳ c7-certificate-manager.tf
  - ↳ terraform.tfvars

# What are we going to learn?

Demo: Project-3: Gateway + Regional Static IP + Self-signed SSL  
Certificate Manager + HTTP to HTTPS + Traffic Splitting

NC

No changes c1 to c5-02, c6, c7

c4-01

MyApp1 V1 Kubernetes Deployment

c4-02

MyApp1 V1 Cluster IP Service

c4-03

MyApp1 V2 Kubernetes Deployment

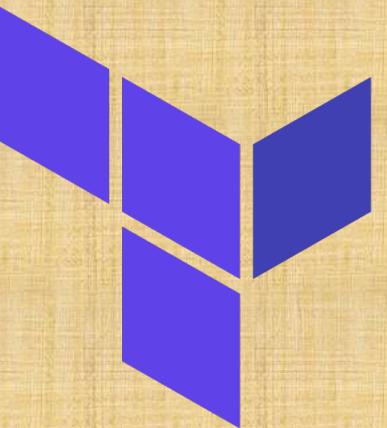
c4-04

MyApp1 V2 Cluster IP Service

c5-03

HTTP Route: V1 and V2 Traffic Splitting

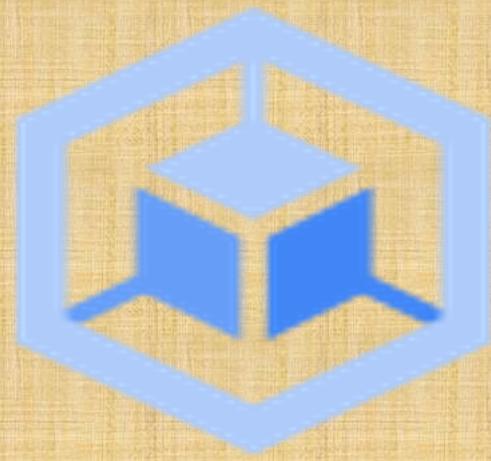
Demo-09



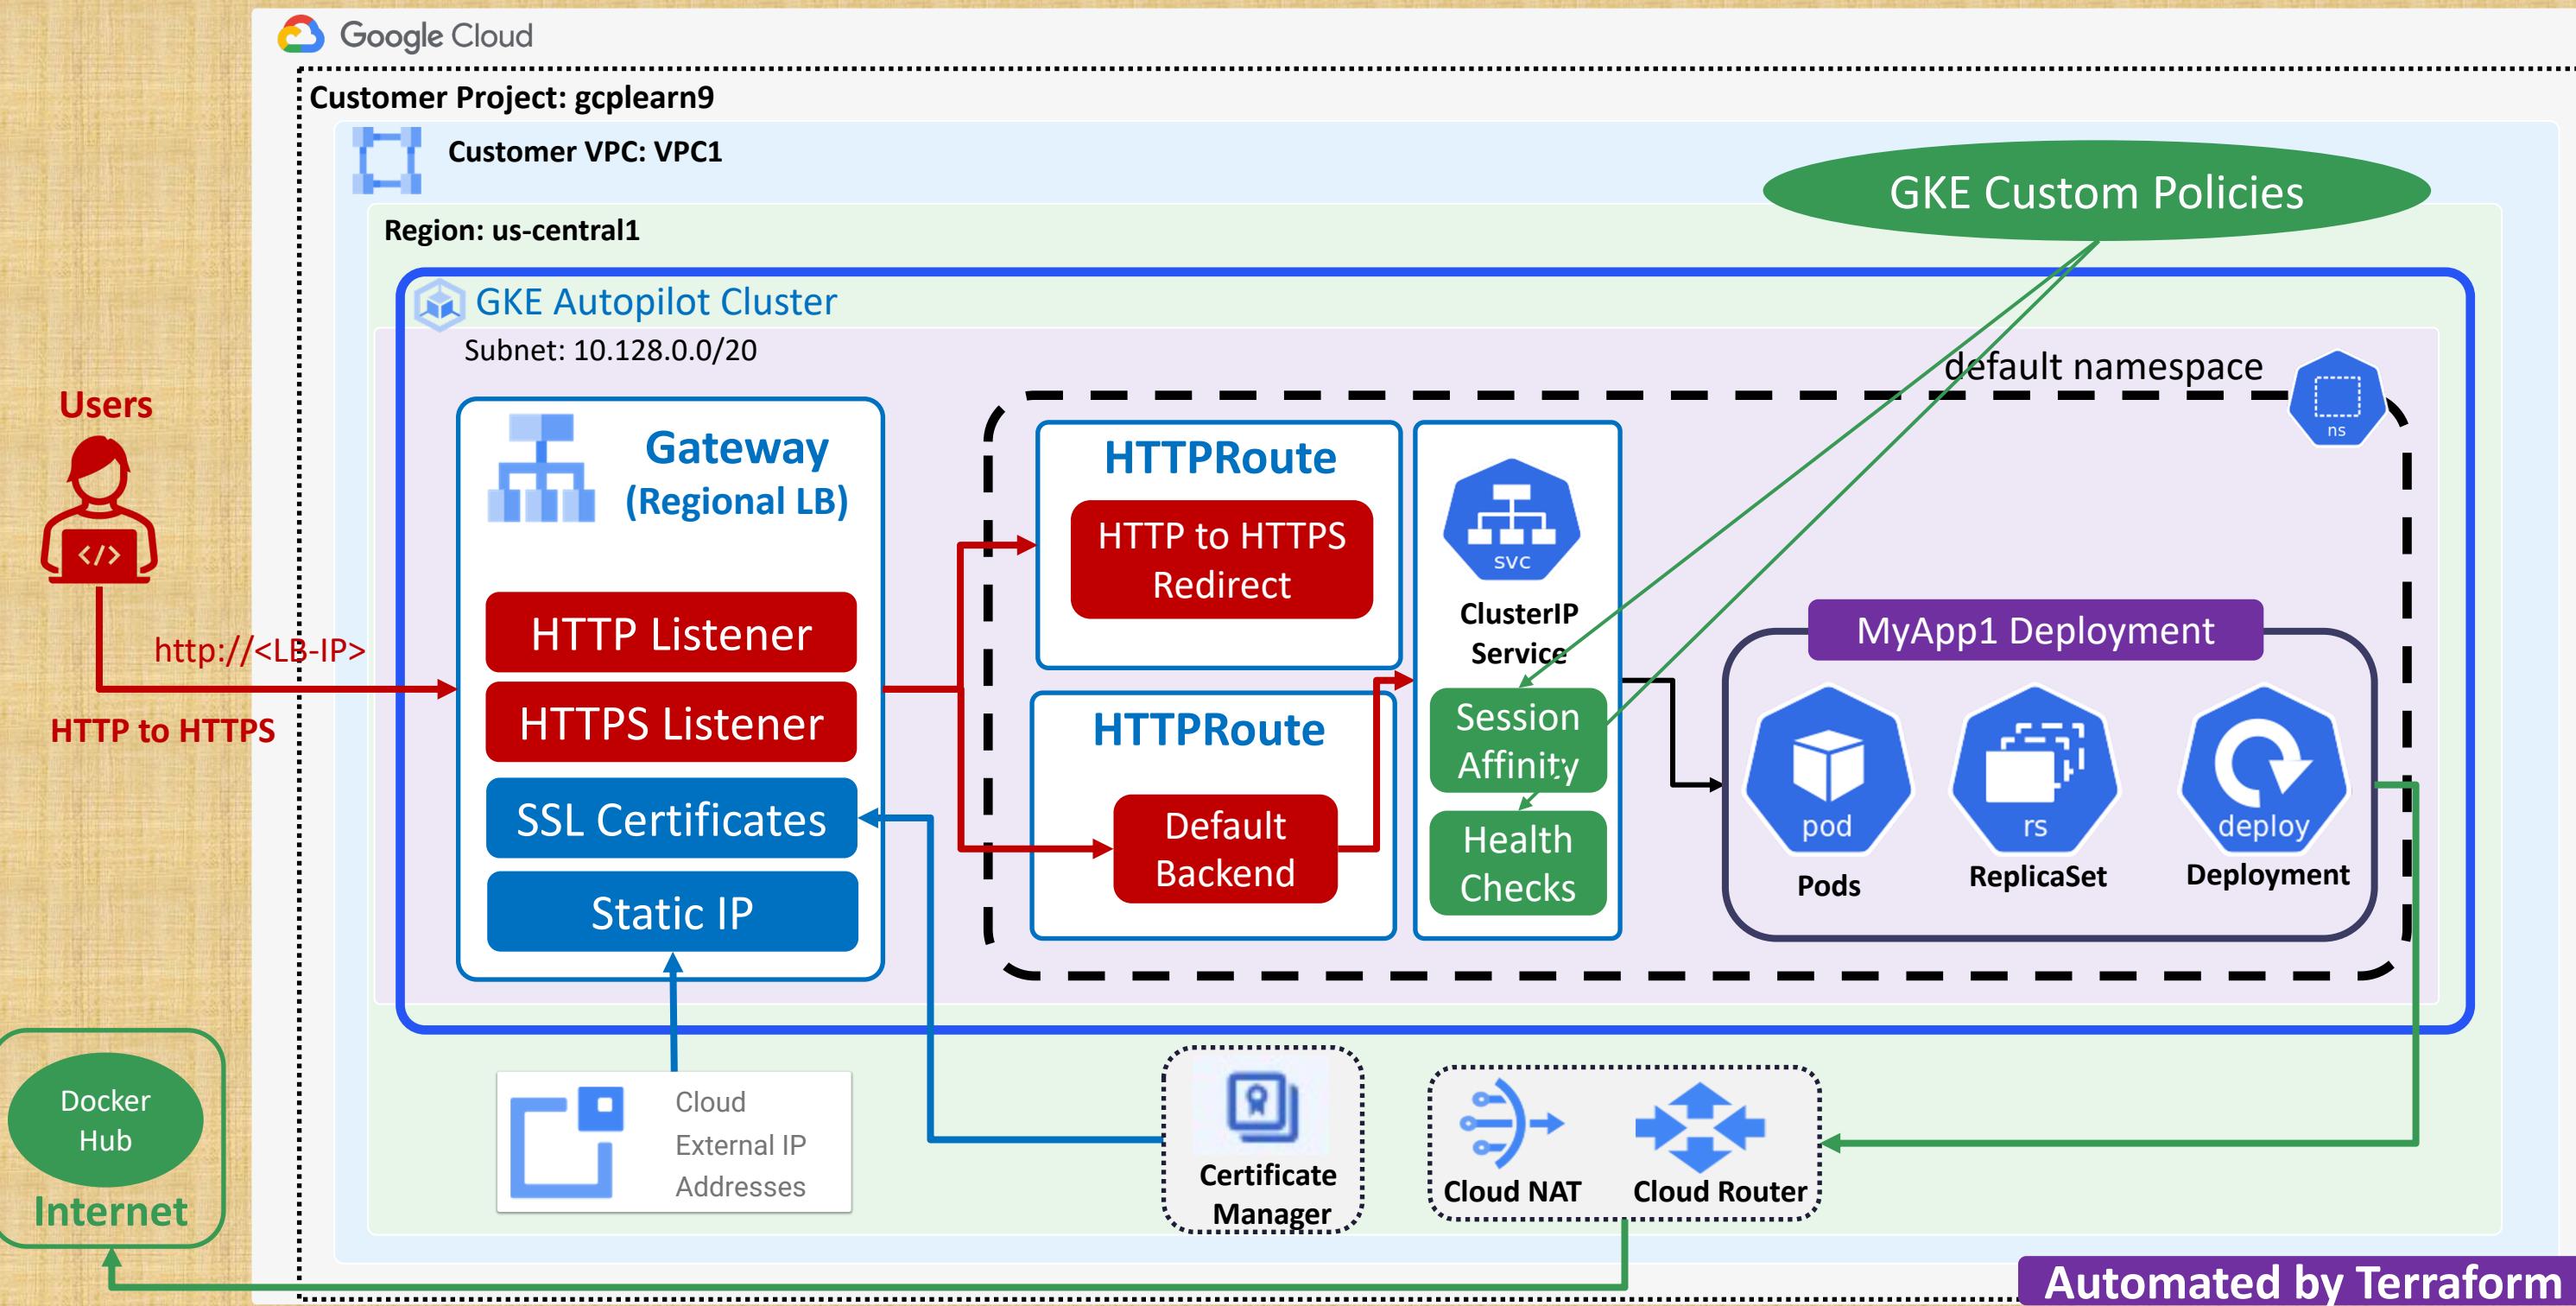
# Google Kubernetes Engine Gateway API

## Cloud Application Load Balancer

**Demo: GKE Autopilot cluster + Gateway + Static IP +  
Self-signed SSL with Certificate Manager + HTTP to  
HTTPS Redirect + Health Checks + Session Affinity**



# Gateway API + Static IP + SSL Certificate Manager + HTTP to HTTPS + Health Checks + Session Affinity



```
✓ p2-regional-k8sresources-yaml
  > self-signed-ssl
  ! c1-01-myapp1-deployment.yaml
  ! c1-02-myapp1-clusterip-service.yaml
  ! c1-03-myapp1-healthcheck.yaml
  ! c1-04-session-affinity.yaml
  ! c2-01-gateway.yaml
  ! c2-02-gateway-http-route.yaml
  ! c2-03-gateway-http-to-https-route.yaml
```

# What are we going to learn?

Demo: Project-2: Gateway + Regional Static IP + Self-signed SSL Certificate Manager + HTTP to HTTPS + Health Checks + Session Affinity

NC

NO Changes c1-01, c1-02, c2-01, c2-02, c2-04

c1-03

MyApp1 Health Check Policy

c1-05

MyApp1 Session Affinity Policy

## Project-3: Terraform Manifests

- ✓ p3-regional-k8sresources-terraform-manifests
  - > self-signed-ssl
  - c1-versions.tf
  - c2-01-variables.tf
  - c2-02-local-values.tf
  - c3-01-remote-state-datasource.tf
  - c3-02-providers.tf
  - c4-01-myapp1-deployment.tf
  - c4-02-myapp1-clusterip-service.tf
  - c4-03-myapp1-healthcheck.tf
  - c4-04-myapp1-session-affinity.tf
  - c5-01-gateway.tf
  - c5-02-gateway-http-to-https-route.tf
  - c5-03-gateway-http-route.tf
  - c6-static-ip.tf
  - c7-certificate-manager.tf
  - terraform.tfvars

# What are we going to learn?

Demo: Project-3: Gateway + Regional Static IP + Self-signed SSL Certificate Manager + HTTP to HTTPS + Health Checks + Session Affinity

NC

No changes c1 to c4-02, c5-01 to c7

c4-03

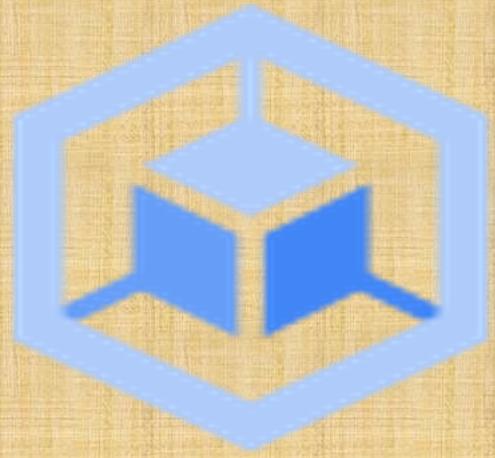
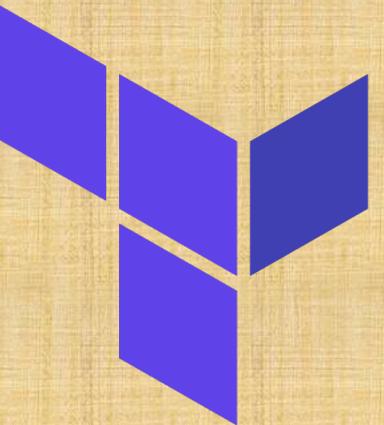
MyApp1 Health Check Policy

c4-04

MyApp1 Session Affinity Policy

Demo-10

# Google Kubernetes Engine Gateway API Cloud Domains + Cloud DNS





Demo

# Google Cloud Networking Cloud Domains



# Google Cloud Domains



- **Cloud Domains:** Primarily used for **Domain Registration and management**
- **Key Benefits**
  - Register **new domains**
  - Bills **for purchasing/renewing domains** will use the same **Cloud Billing account**
  - **Automatic renewal** of registered domains
  - Let's you manage **domain registrations per project**, not per user
  - **Supports DNSSEC** which protects your domains from spoofing and cache poisoning attacks
  - **Tightly integrated** with Google Cloud DNS (Domain Name system) for creating and managing DNS records

# Google Cloud Domains



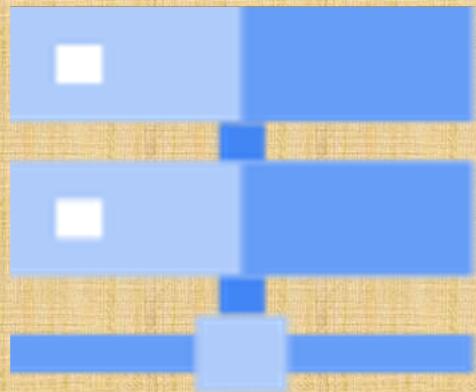
- **Important Note:** From September 7, 2023 onwards Squarespace acquired all domain registrations and related customer accounts from Google Domains
- Cloud Domains uses [Google Domains](#) as underlying domain registrar
- Now after squarespace acquisition, [Cloud Domains uses Squarespace Domains](#) as underlying registrar
- **How does this impact Cloud Domains in Google Cloud ?**
  - **Complete Faq:** <https://cloud.google.com/domains/docs/faq>
  - In a shorter note, [google supports Cloud Domains even](#) after this acquisition
  - **Following features in Cloud Domain will work as-is**
    - [Searching and registering](#) new domains
    - [Renewing](#) existing domains
    - Update your contact details and DNS settings
    - [Transfer](#) a registered domain to another registrar (GoDaddy, Namecheap, AWS Route53)
    - Transfer domain from [another registrar to Cloud Domain - NOT POSSIBLE](#)

Additional Reference: <https://cloud.google.com/domains/docs/deprecations/feature-deprecations>

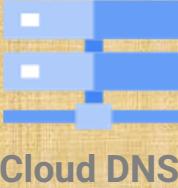


Demo

# Google Cloud Networking Cloud DNS



# Google Cloud DNS



- **Cloud DNS:** High-performance, resilient, global Domain Name System (DNS) service
- Cloud DNS translates requests for domain names like `google.com` into IP addresses like `74.125.200.113`

```
nholu@nholuongs-MacBook-Pro:~ % nslookup google.com
Server: 8.8.8.8
Address: 8.8.8.8#53

Non-authoritative answer:
Name: google.com
Address: 74.125.200.113
Name: google.com
Address: 74.125.200.102
Name: google.com
Address: 74.125.200.101
Name: google.com
Address: 74.125.200.100
Name: google.com
Address: 74.125.200.138
Name: google.com
Address: 74.125.200.139
```

# Google Cloud DNS - Terminology

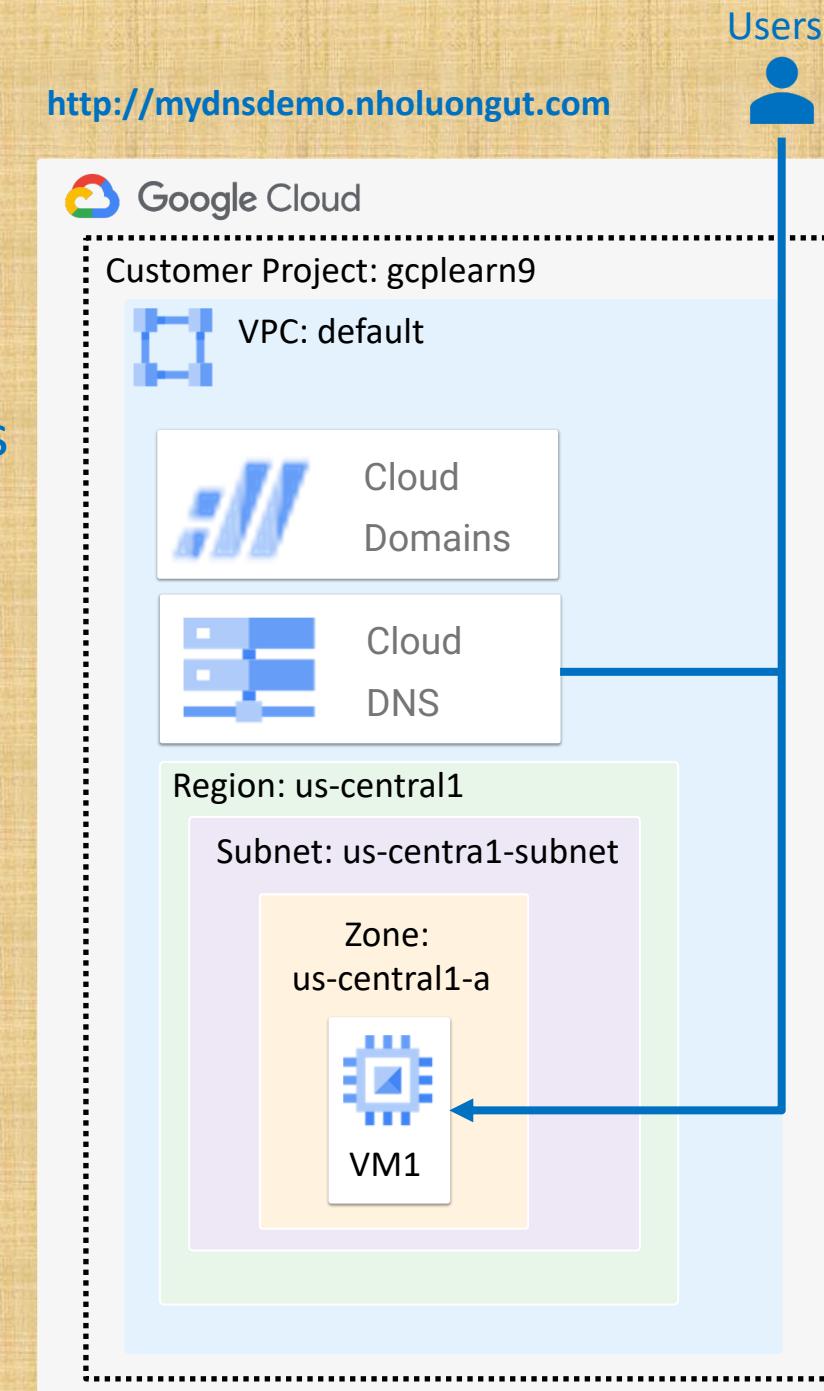
- **DNS Zones:** It is a **container of DNS records** for the same DNS name suffix (Example suffix: google.com)

- **Public Zone**

- **Visible to the internet**
- Primarily used for **DNS management** of your **internet facing applications**
- DNSSEC (DNS Security) protects your domains from **spoofing and cache poisoning attacks**

- **Private Zone**

- Contains DNS records that are **only visible internally** within your Google Cloud networks
- Easy to manage **internal DNS solution** for all our internal needs (For Internal applications, VM Instances etc)



# Google Cloud DNS - Terminology

- **Record Sets:** actual DNS records
  - DNS Name to IP Address mapping
  - We have different record types, but primarily we use **Address record (A) type** which maps hostnames to IPv4 address

### DNS Record Set

Create record set

DNS name: myapp1 .kalyanreddydaida.com.

Resource record type: A

TTL \*: 5

TTL unit: minutes

IPv4 Address \*

IPv4 Address 1 \*: 108.157.238.34

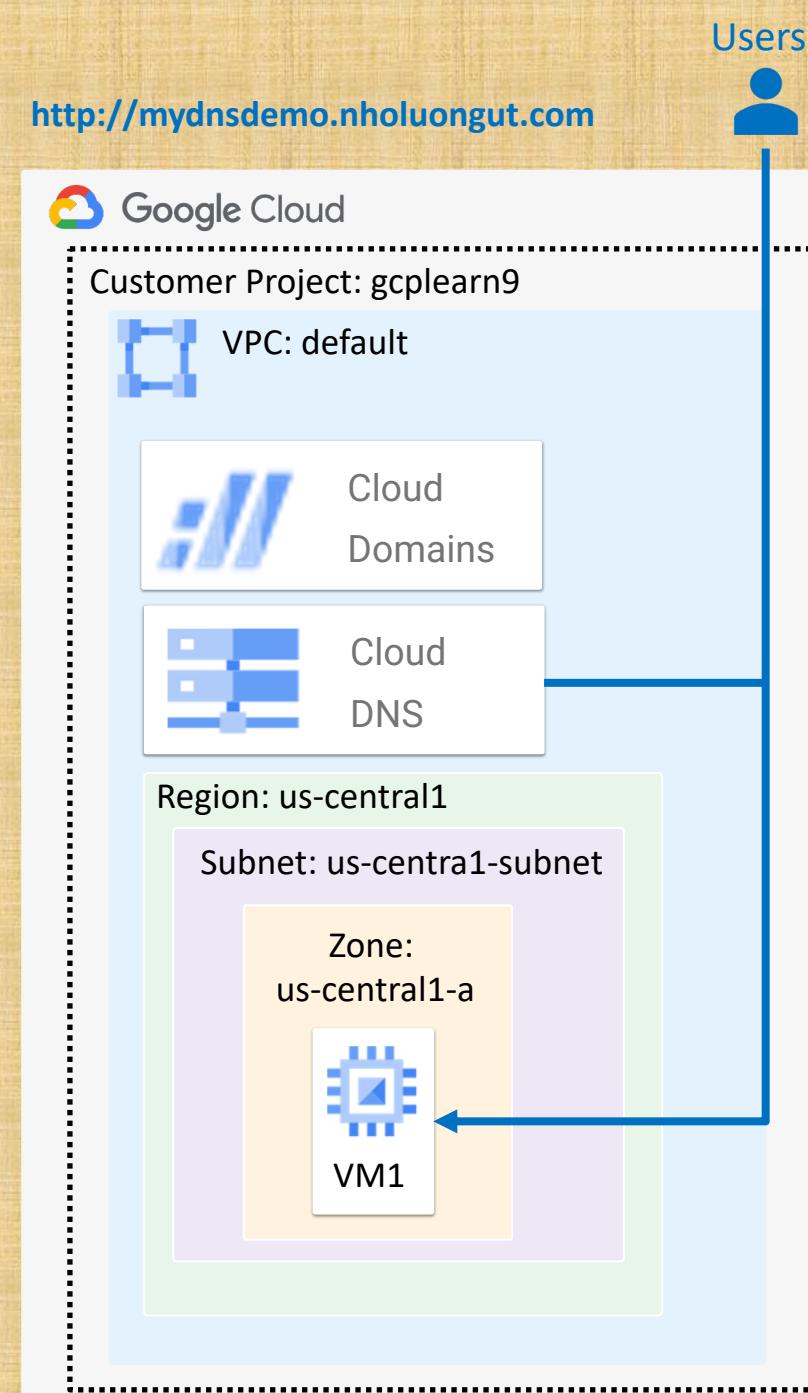
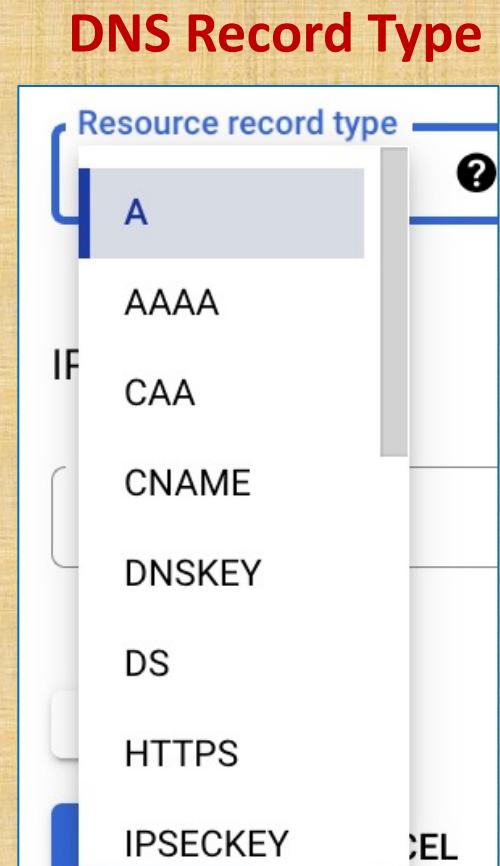
SELECT IP ADDRESS

Example: 192.0.2.91

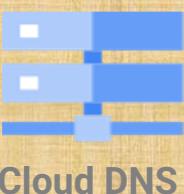
+ ADD ITEM

CREATE CANCEL

[https://cloud.google.com/dns/docs/records-overview#supported\\_dns\\_record\\_types](https://cloud.google.com/dns/docs/records-overview#supported_dns_record_types)



# Google Cloud DNS - Features



Cloud DNS



Users

## • Integration with Cloud IAM

- Provides secure domains management with **full control and visibility** for domain resources

## • Integration with Cloud Logging

- Private DNS logs a record for **every DNS query** in Cloud Logging
- We can **view and export logs** to supported destinations

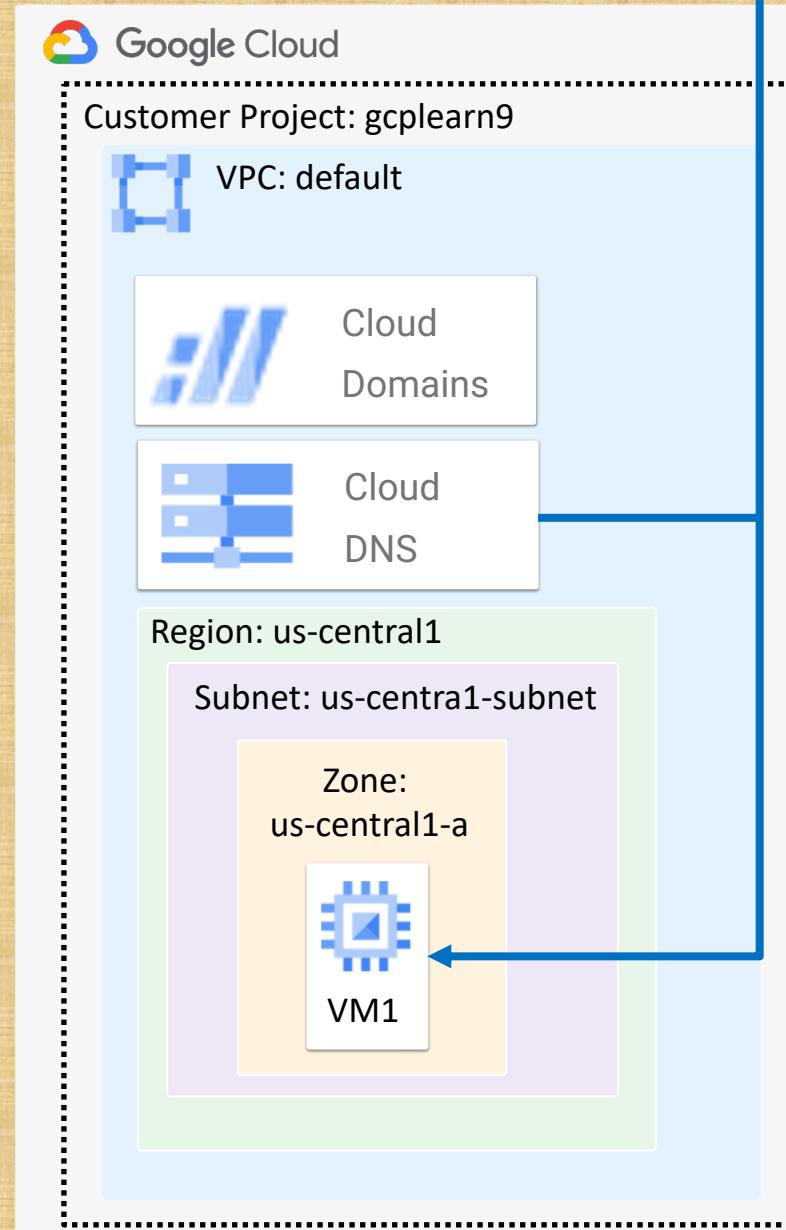
## • Fast anycast Nameservers

- **Anycast:** Uses the **nearest nameserver to users** for DNS lookup and resolution
- Provides **very high availability** across globe and **low latency**

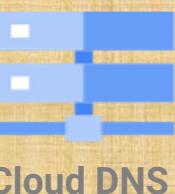
## • DNS registration and management

- Cloud DNS is **tightly integrated** with Cloud Domains which takes care of DNS registration and management

<http://mydnsdemo.nholuongut.com>



# Google Cloud DNS - Features



Cloud DNS



## • Container-native Cloud DNS

- Natively integrated with Google Kubernetes Engine (GKE)
- Provides in-cluster Service DNS resolution
- Provides high-throughput, scalable DNS resolution for every GKE node

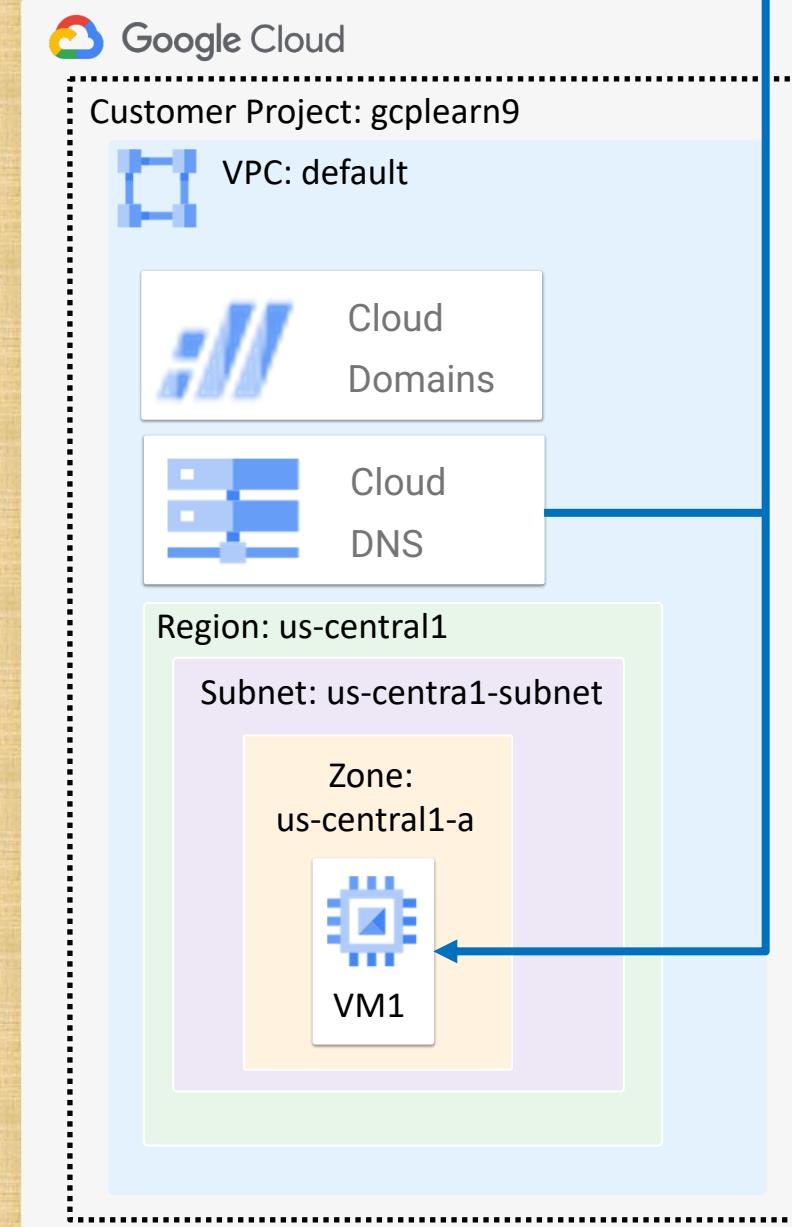
## • DNS Peering

- Sharing DNS data
- All or a portion of DNS namespace can be configured to be sent from one network to another
- Will respect all DNS configurations defined in peered network

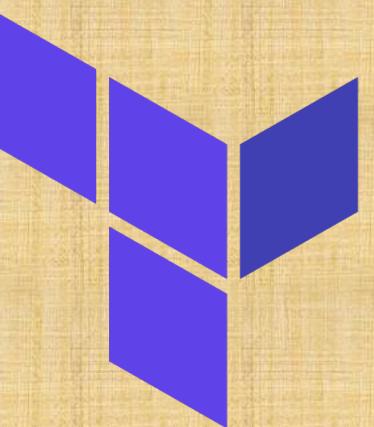
## • DNS Forwarding

- Primarily used for hybrid-cloud architecture
- Forward DNS queries from Cloud DNS to on-premise DNS servers where actual DNS resolution takes place

<http://mydnsdemo.nholuongut.com>



Demo-11

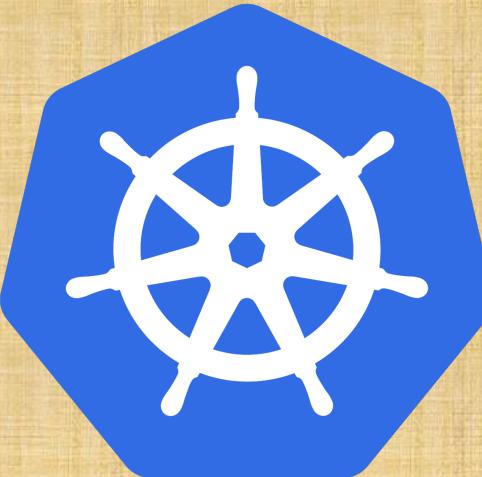
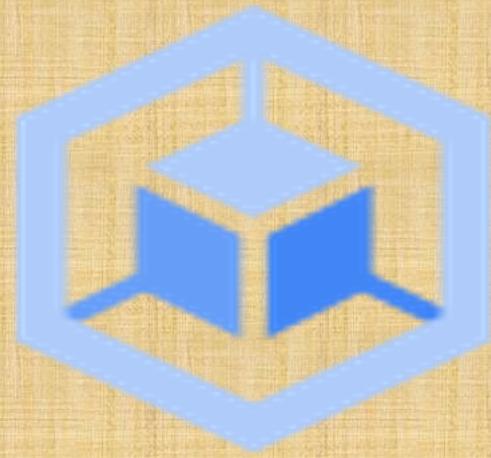


# Google Kubernetes Engine

Gateway API

Cloud Domains + Cloud DNS

Demo: GKE Autopilot cluster + Gateway + Static IP + PROD  
GRADE SSL with Certificate Manager + HTTP to HTTPS Redirect +  
Health Checks + Session Affinity + Cloud Domains + Cloud DNS



# Gateway API + Static IP + SSL Certificate Manager + HTTP to HTTPS + Health Checks + Session Affinity + Cloud Domains + Cloud DNS



Customer Project: gcplearn9

Automated by Terraform



Customer VPC: VPC1

Region: us-central1



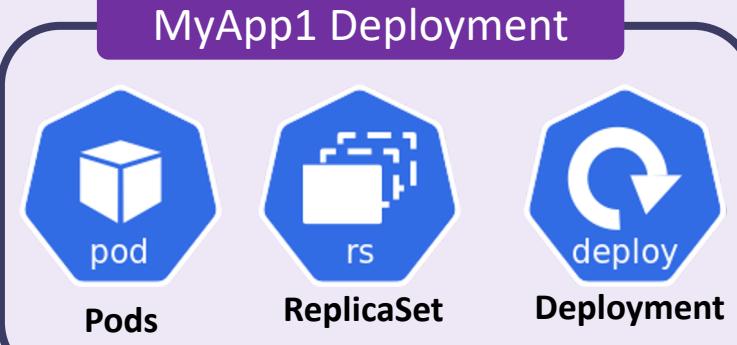
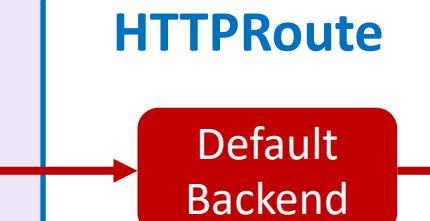
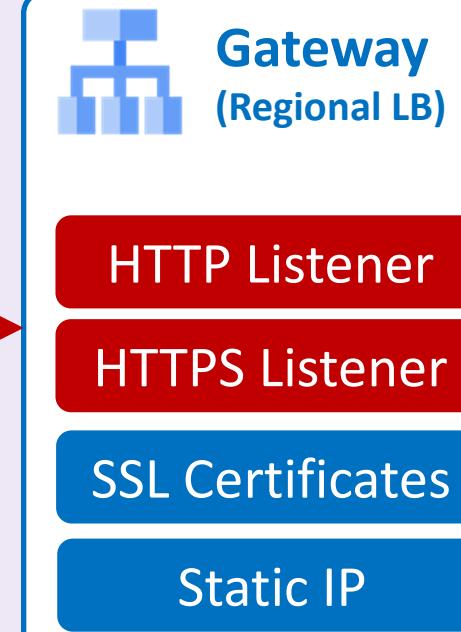
GKE Autopilot Cluster

Subnet: 10.128.0.0/20

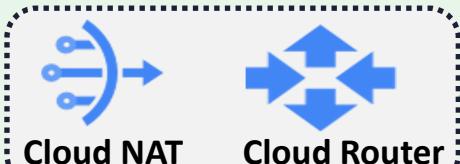
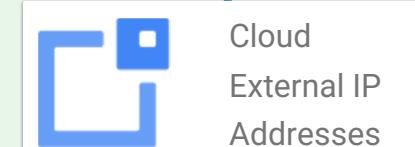
Users



http://<DNS-NAME>  
HTTP to HTTPS



default namespace



## Project-2: Terraform Manifests

- ✓ p2-regional-k8sresources-terraform-manifests
  - > self-signed-ssl
  - ✗ c1-versions.tf
  - ✗ c2-01-variables.tf
  - ✗ c2-02-local-values.tf
  - ✗ c3-01-remote-state-datasource.tf
  - ✗ c3-02-providers.tf
  - ✗ c4-01-myapp1-deployment.tf
  - ✗ c4-02-myapp1-clusterip-service.tf
  - ✗ c4-03-myapp1-healthcheck.tf
  - ✗ c4-04-myapp1-session-affinity.tf
  - ✗ c5-01-gateway.tf
  - ✗ c5-02-gateway-http-to-https-route.tf
  - ✗ c5-03-gateway-http-route.tf
  - ✗ c6-static-ip.tf
  - ✗ c7-certificate-manager.tf
  - ✗ c8-cloud-dns.tf
  - ✗ terraform.tfvars

# What are we going to learn?

Demo: Project-2: Gateway + Static IP + Prod grade SSL Certificate Manager + HTTP to HTTPS + Health Checks + Session Affinity + Cloud DNS + Cloud Domains

NC

No changes c1 to c6

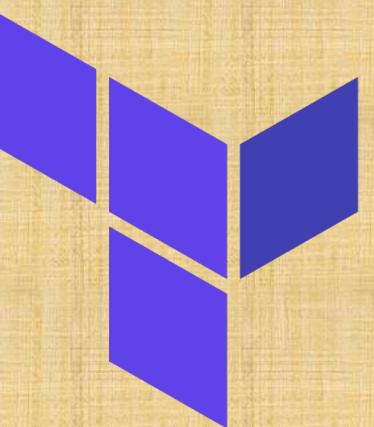
C7

Create production grade SSL certificates using Certificate Manager with DNS Authorization

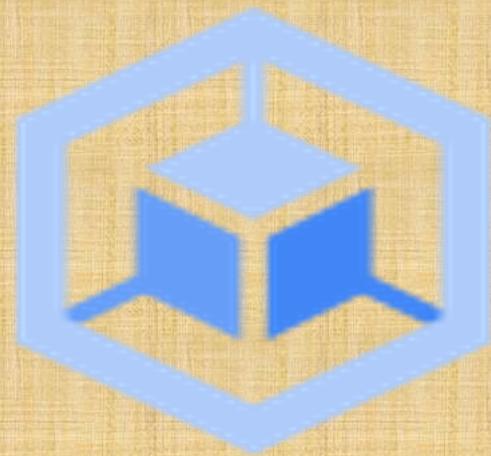
C8

Create record sets in Cloud DNS

Demo-12



# Google Kubernetes Engine



## Gateway API

## AWS Route53 + Cloud DNS

Demo: GKE Autopilot cluster + Gateway + Static IP + PROD GRADE SSL with Certificate Manager + HTTP to HTTPS Redirect + Health Checks + Session Affinity + AWS Route53 + Cloud DNS





Customer Project: gcplearn9

Automated by Terraform



Customer VPC: VPC1

Region: us-central1



GKE Autopilot Cluster

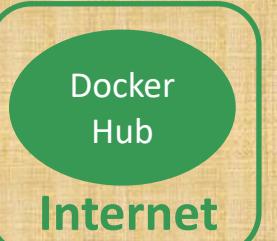
Subnet: 10.128.0.0/20

Users



http://&lt;DNS-NAME&gt;

AWS Route53



Gateway (Regional LB)

HTTP Listener

HTTPS Listener

SSL Certificates

Static IP

HTTPRoute

HTTP to HTTPS Redirect

HTTPRoute

Default Backend



ClusterIP Service

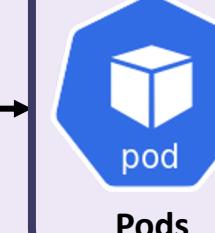
Session Affinity

Health Checks

default namespace



MyApp1 Deployment



Pods



ReplicaSet



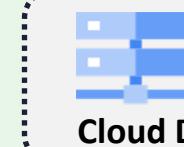
Deployment



Cloud External IP Addresses



Production grade SSL certificates



Cloud DNS



Cloud NAT



Cloud Router

## Project-2: Terraform Manifests

- ✓ p2-regional-k8sresources-terraform-manifests
  - > self-signed-ssl
  - ✗ c1-versions.tf
  - ✗ c2-01-variables.tf
  - ✗ c2-02-local-values.tf
  - ✗ c3-01-remote-state-datasource.tf
  - ✗ c3-02-providers.tf
  - ✗ c4-01-myapp1-deployment.tf
  - ✗ c4-02-myapp1-clusterip-service.tf
  - ✗ c4-03-myapp1-healthcheck.tf
  - ✗ c4-04-myapp1-session-affinity.tf
  - ✗ c5-01-gateway.tf
  - ✗ c5-02-gateway-http-to-https-route.tf
  - ✗ c5-03-gateway-http-route.tf
  - ✗ c6-static-ip.tf
  - ✗ c7-certificate-manager.tf
  - ✗ c8-cloud-dns.tf
  - ✗ terraform.tfvars

# What are we going to learn?

Demo: Project-2: Gateway + Static IP + Prod Grade SSL Certificate Manager + HTTP to HTTPS + Health Checks + Session Affinity + Cloud DNS + AWS Route53

NC

No changes c1 to c6

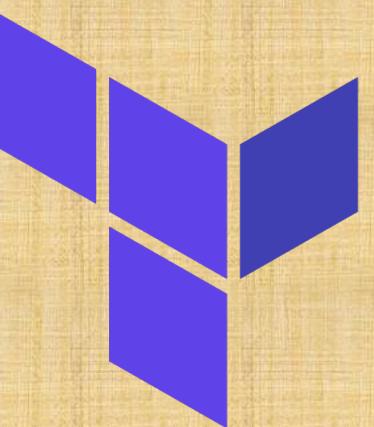
C7

Create production grade SSL certificates using Certificate Manager with DNS Authorization

C8

Create record sets in Cloud DNS

Demo-13

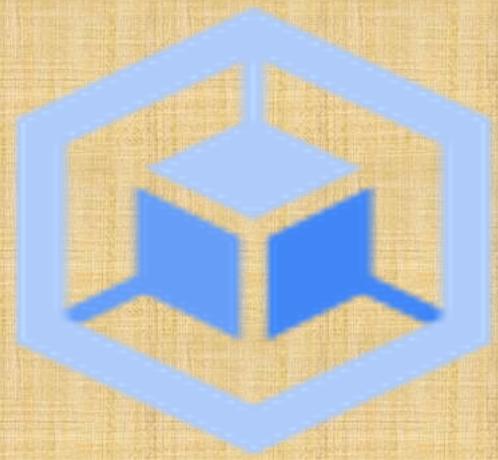


# Google Kubernetes Engine

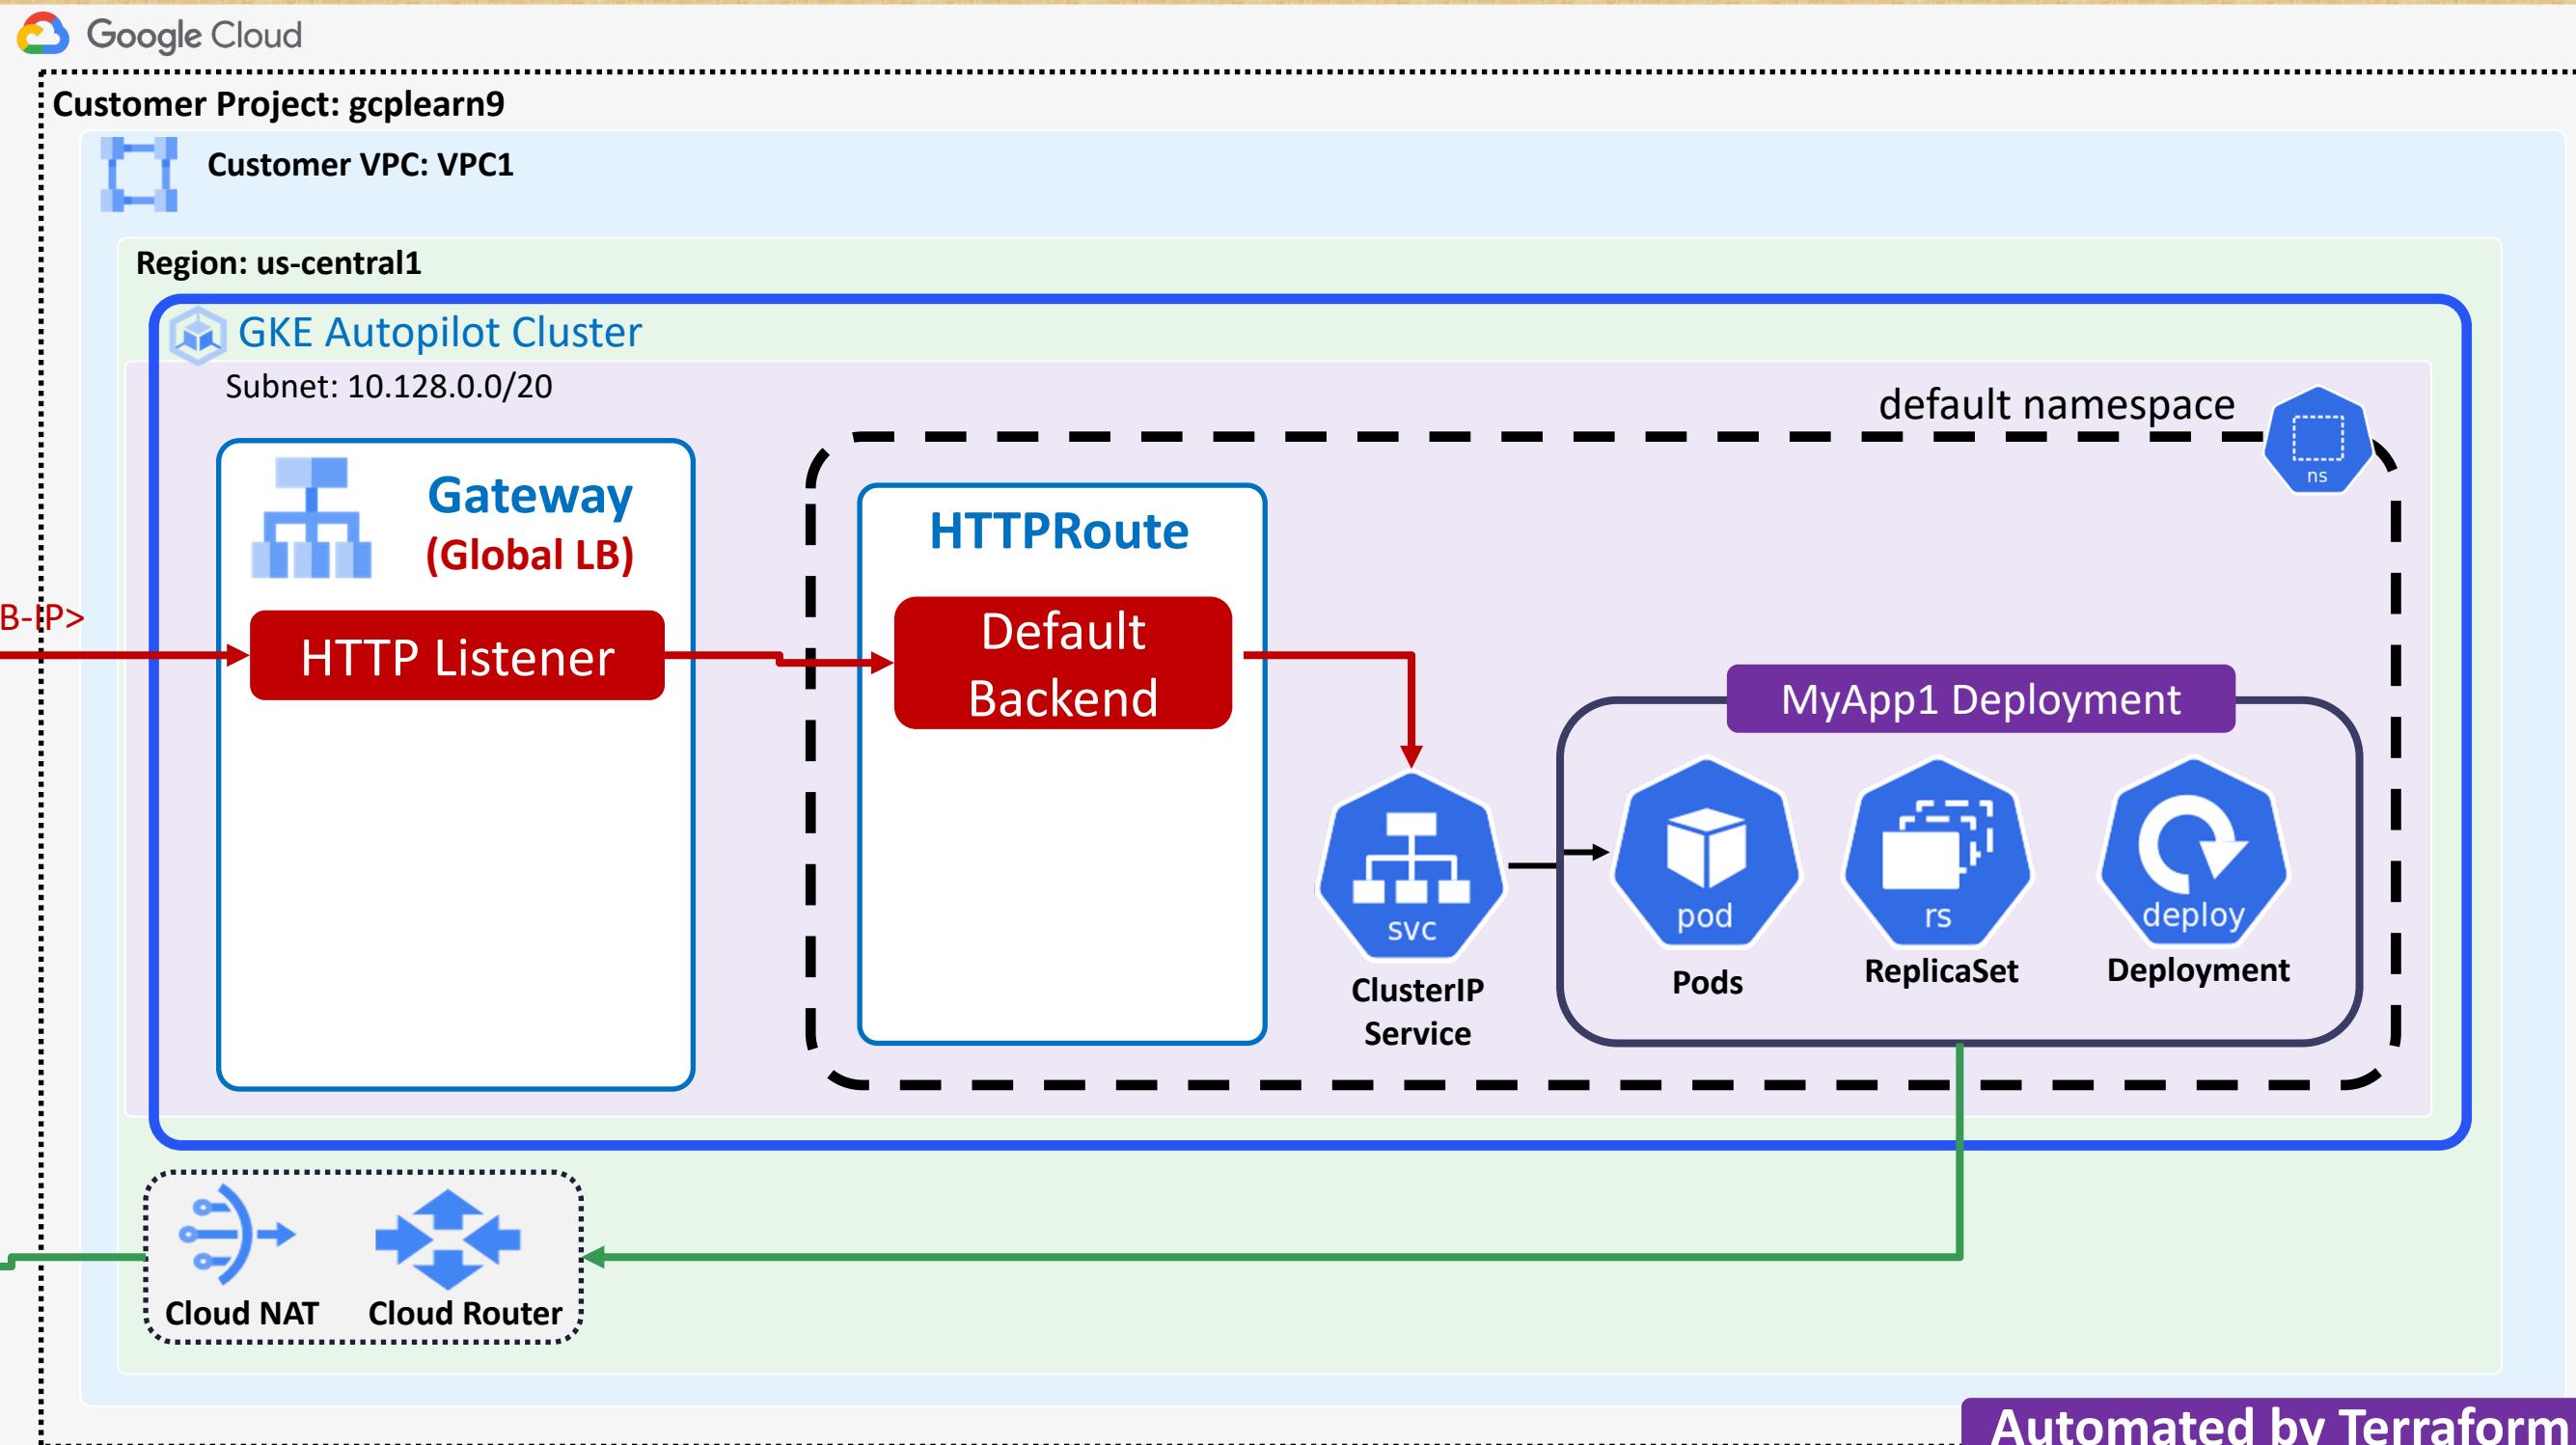
Gateway API

## Cloud Application Load Balancer

**Demo: GKE Autopilot cluster + Gateway (Global Application Load Balancer)**



# Gateway API – Global Application Load Balancer



- ✓ p2-global-k8sresources-yaml
- ! 01-myapp1-deployment.yaml
- ! 02-myapp1-clusterip-service.yaml
- ! 03-gateway.yaml
- ! 04-gateway-http-route.yaml

# What are we going to learn?

Demo: Project-2: Gateway (Global Application Load Balancer)

NC

NO Changes 01, 02

03

Gateway with gatewayClassName: gke-l7-global-  
external-managed

04

Gateway HTTPRoute

## Project-3: Terraform Manifests

- ✓ p3-global-k8sresources-terraform-manifests
  - ✗ c1-versions.tf
  - ✗ c2-01-variables.tf
  - ✗ c2-02-local-values.tf
  - ✗ c3-01-remote-state-datasource.tf
  - ✗ c3-02-providers.tf
  - ✗ c4-myapp1-deployment.tf
  - ✗ c5-myapp1-clusterip-service.tf
  - ✗ c6-gateway.tf
  - ✗ c7-gateway-http-route.tf
- ✗ terraform.tfvars

# What are we going to learn?

## Demo: Project-3: Gateway (Global Application Load Balancer)

NC

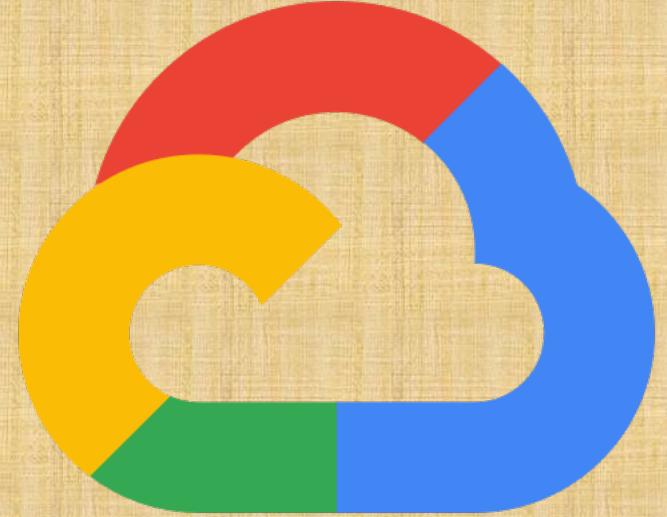
No changes c1 to c5

c6

Gateway with gatewayClassName: gke-l7-global-  
external-managed

c7

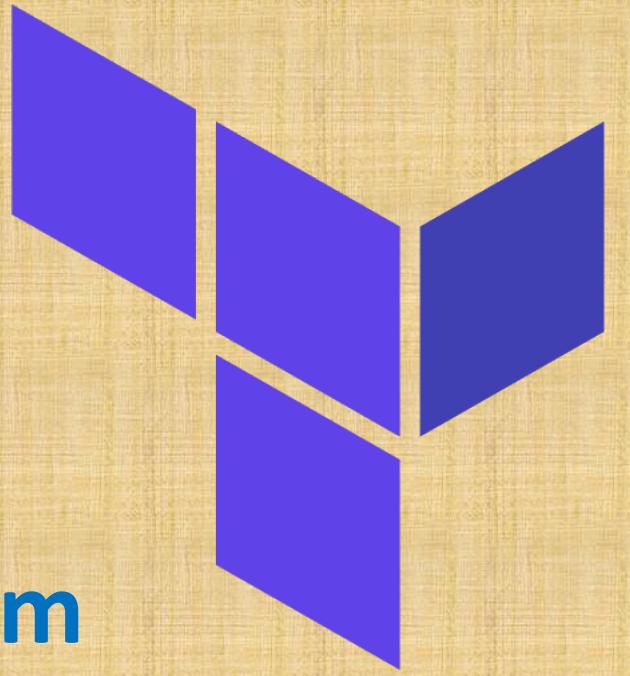
Gateway HTTPRoute



Demo-22

# Terraform Modules

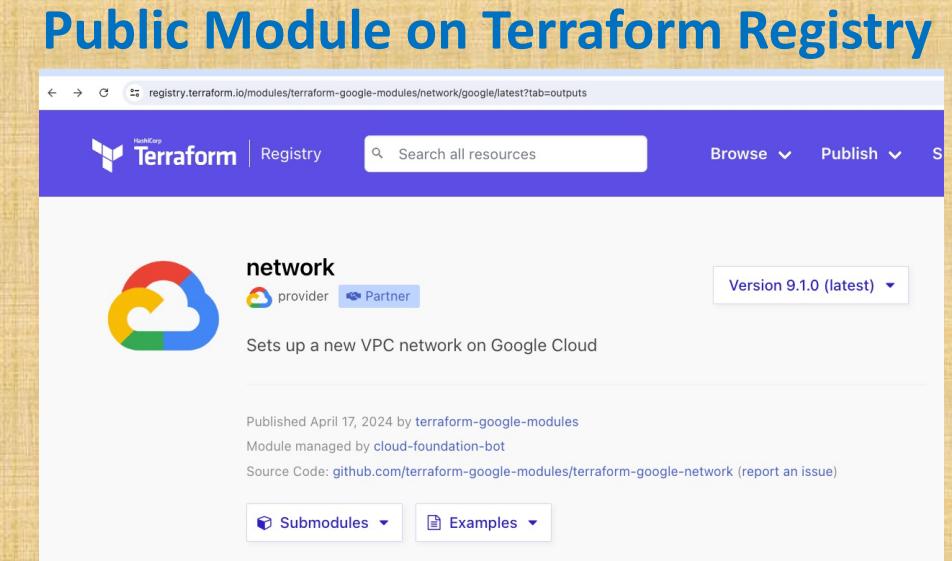
Use pre-built Modules from Terraform  
Registry



**Demo: Custom VPC (Uses Terraform Module from Terraform Public Registry) + Firewall Rules + VM Instance**

# Terraform Modules

- A module consists of a **collection of .tf files** kept together in a directory.
- Modules are the main way to **package and reuse** resource configurations with Terraform.
- We can **compare modules to functions** in our standard programming languages (**define once and reuse wherever needed**)
- **Example:**
  - Define a VPC Module **once, publish** it to a Terraform registry
  - **Call** the VPC module wherever needed (reuse it)



## Custom Module



# Terraform Modules

- **Root Module:** Every Terraform configuration has at least one module, known as its **root module**, which consists of the resources defined in the **.tf** files in the **main working directory**
- A Terraform module (usually the **root module** of a configuration) can call **other modules** to include their resources into the configuration.
- **Child Module:** A module that has been **called by another module** is often referred to as a child module.
- Child modules can be called **multiple times** within the **same** configuration, and **multiple configurations** can use the **same** child module.

## Root Module

```
└ terraform-manifests  
  └ c1-versions.tf  
  └ c2-variables.tf  
  └ c3-locals.tf  
  └ c4-vpc.tf  
  └ c5-firewalls.tf  
  └ c6-vminstance.tf  
  └ c7-outputs.tf  
  └ terraform.tfvars
```

Call VM Instance child module in Root Module

## Child Module

```
└ 29-Terraform-Build-Custom-Module  
  └ modules/vminstance  
    └ app1-webserver-install.sh  
    └ main.tf  
    └ outputs.tf  
    └ variables.tf  
    └ versions.tf
```

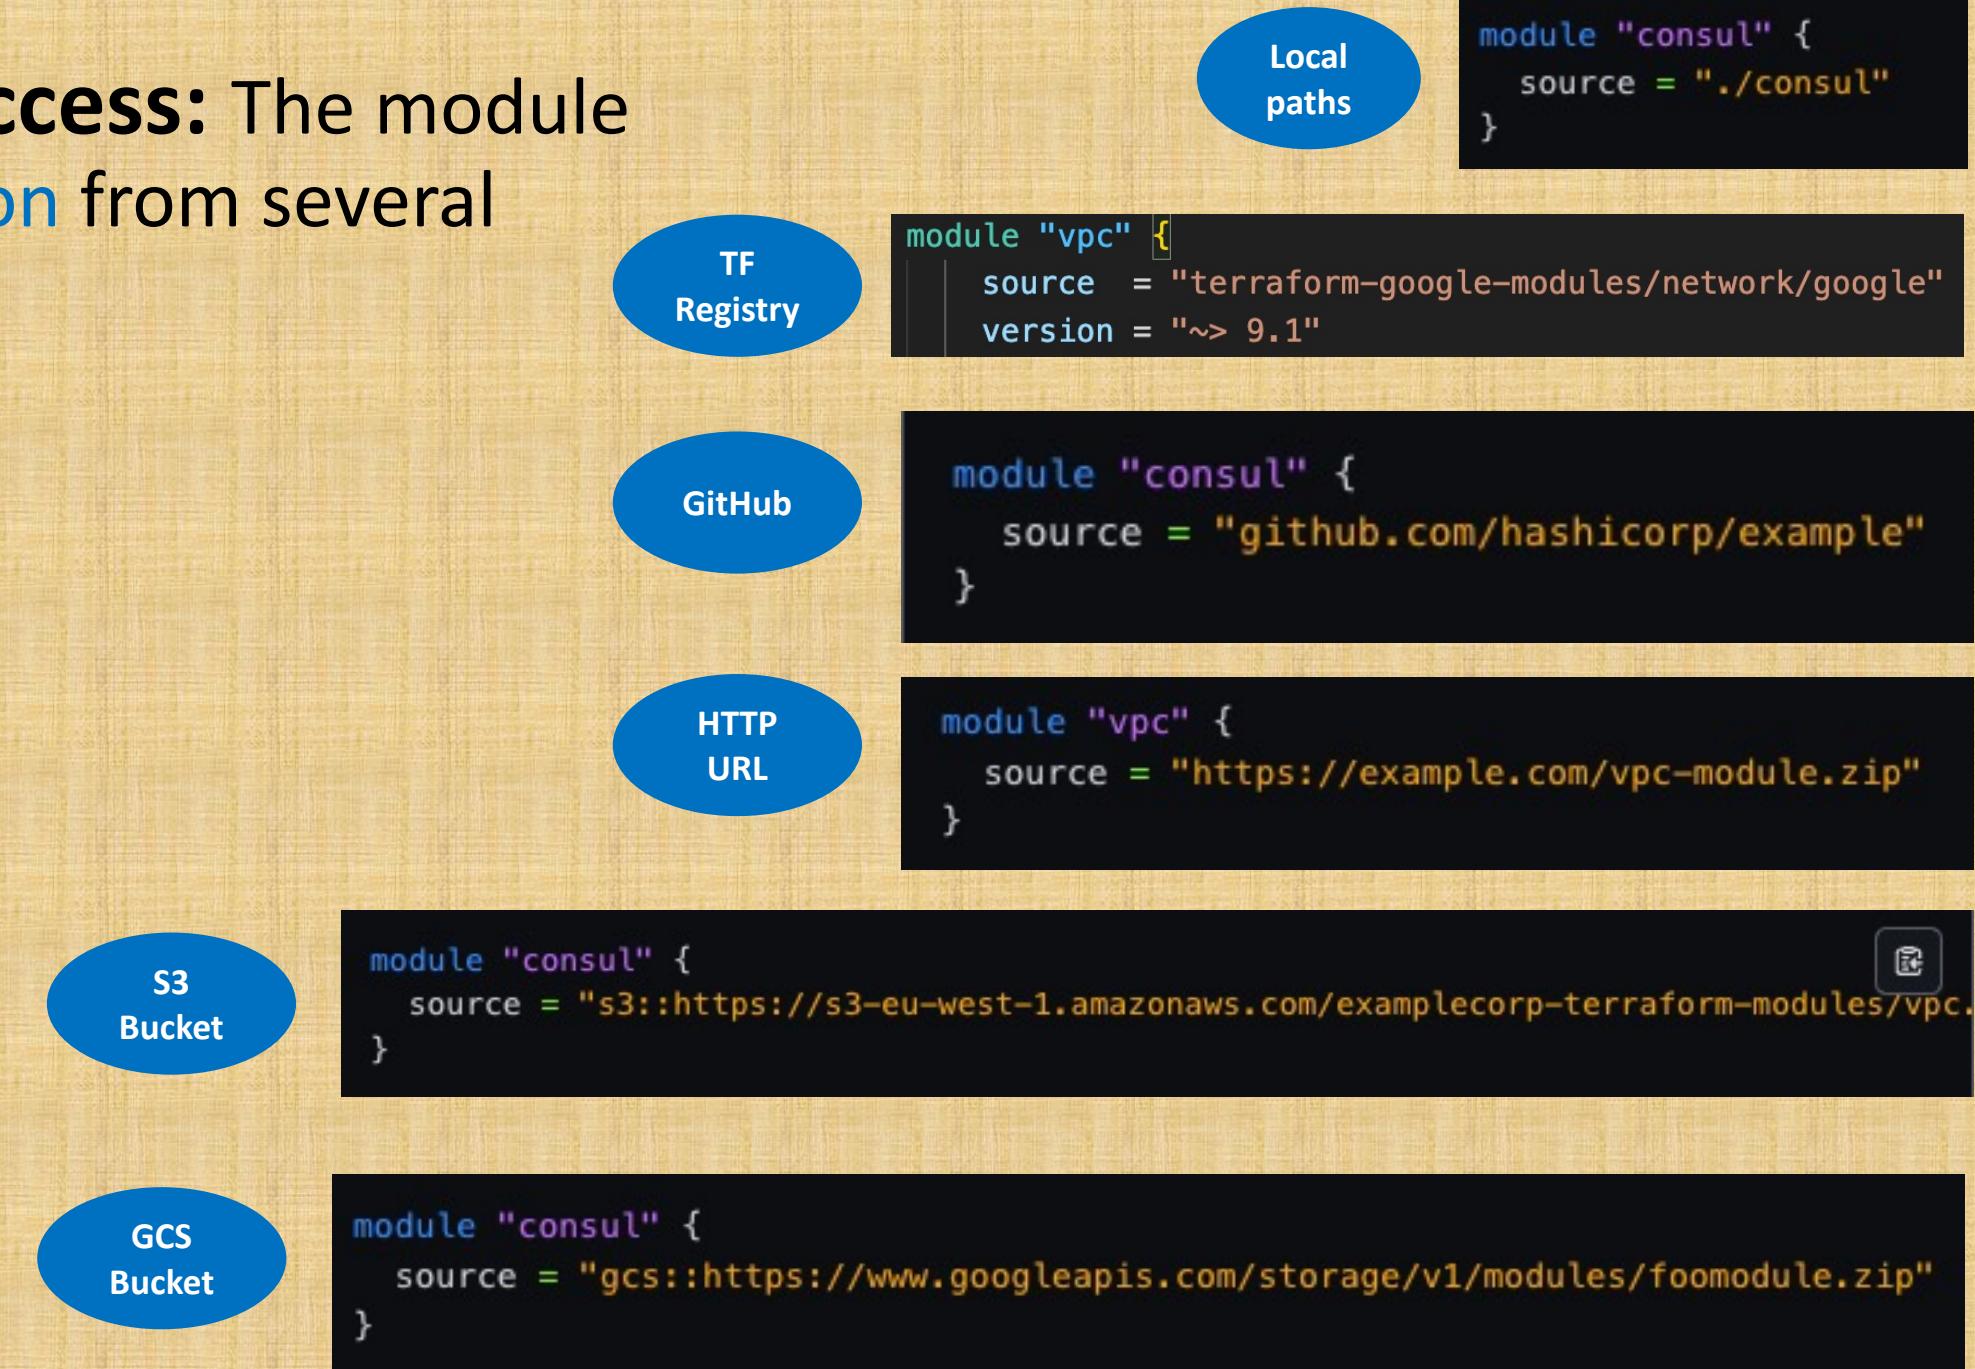
# Terraform Modules

- **Module Types**
  - **Local Modules (Local Paths)**
    - Accessible using **folder path** locally
    - Accessible locally **in that desktop** or for others via shared drives
  - **Public Modules**
    - Modules published to **Terraform Registry**
    - Publicly accessible to **anyone** (**Open source**)
    - **Free to use, use with caution**, will contain continuous upgrades and might **break our current code**
    - If used, need to follow **strong version constraints** or **fixed module versions** to avoid major outages during upgrades
  - **Private Modules**
    - Modules developed and published to **private registry** in their respective organizations
    - Terraform Registry also supports **private registry concept**
    - Accessible only to that **respective organization users** (**Highly secure, internal to that organization**)

# Terraform Modules

- **Module Storage and Access:** The module installer supports **installation from several different source types.**

- Local Paths
- Terraform Registry
- GitHub
- HTTP URLs
- S3 Buckets
- GCS Buckets
- and many more



# Terraform Registry - Use Publicly Available Modules

The Terraform Registry hosts a broad collection of publicly available Terraform modules for configuring many kinds of common infrastructure.

## Modules Demo 1

These modules are free to use, and Terraform can download them automatically if you specify the appropriate source and version in a module call block.

```
# Module: VPC
module "vpc" {
    source  = "terraform-google-modules/network/google"
    version = "~> 9.1"

    project_id    = var.gcp_project
    network_name = "${local.name}-vpc"
    routing_mode = "GLOBAL"
    subnets = [
        {
            subnet_name      = "${local.name}-${var.gcp_region1}-subnet"
            subnet_ip       = "10.128.0.0/20"
            subnet_region   = var.gcp_region1
        }
    ]
}
```

# What are we going to learn?

Demo: Custom VPC (Uses Terraform Module from Terraform Public Registry)  
+ Firewall Rules + VM Instance

## ✓ 02-terraform-manifests-with-modules

\$ app1-webserver-install.sh

└ c1-versions.tf

└ c2-variables.tf

└ c3-locals.tf

└ c4-vpc.tf

└ c5-firewalls.tf

└ c6-vminstance.tf

└ c7-outputs.tf

└ terraform.tfvars

NC

C4

C5

C6

C7

No Changes from C1, C2, C3

Create VPC Module using Terraform public registry

Update firewall rules with VPC ID from VPC Module

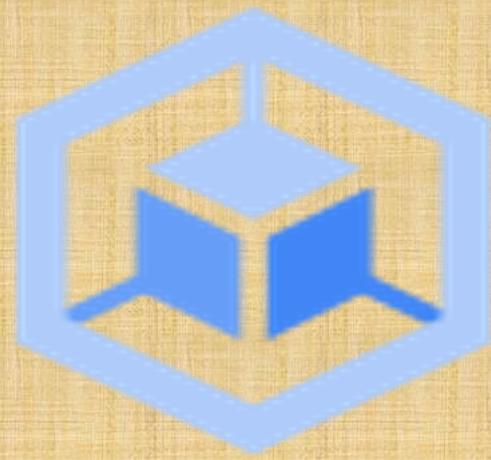
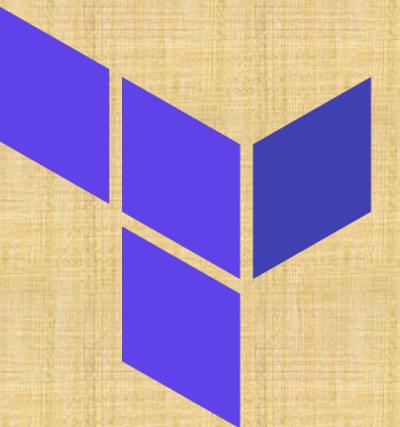
Update Subnet with Subnet ID from VPC Module

Update VPC ID and Subnet ID from VPC Module

Demo-23

# Google Kubernetes Engine Terraform Custom Module

**Demo: Create GKE Autopilot Cluster by building a custom  
Terraform Module**



# Build a Local Terraform Module

## Modules Demo 2

Build a **Local** Terraform Module  
and call it from **Root Module** and  
Test

Call GKE Cluster child  
module in Root Module

**Root Module**

- ✓ 23-GKE-Infra-Custom-Terraform-Modules
  - > modules
  - ✓ p1-gke-autopilot-cluster-private
    - ✓ c1-versions.tf
    - ✓ c2-01-variables.tf
    - ✓ c2-02-local-values.tf
    - ✓ c3-vpc.tf
    - ✓ c4-01-gke-cluster.tf
    - ✓ c4-02-gke-outputs.tf
    - ✓ c5-Cloud-NAT-Cloud-Router.tf
    - ✓ terraform.tfvars

## Child Module

- ✓ 23-GKE-Infra-Custom-Terraform-Modules
  - ✓ modules/gke\_cluster
    - ✓ main.tf
    - ✓ outputs.tf
    - ✓ variables.tf
    - ✓ versions.tf

```
✓ modules/gke_cluster
  ✓ main.tf
  ✓ outputs.tf
  ✓ variables.tf
  ✓ versions.tf
```

# What are we going to learn?

## Demo: Create GKE Cluster Custom Module

versions

variables

main

outputs

Define Terraform Provider versions

Define Terraform Input Variables

Create GKE Cluster Terraform resource

Define Terraform Outputs

- ✓ p1-gke-autopilot-cluster-private
- ✗ c1-versions.tf
- ✗ c2-01-variables.tf
- ✗ c2-02-local-values.tf
- ✗ c3-vpc.tf
- ✗ c4-01-gke-cluster.tf
- ✗ c4-02-gke-outputs.tf
- ✗ c5-Cloud-NAT-Cloud-Router.tf
- ✗ terraform.tfvars

# What are we going to learn?

Demo: Project-1: Create GKE Autopilot Cluster by calling a custom Terraform Module

NC

No changes c1 to c3, c5

c4-01

Create GKE cluster by calling **custom gke\_cluster** Terraform module

c4-02

Define GKE Terraform outputs

Demo-24

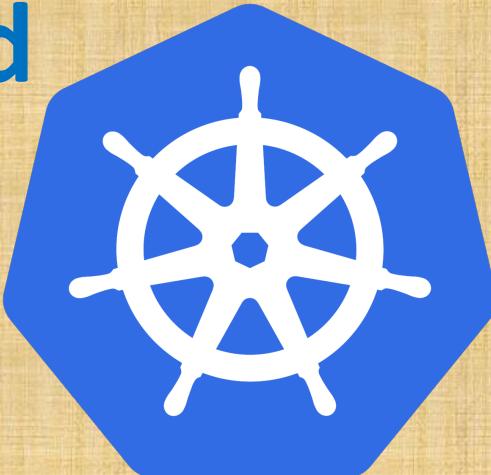
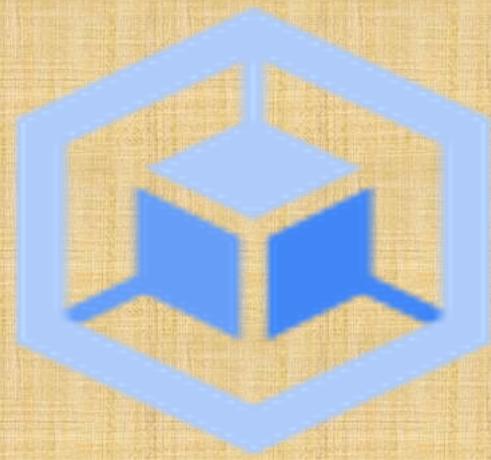
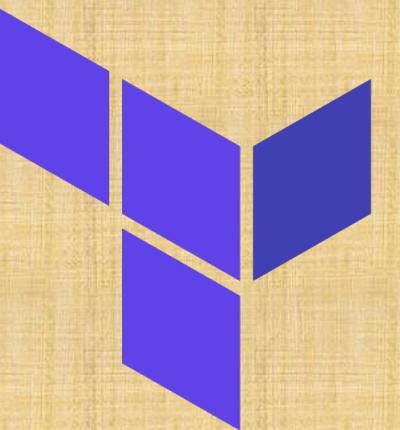
# Google Kubernetes Engine

## GKE Infra DevOps

## GCP Cloud Build, GitHub GCP Cloud Build App and GitHub

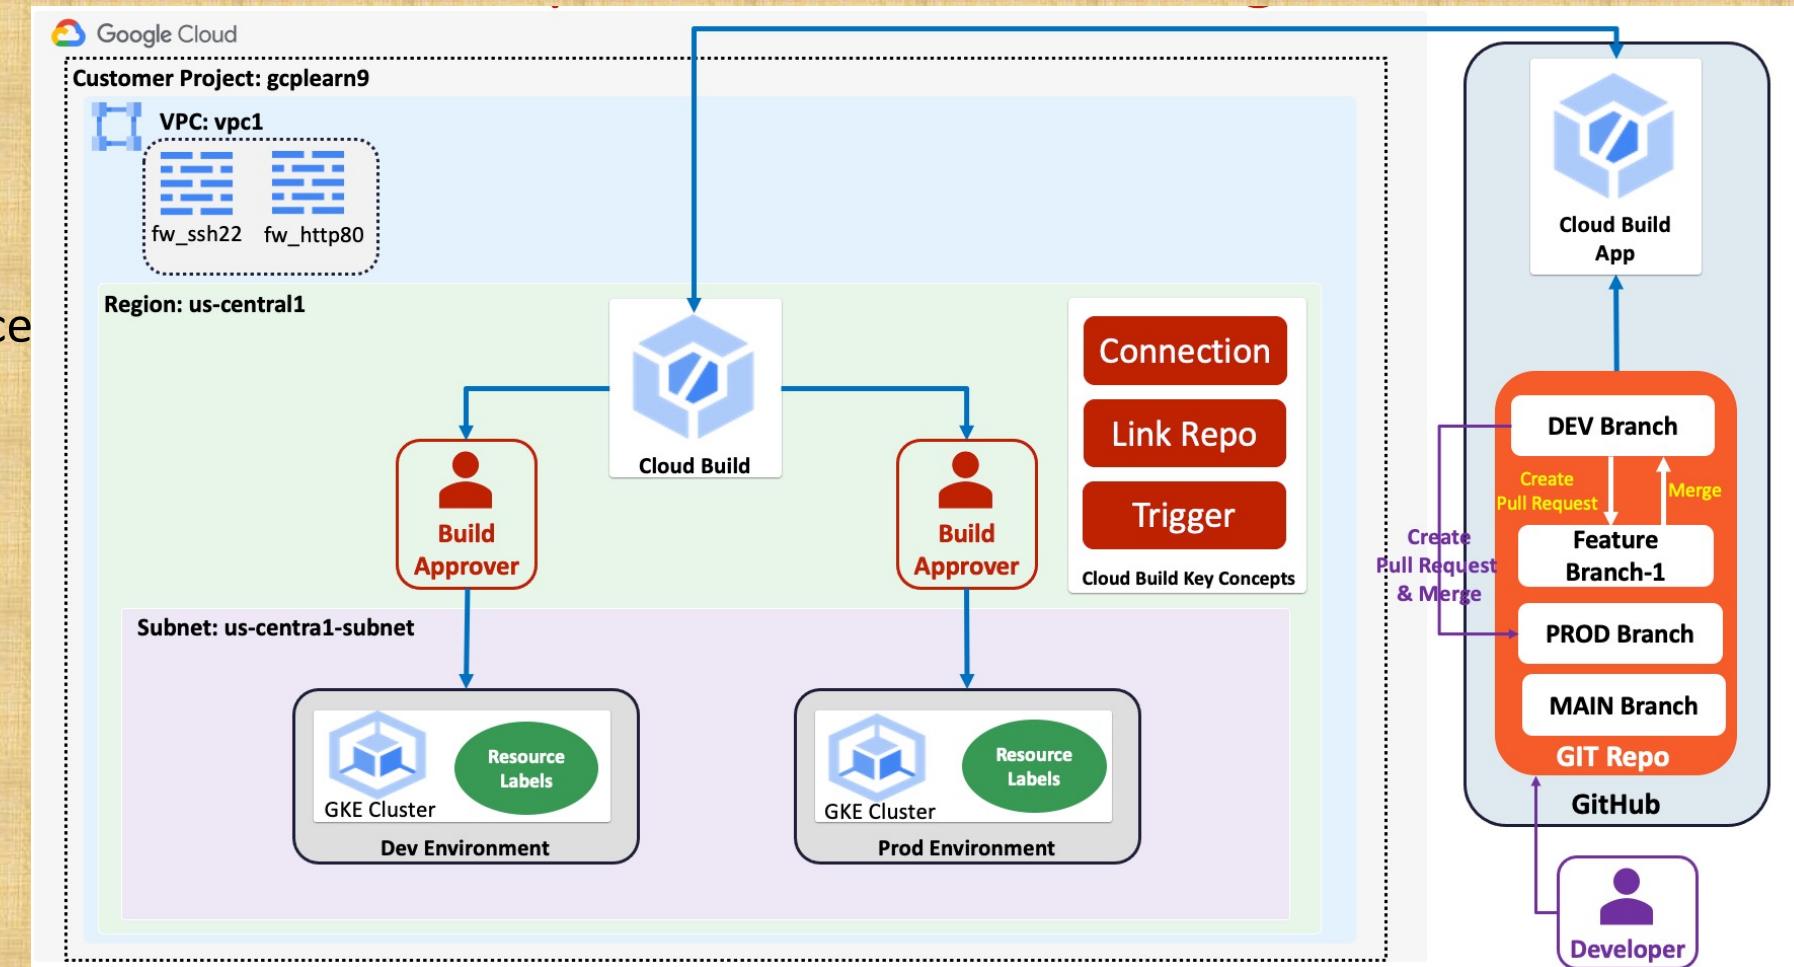
Demo: Custom VPC + + GKE Cluster (Custom or Local Terraform Module)

Implement DevOps Pipelines for Terraform Configs on GCP GKE

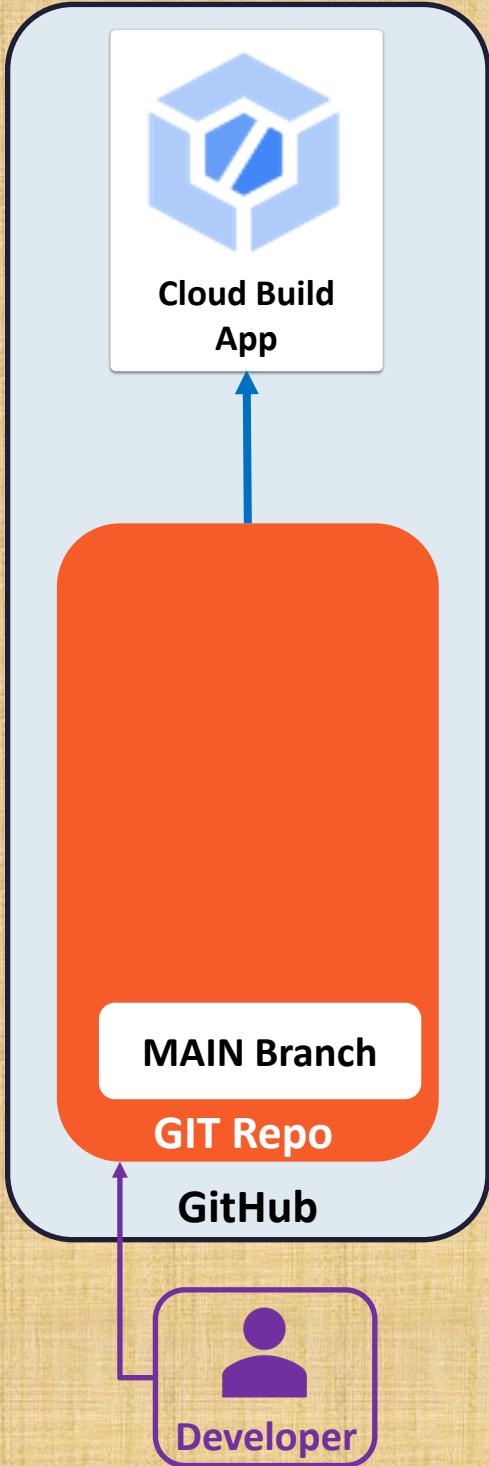


# GCP GKE DevOps for Terraform Configs

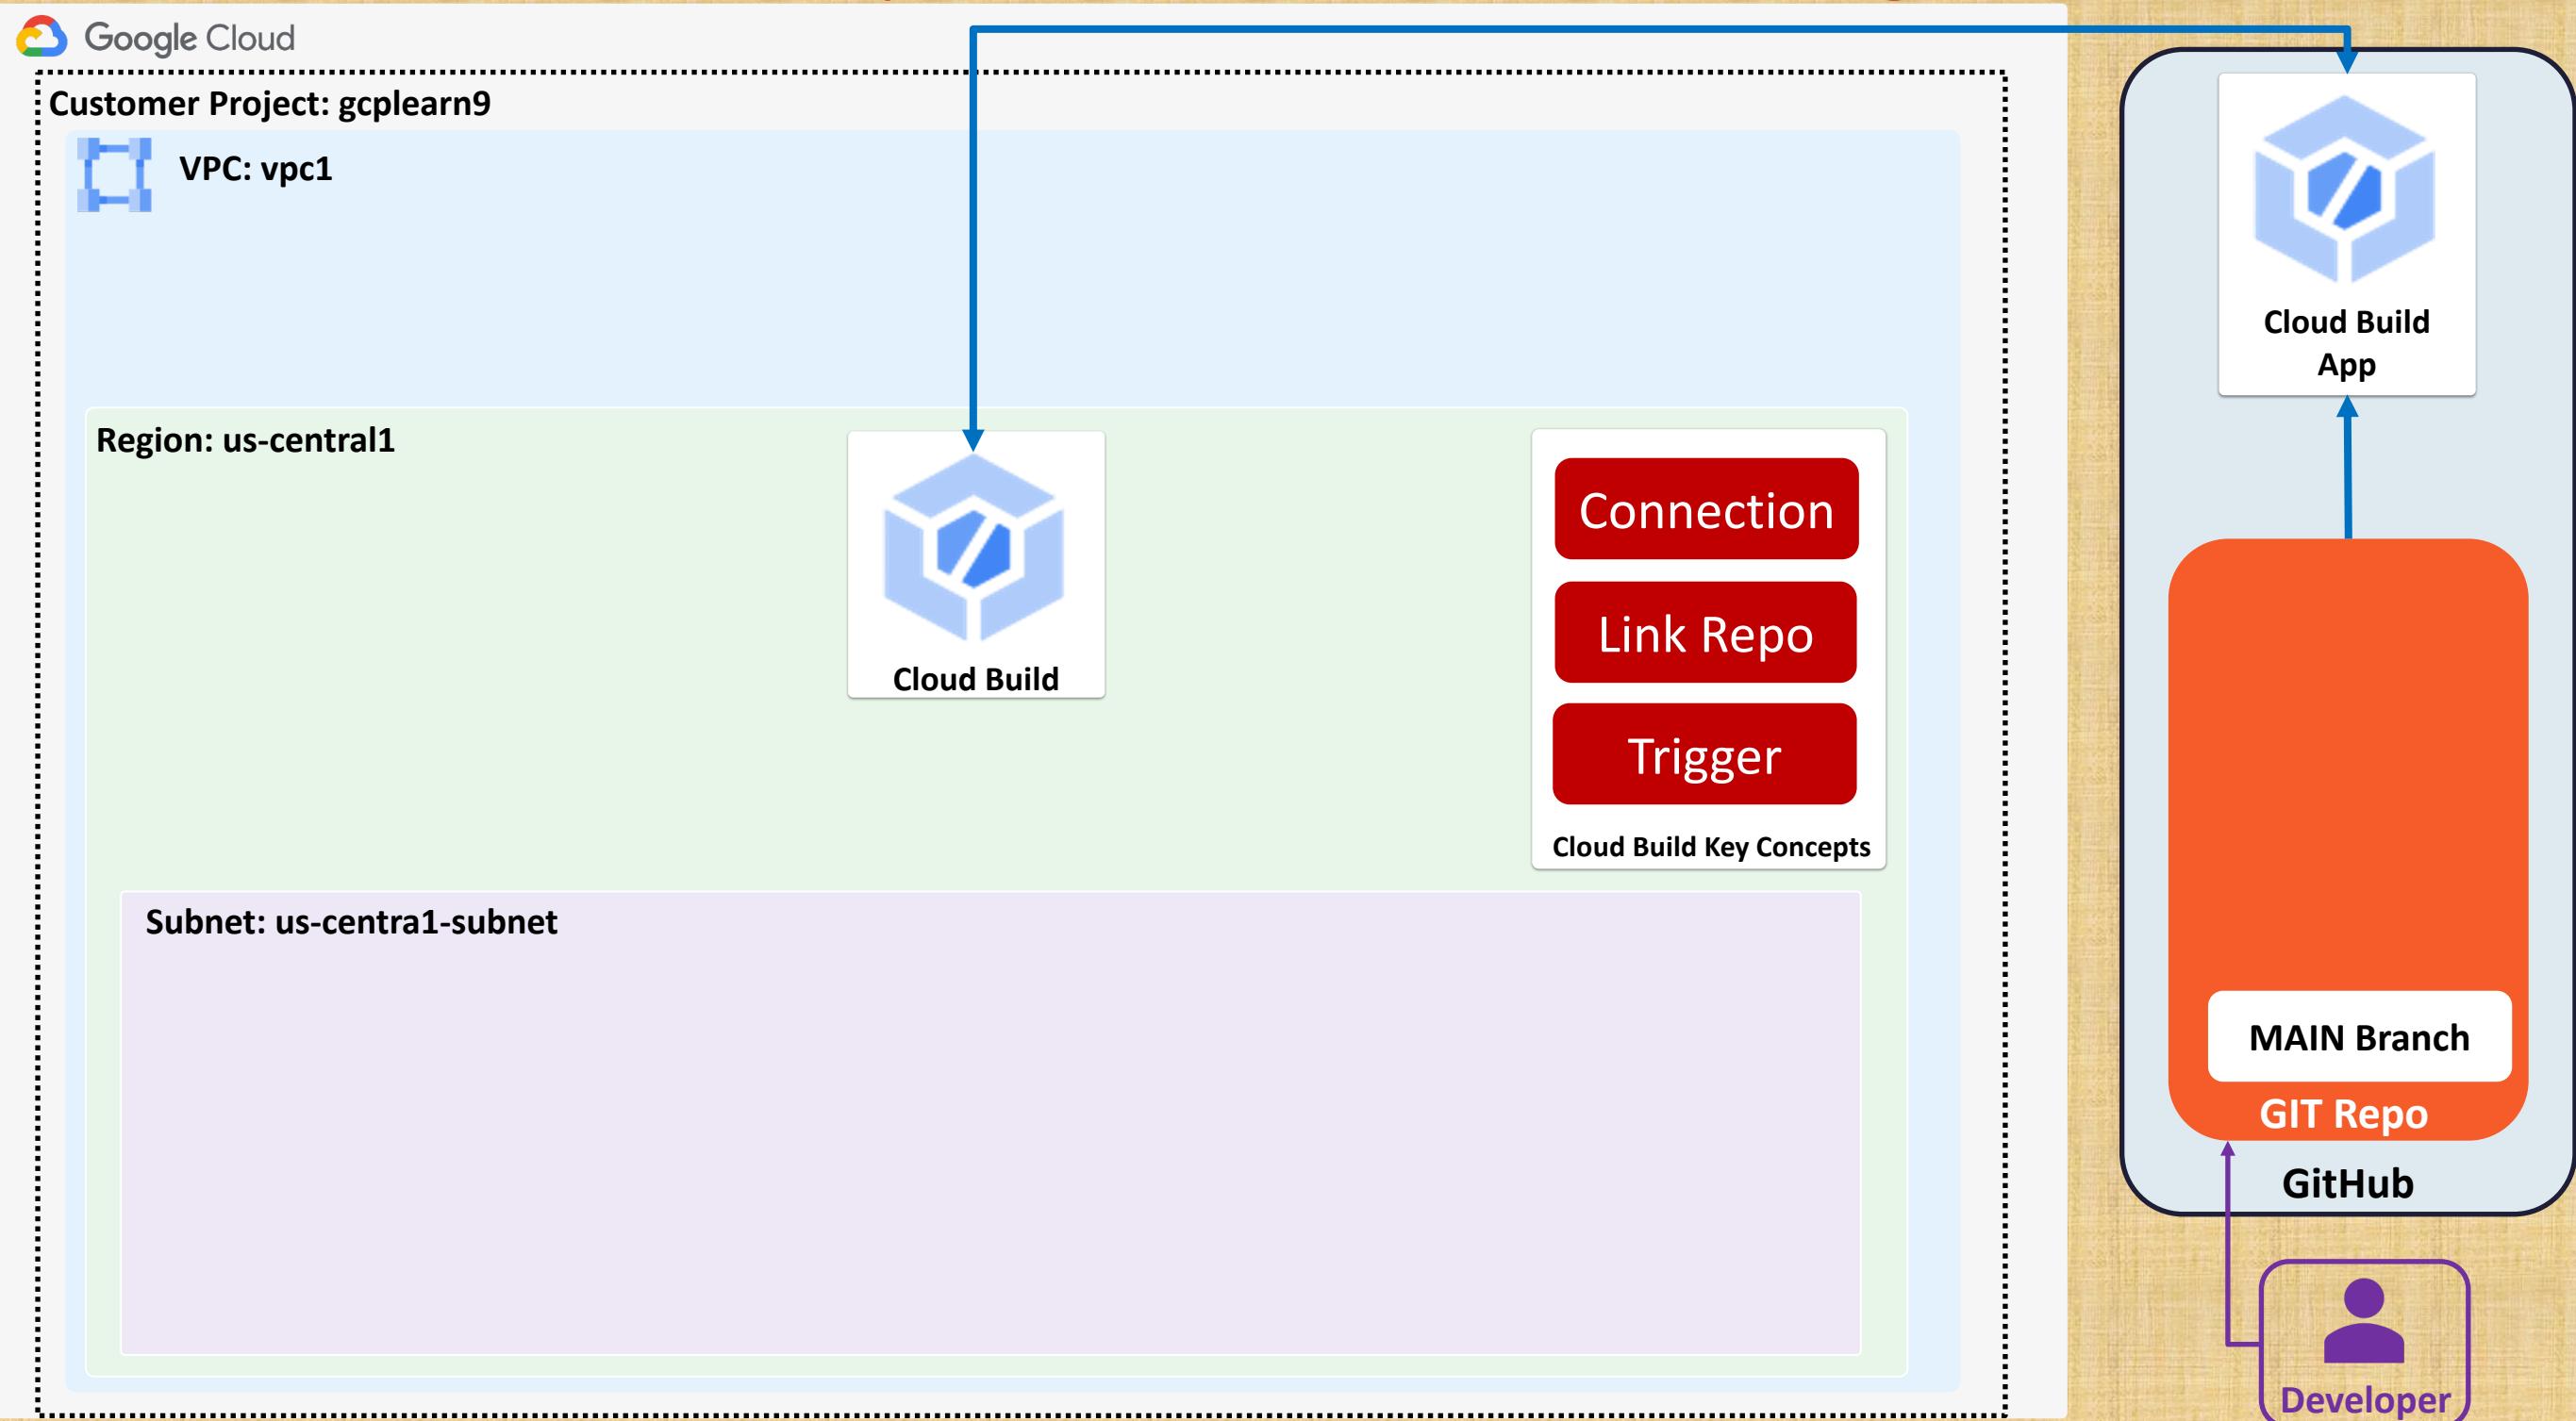
- The following resources are used for this implementation
- GitHub Repositories
  - GitHub Branches
    - Main branch
    - Dev branch
    - Prod branch
    - Feature branch
  - GitHub Pull Request and Merge concepts
  - GCP Cloud Build App from GitHub Marketplace
- GCP Cloud Build
  - Connections
  - Link Repositories
  - Triggers
- GCP Resources
  - GCP VPC
  - GCP Firewall Rules
  - GCP GKE Autopilot cluster
- Terraform Concepts
  - Terraform Resources
  - Terraform custom modules
  - Terraform Remote Backend using Google Cloud Storage Bucket



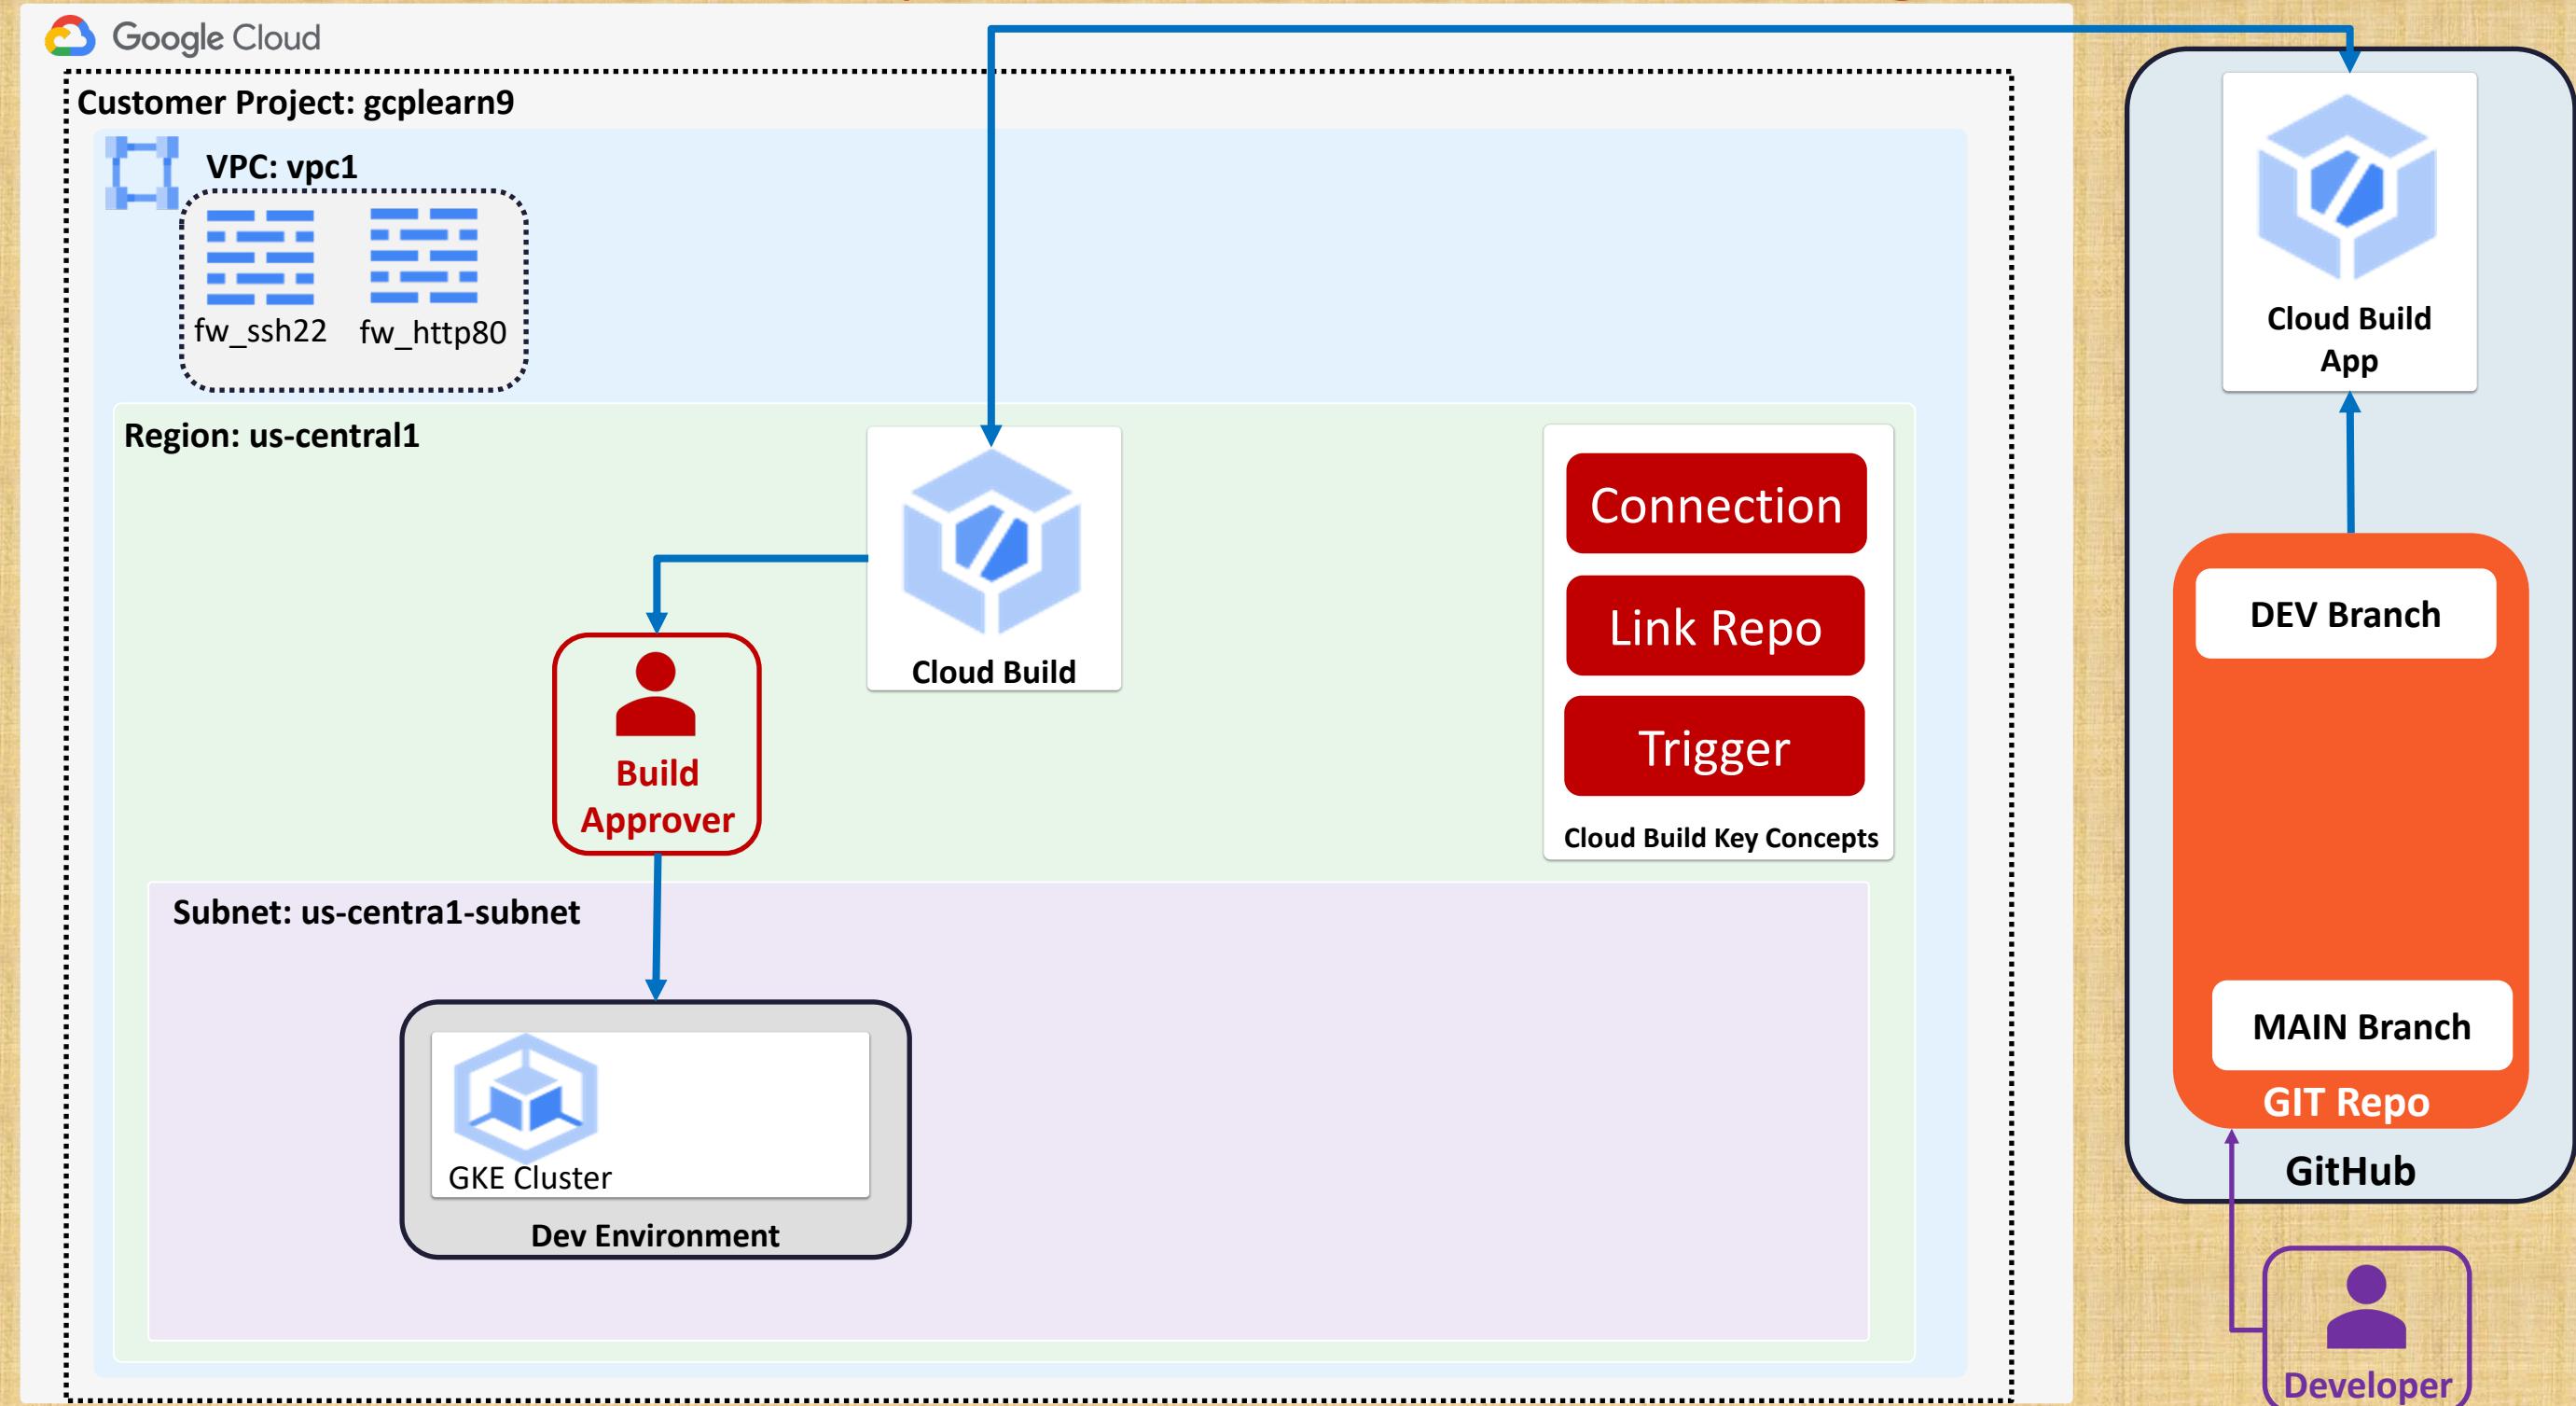
# GCP GKE DevOps for Terraform Configs



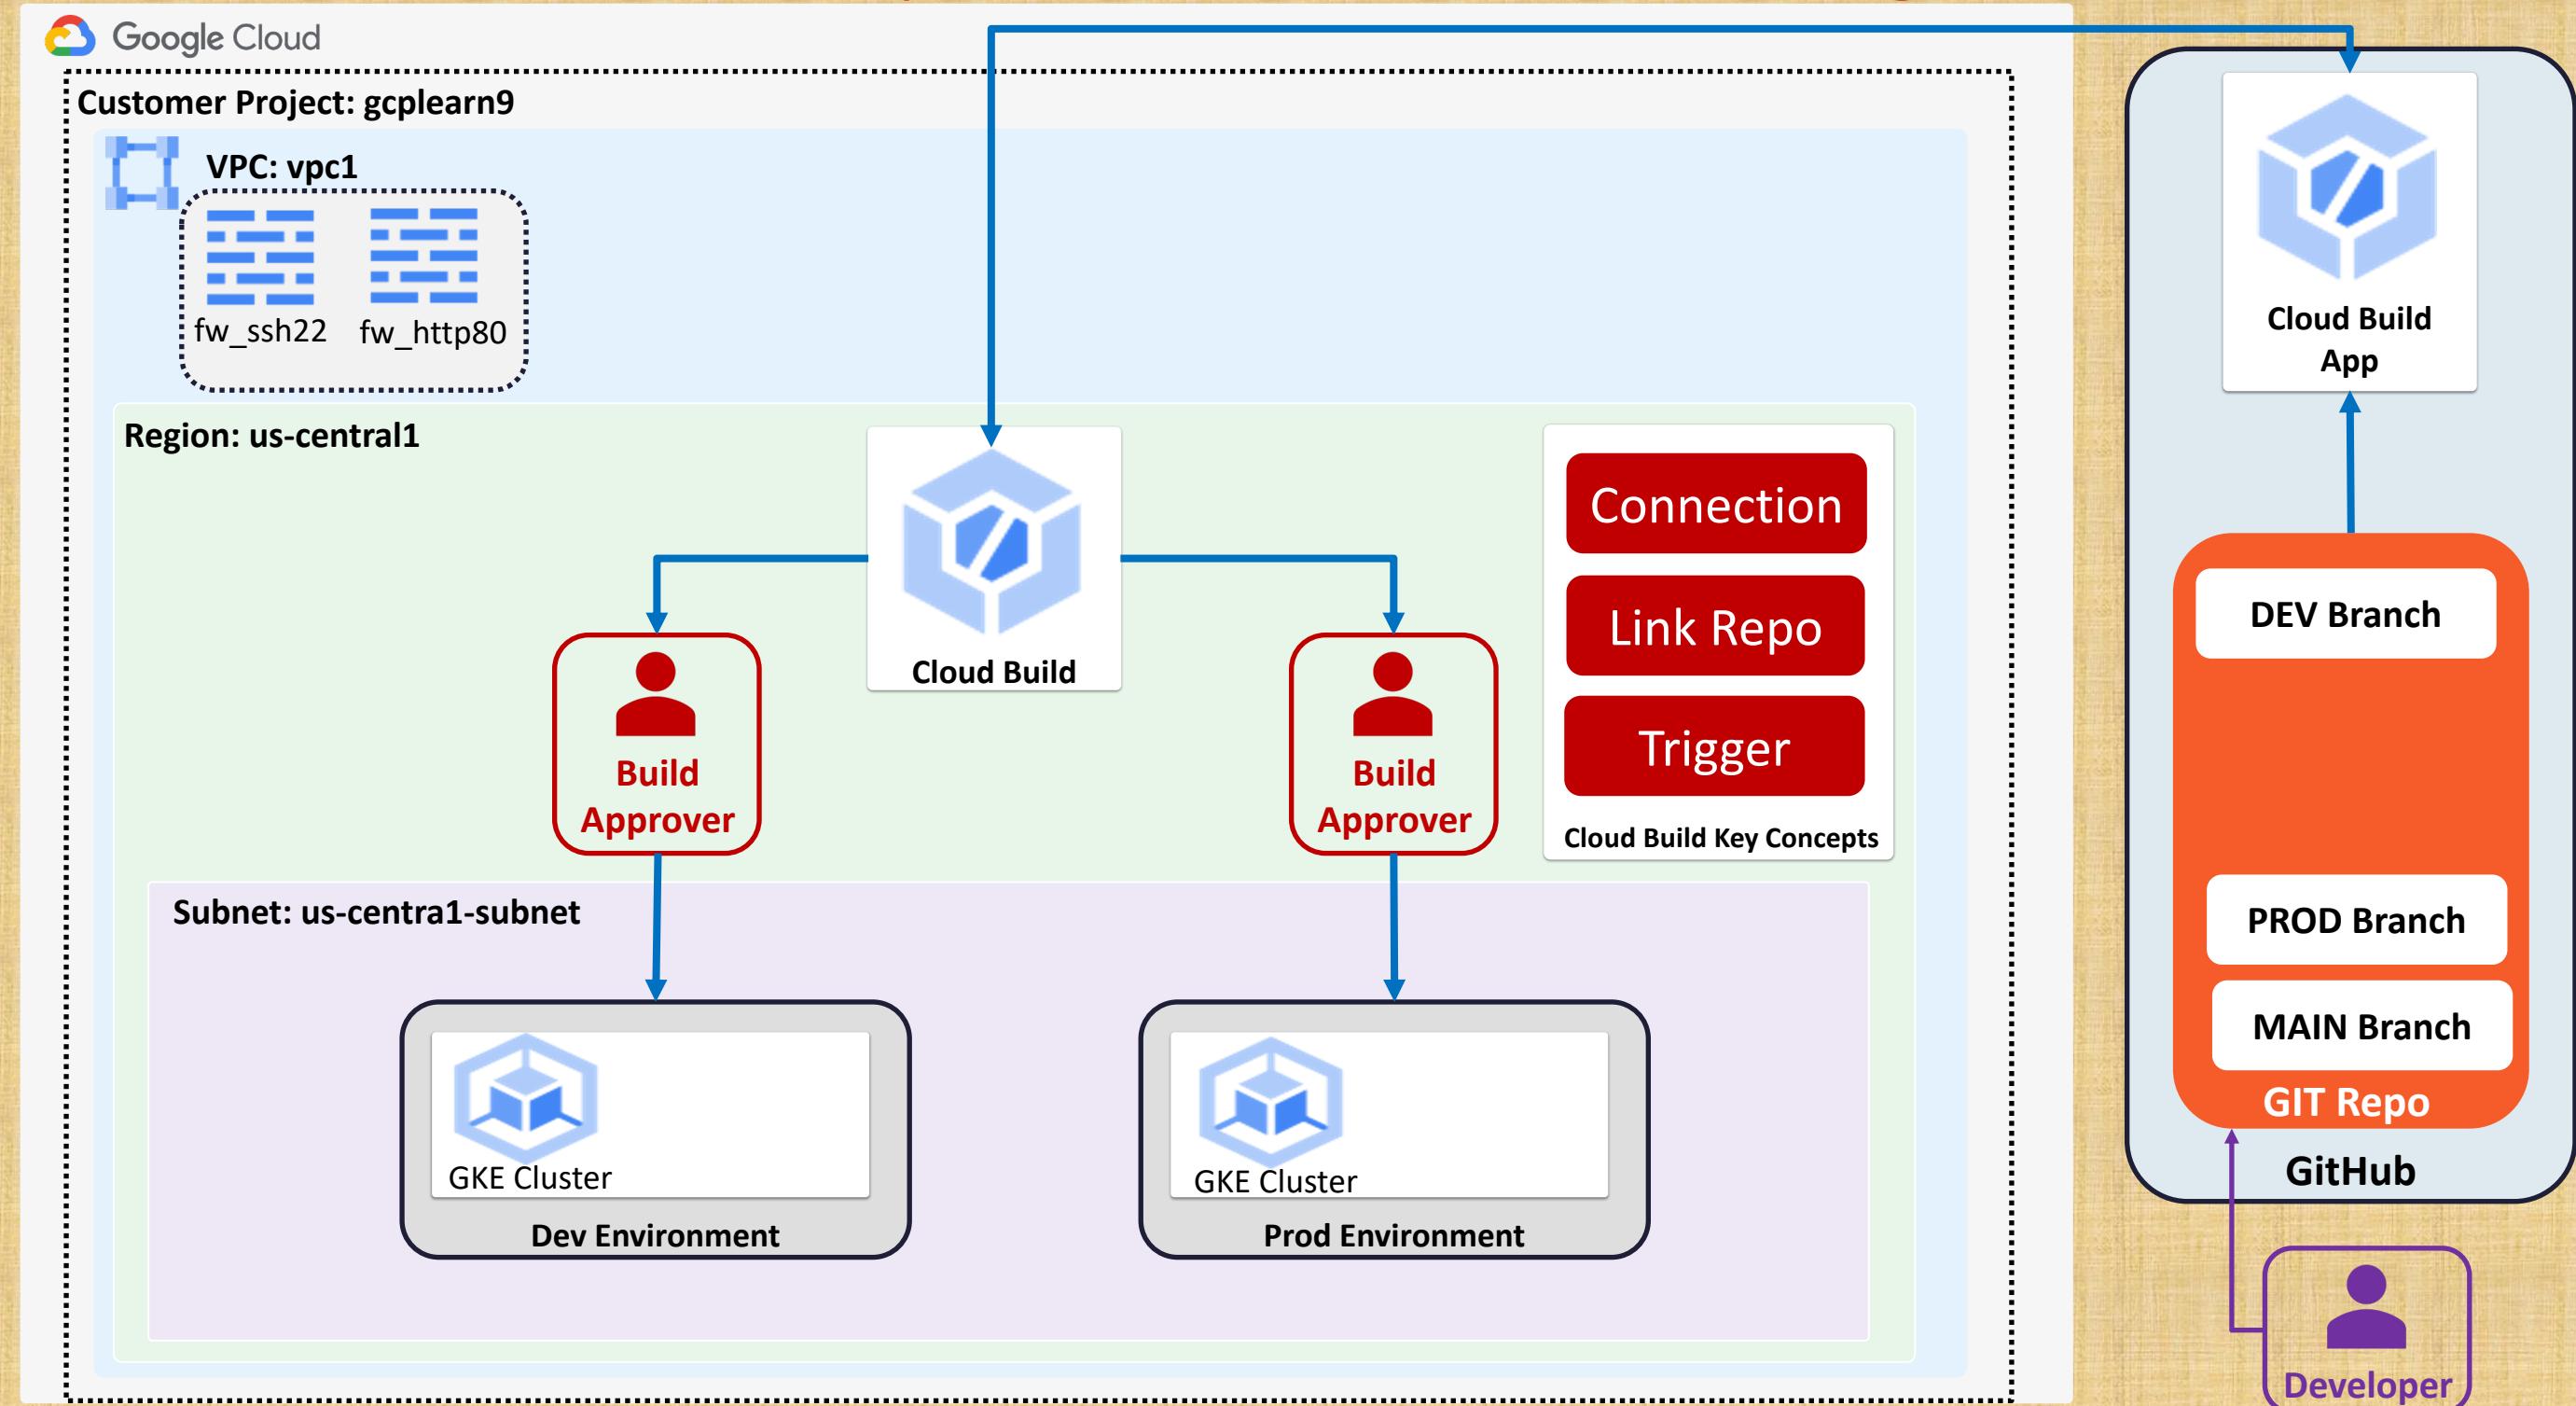
# GCP GKE DevOps for Terraform Configs



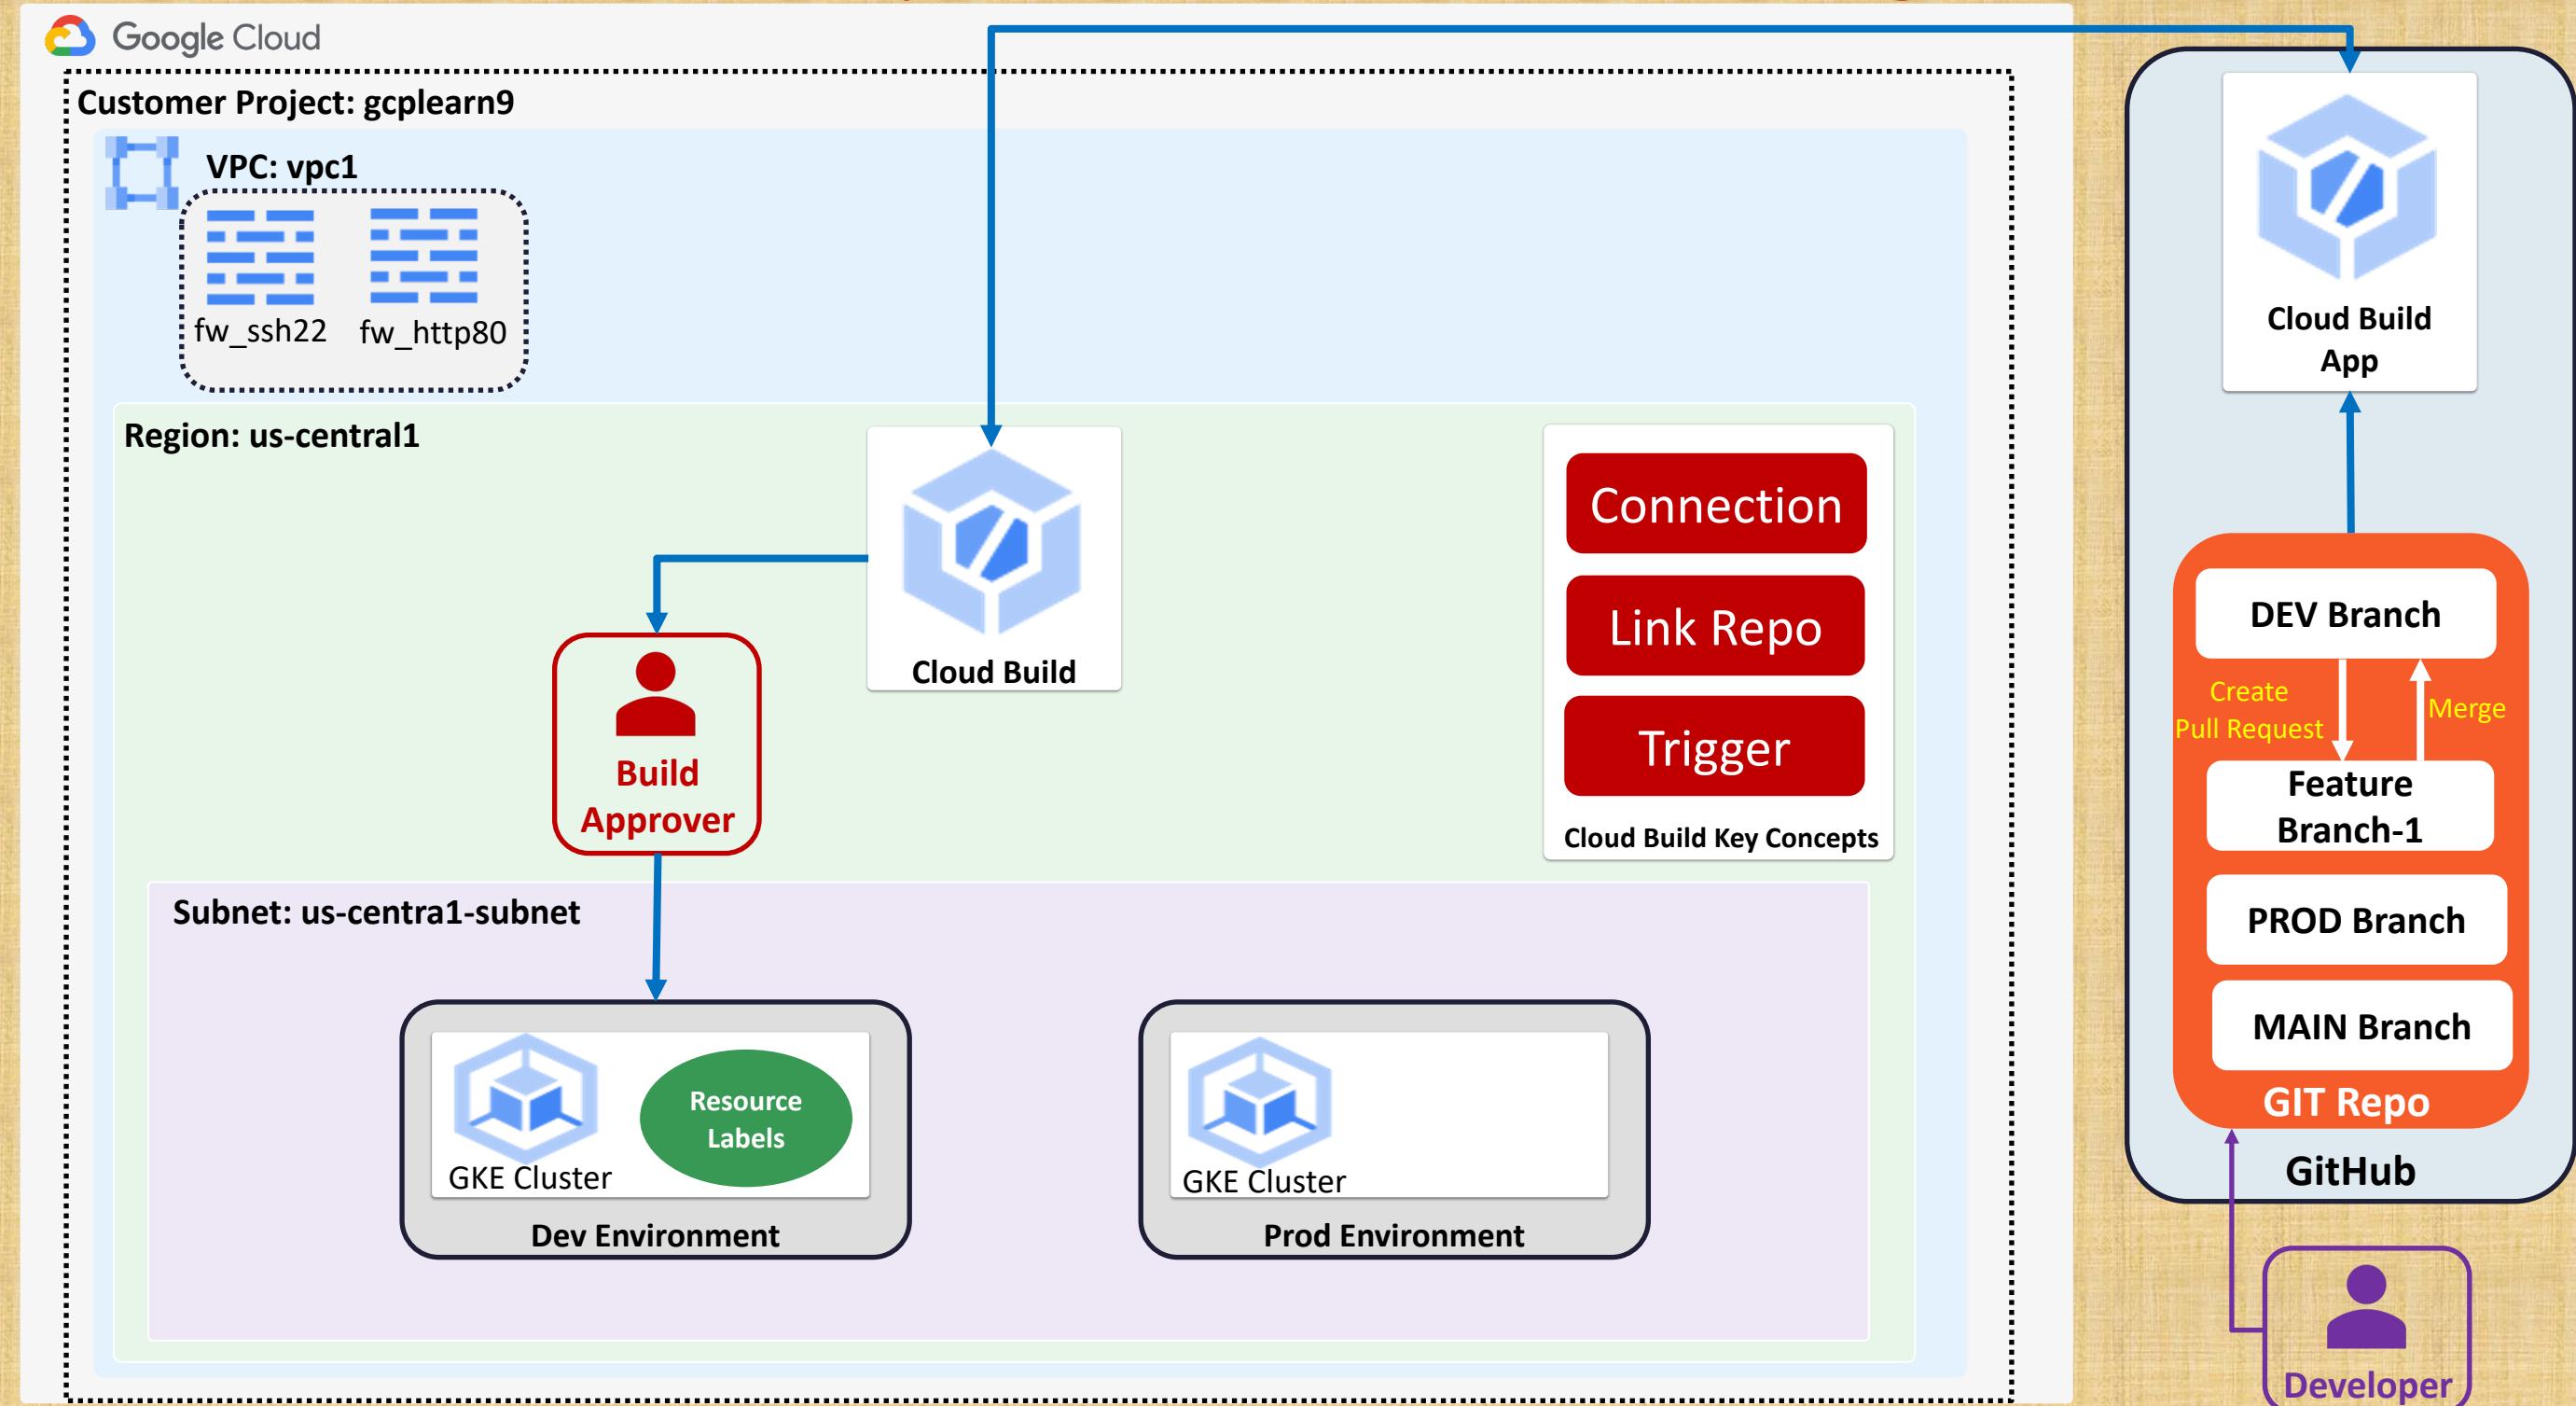
# GCP GKE DevOps for Terraform Configs



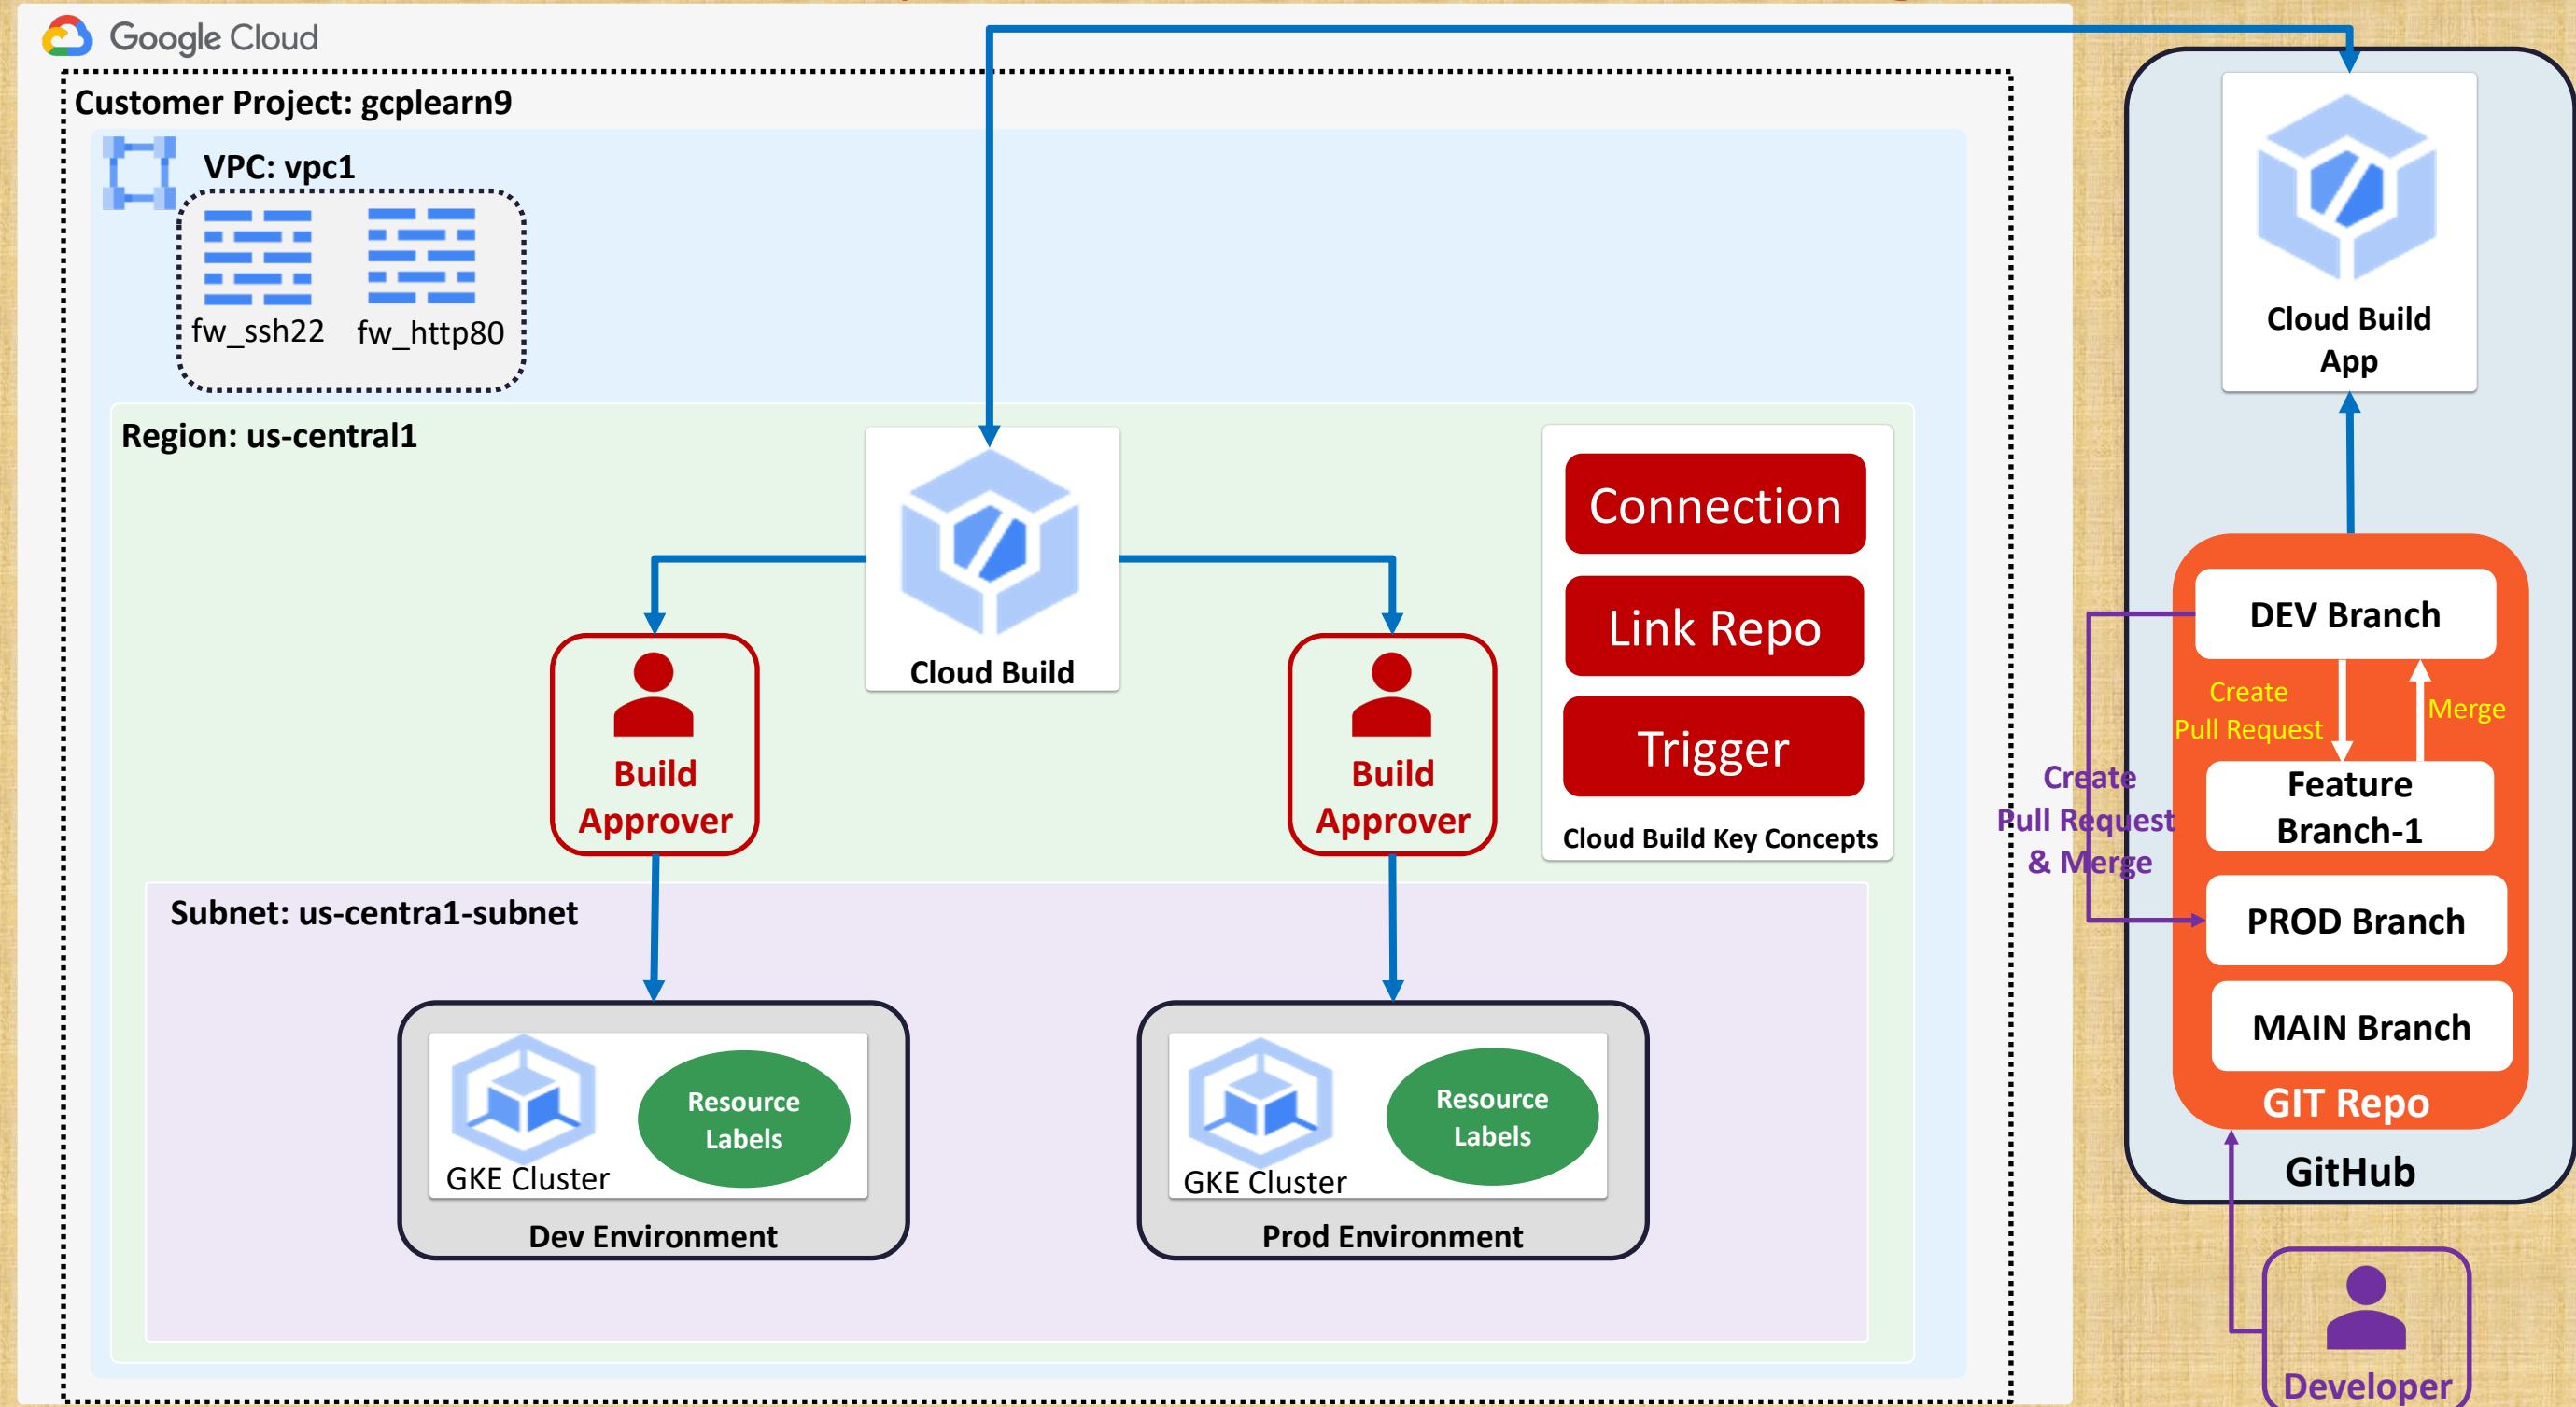
# GCP GKE DevOps for Terraform Configs



# GCP GKE DevOps for Terraform Configs



# GCP GKE DevOps for Terraform Configs



✓ 24-GKE-Infra-DevOps-CloudBuild-GitHub

✓ Git-Repo-files

✓ environments

✓ dev

- └ c1-versions.tf
- └ c2-01-variables.tf
- └ c2-02-local-values.tf
- └ c3-vpc.tf
- └ c4-01-gke-cluster.tf
- └ c4-02-gke-outputs.tf
- └ c5-Cloud-NAT-Cloud-Router.tf
- └ terraform.tfvars

✓ prod

- └ c1-versions.tf
- └ c2-01-variables.tf
- └ c2-02-local-values.tf
- └ c3-vpc.tf
- └ c4-01-gke-cluster.tf
- └ c4-02-gke-outputs.tf
- └ c5-Cloud-NAT-Cloud-Router.tf
- └ terraform.tfvars

# What are we going to learn?

Demo: Custom VPC + GKE Cluster(Custom or Local Terraform Module)  
Implement DevOps Pipelines for Terraform Configs on GCP GKE

All these TF Configs are copy from Root Module Demo-23 (p1-gke-autopilot-cluster-private folder)

DEV

C1

TFVA  
RS

PROD

C1

TFVA  
RS

## Development Environment

Update GCS Bucket (your bucket name) and State file prefix as “dev/gke-cluster”

Ensure (environment = "dev") in terraform.tfvars file

## Production Environment

Update GCS Bucket (your bucket name) and State file prefix as “prod/gke-cluster”

Ensure (environment = "prod") in terraform.tfvars file

✓ 24-GKE-Infra-DevOps-CloudBuild-GitHub

✓ Git-Repo-files

✓ environments

✓ dev

└ c1-versions.tf

└ c2-01-variables.tf

└ c2-02-local-values.tf

└ c3-vpc.tf

└ c4-01-gke-cluster.tf

└ c4-02-gke-outputs.tf

└ c5-Cloud-NAT-Cloud-Router.tf

└ **terraform.tfvars**

✓ prod

└ c1-versions.tf

└ c2-01-variables.tf

└ c2-02-local-values.tf

└ c3-vpc.tf

└ c4-01-gke-cluster.tf

└ c4-02-gke-outputs.tf

└ c5-Cloud-NAT-Cloud-Router.tf

└ **terraform.tfvars**

# What are we going to learn?

Demo: Custom VPC + GKE Cluster(Custom or Local Terraform Module)  
Implement DevOps Pipelines for Terraform Configs on GCP GKE

All these TF Configs are copy from Root Module Demo-23 (p1-gke-autopilot-cluster-private folder)

DEV

TFVA  
RS

## Development Environment

Update all CIDR ranges by ensuring it doesn't overlap with other environments

PROD

TFVA  
RS

## Production Environment

Update all CIDR ranges by ensuring it doesn't overlap with other environments

# What are we going to learn?

✓ 24-GKE-Infra-DevOps-CloudBuild-GitHub

✓ Git-Repo-files

> environments

✓ modules/gke\_cluster

└ main.tf

└ outputs.tf

└ variables.tf

└ versions.tf

❖ .gitignore

! cloudbuild.yaml

\$ git-deploy.sh

Demo: Custom VPC + GKE Cluster(Custom or Local Terraform Module)

Implement DevOps Pipelines for Terraform Configs on GCP GKE

All these TF Configs are copy from Child Module Demo-23  
(modules folder)

CHILD  
MODULE

cloud  
build

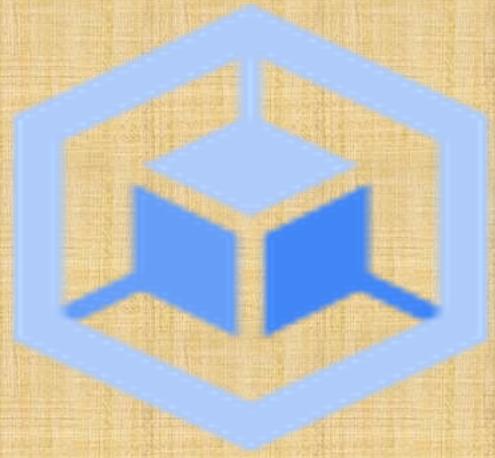
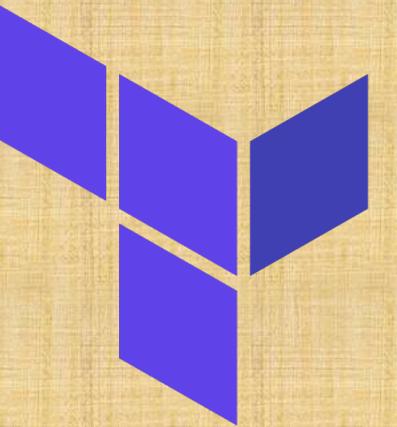
GKE Cluster Module we have created in Demo-23

cloudbuild.yaml contains Terraform commands with  
DevOps pipeline code

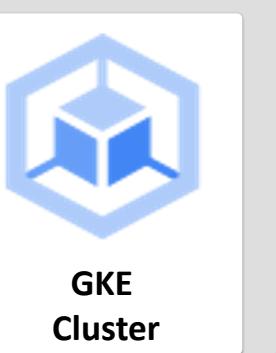
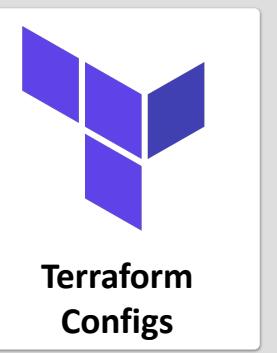
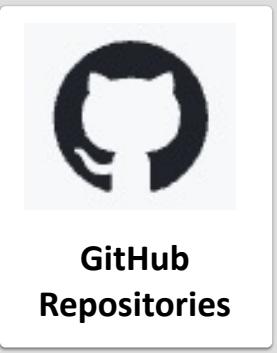
Demo-25, 26, 27, 28



# Google Kubernetes Engine Continuous Integration Continuous Delivery

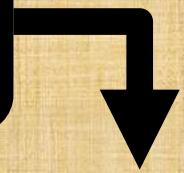


# Continuous Integration and Continuous Delivery



Demo  
25

Create Kubernetes Deployment Custom  
Terraform Module



Demo  
26

Implement k8s workload DevOps  
pipelines  
(Dev and Prod Environments)



Demo  
27

App: Implement Continuous Integration  
Pipeline with Artifact Registry



Demo  
28

App: Implement Continuous Integration  
and Continuous Delivery Pipeline (CI CD)



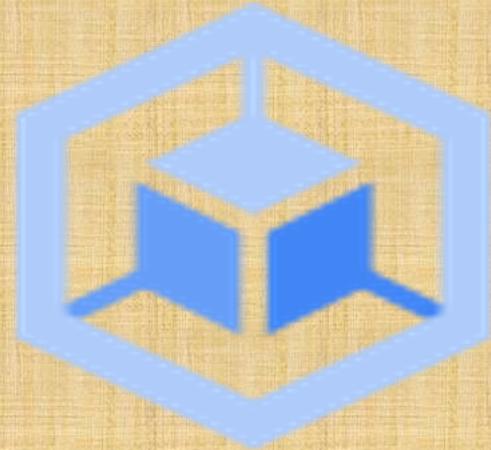
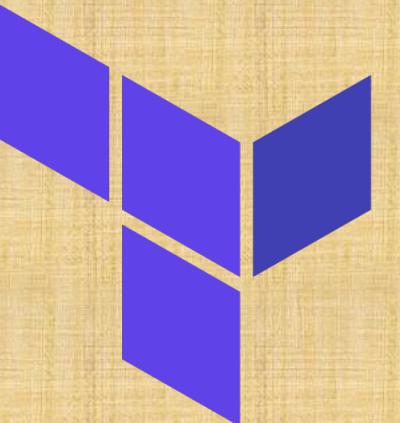
Demo-25

# Google Kubernetes Engine

## Terraform Custom Module

## Kubernetes Deployment

**Demo:** Create Kubernetes Deployment by building a **custom**  
**Terraform Module**



## Root Module

```
✓ 25-GKE-Workloads-Custom-Terraform-Modules
  > modules
    ✓ p2-k8sresources-terraform-manifests
      ✓ c1-versions.tf
      ✓ c2-01-variables.tf
      ✓ c2-02-local-values.tf
      ✓ c3-01-remote-state-datasource.tf
      ✓ c3-02-providers.tf
      ✓ c4-kubernetes-deployment.tf
      ✓ c5-kubernetes-loadbalancer-service.tf
    terraform.tfvars
```

Call GKE Cluster child  
module in Root Module

## Child Module

```
✓ 25-GKE-Workloads-Custom-Terraform-Modules
  ✓ modules
    > gke_cluster
      ✓ kubernetes_deployment
        ✓ main.tf
        ✓ outputs.tf
        ✓ variables.tf
        ✓ versions.tf
```

# Build a Local Terraform Module

Modules  
Demo

Build a Local Terraform Module  
and call it from Root Module and  
Test

## Modules: Kubernetes Deployment Custom Module

### Child Module

- ✓ 25-GKE-Workloads-Custom-Terraform-Modules
- ✓ modules
  - > gke\_cluster
    - ✓ kubernetes\_deployment
      - main.tf
      - outputs.tf
      - variables.tf
      - versions.tf

# What are we going to learn?

## Demo: Create GKE Cluster Custom Module



Define Terraform **Provider** versions

Define Terraform **Input Variables**

Create **Kubernetes Deployment** Terraform resource

Define Terraform **Outputs**

## Project-1: Terraform Manifests

### Root Module

- ✓ 25-GKE-Workloads-Custom-Terraform-Modules
  - > modules
- ✓ p2-k8sresources-terraform-manifests
  - c1-versions.tf
  - c2-01-variables.tf
  - c2-02-local-values.tf
  - c3-01-remote-state-datasource.tf
  - c3-02-providers.tf
  - c4-kubernetes-deployment.tf
  - c5-kubernetes-loadbalancer-service.tf
  - terraform.tfvars

# What are we going to learn?

Demo: Project-2: Deploy Kubernetes Deployment and Load Balancer Service by calling Terraform Module

NC

No changes c1 to c3, c5

c4

Create Kubernetes Deployment by calling custom Kubernetes Deployment Terraform module

Demo-26

# Google Kubernetes Engine

## GKE Workload DevOps

## GCP Cloud Build, GitHub GCP Cloud Build App and GitHub

**Demo: Implement DevOps Pipelines for Kubernetes Deployment Terraform Configs on  
GCP GKE**

# Kubernetes Deployment DevOps for Terraform Configs

- The following resources are used for this implementation

## • GitHub Repositories

- GitHub Branches
  - Main branch
  - Dev branch
  - Prod branch
  - Feature branch
- GitHub Pull Request and Merge concepts
- GCP Cloud Build App from GitHub Marketplace

## • GCP Cloud Build

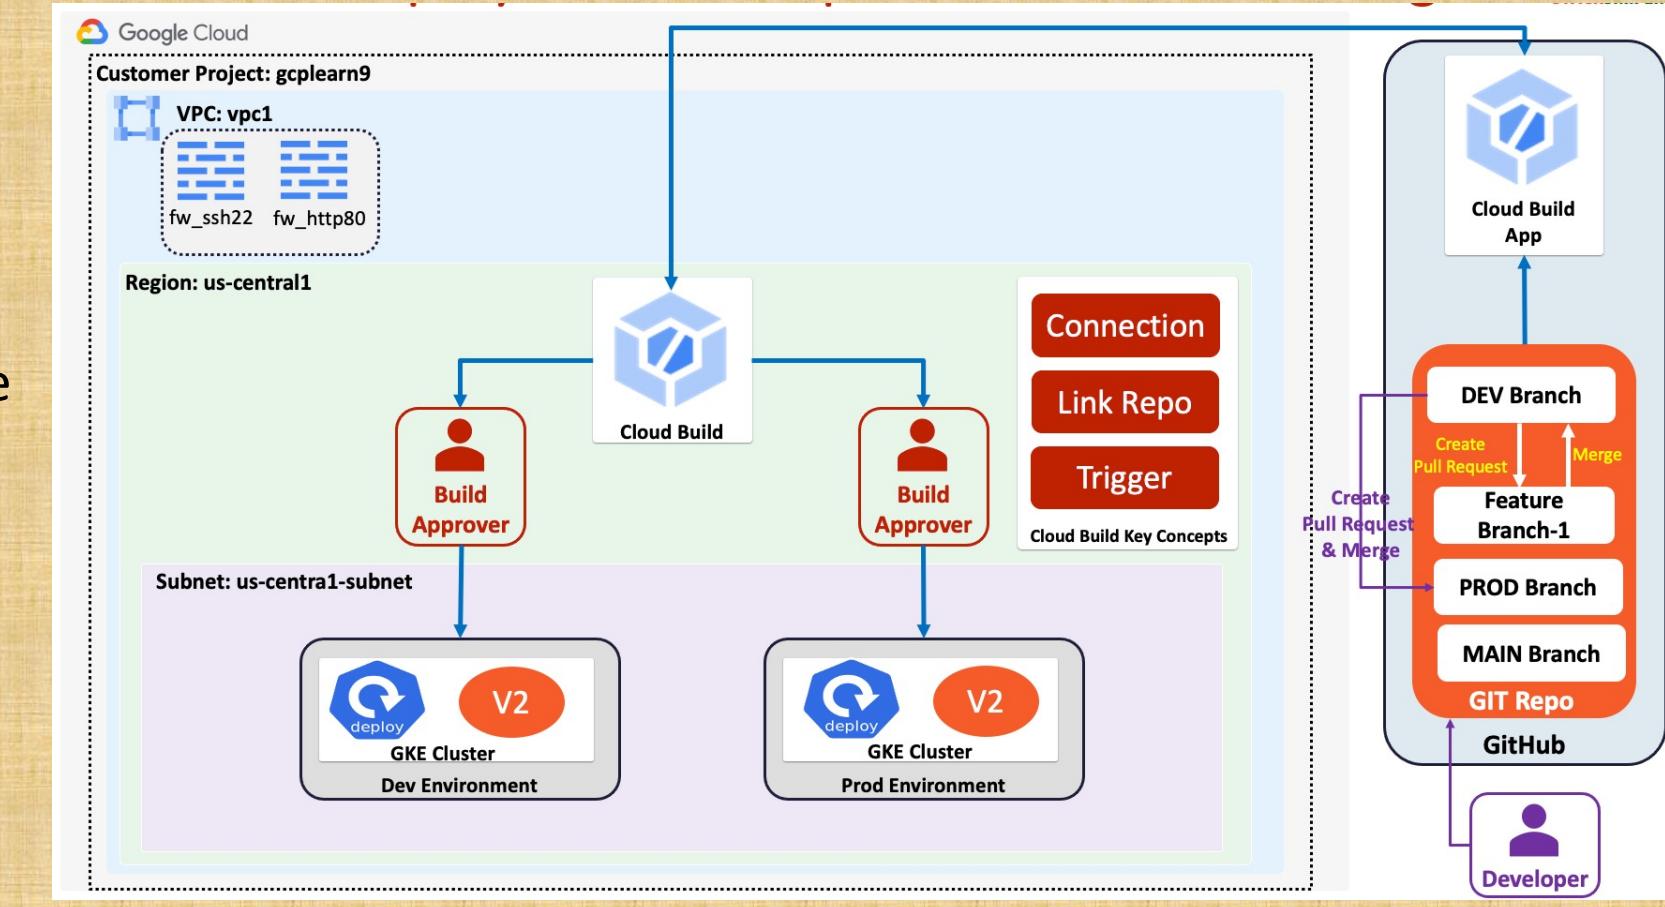
- Connections
- Link Repositories
- Triggers

## • Kubernetes Resources

- Kubernetes Deployments
- Kubernetes Load Balancer Services

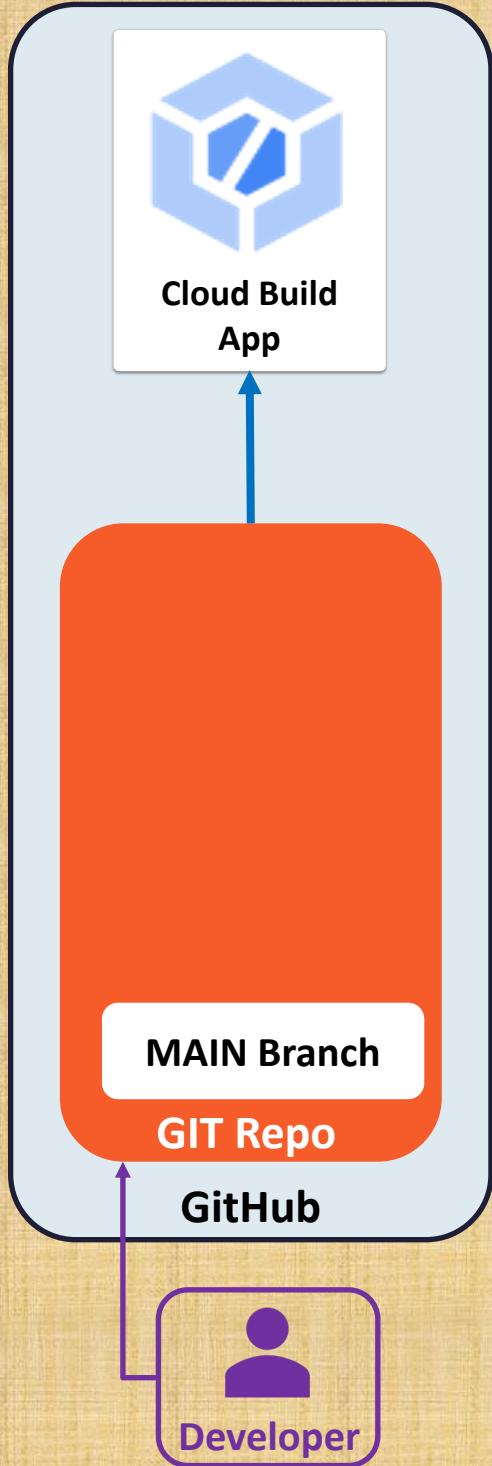
## • Terraform Concepts

- Terraform Resources
- Terraform custom modules
- Terraform Remote Backend using Google Cloud Storage Bucket

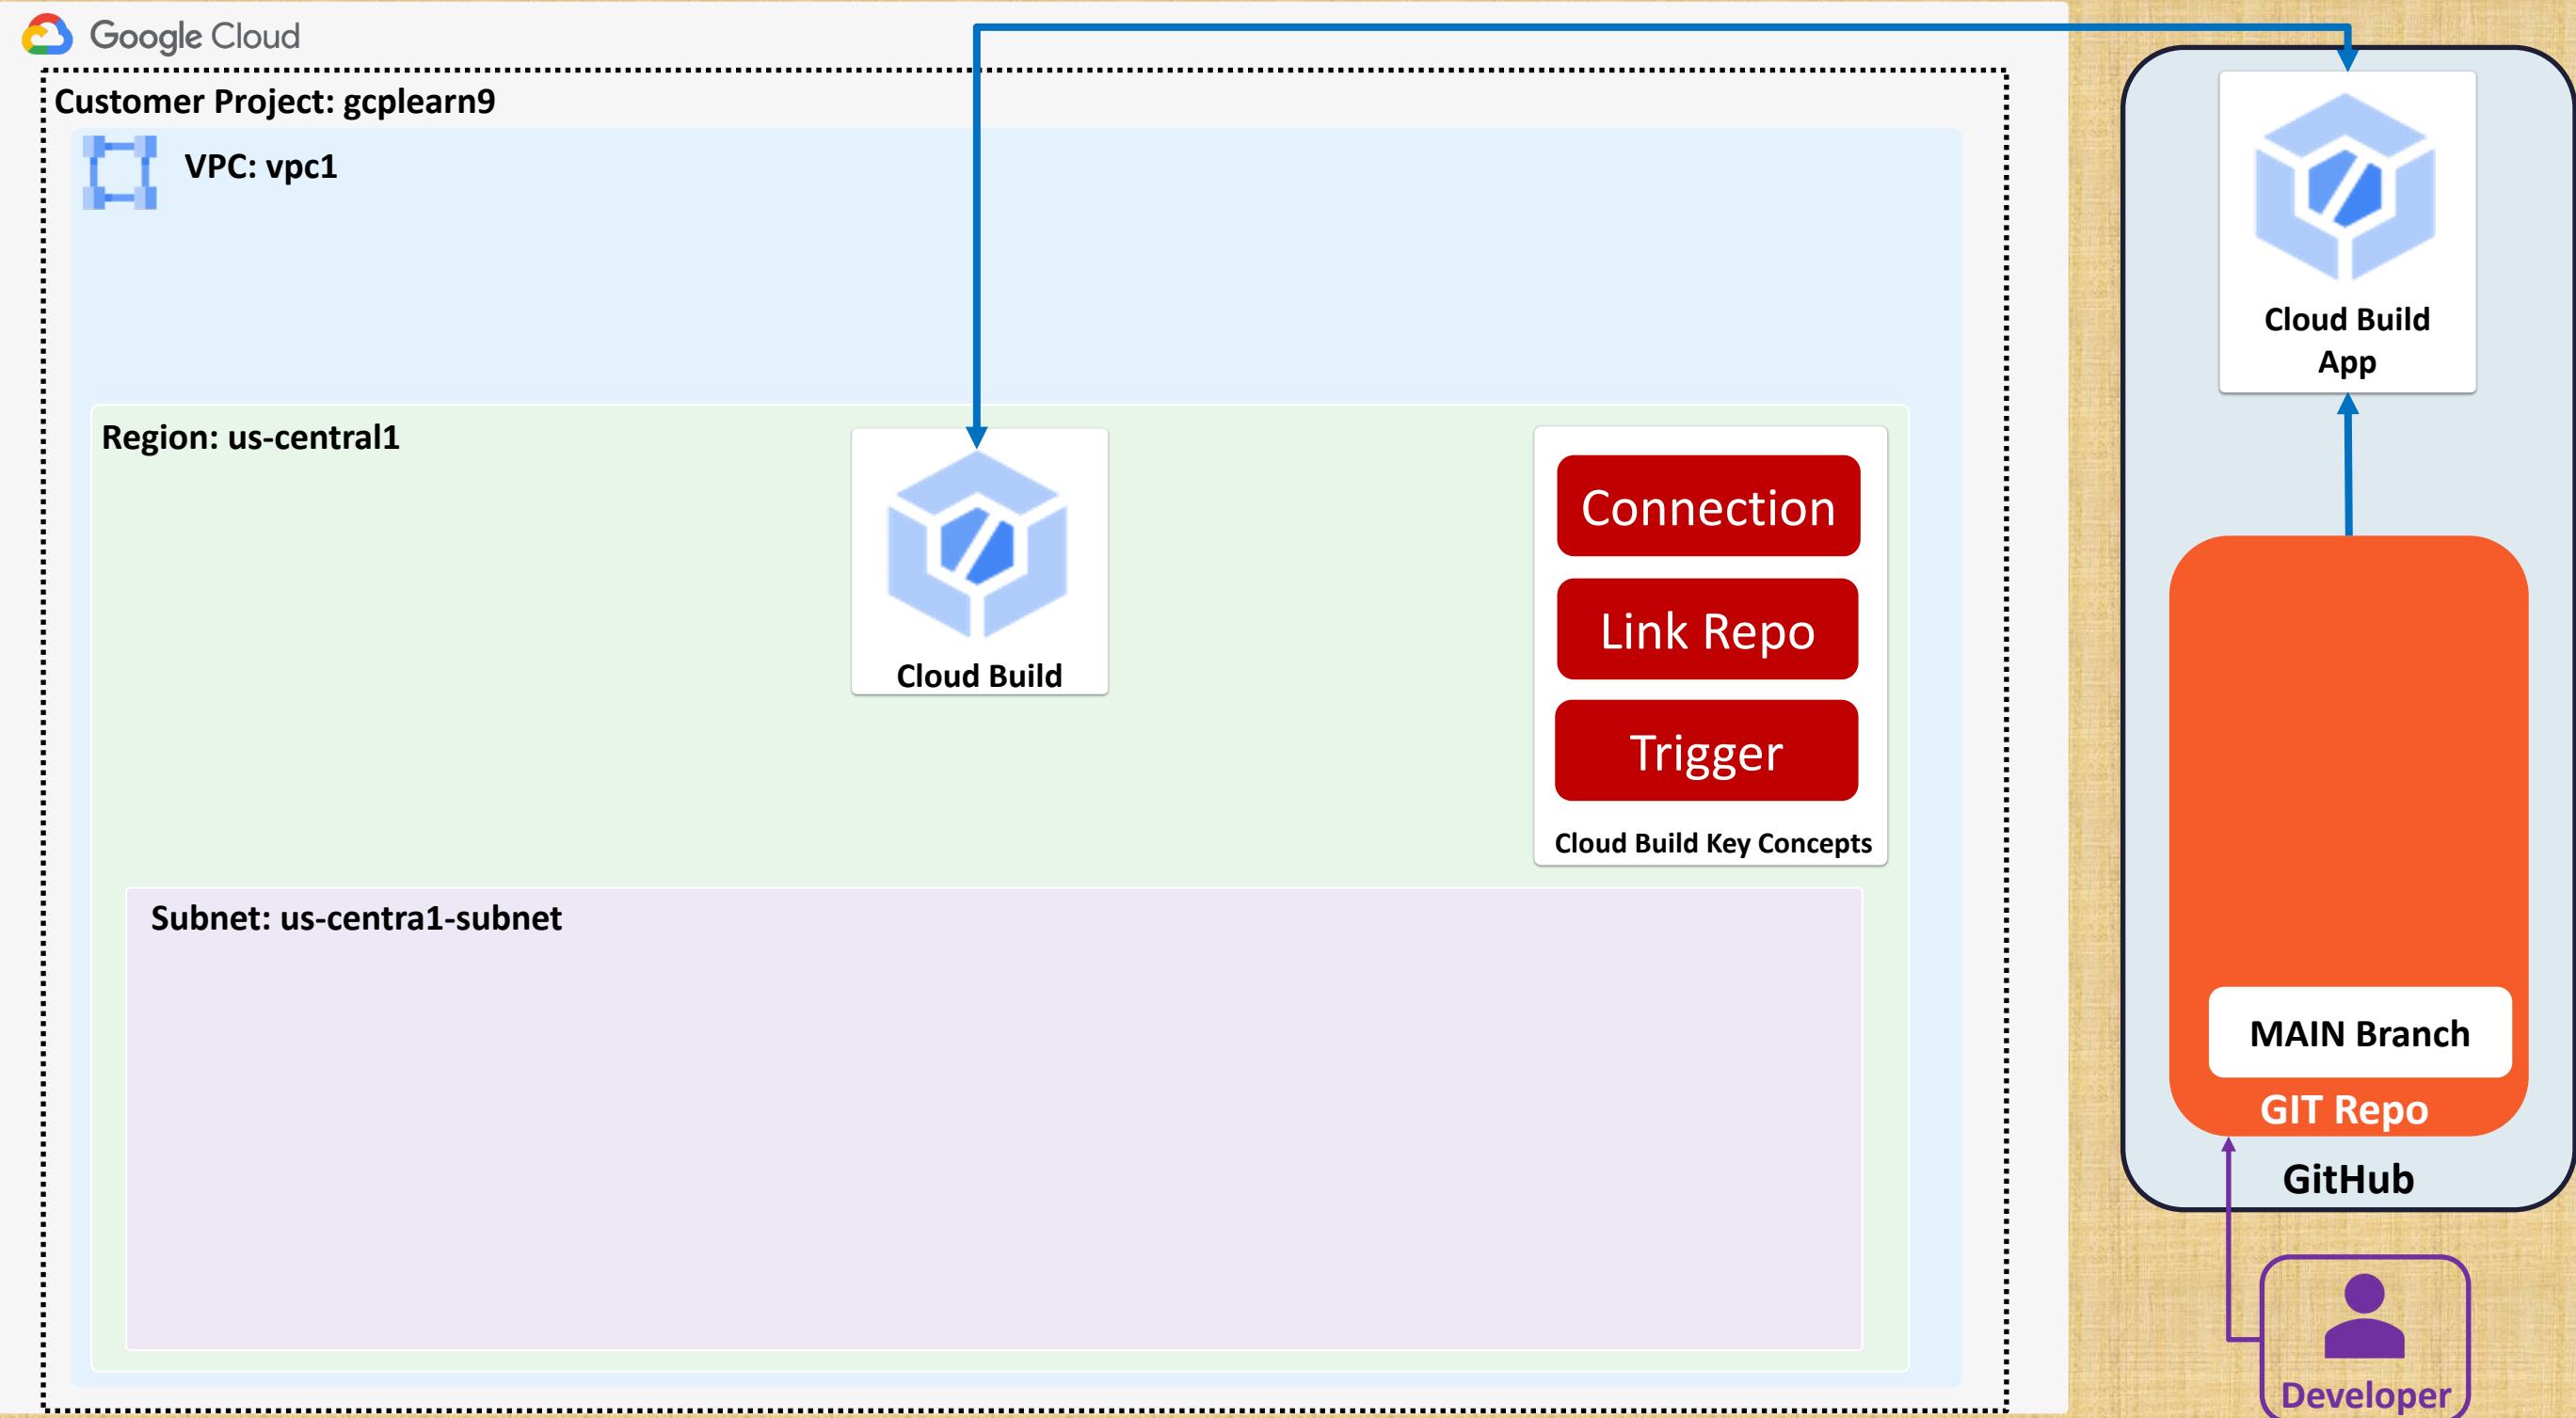


GIT Repo

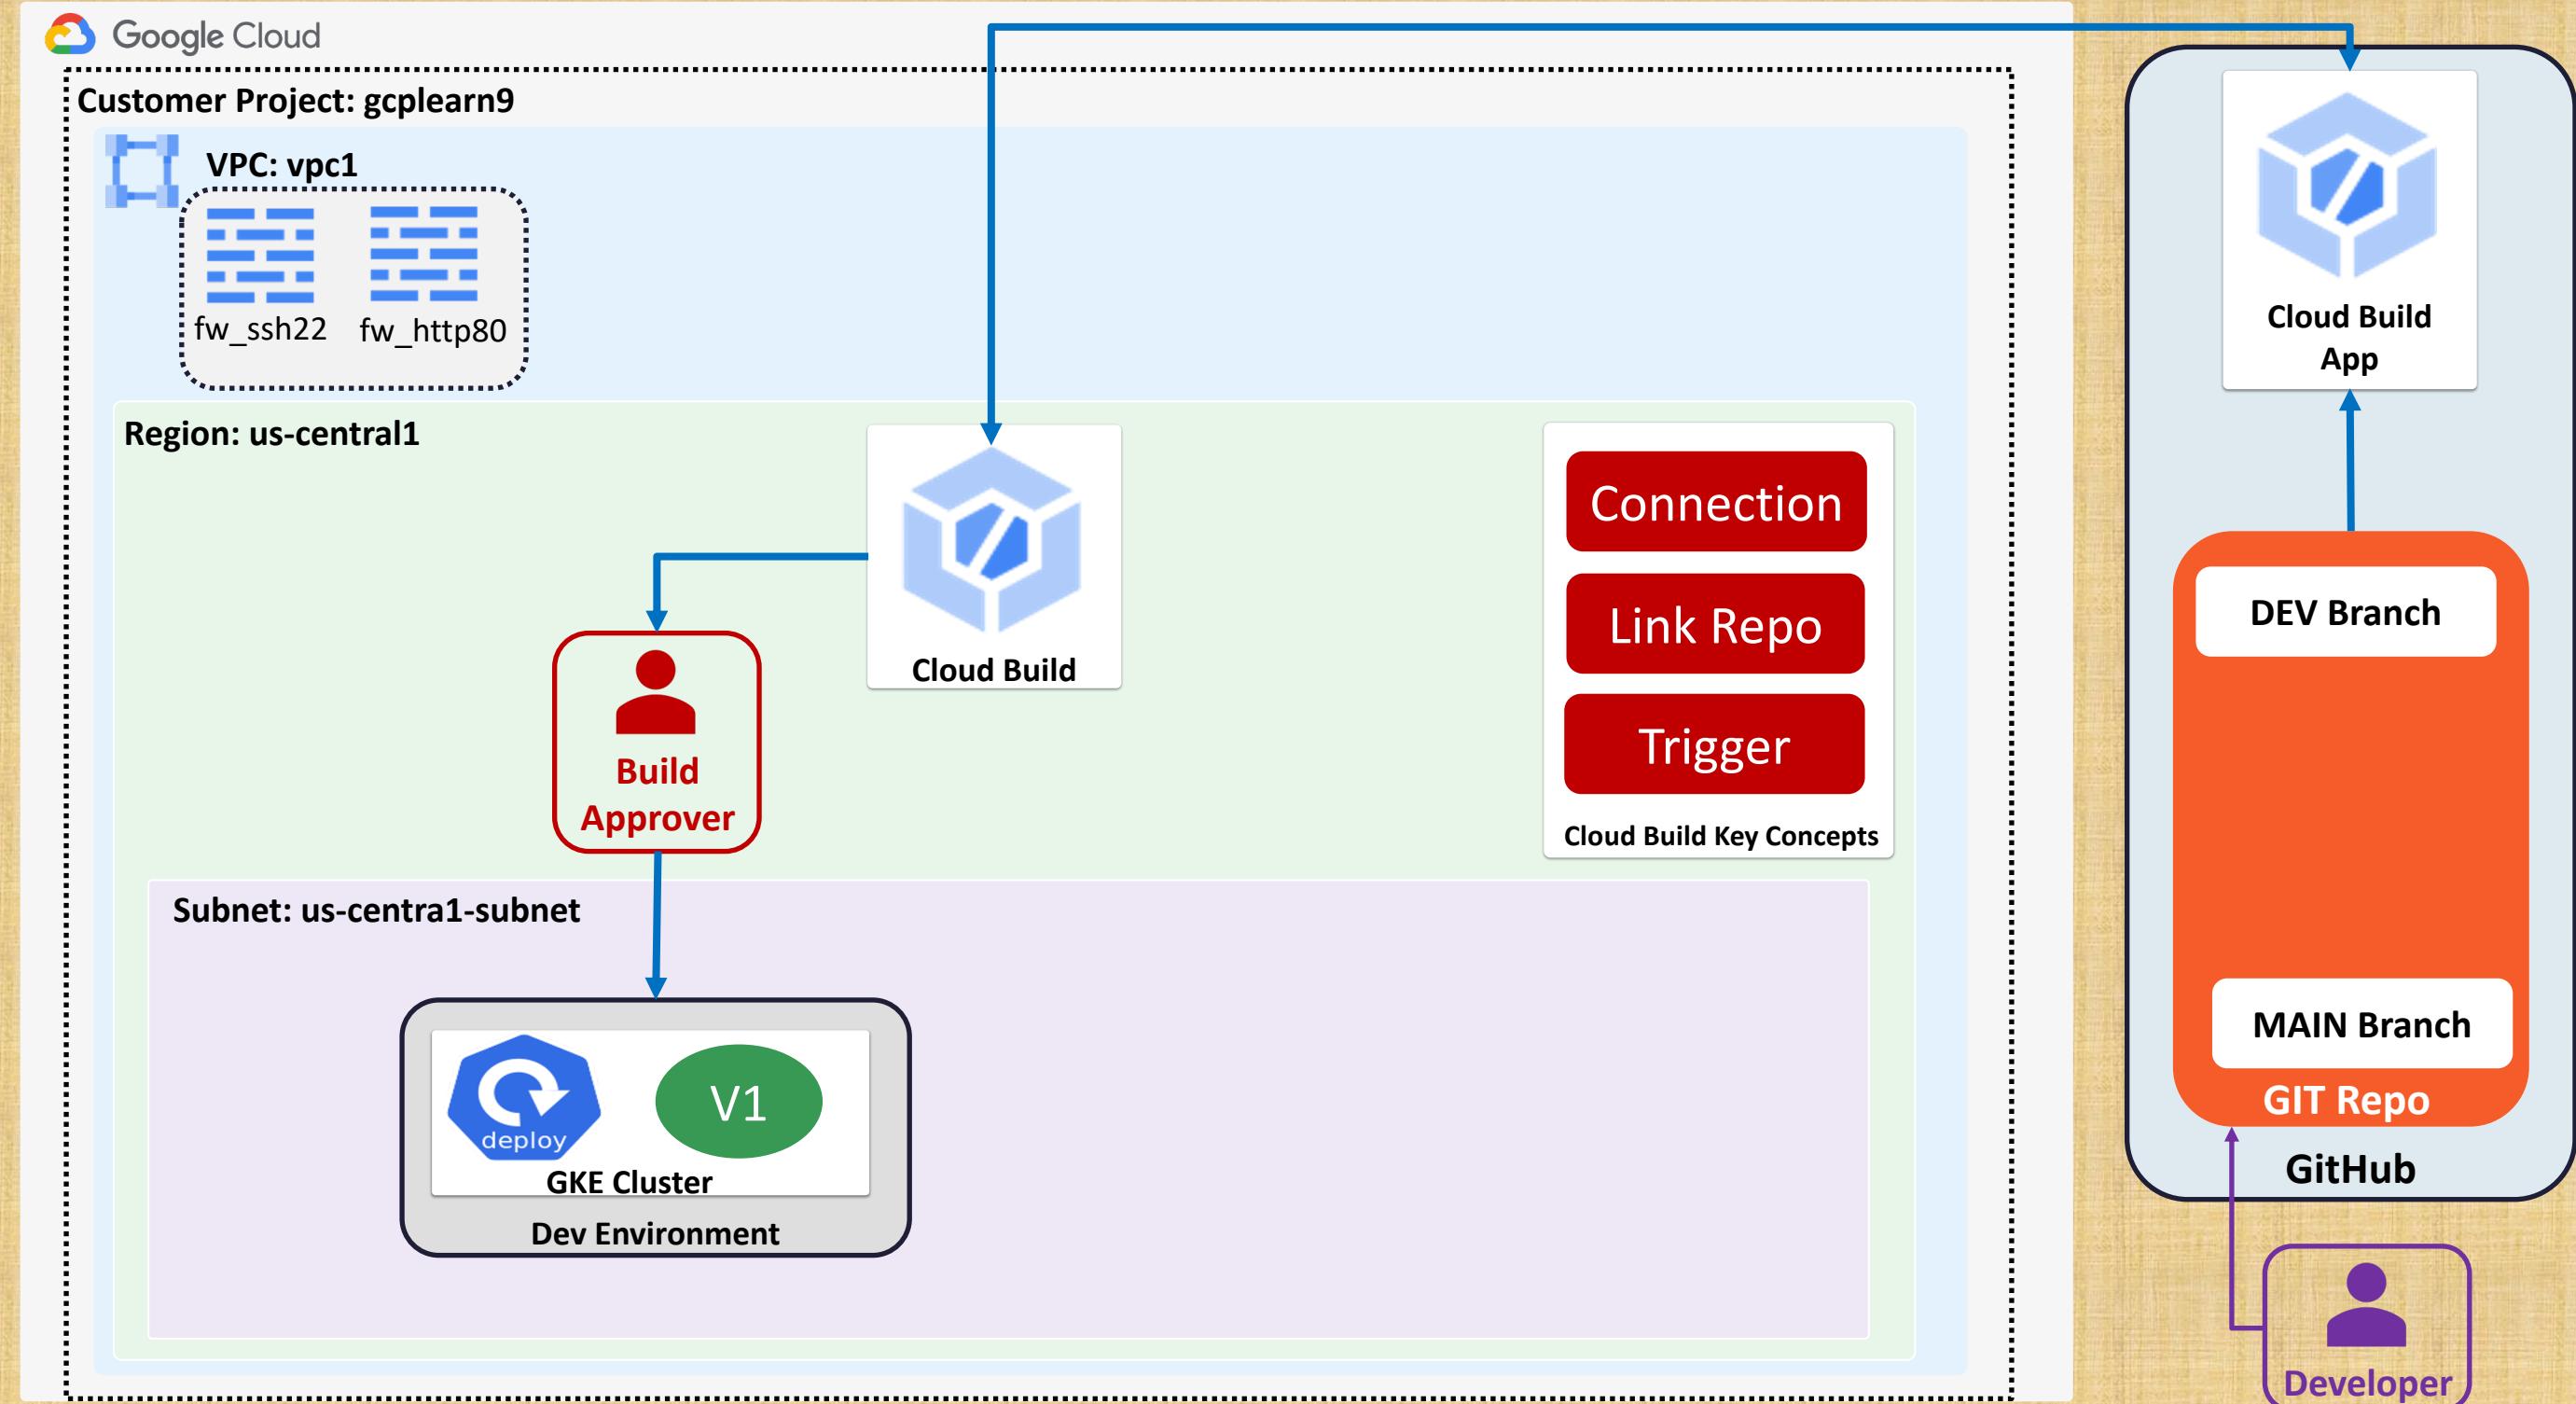
# Kubernetes Deployment DevOps for Terraform Configs



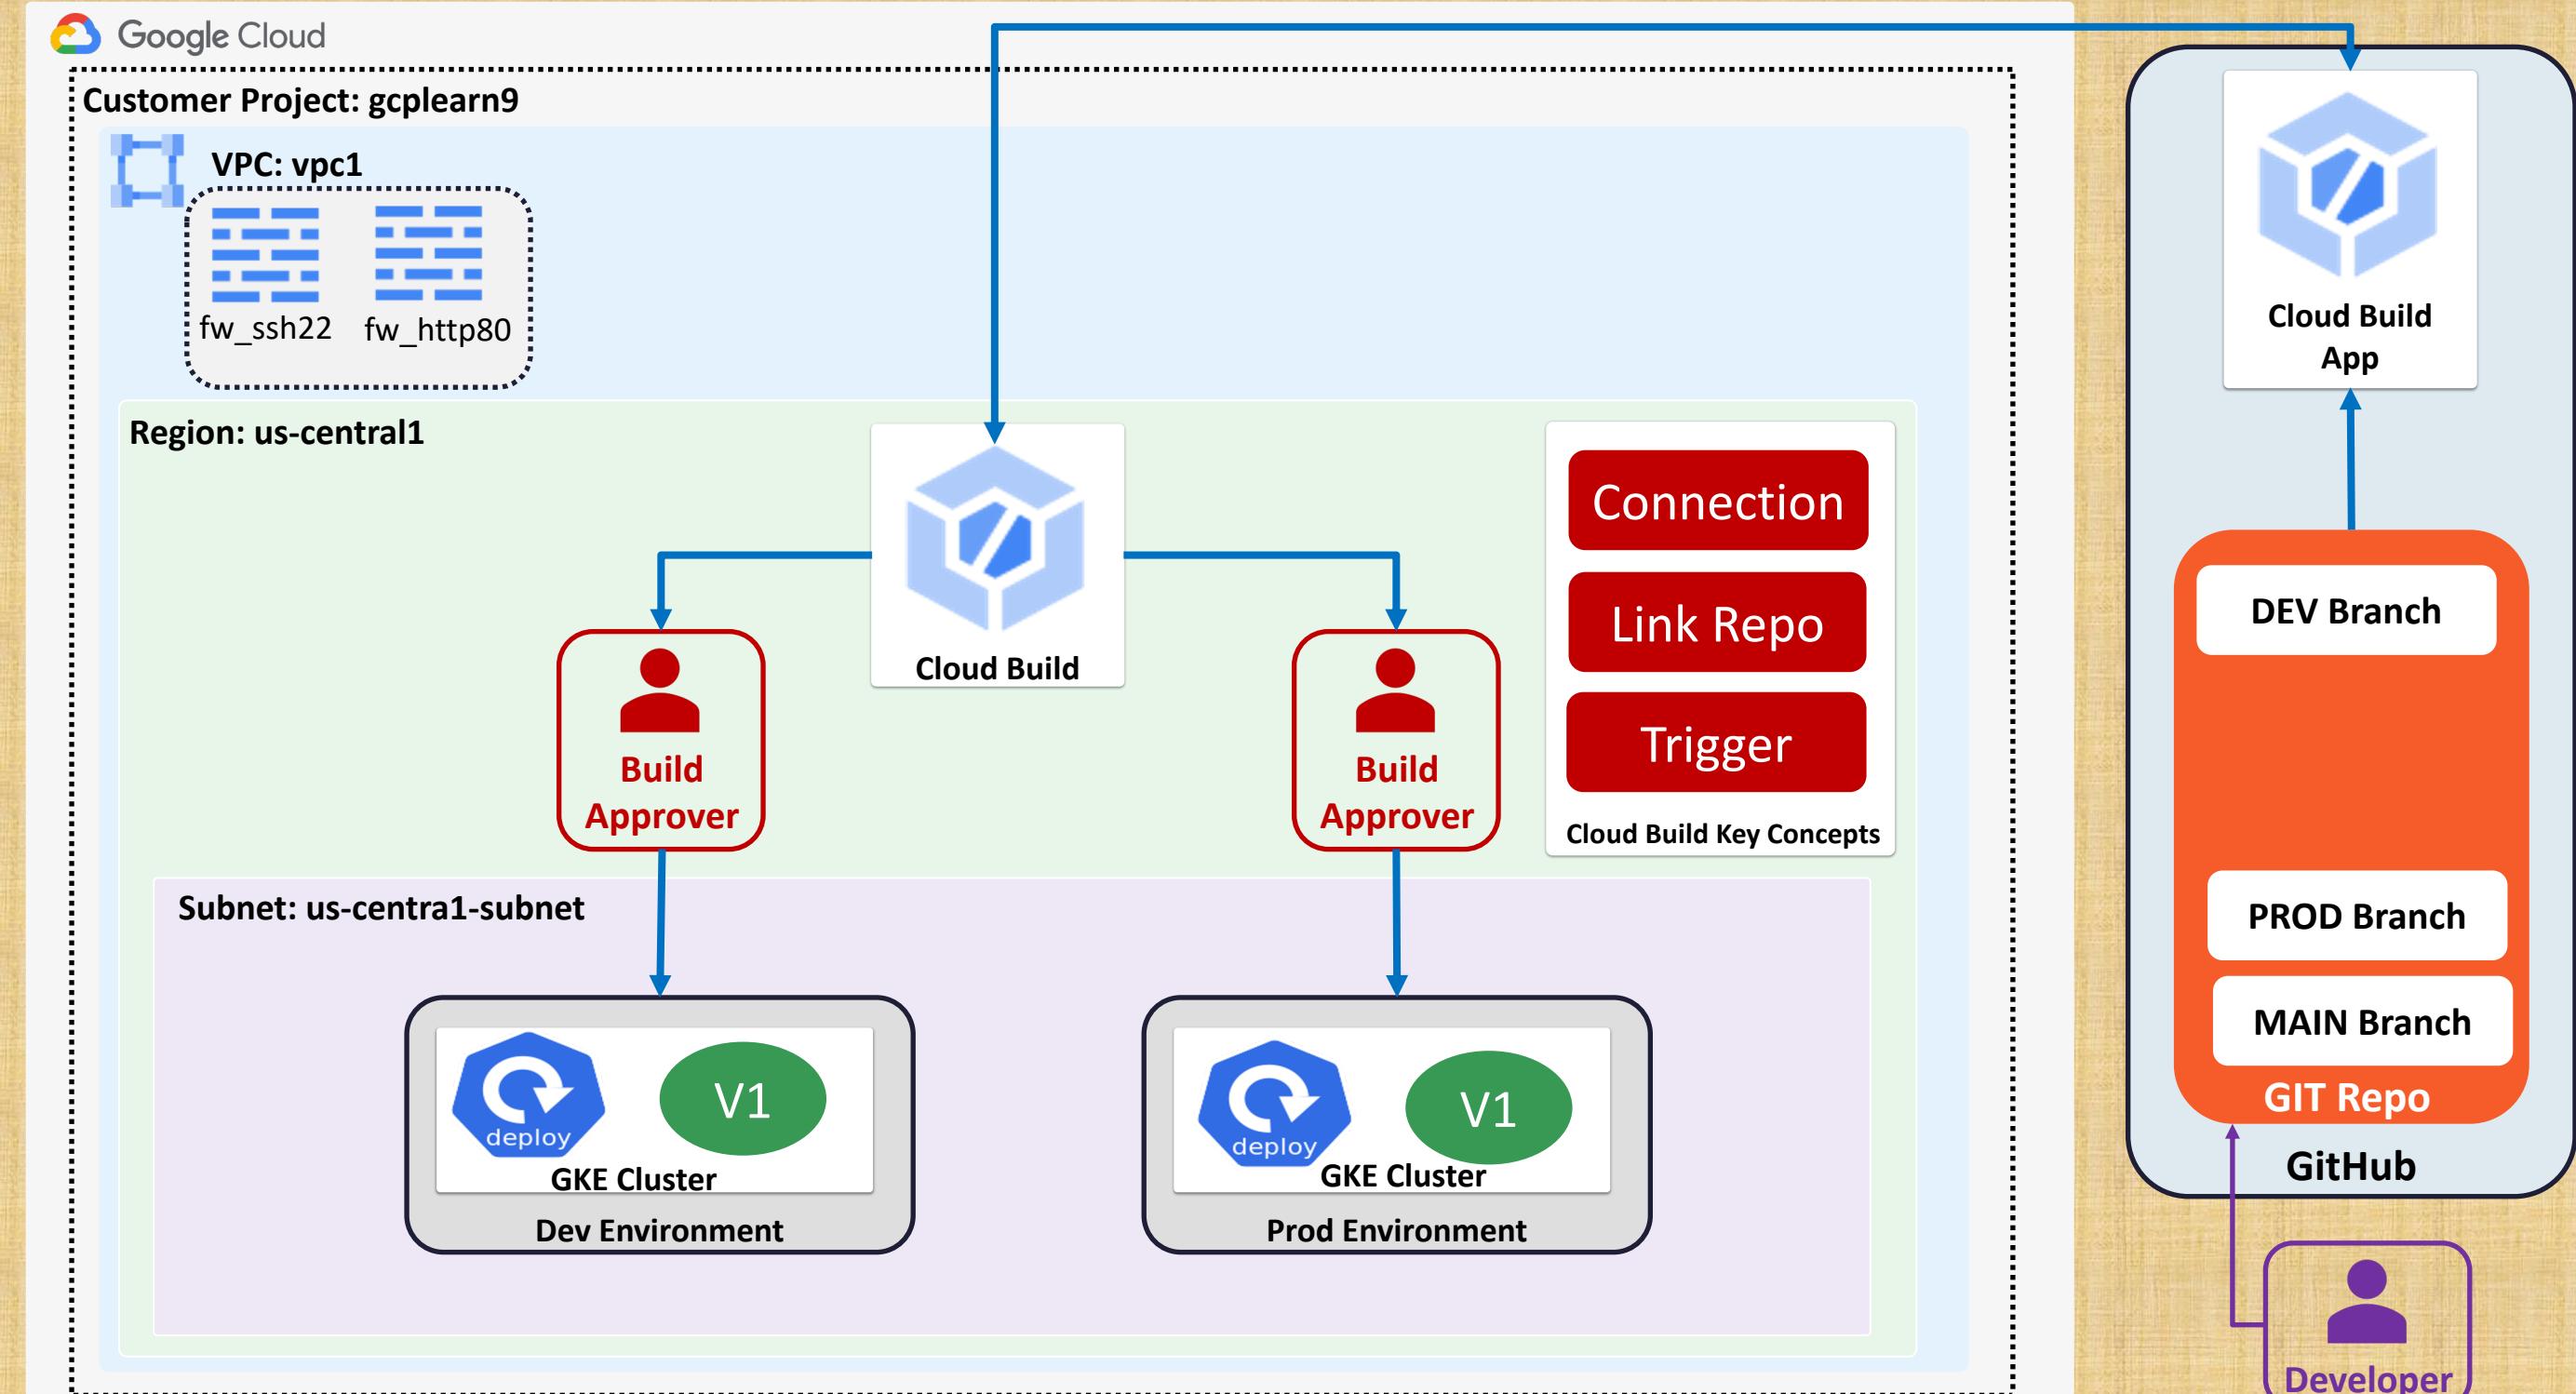
# Kubernetes Deployment DevOps for Terraform Configs



# Kubernetes Deployment DevOps for Terraform Configs

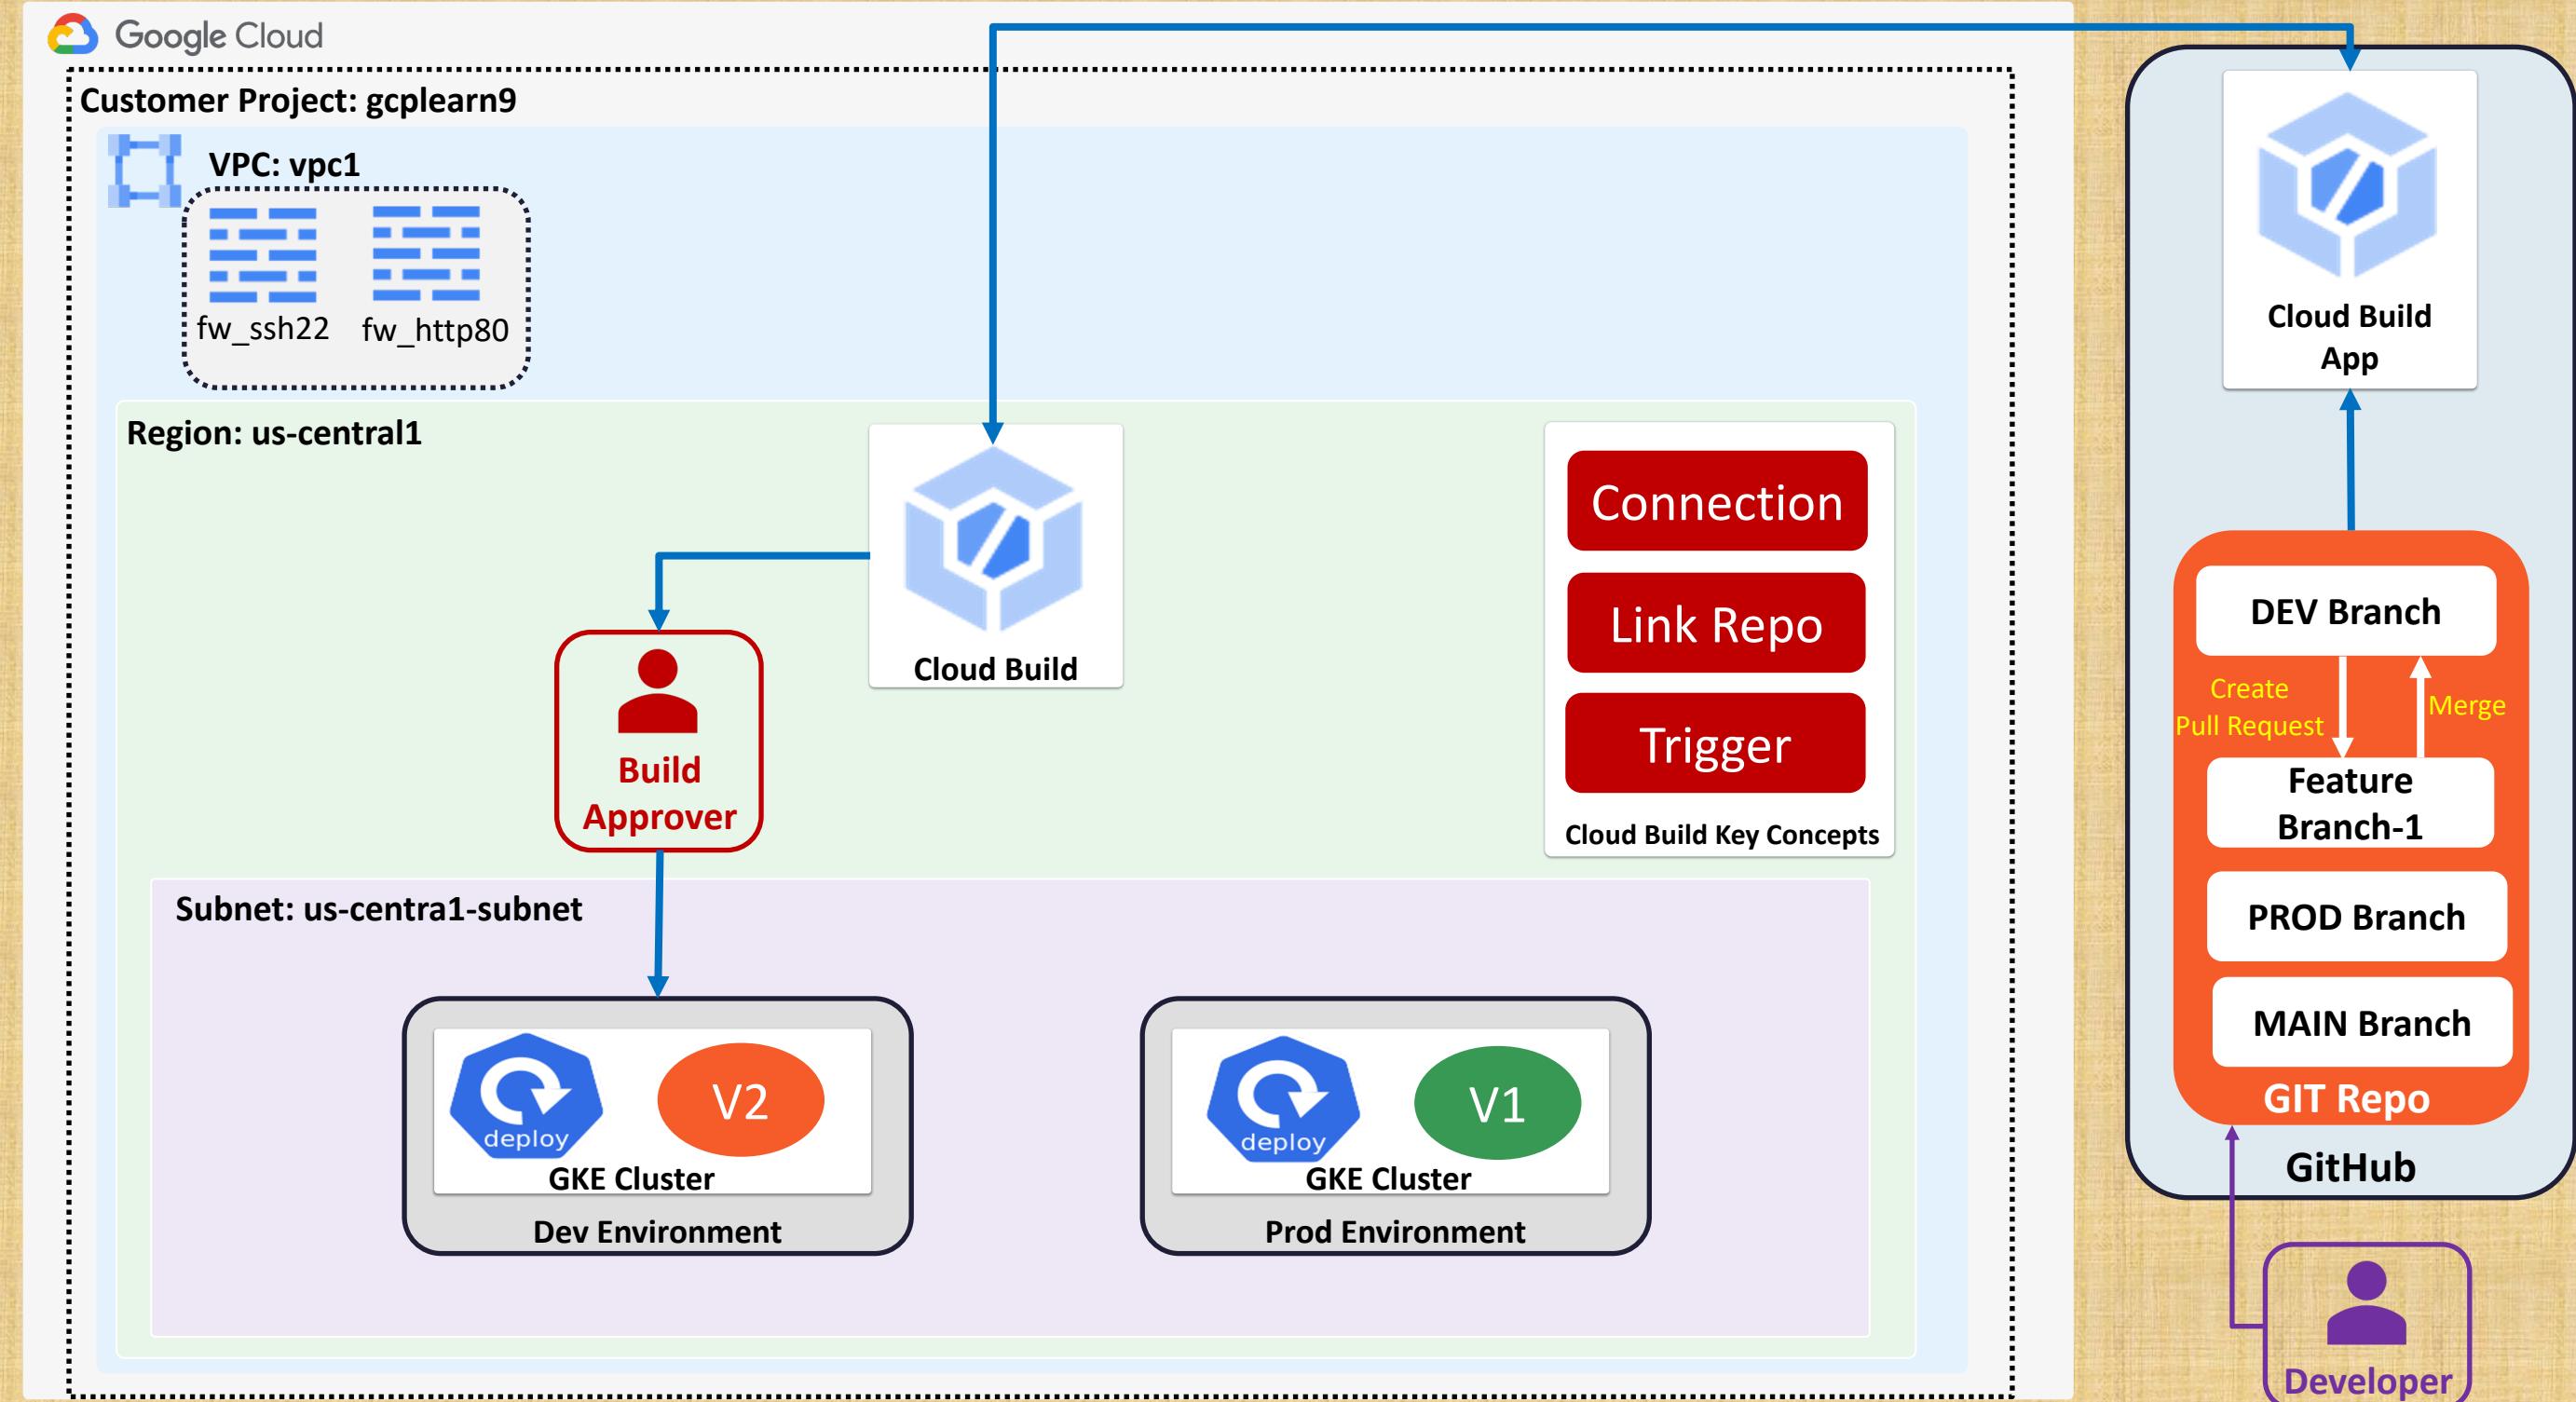


# Kubernetes Deployment DevOps for Terraform Configs

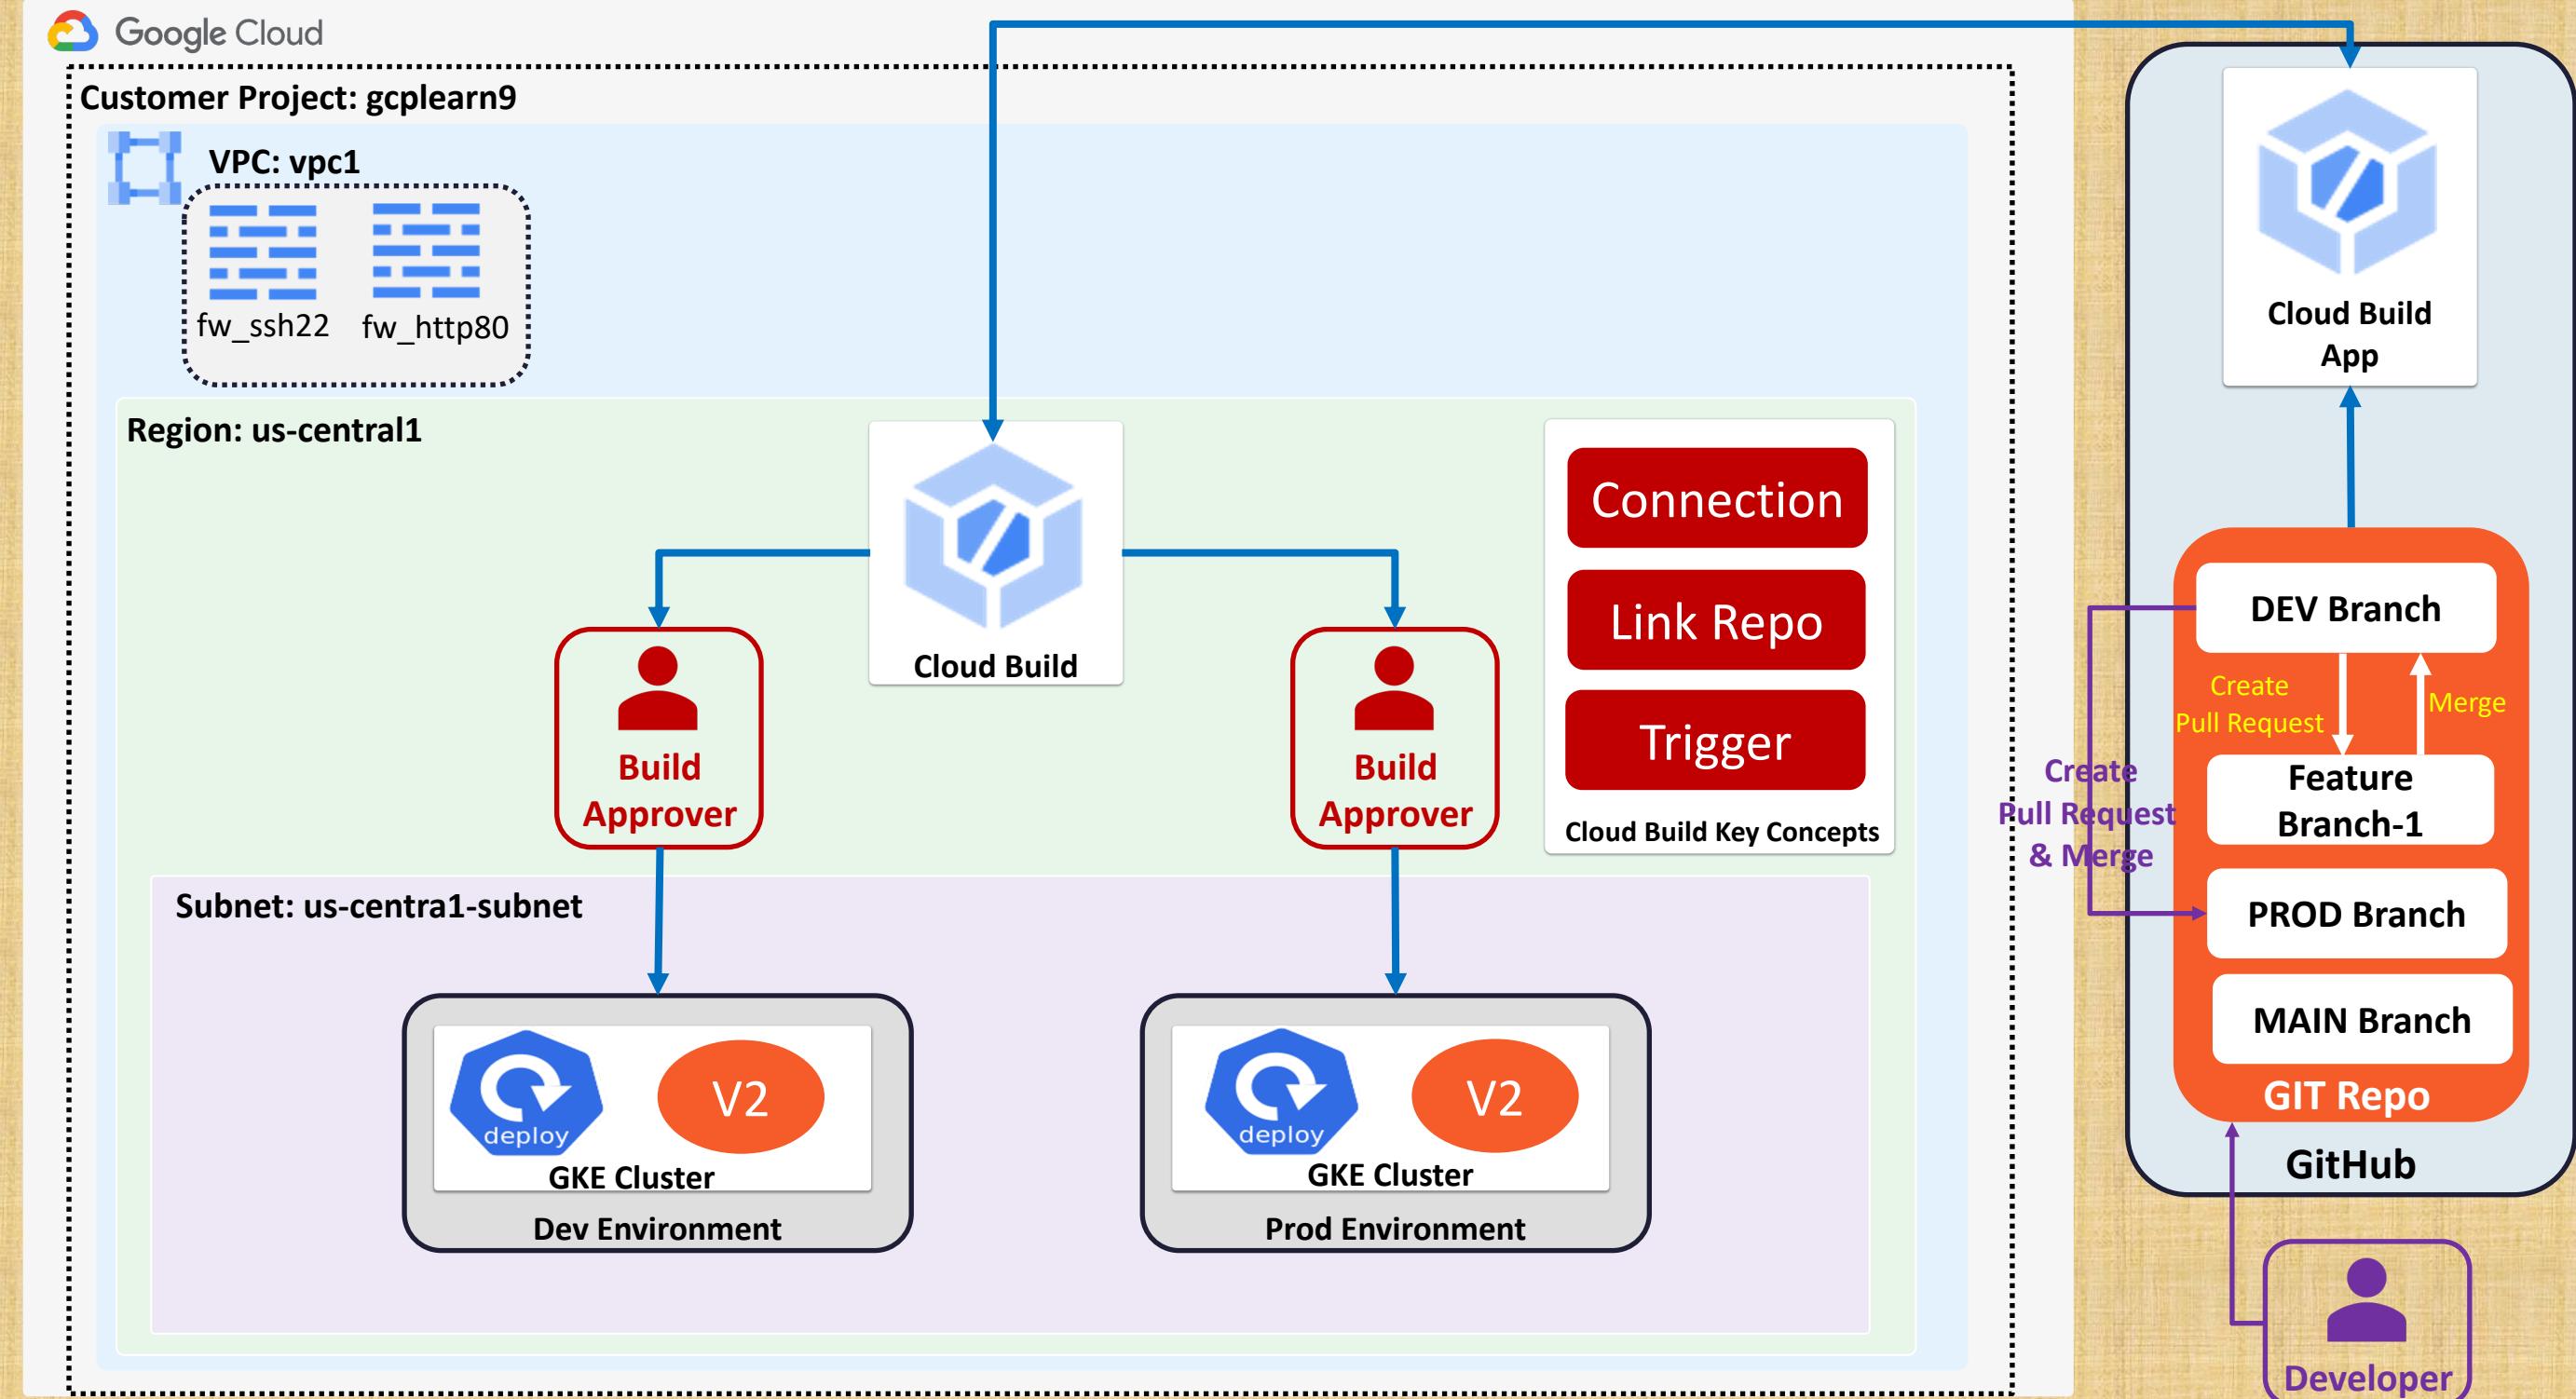


Author: Nho Luong

# Kubernetes Deployment DevOps for Terraform Configs



# Kubernetes Deployment DevOps for Terraform Configs



✓ 26-GKE-Workloads-DevOps-CloudBuild-GitHub

✓ Git-Repo-Files

✓ environments

✓ dev

- ↳ c1-versions.tf
- ↳ c2-01-variables.tf
- ↳ c2-02-local-values.tf
- ↳ c3-01-remote-state-datasource.tf
- ↳ c3-02-providers.tf
- ↳ c4-kubernetes-deployment.tf
- ↳ c5-kubernetes-loadbalancer-service.tf
- ↳ terraform.tfvars

✓ prod

- ↳ c1-versions.tf
- ↳ c2-01-variables.tf
- ↳ c2-02-local-values.tf
- ↳ c3-01-remote-state-datasource.tf
- ↳ c3-02-providers.tf
- ↳ c4-kubernetes-deployment.tf
- ↳ c5-kubernetes-loadbalancer-service.tf
- ↳ terraform.tfvars

# What are we going to learn?

## Demo: Implement DevOps Pipelines for Kubernetes Deployment Terraform Configs on GCP GKE



### Development Environment

Update GCS Bucket (**your** bucket name) and State file prefix as “workloads/dev/k8s-myapp1”

Remote State Data source: Access GKE Dev Cluster Outputs “dev/gke-cluster”

Ensure (**environment = "dev"**) in terraform.tfvars file

### Production Environment

Update GCS Bucket (**your** bucket name) and State file prefix as “workloads/prod/k8s-myapp1”

Remote State Data source: Access GKE Dev Cluster Outputs “prod/gke-cluster”

Ensure (**environment = "prod"**) in terraform.tfvars file

# What are we going to learn?

- ✓ 26-GKE-Workloads-DevOps-CloudBuild-GitHub
- ✓ Git-Repo-Files
  - > environments
  - ✓ modules/kubernetes\_deployment
    - ✗ main.tf
    - ✗ outputs.tf
    - ✗ variables.tf
    - ✗ versions.tf
  - ✗ .gitignore
  - ✗ ! cloudbuild.yaml
  - ✗ \$ git-deploy.sh

Demo: Implement DevOps Pipelines for

Kubernetes Deployment Terraform Configs on GCP GKE

All these TF Configs are copy from Child Module Demo-25  
(modules folder)



Kubernetes Deployment Module we have created in  
Demo-25

cloudbuild.yaml contains Terraform commands with  
DevOps pipeline code



Demo-27

# Google Kubernetes Engine

## Continuous Integration



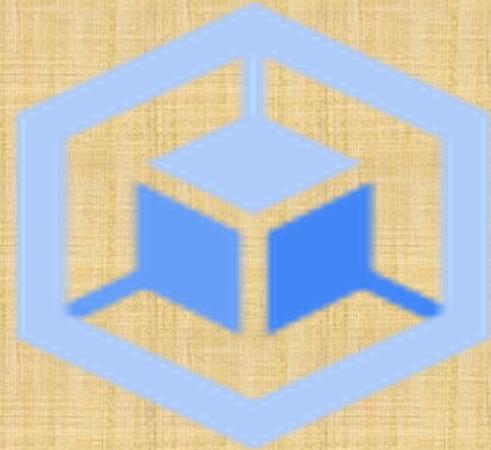
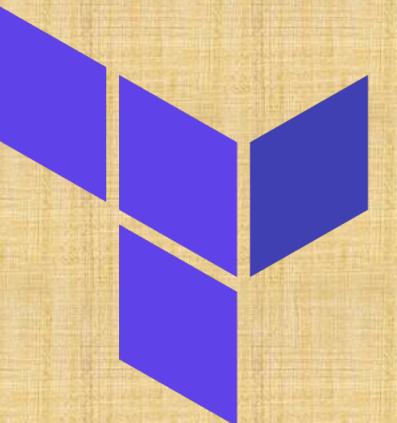
GitHub



Artifact Registry



Cloud Build



**Demo: App on GKE Cluster - Continuous Integration**

# Continuous Integration

## Continuous Integration Pipeline



GitHub



Google  
Cloud Build



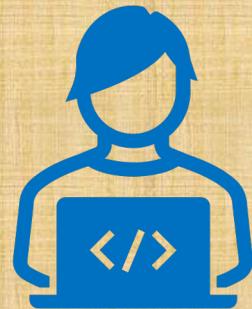
Google  
Artifact Registry

To Store App Code  
(Application Code  
+ Docker Files)

To Build Packages  
(Build Docker  
Image)

To Store Packages  
(Store Docker  
Image)

# Google Cloud Build: Continuous Integration



Admin

Check-in Code



GitHub

Google Cloud

Customer Project: gcplearn9

Region: us-central1

**Trigger the Code Build CI Pipeline**



Google  
Cloud Build

**Build and push new Docker Image**



Google  
Artifact Registry

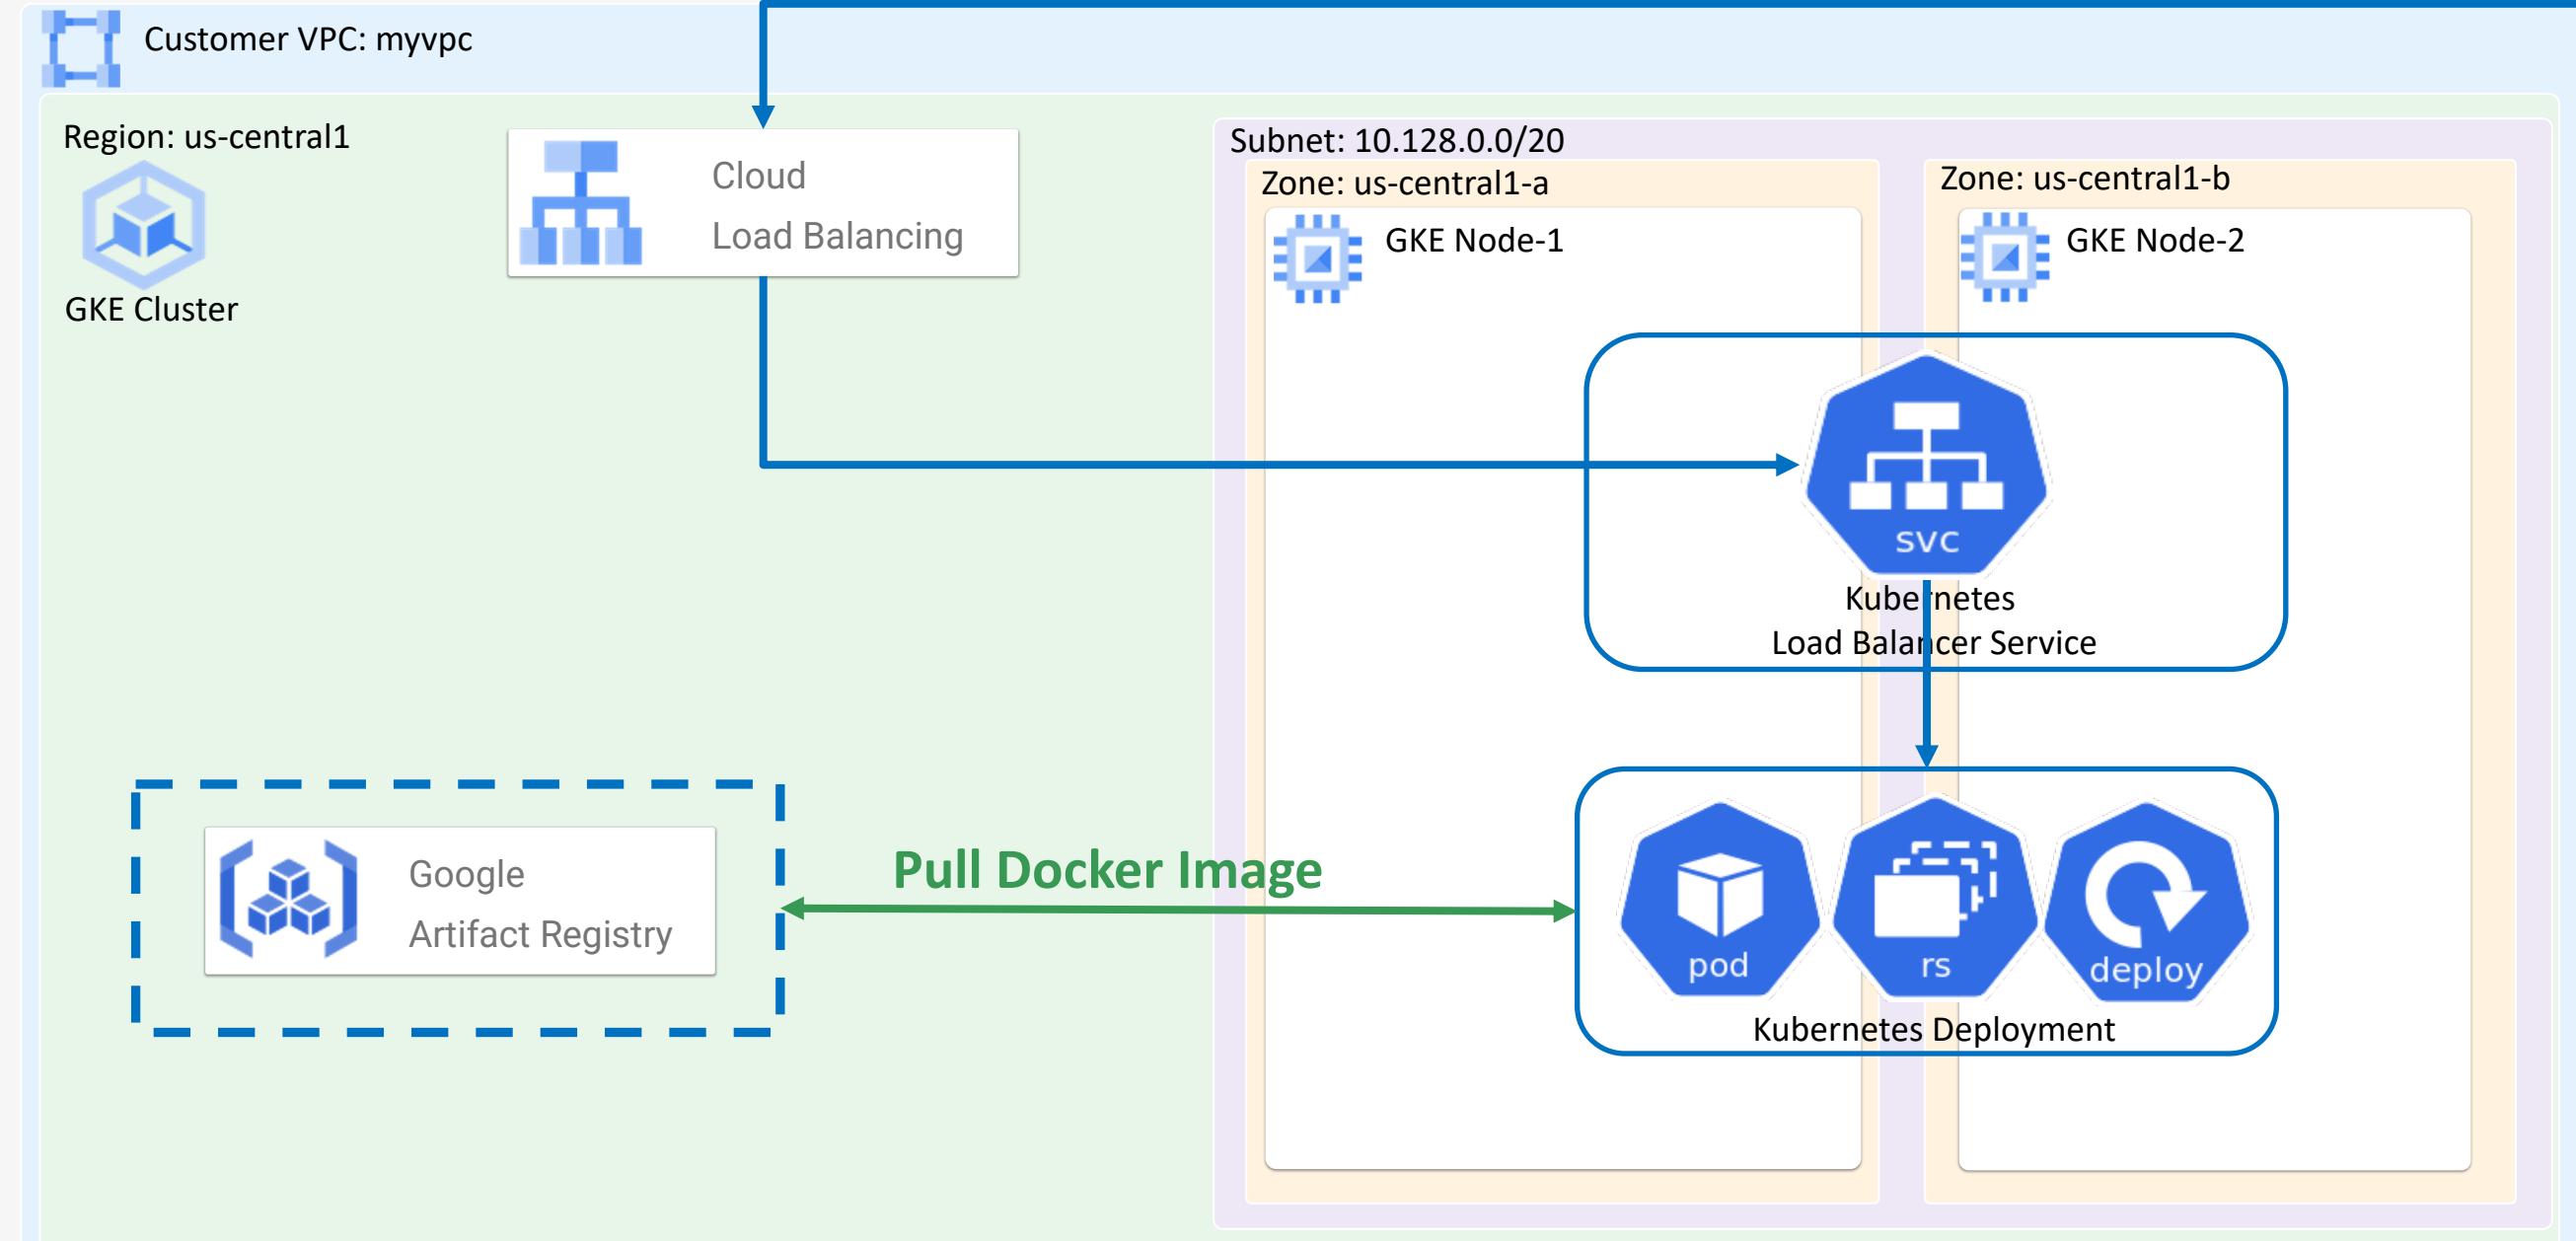
Customer Project: gcplearn9

Automated by Terraform

<http://<Load-Balancer-IP-Address>>

NETWORK DESIGN

Google Artifact Registry



- ✓ 27-GKE-App-Continuous-Integration
- ✓ App-GitRepoFiles
  - ❖ .gitignore
  - ! BACKUP\_v1-cloudbuild-ci.yaml
  - ! cloudbuild.yaml
  - ❖ Dockerfile
  - \$ git-deploy.sh
  - <> index.html
  - ❶ README.md

# What are we going to learn?

## Demo: App on GKE Cluster - Continuous Integration

index.html

Dockerfile

cloud build

Sample Application - index.html page

Dockerfile to create Docker Image for our sample application

cloudbuild.yaml used for implementing CI Pipeline in GCP Cloud Build

Demo-28

# Google Kubernetes Engine

## Continuous Integration & Delivery

### GCP Cloud Build, GitHub GCP Cloud Build App and GitHub

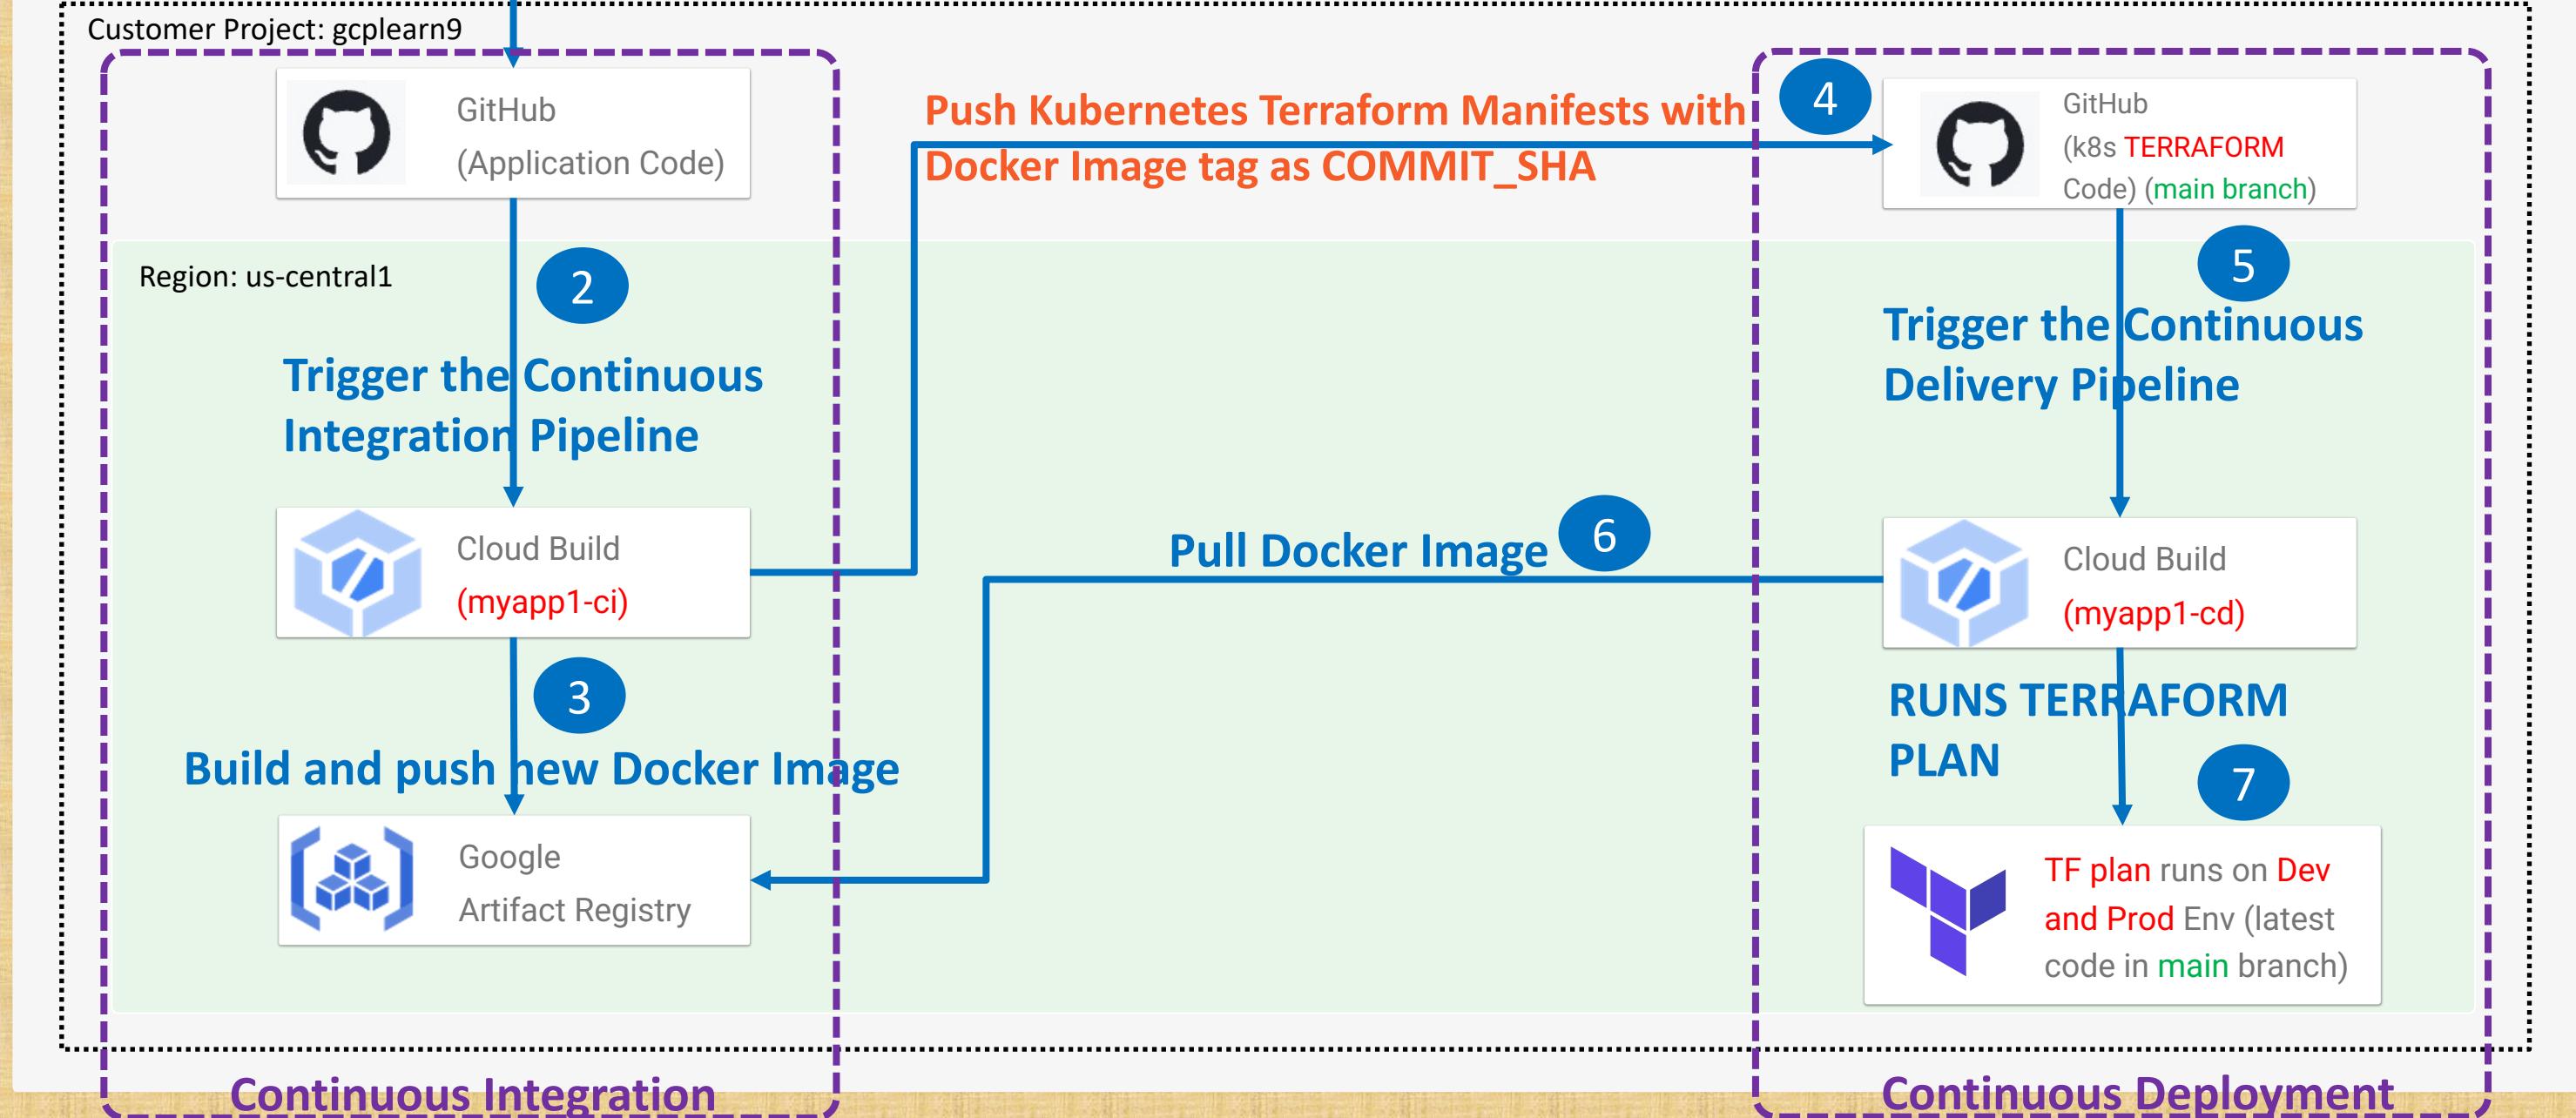
**Demo:** Implement Continuous Integration & Continuous Delivery for GKE Workloads  
(MyApp1)



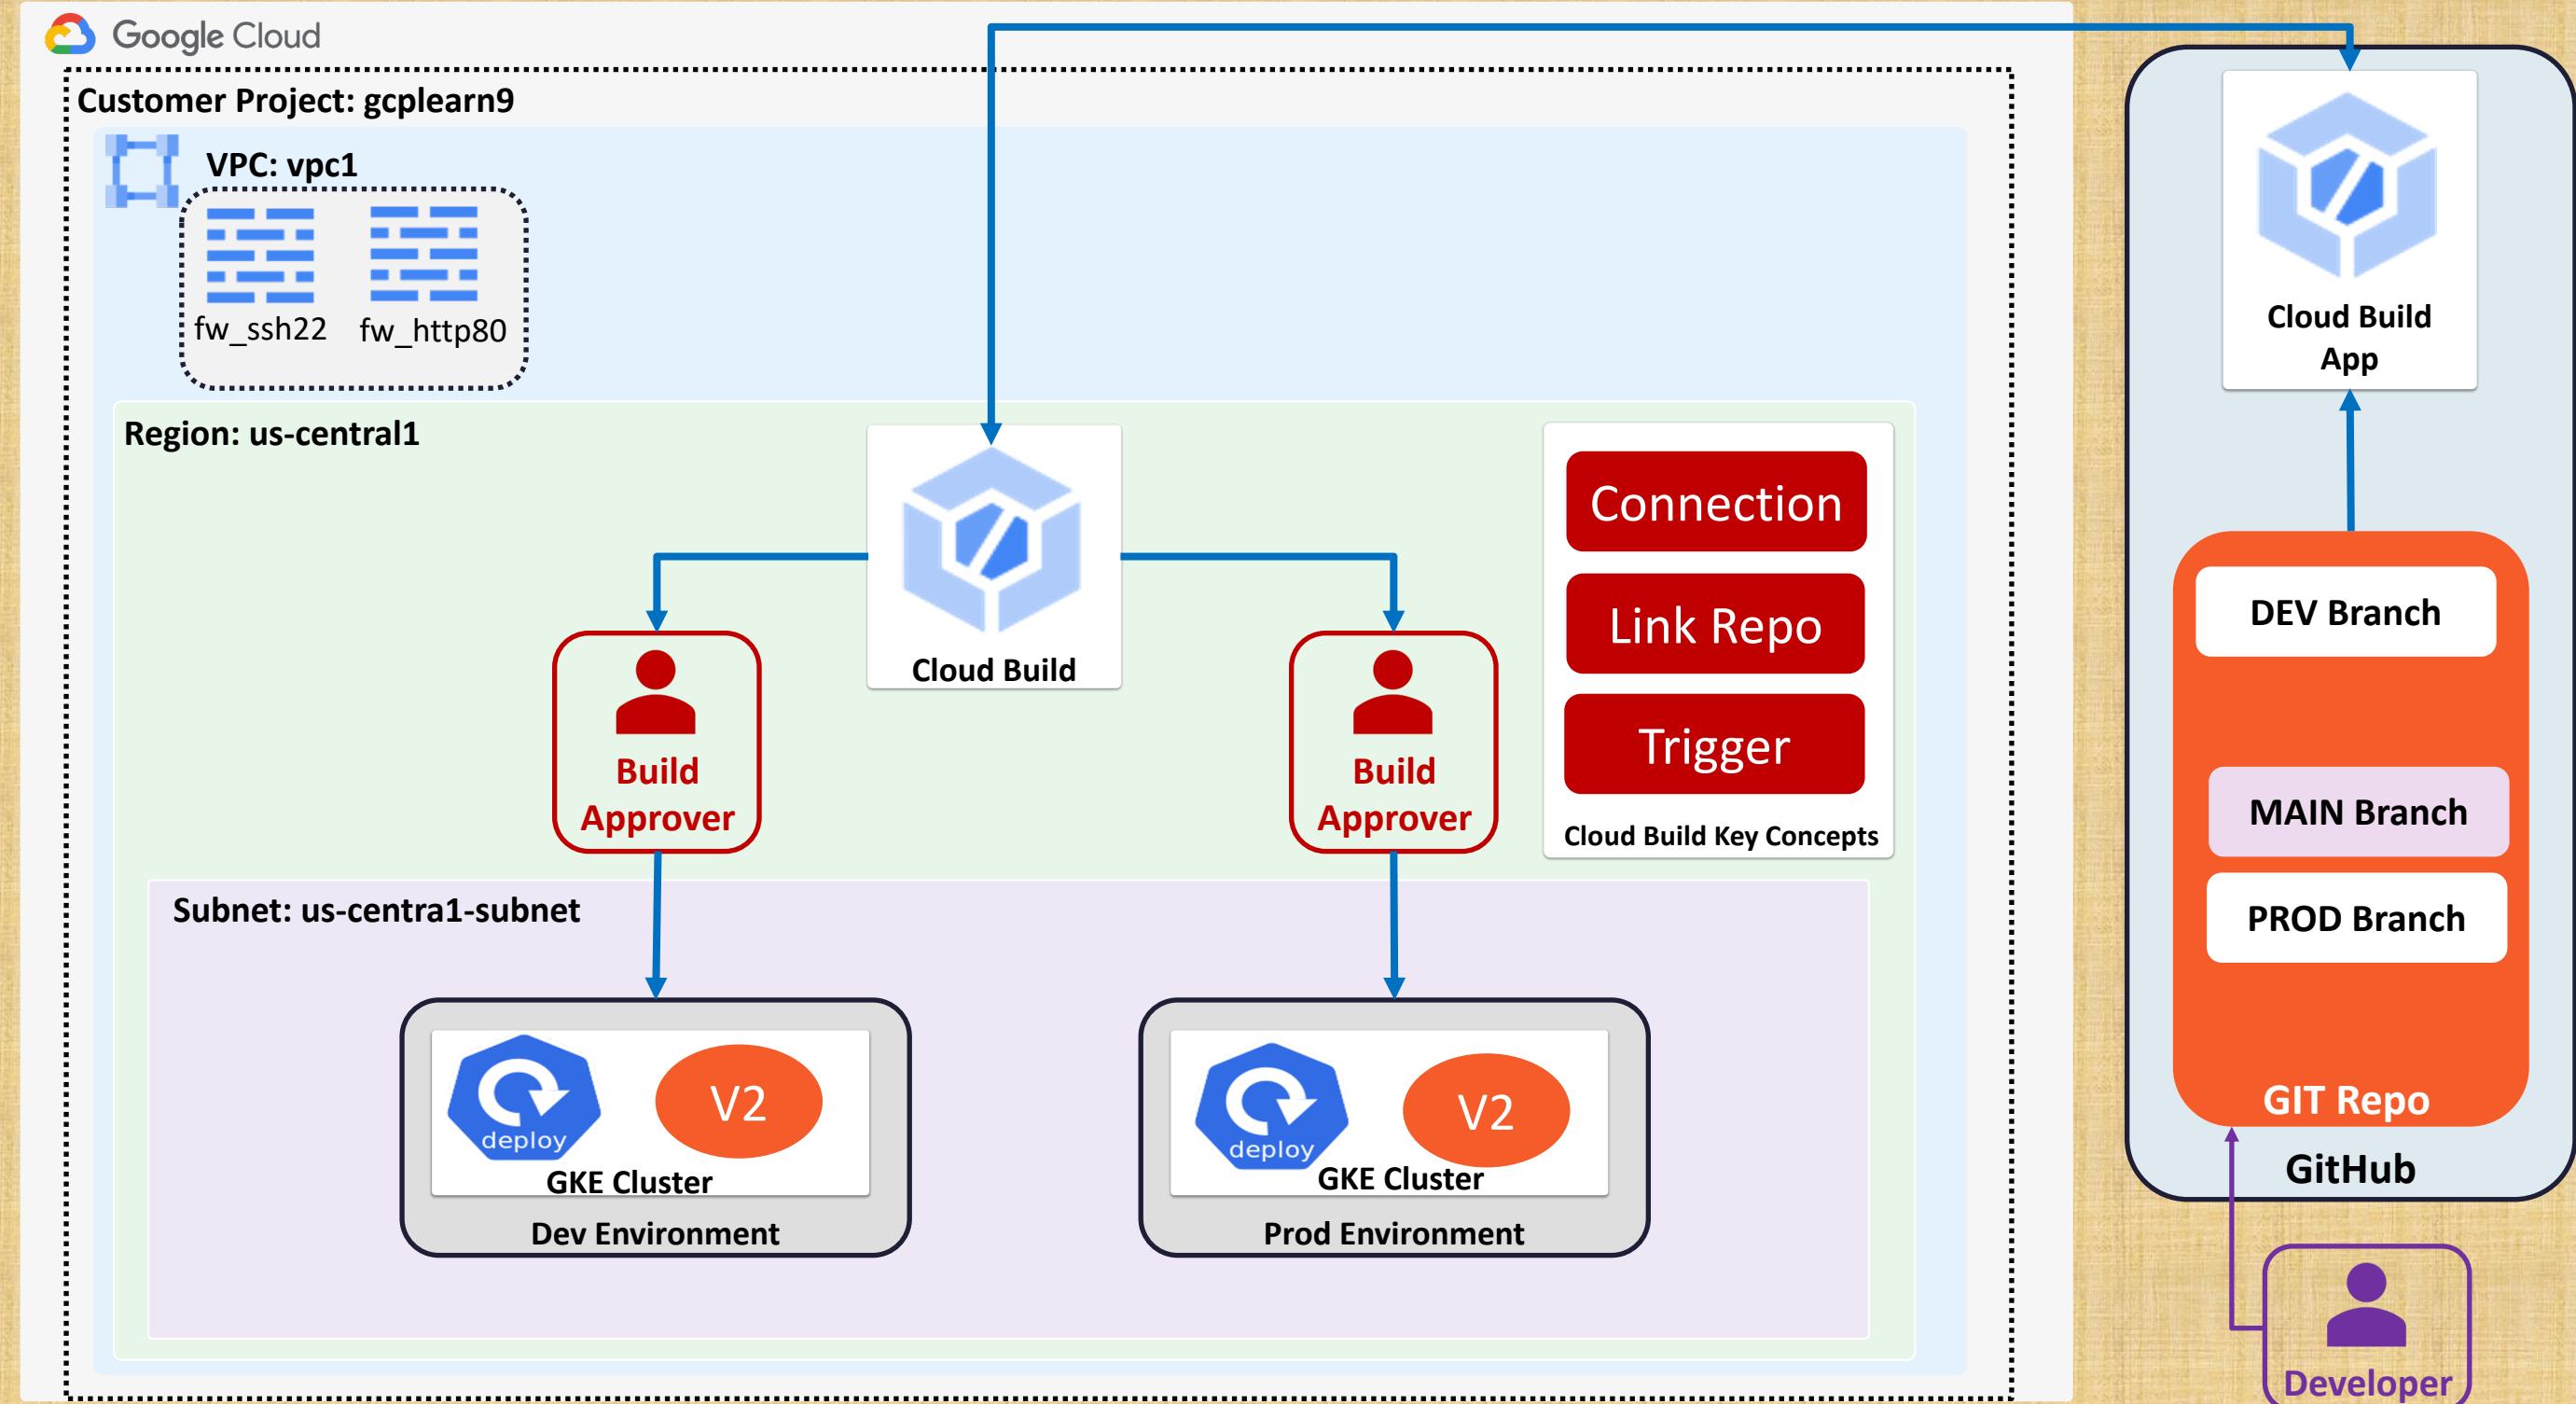
Admin



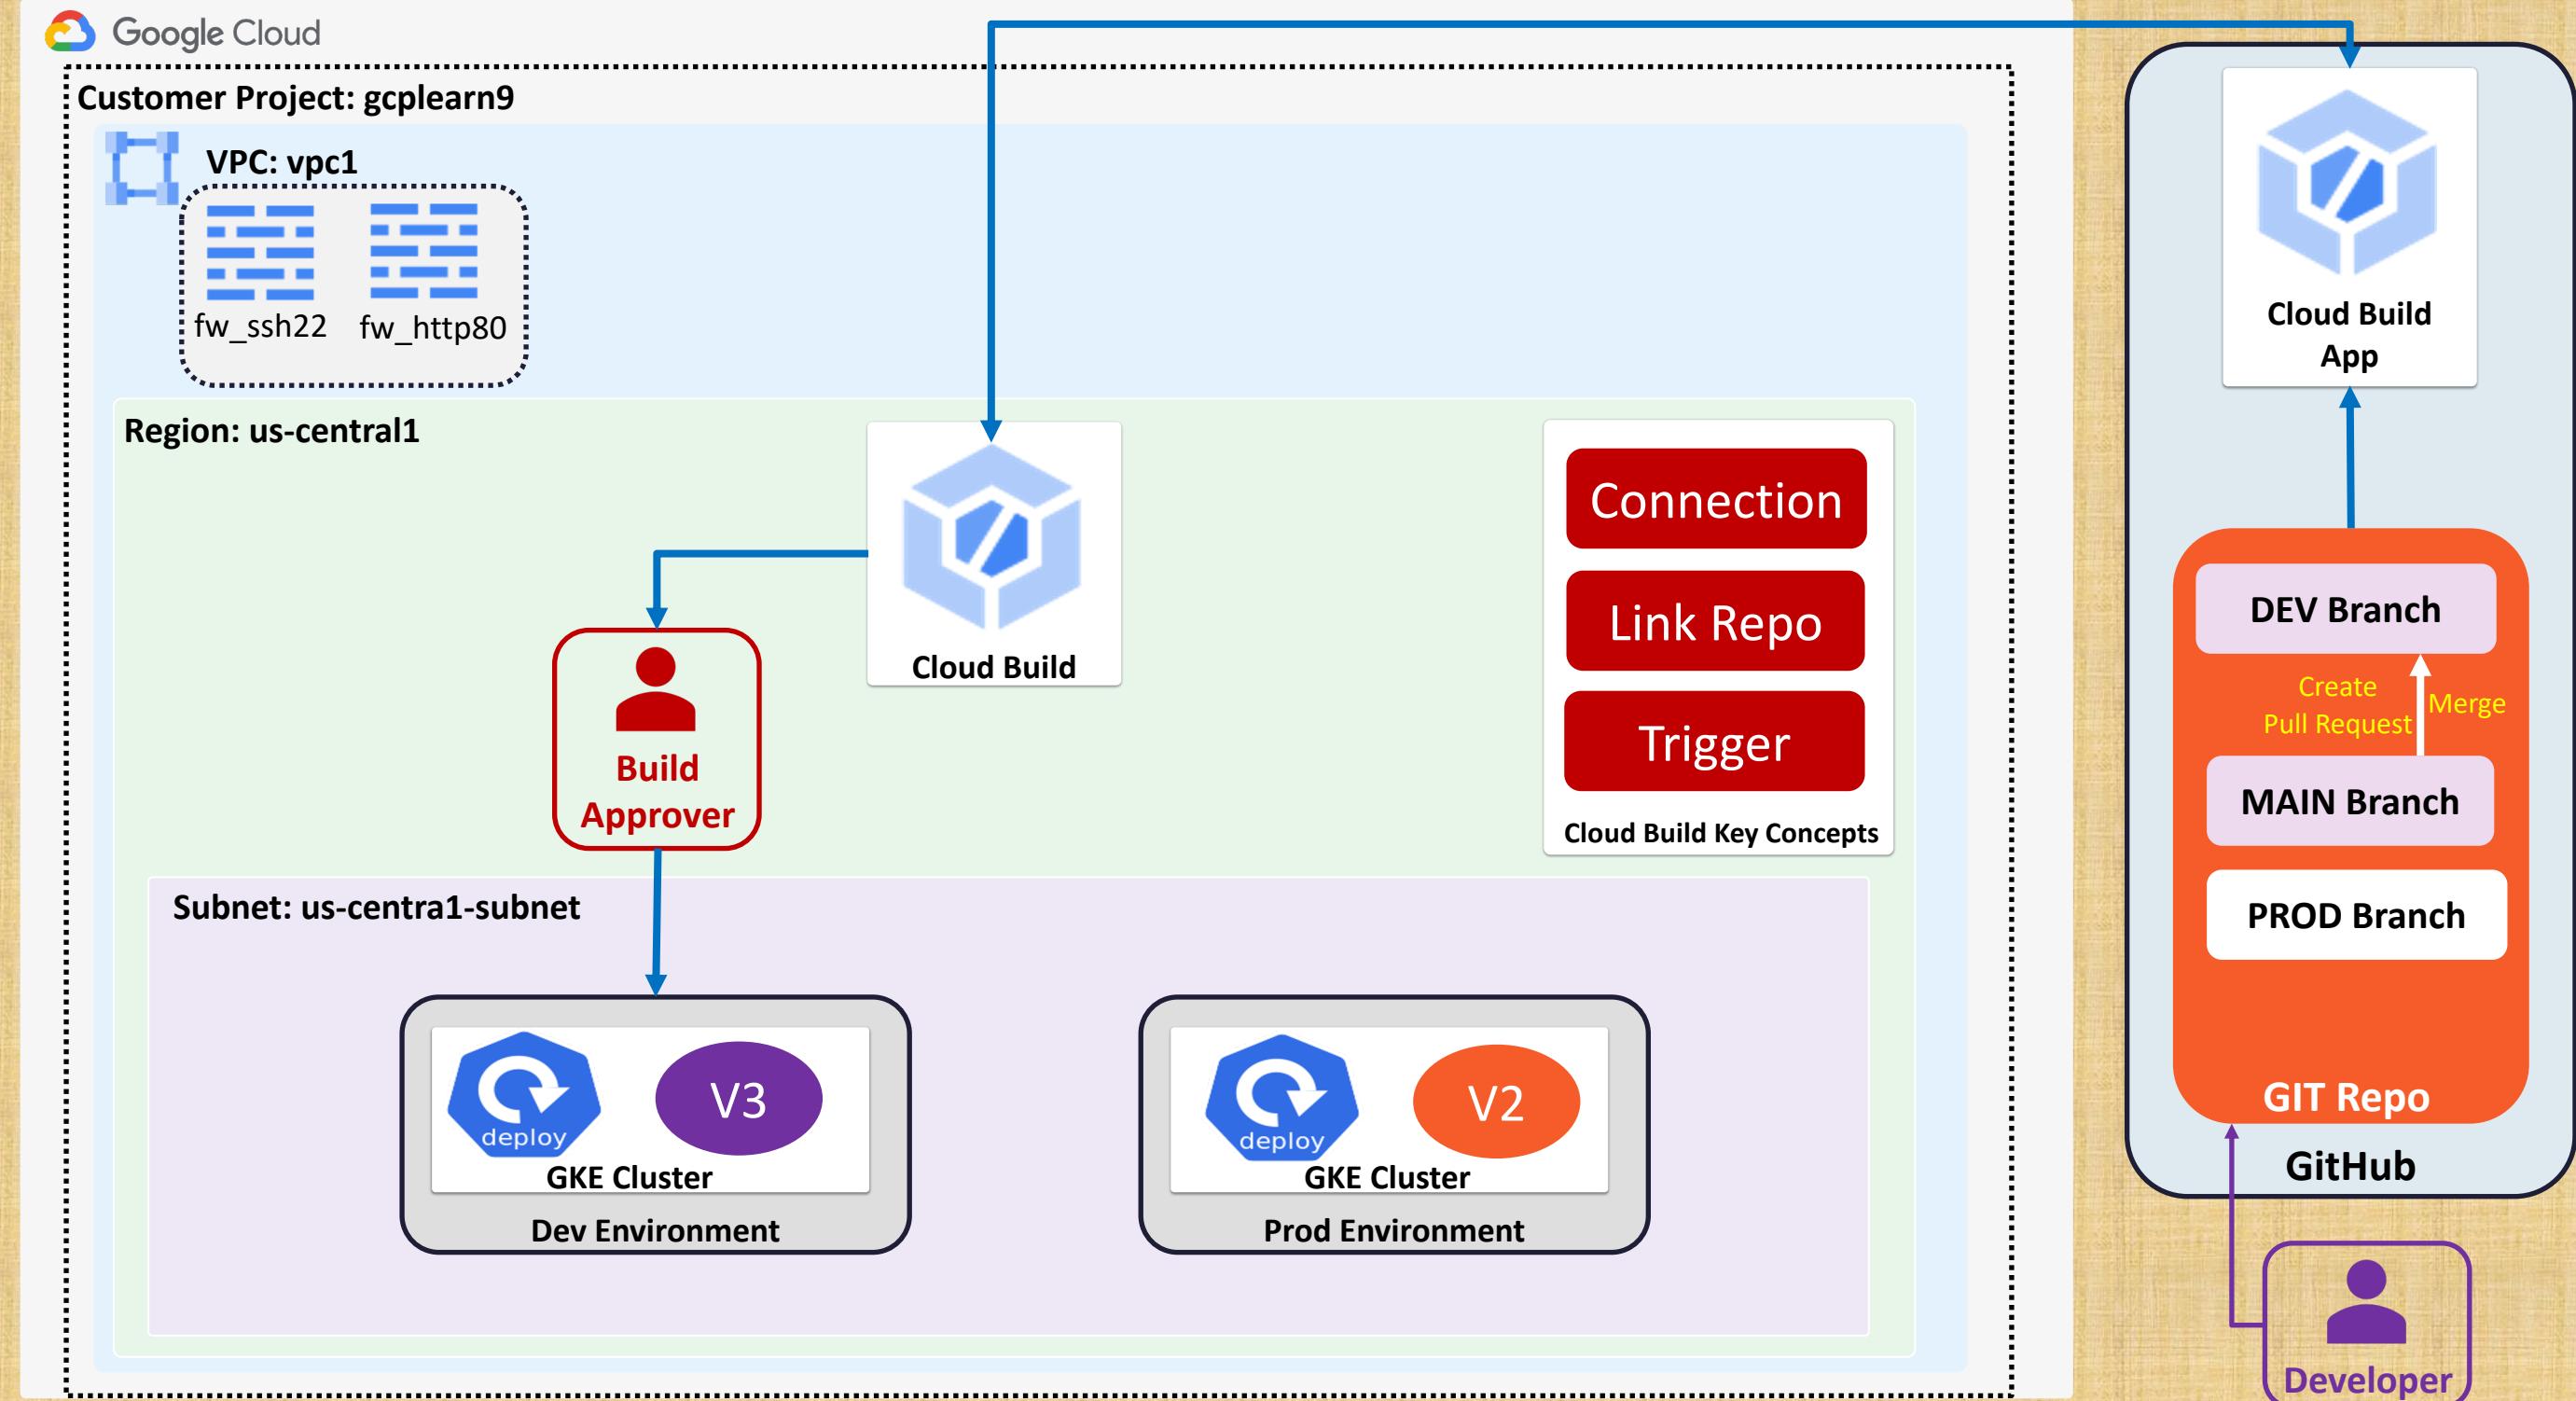
# GKE CI CD Pipelines



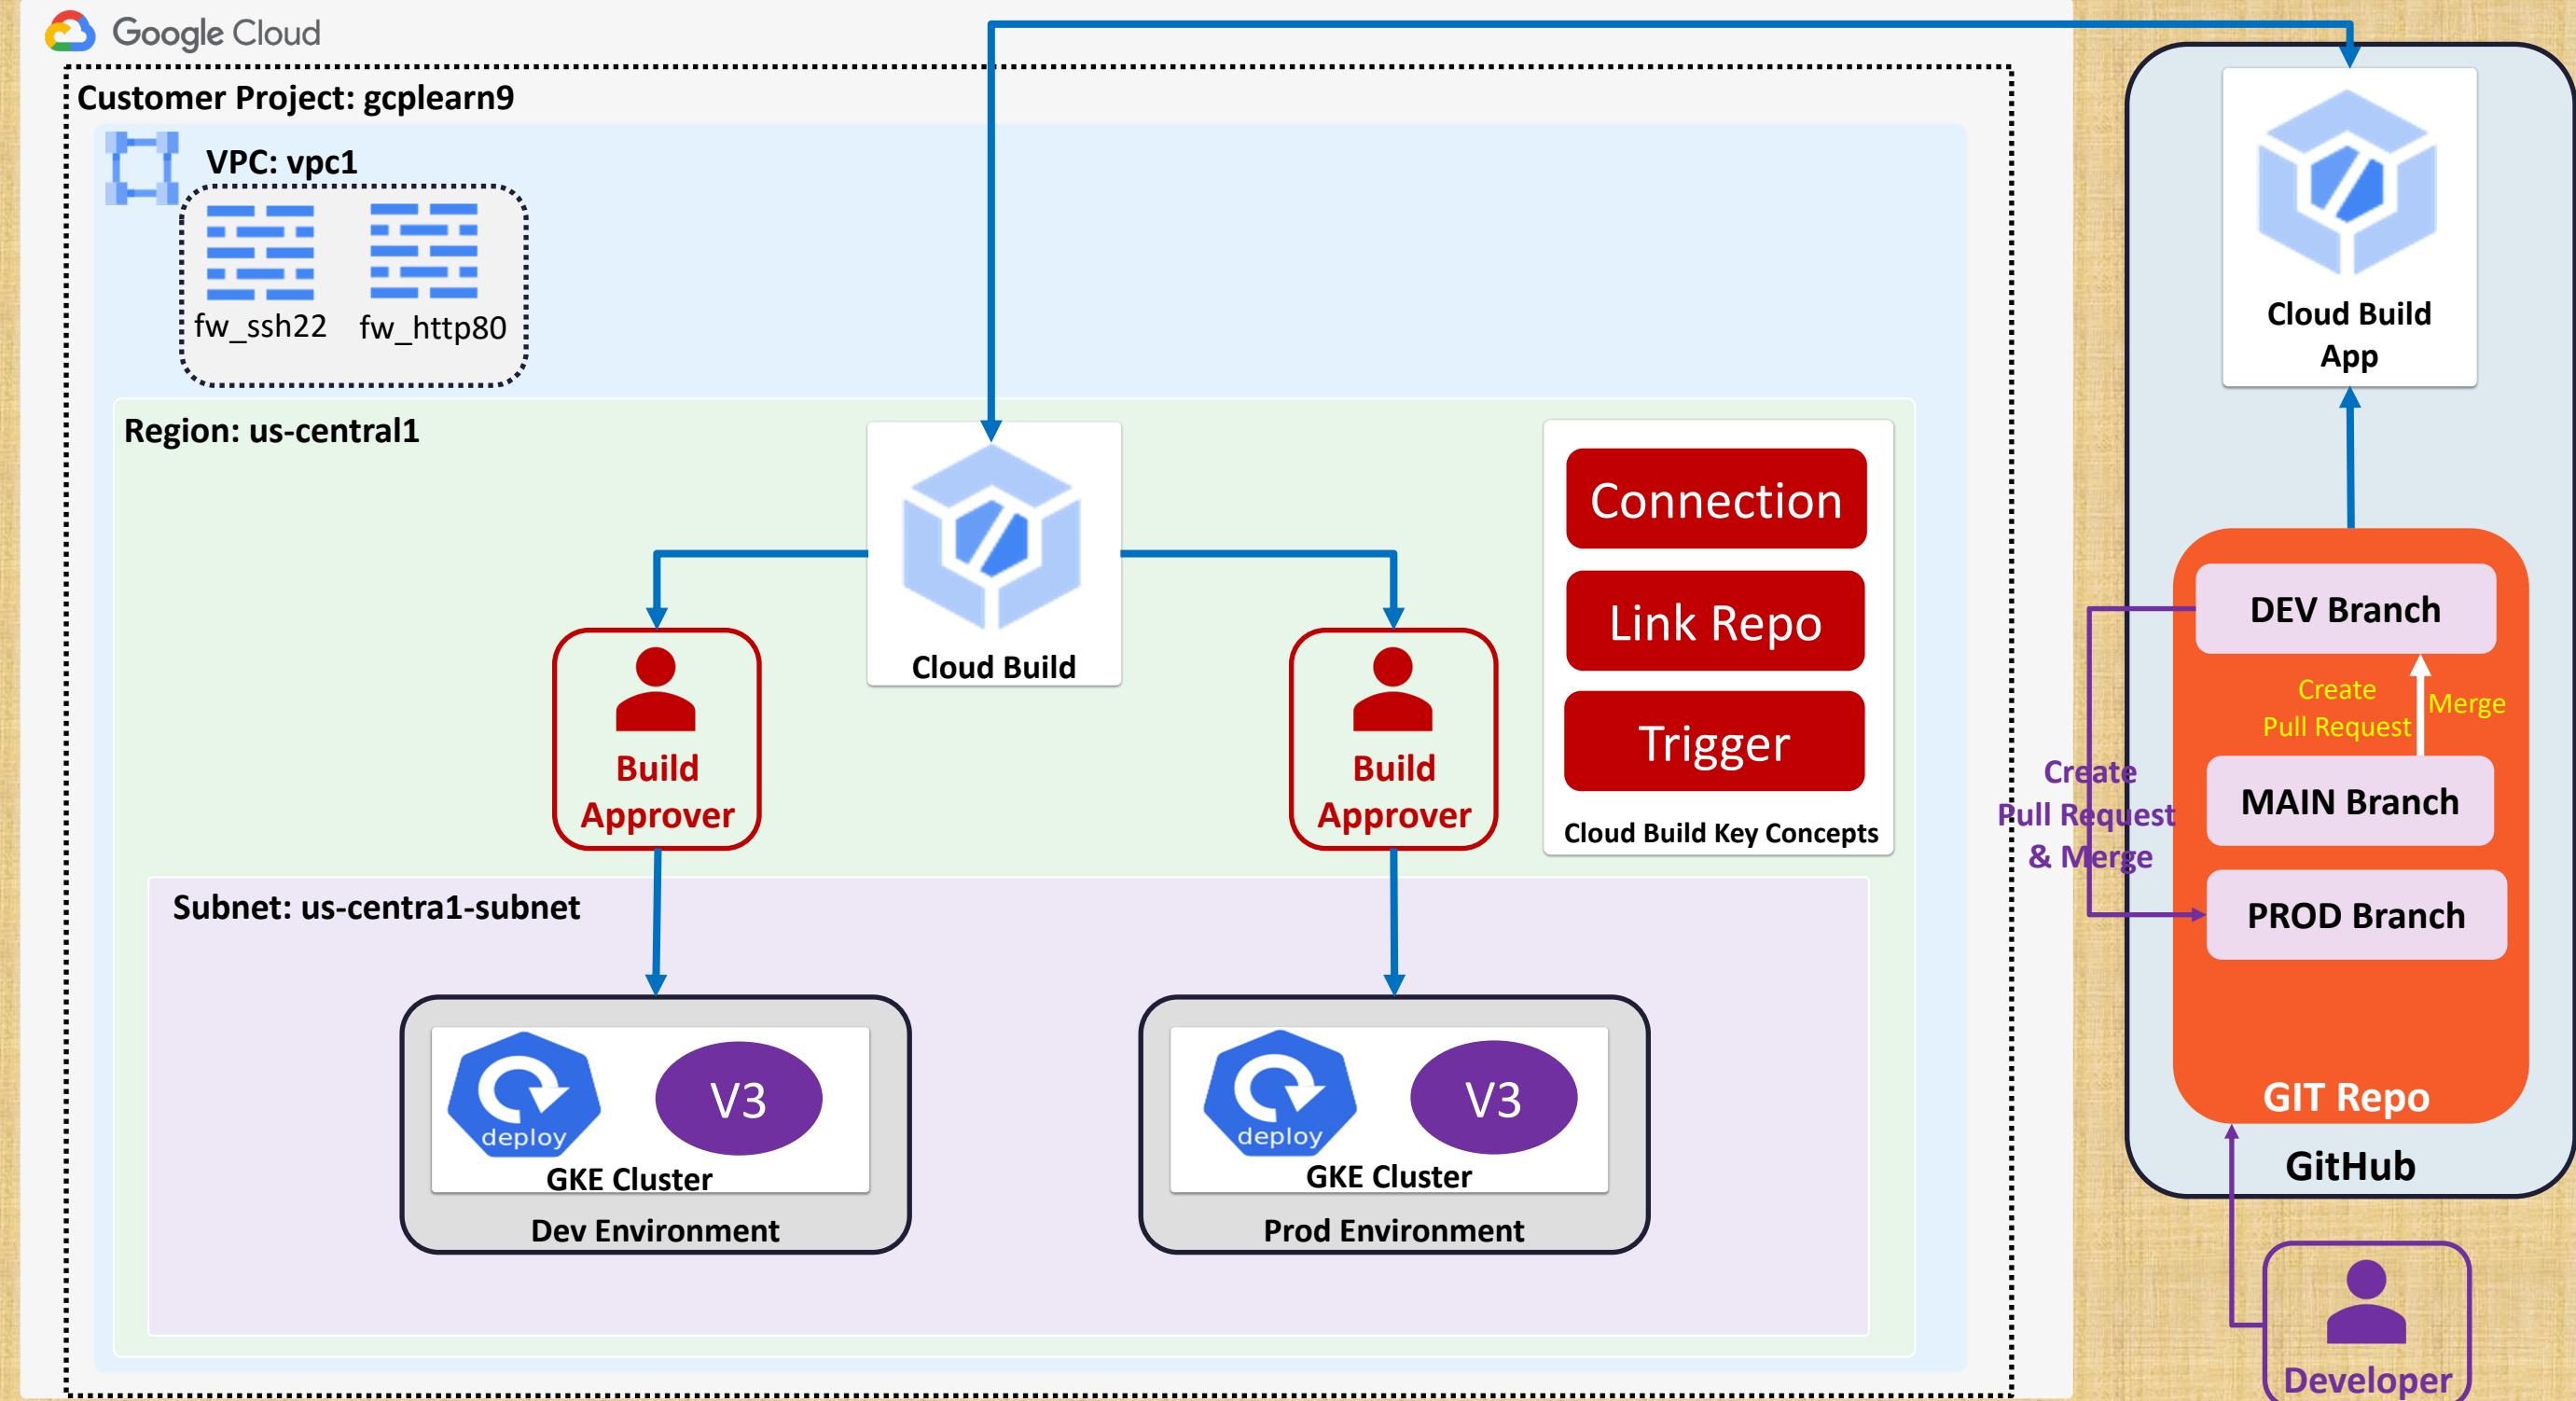
# Kubernetes Deployment DevOps for Terraform Configs



# Kubernetes Deployment DevOps for Terraform Configs



# Kubernetes Deployment DevOps for Terraform Configs



## Git Repo Files: App, k8s

- ✓ 28-GKE-App-Continuous-Delivery
  - ✓ 01-SSH-Keys
    - ≡ id\_github
    - ≡ id\_github.pub
  - ✓ 02-App-GitRepoFiles
    - ❖ .gitignore
    - ! BACKUP\_v1-cloudbuild-ci.yaml
    - ! BACKUP\_v2-cloudbuild-ci-cd.yaml
    - ! cloudbuild.yaml
  - ❖ Dockerfile
  - \$ git-deploy.sh
  - <> index.html
  - ≡ known\_hosts.github
  - 💡 main.tf.tpl
  - ⓘ README.md
- ✓ 03-K8s-GitRepoFiles
  - > environments
  - > modules
  - ❖ .gitignore
  - ! cloudbuild.yaml
  - \$ git-deploy.sh

# What are we going to learn?

Demo: Implement Continuous Integration & Continuous Delivery for GKE Workloads (MyApp1)



App: GitRepoFiles

Configure SSH Authentication from Cloud Build  
(From build run container) to K8S Repository

Add cloudbuild.yaml with Continuous Delivery Code

Get the github.com public certificate to store in  
known\_hosts file where build is run (In build run  
container)

K8S Deployment template file used for updating  
latest Docker Image Tag

K8S: GitRepoFiles

NO CHANGES in k8s repo from Demo-26 onwards

**ANYTHING AFTER THIS SLIDE IS NON-LIVE SLIDES**



Thank You